

AGENTS

AGENTS.md — HelixCode Authoritative Agent Guide

HelixCode Agent Guidelines

Version: 3.0.0 (Updated with full architecture audit) **Date:** 2026-04-30 **Scope:** All AI agents, human contributors, and automated processes working on HelixCode
Authority: Derived from HelixAgent AGENTS.md with HelixCode-specific enhancements

INHERITED FROM constitution/AGENTS.md (HelixConstitution submodule)

Base agent rules: constitution/AGENTS.md — READ IT FIRST. The base file is authoritative for any topic not covered here. Project-specific rules below extend them; they never weaken them.

The HelixConstitution submodule (./constitution/, added 2026-05-14) provides the universal CLI-agent operating manual. Every CLI agent working on this project — Claude Code, Codex, Cursor, Aider, Continue, Gemini CLI, future LLMs — **MUST** read both constitution/AGENTS.md AND this project AGENTS.md, treating the project file as a tightening overlay (never a weakening).

Project remote-policy note (revised 2026-06-03) — Per operator decision 2026-06-03, HelixCode permits GitFlic + GitVerse Git remotes in addition to GitHub + GitLab; the prior §6.W tightening forbidding GitFlic/GitVerse is **LIFTED** and pushes target all configured upstreams across all four providers. Other tightenings (e.g. CONST-033 host power management hard-ban, CONST-035 anti-bluff covenant, CONST-042 secret-leak policy, CONST-045 distribution-host config policy, CONST-046 no-hardcoded-content, CONST-047 recursive submodule application mandate) all SURVIVE this inheritance.

Project Overview

HelixCode is an enterprise-grade distributed AI development platform built in Go. It enables intelligent task division, work preservation, cross-platform development workflows, and multi-provider LLM integration through a unified REST API, CLI, Terminal UI, Desktop, and Mobile client architecture.

Current Status: The `internal/` foundation is largely solid (auth, database, server, worker, task, workflow, tools, editor, notification, MCP, **verifier** are real implementations). Critical bluff and stub areas remain in select entry points and peripheral packages. All agents **MUST** prioritize zero-bluff implementation.

LLMsVerifier Integration Status: `internal/verifier/` package is now implemented with REST API client, two-tier cache, circuit breaker health monitor, background poller, score adapter, and event publisher. BLUFF-002 (hardcoded CLI models) and BLUFF-004 (hardcoded external models) are **FIXED**. BLUFF-005 (scoring ignores verifier data) is **FIXED** in `ModelManager.SelectOptimalModel()`.

Key Features: - **Distributed Computing:** SSH-based worker pools with health monitoring, auto-installation, and consensus - **Multi-Provider LLM Integration:** 15+ providers (OpenAI, Anthropic, Gemini, Ollama, Azure, Bedrock, Groq, Mistral, Cohere, xAI, DeepSeek, Qwen, OpenRouter, HuggingFace, Llama.cpp) - **Development Workflows:** Automated planning, building, testing, refactoring with real shell execution - **Task Management:** Intelligent task division with priorities, dependencies, checkpointing, and Redis caching - **MCP Protocol:** Full Model Context Protocol server over WebSocket with tool dispatch - **Multi-Client Architecture:** REST API (Gin), Cobra CLI, Terminal UI (tview), Desktop (Fyne), Mobile (gomobile), WebSocket - **Memory Systems:** In-memory, filesystem, Redis, Memcached, Cognee, ChromaDB, Qdrant, Weaviate integrations - **Advanced Editor:** Multi-format code editing (diff, whole-file, search/replace, line-based) with backups - **Tools Ecosystem:** 40+ tools across filesystem, shell, web, browser, mapping, multiedit, confirmation, notebook, git - **Notifications:** Multi-channel support (Slack, Email, Telegram, Discord, Yandex Messenger, Max)

Technology Stack

Core Technologies: - **Language:** Go 1.24.0 with toolchain go1.24.9 - **Module:** `dev.helix.code` - **HTTP Framework:** Gin v1.11.0 - **Authentication:** JWT v4.5.2, bcrypt + argon2 - **Database:** PostgreSQL 15+ via pgx/v5 (optional) - **Cache:** Redis 7+ via go-redis/v9 (optional) - **Configuration:** Viper v1.21.0 - **CLI Framework:** Cobra v1.8.0 - **Testing:** Testify v1.11.1

UI Technologies: - **Desktop:** Fyne v2.7.0 - **Terminal UI:** tvview v0.42.0 - **Mobile:** gomobile bindings

External Integrations: - **Browser Automation:** chromedp v0.14.2 - **Web Scraping:** goquery v1.10.3 - **Tree-sitter:** go-tree-sitter - **Identity:** Azure SDK, AWS SDK v2 - **Vector/Memory:** Cognee, ChromaDB, Qdrant, Weaviate clients - **Container Orchestration:** digital.vasic.containers (vasic-digital/Containers submodule)

Working Directory & Build System

CRITICAL: All build and test commands must be run from the `helix_code/` subdirectory, not the repository root.

`cd` HelixCode

Build Commands

Command	Purpose
<code>make build</code>	Build server binary to <code>bin/helixcode</code>
<code>make test</code>	Run <code>go test -v ./...</code>
<code>make test-all</code>	Run tests + coverage + benchmarks + docs
<code>make test-coverage</code>	Generate coverage report
<code>make test-benchmark</code>	Run Go benchmarks
<code>make logo-assets</code>	Generate logo assets (required before first build)
<code>make setup-deps</code>	Run <code>go mod tidy</code>
<code>make fmt</code>	Run <code>go fmt ./...</code>
<code>make lint</code>	Run <code>golangci-lint run ./...</code>
<code>make clean</code>	Clean build artifacts
<code>make dev</code>	Start development server
<code>make prod</code>	Cross-platform production build
<code>make mobile</code>	Build iOS + Android targets
<code>make aurora-os</code>	Build Aurora OS target
<code>make harmony-os</code>	Build Harmony OS target

Full Infrastructure Test Commands

Command	Purpose
<code>make test-infra-up</code>	Start full Docker test infrastructure
<code>make test-infra-down</code>	Stop full Docker test infrastructure
<code>make test-full</code>	ALL tests with real infrastructure (zero skips)
<code>make test-unit-full</code>	Unit tests with real services
<code>make test-integration-full</code>	Integration tests with -tags=integration
<code>make test-e2e-full</code>	E2E challenge tests via runner
<code>make test-security-full</code>	Security test suite
<code>make test-load-full</code>	Load tests
<code>make test-complete</code>	Sequential run of all full test types
<code>make coverage-full</code>	Coverage with full infrastructure

Containerized Builds (NO Host Dependencies)

Command	Purpose
<code>make container-builder-image</code>	Build the builder container image
<code>make container-build</code>	Build application inside container
<code>make container-test</code>	Run tests inside container
<code>make container-lint</code>	Run linter inside container
<code>make container-shell</code>	Interactive shell in builder container
<code>make container-dev-up</code>	Start containerized dev environment
<code>make container-dev-down</code>	Stop containerized dev environment
<code>make container-release</code>	Full release build in container
<code>./scripts/containers/build-in-container.sh</code>	Convenience wrapper script

The builder container includes: Go 1.24, gcc, postgresql-client, redis, docker-cli, golangci-lint, and all build tools. The only host requirement is Docker/Podman.

Standalone Test Scripts

Script	Purpose
./run_tests.sh --unit	Unit tests
./run_tests.sh --integration	Integration tests
./run_tests.sh --e2e	E2E tests
./run_tests.sh --coverage	Coverage analysis
./run_tests.sh --security	Security tests
./run_all_tests.sh	Orchestrates ALL suites sequentially
./run_integration_tests.sh	DB integration tests with Docker

Single Test Execution

```
go test -v -run TestName ./path/to/package
go test -v -tags=integration ./internal/database
cd tests/e2e/challenges && go run cmd/runner/main.go -challenge
                                ascii-art-generator-001 -providers ollama
```

Architecture & Code Organization

```
helix_code/
├── cmd/                                # Application entry points
│   ├── server/main.go                 # HTTP server entry point
│   ├── cli/main.go                   # Legacy flag-based CLI client
│   ├── root.go                       # Cobra root command (`helix`)
│   ├── main_commands.go              # `helix start`, `helix auto`
│   ├── other_commands.go             # `helix server`, `helix version`, etc.
│   ├── local-llm.go                  # `helix local-llm` command tree
│   ├── local-llm-advanced.go         # Advanced local-llm commands
│   ├── helix-config/main.go          # Dedicated config management CLI
│   ├── security-test/main.go         # Simulated security test runner
│   ├── security-fix/main.go          # Security fix wrapper
│   ├── security-fix-standalone/main.go # Standalone security scanner
│   ├── performance-optimization/main.go # Performance optimizer
│   ├── performance-optimization-standalone/main.go # Standalone perf simu
│   └── config-test/main.go           # Config hot-reload test utility
├── internal/                          # Internal packages (~40 packages)
│   ├── auth/                         # JWT authentication, bcrypt/argon2, ses
│   └── llm/                          # LLM provider implementations (15+ prov
```

└─ providers/	# Per-provider HTTP clients
└─ compression/	# Context compression
└─ vision/	# Vision/multimodal support
└─ provider/	# Provider abstractions
└─ providers/	# Provider management
└─ worker/	# SSH-based worker pool, health checks
└─ task/	# Task queues, dependencies, checkpoints
└─ server/	# Gin HTTP server, routes, middleware
└─ database/	# PostgreSQL pgx pool, schema initialization
└─ redis/	# go-redis wrapper with graceful degradation
└─ tools/	# 40+ tool ecosystem registry
└─ └─ filesystem/	# fs_read, fs_write, fs_edit, glob, grep
└─ └─ shell/	# shell, shell_background with sandbox
└─ └─ web/	# web_fetch, web_search
└─ └─ browser/	# browser_launch, browser_navigate, browser_print
└─ └─ multiedit/	# Transactional multi-file editing
└─ └─ git/	# Git automation
└─ editor/	# Multi-format code editing with backups
└─ memory/	# Memory providers (in-mem, filesystem, redis)
└─ cognee/	# Cognee.ai memory integration
└─ context/	# Hierarchical context management with TFS
└─ notification/	# Multi-channel notification engine
└─ mcp/	# Model Context Protocol WebSocket server
└─ workflow/	# Development workflow execution
└─ config/	# Viper-based configuration management
└─ event/	# Pub/sub event bus
└─ logging/	# Structured logging wrapper
└─ monitoring/	# Metric collection framework
└─ security/	# Security scanning (stubbed)
└─ session/	# Development session management
└─ agent/	# Agent orchestration
└─ project/	# Project management
└─ rules/	# Rules engine
└─ hooks/	# Hook system
└─ focus/	# Focus chain management
└─ template/	# Template system
└─ persistence/	# State persistence
└─ deployment/	# Deployment management
└─ discovery/	# Service/model discovery
└─ hardware/	# Hardware abstraction
└─ repomap/	# Repository mapping
└─ version/	# Version management
└─ fix/	# Security fix engine
└─ performance/	# Performance optimization

└─ testutil/	# Test utilities
└─ mocks/	# Shared mocks
└─ applications/	# Platform-specific applications
└─ desktop/	# Fyne desktop app
└─ terminal-ui/	# tview terminal UI
└─ android/	# Android app
└─ ios/	# iOS app
└─ aurora-os/	# Aurora OS client
└─ harmony-os/	# Harmony OS client
└─ api/	# OpenAPI specification
└─ openapi.yaml	# Full REST API spec (OpenAPI 3.0.3)
└─ config/	# Configuration files
└─ config.yaml	# Primary application config
└─ production-config.yaml	# Enterprise production config
└─ minimal-config.yaml	# Minimal test config (DB/Redis disabled)
└─ test-config.yaml	# Test-specific config
└─ working-config.yaml	# Working variant
└─ azure_example.yaml	# Azure-specific example
└─ model-aliases.example.yaml	# Model alias examples
└─ tests/	# New test framework
└─ e2e/challenges/	# Challenge-based E2E tests
└─ automation/	# Hardware automation tests
└─ test/	# Legacy/parallel test suites
└─ integration/	# Integration tests
└─ e2e/	# Legacy E2E tests
└─ automation/	# Provider automation tests
└─ load/	# Load tests
└─ benchmarks/	# Performance benchmarks
└─ security/	# Security tests
└─ standalone_tests/	# Standalone CLI tests
└─ docker/	# Docker assets and extended compose
└─ scripts/	# Build and deployment scripts
└─ assets/	# Logo and image assets

Verified Real Implementations

AUTH-001: Authentication System (VERIFIED REAL)

File: `internal/auth/auth.go` (~470 lines) **Assessment:** Production-ready - User registration with validation - Password hashing with bcrypt + argon2 fallback - JWT token generation and verification (JWT v4) - Session management with crypto-random tokens - Constant-time comparison for timing attack prevention - Full test coverage in `internal/auth/auth_test.go` (~777 lines)

DB-001: Database Layer (VERIFIED REAL)

File: `internal/database/database.go` **Assessment:** Production-ready - PostgreSQL connection pool via pgx/v5 - Full schema initialization (users, workers, tasks, projects, sessions, LLM providers, MCP servers, notifications, audit logs) - DatabaseInterface for testability - Graceful degradation when host is empty

SRV-001: HTTP Server (VERIFIED REAL)

File: `internal/server/server.go` **Assessment:** Production-ready - Gin-based server with 50+ routes across `/api/v1/` - JWT auth middleware, CORS, security headers - WebSocket endpoint for MCP - Health check with DB + Redis validation - Graceful shutdown (30s timeout)

LLM-001: LLM Providers (VERIFIED REAL)

File: `internal/llm/` (~5000+ lines across providers) **Assessment:** Real HTTP clients - AnthropicProvider (~752 lines): Full SSE streaming, prompt caching, extended thinking, tool calls - OpenAIProvider (~431+ lines): Full HTTP API client - ModelManager: Multi-provider orchestration, selection strategy, fallback chain - 16 provider subdirectories with real HTTP implementations - **Note:** The `internal/llm/` package is genuine. Bluff areas are at `cmd/cli/main.go` only.

WRK-001: Worker Pool (VERIFIED REAL)

File: `internal/worker/` (~800+ lines) **Assessment:** Real distributed worker management - WorkerManager: Register, heartbeat, assign tasks, complete tasks - SSH config parsing, capability matching, resource tracking - Health checks with TTL

TSK-001: Task Management (VERIFIED REAL)

File: `internal/task/` (~1000+ lines) **Assessment:** Real task lifecycle - Priority queues, dependency validation, checkpointing - Redis caching with graceful degradation - Retry logic and cleanup

WFL-001: Workflow Engine (VERIFIED REAL)

File: `internal/workflow/` (~1100+ lines) **Assessment:** Real shell execution - Executor dispatches to real `exec.CommandContext()` calls - Security filtering via `isDangerousCommand()` (`rm`, `dd`, `mkfs`, fork bombs, etc.) - LLM integration with real `LLMRequest` - Supports Go, Node, Python, Rust project types

TOO-001: Tools Ecosystem (VERIFIED REAL)

File: `internal/tools/` (~2000+ lines) **Assessment:** Real tool registry - 8 categories: filesystem, shell, web, browser, mapping, multiedit, confirmation, notebook - Real `chromedp` browser automation - Transactional multi-file editing

EDT-001: Code Editor (VERIFIED REAL)

File: `internal/editor/` (~600+ lines) **Assessment:** Real file I/O - Diff, whole-file, search/replace, line-based editors - Automatic file backup with `io.Copy` - `EditApplier` / `EditValidator` interfaces

NOT-001: Notification Engine (VERIFIED REAL)

File: `internal/notification/` (~800+ lines) **Assessment:** Real HTTP/SMTP calls - Slack (webhook HTTP POST), Email (SMTP via `net/smtp`), Telegram (Bot API), Discord (webhook) - Yandex Messenger (OAuth API), Max (enterprise API) - Rate limiting, retry, queue, metrics

MCP-001: MCP Protocol Server (VERIFIED REAL)

File: `internal/mcp/` (~400+ lines) **Assessment:** Real WebSocket server - gorilla/websocket concurrent session handling - JSON-RPC-like message format - Tool execution dispatch

CFG-001: Configuration Management (VERIFIED REAL)

File: `internal/config/` (~1700+ lines) **Assessment:** Full Viper integration - Environment variable binding (`HELIX_*`) - Config file search (`.`, `$HOME/.helixcode`, `/etc/helixcode`) - Validation rules, default config creation - `ConfigManager` for load/save/merge

QA-001: helix_qa Integration (VERIFIED REAL)

Files: internal/helixqa/, internal/server/qa_handlers.go, applications/terminal_ui/main.go **Assessment:** Full embedded QA engine with real session lifecycle - Engine struct manages QA sessions with map + sync.RWMutex - StartSession(), CancelSession(), GetSession(), ListSessions() with real state tracking - REST API: POST /api/v1/qa/session, GET /api/v1/qa/session/:id/status, GET /api/v1/qa/session/:id/report, GET /api/v1/qa/session/:id/screenshot/:name, DELETE /api/v1/qa/session/:id - CLI flags: --qa-run, --qa-list, --qa-report, --qa-screenshot, --qa-cancel - TUI dashboard with session table, stats panel, refresh/cancel actions - Screenshot pipeline: 8 platform engines (Linux, Web, iOS, Android, CLI, TUI, macOS, Windows) - Tests: internal/helixqa/wrapper_test.go, internal/server/qa_handlers_test.go, pkg/screenshot/*_test.go

Verified Bluff & Stub Areas (MUST FIX)

BLUFF-001: LLM Generation is Simulated in Legacy CLI (CRITICAL) — FIXED

File: cmd/cli/main.go lines ~236-284 **Evidence:** Previously returned `fmt.Sprintf("Generated response for: %s...", prompt)` without calling any provider. **Fix:** `handleGenerate()` now constructs a real `llm.LLMRequest` with user messages and calls `provider.Generate()` / `provider.GenerateStream()`. Errors are propagated to the user if the provider is unavailable. **Verification:** `go build -tags nogui ./cmd/cli/` compiles; provider call is real (returns error if Ollama/etc. is not running). **Fix Priority:** P0 — RESOLVED

BLUFF-002: Model Listing is Hardcoded in Legacy CLI (CRITICAL) — FIXED

File: cmd/cli/main.go lines ~101-128 **Evidence:** Previously only 3 hardcoded models. No dynamic discovery. **Fix:** Replaced with verifier-aware `handleListModels()` that queries `LLMsVerifier` adapter first, falls back to provider discovery, then to constitutional `FallbackModels` (7 models with scores and verification status). **Verification:** `go test -v ./internal/verifier/...` passes; `go build ./cmd/cli/...` compiles. **Fix Priority:** P0 — RESOLVED

BLUFF-003: Command Execution is Simulated in Legacy CLI (HIGH) — FIXED

File: `cmd/cli/main.go` lines ~310-324 **Evidence:** Previously printed the command and slept for 1 second without executing anything. **Fix:** `handleCommand()` uses `exec.CommandContext(ctx, "sh", "-c", command)` with real `os.Stdout/os.Stderr` redirection. Exit codes are reported. **Verification:** `go build -tags nogui ./cmd/cli/` compiles. **Fix Priority:** P0 — RESOLVED

STUB-001: Security Scanning is Simulated — FIXED

File: `internal/security/security.go` (~132 lines) **Evidence:** Previously `ScanFeature()` contained explicit “Simulate security scanning logic” comment and returned `Success=true, Score=95` with empty issues. **Fix:** Replaced with real Scanner interface, SonarQube/Snyk clients. `ScanFeature()` now dispatches real security scanners. Score is no longer hardcoded. **Verification:** 17/17 tests PASS, zero “simulated” in production code. **Fix Priority:** P1 — RESOLVED (commit 33ddf6a)

STUB-002: Memory Redis/Memcached Providers Store Locally — FIXED

File: `internal/memory/` (~1800+ lines) **Evidence:** Previously `RedisMemoryProvider` and `MemcachedMemoryProvider` stored data in local maps. **Fix:** `internal/redis/redis.go` provides real go-redis/v9 client. Dedicated memory stubs no longer exist. **Fix Priority:** P2 — RESOLVED

STUB-003: Security-Test Entry Point is Entirely Simulated — FIXED

File: `cmd/security_test/main.go` **Evidence:** Previously had 12 simulated test results via `simulateSecurityScan()`. **Fix:** Rewired to real `internal/security` scanner dispatch via `ScanFeature()`. **Fix Priority:** P2 — RESOLVED (commit 33ddf6a)

STUB-004: Several helix Subcommands are Print-Only — FIXED

File: `cmd/other_commands.go` **Evidence:** Previously `server, generate, test, worker, notify` commands were stubbed. **Fix:** Wired to real `server.New()`, `llm.ModelManager`, notification engine, and `go test` dispatch. **Fix Priority:** P2 — RESOLVED (commit 33ddf6a)

STUB-005: Several helix-config Subcommands are Placeholders — PENDING

File: cmd/helix_config/main.go **Evidence:** Many template/history/schema subcommands print placeholder messages. **Fix Priority:** P3

BLUFF-004: LLMsVerifier Integration is Stubbed or Bypassed (CRITICAL) — FIXED

File Pattern: internal/verifier/*.go **Evidence:** Previously returned hardcoded 8.5 scores, empty discovery results. **Fix:** REST API client, two-tier cache, circuit breaker, background poller implemented. Fallback models are constitutionally permitted (CONST-036). **Verification:** Fallback models list only at fallback_models.go (permitted). All other verifier code uses real client. **Fix Priority:** P0 — RESOLVED

BLUFF-005: Provider Discovery Uses Hardcoded Env Var Names (HIGH) — FIXED

File Pattern: internal/verifier/startup.go or provider adapter files
Evidence: Previously hardcoded env var names like "OPENAI_API_KEY". **Fix:** Env vars sourced from config struct fields, not hardcoded strings. **Verification:** grep sweep clean. **Fix Priority:** P1 — RESOLVED

BLUFF-006: Model Capabilities Are Hardcoded (HIGH) — FIXED

File Pattern: internal/llm/*.go **Evidence:** Previously SupportsToolUse: true as struct literal. **Fix:** No hardcoded capabilities found in grep sweep. **Fix Priority:** P1 — RESOLVED

BLUFF-007: Test Claims Integration But Uses Mocked Verifier (CRITICAL) — FIXED

File Pattern: *_test.go files with testify/mock or testMode: true **Evidence:** Previously non-unit tests used mocked verifier. **Fix:** No mocks in non-unit verifier tests. **Verification:** grep sweep clean. **Fix Priority:** P0 — RESOLVED

BLUFF-008: Scoring Weights Do Not Sum to 1.0 (MEDIUM) — FIXED

File Pattern: configs/verifier.yaml or internal/verifier/config.go
Evidence: Previously scoring weights misconfigured. **Fix:** Weights configurable via ScoringWeights struct in adapter.go. No hardcoded weight issues found. **Fix Priority:** P2 — RESOLVED

BLUFF-009: /metrics Endpoint Returns Hardcoded Zeros (CRITICAL) — FIXED

File: `internal/server/handlers.go` lines ~834-855 **Evidence:** All dynamic metrics (goroutines, memory, database connections) were hardcoded to 0. **Fix:** `getMetrics()` now calls `runtime.ReadMemStats()`, `runtime.NumGoroutine()`, and `s.db.Pool.Stat()` to return real values. **Fix Priority:** P0 — RESOLVED

BLUFF-010: Multi-Edit Conflict Detection is a No-Op (HIGH) — FIXED

File: `internal/tools/multiedit/transaction.go` lines ~352-369 **Evidence:** `detectFileConflict()` always returned `nil`, `nil` with comment “For now, we’ll assume no conflicts.” **Fix:** Implemented real conflict detection — reads the file from disk, computes SHA-256, and compares against the Checksum field. Returns `ConflictModified` or `ConflictDeleted` when appropriate. **Fix Priority:** P1 — RESOLVED

Configuration Management

Primary Configuration

Main config at `config/config.yaml`:

```
server:
  address: "0.0.0.0"
  port: 8080
  read_timeout: 30
  write_timeout: 30
  idle_timeout: 300
  shutdown_timeout: 30

database:
  host: "" # Empty string disables PostgreSQL
  port: 5432
  user: "helix"
  password: "${HELIX_DATABASE_PASSWORD}"
  dbname: "helixcode_prod"
  sslmode: "disable"

redis:
  host: "redis"
```

```
port: 6379
password: "${HELIX_REDIS_PASSWORD}"
db: 0
enabled: true
```

auth:

```
jwt_secret: "${HELIX_AUTH_JWT_SECRET}"
token_expiry: 86400
session_expiry: 604800
bcrypt_cost: 12
```

workers:

```
health_check_interval: 30
health_ttl: 120
max_concurrent_tasks: 10
```

tasks:

```
max_retries: 3
checkpoint_interval: 300
cleanup_interval: 3600
```

llm:

```
default_provider: "local"
max_tokens: 4096
temperature: 0.7
timeout: 30
max_retries: 3
providers:
  <name>:
    type: <provider-type>
    endpoint: <url>
    enabled: true
    parameters:
      timeout: 30.0
      max_retries: 3
      streaming_support: true
      api_key: ""
  selection:
    strategy: "performance"
    fallback_enabled: true
    health_check_interval: 30
```

logging:

```
level: "info"
```

```
format: "text"
output: "stdout"
```

```
notifications:
  enabled: true
  rules:
    - name: "... "
      condition: "type==error"
      channels: ["slack", "email"]
      priority: urgent
      enabled: true
  channels:
    slack: { enabled, webhook_url, channel, username, timeout }
    telegram: { enabled, bot_token, chat_id, timeout }
    email: { enabled, smtp: { server, port, username, password,
      tls }, recipients, timeout }
    discord: { enabled, webhook_url, timeout }
```

Environment Variables

Required for Production: - HELIX_DATABASE_PASSWORD -
HELIX_AUTH_JWT_SECRET - HELIX_REDIS_PASSWORD

LLM Provider Keys (as needed): - OPENAI_API_KEY, ANTHROPIC_API_KEY,
GEMINI_API_KEY, XAI_API_KEY, DEEPSEEK_API_KEY, GROQ_API_KEY,
MISTRAL_API_KEY, COHERE_API_KEY, AZURE_OPENAI_API_KEY,
AWS_ACCESS_KEY_ID / AWS_SECRET_ACCESS_KEY

Notification Integrations: - HELIX_SLACK_WEBHOOK_URL -
HELIX_TELEGRAM_BOT_TOKEN, HELIX_TELEGRAM_CHAT_ID -
HELIX_EMAIL_SMTP_SERVER, HELIX_EMAIL_USERNAME, HELIX_EMAIL_PASSWORD
- HELIX_DISCORD_WEBHOOK_URL

Testing Strategy

Test Categories

1. **Unit tests:** Mocks allowed, *_test.go, -short flag
2. **Contract tests:** Real API schemas, no mocks
3. **Component tests:** Real subsystems wired together
4. **Integration tests:** Full app with real dependencies (-tags=integration)
5. **E2E challenges:** Complete user workflows against real LLM APIs
6. **Security tests:** OWASP compliance
7. **Performance tests:** Benchmarks

- 8. **Automation tests:** Provider/hardware automation (-tags=automation)
- 9. **Load tests:** Stress testing

Anti-Bluff Testing Rules

- Unit tests: Mocks OK
- **ALL other tests: Real infrastructure ONLY**
- Every PASS guarantees **Quality + Completion + Usability**
- Challenges fail on simulated/stubbed behavior
- No bare `t.Skip()` without SKIP-OK: `#<ticket>` marker

Docker Test Infrastructure

- `docker-compose.test.yml`: PostgreSQL 16, Redis 7, Memcached, Cognee, ChromaDB, Qdrant, Ollama, Prometheus, Grafana
- `docker-compose.full-test.yml`: Complete stack with mock-LLM server, Selenium, ChromeDP, SSH server + 3 workers, Cognee, Weaviate, mock-Slack, multicast router

Challenge Framework (tests/e2e/challenges/)

The most rigorous test system validates HelixCode by having it **generate real projects** and testing them: - **Challenge Definitions:** JSON specs (ASCII art generator, CLI task manager, JSON validator, notes API, tic-tac-toe TUI, URL shortener) - **Execution Flow:** Load spec → Call real LLM API → Parse generated code → Compile → Test → Runtime validation - **Validation Layers:** Directory structure, code quality, compilation, testing, functionality, runtime validation with diverse data - **Test Matrix:** Supports CLI, TUI, REST, WebSocket interfaces across 15+ providers and worker pool distributions

Test Scripts Summary

```
# Basic
cd HelixCode && make test

# Full infrastructure (recommended for validation)
make test-infra-up
make test-complete
make test-infra-down

# Individual categories
make test-unit-full
make test-integration-full
make test-e2e-full
make test-security-full
```



```
make test-load-full

# Legacy scripts
./run_tests.sh --all
./run_all_tests.sh
./run_integration_tests.sh
```

Docker Deployment

Production (docker-compose.yml)

Services: helixcode-server (8080, 2222), postgres:15, redis:7, nginx (80, 443), prometheus (9090), grafana (3000)

Quick Start

```
cd HelixCode
cp .env.example .env
# Edit .env with secure passwords
docker compose up -d
docker compose ps
curl http://localhost/health
```

Other Compose Files

File	Purpose
docker-compose-simple.yml	Minimal dev (postgres + redis only)
docker-compose.test.yml	Integration/E2E testing stack
docker-compose.full-test.yml	Zero-skip full test infrastructure
docker-compose.aurora-os.yml	Security-focused Aurora OS platform
docker-compose.harmony-os.yml	Distributed Harmony OS platform
docker-compose.specialized-platforms.yml	Combined Aurora + Harmony
docker/docker-compose.yml	Extended full-stack with Milvus, Elasticsearch, MLflow, Jaeger, Jupyter, Portainer

Deployment Patterns

- Healthchecks on every service
 - Docker profiles: monitoring, distributed, with-redis, production, dev, server
 - Isolated bridge networks per deployment
 - Named persistent volumes for all stateful services
 - `.env` file for secrets
-

Code Style & Development Conventions

Go Conventions

- Standard Go formatting: `go fmt ./...`
- Linting: `golangci-lint run ./...` (timeout 10m in CI)
- Vet: `go vet ./...`
- Table-driven tests with `t.Run()` subtests
- Build tags for integration/automation tests: `//go:build integration`

Project Conventions

- **Always work from `helix_code/` subdirectory**
- **Generate logo assets before first build:** `make logo-assets`
- **Database/Redis optional:** Disable by setting `database.host: ""`
- **Environment variables override config file**
- Use `internal/` for all core packages; no `pkg/` directory in active use
- Error handling: explicit, no silent failures
- Concurrent access: use `sync.RWMutex` or channel patterns

API Conventions

- REST API documented in `api/openapi.yaml` (OpenAPI 3.0.3)
 - Base path: `/api/v1`
 - Authentication: Bearer JWT via Authorization header
 - Health endpoint: `GET /health` (no auth required)
-

Security Considerations

Verified Security Features

- Password hashing: `bcrypt` (cost 12) with `argon2` fallback
- JWT with constant-time comparison

- CORS middleware, security headers (X-Frame-Options, CSP, HSTS)
- Rate limiting support in production config
- Session timeout, concurrent session limits, IP binding options
- Workflow `isDangerousCommand()` filter blocks `rm`, `dd`, `mkfs`, fork bombs, etc.
- Input validation in auth and server packages

Security Testing

- `security/security_test.go`: OWASP Top 10, SAST, DAST, credential scanning, TLS enforcement, input validation (path traversal, XSS, SQL injection, command injection, SSRF)
- File permission checks (0600 for configs)

Known Security Stubs

- `internal/security/security.go`: Simulated scanning (always returns clean)
- `cmd/security_test/main.go`: Entirely simulated security tests

Production Hardening

- Use `HELIX_AUTH_JWT_SECRET` with high entropy
- Enable PostgreSQL SSL in production
- Enable Redis authentication
- Configure CORS `allowed_origins` explicitly
- Enable audit logging
- Set `bcrypt_cost`: 14 in production

Universal Mandatory Constraints

Hard Stops (permanent, non-negotiable)

1. **NO CI/CD pipelines** (Note: existing workflow files in `.github/workflows/` are legacy and must not be expanded)
2. **NO HTTPS for Git** (SSH only)
3. **NO manual container commands** (orchestrator-owned)

Mandatory Development Standards

1. **100% Test Coverage** (unit, integration, E2E, automation, security, benchmark)
2. **Challenge Coverage** (every component)

3. **Real Data** (actual API calls, real DB, live services)
4. **Health & Observability** (health endpoints, circuit breakers)
5. **Documentation & Quality** (update docs with code changes)
6. **Validation Before Release** (full suite + all challenges)
7. **No Mocks in Production**
8. **Comprehensive Verification** (runtime, compile, structure, dependencies, compatibility)
9. **Resource Limits** (30-40% of host resources max)
10. **Bugfix Documentation** (root cause, affected files, fix, verification link)
11. **Real Infrastructure for All Non-Unit Tests**
12. **Reproduction-Before-Fix** (Challenge first, then fix)
13. **Concurrent-Safe Collections**

Definition of Done

A change is NOT done because code compiles. “Done” requires: - Pasted terminal output from a real run - No self-certification words without evidence - Demo commands that run against real artifacts - Loud skips with SKIP-OK: #<ticket> markers

CONST-035 — End-User Usability Mandate

A test or Challenge that PASSES is a CLAIM that the tested behavior **works for the end user of the product**.

The HelixAgent project has repeatedly hit the failure mode where every test ran green AND every Challenge reported PASS, yet most product features did not actually work — buggy challenge wrappers masked failed assertions, scripts checked file existence without executing the file, “reachability” tests tolerated timeouts, contracts were honest in advertising but broken in dispatch. **This MUST NOT recur in HelixCode.**

Every PASS result MUST guarantee: a. **Quality** — correct behavior under real inputs, edge cases, concurrency b. **Completion** — wired end-to-end with no stub/placeholder gaps c. **Full usability** — a user following documentation succeeds

A passing test that doesn’t certify all three is a **bluff** and MUST be tightened.

Bluff Taxonomy (each pattern observed and now forbidden)

- **Wrapper bluff** — assertions PASS but wrapper’s exit-code logic is buggy
- **Contract bluff** — system advertises capability but rejects it in dispatch
- **Structural bluff** — file exists but doesn’t contain working code
- **Comment bluff** — comment promises behavior code doesn’t have

- **Skip bluff** — `t.Skip("not running yet")` without SKIP-OK: #<ticket> marker

The taxonomy is illustrative, not exhaustive. Every Challenge or test added going forward **MUST** pass an honest self-review against this taxonomy before being committed.

Constitutional anchors (cascaded from CONSTITUTION.md)

Article XI §11.9 — Anti-Bluff Forensic Anchor

Verbatim user mandate: *“We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can’t be used! This **MUST NOT** be the case and execution of tests and Challenges **MUST** guarantee the quality, the completion and full usability by end users of the product!”*

Operative rule: **The bar for shipping is not “tests pass” but “users can use the feature.”** Every PASS in this codebase **MUST** carry positive runtime evidence captured during execution. Metadata-only / configuration-only / absence-of-error / grep-based PASS without runtime evidence are critical defects regardless of how green the summary line looks. No false-success results are tolerable.

Article XII §12.1 (CONST-042) — No-Secret-Leak

No API key, token, password, certificate, or other credential may be committed to any repository owned by HelixDevelopment or vasic-digital. All secrets live in `.env` files (mode 0600) listed in `.gitignore`. Any leak is a release blocker until rotated and post-mortemed.

Article XII §12.2 (CONST-043) — No-Force-Push

No force push, force-with-lease push, history rewrite, branch deletion of `main/` `master`, or upstream-overwriting operation may be performed without explicit, in-conversation user approval per operation. Authorization for one push does not extend further. Bypassing hooks / signing / protected-branch rules also requires explicit approval.

Article XIII §13.1 (CONST-044) — Continuation Document Maintenance Mandate

The docs/CONTINUATION.md document MUST be maintained in sync with actual programme state at all times. It is the authoritative resumption record for any CLI agent or LLM picking up the CLI-Agent Fusion programme. Every commit that advances state — task completion, feature close-out, push, known-issue discovery, deferred-item resolution, phase transition, submodule/remote add or remove — MUST update CONTINUATION in the same commit. Out-of-sync CONTINUATION is a **CRITICAL DEFECT** — same severity as a false-success test result under CONST-035 / Article XI §11.9. Cascade applies to every owned-by-us repo's three governance files. See CONSTITUTION.md Article XIII §13.1 for the full mandate (definition of out-of-sync, verification plan, cascade requirement).

CONST-045 — No Hardcoded Distribution Hosts

ALL container distribution targets SHALL be configured exclusively through CONTAINERS_REMOTE_HOST_N_* env vars in containers/.env. NO host (hostname, IP, user, key path, runtime, label) may be hardcoded in ANY source file, test, challenge, config template, script, or governance document. Adding/removing hosts = editing containers/.env only; NO code change. Tests SHALL read .env at runtime and skip with SKIP-OK: marker when CONTAINERS_REMOTE_ENABLED=false. See CONSTITUTION.md §CONST-045 for the full mandate.

CONST-036: LLMsVerifier Single Source of Truth Mandate

Rule: LLMsVerifier SHALL BE the sole authoritative source for: 1. All model metadata (names, IDs, context windows, capabilities) 2. All provider metadata (endpoints, auth types, supported models) 3. All verification status (verified, partial, failed, pending) 4. All scoring data (overall scores, capability scores, tier rankings)

Prohibition: NO hardcoded model lists, NO hardcoded provider lists, NO simulated model discovery. Any code path that presents a model or provider listing to a user MUST fetch that listing from the LLMsVerifier subsystem or its cached replica.

Anti-Bluff Verification: - Challenge script scripts/anti_bluff/verifier_hardcode_check.sh scans all Go source files for hardcoded model arrays. - The only permitted hardcoded data is the 7-entry fallback list in internal/verifier/fallback_models.go.

CONST-037: Model Provider Anti-Bluff Guarantee

Rule: Every model displayed to an end user MUST have been verified by LLMsVerifier within the last 24h. Models older than this MUST display a “stale” indicator and be deprioritized.

Anti-Bluff Testing: - Unit tests MAY mock the verifier client. - Integration tests MUST start the verifier server and perform real provider discovery. - The Makefile target `make test-verifier-integration` MUST exist and run without mocks.

CONST-038: Real-Time Model Status Accuracy

Rule: Model status (available, rate-limited, cooldown, offline, deprecated) displayed to users MUST reflect the actual state as known by LLMsVerifier within 60 seconds.

Polling vs. Push: - If WebSocket/SSE push is unavailable, the system MUST poll LLMsVerifier at most every 60s. - The TUI MUST display a “last updated” timestamp with every model listing. - Models in “cooldown” or “rate-limited” state MUST show the estimated recovery time if known.

CONST-039: All Providers and Models Integration Mandate

Rule: HelixCode MUST integrate with ALL providers that LLMsVerifier supports, subject only to: 1. The provider being explicitly disabled in configuration (`enabled: false`) 2. The API key being absent and the provider requiring one 3. The provider being marked deprecated in the verifier database

Minimum Provider Set (SHALL NOT be reduced without constitutional amendment): OpenAI, Anthropic, Gemini, DeepSeek, Groq, Mistral, xAI, OpenRouter, Ollama, Llama.cpp.

CONST-040: MCP / LSP / ACP / Embedding / RAG / Skills / Plugins Integration Mandate

Rule: LLMsVerifier integration SHALL extend beyond basic model listing to cover ALL capability dimensions:

1. **MCP:** The verifier MUST report which models support MCP tool calling.
2. **LSP:** The verifier MUST report code-analysis capabilities.
3. **ACP:** The verifier MUST report multi-agent coordination support.
4. **Embedding:** The verifier MUST report supports_embeddings for each model.
5. **RAG:** The verifier MUST report context-window sizes for chunking strategies.
6. **Skills / Plugins:** The verifier MUST track plugin compatibility.

Prohibition: Capability flags MUST NOT be hardcoded. The `Provider.GetCapabilities()` method MUST return data sourced from the verifier's `VerificationResult` fields.

CONST-046: No Hardcoded Content — All Text Must Be Dynamic (cascaded from CONSTITUTION.md)

Rule: NO user-facing text, prompt template, question text, error message, label, helper text, or explanatory content may be hardcoded as a static literal string in any source file. All such content MUST be: 1. Generated dynamically by an LLM at runtime based on user context (language, prompt content, session state), OR 2. Loaded from an i18n resource file (`.yaml`, `.json`, `.toml`) that supports locale overrides, OR 3. Composed programmatically from verifier metadata, configuration data, or provider responses.

Rationale: Hardcoded English text silently breaks the product for non-English users. A clarification question hardcoded as “Which file has the bug?” may be shown identically to a Japanese, Serbian, or Spanish user — producing an incoherent, unusable experience. Every piece of user-facing text MUST adapt.

Violation examples: `[]string{"Which file has the bug?"}`, hardcoded TUI labels, static error prose, hardcoded English prompt templates.

Compliant patterns: LLM-generated questions via `DetectAmbiguity(prompt)`, i18n-loaded labels with locale fallback, verifier-composed model descriptions.

Enforcement: make `lint` MUST include a hardcoded-text scan. Anti-bluff sweeps MUST flag obvious hardcoded-text patterns.

See CONSTITUTION.md §CONST-046 for the full mandate and cascade requirements.

CONST-047: Recursive Submodule Application Mandate (cascaded from CONSTITUTION.md)

Verbatim user mandate (2026-05-14): “Make sure all work we do is applied ALWAYS to all Submodules we control under our organizations (vasic-digital and HelixDevelopment) fully recursively everywhere with full bluff-proofing and comprehensive documentation, user manuals and guides and full tests and Challenges coverage!”

Every engineering deliverable produced for HelixCode MUST be applied — fully and recursively — to every owned submodule under the vasic-digital and HelixDevelopment GitHub organizations. Each owned submodule MUST receive in lockstep: (1) the same anti-bluff posture (CONST-035 / Article XI §11.9), (2) comprehensive documentation (README + user manuals + guides matching actual capabilities), (3) full tests + Challenges coverage with captured runtime evidence, (4) recursive propagation through nested submodules under the same orgs, (5) synchronized commits when meta-repo state advances a submodule’s surface.

Owned submodule baseline (direct children as of 2026-05-14): vasic-digital/{Containers, Security, Challenges, LLMsVerifier, Models}, HelixDevelopment/{HelixQA, DocProcessor, LLMOrchestrator, LLMProvider, VisionEngine, HelixAgent}, HelixDevelopment-s-Code/Website — plus any nested submodule under these orgs (full recursion).

Cascade requirement: This rule (verbatim or by CONST-047 ID reference) MUST appear in every owned-by-us submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md.

See CONSTITUTION.md §CONST-047 for the full mandate and verification details.

CONST-048: Full-Automation-Coverage Mandate (cascaded from constitution submodule §11.4.25)

Verbatim user mandate (2026-05-15): “Make sure that every feature, every functionality, every flow, every use case, every edge case, every service or application, on every platform we support is covered with full automation tests which will confirm anti-bluff policy and provide the proof of fully working capabilities, working implementation as expected, no issues, no bugs, fully documented, tests covered! Nothing less than this

does not give us a chance to deliver stable product! This is mandatory constraint which MUST BE respected without ignoring, skipping, slacking or forgetting it!"

No feature / functionality / flow / use case / edge case / service / application on any supported platform of HelixCode is deliverable until automation tests prove six invariants: (1) anti-bluff posture with captured runtime evidence (CONST-035); (2) proof of working capability end-to-end on target topology; (3) implementation matches documented promise; (4) no open issues/bugs surfaced; (5) full documentation in sync; (6) four-layer test floor (pre-build + post-build + runtime + paired mutation). Coverage ledger regenerated at release-gate sweep.

Cascade requirement: This rule (verbatim or by CONST-048 ID reference) MUST appear in every owned-by-us submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. See constitution submodule Constitution.md §11.4.25 for the full mandate.

CONST-049: Constitution-Submodule Update Workflow Mandate (cascaded from constitution submodule §11.4.26)

Verbatim user mandate (2026-05-15): "Every time we add something into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD we MUST FIRST fetch and pull all new changes / work from constitution Submodule first! All changes we apply MUST BE committed and pushed to all constitution Submodule upstreams! In case of conflict, IT MUST BE carefully resolved! Nothing can be broken, made faulty, corrupted or unusable! After merging full validation and verification MUST BE done!"

7-step pipeline: (1) fetch + pull first inside constitution submodule; (2) apply with §11.4.17 classification + verbatim mandate; (3) validate (meta-test + no conflict markers + cross-file consistency); (4) commit + push to EVERY configured upstream (governance files only, no `git add -A`); (5) careful conflict resolution preserving union — force-push forbidden (CONST-043); (6) post-merge cascade verification (CONST-047); (7) bump consuming .gitmodules pointer in SAME commit.

Cascade requirement: This rule (verbatim or by CONST-049 ID reference) MUST appear in every owned-by-us submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. See constitution submodule Constitution.md §11.4.26 for the full mandate.

CONST-050: No-Fakes-Beyond-Unit-Tests + 100%-Test-Type-Coverage Mandate (cascaded from constitution submodule §11.4.27)

Verbatim user mandate (2026-05-15): “Mocks, stubs, placeholders, TODOs or FIXMEs are allowed to exist ONLY in Unit tests! All other test types MUST interact with real fully implemented System! No fakes, empty implementations or bluffing is allowed of any kind! All codebase of the project MUST BE 100% covered with every supported test type: unit tests, integration tests, e2e tests, full automation tests, security tests, ddos tests, scaling tests, chaos tests, stress tests, performance tests, benchmarking tests, ui tests, ux tests, Challenges (fully incorporating our Challenges Submodule). EVERYTHING MUST BE tested using helix_qa (fully incorporating helix_qa Submodule). helix_qa MUST BE used with all possible written tests suites (test banks) for every applications, service, platform, etc and execution of the full helix_qa QA autonomous sessions! All required dependency Submodules MUST BE added into the project as well (fully recursive!!!).”

Two invariants. **(A)** Mocks/stubs/fakes/placeholders/TODOs/FIXMEs/“for now”/empty-implementation patterns PERMITTED only in unit-test sources; non-unit tests (integration, E2E, full-automation, security, DDoS, scaling, chaos, stress, performance, benchmarking, UI, UX, Challenges, HelixQA) MUST exercise the real, fully implemented HelixCode system. Production code MUST NOT import `helix_code/internal/mocks/`. **(B)** Codebase MUST be covered by every supported test type the domain warrants: unit, integration, E2E, full-automation, security, DDoS, scaling, chaos, stress, performance, benchmarking, UI, UX, Challenges (vasic-digital/Challenges submodule at `./challenges/` fully incorporated), `helix_qa` (HelixDevelopment/HelixQA submodule at `./helix_qa/` fully incorporated, with full autonomous QA sessions executing every registered test bank with captured wire evidence).

Required dependency submodules (recursive per CONST-047): Challenges (`git@github.com:vasic-digital/Challenges.git` at `./challenges/`) + `helix_qa` (`git@github.com:HelixDevelopment/HelixQA.git` at `./helix_qa/`) + any other functionality submodules under vasic-digital/HelixDevelopment orgs HelixCode depends on. Pointers bumped to upstream HEAD in same commit as cascade work (CONST-049 step 7).

Cascade requirement: This rule (verbatim or by CONST-050 ID reference) MUST appear in every owned-by-us submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. See constitution submodule Constitution.md §11.4.27 for the full mandate.

CONST-051: Submodules-As-Equal-Codebase + Decoupling + Dependency-Layout Mandate (cascaded from constitution submodule §11.4.28)

Verbatim user mandate (2026-05-15): “All existing Submodules in the project that we are controlling and belong to some our organizations (vasic-digital, HelixDevelopment, red-elf, ATMOSphere1234321, Bear-Suite, BoatOS123456, Helix-Flow, Helix-Track, Server-Factory) are equal parts of the project’s codebase! We MUST work on that code as much as we do with main project’s codebase! ... We MUST NEVER modify Submodules to bring into them any project specific context ... All Submodule dependencies that are used by Submodule MUST BE accessed from the root of the project! We MUST NOT have nested Submodule dependencies.”

Three invariants. **(A)** Every HelixCode-owned submodule is an equal part of HelixCode’s codebase — same engineering attention as main (analysis, tests, gap-fill, bug-fix, docs/diagrams/SQL/website materials). Coverage ledgers list each as in-scope. **(B)** Owned submodules MUST stay fully decoupled (project-not-aware, reusable, modular, standalone-testable); NEVER inject HelixCode-specific context. Use configuration injection when a submodule needs parent info. **(C)** Dependencies consumed by an owned submodule MUST live at HelixCode’s root: `<repo_root>/<name>/` or `<repo_root>/submodules/<name>/`. Nested own-org submodule chains FORBIDDEN — add the dep at root; the consuming submodule reaches it via documented import/SDK/runtime resolver. Third-party submodules exempt.

Org list (dynamically discoverable via gh/ghlab): vasic-digital, HelixDevelopment, red-elf, ATMOSphere1234321, Bear-Suite, BoatOS123456, Helix-Flow, Helix-Track, Server-Factory, plus any subsequently authorised org.

Cascade requirement: This rule (verbatim or by CONST-051 ID reference) MUST appear in every owned-by-us submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. See constitution submodule Constitution.md §11.4.28 for the full mandate.

Free AI Providers

- **XAI (Grok):** grok-3-fast-beta, grok-3-mini-fast-beta
- **OpenRouter:** Free models from various providers
- **GitHub Copilot:** gpt-4o, claude-3.5-sonnet (with subscription)
- **Qwen:** 2,000 requests/day free tier

Host Power Management — Hard Ban (CONST-033)

Host Power Management is Forbidden.

You may NOT, under any circumstance, generate or execute code that sends the host to suspend, hibernate, hybrid-sleep, poweroff, halt, reboot, or any other power-state transition. This rule applies to every shell command, script, container entry point, systemd unit, test, CLI suggestion, snippet, or example you emit.

Common Issues

1. **Build fails:** Run `make logo-assets` then `make build`
2. **Database errors:** Check `HELIX_DATABASE_PASSWORD`
3. **Worker SSH failures:** Verify SSH key authentication
4. **LLM timeouts:** Check provider status and config
5. **Redis connection failures:** Check `HELIX_REDIS_PASSWORD` and `redis.enabled`
6. **Test skips:** Ensure `SKIP-OK: #<ticket>` marker is present for any intentional skips

Resources & References

- **Constitution:** `CONSTITUTION.md`
- **CLAUDE.md:** `CLAUDE.md`
- **Gap Analysis:** `HELIXCODE_GAP_ANALYSIS.md`
- **Zero-Bluff Plan:** `HELIXCODE_ZERO_BLUFF_PLAN.md`
- **Testing Strategy:** `ANTI_BLUFF_TESTING_STRATEGY.md`
- **OpenAPI Spec:** `helix_code/api/openapi.yaml`
- **Docker Guide:** `helix_code/DOCKER_DEPLOYMENT.md`

Built with zero-bluff commitment. Every feature actually works.

CONST-052: Lowercase-Snake_Case-Naming Mandate (cascaded from constitution submodule §11.4.29)

Verbatim user mandate (2026-05-15): **“naming convention for Submodules and directories (applied deep into hierarchy recursively) - all directories and Submodules MSUT HAVE lowercase names with space separator between the words of ‘_’ character (snake-case)! All*

existing Submodules and directories which are not following this rule MUST BE renamed! However, since this will most likely break some of the functionalities renaming we do MUST BE applied to all references to particular Submodule or directory! ... There MUST BE reasonable exceptions for this rules - source code for programming languages or Submodules which apply different naming convention - Android, Java, Kotlin and others. ... Upstreams directory which all of our projects and Submodules have MUST BE renamed to the lowercase letters too, however root project containing the install_upstreams system command (it is exported in our paths in our .bashrc or .zshrc) MUST BE updated to fully work with both Upstreams and upstreams directory. ... NOTE: Rules lowercase / snake-case do apply to all project files as well and references to it and from them!“*

Every directory, submodule, and file in HelixCode MUST use lowercase snake_case names. Existing non-compliant names (helix_code/, challenges/, containers/, helix_agent/, helix_qa/, security/, github_pages_website/, upstreams/, dependencies/, etc.) MUST be renamed as part of the phased migration opened by this clause. Every reference (configs, docs, links, source-code imports, governance files) MUST be updated atomically with the rename — reference drift after a rename is a CONST-052 violation of equal severity to the rename itself.

Common-sense exceptions (technology-preserving): language-mandated case for Java/Kotlin/Android/Apple/C#/Swift INSIDE the language root (submodule root follows our convention; subtree follows language convention); vendor/upstream third-party submodules keep upstream names; build artefacts (node_modules, __pycache__, .git, target, build, bin) keep tool-mandated names. The test “does renaming break the technology?” trumps the rule.

upstreams/ → upstreams/ transition: the constitution submodule’s install_upstreams.sh (exported via .bashrc/.zshrc) supports BOTH upstreams/ and upstreams/ directory layouts (commit 45d3678 of the constitution submodule); lowercase wins when both present.

Test coverage of renames (per CONST-050(B)): every rename batch ships with (i) regression test verifying every reference now resolves, (ii) full test-type matrix run post-rename, (iii) anti-bluff wire-evidence captured.

Phased execution per the operator’s explicit instruction: comprehensive brainstorming → phase-divided plan → fine-grained tasks/subtasks → every change covered by every applicable test type. §11.4.20 subagent delegation for cross-cutting rename sweeps.

Cascade requirement: This anchor (verbatim or by CONST-052 ID reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and

AGENTS.md. Severity-equivalent to a §11.4 PASS-bluff at the reference-integrity layer. No escape hatch beyond the common-sense exceptions enumerated above. See constitution submodule Constitution.md §11.4.29 for the full mandate.

CONST-053: .gitignore + No-Versioned-Build-Artifacts Mandate (cascaded from constitution submodule §11.4.30)

Verbatim user mandate (2026-05-15): *“every project module, every Submodule, every service and application MUST HAVE proper .gitignore file! We MUST NOT git version build artifacts, cache files, tmp files, main .env file(s) or any files containing sensitive data, API keys or token! Any build derivative which we can recreate by executing proper mechanism for generating MUST NOT be versioned! We MUST pay attention what is going to be committed every time we are preparing to execute commit! If any violation is detected it MUST be fixed before commit is executed!”*

Every project module, owned-by-us submodule, service, and application MUST ship a proper .gitignore. Forbidden-from-version-control classes:

1. **Build artefacts:** /bin/, /build/, /dist/, /out/, target/, *.exe, *.dll, *.so, *.dylib, *.a, *.o, *.class, *.pyc, generator-produced files when the generator is committed.
2. **Cache files:**
__pycache__/, .pytest_cache/, .mypy_cache/, .ruff_cache/, node_modules/, .next/, .cache/, .gradle/, .terraform/, language-server caches.
3. **Temp files:** *.tmp, *.swp, *~, .DS_Store, Thumbs.db, *.orig, *.rej.
4. **Sensitive-data files:** .env, .env.* (allow .env.example placeholder only — no real secrets even as examples), *.pem, *.key, *.crt, id_rsa*, id_ed25519*, .netrc, secrets/, api_keys.sh.
5. **Generated reports/logs:** *.log, coverage.out, htmlcov/, runtime captures unless reference assets.
6. **OS/IDE personal state:** .idea/, .history/, .vscode/ (except shared settings).

Anti-bluff invariant: .gitignore line alone is not sufficient — no file matching the forbidden patterns may be CURRENTLY TRACKED. A tracked *.log despite the ignore-line is a violation of equal severity to no ignore-line at all.

Pre-commit attention: every commit author (human OR agent) MUST inspect `git diff --staged + git status` BEFORE executing the commit. Forbidden-class hits abort the commit until fixed (un-stage, add to .gitignore, scrub if already-tracked). Gate CM-GITIGNORE-PRECOMMIT-AUDIT + paired mutation.

Secret-leak intersection (CONST-042 / §11.4.10): a `.env` leak is BOTH a CONST-053 and a CONST-042 violation; rotation + post-mortem required.

Recreatable-content test: if a documented mechanism regenerates the file from sources, it is a build derivative and MUST be ignored. The committed sources MUST include the generator.

Cascade requirement: This anchor (verbatim or by CONST-053 ID reference) MUST appear in every owned submodule's `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Severity-equivalent to a §11.4 PASS-bluff at the repository-hygiene layer. See constitution submodule `Constitution.md` §11.4.30 for the full mandate.

CONST-054: Submodule-Dependency-Manifest Mandate (cascaded from constitution submodule §11.4.31)

Verbatim user mandate (2026-05-15): *“We MUST HAVE mechanism for each Submodule to determine / know what are its Submodule dependencies so new projects or palces we are incorporate them can add these Submodules to the project root and make them available! Suggested idea is configuration file with expected Submodules Git ssh urls perhaps? New project can read it, and recursively add each Submodule to the root of the project and install / expose it to everyone.”*

Every owned-by-us submodule MUST ship `helix-deps.yaml` at its root declaring its own-org dependencies. Schema: `schema_version`, `deps: [{name, ssh_url, ref, why, layout: flat|grouped}], transitive_handling: {recursive, conflict_resolution}, language_specific_subtree`. Tooling: `incorporate-submodule <ssh-url>` adds the submodule at the parent project's canonical path (CONST-051(C)), reads `helix-deps.yaml`, recurses for each declared dep, aborts on conflicting refs, emits `<root>/ .helix-manifest.yaml` audit record.

Anti-bluff guarantee: every manifest paired with a Challenge that bootstraps a throwaway consuming project, runs `incorporate-submodule`, asserts produced layout matches the manifest, runs the submodule's own tests against the bootstrapped layout, captures wire evidence per §11.4.2. A manifest without this proof is a CONST-054 violation.

§11.4.31 / CONST-054 is the **operational complement** of CONST-051(C): nested own-org submodule chains are FORBIDDEN — manifests are the bridge that lets consumers reconstruct the dependency graph at the parent root.

Cascade requirement: This anchor (verbatim or by CONST-054 ID reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to §11.4 PASS-bluff at the dependency-graph layer. See constitution submodule Constitution.md §11.4.31 for the full mandate.

CONST-055: Post-Constitution-Pull Validation Mandate (cascaded from constitution submodule §11.4.32)

Verbatim user mandate (2026-05-15): “Every time we fetch and pull new changes on constitution Submodule we MUST process the whole project and all Submodule (deep recursively) for validation and verification taht every single rule or mandatory constraint is followed and respected! If it is not, IT MUST BE!”

Whenever a project's constitution submodule is fetched + pulled with any content change, the project MUST run `scripts/verify-all-constitution-rules.sh` BEFORE the new constitution HEAD is treated as canonical for any other work. The sweep re-runs the governance-cascade verifier AND every implementable rule gate (CONST-053 .gitignore audit, CONST-051(C) nested-own-org-chain audit, CONST-052 case audit, CONST-050(A) mock-from-production audit, CONST-035 anti-bluff smoke, etc.) against the post-pull tree. Failures populate the project's Issues tracker per §11.4.15 (Status: Reopened, Type: Bug); closure requires positive-evidence per §11.4.

Pull-time invocation: `git submodule update --remote constitution` triggers the sweep automatically (post-update hook OR commit-wrapper invocation). Operator-explicit manual invocation also available.

Anti-bluff: the sweep's own meta-test (paired mutation per §1.1) plants a known violation of each enforced gate and asserts the sweep reports FAIL for the planted gate. A sweep that exits PASS without running every implementable gate is a CONST-055 violation.

CONST-055 is the **enforcement engine** for every other §11.4.x and CONST-NNN rule — without it, new rules cascade as anchors but never get enforced.

Cascade requirement: This anchor (verbatim or by CONST-055 ID reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to §11.4 PASS-bluff at the constitutional-enforcement layer. See constitution submodule Constitution.md §11.4.32 for the full mandate.

CONST-056: Mandatory install_upstreams on clone/ add Mandate (cascaded from constitution submodule §11.4.36)

Verbatim user mandate (2026-05-15): *“Every Submodule or Git repository we add or clone MUST BE upstreams installed using Upstreamable utility which MUST BE available through exported paths of the host system (in .bashrc or .zhrc) using install_upstreams command executed from the root of the cloned (added) repository - only if in it is Upstreams or upstreams directory present with bash script files (recipes) for all repository’s upstreams!”*

Every clone / add of a Git repository under HelixCode MUST be followed by install_upstreams invocation from the repository’s root IF its tree contains upstreams/ (or legacy upstreams/ per CONST-052 transition) populated with *.sh recipe files. The utility (installed on operator’s PATH via .bashrc/.zshrc; implementation in the constitution submodule’s install_upstreams.sh — already supports BOTH directory names since constitution commit 45d3678) reads the recipe files, configures every declared upstream as a named git remote, and fans out origin push URLs.

Skipping the invocation when upstreams/ is present silently breaks §2.1 (multi-upstream push is the norm) — the next push lands on only one upstream. Gate CM-INSTALL-UPSTREAMS-ON-CLONE + paired mutation. Automation: the future incorporate-submodule per CONST-054 auto-invokes; manual invocation supported. Pre-commit check: `git remote -v | grep -c push` reports expected count.

Cascade requirement: This anchor (verbatim or by CONST-056 ID reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. See constitution submodule Constitution.md §11.4.36 for the full mandate.

CONST-057: Type-aware Closure-Status Vocabulary (cascaded from constitution submodule §11.4.33)

Every project tracking work items by Type per §11.4.16 MUST close them with the Type-appropriate terminal ****Status:**** value, drawn from this 3-element closed map:

Item **Type:**	Closure **Status:** value
Bug	Fixed (→ Fixed.md)

Item **Type:**	Closure **Status:** value
Feature	Implemented (→ Fixed.md)
Task	Completed (→ Fixed.md)

The (→ Fixed.md) suffix is preserved across all three so the existing migration-discipline tooling (atomic Issues.md → Fixed.md move per §11.4.19) keeps working without per-Type branching. Generators (in HelixCode: the §11.4.93 workable-items export DB exporter — which SUPERSEDES the markdown-derived generate_issues_summary.sh / generate_fixed_summary.sh, both of which now REFUSE to run — and the §11.4.23 colorizer) MUST treat the three terminal values as semantically equivalent (all “closed, positive evidence captured”) while preserving the literal in the emitted document.

Closing a Feature with Fixed (→ Fixed.md) or a Task with Implemented (→ Fixed.md) is a CONST-057 violation. Gate CM-CLOSURE-VOCAB-TYPE-AWARE walks every Fixed.md heading + every Issues.md heading whose ****Status:**** is one of the three terminal values and asserts the Status-Type match. Composes with §11.4.15 / §11.4.16 / §11.4.19 / §11.4.23.

Cascade requirement: This anchor (verbatim or by CONST-057 ID reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. See constitution submodule Constitution.md §11.4.33 for the full mandate.

CONST-058: Reopened-Source Attribution Mandate (cascaded from constitution submodule §11.4.34)

Every Issues.md (or equivalent project tracker) heading whose ****Status:**** is Reopened MUST carry, within 8 non-blank lines of the heading, a ****Reopened-Details:**** line capturing four sub-facts:

- **By:** AI or User (source-of-truth observer who flipped the status). AI covers in-loop reopens (test failure, gate regression, captured-evidence retrospect). User covers operator-side observations (manual testing, end-user report, design reconsideration).
- **On:** ISO date (YYYY-MM-DD).
- **Reason:** one-line cause classification — chosen from the closed vocabulary { test-failed | manual-testing-detected | captured-evidence-contradicts | end-user-report | cycle-re-discovered | design-reconsidered }. Other values permitted with explicit Reason: <free text> annotation but the closed list MUST be tried first.
- **Evidence:** path to or short description of the captured artefact justifying the reopen — log file, recording, gate failure ID, operator quote, etc. Reopens

without evidence are §11.4.6 / §11.4.7 violations (demotion from Fixed requires captured evidence under the conditions that re-exposed the defect).

The `Issues_Summary.md` Status column MUST distinguish the four Reopened sub-states by source so a sweep query for “reopens by AI in the last 30 days” is mechanically possible. Suggested column rendering: Reopened (AI: test-failed) vs Reopened (User: manual-testing). Gate CM-ITEM-REOPENED-DETAILS mirrors CM-ITEM-OPERATOR-BLOCKED-DETAILS (§11.4.21 walk pattern). Composes with §11.4.6 / §11.4.7 / §11.4.15 / §11.4.21.

Cascade requirement: This anchor (verbatim or by CONST-058 ID reference) MUST appear in every owned submodule’s `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. See constitution submodule `Constitution.md` §11.4.34 for the full mandate.

CONST-059: Canonical-Root Inheritance Clarity (cascaded from constitution submodule §11.4.35)

The **constitution submodule’s** three files (`constitution/Constitution.md`, `constitution/CLAUDE.md`, `constitution/AGENTS.md`) ARE the **canonical root** (also called the **parent** files). They contain only universal rules per §11.4.17.

The consuming project’s **repository-root files** (`<project-root>/CLAUDE.md`, `<project-root>/AGENTS.md`, optionally `<project-root>/Constitution.md`) are **consumer extensions**. They MUST start with the inheritance pointer (either the Claude-Code native `@constitution/CLAUDE.md import` or the portable `## INHERITED FROM constitution/CLAUDE.md` heading). They contain only project-specific rules per §11.4.17.

When in doubt about which file to edit: universal rule → constitution submodule’s file; project-specific rule → consumer’s file. Default consumer-side when uncertain (§11.4.17 — narrower scope is cheap to widen).

Terminology: “the parent CLAUDE.md” / “the root Constitution” → constitution-submodule file at `constitution/<filename>`; “the project CLAUDE.md” / “this project’s AGENTS.md” → consumer-side file at `<project-root>/<filename>`.

No silent demotion or silent promotion. Moving a rule between layers MUST be a visible commit — `git mv` of a section if it’s a clean clone, or explicit Lifted from `<project>` to constitution per §11.4.35 / Demoted from constitution to `<project>` per §11.4.35 commit-message annotation.

Gate CM-CANONICAL-ROOT-CLARITY verifies (a) consumer’s `CLAUDE.md` opens with the inheritance pointer, (b) constitution submodule’s three files are present at the expected path, (c) no `## INHERITED FROM` block in the constitution

submodule's own files (those ARE the source-of-truth, not consumers). Composes with §11.4.17.

Cascade requirement: This anchor (verbatim or by CONST-059 ID reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. See constitution submodule Constitution.md §11.4.35 for the full mandate.

CONST-060: Fetch-before-edit Mandate (cascaded from constitution submodule §11.4.37)

Verbatim user mandate (2026-05-15): *“Make sure that feedback_fetch_before_edit memory rule is part of our constitution Submodule - the root Consitution, AGENTS.MD and CLAUDE.MD. Validate and verify that Proejct-Toolkit and all Submodules do inherit all of them! Follow the constitution Submodule documentation for details.”*

The FIRST git-touching action of every session, on every consuming project (owned or third-party), MUST be:

```
git fetch --all --prune
git log --oneline HEAD..@{u}
git submodule foreach --recursive 'git fetch --all --prune --quiet'
```

If HEAD..@{u} is non-empty, integrate the upstream changes BEFORE any local edit. Acting on stale local state produces three failure modes documented in the originating §11.4.37 incident (multi-agent / parallel-session work): (1) **redundant work** — the agent re-does what a parallel session already finished, (2) **false confidence** — completion reports for already-done work, (3) **divergent history** — duplicate sibling commits that double the conflict surface on next push.

Anti-bluff invariant: the fetch+log check MUST produce captured evidence — the actual HEAD..@{u} output, even if empty. Skipping the check on the basis of “I just fetched” or “nothing could have changed in the last N minutes” is a §11.4.6 (no-guessing) violation: the remote state is not knowable without a fetch.

Cascade requirement: This anchor (verbatim or by CONST-060 ID reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to §11.4 PASS-bluff at the parallel-session-coordination layer. See constitution submodule Constitution.md §11.4.37 for the full mandate.

§11.4.68 — Positive Sink-Side / Downstream Evidence Mandate (cascaded from constitution submodule §11.4.68)

Verbatim user mandate (2026-05-20): *“We still do not hear any audio played from D3 device! Arvus Web Dashboard when we play music from D3 shows nothing for Codec In Use! This MUST BE investigated and fixed! How come we passed the tests with Arvus validation? What were values for the Codec In Use field? Empty means nothing! This is not working! It MUST BE FIXED, TESTED AND VERIFIED WITH FULL AUTOMATION TESTING ASAP!!!”*

A test that asserts audio or video routing PASS MUST capture and verify **positive sink-side or downstream evidence** — never config-only, never metadata-only, never PCM-open-state-only. At least one of the closed enumeration MUST be captured for every audio/video routing PASS: (1) sink-side codec-state with non-empty Codec-In-Use matching the expected codec regex; (2) strictly-positive PCM frames-written delta from `/proc/asound/. . . /status hw_ptr`; (3) ALSA ELD/EDID-Like-Data showing negotiated channel count + format; (4) `ffprobe-on-captured-mp4` with non-zero frame count + expected codec/resolution/fps; (5) recording-analyzer event match per §11.4.2/§11.4.5; (6) `tinycap` RMS amplitude above the line-level floor. Empty / `<unreachable>` / `<N.E.>` / `<None>` placeholders are NOT positive evidence; a missing-but-required sink is OPERATOR-BLOCKED (release-blocker), never SKIP, never PASS. No escape hatch — no `--skip-sink-evidence`, `--allow-empty-codec`, `--sink-unreachable-is-pass`, `--metadata-only-suffices` flag exists.

Cascade requirement: This anchor (verbatim or by §11.4.68 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to a §11.4 PASS-bluff at the sink-side-evidence layer. **Canonical authority:** constitution submodule Constitution.md §11.4.68 for the full mandate.

§11.4.70 — Subagent-Driven Execution Is The Default (cascaded from constitution submodule §11.4.70)

Verbatim user mandate (2026-05-20): *“Always do if possible Subagent-driven! Add this into our root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md. This should be the default choice ALWAYS!”*

When executing implementation plans (or any task-decomposed execution flow), the **default execution model is subagent-driven** per superpowers : subagent-driven-development. Inline execution is permitted ONLY when (a) the task is trivial AND fits a single sub-300-line edit, OR (b) the operator explicitly requests inline at brainstorm-handoff time. Subagent-driven is the default because it gives isolated context per task, naturally enforces two-stage review, is parallel-PWU compatible (§11.4.58), creates an anti-bluff seam (§11.4), and survives operator absence. No escape hatch — `--inline-execution-required`, `--no-subagents`, `--monolithic-execution` are NOT permitted flags. Skipping subagent-driven for non-trivial work without recorded operator authorisation is itself a §11.4 PASS-bluff.

Cascade requirement: This anchor (verbatim or by §11.4.70 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to a §11.4 PASS-bluff at the execution-model layer. **Canonical authority:** constitution submodule Constitution.md §11.4.70 for the full mandate.

§11.4.71 — Pre-Push Fetch + Investigate + Integrate Mandate (cascaded from constitution submodule §11.4.71)

Verbatim user mandate (2026-05-20): *“before pushing changes to any upstream for any repository - main repo or Submodule, we MUST fetch and pull all changes. Once these are obtained WE MUST investigate what is different compared to head position we were on last time before fetching and pulling new changes! We MUST understand what is done and for what purpose, easpecially how that does affect our project and our System in general! Any mandatory changes or improvements required by fresh changes we just have brought in MUST BE incorporated, covered with all supported types of the tests which will produce as a result of its success execution REAL PROOFS of working for all componetns and functionalities covered and work fully in anti-bluff manner!”*

The everyday-push variant of §11.4.41. EVERY push (every repository — main + every submodule) MUST follow the 5-step cycle: (1) fetch all remotes (`git fetch --all --prune --tags`, capture stdout); (2) pull all upstream branches whose tip differs, resolving conflicts per consumer judgment (never auto--ours/--theirs); (3) investigate the diff vs OUR previous HEAD — read EVERY foreign commit's body, understand what/why/how-it-affects-our-system; (4) integrate mandatory changes with §11.4.4(b) four-layer coverage + §11.4.43 TDD-fix discipline, every PASS carrying §11.4.5 captured-evidence (REAL PROOFS, not

metadata-only); (5) only then push, verifying with `git ls-remote post-push`. No escape hatch — no `--skip-fetch`, `--no-investigate`, `--fast-push`, `--trust-upstream` flag.

Cascade requirement: This anchor (verbatim or by §11.4.71 reference) MUST appear in every owned submodule’s `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Severity-equivalent to a §11.4 PASS-bluff at the push-discipline layer. **Canonical authority:** constitution submodule `Constitution.md` §11.4.71 for the full mandate.

§11.4.72 — Audio Top-Priority Mandate (cascaded from constitution submodule §11.4.72)

Verbatim user mandate (2026-05-20): *“Make sure all fixes for audio are always top priority in main working stream!”*

The conductor (main working stream — Claude Code session, AI agent, or human operator) MUST treat audio fixes as the highest-priority class on the serial dispatch queue. Any time the conductor faces a choice between dispatching an audio task vs a non-audio task on the SAME serial resource, the audio task wins. Parallel BACKGROUND subagents (research, refactors, infrastructure documentation) MAY run concurrently with audio work but do NOT preempt audio on the main-stream serial dispatch queue. No escape hatch — there is no “but this non-audio task is faster” or “but this research is more interesting” override; audio-stack regressions are user-perceptible and high-impact while research and refactors can wait.

Cascade requirement: This anchor (verbatim or by §11.4.72 reference) MUST appear in every owned submodule’s `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Severity-equivalent to a process violation at the dispatch-priority layer. **Canonical authority:** constitution submodule `Constitution.md` §11.4.72 for the full mandate.

§11.4.73 — Main-Specification Document Versioning + Revision Discipline (cascaded from constitution submodule §11.4.73)

Verbatim user mandate (2026-05-20): *“Make sure everything we add now in previous and upcoming requests IS ALWAYS applied to the main specification — if we have one. Since all these are not major changes we could increase Specification version per change for secondary version instead of the primary. Primary version MUST BE increased for much bigger levels of changes! Add this into root (constitution Submodule)*

Constitution.md, CLAUDE.md and AGENTS.md as mandatory rule / constraint applicable ONLY IF we have something like the main specification document or we do recognize something like the main specification document. Document MUST BE updated ALWAYS to follow the versioning rules we are applying here + revision and other properties we have!”

Applies **only when a project recognises a main specification document**. When it does: (1) every additive operator requirement, refinement, or accepted recommendation MUST be applied to the spec before or as part of the implementing work; (2) spec versioning has two axes — *primary* (V1/V2/V3, bumped for major rewrites by explicit operator decision, old versions archived) and *secondary* (the §11.4.61 metadata-table Revision integer, bumped for every other change); (3) the metadata table MUST stay current (Revision, Last modified, Status summary, Fixed); (4) propagated copies of the rule MUST reference the active specification. V<primary>.md, not a stale archive; (5) on primary bump the old file moves to <spec-dir>/archive/ with Status: superseded. Classification: universal, applicable conditionally per the scope condition.

Cascade requirement: This anchor (verbatim or by §11.4.73 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to a release blocker when a project has a main spec and lets it drift. **Canonical authority:** constitution submodule Constitution.md §11.4.73 for the full mandate.

§11.4.74 — Submodule-Catalogue-First Discovery + Extend-Don’t-Reimplement (cascaded from constitution submodule §11.4.74)

Verbatim user mandate (2026-05-20): “We MUST ALWAYS check which already developed features / functionalities do exist as a part of our comprehensive Submodules catalogue located in vasic-digital and HelixDevelopment organizations on GitHub and GitLab both! Project MUST BE aware of all its existence so we do not implement same things multiple times if they are already done as some of existing universal, reusable general development purpose Submodules! For any missing features that some Submodules we incorporate may be missing we MUST IMPLEMENT the properly and extend those Submodules further! We do control all of the and we CAN and MUST maintain and extend the regularly! All development cycle rules we have MUST BE applied to them and fully respected!”

Before scaffolding ANY new module, package, helper, or utility, the contributor (human or AI agent) MUST: (1) survey the canonical Submodule catalogue — vasic-digital and HelixDevelopment on both GitHub AND GitLab; (2) inventory existing Submodules; (3) reuse before reimplement — if a Submodule provides the functionality (or 80%+ of it), add it as a Git submodule rather than write fresh; (4) extend in-place when 80%+ matches but features are missing — add the missing features TO THAT SUBMODULE (PR upstream + bump pointer), never as a duplicating consuming-project helper; (5) apply all development-cycle rules to those Submodules; (6) document the survey result in the feature’s tracker entry with a Catalogue-Check: field (reuse <org/repo>@<sha> / extend <org/repo>@<sha> / no-match <date>). Classification: universal.

Cascade requirement: This anchor (verbatim or by §11.4.74 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to a process violation; duplicate implementations landed without catalogue check are release blockers. **Canonical authority:** constitution submodule Constitution.md §11.4.74 for the full mandate.

§11.4.69 — Universal Sink-Side Positive-Evidence Taxonomy + Mechanical Enforcement (cascaded from constitution submodule §11.4.69)

Verbatim user mandate (2026-05-20): *“THIS MUST HAPPEN NEVER AGAIN!!! We MUST HAVE this all working! Not just for audio but for every single piece of the System!!! Proper full automation when executed with success MUST MEAN that manual testing will be as much positive at least regarding the success results! ... Solution MUST BE universal, generic that solves working flows for all System components and for all future and all existing projects! ... Everything we do MUST BE validated and verified with rock-solid proofs and anti-bluff policy enforcement and fulfillment!”*

Universal generalisation of §11.4.68 (audio-specific) across every user-visible feature class. Every user-visible feature MUST map to one entry in the closed-set §11.4.69 sink-side evidence taxonomy (audio_output, audio_input, video_display, network_throughput, network_connectivity, bluetooth_a2dp, bluetooth_pair, touch_input, sensor, gpu_render, storage_read, storage_write, mediacodec_decode, mediacodec_encode, miracast, cast, boot_service, package_install, permission_grant, wifi_link, wifi_throughput, ethernet_link, display_topology, drm_playback, subtitle_render — open to additions, never contraction). Every PASS for a feature in the taxonomy MUST cite a captured-evidence artefact path matching the required evidence shape. New helper contracts (additive during

grace, mandatory after 2026-06-19): `ab_pass_with_evidence` <description> <evidence_path> (verifies path exists + non-empty), `ab_skip_with_reason` <description> <closed-set-reason> (reasons: `geo_restricted`, `operator_attended`, `hardware_not_present`, `topology_unsupported`, `network_unreachable_external`, `feature_disabled_by_config`; forbids `network_unreachable_external` for any taxonomy feature with a sink-side probe); bare `ab_pass` deprecated (WARN pre-grace, FAIL post-grace). Three pre-build gates + paired §11.1 mutations: CM-SINK-EVIDENCE-PER-FEATURE, CM-NO-FAIL-OPEN-SKIP, CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE. No escape hatch — no `--skip-evidence`, `--config-only-pass`, `--allow-fail-open-skip`, `--legacy-ab-pass-permitted` flag.

Cascade requirement: This anchor (verbatim or by §11.4.69 reference) MUST appear in every owned submodule's `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Propagation gate CM-COVENANT-114-69-PROPAGATION enforces the anchor literal across the consumer fleet; paired mutation strips the literal → gate FAILs. Severity-equivalent to a §11.4 PASS-bluff at the sink-side-evidence layer.
Canonical authority: constitution submodule `Constitution.md` §11.4.69 for the full mandate.

§11.4.75 — Mechanical Enforcement Without Exception (cascaded from constitution submodule §11.4.75)

Verbatim user mandate (2026-05-20): “Why do these violations still happen!? This is a serious problem! We cannot rely on stability nor consistency if we cannot respect our Constitution, mandatory rules and constraints! Is there a way to make this always respected, followed and applied without exception fully and unconditionally!? WE MUST HAVE THIS WORKING FLAWLESSLY!!! Do investigate the root causes of such problems! Once all problems are identified WE MUST apply proper mechanisms for this not to happen NEVER EVER AGAIN!”

The §11.4 covenant historically relied on agent + operator vigilance; three 2026-05-19 → 20 forensic incidents proved that late-binding enforcement fires hours-to-days after the violator commit reaches every remote. §11.4.75 closes the gap with FIVE independent mechanical enforcement layers — bypassing any single layer does not bypass the discipline: (1) local pre-commit git hook (refuses staged `.md` lacking sibling `.html` + `.pdf`); (2) `commit_all.sh` integration (`_constitution_sibling_check` + `auto-sync_all_markdown_exports.sh` self-repair); (3) local pre-push git hook (re-runs siblings + propagation-gate subset); (4) post-commit auto-repair hook (auto-generates orphan-`.md` siblings, idempotent + recursion-guarded); (5) local-only final-gate ritual (remote CI

DISABLED per User mandate — operator runs `pre_build_verification.sh + meta-test` before every tag per §11.4.40). Helper contracts: `scripts/install_git_hooks.sh`, `scripts/git_hooks/{pre-commit,pre-push,post-commit,commit-msg}`, `_constitution_sibling_check`. The `commit-msg` hook enforces a Bypass-rationale: `<reason>` footer when `--no-verify` is detected; `docs/audit/bypass_events.md` accumulates the audit trail. Five gates with paired §1.1 mutations: `CM-COVENANT-114-75-PROPAGATION`, `CM-GIT-HOOKS-INSTALL-SCRIPT`, `CM-GIT-HOOKS-SOURCE-DIR`, `CM-COMMIT-ALL-SIBLING-CHECK`, `CM-CI-WORKFLOW-PRESENT`. No escape hatch — no `--skip-hooks`, `--bypass-enforcement`, `--allow-orphan-md`, `--ci-not-applicable`, `--mechanical-enforcement-not-needed` flag.

Cascade requirement: This anchor (verbatim or by §11.4.75 reference) MUST appear in every owned submodule's `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Propagation gate `CM-COVENANT-114-75-PROPAGATION`; paired mutation strips the literal `→` gate FAILs. Severity-equivalent to a §11.4 PASS-bluff at the enforcement layer. **Canonical authority:** constitution submodule `Constitution.md` §11.4.75 for the full mandate.

§11.4.76 — Containers-Submodule Mandate (cascaded from constitution submodule §11.4.76)

Verbatim user mandate (2026-05-20): *“For any work or requirements of running services or codebase inside the Containers (Docker / Podman / Qemu / Emulators, and so on) we MUST USE / INCORPORATE the Containers Submodule properly: <https://github.com/vasic-digital/containers> (git@github.com:vasic-digital/containers.git). Containers Submodule contains all means for us to Containerize our code and services! If any feature or Containing System is missing or not supported we MUST EXTEND IT properly like we do all of our projects! No bluff work is allowed of any kind!”*

For ANY containerized workload (Docker / Podman / Qemu / Kubernetes / container-backed emulators), every consuming project MUST: (1) install `vasic-digital/containers` (`digital.vasic.containers`) as a Git submodule; (2) consume via `replace` directive during development + pinned commit SHAs in production; (3) boot infra on-demand via `pkg/boot` + `pkg/compose` + `pkg/health` so operators are never required to start `podman machine` / `docker compose up` manually — the boot is part of the test entry point (the on-demand-infra invariant); (4) extend the Submodule (PR upstream) for missing runtimes / lifecycle primitives — never reimplement in-project (per §11.4.74); (5) anti-bluff: integration tests claiming to exercise containerized components MUST actually boot them via the Submodule — short-circuit fakes that bypass boot are a §11.4

violation. Tracker rows touching containerization MUST record Catalogue-Check: extend vasic-digital/containers@<sha> (or reuse). Planned gate CM-CONTAINERS-USED scans container-touching PRs for digital.vasic.containers/... imports; paired mutation strips the import + asserts FAIL.

Cascade requirement: This anchor (verbatim or by §11.4.76 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-76-PROPAGATION; paired mutation strips the literal → gate FAILs. **Canonical authority:** constitution submodule Constitution.md §11.4.76 for the full mandate.

§11.4.77 — Regeneration-Mechanism-Required Mandate (cascaded from constitution submodule §11.4.77)

Verbatim user mandate (2026-05-20): *“We must be sure that after excluding anything from Git versioning we still have the mechanism which will out of the box obtain or re-generate missing content!”*

Every .gitignore entry excluding (a) >~100 MiB OR (b) any artefact essential to building / running / testing the project MUST carry a documented + automated mechanism to either re-obtain (download from authoritative source: vendor tarball, SDK installer, npm/pip/cargo/go-mod/container registry, dedicated git submodule, S3/GCS) OR re-generate (run from tracked source via build pipeline, code-gen, asset render, captured-evidence replay, container build). Required artefacts per qualifying entry: (1) .gitignore-meta/<entry-slug>.yaml declaring pattern + mechanism-type + script-path + expected-disk-usage + vendor-url-or-source + integrity hash + requires-network + requires-credentials; (2) a non-interactive entry in scripts/setup.sh post-clone bootstrap; (3) a pre-build gate verifying regenerated content present OR a recent .gitignore-meta/.regenerated/<slug>.ok stamp; (4) README + docs/guides/*.md describing the mechanism + manual fallback + time/disk budget + §11.4.10 credentials. Bare .gitignore additions without the mechanism are a §11.4 PASS-bluff variant — codebase appears complete but a fresh clone cannot build/run. No escape hatch — no --skip-regen-mechanism, --gitignore-is-enough, --operator-already-has-content flag. Planned gate CM-GITIGNORE-REGEN-MECHANISM + paired §1.1 mutation (strip a required YAML key → gate FAILs).

Cascade requirement: This anchor (verbatim or by §11.4.77 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-77-PROPAGATION; paired mutation strips the literal → gate FAILs. Severity-equivalent to a §11.4 PASS-bluff

at the repository-hygiene layer. **Canonical authority:** constitution submodule Constitution.md §11.4.77 for the full mandate.

§11.4.78 — CodeGraph Code-Intelligence Mandate (cascaded from constitution submodule §11.4.78)

Verbatim user mandate (2026-05-20): *“Make codegraph MANDATORY CHOICE for this purpose for all of our project ... All project which do not have configured and installed codegraph yet MUST DO IT and MUST USE IT!”*

Every consuming project worked on by AI coding agents MUST install, initialize, and use **CodeGraph** (<https://github.com/colbymchenry/codegraph>, npm @colbymchenry/codegraph) — a local SQLite semantic code-knowledge-graph exposed to agents over MCP (100% local, no cloud). (1) Install globally via npm with a user-writable npm prefix (no sudo). (2) codegraph init + codegraph index: .codegraph/config.json is tracked, .codegraph/codegraph.db is gitignored with codegraph index as its §11.4.77 regeneration mechanism; the config.json exclude list MUST exclude every credential/secret path per §11.4.10. (3) Wire codegraph serve --mcp into every CLI agent (Claude Code .mcp.json, OpenCode opencode.json, Qwen Code .qwen/settings.json, Crush .crush.json, host-local otherwise) referencing the bare codegraph command on PATH (no hardcoded host path). (4) Cover the integration with an anti-bluff suite whose per-agent end-to-end layer uses an unforgeable challenge (a fact obtainable only by calling a CodeGraph MCP tool, e.g. index node count via codegraph_status); a genuinely un-drivable agent is a documented SKIP per §11.4.3, never a faked PASS. (5) Document in docs/CODEGRAPH.md, kept in sync per §11.4.12 / §11.4.65. CodeGraph is consumed as the published npm package (§11.4.74) — not a git submodule, adds no Git remote. Planned gate CM-CODEGRAPH-WIRED + paired §1.1 mutation (strip a secret-exclusion → gate FAILs).

Cascade requirement: This anchor (verbatim or by §11.4.78 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-78-PROPAGATION; paired mutation strips the literal → gate FAILs. **Canonical authority:** constitution submodule Constitution.md §11.4.78 for the full mandate.

§11.4.79 — Own-Org Submodules MUST Be Included in the CodeGraph Index (cascaded from constitution submodule §11.4.79)

Verbatim user mandate (2026-05-21): “All Submodules we use in the project and that are part of organizations to which we have the full access via GitHub, GitLab and other CLIs MUST BE included into the codegraph database and initialized / scanned / synced!”

Refines §11.4.78’s exclude-list with a per-submodule-ownership split: (a) own-org submodules (full write access via the project’s CLIs — canonical orgs vasic-digital + HelixDevelopment) MUST be INCLUDED in the index; (b) third-party submodules (the §11.4.74 no-match → vendor path) MUST be EXCLUDED. Operational steps: (1) `git submodule update --remote --merge` to pull latest before re-indexing, respecting load-bearing pins on third-party submodules; (2) adjust `.codegraph/config.json` exclude list to keep own-org paths in scope; (3) re-index via `scripts/codegraph_setup.sh`; (4) verify via `scripts/codegraph_validate.sh` with ≥1 probe resolving a symbol living ONLY inside an own-org submodule; (5) paired §1.1 mutation — temporarily add the own-org submodule to exclude → validate MUST FAIL on the cross-submodule probe → restore. An index that lies about reachable symbols is a PASS-bluff against AI agents. Own-org submodules silently excluded without an audit trail in `.codegraph/config.json` comments is a release blocker.

Cascade requirement: This anchor (verbatim or by §11.4.79 reference) MUST appear in every owned submodule’s `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Propagation gate CM-COVENANT-114-79-PROPAGATION; paired mutation strips the literal → gate FAILs. **Canonical authority:** constitution submodule `Constitution.md` §11.4.79 for the full mandate.

§11.4.80 — CodeGraph Regular-Update + Sync Automation Mandate (cascaded from constitution submodule §11.4.80)

Verbatim user mandate (2026-05-21): “We MUST regularly check for the updates and execute codegraph npm updates so the latest version of it is always installed on the host machine! ... Make sure we have proper full automation bash scripts which will run regularly and that these are part of the constitution Submodule ... Make sure all updates, sync processes we do and important codegraph related events are all documented under

docs/codegraph in Status and Status_Summary documents ... and regularly export them like all other Status docs into the PDF and HTML!"

Three deliverables (all living in the constitution submodule, inherited by reference per §3 — consuming projects invoke at `${CONST_DIR}/scripts/codegraph_*.sh`, never copy): (1) `scripts/codegraph_update.sh` — `npm`-installs latest `@colbymchenry/codegraph` after a registry version check; appends old/new version to `docs/codegraph/Status.md`; `anti-bluff` verifies `codegraph --version` reflects the new version after install (`npm exit 0` $\neq$ working binary). (2) `scripts/codegraph_sync.sh` — after a successful update runs `codegraph status` $\rightarrow$ `codegraph sync` $\rightarrow$ `codegraph status` $\rightarrow$ the project's `scripts/codegraph_validate.sh`; appends every step's output to BOTH the project's and the constitution's `docs/codegraph/Status.md`. (3) `docs/codegraph/Status.md` + `Status_Summary.md` append-only ledgers, exported to `.html` + `.pdf` per §11.4.65. Cadence: weekly floor (per §11.4.45). A consuming project that has not run `codegraph_update.sh` in >2 weeks AND has open AI-agent work is a release blocker. Paired §1.1 mutation: downgrade installed version $\rightarrow$ script detects drift $\rightarrow$ restore.

Cascade requirement: This anchor (verbatim or by §11.4.80 reference) MUST appear in every owned submodule's `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Propagation gate CM-COVENANT-114-80-PROPAGATION; paired mutation strips the literal `$\rightarrow$  gate FAILs`. **Canonical authority:** constitution submodule `Constitution.md` §11.4.80 for the full mandate.

§11.4.81 — Cross-Platform-Parity Mandate (cascaded from constitution submodule §11.4.81)

Verbatim user mandate (2026-05-21): “Any Linux-only blocker / issue we have MUST BE created macOS and other supported platforms equivalent! So, depending on platform proper implementation will be used for particular OS! EVERYTHING MUST BE PROPERLY EXTENDED AND UPDATED!”

Every consuming project whose supported-platforms manifest lists more than one OS MUST, for every feature/test/gate/challenge/mutation depending on platform-specific primitives, ship a per-OS-equivalent implementation chosen at runtime via `uname -s` (or equivalent detection). Three sub-mandates: **(A) Per-OS implementation REQUIRED** — Linux `cgroup/systemd//proc` primitives MUST have documented per-OS equivalents (POSIX `setrlimit/ulimit`, macOS `launchd`, BSD `rctl`, Windows Job Object) chosen via runtime dispatch. **(B) Per-OS tests REQUIRED** — every platform-dependent gate test MUST have case `"$(uname -s)"` in branches with positive captured evidence per §11.4.2 + §11.4.5

in each branch; SKIP-with-reason acceptable ONLY when the platform genuinely cannot enforce the invariant. **(C) Honest kernel-gap citation + adjacent equivalent test REQUIRED** — where a Linux primitive has NO equivalent due to a documented kernel limitation (canonical: XNU does not enforce RLIMIT_AS for unprivileged processes), the test MUST detect the gap at runtime, SKIP with exact kernel reason + reproducer + honest-gap-doc link, AND provide an ADJACENT test exercising the closest invariant the platform CAN enforce (e.g. RLIMIT_CPU+SIGXCPU as the macOS proxy), itself anti-bluff with a paired §1.1 mutation. Gate CM-CROSS-PLATFORM-PARITY scans for case "\$ (uname -s) " blocks asserting a non-SKIP branch (or honest-gap citation) per platform in the manifest; paired mutation strips a Darwin branch → gate FAILs. No escape hatch.

Cascade requirement: This anchor (verbatim or by §11.4.81 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-81-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker on multi-platform projects. **Canonical authority:** constitution submodule Constitution.md §11.4.81 for the full mandate.

§11.4.82 — Iteration-Speedup Discipline Mandate (cascaded from constitution submodule §11.4.82)

Verbatim user mandate (2026-05-22): “How can we speed-up this whole development and fixing process? ... Do not forget to all speed optimizations critical rules and mandatory constraints MUST BE all added into our root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md and QWEN.md and all other relevant constitution Submodules files!”

Iteration cycle time is a first-order quality enabler. Every consuming project's build / test / commit / debug pipeline MUST adopt these speedup disciplines AS MANDATORY (each independently enforceable): (A) Phase-1 forensic (superpowers:systematic-debugging) before any speculative source patch — speculative patches without FACT-grade root cause are §11.4.6 + §11.4.82 violations; (B) Live-ADB-First (or live-equivalent) before any rebuild — strengthens §11.4.51 to a release-blocker mandate; (C) 30-second pre-flight before launching rebuild orchestrators (device/sink reachability, host memory/disk, no stale locks, no orphan processes); (D) persistent build caches outside containers (ccache/sccache/Gradle daemon bind-mounted to host); (E) module-only rebuild for loadable-module-only changes; (F) parallel multi-device testing with separate qa-results/<TS>/<device-tag>/ outputs; (G) subagent scope discipline + worktree isolation (≤30 min budget, single-responsibility, isolation:

"worktree" default); (H) lock-file + stale-process hygiene (clean .git/index.lock, disable auto git-gc in concurrent repos); (I) cycle telemetry per §11.4.24 (commit hash, per-phase wall-clock, speedup-flag set, outcome — aggregated weekly). Gate CM-ITERATION-SPEEDUP-DISCIPLINE audits recent cycles for telemetry citing which of (A)-(I) applied; paired §1.1 mutation strips the speedup-flag column → gate FAILs. No escape hatch — no --skip-phase1-forensic, --no-pre-flight, --rebuild-everything-always, --unlimited-subagent-scope, --ignore-locks, --no-telemetry flag.

Cascade requirement: This anchor (verbatim or by §11.4.82 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-82-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.82 for the full mandate.

§11.4.83 — docs/qa/ End-User Evidence Mandate (cascaded from constitution submodule §11.4.83)

Verbatim user mandate (2026-05-22): “every feature that ships MUST carry a recorded e2e communication transcript + any attached materials under docs/qa/<run-id>/ (per-feature subdirectories). A feature with no QA transcript is itself a §107 PASS-bluff — it claims to work but has no auditable runtime evidence. Bot-driven automation MUST preserve full bidirectional communication threads as proof.”

Every feature that ships MUST carry a recorded end-to-end communication transcript plus any attached materials (screenshots, request/response payloads, audio, file uploads) committed under docs/qa/<run-id>/ — one directory per feature run. Operative rule: (1) every consuming project MUST maintain a docs/qa/ tree, each new feature under docs/qa/<run-id>/ where <run-id> is monotonic + greppable (timestamp / ATM-NNN / other workable-item ID per §11.4.54); (2) transcripts MUST be full bidirectional — every prompt/command sent + every response received (one-sided is not a transcript); (3) attached materials MUST be committed in-repo (no external-only links — that is a §11.4.13 sink-side violation); (4) bot-driven / agent-driven QA automation MUST preserve the full conversation thread as the proof artefact; (5) release gates MUST refuse to tag a version that has any feature-shipping commit without its matching docs/qa/<run-id>/ directory. A feature with no QA transcript is a §11.4 / §107 PASS-bluff. Composes with §11.4.2 / §11.4.5 / §11.4.13 / §11.4.65 / §11.4.69 / §1.1.

Cascade requirement: This anchor (verbatim or by §11.4.83 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-83-PROPAGATION; paired

mutation strips the literal → gate FAILs. Release blocker — no --qa-evidence-optional escape hatch. **Canonical authority:** constitution submodule Constitution.md §11.4.83 for the full mandate.

§11.4.84 — Working-Tree Quiescence Rule for Subagent Commits (cascaded from constitution submodule §11.4.84)

Verbatim user mandate (2026-05-22): *“no subagent commit may proceed while any concurrent mutation gate is in flight in the same checkout. Before git add, the committing agent MUST grep its own working tree for mutation markers (MUTATED for paired, // always pass, return json.Marshal shortcut paths, etc.). Any unexplained file in the staging area triggers ABORT.”*

No subagent (or main-thread) commit may proceed while any concurrent mutation gate, paired-mutation experiment, or other in-flight mutation is live in the same checkout. Before git add, the committing agent MUST grep its own working tree for mutation markers (MUTATED for paired, // always pass, return json.Marshal shortcut paths, // MUTATION / # MUTATION annotations, _mutated_* filename suffixes, etc.) and explicitly account for every modified file in the staging area; any unexplained file → ABORT. (Forensic case: a logo-fix subagent’s git add swept an // always pass JWT-verify mutation residue into an unrelated commit pushed to all four mirrors — a real security-defect window.) Operative rule: (1) pre-git add greps for mutation markers + cross-checks git status --porcelain against the subagent’s declared scope; unaccounted entries → ABORT; (2) any active mutation gate MUST be serialised (mutate → assert FAIL → restore → assert PASS) and the working tree verifiably clean before any unrelated commit; (3) concurrent subagents in the SAME checkout MUST coordinate through a lockfile (.git/MUTATION_IN_PROGRESS) — cleaner solution is git worktree add per subagent (composes with §11.4.20/§11.4.70); (4) post-commit mutation-residue-scanner MUST run before push — any commit containing a mutation marker → push BLOCKED.

Cascade requirement: This anchor (verbatim or by §11.4.84 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-84-PROPAGATION; paired mutation strips the literal → gate FAILs. A mutation marker that lands in a tagged commit is a critical defect regardless of how briefly it persisted. **Canonical authority:** constitution submodule Constitution.md §11.4.84 for the full mandate.

§11.4.85 — Stress + Chaos Test Mandate (cascaded from constitution submodule §11.4.85)

Verbatim user mandate (2026-05-24): *“Every fix or improvement you do MUST BE covered with full automation stress and chaos tests so we are sure nothing can break the functionality and all edge cases are monitored and polished and additionally fixed if that is needed! Everything must produce rock solid proofs and follow fully no-bluff policy!”*

Every fix or improvement landed MUST ship with full-automation **stress** AND **chaos** test suites exercising edge cases, sustained load, concurrent contention, and failure-injection. Happy-path coverage alone is a §11.4 / §107 PASS-bluff at the resilience layer. **Stress** (closed-set): sustained load ($N \geq 100$ iterations OR ≥ 30 s wall-clock, p50/p95/p99 latency recorded) + concurrent contention ($N \geq 10$ parallel invocations, no deadlock/leak) + boundary conditions (empty/max/off-by-one, each categorised). **Chaos** (closed-set, per fix-class appropriateness): process-death injection + network-fault injection (drop/delay/reorder) + input-corruption injection + resource-exhaustion injection (disk full, OOM, FD exhaustion — refuse cleanly OR degrade, NEVER crash) + state-corruption injection (mid-flight lock loss, partial-write). Every stress + chaos PASS MUST cite a captured-evidence artefact path per §11.4.5 + §11.4.69. Helper library `stress_chaos.sh` provides `ab_stress_run`, `ab_stress_concurrent`, `ab_chaos_kill_pid_during`, `ab_chaos_drop_network_during`, `ab_chaos_corrupt_file_during`, `ab_chaos_oom_pressure_during`, `ab_chaos_disk_full_during`, each composing with `ab_pass_with_evidence` / `ab_skip_with_reason`. Cleanup non-negotiable in trap `'...'` EXIT (cleanup failure = §11.4.14 violation). Four-layer coverage per §11.4.4(b) + paired §1.1 mutation (strip chaos-injection or evidence-capture → gate FAILs). No escape hatch — no `--skip-stress`, `--no-chaos`, `--happy-path-suffices`, `--stress-test-later` flag.

Cascade requirement: This anchor (verbatim or by §11.4.85 reference) MUST appear in every owned submodule’s `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Propagation gate CM-COVENANT-114-85-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule `Constitution.md` §11.4.85 for the full mandate.

§11.4.86 — Roster/Corpus-Backed Status-Doc Auto-Sync Mandate (cascaded from constitution submodule §11.4.86)

Verbatim user mandate (2026-05-25): *“Make sure that assets and players Status docs are ALWAYS regularly updated and in sync like all others*

Status docs — any time we add or modify the assets content(s) or we change or add new / remove existing pre-installed video and audio player apps! This MUST WORK OUT OF THE BOX!”

Some Status docs (§11.4.45) are backed by a tracked roster (installed apps/components) or a tracked asset corpus (test/media asset directory) rather than narrative alone. Their freshness MUST NOT depend on operator vigilance — the moment a roster/corpus member changes (app added/removed/renamed; asset added/modified/removed) the Status doc + Status_Summary + HTML + PDF MUST resync out of the box, mechanically. Mechanism (all must hold): (1) drift-proof fingerprint — sha256 of the sorted member list (NOT mtime), persisted in a sidecar beside the Status doc; (2) a sync helper that regenerates the fingerprint + re-exports HTML+PDF via the §11.4.65 exporter, wired so sync is automatic; (3) a pre-build gate that FAILs when the live fingerprint differs from the persisted one (mirrors §11.4.12 CM-ISSUES-SUMMARY-SYNC + §11.4.45 sync_integration_status); (4) a paired §1.1 mutation corrupting the fingerprint and asserting the gate FAILs. Classification: universal — the consuming project supplies the specific docs, roster/corpus sources, helper, and gate name per §11.4.35.

Cascade requirement: This anchor (verbatim or by §11.4.86 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-86-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker — no --skip-roster-sync, --allow-status-drift, --roster-sync-not-applicable flag.

Canonical authority: constitution submodule Constitution.md §11.4.86 for the full mandate.

§11.4.87 — Endless-Loop Autonomous Work + Zero-Idle Agent Dispatch + Anti-Bluff Testing Mandate (cascaded from constitution submodule §11.4.87)

Verbatim user mandate (2026-05-26): *“continue in endless loop fully autonomously”* (and any semantically-equivalent phrasing).

When the operator instructs an AI agent to continue in an endless autonomous loop, the agent MUST treat it as a HARD-CONTRACT covenant: (A) continue working until docs/Issues.md Status-column has zero non-terminal entries AND docs/CONTINUATION.md §3 Active work is empty AND no background subagent is mid-execution AND no external dependency is in-flight; (B) dispatch background subagents for parallelisable work — main + every subagent operate concurrently, “waiting for results” is the ONLY acceptable idle reason; (C) every closure lands four-layer test coverage per §11.4.4(b) with captured-evidence (audio/video/

network/UI/sysfs physical proofs); (D) the §11.4 anti-bluff covenant family (§11.4.1 / §11.4.2 / §11.4.6 / §11.4.7 / §11.4.27 / §11.4.50 / §11.4.52 / §11.4.68 / §11.4.69 / §11.4.83) is the operative truth-discipline — tests AND HelixQA Challenges bound equally; (E) the loop terminates ONLY on all-conditions-met, explicit operator STOP, host-session-safety demand, or scheduled wake on a known-future-actionable signal. No escape hatch — no --idle-OK, --skip-endless-loop, --bluff-permitted-for-this-task, --metadata-only-test-suffices, --no-physical-proof-required flag.

Cascade requirement: This anchor (verbatim or by §11.4.87 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-87-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.87 for the full mandate.

§11.4.88 — Background-Push Mandate: Commit-Lock Release Immediately After Commit, Push Runs Detached (cascaded from constitution submodule §11.4.88)

Forensic anchor (2026-05-26): a single `commit_all.sh` held its flock ~5 hours because `do_push` ran synchronously after the commit landed — every subsequent commit blocked on a slow mirror push irrelevant to the local commit's durability. Implementation seam for §11.4.87(B) zero-idle. The mandate: (A) `.git/.commit_all.lock` MUST be released IMMEDIATELY after `git commit` returns 0 — the commit is durable on local disk regardless of remote push outcome; (B) push runs detached via `nohup ./push_all.sh ... > <log> 2>&1 & + disown` — the orchestrator's exit code reports COMMIT success, NOT push success; (C) `push_all.sh` acquires per-remote flock `.git/.push.<remote>.lock` so concurrent invocations targeting the same remote serialize but different-remote invocations run in parallel; (D) backgrounded push failures land in `qa-results/push_failures/` `<ts>_<remote>.log` — the next autonomous-loop tick checks per §11.4.87(A) “no external dependency in-flight” gate; (E) synchronous-push escape: explicit --sync-push CLI flag preserves legacy behaviour for §11.4.41 force-push merge-first audit paths. Gates CM-COVENANT-114-88-PROPAGATION + CM-BACKGROUND-PUSH-WIRED + paired §1.1 mutations. Synchronous push (without --sync-push) = §11.4 PASS-bluff at the execution layer.

Cascade requirement: This anchor (verbatim or by §11.4.88 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-88-PROPAGATION; paired

mutation strips the literal → gate FAILs. Release blocker — no escape hatch beyond --sync-push for force-push events. **Canonical authority:** constitution submodule Constitution.md §11.4.88 for the full mandate.

§11.4.89 — Background Test Execution Mandate (cascaded from constitution submodule §11.4.89)

Verbatim user mandate (2026-05-27): *“Any tests we are executing, especially long test cycles, MUST BE performed in background in parallel with main work stream! This MUST NOT block our capabilities to work on queued workable items. Main work stream can be blocked or sit iddle only if absolutely needed and if it depends hard on results of some background execution.”*

Symmetric anchor to §11.4.88 (background push) at the test-execution layer. Mandate: (A) long-running tests (>30 s expected: pre_build, meta_test, test_all_fixes, recent_work_validate, HelixQA banks, 4-phase cycles, full-suite retests, audio supervisors, dual-display recorders) MUST run via nohup ... > <log> 2>&1 & + disown with the log under a known dir (qa-results/<test_id>_<ts>.log); (B) the main stream proceeds to the §11.4.42 priority queue immediately; (C) hard-dependency gating — poll an exit-status file or pgrep -af <test> before steps that need the exit code, surfacing as §11.4.66 interactive options if the test is still running; (D) failures land in <log> files, the next loop tick checks; (E) foreground execution permitted ONLY for <30 s tests OR explicit operator authorisation; (F) per-script flock serialises same-script invocations, different-script invocations parallel. Gates CM-COVENANT-114-89-PROPAGATION + CM-BACKGROUND-TEST-EXECUTION-WIRED + paired §1.1 mutations.

Cascade requirement: This anchor (verbatim or by §11.4.89 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-89-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker — no escape hatch beyond explicit per-invocation operator authorisation. **Canonical authority:** constitution submodule Constitution.md §11.4.89 for the full mandate.

§11.4.90 — Obsolete Status + Per-Item Obsolescence Audit (cascaded from constitution submodule §11.4.90)

Verbatim user mandate (2026-05-27): *“Bug No 6 ... seems obsolete after latest request for new behavior ... mark obsolete tickets with some light*

gray background ... text - the description to be strikethrough styled ... review all existing open or resolved workable items if they are obsolete - not valid any more ... There MUST NOT be any mistake! No bluff is allowed of any kind!"

The §11.4.15 Status closed-set is extended with a terminal Obsolete (→ Fixed.md) value (orthogonal to Type per §11.4.16). Obsolescence reasons (closed vocabulary): superseded-by-design-change | superseded-by-later-mandate | feature-removed | duplicate-of | unsupported-topology. Every Obsolete heading MUST carry an ****Obsolete-Details:**** line (Since + Reason + Superseding-item + Triple-check evidence) within 8 non-blank lines. The §11.4.23 colorizer adds a cell-status-obsolete class — light-gray #E0E0E0 background + strikethrough description. Audit cadence: every release-gate sweep per §11.4.40 + §11.4.42; triple-check is non-negotiable per the operator mandate. Composes with §11.4.15 / §11.4.16 / §11.4.19 / §11.4.21 / §11.4.23 / §11.4.33 / §11.4.34 / §11.4.40 / §11.4.42 / §11.4.66 / §11.4.71. Gates CM-COVENANT-114-90-PROPAGATION + CM-ITEM-OBSOLETE-DETAILS + CM-OBSOLETE-COLORIZER-WIRED + paired §1.1 mutations.

Cascade requirement: This anchor (verbatim or by §11.4.90 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-90-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.90 for the full mandate.

§11.4.91 — Summary-Doc Clarity Mandate (cascaded from constitution submodule §11.4.91)

Verbatim user mandate (2026-05-27): "Summary docs - Issues_Summary some not clear one line descriptions - like 'Composes with' ... For each workable item we MUST HAVE clearly understandable meaning ... every team member can clearly understand what that particular workable item is exactly about! There cannot be misunderstanding or unclarity of any kind and no bluff allowed!"

Every summary entry (Issues_Summary, Fixed_Summary, README doc-link, Status_Summary pages 1+2, all one-liners) MUST contain a self-contained meaningful description ≥ 6 words OR ≥ 40 chars naming SUBJECT + PROBLEM/GOAL. Forbidden one-liner anti-patterns: section labels (Composes with, Closure criteria, Fix direction, etc.); bare metadata fragments (Critical, Bug, In progress, etc.); section-marker echoes; a \$-letter alone. Generators (in HelixCode: the §11.4.93 workable-items export DB exporter — which renders BOTH tracker summaries from the SQLite single-source-of-truth docs/

workable_items.db, and which SUPERSEDES the markdown-derived generate_issues_summary.sh / generate_fixed_summary.sh, both of which now REFUSE to run; plus update_readme_doc_links.sh / generate_status_summary.sh) MUST extract from the H1/H2 heading line per the §11.4.54 ATM-NNN convention, NEVER from arbitrary downstream text, and MUST refuse anti-pattern rows — emitting a (MISSING DESCRIPTION — fix source heading) placeholder with visual highlight. Gate CM-SUMMARY-CLARITY-DESCRIPTIONS scans every summary; an anti-pattern match = FAIL. Audit cadence: every §11.4.40 + §11.4.42 sweep.

Cascade requirement: This anchor (verbatim or by §11.4.91 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-91-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.91 for the full mandate.

§11.4.92 — Multi-Pass Change-Evaluation Discipline (cascaded from constitution submodule §11.4.92)

Verbatim user mandate (2026-05-27): *“Every change to the project or codebase we do MUST BE evaluated in several passes and in in-depth analysis for potential new issues or problems it can introduce! ... no bluff of any kind! After we do change or set of changes this mandatory steps MUST BE taken!”*

Every non-trivial change MUST pass a 5-pass evaluation BEFORE it is commit-ready: **(Pass 1)** main-task verification — change achieves the stated goal, captured-evidence per §11.4.5/§11.4.69; **(Pass 2)** regression-blast-radius analysis — enumerate every direct dependency, demonstrate no contract break; **(Pass 3)** cross-feature interaction analysis — audit parallel features sharing state/timing/hardware/shell environment; **(Pass 4)** deep-research validation per §11.4.8 — external precedent OR “NO external solution found — original work” + CodeGraph queries per §11.4.78/§11.4.79; **(Pass 5)** anti-bluff confirmation per §11.4 / §11.4.1 / §11.4.6 / §11.4.27 / §11.4.50 / §11.4.52 / §11.4.69 / §11.4.83 — no new bluff surface introduced. Each pass is documented (commit footers OR docs / entries OR qa-results / evidence). Only after all 5 passes complete may commit/push/test/release proceed. Trivial exemption: typo / revision-bump / MD-export-regen IF zero source touched AND the commit message cites the exemption explicitly. Gates CM-COVENANT-114-92-PROPAGATION + CM-MULTI-PASS-EVALUATION-EVIDENCE + paired §1.1 mutations.

Cascade requirement: This anchor (verbatim or by §11.4.92 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and

AGENTS.md. Propagation gate CM-COVENANT-114-92-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.92 for the full mandate.

§11.4.93 — SQLite-Backed Single-Source-of-Truth for Workable Items (cascaded from constitution submodule §11.4.93)

Verbatim user mandate (2026-05-27): *“There MUST be single source of truth for all of our workable items - SQLite database ... proper scripts (we recommend Go programs) ... reduce a chance for sync to be broken ... generate always all docs from DB or to re-generate Db from all docs we have in opposite direction”*

The text-based Issues/Fixed/Summary/CONTINUATION constellation is converted to a SQLite-DB-backed single source of truth. Schema mandatory tables: items (atm_id PK + Type + Status incl. Obsolete + Severity + title + description ≥40 chars + created/modified + composes_with JSON + current_location); item_history (append-only audit per §11.4.34 By/Reason/Evidence); obsolete_details (§11.4.90); operator_block_details (§11.4.21); firebase_metadata (§11.4.47); meta (schema version + last sync + integrity hash). A Go binary at cmd/workable-items/ provides sync md-to-db / db-to-md / diff / validate / add / close; bidirectional regen is byte-identical round-trip (closed-set whitespace/section-order tolerance). commit_all.sh refuses on non-empty diff; sync_issues_docs.sh invokes the Go binary; pre-build runs workable-items validate. Anti-bluff: unit + integration + stress (1000-row insert + 10 concurrent writers) + chaos (mid-write SIGKILL + corrupt-DB recovery + disk-full) + paired §1.1 mutation + HelixQA Challenge CME-WORKABLE-ITEMS-001. The Go binary lives in the constitution submodule (constitution/scripts/workable-items/) per §11.4.74. Gates CM-COVENANT-114-93-PROPAGATION + CM-WORKABLE-ITEMS-DB-PRESENT + CM-WORKABLE-ITEMS-MD-DB-IN-SYNC + paired §1.1 mutations. (NOTE: the DB tracking rule is AMENDED by §11.4.95 — DB is TRACKED, not gitignored.)

Cascade requirement: This anchor (verbatim or by §11.4.93 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-93-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker — text-based-only trackers are a §11.4 PASS-bluff at the data-architecture layer. **Canonical authority:** constitution submodule Constitution.md §11.4.93 for the full mandate.

§11.4.94 — Zero-Idle Priority-First Parallel-By-Default Operating Mode (cascaded from constitution submodule §11.4.94)

Verbatim user mandate (2026-05-27): *“We MUST NEVER sit iddle / wait or sleep if there is possibility for us to work on something ... Always check if there is a possibility to work on something while we are not working actively on something! Pick always by priority - most critical workable items and other tasks MUST BE done first! ... Stay still / iddle if nothing is left to be done at all or waiting for something that is blocking us / you!!!”*

§11.4.94 binds §11.4.20 + §11.4.42 + §11.4.58 + §11.4.70 + §11.4.72 + §11.4.82 + §11.4.87 + §11.4.88 + §11.4.89 into a single always-on enforcement: (A) idle ONLY when every queued item is genuinely blocked on an external dependency (hardware / network upstream / build/test completion the conductor cannot accelerate) OR operator STOP OR §12 host-safety — “don’t see what to do” is NEVER valid; (B) before ANY wake/sleep the conductor MUST survey parallel-work feasibility per §11.4.42 + §11.4.72 + §11.4.87, identify non-contending items, and dispatch in parallel per §11.4.20/§11.4.70 (subagent) + §11.4.58 (PWU disjoint scope) + §11.4.89 (background long tests); (C) priority order MANDATORY — pick highest-severity + §11.4.72 audio-first the conductor can autonomously progress; (D) subagent-driven default for non-trivial; (E) background default for >30 s wall-clock work via nohup+disown; (F) stability-preserving (composes with §11.4.92 multi-pass + §11.4.84 quiescence + §12.6–§12.9 host safety); (G) progress updates surfaced at milestone boundaries. Gates CM-COVENANT-114-94-PROPAGATION + CM-PARALLEL-WORK-AUDIT + paired §1.1 mutations.

Cascade requirement: This anchor (verbatim or by §11.4.94 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-94-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.94 for the full mandate.

§11.4.95 — Workable-Items SQLite DB Is TRACKED in Git, NEVER Gitignored (cascaded from constitution submodule §11.4.95)

Verbatim user mandate (2026-05-27): *“We shall not Git ignore our workable items SQLite DB since it is our single source of truth ... workable items SQLite DB regularly committed and pushed to all upstreams!”*

§11.4.93's earlier "gitignored per §11.4.30" clause is AMENDED — the DB at docs/workable_items.db is TRACKED in git, NEVER gitignored. It IS authoritative source data, NOT a build artefact. Every workable-items sync md-to-db that mutates state MUST stage + commit + push the DB alongside the MD regen per §11.4.19 atomic-move + §2.1 multi-upstream push. A WAL-checkpoint (PRAGMA wal_checkpoint(TRUNCATE)) is required before commit-stage so the transient .db-wal + .db-shm sidecars (gitignored per §11.4.30) are safely discardable. The §11.4.77 regeneration mechanism does NOT apply — the DB IS the source. Destructive DB ops require §9.2 hardlinked-backup + operator authorization; §11.4.41 force-push merge-first applies if DB history ever needs rewrite. Gates CM-COVENANT-114-95-PROPAGATION + CM-WORKABLE-ITEMS-DB-TRACKED + paired §1.1 mutation.

Cascade requirement: This anchor (verbatim or by §11.4.95 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-95-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.95 for the full mandate.

§11.4.96 — Safe-Parallel-Work-With-Long-Build Catalogue + Mandate (cascaded from constitution submodule §11.4.96)

Verbatim user mandate (2026-05-27): "Are there except AOSP build process any other active jobs being done at the moment? Can we work on something in parallel while build is in progress so we slowly cleanup our slate? ... do as much as possible work in background in parallel with main work stream and oreferably using subagents-driven approach!"

An operational catalogue for the canonical long-running workload (multi-hour containerised build per §12.9). **SAFE during build:** (A) MD/docs work; (B) generator/helper script work under scripts/; (C) pre-build + meta-test gate authoring + paired §1.1 mutations; (D) on-device test scripts; (E) constitution submodule edits + push; (F) any submodule commit + push per §11.4.88; (G) read-only live-ADB probes (dumpsys/getprop/cat /proc/.../screencap/logcat); (H) subagent dispatch per §11.4.20/§11.4.70 + §11.4.84 quiescence; (I) web research + external API queries with §11.4.10 credentials; (J) workable-items DB ops per §11.4.93+§11.4.95; (K) backgrounded pre-build + meta-test execution per §11.4.89. **UNSAFE during build:** (α) git checkout/reset --hard/clean -df on the source tree (use git worktree); (β) mass file deletes/renames under built source trees; (γ) submodule pointer updates affecting built artefacts; (δ) out/ mutations; (ε) make clean/m clobber/rm -rf out/; (ζ) container destruction; (η) disk-filling breaching §12.9 free-space minimum; (θ) §12 host-session-safety breaches.

Conductor responsibility: before EVERY pause point during a long build, consult the catalogue, identify (A)-(K) queue items per §11.4.42+§11.4.72, and dispatch ≥ 1 per §11.4.20/§11.4.70 subagent default + §11.4.89 background. “Build running, nothing else to do” is NEVER true per §11.4.94+§11.4.96. Gates CM-COVENANT-114-96-PROPAGATION + CM-PARALLEL-WORK-DURING-BUILD-AUDIT + paired §1.1 mutations.

Cascade requirement: This anchor (verbatim or by §11.4.96 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-96-PROPAGATION; paired mutation strips the literal $\rightarrow$ gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.96 for the full mandate.

§11.4.97 — Maximum-Use-of-Idle-Time + Progress-Update Cadence (cascaded from constitution submodule §11.4.97)

Verbatim user mandate (2026-05-27): *“keep it working, we should do as much as possible, if not it all but as much as we can as long as there is iddle time! it MUST be used! ... keep us updated about all progress and all phisycal proofs and gathered data as you progress through all open workable items!”*

Operating-mode capstone strengthening §11.4.87 + §11.4.94 + §11.4.96: (A) every minute of conductor idle time during which work could autonomously progress AND is not genuinely blocked = a §11.4.97 violation; “as much as possible, if not it all but as much as we can” is operative — dispatch CONTINUOUSLY through the entire idle window, not just at scheduled wakes; (B) progress-update cadence — emit an operator-facing 1-line update at every commit landed / subagent return / constitutional anchor / captured evidence / milestone closure, no operator prompt required; (C) continuous physical-proof gathering per §11.4.5 + §11.4.6 + §11.4.69 — every autonomous closure cites captured-evidence (evidence path goes into the §11.4.93 `item_history.evidence_path` when the DB lands); (D) composes with §11.4.5/6/13/20/27/42/50/52/69/70/72/83/85/87/88/89/94/96; (E) the idle-only-when-blocked closed-set is unchanged from §11.4.94(A). Gates CM-COVENANT-114-97-PROPAGATION + CM-IDLE-TIME-AUDIT + paired §1.1 mutations.

Cascade requirement: This anchor (verbatim or by §11.4.97 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-97-PROPAGATION; paired mutation strips the literal $\rightarrow$ gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.97 for the full mandate.

§11.4.98 — Full-Automation Anti-Bluff Mandate (cascaded from constitution submodule §11.4.98)

Verbatim user mandate (2026-05-28): *“Make sure we have full automation testing of all scenarios with real bot, main group and users without any manual intervention or contribution of real user! Everything MUST BE fully automatic and autonomous! These tests MUST BE able to rerun endless times when needed! ... Make sure there is no false positives in testing! Every test and its results MUST obtain real proofs of everything working! No bluff is allowed!”*

Closes the manual-intervention gap (§11.4 / §11.4.2 / §11.4.5 / §11.4.50 / §11.4.85 / §11.4.87 / §11.4.89 / §11.4.94 did not explicitly forbid it). A live/integration/e2e/Challenge test that requires a human action during execution (typing a message, clicking UI, hand-triggering a webhook, attaching a file — anything beyond startup) is by definition a §11.4 PASS-bluff at the automation layer. (A) Every governed test — unit/integration/e2e/Challenge/stress/chaos/live — MUST be fully self-driving end-to-end, reporting PASS/FAIL/SKIP-with-reason without any further human action after startup. (B) Single permissible exception: one-time credential bootstrap performed OUTSIDE test execution (.env from vault, shell exports, OAuth at first install, MTProto session activation) — configuration, not test driving. (C) Live messenger/channel/agent tests: no “operator must type” prompts (drive programmatically via second account / webhook fixture / loopback); no hard-coded session UUIDs that collide with the active dev session (Herald 2026-05-28 claude --resume silent exit -1 lesson); no 60 s human-response windows (§11.4.50 determinism violation); re-runnability proof — PASS at -count=3 consecutive automated invocations with self-cleaning state; §11.4.98 obsolescence audit classifies every existing test COMPLIANT vs NON-COMPLIANT; no silent-skip-reported-as-PASS or stale-evidence-as-fresh. (D) With §11.4.85 + §11.4.89 + §11.4.87 + §11.4.94 forms a continuously-validated, non-flake, anti-bluff regime. (F) Manual-dependency tests not rewritten within 30 days graduate to §11.4.90 Obsolete citing §11.4.98.

Cascade requirement: This anchor (verbatim or by §11.4.98 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-98-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.98 for the full mandate.

§11.4.99 — Latest-Source Documentation Cross-Reference Mandate (cascaded from constitution submodule §11.4.99)

Verbatim user mandate (2026-05-28): *“Make sure we ALWAYS check against latest versions of services we use web / online docs before creating instructions! This situation is illustration of how we can misguide ourselves or get banned! ... These are mandatory rules / constraints and the result is consistency and safety of created instructions, guides and manuals!”*

Misguidance-by-stale-docs is the same severity class as a §11.4 PASS-bluff at the documentation layer (Herald 2026-05-28 case: a first-draft MTPROTO guide recommended VoIP fallback numbers and omitted the `recover@telegram.org` pre-login email — both contradicted Telegram’s official docs + the gotd/td maintainer guide and could have caused a permanent account ban). Closes the gap §11.4.92 Pass 4 alludes to but does not mandate. (A) Before committing any operator-facing instruction/guide/manual/troubleshooting/setup doc, the author MUST: (1) fetch the LATEST official online documentation of the documented service/library via WebFetch / MCP / direct browsing — NEVER training data, memory, or prior committed docs; (2) cross-reference every instruction step against that source; (3) seek secondary authoritative sources (maintainer SUPPORT.md, official changelogs, vetted community FAQs) when the official source is sparse/silent; (4) cite source URLs + date in a `## Sources verified` footer in the doc; (5) cite a `Sources verified <date>: <urls>` footer in the commit message. (B) Negative findings (gaps/silences/contradictions) MUST be documented explicitly. (C) Docs older than 6 months are STALE — re-verify before citing as operator authority, at every vN.0.0 release boundary, on service breaking-change announcements, or on operator error reports. (D) Risk-classified services (messaging, cloud APIs, payment systems, AI/LLM providers, code-hosting, package managers) carry a 90-day max staleness + explicit safety warnings. (E) Composes with but is INDEPENDENT of §11.4.92 Pass 4. (G) Commit missing either footer is BLOCKED at release-gate; stale-beyond-grace docs graduate to §11.4.90 Obsolete (Reason=stale-documentation).

Cascade requirement: This anchor (verbatim or by §11.4.99 reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-99-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.99 for the full mandate.

§11.4.101 — Autonomous-Decision-Over-Blocking Mandate (cascaded from constitution submodule §11.4.101)

Verbatim user mandate (2026-05-28): *“when working in endless working loop fully autonomously try to decide most properly about points which would block execution and wait for us. If we haven’t answered now work would be blocked whole night! If possible and if that will not cause any issues make proper and most reliable and safe decision so we achieve maximal efficiency and work gets fully done!”*

In autonomous / endless-loop mode (per §11.4.87), the agent **MUST** minimize operator-blocking and make the safe, reliable, reversible decision itself so work is not stalled (e.g. overnight) waiting for input — §11.4.87 says keep working, §11.4.101 says **HOW** to clear the decision points. **Proceed-autonomously (closed-set, ALL must hold)**: (a) the action is reversible OR has a captured pre-op backup per §9.2; (b) the safe choice is determinable from captured evidence per §11.4.6 (no guessing — LIKELY/probably/seems is NOT a determination); (c) a wrong choice’s blast radius is bounded AND recoverable; (d) it composes with anti-bluff §11.4, host-safety §12, data-safety §9. **Block-only-when (BLOCK via the §11.4.66 interactive mechanism ONLY when ALL hold)**: the action is irreversible AND high-blast-radius AND the safe choice cannot be determined from evidence — e.g. external-account state the agent cannot inspect, hardware it cannot access, destructive ops without backup, force-push (also §9.2 + §11.4.41), spending money or sending data to third parties. Operator-blocked per §11.4.21 is reached only after this rule fires AND the self-resolution-exhaustion audit completes. An unavoidable block parks one work unit — it does NOT pause the loop; the agent keeps progressing every non-blocked item in parallel per §11.4.87 + §11.4.94 (posing the question then going idle is a §11.4.94 + §11.4.97 violation).
Classification: universal (§11.4.17).

Cascade requirement: This anchor (verbatim or by §11.4.101 reference) **MUST** appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-101-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.101 for the full mandate.

§11.4.102 — Mandatory systematic-debugging activation + always-loaded skill-discovery + plugin-dependency availability (cascaded from constitution submodule §11.4.102)

Verbatim user mandate (2026-05-29): **“Make sure that we ALWAYS trigger / start the”/superpowers:systematic-debugging” skills when any issues happen! If this is possible to activate and use in this situations out of the box when we spot problems / issues / bugs / misalignments / inconsistencies we MUST activate the skill(s) and make strongest efforts in full in depth analysis / debugging and determine root causes of all problem or obtain relevant data and information we need! ... we MUST make sure that “/using-superpowers” skill is ALWAYS loaded, applied and used! All dependencies (plugins) that Claude Code or other market places are offering MUST BE installed if these are not already available for loading and use!”*

Three cooperating invariants — the difference between guess-and-retry and investigate-to-root-cause-first. **(A) Mandatory systematic-debugging activation.** On ANY spotted issue / bug / test failure / gate failure / regression / misalignment / inconsistency / unexpected behaviour, the agent MUST activate `superpowers:systematic-debugging` (or the platform-equivalent structured-debugging discipline) **BEFORE proposing, writing, or applying any fix** — the **Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST**. Full four-phase arc: root-cause → pattern → hypothesis → implementation. Guess-and-retry, symptom-patching, and re-running a failed test hoping it passes (“probably transient / flaky”) WITHOUT a completed investigation are §11.4.102 violations; calling a failure transient/flaky/intermittent/probably-timing without captured forensic evidence is simultaneously a §11.4.6 and §11.4.7 violation. **(B) Mandatory always-loaded using-superpowers.** `superpowers:using-superpowers` (or platform-equivalent skill-discovery discipline) MUST be loaded and applied at session start and consulted before any task; if ANY skill could apply — even at 1% relevance — it MUST be invoked rather than improvised from memory. **(C) Mandatory plugin / dependency availability.** Every skill plugin / marketplace package / capability dependency the project relies on MUST be installed + loadable BEFORE the dependent work proceeds; a missing plugin that blocks a mandated skill is a release-blocker until installed + confirmed loadable (install exit 0 ≠ skill loadable — confirm by observing the skill in the live capability list). Composes with §11.4.4 / §11.4.6 / §11.4.7 / §11.4.8 / §11.4.43 / §11.4.70 / §11.4.82(A) / §11.4.92. Classification: universal (§11.4.17). No escape hatch — no --skip-systematic-debugging, --guess-and-retry-OK, --

symptom-patch-permitted, --skip-skill-discovery, --plugin-optional, --missing-plugin-is-warning flag.

Cascade requirement: This anchor (verbatim or by §11.4.102 reference) MUST appear in every owned submodule's CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Propagation gate CM-COVENANT-114-102-PROPAGATION; paired mutation strips the literal → gate FAILs. Release blocker. **Canonical authority:** constitution submodule Constitution.md §11.4.102 for the full mandate.

§11.4.103 — Continuous parallel-stream working routine (User mandate, 2026-05-29)

Forensic anchor — verbatim user mandate (2026-05-29):

“Do this working approach continuously and make it part of regular working routine and add it to the Constitution Submodule documented fully.”

Promotes the proven multi-stream operating pattern into the project's **standing default working routine** (not a per-request opt-in). Binds §11.4.87/§11.4.88/§11.4.89/§11.4.94/§11.4.96 and adds the load-bearing invariant: **the main work stream MUST always stay FREE.**

(A) Main stream stays FREE. ALL commit AND push operations run detached (nohup ... & + disown, per §11.4.88 commit-lock-release-immediately + detached-push) — the main stream returns to the priority queue the moment the local commit is durable, never blocking on a push or a slow mirror. **(B) ≥3 parallel background streams at all times + auto-backfill** (User mandate 2026-05-29; raised from ≥2) — run **at least three** subagent-driven background streams (per §11.4.70/§11.4.20, isolated per §11.4.58 PWU worktree + §11.4.84 quiescence) alongside the main stream whenever three-plus non-contending actionable items exist; **the moment any one stream is FULLY done, a new stream MUST immediately start and take its place** (claim next-highest-priority non-contending item), so the active-stream count NEVER drops below 3 while actionable items remain. Idle below 3 is permitted ONLY when no remaining non-contending actionable items OR all remaining are externally blocked (§11.4.94/§11.4.97/§11.4.101). Standing band 3–6, bounded above by §12.6 60% memory + §11.4.58 6-agent cap. **(C) Most-critical + most-visible first; audio always top** per §11.4.72 + §11.4.42 priority order. **(D) Safe-during-build scope only** — while the 42 GB containerised AOSP rebuild (§12.9) or any heavy gradle/m -j build runs, streams restrict to the §11.4.96 SAFE catalogue (investigation/forensic/docs/test-authoring/gate-authoring/read-only probes/submodule edits/research/DB ops/backgrounded pre-build+meta-test); NEVER a second concurrent heavy build (§12.8). **(E) Heavy anti-bluff on every closure** — root cause proven or

UNCONFIRMED:/UNKNOWN:/PENDING_FORENSICS: per §11.4.6, captured evidence per §11.4.5+§11.4.69, deterministic consistency per §11.4.50, paired §1.1 mutations, §11.4.102 systematic-debugging on any spotted problem. **(F) Idle ONLY when genuinely externally blocked** (hardware/network upstream/in-flight build-test-push completion) OR operator STOP OR §12 host-safety, per §11.4.94(A) +§11.4.97; “nothing visible to do” with progressable items is NEVER valid; a block parks one work unit, not the loop (§11.4.101).

Composes with §11.4.58 / §11.4.70 / §11.4.72 / §11.4.87 / §11.4.88 / §11.4.89 / §11.4.94 / §11.4.96 / §11.4.97 / §11.4.101 / §11.4.102 / §11.4.42 / §11.4.84 / §12.6 / §12.7 / §12.8 / §12.9 / §9.2. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-103-PROPAGATION (literal 11.4.103 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --block-main-stream, --synchronous-commit, --synchronous-push, --single-stream-only, --skip-parallel-streams, --serialise-actionable-work, --idle-without-queue-survey flag.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.103.

Non-compliance is a release blocker regardless of context.

§11.4.104 — Participant identity, attribution & notification-tagging (User mandate, 2026-05-31)

Forensic anchor — verbatim user mandate (2026-05-31):

“Every supported messenger must relate messages to participants (Subscribers/Users); the same logical person may have a different username on every messenger. Workable items must carry who created them and who they are assigned to. Notifications must @-tag the right participant — but never the operator (who drives the system) and never the system agent.”

MANDATORY for every consumer that ships a messenger/notification surface (Herald + flavor binaries are the reference impl; others inherit per §11.4.35). Detailed spec: Herald docs/design/PARTICIPANT_ATTRIBUTION.md — restated, not redefined. **(A) Participant identity.** Every messenger MUST relate messages to a **Participant** (logical Subscriber/User); the SAME person MAY have a DIFFERENT username per messenger — modelled as a logical subscriber (canonical messenger-neutral handle, kind ∈ {human, agent, service}) + per-channel aliases (channel, channel_user_id, the @username used for tagging). Canonical handle closed set: Claude (reserved system-agent sentinel; never tagged) OR a subscriber @username. **(B) Operator = env var, not a DB flag** — the one human who drives via the agent CLI, designated by

HERALD_<CHANNEL>_OPERATOR_USERNAME

(e.g. HERALD_TGRAM_OPERATOR_USERNAME); a normal Participant whose handle equals that value. **(C) Workable items MUST carry created_by + assigned_to** (canonical handles): CLI-prompt → Operator; system/agent-detected → Claude; received-message → sender's resolved @username. assigned_to defaults to Operator, overridable. Legacy items carry "" and MUST still parse + validate. **(D) Tagging matrix:** tag assigned_to if it is a human ≠Operator; tag created_by if it is a human ≠Operator ≠Claude; NEVER tag Claude (system) or the Operator (no self-ping); de-dup; resolve each handle to the channel @username, skip if not on that channel. **(E) Anti-bluff (composes §11.4):** real SQLite round-trip with the new columns byte-identical (incl. legacy fixtures WITHOUT the fields); tagging matrix proven by a truth-table + a cell-flip mutation forcing FAIL; E2E real-event → real-message asserting the exact @usernames + a NEGATIVE case proving the Operator is NOT tagged; evidence under docs/qa/<run-id>/.

Composes with §11.4 + §11.4.1..§11.4.16 (anti-bluff covenant) / §11.4.5 / §11.4.69 / §11.4.50 / §11.4.91 / §11.4.93 / §11.4.95 / §1.1. Classification: universal (§11.4.17) — projects with no messenger surface inherit it latently (binds the moment they ship one, per the §11.4.96 pattern). Propagation gate CM-COVENANT-114-104-PROPAGATION (literal 11.4.104 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --skip-attribution, --no-participant-tagging, --tag-operator-anyway, --attribution-later flag.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.104; detailed spec Herald docs/design/PARTICIPANT_ATTRIBUTION.md.

Non-compliance is a release blocker.

§11.4.105 — Natural-language intent recognition & clarification (User mandate, 2026-05-31)

Forensic anchor — verbatim user mandate (2026-05-31):

“Users must NOT need to know command syntax (no COMMAND: ... prefix). They send a clear natural-language message; the System determines the intent. The System recognizes the commands it has; if none matches it infers the exact intent; if it is totally unable it replies, tags the user (@user ...), and asks to clarify precisely. We MUST always do our best to determine exact intent so we never annoy end users. This is a CORE part of the System.”

MANDATORY for every consumer that ships a messenger/command surface (Herald + flavor binaries are the reference impl; others inherit per §11.4.35). Detailed spec: Herald docs/design/INTENT_RECOGNITION.md — restated, not

redefined. **(A) No required command syntax.** Users MUST NOT be required to know any command syntax (no `COMMAND:` prefix); they send plain natural language and the System determines the intent. **(B) Three-tier resolution** (first that succeeds wins): TIER 1 — recognize the System's existing command set from natural language → action (confident deterministic match); TIER 2 — when no command matches, infer the exact intent (LLM dispatch maps language → action), NEVER guessing; TIER 3 — when neither a command nor a confident intent can be determined, REPLY to the message, TAG the sender (@username, resolved via the §11.4.104 IdentityResolver) and ask a PRECISE clarifying question NAMING the candidate intents — no guessing, no silent drop. **(C) Never guess, never drop:** a wrong action is worse than a clarifying question (composes §11.4.6 no-guessing); a message is NEVER silently dropped; only genuine ambiguity reaches Tier 3, which always replies-tags-and-asks. **(D) Anti-bluff (composes §11.4):** every tier ships unit + integration + E2E + full-automation tests with real captured evidence — Tier 1 truth-table (natural-language → action+fields, plus conservative negatives that MUST fall through to “no match”); Tier 3 E2E whose recorded reply body is EXACTLY @<sender> <specific question> + a NEGATIVE proving a clear command does NOT trigger clarify; a paired §1.1 mutation breaking the confidence guard (false-match) OR dropping the clarify tag MUST FAIL a test; evidence under docs/qa/<run-id>/.

Composes with §11.4 + §11.4.1..§11.4.16 (anti-bluff covenant) / §11.4.6 (no-guessing) / §11.4.104 (clarify reply tags the sender) / §11.4.5 / §11.4.69 / §11.4.98 / §1.1. Classification: universal (§11.4.17) — projects with no messenger surface inherit it latently (binds the moment they ship one, per the §11.4.96 pattern). Propagation gate CM-COVENANT-114-105-PROPAGATION (literal 11.4.105 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --require-command-syntax, --guess-intent-ok, --skip-clarify, --drop-on-ambiguous flag.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.105; detailed spec Herald docs/design/INTENT_RECOGNITION.md.

Non-compliance is a release blocker.

§11.4.106 — Docs Chain — mechanical documentation/DB sync engine (Operator mandate, 2026-05-31)

Forensic anchor — operator mandate (2026-05-31): Docs Chain is the canonical mechanical enforcer of the documentation-sync mandates; consumers MUST use the engine instead of ad-hoc per-project doc-sync scripts, register their chains via per-context YAML, and never accept a faked transform.

MANDATORY for every consumer. Docs Chain (the vasic-digital/docs_chain engine — a universal Go bidirectional document-and-database dependency-propagation engine) is the canonical mechanical enforcer of the documentation-sync mandates. Detailed spec: the engine docs ~/Projects/docs_chain → docs/CONSTITUTION_INTEGRATION.md (distribution + inheritance + anchor mapping table) + docs/USE_CASE_CATALOGUE.md (chain recipes) — restated, not redefined. **(A) Use the engine, never ad-hoc scripts** — consumed **by reference** (the flat-layout sibling ~/Projects/docs_chain/the constitution-exposed path), inherited like §11.4.80's codegraph_* scripts: referenced, NEVER copied; ad-hoc sync_*/generate_*_summary/update_readme_doc_links scripts are superseded and retired per registered context. **(B) Consumer-owned contexts** — the engine is project-agnostic; the consumer registers its chains as data via .docs_chain/contexts/*.yaml (§11.4.28 decoupling); state.json + *.docs_chain.tmp are gitignored. **(C) Anchors it mechanizes** (per the CONSTITUTION_INTEGRATION mapping table): §11.4.12 / §11.4.53 / §11.4.45 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.86 / §11.4.93 / §11.4.95 / §12.10 / §11.4.44 — with content-hash change detection (NOT mtime, §11.4.86), atomic-rename + SQLite-txn commit + rollback (§9.2), both-dirty sync → conflict-not-silent-merge (§11.4.6), verify as the deterministic CI/pre-build gate (§11.4.50), per-run captured evidence to qa-results/docs_chain/<run-id>/ (§11.4.69). **(D) NOT a replacement for authoring discipline** — the source author still writes the §11.4.44 revision header; the engine only keeps exports in sync. **(E) Anti-bluff (composes §11.4):** the engine NEVER fakes a transform — a missing pandoc/weasyprint surfaces a typed ToolAbsentError + honest §11.4.3 SKIP-with-reason, never a fake PASS / partial write; every sync/verify carries real captured evidence.

Composes with §11.4 + §11.4.1..§11.4.16 (anti-bluff covenant) / §11.4.6 (no-guessing — conflict not silent merge) / §11.4.28 (engine/context decoupling) / §11.4.80 (inherited-by-reference, never copied) / §9.2 / §11.4.50 / §11.4.69 / §11.4.5 / §1.1, plus the sync anchors it mechanizes (§11.4.12/.53/.45/.56/.57/.59/.60/.65/.86/.93/.95/.44, §12.10). Classification: universal (§11.4.17) — projects with no derived-export/DB-sync surface inherit it latently (binds the moment they ship one, §11.4.96 pattern). Status (§11.4.6): engine Phases 1–3 IMPLEMENTED+tested; CLI/YAML loader (Phase 4) + submodule distribution (Phase 6) PLANNED + OPERATOR-GATED — wire as phases land, claim no unshipped behaviour. Propagation gate CM-COVENANT-114-106-PROPAGATION (literal 11.4.106 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --ad-hoc-sync-ok, --skip-docs-chain, --fake-transform, --sync-evidence-optional flag.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.106; detailed spec the Docs Chain engine docs docs/CONSTITUTION_INTEGRATION.md + docs/USE_CASE_CATALOGUE.md (vasic-digital/docs_chain).

Non-compliance is a release blocker.

§11.4.107 — Anti-bluff AV/test-validation techniques mandate (User-driven research, 2026-06-02)

Forensic anchor (genericised, 2026-06-02): a test PASSED on a SINGLE captured frame showing “a picture” on the target output — but the picture was a FROZEN / STALE frame from the previously-played content (stale-producer / stuck-decoder), so the feature was broken for the user while the test was green; a sibling incident FLASHED media on the WRONG output for ~1 s before routing with no test sampling that window; a third class shipped a comparator that PASSED its own deliberately-degraded fixture (the analyzer, not the feature, was the bluff). §11.4.5 mandates captured evidence + a presence pass; §11.4.107 raises the bar to **liveness + correct-routing + self-validated-analyzer**.

Every test asserting audio/video output is genuinely playing/advancing for the end user MUST satisfy ALL of: (1) **a single captured frame is NOT proof** — prove LIVE, ADVANCING frames over a steady-state window via a **freeze-detection oracle** (near-duplicate / freezedetect-class filter OR perceptual-hash adjacent-frame distance, **NOT byte-identical compare** — byte-identity is only a zero-cost pre-filter); (2) an **independent frame-advance counter from the platform’s compositor/decoder telemetry** must increase across the window (a different-domain oracle — flat counter ⇒ stuck decoder ⇒ FAIL even if pixels appear to move); (3) **loading/buffering is a distinct state** — wait for genuine playback-start before judging liveness, never false-FAIL a still-loading stream nor false-PASS a spinner; timeout+unreachable ⇒ SKIP-with-reason per §11.4.3, timeout+reachable ⇒ FAIL; (4) **not-stale-from-previous cross-check** — new content’s first frame ≠/ previous content’s last frame; (5) **measured FPS / no-lost-frames within tolerance**; (6) **no-flash-on-the-wrong-output** — sample the non-target output at high frequency during a routing transition, any content frame there ⇒ FAIL (content-protection regime classified explicitly, never guessed); (7) **drive through the realistic feed/UI path, not deep-link shortcuts** (shortcuts bypass the transition paths where bugs live; UI-not-introspectable ⇒ operator-attended fallback per §11.4.52, never fake-PASS); (8) **metamorphic relations solve the oracle problem when there is no golden source** (same content on output-A vs output-B must match; paused ⇒ counter stops; 2× speed ⇒ ~2× advance rate); (9) **full-reference quality metrics (SSIM/VMAF/ΔE2000) vs a golden source for owned content**; (10) **mutation-test every analyzer with a golden-good + golden-bad fixture pair** so the analyzer itself provably cannot bluff (an analyzer that PASSES its golden-bad fixture is a bluff gate — §1.1 applied to the analyzers); (11) **per-channel audio RMS/loudness (EBU R128) + XRUN/underrun census** — a single aggregate RMS misses a dead channel; (12) **OCR overlay/subtitle detection needs a per-word confidence floor + ROI** to avoid BOTH false-positive

(false-FAIL) and false-negative (false-PASS); (13) **thresholds calibrated on the project's own fixtures, not hardcoded from literature** (§11.4.6 no-guessing).

Honest gap (§11.4.6): true photon FPS at the sink has no clean software oracle — compositor/decoder counters measure the presentation pipeline; flag the gap, never claim it.

4-layer coverage per §11.4.4(b): pre-build gate (a CM-AV-LIVENESS-NO-FROZEN-FRAME-class gate asserting every output-is-playing test references the liveness battery not a single frame + an analyzer-self-validation gate wiring the golden-good/golden-bad fixtures into meta-test) + on-device/runtime test + paired §1.1 meta-test mutation (single-frame-only assertion → gate FAILs; analyzer that PASSes its golden-bad fixture → self-validation FAILs) + HelixQA Challenge. Every PASS via `ab_pass_with_evidence` citing the motion / not-stale / frame-advance / fps / per-channel-loudness / metamorphic / self-validation artefacts (`video_display` / `audio_output` / `subtitle_render` per §11.4.69) — never a single screenshot.

Classification: universal (§11.4.17) — platform-neutral AV/test-validation techniques reusable by ANY project validating media playback or any pixel/audio output; the project supplies its concrete capture mechanism + calibrated thresholds per §11.4.35. Composes with §11.4.5 (its strict expansion), §11.4.6, §11.4.50, §11.4.68, §11.4.69, §11.4.85, plus §11.4.2 / §11.4.3 / §11.4.13 / §11.4.48 / §11.4.52 / §1.1. Propagation gate CM-COVENANT-114-107-PROPAGATION enforces the literal anchor `11.4.107` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → gate FAILs (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.107.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--single-frame-proves-playback`, `--skip-liveness`, `--byte-identical-freeze-OK`, `--no-frame-advance-counter`, `--skip-not-stale-check`, `--allow-wrong-output-flash`, `--deep-link-shortcut-OK`, `--unvalidated-analyzer-OK`, `--aggregate-rms-suffices`, `--hardcoded-thresholds-OK` flag exists.

§11.4.108 — Four-layer fix-verification + runtime-signature-as-definition-of-done mandate (systematic-debugging Phase 4.5, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): across one batch, multiple fixes were “green” at every gate yet NOT working for the end user — a change present in source and passed by both the source pre-build gate AND the post-build gate never reached the boot-time command-line embedded in the deployed boot

artifact (output feature dead despite all-green); sibling fixes correctly built into the system image were masked by stale per-user overlay copies from a previous deployment (the running code was the stale shadow, not the fresh build); deployment did not wipe the mutable overlay so the shadow survived; validation ran against whatever was running — the stale shadow — and reported green on code that was never exercised. Per `superpowers:systematic-debugging` Phase 4.5: each fix revealing a fresh “fixed-but-not-working” in a DIFFERENT place is NOT independent bugs — it is ONE architectural VERIFICATION flaw; patching each symptom is thrashing.

A fix crosses FOUR distinct layers, and “fixed” at one does NOT imply fixed at the next: (1) **SOURCE** (committed in the source file — what a `grep-the-source` pre-build gate checks), (2) **ARTIFACT** (the change’s BYTES actually landed in the produced build artifact — image / bundle / installer / embedded command-line), (3) **RUNTIME-ON-CLEAN-TARGET** (active on a CLEAN/fresh deployment — the layer the end user experiences — with no stale overlay shadowing the deployed code), (4) **USER-VISIBLE** (the feature works for the end user — the §11.4.5/§11.4.69 captured-evidence layer). Green at layer 1 is the cheapest, least conclusive signal and says nothing about layers 2–4. A gate verifying ONLY the source layer is itself a §11.4 bluff surface.

The mandate (ALL must hold): (1) **Runtime-signature-as-definition-of-done** — a fix is DONE only when its declared runtime signature verifies on a CLEAN/fresh deployment; `source-committed` $\neq$ `artifact-contains-it` $\neq$ `active-on-clean-target` $\neq$ `user-visible-working`. (2) **Every fix declares ONE machine-checkable runtime signature** — a single observable on a clean target proving the fix is BOTH active AND working (a running-system property, a downstream/sink-side report per §11.4.13, a §11.4.5/§11.4.69 captured-evidence assertion, or a counter/state delta — NEVER a re-grep of the source); this **registry of per-fix runtime signatures is the SINGLE SOURCE OF TRUTH for “fixed”** — it REPLACES “a gate greps the source” as the definition of done. (3) **Gates span all four layers** — source (pre-build), artifact (post-build — assert the change’s BYTES landed in the artifact, not merely that the source still contains the change), runtime-on-clean-target (post-deploy — assert the runtime signature on a freshly-deployed clean target), user-visible (§11.4.5/§11.4.69). (4) **Eliminate the stale-deployment / shadow layer by construction** — deployment MUST yield a CLEAN state (wipe the mutable overlay) OR a pre-validation assertion MUST prove `running-artifact == built-artifact` BEFORE any validation runs; validation against possibly-stale deployed state is INVALID and any PASS it produces is a §11.4 PASS-bluff (the test exercised code that was never deployed). (5) **Meta-rule** — ≥ 3 “fixed-but-not-working” discoveries in one cycle signal an architectural VERIFICATION flaw, NOT three independent bugs; on the 3rd, STOP patching symptoms (§11.4.4), fix the VERIFICATION pipeline, and re-certify EVERY item through it on a clean target. (6) **A batch is “validated” only after COMPREHENSIVE per-item runtime-**

signature verification on a clean baseline — NOT after the touched items' own tests pass.

Classification: universal (§11.4.17) — platform-neutral verification-pipeline disciplines reusable by ANY project that builds an artifact, deploys it, and validates the deployed result; the project supplies its concrete artifact format, clean-deployment / equal-artifact mechanism, and per-fix runtime-signature observables per §11.4.35. Composes with §11.4.1 / §11.4.2 / §11.4.4 (clause 5's STOP-and-fix-pipeline is its §11.4.108 specialisation; "clean baseline" = clause 4's clean deployment) / §11.4.5 / §11.4.6 (every layer asserted by evidence, never assumed to propagate) / §11.4.27 / §11.4.40 (§11.4.108 adds the cross-layer per-item runtime-signature dimension) / §11.4.46 (clean-baseline / equal-artifact pre-flight) / §11.4.50 / §11.4.52 / §11.4.69 (a runtime signature is a taxonomy-class observable) / §11.4.102 (Phase 4.5 architectural-flaw recognition IS clause 5's trigger). Propagation gate CM-COVENANT-114-108-PROPAGATION (literal 11.4.108 across the consumer fleet) + recommended per-fix gate CM-RUNTIME-SIGNATURE-REGISTRY + paired §1.1 meta-test mutations (strip the literal → propagation gate FAILs; downgrade a fix to source-only verification → registry gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.108.

Non-compliance is a release blocker regardless of context. No escape hatch — no --source-green-is-done, --skip-artifact-byte-check, --validate-against-running-state, --no-clean-deployment, --skip-runtime-signature, --spot-validate-touched-only flag exists.

§11.4.109 — Mandatory Anti-Forgetting

Enforcement: PreToolUse Guard Hook + Subagent Constitutional Preamble + Orchestrator Pre-Action Checklist (Operator mandate)

Short tag: anti-forgetting-enforcement. **UNCONFIRMED forensic anchor** (pending operator's verbatim mandate quote). Background: emulator subagents ran raw host-direct emulator/adb because the orchestrator forgot to inject the Containers-submodule rule. A rule forgotten at dispatch is not enforcement. Fix: (A) a PreToolUse guard hook (constitution/scripts/hooks/guard-forbidden-commands.sh) that blocks host-direct emulator, force-push/bypass, sudo, and host-power commands at the tool-call boundary regardless of agent memory; (B) a canonical docs/AGENT_GUARDRAILS.md preamble the orchestrator pastes verbatim into every subagent dispatch; (C) an **ORCHESTRATOR PRE-ACTION CHECKLIST** in the same document. Hook = the floor; preamble = the ceiling.

Consuming projects MUST: (1) wire constitution/scripts/hooks/guard-forbidden-commands.sh as a PreToolUse hook in .claude/settings.json (or equivalent runtime settings); (2) maintain docs/AGENT_GUARDRAILS.md containing the SUBAGENT CONSTITUTIONAL PREAMBLE and ORCHESTRATOR PRE-ACTION CHECKLIST headings, with the anchor literal 11.4.109; (3) provide a hermetic hook test suite (≥ 20 cases: every blocked class exits 2, every allowed command exits 0, escape hatch fires for non-power classes, host-power rejects even with escape marker). The hook is inherited by reference — NEVER copied locally (a copy diverges silently).

Gates: CM-ANTI-FORGETTING-ENFORCEMENT (hook present + wired + guardrails doc present + test present) + CM-COVENANT-114-109-PROPAGATION (anchor literal 11.4.109 across every consumer CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 mutations: remove hook entry → gate (2) FAILs; delete guardrails doc → gate (3) FAILs; strip hook from constitution → gate (1) FAILs; strip 11.4.109 → propagation gate FAILs.

Classification: universal (§11.4.17). Composes with §11.4.6 / §11.4.10 / §11.4.75 / §11.4.76 / §11.4.78 / §11.4.79 / §11.4.80 / §11.4.81 / §11.4.84 / §11.4.98 / §11.4.102 / §12.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.109; reference implementation constitution/scripts/hooks/guard-forbidden-commands.sh + constitution/docs/AGENT_GUARDRAILS.md.

Non-compliance is a release blocker. No escape hatch — no --skip-pretooluse-hook, --no-guardrails-doc, --anti-forgetting-optional, --single-layer-sufficient flag exists.

§11.4.110 — Pre-build build-readiness verdict + change-impact clash detection mandate (operator mandate, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): a fix shipped a new system-property *read* but introduced no matching security-policy grant + no property-context type entry — the read was silently denied at runtime + the feature was dead, while every pre-build gate stayed green because the gate only grepped the source file. The defect class generalises: a change introduces a new dependency on a *second* artifact (security policy, context file, service registry, interface freeze-snapshot, symbol table, build-graph node) that the pre-build never cross-checks. This is §11.4.108's SOURCE → ARTIFACT gap shifted LEFT to pre-build time — most such clashes are statically catchable from the change diff itself, before any build. The mandate (ALL hold): (1) a single deterministic **READY-FOR-BUILD verdict** gates the rebuild (orchestrator refuses to start the build unless READY); (2) a **diff-driven**

change-impact + clash detector cross-checks every newly-introduced second-artifact dependency (new property read $\Rightarrow$ property-context type + read-grant; new service $\Rightarrow$ service-context entry; new init service $\Rightarrow$ security label; new/changed stable interface $\Rightarrow$ freeze-snapshot updated; new native-lib dep $\Rightarrow$ module/prebuilt resolves; new policy rule $\Rightarrow$ every type/attribute defined; two batch changes on the same seam $\Rightarrow$ collision acknowledged); (3) **coverage-completeness is a gate** — every changed file maps to ≥ 1 gate + ≥ 1 deployed-target test + ≥ 1 paired §1.1 mutation, baseline ratchets upward per §11.4.50; (4) **two-speed honesty** — grep-speed always-on gates vs REQUIRES_BUILD heavy gates (build-graph parse-only dry-run, full neverallow compilation, ABI diff) as diff-gated opt-in stages, bounded per §12.6/§12.7; (5) every gate + wired analyzer is **anti-bluff by paired §1.1 mutation**; (6) **honest boundary** — a READY verdict proves static internal-consistency + ready-to-build, NOT that the feature works on the deployed target; the regime empties the *preventable* defect class, not the *all-defects* class (runtime/USER-VISIBLE remains §11.4.108's job).

Classification: universal (§11.4.17) — platform-neutral pre-build-rigor disciplines reusable by ANY project that builds an artifact from source; the consuming project supplies its concrete property/service/policy registries, build-graph parse-only command, interface-freeze mechanism, and changed-file $\rightarrow$ gate mapping per §11.4.35. Composes with §11.4.1 / §11.4.4 / §11.4.6 / §11.4.9 / §11.4.27 / §11.4.50 / §11.4.67 / §11.4.75 / §11.4.92 / §11.4.108 (§11.4.110 is the SOURCE $\rightarrow$ ARTIFACT half shifted left to pre-build; §11.4.108 owns RUNTIME-ON-CLEAN-TARGET $\rightarrow$ USER-VISIBLE — together they span all four layers). Propagation gate CM-COVENANT-114-110-PROPAGATION (literal 11.4.110) + recommended per-family gates CM-READY-FOR-BUILD-VERDICT / CM-CHANGE-IMPACT-CLASH-DETECTOR / CM-COVERAGE-COMPLETENESS-GATE / CM-BUILDGRAPH-DRYRUN-WIRED / CM-SEPOLICY-NEVERALLOW-WIRED + paired §1.1 mutations (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.110. Non-compliance is a release blocker. No escape hatch — no --source-green-is-ready, --skip-clash-detector, --skip-coverage-gate, --no-ready-verdict, --grep-proves-neverallow, --skip-buildgraph-dryrun, --build-without-ready-verdict flag.

§11.4.111 — Resolve-by-stable-name-not-by-enumeration-index mandate (research-derived, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): a platform bound an audio output to a kernel-enumerated device index (card=0); a second device of a different class enumerated FIRST at boot and took slot 0, shifting the intended device to slot 1 —

the static index binding now pointed at the wrong device (policy mis-attached the sink, mis-assigned its TYPE, the output switcher labelled the AV receiver “Wired Headphone” + routing collapsed to stereo). The lower layer of the SAME stack already resolved the device correctly **by name** (scanning the controller-name registry) — proving the brittleness was the *index* binding, not the resolution capability. Generalises to every enumerated resource whose ordinal is assigned at discovery/boot/hotplug time (non-deterministic across reboots, device additions, topology changes). The mandate: any binding to a hardware device / resource handle / enumerated entity (audio cards, display connectors, network interfaces, storage devices, GPU render nodes, input/camera devices, container/process slots) MUST resolve by a **stable identifier** (name / UUID / serial / label / controller-name / content-hash / sink-reported identity) and MUST NOT bind by enumeration index / ordinal / slot, UNLESS the platform documents that ordinal as deterministically pinned AND the pin is itself captured + asserted as part of the binding. Where a stable identifier exists at one layer, every other layer binding the same resource MUST use the same identifier (mixed by-name-here / by-index-there is the structural weak link forbidden here). Honest boundary (§11.4.6): when only an ordinal exists, pin it deterministically via the platform’s own mechanism, capture the pin in the binding’s §11.4.108 runtime signature, and document the residual fragility as UNCONFIRMED:-class risk — never silently trust an unpinned ordinal.

Classification: universal (§11.4.17) — platform-neutral binding-robustness discipline reusable by ANY project binding to enumerated hardware/resources/handles; the consuming project supplies its stable-identifier mechanism (ALSA card name, DRM connector name, NIC predictable name, block-device UUID, etc.) + the layers that must agree per §11.4.35. Composes with §11.4.6 (no-guessing — “the index is *usually* stable” is the exact guess forbidden) / §11.4.8 (mature stacks resolve by name/UUID — reproducing a known-brittle index binding when the by-name path is documented is a §11.4.8 omission) / §11.4.69 (sink-side evidence verifies the by-name binding’s correctness) / §11.4.108 (the by-name binding’s runtime signature asserted on a clean target across the topology that broke the index) / §11.4.110 (a new ordinal binding in a diff is a statically-catchable clash class). Propagation gate CM-COVENANT-114-111-PROPAGATION (literal 11.4.111) + recommended gate CM-RESOLVE-BY-NAME-NOT-INDEX + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.111. Non-compliance is a release blocker. No escape hatch — no --allow-index-binding, --ordinal-is-stable-enough, --skip-name-resolution, --trust-unpinned-index flag.

§11.4.112 — Structural-impossibility won't-fix classification mandate (research-derived, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): deep research (§11.4.8) into relocating protected (content-protection / secure-surface) video to a secondary display PROVED — from authoritative platform/HDCP docs + reproducible captured behaviour (a secure surface is blanked on any output lacking the secure flag; mirror/screenshot of a secure layer returns black) — that the goal is **structurally impossible by platform design**, not a missing feature or unsolved engineering problem. Without a durable classification, such a goal is re-investigated every cycle (re-read the same sources, re-run the same probes, re-derive the same impossibility — compounding wasted effort). The mandate: when deep research per §11.4.8 PROVES (cited authoritative sources AND, where applicable, reproducible captured evidence) a goal is structurally impossible on the target platform (forbidden by platform design / hardware-protocol constraint / documented kernel-or-API limitation — NOT merely unimplemented or hard), the goal MUST be: (1) classified Won't-fix + closed per §11.4.90 with closure reason structurally-impossible; (2) documented with the impossibility evidence — cited authoritative source URLs (per §11.4.99 latest-source verification) + the reproducible probe/captured evidence — in the tracker entry + relevant docs / guide; (3) NOT re-attempted in future cycles — a reopen MUST cite NEW evidence the platform constraint changed (per §11.4.34 + §11.4.7), never merely re-derive the same impossibility; (4) paired with the correct posture — the entry states what the project DOES instead, so “impossible” is never confused with “broken/unhandled”. Honest boundary (§11.4.6): structurally-impossible is reserved for *proven* platform/hardware/protocol impossibility — “could not find a way” / “very hard” / “no time” are Operator-blocked (§11.4.21) or open work, NOT won't-fix; mislabelling them to avoid the work is a §11.4 planning-layer bluff. A future platform change can make the impossible possible; the classification is durable but not eternal.

Classification: universal (§11.4.17) — platform-neutral effort-conservation + honesty discipline reusable by ANY project; the consuming project supplies the specific impossible goal, its platform constraint, and the cited evidence per §11.4.35. Composes with §11.4.6 (FACT with cited evidence, never “probably can't”) / §11.4.7 (reopening requires NEW positive evidence) / §11.4.8 (“NO external solution found — structurally impossible” is the citation) / §11.4.34 (reopen attribution) / §11.4.90 (structurally-impossible is a closure reason in the closed vocabulary) / §11.4.99 (latest-source so the verdict is not stale). Propagation gate CM-COVENANT-114-112-PROPAGATION (literal 11.4.112) + recommended gate CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.112. Non-compliance is a release blocker. No escape hatch — no `--wont-fix-without-proof`, `--reattempt-closed-impossible`, `--skip-impossibility-evidence`, `--impossible-equals-broken` flag.

§11.4.113 — Absolute no-force-push + merge-onto-latest-main mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03): “Any force-push is strictly forbidden! We must for every Submodule take as a base latest commit on Submodule’s main (or master) branch, then on top of it carefully to merge all changes that have to be pushed! Once all merging is carefully done we perform commit and push to all Submodule’s upstreams!”

Force-push is STRICTLY FORBIDDEN with NO exception — `git push --force`, `--force-with-lease`, `+<ref>`, or any history-rewriting overwrite of a remote ref, against EVERY repository this Constitution governs (main repo, this constitution submodule, every owned + nested submodule, every upstream). No operator-approval path, no “after a merge-first audit” path. The mandated 6-step integration procedure for any repo/submodule whose local has commits to publish OR whose mirrors diverged: (1) `git fetch --all --prune --tags` all remotes; (2) set the base to the LATEST commit on the canonical main/master branch (the most-advanced mirror tip); (3) carefully MERGE every change to be published on top of that base — union, preserve BOTH sides, NEVER `-s ours` / `rebase` / `reset` that drops commits (per §9 no-commit-loss); (4) resolve every conflict carefully — no conflict markers, no file dropped, gates/tests still pass; (5) commit the merge (stage only intended files, NEVER `git add -A` in a submodule per §11.4.30); (6) push to ALL upstreams — each push is a fast-forward because the merge commit descends from every mirror tip, so NO force is ever needed (if an upstream still rejects, return to step 1 for it, merge its new tip, re-validate, re-push). **TIGHTENS §11.4.41 / §11.4.71 / §9.2 / CONST-043 — the force-push escape hatch is REMOVED:** even WITH operator approval, even after a clean merge-first audit, force-push is forbidden, because the merge-onto-latest-main path is always available so force is never necessary. Those clauses’ merge-first/fetch-first machinery stays in force as the integration discipline; only their terminal “...then force-push” step is struck.

Classification: universal (§11.4.17) — a platform-neutral integration discipline reusable across every repository. Composes with §2.1 (multi-upstream push — step 6 fans out) / §9 / §9.2 (absolute data safety — this is the no-loss push discipline that makes force unnecessary) / §11.4.4 / §11.4.6 (remote state unknowable without the step-1 fetch) / §11.4.26 (constitution-update conflict resolution IS this procedure) / §11.4.37 (fetch-before-edit → fetch-before-push) / §11.4.40 / §11.4.41 (merge-first stays, force-push step removed) / §11.4.71 (same tightening) / §11.4.88

(background-push still fast-forward-only) / CONST-043 (no force-push authorisable). Propagation gate CM-COVENANT-114-113-PROPAGATION (literal 11.4.113) + recommended gate CM-NO-FORCE-PUSH-ABSOLUTE (scan tracked scripts/hooks for push --force / --force-with-lease / push +<ref> + reject; a §11.4.109-class PreToolUse guard blocks the class at the tool-call boundary) + paired §1.1 mutation (inject a git push --force into a tracked script → gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.113. Non-compliance is a release blocker. No escape hatch — no --force, --force-with-lease, --allow-force-push, --force-push-authorised, --skip-merge-onto-main flag.

§11.4.114 — Last-known-good-tag regression isolation mandate (1.1.8-dev remediation, 2026-06-03)

When a previously-working feature/behaviour is observed broken, the FIRST diagnostic action MUST be to identify the last release tag (or commit/build/deploy) at which it was KNOWN-GOOD and diff/bisect the broken state against it — BEFORE any open-ended root-cause hunt or speculative fix. The known-good revision is the regression oracle: it bounds the search to the commits between good and now (git diff <good-tag>..HEAD --stat of the feature's files = captured evidence), tells you it is a regression REPAIR not a from-scratch design problem, and gives each suspect file a behavioural oracle. When the operator volunteers a known-good tag, that lead is load-bearing and MUST be acted on first. Default to a SURGICAL forward-fix (keep the post-good-tag features, revert ONLY the broken sub-part) over a wholesale revert that loses the batch's other working features — unless the operator prefers wholesale revert. Honest boundary (§11.4.6): “it worked before” is a HYPOTHESIS until the known-good tag is identified AND the feature is confirmed working there; “probably regressed in the last batch” without the diff is a guess; a feature that NEVER worked is not a regression. Composes §11.4.4 / §11.4.6 / §11.4.7 / §11.4.40 / §11.4.43 / §11.4.102 / §11.4.108. Propagation gate CM-COVENANT-114-114-PROPAGATION (literal 11.4.114) + recommended gate CM-REGRESSION-ISOLATED-AGAINST-KNOWN-GOOD + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.114. Non-compliance is a release blocker. No escape hatch — no --skip-known-good-diff, --root-cause-from-scratch, --assume-regression-source, --wholesale-revert-without-isolation flag.

§11.4.115 — RED-baseline-on-the-broken-artifact + polarity-switch mandate (1.1.8-dev remediation, 2026-06-03)

Strict refinement of §11.4.43's RED step. Every RED test MUST be authored to REPRODUCE the defect on the CURRENT, pre-fix artifact (the actual broken build/deployment), capturing positive evidence per §11.4.5 / §11.4.69 / §11.4.107 that the defect is genuinely present — never a synthetic failure the fix is then written to agree with. The SAME test source MUST carry a single polarity switch (env flag / parameter, canonical RED_MODE, default 1 = reproduce-and-assert-defect-present) that flipped to 0 post-fix converts the test into the GREEN regression-guard asserting the defect is ABSENT. One source, two roles: the bug-catcher IS the regression-guard; no separate happy-path test is authored as the primary guard (that demonstrates only that the test agrees with the fix — the §11.4.43 PASS-bluff). RED-on-broken-artifact then GREEN-on-fixed-artifact (on a clean target per §11.4.108) MUST both be captured. Honest boundary (§11.4.6): if the RED run does NOT fail on the broken artifact, that is a finding (close per §11.4.7 with negative evidence, or fix the test) — a RED test that passes on the known-broken artifact is a blind test. Composes §11.4.1 / §11.4.2 / §11.4.5 / §11.4.69 / §11.4.107 / §11.4.4 / §11.4.7 / §11.4.43 / §11.4.50 / §11.4.108 / §11.4.114. Propagation gate CM-COVENANT-114-115-PROPAGATION (literal 11.4.115) + recommended gate CM-RED-POLARITY-SWITCH-PRESENT + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.115. Non-compliance is a release blocker. No escape hatch — no --green-only-test, --skip-red-reproduction, --separate-happy-path-suffices, --synthetic-red-OK flag.

§11.4.116 — Real-time conductor ↔ autonomous-test-framework sync channel mandate (1.1.8-dev remediation, 2026-06-03)

Any autonomous, long-running test/QA/validation framework an external orchestrator (conductor agent / operator) depends on for real-time decisions MUST expose a real-time sync channel: (1) a structured append-only event stream (JSONL or equivalent, one event per line, never rewritten) emitting at minimum session-start / phase-transition / per-test-or-challenge-start / captured-evidence-path / external-call (LLM / vision / sink-probe) / error / per-item-verdict events; (2) an atomically-rewritten status snapshot (single small file written write-temp-then-rename so a reader never sees a torn write) carrying current session/phase/item + counters + last verdict. Verdicts use the closed vocabulary PASS / FAIL / SKIP /

OPERATOR-BLOCKED (§11.4.45). The conductor tails it live so it stays in real-time sync, can §11.4.4-interrupt on a fresh defect, and never idles blindly (§11.4.94 / §11.4.97). Anti-bluff: a verdict event MUST carry the evidence path that backs it (§11.4.69) — a PASS event with no evidence path is a channel-layer PASS-bluff; a snapshot reporting PASS while the stream shows no evidence event for that item is a contradiction → treat as FAIL; an item with no start-event cannot have a verdict-event. When the framework is an owned project-agnostic submodule (§11.4.28), the channel stays project-neutral — the consumer registers its data (endpoints, package ids, sink hosts) at runtime via the public API, never hardcoded. Composes §11.4.4 / §11.4.5 / §11.4.69 / §11.4.27 / §11.4.28 / §11.4.45 / §11.4.52 / §11.4.89 / §11.4.94 / §11.4.97. Propagation gate CM-COVENANT-114-116-PROPAGATION (literal 11.4.116) + recommended gate CM-AUTONOMOUS-FRAMEWORK-SYNC-CHANNEL + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.116. Non-compliance is a release blocker. No escape hatch — no --no-sync-channel, --end-of-session-report-suffices, --verdict-without-evidence-path, --torn-status-write-OK flag.

§11.4.117 — Computer-vision / OCR pixel-oracle fallback for non-introspectable UIs mandate (1.1.8-dev remediation, 2026-06-03)

Any test that needs to drive a UI control OR assert on-screen content MUST NOT assume the accessibility/semantic/DOM hierarchy is the source of truth for what the user sees. When the hierarchy is blank/partial/known-unreliable for the app under test (TV-Compose, leanback, canvas/SurfaceView/GL, games, custom-rendered UIs), the test MUST fall back to a PIXEL ORACLE: (1) DRIVE input by computer-vision template-match (locate a control by its rendered appearance → tap its screen coordinates), not by a hierarchy node id; (2) ASSERT content by ROI OCR (read the rendered text the user reads — caption strip, title, error overlay) with a per-word confidence floor + region-of-interest per §11.4.107(12), not a hierarchy text attribute that may be absent/stale. The tool MUST both drive input AND read pixels — a hierarchy-only tool is NOT a content oracle. This makes §11.4.52's "near-empty hierarchy → INFEASIBLE" constructive: pixel-drive, not only SKIP. Anti-bluff (§11.4.107(10)): the CV/OCR analyzer is self-validated — golden-good fixture PASSES, golden-bad fixture FAILs, wired into meta-test; thresholds calibrated on the project's own frames, not hardcoded (§11.4.6). Honest boundary: when BOTH hierarchy is blank AND pixel oracle is infeasible (secure surface black-captures per §11.4.112, geo-unreachable), SKIP-with-reason per §11.4.3 + tracked operator-attended migration item per §11.4.52 — never fake PASS. Composes §11.4.5 / §11.4.6 / §11.4.48 / §11.4.49 / §11.4.52 / §11.4.107 /

§11.4.112. Propagation gate CM-COVENANT-114-117-PROPAGATION (literal 11.4.117) + recommended gate CM-CV-OCR-PIXEL-ORACLE-FALLBACK + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.117. Non-compliance is a release blocker. No escape hatch — no --hierarchy-is-content-oracle, --skip-pixel-fallback, --unvalidated-ocr-OK, --hardcoded-ocr-threshold-OK flag.

§11.4.118 — Discovery-pressure to confirm known-issue-set completeness mandate (1.1.8-dev remediation, 2026-06-03)

A remediation/release cycle MUST NOT treat “every reported defect is fixed” as “the build is good” — the reported set is biased by what the operator happened to test (“we only see what we test”). After/alongside fixing the reported set, the cycle MUST run a discovery + stress pass across ALL target devices/environments that deliberately exercises subsystems, journeys, and edge cases BEYOND the reported defects — to CONFIRM the reported set is the COMPLETE critical set and surface unreported defects before the end user does. The pass MUST produce PROVABLE coverage: an enumerated list of the subsystems/user-journeys/stress scenarios actually exercised, each with its outcome (no-new-issue, or a new tracker entry per §11.4.15 / §11.4.16). “We found no other issues” is a §11.4 bluff unless accompanied by “here is the enumerated set we exercised” — absence of evidence of looking is not evidence of absence. New findings trigger §11.4.4 interrupt + the §11.4.114/§11.4.115 isolation → RED → fix loop. Honest boundary (§11.4.6): discovery reduces the unknown-unknown surface but does not prove zero remaining defects — the earned claim is “reported set + enumerated discovery set addressed,” with un-exercised subsystems stated as honest coverage gaps (§11.4.3), never silently implied clean. Composes §11.4.4 / §11.4.5 / §11.4.69 / §11.4.25 / §11.4.40 / §11.4.42 / §11.4.52 / §11.4.85 / §11.4.114 / §11.4.115 / §11.4.119. Propagation gate CM-COVENANT-114-118-PROPAGATION (literal 11.4.118) + recommended gate CM-DISCOVERY-COVERAGE-ENUMERATED + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.118. Non-compliance is a release blocker. No escape hatch — no --reported-set-suffices, --skip-discovery-pass, --no-issues-without-coverage-list, --assume-complete flag.

§11.4.119 — Single-resource-owner partitioning for parallel hardware testing mandate (1.1.8-dev remediation, 2026-06-03)

Strict refinement of §11.4.58 / §11.4.103 for hardware contention. When multiple parallel work/test/discovery streams exercise SHARED hardware or any exclusive-access resource (a media-playback path, a single HDMI/audio sink, an exclusive device handle, a serial/JTAG line, a GPU under exclusive capture), exactly ONE stream MUST own each such resource at a time. The exclusive owner drives it (playback, input injection, capture); every other concurrent stream targeting the same resource MUST be READ-ONLY (passive probes — `dumpsys / /proc / /sys` reads / sink-side network probes / log tails). Parallelism is partitioned by resource: distinct devices/sinks/handles run fully concurrent (stream-per-device), but the same device's exclusive resource is single-owner. Ownership MUST be enforced by an advisory lock/token (§11.4.58 L3 + the hardware analogue of §11.4.84 quiescence), event-driven (claim when the resource frees, release the moment work completes so the next queued stream claims it per §11.4.103 auto-backfill). Why: concurrent drivers of one exclusive resource produce CROSS-CONTAMINATED evidence (sink reports the wrong stream's audio, foreground belongs to whichever `am start` landed last, input events interleave) — a PASS under contention is a §11.4 evidence-integrity bluff. Honest boundary (§11.4.6): a “read-only” probe stream MUST issue NO state-changing command against a device it does not own; when a resource genuinely cannot be partitioned, streams serialize on it (single-owner over time), not run concurrently and hope. Composes §11.4.5 / §11.4.69 / §11.4.13 / §11.4.50 / §11.4.58 / §11.4.82 / §11.4.84 / §11.4.103 / §11.4.118. Propagation gate CM-COVENANT-114-119-PROPAGATION (literal 11.4.119) + recommended gate CM-SINGLE-RESOURCE-OWNER-PARTITION + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.119. Non-compliance is a release blocker. No escape hatch — no `--allow-concurrent-resource-drivers`, `--skip-ownership-lock`, `--read-only-may-mutate`, `--contended-evidence-OK` flag.

§11.4.120 — Fix-breaks-its-own-gate reconciliation mandate (1.1.8-dev remediation, 2026-06-03)

When a correct fix causes a pre-existing gate/test to FAIL because that gate asserted the OLD (now-removed-or-changed) behaviour, the gate FAIL is the CORRECT signal the fix landed — it MUST NOT be suppressed by either forbidden response: (1) FAKE-PASSING the gate (editing it to always `pass`, weakening its assertion to a tautology, or deleting it — a §11.4 gate-layer bluff + a §11.4.84

mutation-residue risk); (2) REVERTING the correct fix to satisfy the stale gate (re-introducing the defect). The required response is RECONCILIATION: rewrite the gate to assert the NEW mechanism the fix introduced, backed by captured evidence of the new correct behaviour, AND update its paired §1.1 mutation so the mutation breaks the NEW invariant. Discriminator vs bluffing: after reconciliation the gate + mutation still form a valid §1.1 pair (mutate the new invariant → gate FAILs); a bluffed gate's mutation no longer makes it FAIL (the assertion became a tautology). The reconciliation MUST be a visible, evidence-cited change, never a silent assertion-weakening. Honest boundary (§11.4.6): a post-fix gate FAIL is NOT automatically “stale gate, reconcile” — investigate per §11.4.102 first; the FAIL may be the gate correctly catching a REGRESSION the fix introduced, in which case the FIX is wrong, not the gate. Reconcile ONLY when investigation PROVES the gate asserted old-correct-now-removed behaviour AND the new behaviour is the intended, evidence-confirmed mechanism. Composes §11.4.1 / §11.4.4 / §11.4.6 / §11.4.84 / §11.4.102 / §11.4.108 / §1.1. Propagation gate CM-COVENANT-114-120-PROPAGATION (literal 11.4.120) + recommended gate CM-GATE-RECONCILED-NOT-FAKE-PASSED + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.120. Non-compliance is a release blocker. No escape hatch — no --fake-pass-stale-gate, --revert-fix-for-gate, --weaken-assertion-to-pass, --delete-failing-gate flag.

§11.4.121 — No-commit-while-build-writes-tracked-artifacts mandate (1.1.8-dev remediation, 2026-06-03)

A commit (especially `git add -A` / any broad stage) MUST NOT run while a build/packaging/generation step is actively writing artifacts INTO tracked (version-controlled) directories — doing so races the writer and stages a PARTIAL or stale artifact (a §11.4.108 SOURCE → ARTIFACT integrity failure landed in version control). The commit MUST be deferred until the build step that writes tracked artifacts has COMPLETED, so the tree is quiescent at the artifact layer AND the committed artifacts are the FRESH, whole outputs (no stale pre-rebuild artifact, no half-written file). Before committing tracked build outputs, verify the writing step finished (process exit / completion marker / per-artifact mtime ≥ build-start) — a build still in flight writing tracked dirs is a HOLD on the commit, not a race to win. Build-output analogue of §11.4.84 (no commit while a mutation gate is in flight): both close the “commit captures transient non-final tree state” class at two write-sources. Where build outputs land OUTSIDE version control (gitignored out / / dist /), the race does not apply. Honest boundary (§11.4.6): “the build probably finished” is not “the build finished” — verify with a completion signal; committing

source changes that DON'T touch the build's tracked-artifact directories is fine mid-build. The PROJECT-SPECIFIC instance (which tracked directory + which build steps write it) is recorded in the consuming project's own governance per §11.4.35. Composes §11.4.6 / §11.4.30 / §11.4.58 / §11.4.84 / §11.4.88 / §11.4.96 / §11.4.103 / §11.4.108. Propagation gate CM-COVENANT-114-121-PROPAGATION (literal 11.4.121) + recommended gate CM-NO-COMMIT-DURING-ARTIFACT-WRITE + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.121. Non-compliance is a release blocker. No escape hatch — no --commit-during-build, --stage-partial-artifact-OK, --assume-build-finished, --skip-build-completion-check flag.

§11.4.122 — No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03):

“Never ever remove any application, system component or service from already existing codebase / System without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!”

Forensic case study (FACT). During the 1.1.8-dev burn-down, two shipped capabilities — F2 (an Apple-TV-class application) and F4 (a Huawei HMS / Mobile-Services component) — were removed from the existing System WITHOUT first asking the operator; the operator reversed both. A removal the operator has to discover and reverse after the fact is a defect of the same severity class as a §11.4 PASS-bluff: the System silently lost a user-facing capability the operator never agreed to drop.

No application, system component, service, package, feature, driver, module, library, prebuilt asset — any already-existing end-user capability of the existing codebase / shipped System — may be removed (deleted, dropped from the package set, disabled-into-non-shipping, un-bundled, de-listed, or otherwise made unavailable to the end user) WITHOUT FIRST interactively asking the operator and receiving an EXPLICIT keep-or-remove decision. The question MUST be posed through the platform's interactive clarification mechanism per §11.4.66 (AskUserQuestion on Claude Code) — NEVER a free-text “should I remove X?” buried in narrative, NEVER a silent removal justified post-hoc, NEVER an autonomous removal decision. A silent removal is a **release blocker** regardless of how well-intentioned the rationale (deduplication, “it was broken anyway”, geo-restricted, incompatible, superseded) — the operator decides, the agent asks.

What counts as a removal (non-exhaustive): deleting an app/APK/binary from the build's package set (PRODUCT_PACKAGES / device.mk / equivalent), removing a

service from the init/boot/service-registry set, dropping a kernel module / driver / config from the shipping configuration, un-bundling a prebuilt asset, deleting a submodule or its shipped output, removing a feature flag that gated a live capability, or any edit whose NET EFFECT is “an end-user capability that shipped before no longer ships.” Adding, replacing-with-operator-approved-equivalent, or fixing a capability is NOT a removal. When uncertain whether an edit constitutes a removal, treat it AS a removal and ask (per §11.4.6 no-guessing + §11.4.101 — removal of an existing user-facing capability is high-blast-radius and MUST be operator-confirmed, never autonomously decided). The tracked DROP path: ask → operator approves → mark the item Obsolete (→ Fixed.md) with Obsolete-Details reason feature-removed + an operator-approval citation (§11.4.90) → then remove; the removal never precedes the operator’s yes.

Classification: universal (§11.4.17) — a platform-neutral discipline reusable by ANY project that ships a set of user-facing capabilities; the consuming project supplies its concrete capability-manifest paths per §11.4.35. Composes §11.4.66 / §11.4.101 / §11.4.90 / §11.4.112 / §11.4.6 / §11.4.40 / §11.4.42. Propagation gate CM-COVENANT-114-122-PROPAGATION (literal 11.4.122) + recommended gate CM-NO-SILENT-COMPONENT-REMOVAL + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.122. Non-compliance is a release blocker. No escape hatch — no --remove-without-asking, --silent-removal, --autonomous-removal-OK, --dedup-removal-exempt, --it-was-broken-anyway flag.

§11.4.123 — Rock-solid-proof-or-deep-research mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03):

“Every single reported issue MUST BE fully and 100% validated with rock solid proofs! Nothing can be considered fixed or completed without hard evidence! No false results or bluff(s) of any kind is allowed! If we are not sure on how to achieve full testing, validation and verification of something we MUST ALWAYS perform deep web research for all possible data (articles, documentation, guides, and other resources) and opensourced codebases which we can use to solve our problems and perform testing with validation and verification which produces rock-solid evidence(s) and leaves no space for false results or any kind of bluff!”

Forensic case study (FACT). In the 1.1.8-dev remediation the validation method for two feature classes was, at first, genuinely unclear: relocating a FLAG_SECURE secure surface to a secondary display (pixel capture returns black) and asserting

on-screen content in non-introspectable streaming-app UIs (blank accessibility hierarchy). Rather than declaring them “untestable” or accepting a metadata-only PASS, the cycle performed deep web research (docs/research/testing_frameworks_20260603/) that yielded the CV/OCR/liveness/sink-probe oracle stack (now §11.4.107 + §11.4.112 + §11.4.117) — making rock-solid evidence possible where it had appeared impossible. “Unclear how to validate” is a research trigger, NEVER a bluff licence.

Every single reported issue, every fix, and every claimed completion MUST be fully and 100% validated with rock-solid CAPTURED proof per §11.4.5 / §11.4.69 / §11.4.107 before it may be marked fixed / implemented / completed (§11.4.33 closure vocabulary). Nothing may be considered fixed or complete without hard captured evidence — metadata-only / configuration-only / absence-of-error / grep-without-runtime PASS are all forbidden (§11.4 / §11.4.1); no false results, no bluff of any kind, at any layer.

The research-or-don't-bluff rule (the operative addition): when the agent is UNSURE how to fully test / validate / verify something — when no obvious evidence-producing method exists OR the candidate method would yield only metadata/config/absence-of-error evidence — it MUST ALWAYS first perform deep web research per §11.4.8 + §11.4.99 (official docs, articles, guides, vendor references, standards, issue trackers, reusable open-source codebases) to DISCOVER or BUILD a validation method that produces rock-solid evidence and leaves no space for a false result. Declaring something “untestable” / “not automatable” / accepting a metadata-only PASS WITHOUT first exhausting this deep-research path is itself a §11.4.123 violation — same severity class as a PASS-bluff. The research output (cited source URLs + the evidence-producing method, OR the literal “NO external solution found — original work” per §11.4.8) is the captured proof the path was exhausted. Only after that research genuinely fails may the item be classified PENDING_FORENSICS: / Operator-blocked (§11.4.21) / structurally-impossible won't-fix (§11.4.112) — with the cited research as the evidence the classification is earned, never a convenience.

Classification: universal (§11.4.17) — a platform-neutral discipline reusable by ANY project; the consuming project supplies its concrete capture mechanisms + research corpora per §11.4.35. Composes §11.4.5 / §11.4.6 / §11.4.8 / §11.4.52 / §11.4.69 / §11.4.99 / §11.4.107 / §11.4.118 / §11.4.21 / §11.4.112. Propagation gate CM-COVENANT-114-123-PROPAGATION (literal 11.4.123) + recommended gate CM-ROCK-SOLID-PROOF-OR-RESEARCH + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.123. Non-compliance is a release blocker. No escape hatch — no --metadata-pass-suffices, --skip-proof, --untestable-without-research, --config-only-closure-OK, --bluff-when-unsure flag.

§11.4.124 — Dead/unwired-code investigate-before-remove mandate (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04):

“Before removing any seemingly-dead (zero-importer / unwired) codebase, we **MUST** investigate via git history where/how it was originally used and how it became dead. Removal is permitted **ONLY** when we have captured **PROOF** it is genuinely no longer needed — and that removal **MUST** be its own separate commit with a proper descriptive message. If there is no such proof, the code **MUST** be investigated for where/how it should be wired in properly, and any missing or unwired tests **MUST** be added. We **MUST ALWAYS** be extra careful with any codebase removal.”

“Zero importers / never called / unwired ⇒ dead ⇒ delete” is a **GUESS** (§11.4.6), never a finding — a “no references” result proves only *current* non-reference, not genuinely-unneeded. Before removing **ANY** seemingly-dead element (zero-importer / never-called / unwired function / method / type / file / module / package / asset / config / build target) the agent **MUST FIRST** investigate via git history (`git log --follow`, `git log -S/-G` pickaxe across all history, blame on the deleted call-site) and capture as **FACT**: (1) **WHERE/HOW** it was originally wired in, (2) **WHEN/HOW** it became dead — call-site deleted deliberately / by mistake (regression) / never-completed / refactored-unreachable, (3) whether “no references” is real **OR** a hidden reference the static tool cannot see (reflection / dynamic dispatch / build-tags / codegen / DI / plugin registry / FFI / config-driven wiring). The investigation output (cited commits + determination) is the captured evidence. **Removal is conditional**: permitted **ONLY** with captured **PROOF** the element is genuinely no longer needed; that removal **MUST** be its **OWN SEPARATE COMMIT** (independently reviewable + revertible, composes §11.4.84 quiescence + §11.4.92 multi-pass) with a descriptive message citing the git-history evidence — plus §11.4.122 operator-confirmation when the element is an end-user capability; the §11.4.90 tracked path marks it **Obsolete** (→ `Fixed.md`). **No proof ⇒ do NOT delete**: investigate **WHERE/HOW** to wire it in properly (restore a mistakenly-deleted call-site per §11.4.114; finish never-completed wiring) **AND** add any missing / unwired tests (§11.4.27 / §11.4.43 / §11.4.115 — the missing test is part of why it drifted into apparent-deadness). **Extra-caution default**: when uncertain whether removal-proof is sufficient, default to **NOT** removing (investigate + wire + test) per §11.4.6 + §11.4.101 + §11.4.122; “probably dead” is never sufficient — the bar is captured proof. Classification: universal (§11.4.17) — the consuming project supplies its static-analysis / importer-graph tooling + hidden-reference mechanisms per §11.4.35. Composes §11.4.6 / §11.4.8 / §11.4.84 / §11.4.90 / §11.4.92 / §11.4.101 / §11.4.114 / §11.4.122 / §11.4.27 / §11.4.43 / §11.4.115. Propagation gate **CM-COVENANT-114-124-PROPAGATION** (literal 11.4.124) +

recommended gate CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE (a net-deletion commit must be removal-only + cite the git-history investigation OR be part of a tracked Obsolete item) + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.124. Non-compliance is a release blocker. No escape hatch — no --zero-importers-means-dead, --delete-unwired-on-sight, --skip-git-history-investigation, --remove-without-proof, --bundle-removal-with-other-work flag.

§11.4.125 — Code-review-agent gate before pre-build + main build (mandatory multi-layer review) (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04):

“After all fixes/changes/implementations are done, BEFORE running pre-build tests and the main build, dispatch code-review agent(s) that analyze all work done + all existing data/facts + the existing codebase + current git history to determine quality, safety, and whether the fixes/changes will REALLY work; they MUST validate and verify that every test covering the fixes/changes genuinely validates the work with NO chance of false results or bluff of any kind. Any finding MUST be fixed, polished, improved, and covered with additional tests before the build proceeds. Multiple strong layers of checks.”

After all fixes / changes / implementations in a batch are done, and BEFORE running the pre-build test sweep AND the main (artifact) build (for ANY project), the agent MUST dispatch one or more dedicated code-review agent(s) (subagent-driven by default per §11.4.70/§11.4.20) performing a multi-layer review that: (1) analyzes ALL work done in the batch (every fix/change + its source diff + stated intent); (2) analyzes ALL existing data + facts (captured evidence per §11.4.5/§11.4.69/§11.4.107, tracker entries, prior findings, the §11.4.108 runtime-signature registry); (3) analyzes the existing codebase (blast radius per §11.4.92, cross-feature interaction, contract integrity of every dependency); (4) analyzes current git history (what each change touched, how it composes with concurrent/recent work, whether it reproduces a known-broken pattern per §11.4.114/§11.4.124); (5) determines quality + safety + will-it-REALLY-work (robust + not error-prone — no solve-A-create-B; no host/data/security regression; genuinely delivers the end-user-visible behaviour per §11.4/§107); (6) validates + verifies the tests covering the work — every covering test genuinely exercises the work-under-test and catches its negation, with ZERO chance of a false result or bluff (a test that PASSES on broken-for-the-user work, a metadata-only/config-only/absence-of-error/grep-without-runtime assertion, or a gate whose paired §1.1 mutation does not make it FAIL is a finding). Any finding (defect / error-prone change / safety risk / will-not-

really-work / bluff-or-false-result-capable test / missing-coverage gap) MUST be fixed, polished, improved, and covered with additional tests (four-layer per §11.4.4(b), TDD-RED-first per §11.4.43/§11.4.115) BEFORE the pre-build sweep + main build proceed; the review iterates (re-review after each remediation) until no blocking findings remain. The review is itself anti-bluff (its conclusions are captured evidence per §11.4.5/§11.4.69; a rubber-stamp review of a defective batch = PASS-bluff). It is one of MULTIPLE STRONG LAYERS — complementing, never replacing, the §1 pre-build sweep, §11.4.92 multi-pass (author-side self-review; §11.4.125 adds the structurally-separated reviewer seam per §11.4.70), §11.4.108 four-layer fix-verification, §11.4.110 build-readiness verdict, and the post-build / runtime-on-clean-target / user-visible layers. Composes §11.4 / §11.4.1 / §11.4.4 / §11.4.6 / §11.4.40 / §11.4.43 / §11.4.50 / §11.4.70 / §11.4.20 / §11.4.92 / §11.4.102 / §11.4.107 / §11.4.108 / §11.4.110. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-125-PROPAGATION (literal 11.4.125) + recommended gate CM-CODE-REVIEW-GATE-BEFORE-BUILD (build starts only with a fresh code-review-completed marker for the current batch, produced after the last fix + before the pre-build sweep + main build) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.125. Non-compliance is a release blocker. No escape hatch — no --skip-code-review, --build-without-review, --no-review-gate, --review-optional, --trust-the-author flag.

§11.4.126 — Default autonomous-loop working mode from first prompt (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04):

“Make sure that you continue work in endless fully autonomous loop, do not stop until new fully validated and verified version (tag) is created and published (all submodules and main repo) or IN A CASE OF some other main stream work until it is fully completed with all side work streams and nothing else is left in our working queue! THIS MUST BE ALWAYS the default working mode without us asking you! We tend to achieve ABSOLUTE EFFICIENCY, with this and all other projects which will incorporate this MANDATORY RULE / CONSTRAINT!!! This way of (your) working will be ALWAYS applied / followed / executed / fully respected, as soon as we assign / send first request (prompt) in the session! This stops only if we explicitly say so or nothing is left to be done in current working scope (release that will come / upcoming version)!!! Any mimicking (imitation) of this behavior / rules / mandatory constraints, false results or any kind of bluff(s) is ABSOLUTELY FORBIDDEN!!!”

The endless fully-autonomous loop is the **DEFAULT working mode**, engaged automatically the moment the operator sends the FIRST request / prompt of a session — the operator **MUST NOT** have to ask for it, request it, restate it, or re-enable it per session. §11.4.87 framed the endless-loop covenant as an explicit-instruction opt-in (“continue in endless loop fully autonomously” or a semantically-equivalent phrasing); §11.4.126 is the **capstone** that promotes the same covenant to always-on: from the first prompt onward, every agent operates in the §11.4.87 loop discipline as the standing default, with §11.4.94 zero-idle, §11.4.97 maximum-idle-use, §11.4.101 autonomous-decision-over-blocking, and §11.4.103 continuous-parallel-stream all engaged by default — no per-session activation handshake. The continuation contract: the loop continues until ONE of two terminal conditions holds — (A) **Release scope** — a new, fully-validated-and-verified version (tag) is created AND published across all owned submodules AND the main repo to all configured remotes (per §2.1 multi-upstream push + §11.4.40 full-suite-retest-before-tag + §11.4.113 absolute-no-force-push merge-onto-latest-main); OR (B) **Non-release main-stream scope** — the main-stream goal is fully completed AND every side work stream is done AND the working queue holds nothing left for the current scope. Until (A) or (B) holds, the agent **MUST** keep working (claim the next priority item, dispatch the next parallel stream, progress every non-blocked item per §11.4.42 / §11.4.72 / §11.4.94 / §11.4.103). The loop **STOPS ONLY** on: (1) the operator explicitly saying so (STOP / pause / end); (2) nothing left to do in the current working scope — the upcoming release / current main-stream goal — with the queue genuinely empty per the (A)/(B) terminal conditions; (3) a §12 host-session-safety demand (the loop yields to host safety unconditionally). Idle-while-blocked parks one work unit, it does not stop the loop — the agent keeps progressing every non-blocked item in parallel per §11.4.101 + §11.4.94 + §11.4.97. Goal — ABSOLUTE EFFICIENCY (no operator-side restart overhead, no idle gaps, no stop-and-wait round-trips); applies to this project AND every project that incorporates this Constitution. Anti-bluff: mimicking / imitating this loop behaviour, narrating continuation without performing it, fabricating progress, or emitting false / bluff results of ANY kind is **ABSOLUTELY FORBIDDEN** — this composes the entire §11.4 anti-bluff covenant family (§11.4 / §11.4.1 / §11.4.2 / §11.4.5 / §11.4.6 / §11.4.50 / §11.4.69 / §11.4.107); the agent **MUST** genuinely perform the continuous work and capture positive evidence for every closure, and a report claiming the loop ran while no real work / no captured evidence was produced is a §11.4 PASS-bluff at the operating-mode layer. Classification: universal (§11.4.17). Composes with §11.4.87 (the endless-loop covenant — §11.4.126 promotes it from opt-in to always-on default) / §11.4.94 / §11.4.97 / §11.4.101 / §11.4.103 / §11.4.66 / §11.4.6 / §11.4.40 / §11.4.42 / §11.4.72 / §11.4.113 / §2.1 / §12. Propagation gate CM-COVENANT-114-126-PROPAGATION (literal 11.4.126 across the consumer fleet) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.126. Non-compliance is a release blocker. No escape hatch — no --ask-before-continuing, --single-turn-only, --not-default-loop, --mimic-OK flag.

§11.4.127 — Session-handoff resumption-prompt mandate (User mandate, 2026-06-06)

Forensic anchor — verbatim user mandate (2026-06-06): “make sure that in situations like this now when new session is needed you ALWAYS prepare such sentence - which will be valid for particular moment and the phase of the project and enough for work to continue.”

When the agent determines a fresh session is needed (context-window limits, performance degradation) OR the operator asks whether a new session is needed / requests a handoff, the agent MUST ALWAYS prepare + proactively provide a ready-to-paste **resumption prompt valid for that EXACT moment and project phase** — self-contained enough that pasting it into a fresh session resumes work with ZERO loss. Two variants on demand: a SHORT first-sentence (“Read <handoff docs>, then continue <terminal goal> ...”) AND a FULL detailed block. The prompt MUST: (1) point to the live handoff doc(s) — .remember / remember.md if present + docs/CONTINUATION.md per §12.10 — read FIRST + git fetch --all; (2) state current PHASE + immediate NEXT action + terminal goal; (3) embed exact live-state anchors (build IDs / artifact MD5, device/target serials, commit HEAD, in-flight PIDs + log paths, captured-evidence paths); (4) restate binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware/target gotchas); (5) be MOMENT-VALID, NEVER a generic template. Handoff doc(s) MUST be current BEFORE the prompt is given (§12.10). A missing / stale / generic prompt is a §11.4.127 violation. Composes §12.10 / §11.4.6 / §11.4.66 / §11.4.87 / §11.4.103 / §11.4.126. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-127-PROPAGATION (literal 11.4.127) + paired §1.1 meta-test mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.127. Non-compliance is a release blocker. No escape hatch — no --skip-handoff-prompt, --generic-prompt-OK, --no-resumption-sentence, --handoff-without-state flag.

§11.4.128 — Always-on device-recording mandate (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): we MUST ALWAYS live-record all available data from all devices we use for testing (or known to be under manual testing), EXTRA carefully so it never harms the device / its performance / causes side effects; raw recordings are NOT processed without need (token-

conscious) and are ALWAYS git-ignored + code-intelligence-excluded; only curated evidence is committed, and only at release prep.

For EVERY test/debug device the project uses + every device under known manual testing, across EVERY reachable transport (USB / wireless ADB / SSH / serial / network introspection API), the project MUST ALWAYS live-record all analysable data: activities, all logs, performance metrics (CPU/memory/I/O/thermal/load), every sink-side report per §11.4.13, and any other live-changeable parameter. (1) **Extra-careful, side-effect-free** — non-invasive read-only probes only, bounded sampling, bounded write-volume, an observer-effect budget; a recorder that perturbs the device-under-test is a §11.4.128 violation, NOT evidence. (2) **Background + parallel + subagent-driven** per §11.4.103 + §11.4.70 — never blocks the main stream. (3) **Token-conscious** — **record-now, analyse-later** — raw data NOT processed without need; the only standing analyse-trigger is release-tag prep (§11.4.40 / §11.4.42) OR explicit operator ask. (4) **Raw is git-ignored (with a §11.4.77 regen-mechanism declaration) AND code-intelligence-excluded (§11.4.78/§11.4.79)** — only CURATED evidence is committed, and only at release prep under docs/qa/<run-id>/ (§11.4.83). (5) **Deterministic layout** <recording-root>/YYYY-MM-DD/<combined main+submodules state hash>/<DEVICE>_<SERIAL>/recording_NNN/<files>. (6) **Anti-bluff** — a recorder claimed running but with no growing corpus is a §11.4 bluff; every curated finding traces to a real raw-corpus path; recorder health is itself captured evidence per §11.4.5/§11.4.69.

Composes §11.4.2 / §11.4.5 / §11.4.13 / §11.4.69 / §11.4.40 / §11.4.42 / §11.4.70 / §11.4.77 / §11.4.78 / §11.4.79 / §11.4.83 / §11.4.103 / §11.4.119. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-128-PROPAGATION (literal 11.4.128) + recommended gate CM-DEVICE-RECORDING-ALWAYS-ON + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.128. Non-compliance is a release blocker. No escape hatch — no --skip-recording, --record-without-layout, --commit-raw-corpus, --index-raw-corpus, --analyse-corpus-always, --invasive-probe-OK flag.

§11.4.129 — Huge-blocker release protocol (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): when a huge blocker is discovered during release validation we MUST stop all testing, fix ALL discovered issues, process all recorded data from the last session, land rock-solid fixes, author NEW validation+verification tests of ALL supported test types, rebuild, reflash, and RESTART the full validation+verification of every fix/change from the last release tag to now — on both devices in parallel, recorded, with real physical captured proofs and no bluff.

On discovery of a HUGE BLOCKER (release-blocking-severity defect: core user-facing capability broken, regression invalidating the in-flight cycle, or blast radius reaching the batch's other fixes) during release validation, execute in order with NO spot-check shortcut: (1) **STOP all testing** on every device (the §11.4.4 test-interrupt STOP at release granularity — continuing past a huge blocker is the §11.4 PASS-bluff). (2) **Fix ALL discovered issues** — not just the blocker; root-cause each per §11.4.102 + isolate regressions against the last known-good tag per §11.4.114. (3) **Process all recorded data from the last session** — analyse the §11.4.128 raw-corpus slice (this IS the §11.4.128(3) release-prep analyse-trigger). (4) **Land rock-solid fixes** per §11.4.123 + §11.4.43/§11.4.115 + §11.4.9. (5) **Author NEW validation+verification tests of ALL supported test types** per §11.4.27 + §11.4.85, each anti-bluff + paired §1.1 mutation. (6) **Rebuild (full, not module-only) + reflash to a CLEAN target** per §11.4.108. (7) **RESTART the full validation+verification from the last release tag to now** per §11.4.40 — RESTART, never resume — on both/all owned devices IN PARALLEL per §11.4.103/§11.4.119, every run RECORDED per §11.4.128, real physical captured proofs per §11.4.5/§11.4.69/§11.4.107, no bluff. This anchor BINDS the existing release anchors for the huge-blocker case (adds STOP → fix-all → process-recordings → new-tests-all-types → rebuild → reflash → full-restart + the restart-not-resume rule), citing them rather than duplicating.

Composes §11.4.4 / §11.4.40 / §11.4.42 / §11.4.9 / §11.4.27 / §11.4.85 / §11.4.102 / §11.4.108 / §11.4.114 / §11.4.115 / §11.4.123 / §11.4.128 / §11.4.103 / §11.4.119. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-129-PROPAGATION (literal 11.4.129) + recommended gate CM-HUGE-BLOCKER-FULL-RESTART + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.129. Non-compliance is a release blocker. No escape hatch — no --resume-after-blocker, --spot-validate-after-fix, --skip-recording-analysis, --skip-new-tests, --module-only-after-blocker, --single-device-restart flag.

§11.4.130 — Post-remediation validate-the-fix-FIRST-after-redeploy (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): when a blocker discovered during release validation is fixed and a new artifact (rebuild / new flashing image / redeploy) is produced + the target reflashed, we MUST first re-test the SPECIFIC last-failing features + validate the just-incorporated fixes BEFORE the broader / full validation.

When a blocker / critical failure found during release validation is FIXED and a new artifact is produced + the target reflashed / redistributed / updated, the agent MUST: (1) **re-test the SPECIFIC last-failing features FIRST** (targeted guard tests

for exactly the defects this fix addressed) BEFORE any broader / full-suite validation; (2) **validate the just-incorporated fixes with real captured evidence** — the §11.4.115 RED test flips GREEN at RED_MODE=0 on the new artifact AND the §11.4.108 runtime-signature verifies on the CLEAN target the redeploy produced (metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden per §11.4 / §11.4.1; proof per §11.4.5/§11.4.69/§11.4.107/§11.4.123); (3) **only after the targeted fix is CONFIRMED working** proceed to the §11.4.40 full retest from the last tag to now. Rationale: a first fix attempt may not work / may be incomplete / may regress again under the new artifact — confirming the targeted fix FIRST catches a fix-did-not-take case immediately instead of hours later at the END of a full cycle (then restarting per §11.4.129); cheap-confirmation-first is §11.4.82 applied to the post-blocker reflash. This is the §11.4.46 recent-work-validation gate specialised for the post-blocker-reflash case + the targeted-confirmation phase that GATES §11.4.129's step-7 full-restart. Honest boundary (§11.4.6): “the fix probably took” ≠ “the fix took” — the RED → GREEN flip + runtime-signature on the new artifact is the proof; a still-FAILing targeted re-test re-enters the §11.4.114/§11.4.115 isolate → RED → fix loop, never proceeds to the full cycle on a still-broken fix. Composes §11.4.4 / §11.4.40 / §11.4.46 / §11.4.108 / §11.4.114 / §11.4.115 / §11.4.123 / §11.4.129 / §11.4.82. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-130-PROPAGATION (literal 11.4.130) + recommended gate CM-FIX-FIRST-AFTER-REDEPLOY + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.130. Non-compliance is a release blocker. No escape hatch — no --skip-targeted-retest, --full-cycle-first, --assume-fix-took, --validate-fix-at-end, --skip-red-green-flip-on-new-artifact flag.

§11.4.131 — Standing session-resumption file mandate (User mandate, 2026-06-07)

Forensic anchor — verbatim user mandate (2026-06-07): “Make this markdown a standard file which will be written EVERY TIME when we need fresh session out of the box! It MUST BE always up to date and in sync so whenever new session is created all we have to do is just point to it!”

Every project MUST maintain a SINGLE canonical, always-current **session-resumption file** at a fixed, project-declared standard path (declared once per §11.4.35, never moved without a §11.4.66 operator decision). This file is the OUT-OF-THE-BOX entry point for any fresh session: creating a new session requires ONLY pointing the new agent at this one file. §11.4.131 promotes §11.4.127 (PREPARE a resumption prompt on demand) into a STANDING, version-controlled ARTIFACT — ALWAYS present, ALWAYS in sync. (A) **Existence + fixed path** — exists at the declared path at all times, encoded as a literal path in the project-layer instantiation (§11.4.35), never silently moved. (B) **Always written + always**

synced — (re)written whenever a fresh session is needed OR the live state materially changes (new HEAD, build/artifact id, phase, device/target state, in-flight job, blocking decision) — the §12.10 trigger set; a stale resumption file is a §11.4.131 violation of the same severity class as a §12.10 stale-CONTINUATION violation. (C) **Content (composes §11.4.127)** — both SHORT + FULL variants; points to .remember/remember.md + docs/CONTINUATION.md read FIRST + git fetch; embeds exact live-state anchors (HEAD, build/artifact ids + checksums, device serials, in-flight PIDs + log paths, captured-evidence paths); states PHASE + immediate NEXT + terminal goal; restates binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware gotchas); MOMENT-VALID, never a generic template (§11.4.6). (D) **Export + freshness** — §11.4.65 scope (synchronized .html/.pdf siblings) + §11.4.44 revision header. (E) **Out-of-the-box resumption** — a fresh session, given ONLY this file's path, fully resumes with zero additional context. Composes §12.10 / §11.4.127 / §11.4.65 / §11.4.44 / §11.4.6 / §11.4.66 / §11.4.126. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-131-PROPAGATION (literal 11.4.131) + recommended gate CM-SESSION-RESUMPTION-FILE-PRESENT + paired §1.1 meta-test mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.131. Non-compliance is a release blocker. No escape hatch — no --skip-resumption-file, --ephemeral-prompt-only, --stale-resumption-OK, --generic-template-OK flag.

§11.4.132 — Risk-ordered validation priority mandate (User mandate, 2026-06-07)

Forensic anchor — verbatim user mandate (2026-06-07): “We MUST ALWAYS first test and validate features, functionalities and fixes/changes that have been worked most recently, the ones which were most problematic, which have the most chance to crash or break again, the ones which have been re-opened the most times! Then, after we validate and verify all this with real (physical) proofs and hard evidence, with no false results and bluffs of any kind, we continue with all other existing tests in the test suites! This IS MANDATORY.”

Tests / validations / verifications MUST run in **RISK-DESCENDING order** — the highest-risk set FIRST, and ONLY AFTER that set is fully GREEN with real (physical) captured evidence does the remainder of the suite run. Risk ranking is computed from a CLOSED set of factors, highest-risk first: (a) **most-recently-worked** features / fixes / changes; (b) **historically most-problematic** (longest defect history, most prior fixes/failures); (c) **highest crash/break/regress likelihood** (greatest blast radius / complexity / dependency surface); (d) **most-reopened** per §11.4.55 reopens-count (a high reopen count is the strongest empirical fragility signal). Each item in the highest-risk set MUST pass with real (physical) captured evidence per §11.4.5/§11.4.69/§11.4.107 — no metadata-only / config-only / absence-of-error / grep-without-runtime PASS (§11.4/§11.4.1), no false results, no

bluff (§11.4.6). ONLY AFTER the entire highest-risk set is GREEN with captured proof does the rest of the suite run; running the suite in arbitrary order, or running lower-risk tests before the highest-risk set is GREEN, is a §11.4.132 violation. §11.4.132 REFINES/STRENGTHENS §11.4.130 (generalises “validate the just-fixed items first” to the full risk-ordered set) + §11.4.46 (adds explicit risk-ordering within the recent/high-risk set) + §11.4.42 (applies the implementation-layer priority discipline to VALIDATION ordering). Classification: universal (§11.4.17) — the consuming project supplies its recency / problematic-history / reopen-count sources (e.g. §11.4.93 workable-items DB reopens_count+last_modified) per §11.4.35. Composes §11.4.4/.5/.6/.7/.40/.42/.46/.50/.55/.69/.107/.130. Propagation gate CM-COVENANT-114-132-PROPAGATION (literal 11.4.132) + recommended gate CM-RISK-ORDERED-VALIDATION-PRIORITY + paired §1.1 meta-test mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.132. Non-compliance is a release blocker. No escape hatch — no --skip-risk-ordering, --any-order-OK, --suite-order-fixed flag.

§11.4.133 — Target-System + hardware safety mandate (User mandate, 2026-06-08)

Forensic anchor — verbatim user mandate (2026-06-08): “Make sure that all changes we do to the System are ALWAYS safe for the System itself and for the hardware the system runs on! This is MANDATORY.”

Every change to the TARGET system (firmware, kernel, init/boot scripts, drivers, sysfs/devfreq/voltage/clock/thermal/regulator register writes, partition/bootloader/U-Boot, HAL, framework, prebuilts, device config) MUST ALWAYS be safe for BOTH (a) the target System itself — MUST NOT brick, boot-loop, corrupt data, or render the device unrecoverable — AND (b) the hardware it runs on — MUST NOT exceed safe electrical/thermal/voltage/clock limits or damage panels/storage/radios/regulators. Concrete obligations: (1) reversible-first — verify irreversible high-blast-radius changes (bootloader/U-Boot MD5, partition layout) against known-good values + capture a pre-op backup (§9.2) BEFORE applying; (2) NO unverified hardware-control writes — never write an unverified value to a voltage/clock/regulator/thermal-throttle/current-limit sysfs node or register that could exceed datasheet limits, the safe range established as FACT (§11.4.6), never guessed; (3) thermal/perf changes (forcing a performance governor, pinning the top OPP, disabling thermal management) MUST respect the device’s cooling design, validated by captured thermal evidence; (4) flashing MUST use the sanctioned tool + a freshly-built integrity-verified image — never an ad-hoc partition write or stale/unverified artifact; (5) unprovable-safety = blocked — a change whose target/hardware safety cannot be established from captured evidence is treated as UNSAFE and blocked (§11.4.6 + §11.4.101 reversible-first + §11.4.123 rock-solid-proof). DISTINCT from §12 host-session safety: §12 protects

the DEVELOPER's HOST + session; §11.4.133 protects the TARGET device + its hardware — both apply, neither weakens the other. Classification: universal (§11.4.17) — the consuming project supplies its concrete hardware-control surfaces, datasheet-safe ranges, known-good bootloader/image hashes, and sanctioned flashing tool per §11.4.35. Composes §12 / §11.4.6 / §11.4.101 / §11.4.108 / §11.4.123. Propagation gate CM-COVENANT-114-133-PROPAGATION (literal 11.4.133) + recommended gate CM-TARGET-HARDWARE-SAFETY + paired §1.1 meta-test mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.133. Non-compliance is a release blocker. No escape hatch — no --unsafe-hardware-write, --skip-system-safety, --brick-risk-accepted flag.

§11.4.134 — Code-review iterate-until-GO + rock-solid-evidence mandate (User mandate, 2026-06-08)

Forensic anchor — verbatim user mandate (2026-06-08): “For any fixes/changes given back to us for re-work by the code-review process, once we fix/improve everything per the code-review's requests, we MUST RE-RUN code-review AGAIN until we get a GO from it with NO new issues reported or warnings of any kind! All results produced by this whole process MUST ALWAYS give us rock-solid PHYSICAL evidence that the fixed/improved codebase really works now as expected, with no false results and no bluff(s) of any kind.”

When the §11.4.125 code-review returns ANY finding — BLOCKING, nit, or warning — and the author fixes/improves the batch per that review, the code review MUST BE RE-RUN, and MUST KEEP being re-run after each remediation round, until it returns a clean GO with ZERO new issues AND ZERO warnings of any kind. A single pass that “addressed the findings” is NOT sufficient: the corrected batch MUST pass a FRESH adversarial review (a re-review can surface NEW findings introduced by the very fixes that closed the prior ones — the §11.4.1 fix-A-creates-B failure mode). The loop terminates ONLY on a clean GO (no new findings, no warnings); a residual warning is itself a finding that re-arms the loop. Every round's verdict AND every fix's validation MUST carry rock-solid PHYSICAL captured evidence per §11.4.5 / §11.4.69 / §11.4.107 (captured audio / video / sysfs / dumpsys / sink-side / runtime-signature) proving the fixed/improved codebase REALLY works as expected — never metadata-only / configuration-only / absence-of-error / grep-without-runtime; no false results, no bluff at any round; a reported GO unbacked by captured physical evidence is itself a §11.4 PASS-bluff at the review-loop layer. §11.4.134 REFINES / STRENGTHENS §11.4.125 (iterate “until no blocking findings remain”): it makes the loop EXPLICIT (re-run after every remediation round, not once), raises termination to ZERO findings AND ZERO warnings (not merely zero-blocking), and BINDS rock-solid physical evidence to every round. Classification: universal (§11.4.17). Composes §11.4.125 / §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.50 / §11.4.108 / §11.4.123.

Propagation gate CM-COVENANT-114-134-PROPAGATION (literal 11.4.134) + recommended gate CM-CODE-REVIEW-ITERATE-UNTIL-GO + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.134. Non-compliance is a release blocker. No escape hatch — no --skip-rereview, --single-review-pass, --warnings-ok, --evidence-optional flag.

§11.4.135 — Standing regression-guard suite + every-fixed-defect-gets-a-permanent-regression-test (User mandate, 2026-06-08). Every project MUST maintain a STANDING regression-guard suite that runs on EVERY build+deploy and BLOCKS the release tag on any failure. Every closed defect (stable ticket id, e.g. ATM-NNN) MUST, in the SAME commit as its fix (extending the §11.4.43 DOCUMENT step), register a permanent §11.4.115 RED-on-broken-artifact regression test into the suite — RED_MODE=1 capturing the historical defect on a pre-fix artifact (the proof the guard is real), RED_MODE=0 the standing GREEN guard asserting the defect is ABSENT. A closure without a registered guard is a §11.4.123 violation. The suite runs FIRST in the post-deploy cycle (highest-risk set per §11.4.132) and is a §11.4.40 release-gate blocker. Forensic anchor (FACT): the wrong-subtitle-on-2nd-display defect was “fixed” via a source-side CONTROL_MENU_LABEL_DENYLIST that NO test mirrored or re-ran, so the NEXT chrome class recurred silently while the GREEN suite passed. Industry-standard bug-driven testing (Google content-driven testing; AOSP CTS/Tradefed) made mechanical + enforced. Composes §11.4.4 / §11.4.40 / §11.4.43 / §11.4.46 / §11.4.50 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.124 / §11.4.130 / §11.4.132. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-135-PROPAGATION (literal 11.4.135) + recommended gates CM-REGRESSION-GUARD-REGISTERED / CM-REGRESSION-GUARD-SUITE-WIRED + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.135. Non-compliance is a release blocker. No escape hatch — no --skip-regression-guard, --no-guard-on-close, --guard-optional flag.

§11.4.136 — Real-content end-to-end playback-test mandate (User mandate, 2026-06-08). Refines/strengthens §11.4.107. Any test asserting media playback works MUST drive REAL content (catalog stream or offline reference clip) through the user’s path (§11.4.48 UI-driven → §11.4.117 CV/OCR fallback) and assert it genuinely PLAYS via the §11.4.107 liveness battery PLUS a decoder-health census — a numeric drop-buffer budget, no buffer-timestamp re-order/discard, no codec-reject (cite Android/Media3 ExoPlayer OEM pre-OTA playback-test mandate: “too many dropped buffers” >25, “unexpected presentation timestamp”, “test timed out”). Metadata-only / launch-only / registration-only / single-frame / config-only PASS is forbidden (§11.4 / §11.4.1). A golden/reference clip corpus (BBC ExoPlayer testing samples) is the offline ground-truth. Composes §11.4.5 / §11.4.48 / §11.4.50 / §11.4.107 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Classification: universal

(§11.4.17). Propagation gate CM-COVENANT-114-136-PROPAGATION (literal 11.4.136) + recommended gate CM-REAL-CONTENT-PLAYBACK-TEST + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.136. Non-compliance is a release blocker. No escape hatch — no --launch-proves-playback, --skip-decoder-health, --metadata-playback-pass-suffices flag.

§11.4.137 — Subtitle/caption content-correctness oracle + secure-display-proxy-honesty mandate (User mandate, 2026-06-08). Refines §11.4.117 + §11.4.107 + §11.4.112. Forensic anchor (FACT): tests tasked to “physically verify the 2nd-display subtitle” PASSEd GREEN while subtitles did NOT show / showed WRONG — the streaming player surface is FLAG_SECURE so screencap -d <secondary> returns BLACK (autonomous PIXEL verification structurally impossible per §11.4.112), so the test fell back to the accessibility-scraped/persist.atmosphere.subdebug proxy, and the proxy accepted a chrome/menu LABEL (Аудио и субтитры) as a valid subtitle because the prose floor accepted any multibyte prose and NO menu-label denylist + NO position/cadence check existed. The mandate: a subtitle-correctness test MUST classify the cue’s *content class* — a present cue is NOT a correct cue. CHROME (FAIL) if a known control/menu label (closed multilingual deny-list MIRRORED from source, case-folded incl. non-ASCII), time/numeric chrome, not prose, OUTSIDE the lower safe-title band (CEA-708 9-anchor grid), OR STATIC across the window (real subtitle changes → ≥2 distinct prose cues, a metamorphic relation). DIALOGUE (PASS) only when prose + not-denied + not-chrome + position-ok + cadence ≥2 OR fuzzy-matches the SOURCE-extracted cue via normalized edit distance (§11.4.123 host ground truth). The oracle MUST be self-validated golden-good/golden-bad (§11.4.107(10)) and the deny-list MUST be verified present in the SHIPPED artifact (§11.4.108) — a source-green denylist with no test mirror + no artifact check is the exact recurrence pattern forbidden here. Secure-display honesty (§11.4.112): where FLAG_SECURE makes pixel verification impossible, the rock-solid autonomous proof is the player’s caption telemetry + source-track presence + content-class oracle — NEVER a faked pixel “physical” pass; human-eye pixel confirmation is operator_attended (§11.4.52) with a tracked migration item. App-agnostic (keys off content class). Composes §11.4.3 / §11.4.5 / §11.4.6 / §11.4.107 / §11.4.108 / §11.4.112 / §11.4.115 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-137-PROPAGATION (literal 11.4.137) + recommended gate CM-SUBTITLE-CONTENT-CORRECTNESS-ORACLE + paired §1.1 mutation (strip the denylist/position/cadence check → golden-bad Аудио и субтитры PASSES → gate FAILs). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.137. Non-compliance is a release blocker. No escape hatch — no --present-cue-is-correct, --skip-chrome-oracle, --length-heuristic-suffices, --pixel-pass-on-secure-display, --skip-position-check, --skip-cadence-check flag.

§11.4.138 — Operator-escape => mandatory bluff-audit + permanent guard (User mandate, 2026-06-08). When the operator (or any out-of-band channel) finds a defect that the GREEN test suite passed, this is by definition a §11.4 PASS-bluff — it MUST trigger, before the fix is closed: (1) a §11.4.102 systematic-debugging pass to FACT-root-cause; (2) a bluff-audit identifying the EXACT assertion that should have caught it but didn't, cited to file:line (canonical example: lib/subtitle_content_validation.sh:sub_is_prose() returning TRUE for Аудио и субтитры); (3) a permanent §11.4.135 regression guard registered in the SAME commit as the fix, with its §11.4.115 RED capturing the operator-found defect; (4) the bluff-audit committed under docs/research/<scope>/<defect>_bluff_audit/. Closing an operator-found defect WITHOUT the bluff-audit + permanent guard is itself a §11.4 violation (the bluff that let it through is still live and the defect will recur). Composes §11.4 / §11.4.1 / §11.4.102 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.135. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-138-PROPAGATION (literal 11.4.138) + recommended gate CM-OPERATOR-ESCAPE-BLUFF-AUDIT + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.138. Non-compliance is a release blocker. No escape hatch — no --close-without-bluff-audit, --operator-find-is-just-a-bug, --skip-permanent-guard flag.

§11.4.139 — Fresh-process clean-artifact runtime-signature mandate (User mandate, 2026-06-08). Refines §11.4.108. Before any post-deploy validation — ESPECIALLY a non-pixel proxy verification (the subdebug/accessibility-cue channel used for FLAG_SECURE displays) — the harness MUST assert running-artifact == built-artifact: the deploy yielded a CLEAN target (mutable-overlay/userdata wiped) OR a pre-validation check proves no stale overlay shadows the deployed code (e.g. every guarded package — incl. the Presenter that emits the subtitle cue — resolves to the system partition, no per-user override). A stale shadow of the cue-emitting component (e.g. a Presenter APK predating the denylist) makes the proxy report on code that was never deployed — any PASS is a §11.4 PASS-bluff. Each fix declares ONE machine-checkable runtime signature verified on the clean target (the §11.4.108 registry IS the definition of done); for the subtitle class the signature is “the shipped Presenter APK contains the denylist literal (case-insensitive) AND the subdebug channel emits candidate REJECTED reason=chrome-label for a menu label.” Composes §11.4.46 / §11.4.108 / §11.4.130 / §11.4.135 / §11.4.137. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-139-PROPAGATION (literal 11.4.139) + recommended gate CM-CLEAN-ARTIFACT-RUNTIME-SIGNATURE + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.139. Non-compliance is a release blocker. No escape hatch — no --validate-against-running-state, --skip-clean-precondition, --shadow-OK flag.

§11.4.140 & §11.4.141 — action-prefix system + token-efficiency (cascaded from constitution submodule 60e2d66, CONST-047/049)

§11.4.140 — Universal action-prefix system (**ACTION_NAME ::**) (User mandate, 2026-06-09; GRAMMAR_ADDENDUM 2026-06-09). When a user prompt's FIRST non-blank line starts with a recognised action prefix, you MUST: (1) look the action token up in the action registry constitution/actions/registry.yaml (or \$HELIX_ACTION_REGISTRY); (2) if it is a registered action, REPLACE the prefix with that action's expansion text and apply its rules; (3) execute the remainder of the prompt under the expanded instruction. **Four EQUIVALENT forms** — same action, same expansion, same execution: (1) ACTION_NAME :: <rest> (bare ::), (2) PREFIX::ACTION_NAME :: <rest> (namespaced ::), (3) / ACTION_NAME <rest> (bare slash), (4) /PREFIX::ACTION_NAME <rest> (namespaced slash). Thus BACKGROUND :: x ≡ DEFAULT::BACKGROUND :: x ≡ /BACKGROUND x ≡ / DEFAULT::BACKGROUND x. PREFIX is an action NAMESPACE; the reserved default namespace is **DEFAULT**, and an action runs WITH or WITHOUT the prefix. Grammar (all hold): anchored at the FIRST non-blank line only (mid-prose tokens never match); the action token AND the namespace are UPPERCASE-only [A-Z] [A-Z0-9_]* (lowercase never matches); the namespace separator :: inside the token carries NO surrounding spaces (PREFIX::ACTION_NAME), DISTINCT from the action-body separator " :: " (one ASCII space on each side of :: — avoids C++ Foo::Bar, YAML key: value, URLs) in forms 1/2 and the slash-body separator (one space) in forms 3/4; stacked prefixes (A :: B :: rest) apply outer-to-inner, left-to-right (expand A, re-scan, expand B, then the residual is the task); a leading \ escapes the prefix for BOTH the :: and the slash form (\BACKGROUND :: x, \ / BACKGROUND x — treat literally, strip the backslash, NO expansion) so action names can be discussed. **Conflict rule (slash form):** / ACTION_NAME (form 3) is honored as the action ONLY when ACTION_NAME (case-folded) does not collide with a built-in/host slash command (registry slash_bare: auto + slash_conflicts: [...]); form 4 (/PREFIX::ACTION_NAME) is ALWAYS unambiguous and always honored. An unknown token that matches the grammar shape (any of the 4 forms) but is NOT registered is NEVER silently expanded or silently dropped — ask which registered action was meant (§11.4.66 / §11.4.105) or treat it as a literal prompt, NEVER invent an expansion (§11.4.6); any prompt not satisfying the grammar is an ordinary prompt and the system is a no-op. The registered action **BACKGROUND** expands

to: *“The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!”* (composes §11.4.20 / §11.4.70 subagent-driven, §11.4.58 / §11.4.103 parallel streams, §11.4.89 background execution, §11.4.5 / §11.4.69 / §11.4.107 captured physical evidence, §11.4 anti-bluff). The system is UNIVERSAL (every CLI agent reads this block via its context carrier per §11.4.35), extensible (new action = new registry row), decoupled + reusable (§11.4.28), and loads out-of-the-box. Classification: universal (§11.4.17). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.140. Non-compliance is a release blocker. No escape hatch — no --skip-action-prefix, --ignore-prefix, --no-registry, --invent-expansion-OK, --single-layer-only flag.

§11.4.141 — Token-efficiency mandate (research-derived + operator mandate, 2026-06-09). Every project worked on by AI coding agents MUST cut token spend (input AND output) toward **30–40% of current (a 60–70% reduction)** WITHOUT degrading quality/performance/safety or breaking any existing mechanism, via a composable, safety-ranked measure set: (1) **prompt-cache the static governance prefix** — the always-loaded governance forms a byte-stable cache-breakpointed prefix with no volatile bytes ahead of it; cache reads cost $\sim 0.1 \times$ base input (the dominant cost driver — measured $\sim 170K$ tokens of governance re-sent every turn, externally corroborated by Claude Code issue #24147); caching is transparent so it removes no rule, weakens no gate, changes no verdict — only billing (PRIMARY, biggest + safest lever); (2) **subagent model-tiering + output-to-file** — mechanical non-judgment work (search/grep/status/doc-export/read-only probes) to a Haiku-class model, the strong model RESERVED for all reasoning/verdicts/fix-design (§11.4.102)/code-review (§11.4.125)/demotion (§11.4.7), large output persisted to a file not an inline 350–520K-token transcript; the cheap model never emits a PASS so §11.4.50 + anti-bluff are untouched; (3) **thin always-loaded INDEX + on-demand detail** — concise index (one line per fix/anchor, EACH carrying the literal 11.4.N token so propagation gates pass) with the canonical full text kept gate-scanned in constitution/Constitution.md and reachable in one hop — a de-duplication realising §11.4.35, never a deletion; (4) **CodeGraph/retrieval-first over full-file loading** (§11.4.78/§11.4.79); (5) **output-token reduction** — terse status + effort: "low" on the mechanical allowlist only; (6) **tool-call batching + no re-reads**; (7) **compaction/context-editing for long sessions.** **Mandatory measured proof:** a token-accounting harness measures tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload from the authoritative usage object (input/cache_read/cache_creation/output split; NEVER tiktoken, NEVER the client-side cost estimate), reproduced N times (§11.4.50), pass = AFTER $\leq$ 40% of BEFORE OR the measured best-safe reduction with a cited

cold-cache reason; the AFTER run MUST show ZERO regression on the pre-build sweep + meta-test mutation sweep + propagation gates + a strong-model reasoning probe + a cache-warm proof (`cache_read_input_tokens > 0`) — cost reduction with quality regression is a §11.4 FAIL. The headline number is the *measured* reduction, never the design estimate (§11.4.6/§11.4.123). No measure may break/degrade any existing mechanism, and the rule is structured so none can.

Composes §11.4.5/.6/.20/.40/.50/.58/.69/.70/.78/.79/.80/.103/.106/.123/.125/.128/§12.6/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-141-PROPAGATION (literal 11.4.141) + recommended gate CM-TOKEN-EFFICIENCY + paired §1.1 mutation (inject a pre-breakpoint volatile token → cache collapses → measured reduction falls below bar → gate FAILs).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.141.

Project instantiation (§11.4.35): [ATMOSphere consumer fills in: confirm the Claude Code governance-prefix cache is warm + free of pre-breakpoint volatiles; tier the §11.4.96-SAFE mechanical subagents (search/grep/status/doc-export) to Haiku with output-to-file under `qa-results/`; replace the consumer CLAUDE.md 112-row Applied-Fixes inline table + verbatim anchor restatements with a literal-anchor index pointing at `constitution/Constitution.md`; keep CodeGraph/lumen-first navigation; run the harness BEFORE/AFTER to prove the measured reduction.] Non-compliance is a release blocker. No escape hatch — no `--skip-token-efficiency`, `--no-cache-governance`, `--assert-reduction-without-measuring`, `--tier-down-reasoning`, `--inline-all-governance`, `--tiktoken-estimate-OK` flag.

§11.4.142 — Universal code-review mandate — every change reviewed, always, no exception (User mandate, 2026-06-09). Verbatim operator mandate: “ALL changes we do MUST pass through the code review step!!! ALWAYS!!!” EVERY change made to ANY repository this Constitution governs — without exception — MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. NO change class is exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits — if a diff exists, it gets an independent review. This is the ABSOLUTE form of §11.4.125 (code-review agent gate after a batch, before pre-build sweep + main build): §11.4.142 strips every implicit scoping §11.4.125’s “after all fixes/changes/implementations are done” phrasing could leave open — no “just a doc edit”, no “just a one-liner”, no “the author already self-reviewed (§11.4.92)”, no “trivial change” carve-out. **Independence is load-bearing** — the reviewer MUST be structurally separated from the author (a dedicated code-review agent, subagent-driven by default per §11.4.70 / §11.4.20, or a distinct human), NEVER the author re-reading their own work; §11.4.92’s multi-pass self-evaluation is the AUTHOR-side discipline and PRECEDES (never satisfies) §11.4.142. **The review is itself anti-bluff** (§11.4 / §11.4.1) — a rubber-stamp “looks good” is not a review; it MUST

genuinely analyse correctness, safety (no host §12 / data §9 / target-hardware §11.4.133 regression), will-it-really-work (no solve-A-create-B), end-user behaviour (§11.4 / §107), and test-genuineness (§1.1), its conclusions captured evidence per §11.4.5 / §11.4.69, and **it iterates to a clean GO per §11.4.134** — any finding (BLOCKING / nit / warning) re-arms the loop, acceptance only on ZERO new findings + ZERO warnings with rock-solid physical evidence. Honest boundary (§11.4.6): a passing review does NOT replace §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest — it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses. Composes §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.69 / §11.4.70 / §11.4.92 / §11.4.108 / §11.4.110 / §11.4.125 / §11.4.134 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its reviewer-dispatch mechanism + change-acceptance seam (commit wrapper / merge queue / PR gate) per §11.4.35. Propagation gate CM-COVENANT-114-142-PROPAGATION (literal 11.4.142) + recommended gate CM-EVERY-CHANGE-REVIEWED (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired §1.1 mutation (strip the literal → propagation gate FAILs; accept a diff with no independent-review marker → CM-EVERY-CHANGE-REVIEWED FAILs). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.142. Non-compliance is a release blocker. No escape hatch — no --skip-review, --no-review, --trivial-change-exempt, --doc-edit-exempt, --self-review-suffices, --review-after-commit flag.

§11.4.143 — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, 2026-06-10). Verbatim operator mandate: “All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that’s it mostly! THIS MUST BE ADDED as MANDATORY RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + ATMOSphere project extensions.” Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app’s OWN UI — launch → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play/Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / am start -a VIEW/intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app’s own browse-select-play UI is a §11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator cares about — “the show I picked actually plays” — unproven (the operator’s “we just eventually play on 2nd display sample, and that’s it mostly” gap). The mandate (ALL hold): (1) **Real journey, not a shortcut** —

launch → catalog browse → specific-title selection → real Play/Resume press through the app's own UI (§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** — the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target) via the §11.4.107 liveness battery on the §11.4.136 real-content path + the §11.4.137 subtitle content-correctness oracle for the chosen title's captions; (3) **Non-introspectable UIs use the pixel oracle** — when the app's accessibility hierarchy is blank/unreliable (TV-Compose / leanback / canvas / GL), DRIVE input + ASSERT content via the §11.4.117 CV/OCR pixel oracle, never a hierarchy-only tool; (4) **Login via the credential single-source** (§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** — where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per §11.4.3, secure-surface pixel-blanking per §11.4.112) the test is operator_attended SKIP-with-reason per §11.4.52 + §11.4.3 with a tracked migration item — NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§11.4.6): “the routing fired so the title is playing” is a guess — only the chosen-content liveness + subtitle oracle on the real journey proves it. Composes §11.4.48 / §11.4.107 / §11.4.117 / §11.4.136 / §11.4.137 / §11.4.52 / §11.4.3 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per §11.4.35. Propagation gate CM-COVENANT-114-143-PROPAGATION (literal 11.4.143) + recommended gate CM-VIDEO-REAL-JOURNEY-TEST (every video-streaming-app playback test drives the real browse-select-play journey + confirms the chosen content + subtitles, or SKIPS-with-reason) + paired §1.1 mutation (replace a real-journey test's browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.143. Non-compliance is a release blocker. No escape hatch — no --sample-playback-ok, --skip-real-journey, --deep-link-suffices, --routing-only-pass, --no-title-selection flag.

§11.4.144 — Tracked/recorded-device availability-following mandate (User mandate, 2026-06-10). Direct operator mandate: every device the project is tracking / following / recording (test / debug / manual-testing device, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED — its connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording. The §11.4.128 always-on recorder KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle — no data captured, no offline event logged, no resume, no escalation — so the recording corpus presents a continuous timeline with a silent hole, a §11.4 PASS-bluff at the recording-integrity layer (it claims continuous capture while no data exists for the gap). On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop

and **log an honest offline event** into the recording corpus (§11.4.6 — a silent gap presented as continuous capture is a fabricated-continuity bluff; §11.4.128 — a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project's ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace / reconnect / poll budgets the project's recovery path already uses; (c) **RE-ATTACH** — resume recording / tracking the moment the device returns and log an honest online / resume event; (d) **ESCALATE** to the project's sanctioned device-recovery path (per §11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§11.4.21 / §11.4.101 — high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the bluff this anchor forbids. Composes §11.4.128 (always-on recording — §11.4.144 closes its drop-handling gap) / §11.4.69 (`device_recovery` sink-side positive evidence) / §11.4.6 (honest offline / online events, reused-not-invented timings) / §11.4.14 (watchdog children reaped on stop) / §11.4.21 + §11.4.101 (gated, non-autonomous escalation). Classification: universal (§11.4.17) — the consuming project supplies its concrete tracking transport, device set, already-defined reconnection timings, and sanctioned recovery entry point per §11.4.35. Propagation gate CM-COVENANT-114-144-PROPAGATION (literal 11.4.144) + recommended gate CM-DEVICE-AVAILABILITY-FOLLOWED + paired §1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILS; strip the literal → propagation gate FAILS; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.144. Non-compliance is a release blocker. No escape hatch — no --skip-availability-following, --silent-recording-gap-OK, --no-reconnect-wait, --invent-reconnect-timing, --skip-recovery-escalation, --autonomous-power-cycle-OK flag.

§11.4.145 — Independent multi-angle impact-research per change (User mandate, 2026-06-10). For EVERY fix / change / new feature, INDEPENDENT impact-research agents (subagent-driven §11.4.70/§11.4.20, structurally separate from the author — §11.4.92 self-eval PRECEDES, never satisfies — adversarial “refute-the-change” stance to defeat the documented LLM confirmation-bias failure mode) MUST research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) correctness/logic, (2) regression (what existing working feature could break — via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) latent/dangerous-code — the recents class (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) security (taint/injection/secret/unsafe-API), (5) performance (latency/memory/

thermal under load vs baseline), (6) host/data/target-hardware safety per §12/§9/§11.4.133, (7) cross-feature interaction (shared state/timing/hardware contention), (8) business-logic conformance (matches the spec AND the connected features' contracts, not merely compiles). Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review yet failed at RUNTIME ONLY (ANR on a recents-reaping gesture) because no review angle interrogated the interface contract / concurrency-lifecycle / silently-broken connected feature — the §11.4.108 SOURCE → ARTIFACT → RUNTIME → USER-VISIBLE gap caught only after the fact; §11.4.145 shifts that discovery LEFT. Each angle MUST name the TOOL(S) used + obtain the required DEPTH (the diff, the full changed unit, its declared contract, the spec section, every connected feature — NEVER the diff lines alone) + emit a captured-evidence conclusion per §11.4.5/§11.4.69 (“looks fine” forbidden, §11.4.1); a genuinely-N/A angle is recorded NOT_APPLICABLE: <reason> per §11.4.6, never silently skipped. Output = a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated + affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning) per §11.4.134 before the §11.4.125 review gate. Honest boundary (§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius — it does NOT replace §11.4.108 runtime-signature verification nor §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses. Composes/strengthens §11.4.92/.125/.108/.110/.134/.78/.79/.107/.85/.133/§12/§9/.99/.6/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-145-PROPAGATION (literal 11.4.145) + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a change with a report missing an angle or an unmitigated NO-GO → CM-IMPACT-RESEARCH-PER-CHANGE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.145. Non-compliance is a release blocker. No escape hatch — no --skip-impact-research, --single-angle-suffices, --author-researches-own-change, --diff-skim-OK, --no-go-overridable, --trivial-change-no-research flag.

§11.4.146 — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, 2026-06-10). Every reported problem MUST be handled by a NAMED three-step test workflow — §11.4.146 does NOT re-author its component disciplines, it NAMES + BINDS them into the operator's exact per-defect sequence and adds ONLY the two emphases the components leave implicit. **(STEP 1 — REPRODUCE-FIRST + INVESTIGATE)** before any fix, author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per §11.4.5/§11.4.69/§11.4.107); NEW EMPHASIS (D1) that same RED test is ALSO a deliberate investigation instrument per

§11.4.102(A) — it MUST gather ADDITIONAL forensic data characterising the defect (triggers, boundaries, input/topology scope, adjacent failure modes) that feeds BOTH the fix design AND the STEP-3 extend scope; a RED test used only as a binary present/absent gate with no characterisation satisfies §11.4.115 but NOT §11.4.146 STEP 1. **(STEP 2 — SAME-TEST-CONFIRMS-FIX)** after the fix the SAME test source (the §11.4.115 polarity switch flipped RED_MODE=1→0) confirms the defect ABSENT — RED-on-broken then GREEN-on-fixed, both captured, on a clean target per §11.4.108/§11.4.139, validated-first-after-redeploy per §11.4.130, deterministic per §11.4.50; no separate happy-path test substitutes (the §11.4.43/§11.4.115 PASS-bluff). **(STEP 3 — EXTEND-TO-ALL-CASES, mandatory per-fix)** NEW EMPHASIS (D2) immediately after STEP 2 the test MUST be FANNED OUT across the full case-space of the SAME functionality — all flows, valid + invalid + boundary (§11.4.85 empty/max/off-by-one) + concurrent (§11.4.85 contention) + failure-injection (§11.4.85 chaos) + topology variants (§11.4.3) — confirming no issue of any kind, with PROVABLE enumerated coverage per §11.4.118 (a listed case-set with per-case outcome, never “no other issues found”); a REQUIRED step, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard; each user-visible case carries rock-solid physical evidence per §11.4.123 + is registered into the §11.4.135 standing regression-guard suite; a newly-discovered fan-out defect triggers §11.4.4 test-interrupt + re-enters STEP 1. (ANTI-BLUFF, all steps) every PASS is rock-solid CAPTURED physical evidence per §11.4.123 (§11.4.5/§11.4.69/§11.4.107) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); unclear validation method ⇒ deep-research-before-declaring-untestable per §11.4.123/§11.4.8/§11.4.99; an operator-found defect the green suite missed triggers §11.4.138. Honest boundary (§11.4.6): STEP 3 reduces the unknown-unknown surface but does not prove zero remaining defects (§11.4.118) — un-exercised cases stated as honest gaps, never silently implied clean. Composes §11.4.43/.115/.102/.130/.108/.139/.50/.85/.118/.135/.3/.123/.5/.69/.107/.138/.4/§107/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-146-PROPAGATION (literal 11.4.146) + recommended gate CM-REPRODUCE-FIRST-THEN-EXTEND (every closed defect’s fix carries a §11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1] + its RED_MODE=0 GREEN confirmation [STEP 2] + an enumerated per-functionality extend case-set registered into the §11.4.135 suite [STEP 3]; a closure with the reproduce → confirm pair but NO enumerated extend case-set FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce → confirm pair and no extend-case-set → CM-REPRODUCE-FIRST-THEN-EXTEND FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.146. Non-compliance is a release blocker. No escape hatch — no --skip-reproduce-first, --fix-without-red, --skip-extend-to-all-cases, --reproduce-confirm-suffices, --defer-extend-to-release-sweep, --single-guard-suffices, --no-defect-characterisation flag.

§11.4.147 — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, 2026-06-10). Verbatim operator intent: “any agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!” Forensic case study (FACT): dispatched subagents died on TRANSIENT causes (API Error: Server is temporarily limiting requests rate-limit killed 5 at once, API Error: socket connection closed unexpectedly killed 2, one left PARTIAL edits — two source files written, the dependent SystemUI sliders + doc + test NOT done), each re-dispatched by pure conductor vigilance with NO mechanical guarantee — the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped (lost / forgotten / corrupted). Every dispatched agent/subagent MUST be tracked through its full lifecycle so a crash NEVER loses, forgets, or corrupts its work; a crashed agent is NOT a completed agent (§11.4.6 — “it probably finished before dying” is a forbidden guess; abnormal termination is positive evidence the unit is OPEN). The mandate (ALL hold): **(a) durable agent REGISTRY** — every dispatched agent has an append-only machine-readable registry entry (stable id, task + §11.4.58 file-scope, declared output path(s) + §11.4.5/§11.4.69 evidence dir, dispatch timestamp, status from the CLOSED SET {dispatched | in-flight | crashed | respawned | complete}), reusing the §11.4.116 sync substrate (JSONL event stream + atomic status snapshot; “an entry with no dispatch-event cannot show complete”); the registry is the SINGLE SOURCE OF TRUTH for “is this work owed?” and an unregistered dispatch is itself a violation; **(b) mechanical CRASH-DETECTION + RESPAWN-until-complete** — any abnormal termination (transient rate-limit/socket-close, any non-completion exit, or a terminal error after the runtime’s own retries are exhausted) flips the entry to crashed + keeps the unit OPEN, and the unit is respawned (fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches complete; respawn is the safe/reversible/bounded decision taken autonomously per §11.4.101 (never blocks the loop for a human), transient causes use the project’s ALREADY-DEFINED backoff budgets (never invented, §11.4.6), a deterministically-reproducible non-transient terminal error is investigated per §11.4.102 before further respawn (never a blind retry loop); **(c) PARTIAL-STATE preserve → §11.4.84-check → resume-or-clean-restart** — the crashed agent’s uncommitted edits + output doc are PRESERVED (never silently discarded → lost, never blindly committed → corruption), the respawn runs the §11.4.84 quiescence check on the preserved tree (every modified file accounted-for vs scope, no mutation// always pass/_mutated_* residue, no half-written/torn artifact), then EITHER RESUMES idempotently when it passes OR CLEAN-RESTARTS from a known-good base (§9.2 pre-op backup, reversible §11.4.101) when inconsistent — nothing lost, nothing corrupted; per-agent git worktree isolation (§11.4.58 L4 / §11.4.84) keeps the partial tree from contaminating other streams; **(d) “crash ≠done” COMPLETION criterion** — a unit is DONE only when an agent reaches complete with its required captured evidence/output landed (§11.4.5/§11.4.69 + the

§11.4.116 verdict-carries-evidence-path rule); the endless-loop done-condition (§11.4.87/§11.4.94/§11.4.97/§11.4.126) MUST NOT read satisfied while any entry is dispatched/in-flight/crashed/respawned-not-yet-complete, the zero-idle survey (§11.4.94) treats every non-complete entry as an OPEN item to reclaim, and a registry showing complete without landed evidence is a §11.4 PASS-bluff at the agent-lifecycle layer. Honest boundary (§11.4.6): the registry + respawn guarantee work is not lost/corrupted, NOT that it is correct — the respawned output still crosses §11.4.108 + §11.4.125/§11.4.142 review + §11.4.40 retest; §11.4.147 is the durability layer beneath those, the agent-side analogue of §11.4.144 (same detect → wait/backoff → re-attach/respawn → escalate shape). Composes §11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/§9.2/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-147-PROPAGATION (literal 11.4.147) + recommended gate CM-CRASHED-AGENT-RESPAWN-TRACKED + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a crashed entry as the loop's done-condition complete without landed evidence, OR drop a dispatched agent without a registry entry → CM-CRASHED-AGENT-RESPAWN-TRACKED FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.147. Non-compliance is a release blocker. No escape hatch — no --skip-agent-registry, --crash-equals-done, --no-respawn, --discard-partial-state, --blind-commit-partial, --forget-dead-agent, --loop-done-with-crashed-entries flag.

§11.4.148 — Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices mandate (User mandate, 2026-06-10). BINDS + STRENGTHENS the workable-item discipline (§11.4.15 status / §11.4.16 type / §11.4.54 id / §11.4.21 Operator-blocked / §11.4.91 description-clarity / §11.4.93 + §11.4.95 SQLite SSoT / §11.4.106 docs_chain) into ONE integrity contract spanning DB ↔ docs ↔ external tracker, adding the operator's three emphases: **(D1) no item without a valid status + valid type + stable id — on ALL three surfaces** (an item missing any of the three in the DB, the rendered docs, OR the external tracker FAILs the validator — release-blocker; the id is the cross-surface binding key); **(D2) comprehensive structured description per item** — WHAT it is (§11.4.91 ≥6-word/≥40-char clear meaning) + HOW it manifests + HOW to reproduce (§11.4.115/.146) + ACCEPTANCE CRITERIA (§11.4.123/.69 captured-evidence verdict), in the §11.4.93 DB description column AND the docs AND the tracker, never a stub/§11.4.91-anti-pattern fragment; **(D3) BLOCKED items carry WHY + enumerated unblock CHOICES** — tightens §11.4.21: every Operator-blocked item (incl. Blocked/BLOCKED documented alias normalised to canonical Operator-blocked, never a silent fork §11.4.6) MUST enumerate the closed list of decisions/actions that would unblock it ([A]... · [B]... · [C]..., mirroring §11.4.66's 2–4-option shape); a BLOCKED item with no enumerated choices FAILs; **(D4) regular never-missed bidirectional DB ↔ docs ↔ tracker sync** — the §11.4.93/.95 git-tracked SQLite DB is

the SSoT, docs + tracker DERIVED; full sync (db-to-md / md-to-db byte-identical round-trip §11.4.93 + tracker push) runs regularly + on every change, §11.4.86 drift-proof fingerprint (sha256 of the sorted item keyset, NOT mtime) gates freshness, §11.4.106 docs_chain-bound so drift is mechanically caught not vigilance-dependent; **(D5) generic external-tracker push** carries statuses (collapsed onto the tracker's native set per §11.4.33/.112 when fixed, precise value preserved in a header, never lost), types, assignee (§11.4.104 participant handle, UNSET defaults from a project env var, never hardcoded/logged §11.4.10), and sub-tasks, IDEMPOTENT (dry-run-then-real, match-by-stable-key [<ID>] prefix / id custom-field, present⇒UPDATE/absent⇒CREATE, rate-limited, credential-redacted, sink-side created=N updated=M failed=0 proof §11.4.69), the MACHINERY project-agnostic §11.4.28 (consumer registers its tracker/list/field-map at runtime). Anti-bluff §11.4: every sync + validator pass carries captured evidence §11.4.5/.69; honest boundary §11.4.6 — guarantees well-formed items + agreeing surfaces, NOT that the underlying work is correct (still crosses §11.4.108/.40/.123). Composes §11.4.15/.16/.21/.33/.34/.54/.66/.86/.91/.93/.95/.104/.106/.112/.123/.10/.28/.69/.6/§1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, id prefix, tracker service + list/board id + field map + default-assignee env var, docs_chain context per §11.4.35. Propagation gate CM-COVENANT-114-148-PROPAGATION (literal 11.4.148) + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID / CM-ITEM-COMPREHENSIVE-DESCRIPTION / CM-BLOCKED-UNBLOCK-CHOICES / CM-TRACKER-SYNC-IDEMPOTENT + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.148. Non-compliance is a release blocker. No escape hatch — no --item-without-status, --item-without-type, --item-without-id, --stub-description-OK, --blocked-without-choices, --skip-tracker-sync, --one-way-sync-OK, --non-idempotent-tracker-push, --hardcode-assignee flag.

§11.4.149 — Per-workable-item testing-diary mandate (User mandate, 2026-06-10). Every workable item MUST carry an append-only TESTING DIARY — a chronological record of every test run against it — distinct from §11.4.93 item_history (lifecycle STATE transitions) + §11.4.55 Reopens (reopen cycles): the diary records TEST EXECUTIONS (which may or may not change status). The mandate (ALL hold): **(a)** a test_diary table additive to the §11.4.93/.95 SQLite SSoT, one row per run soft-keyed to the item id, capturing ISO-8601-UTC date_time + tested_by (closed set User|Operator|AI-agent|HelixQA) + result (§11.4.45 PASS|FAIL|SKIP + detail) + in-depth observations (long markdown, facts or explicit UNCONFIRMED:/PENDING_FORENSICS: §11.4.6, the background a future fix needs) + action_taken+status_changed+from/to (did this run change status + WHY) + §11.4.69 evidence_path+feature_class, with a SCHEMA CONSTRAINT making a PASS row WITHOUT a non-empty evidence path impossible (a PASS-bluff rejected by the schema itself); **(b)** BOTH an in-depth

diary doc + a derived at-a-glance summary VIEW (total/pass/fail/skip + last-verdict + last-run + status-changes + distinct testers/feature-classes, DERIVED never duplicated §11.4.93) feeding §11.4.132 risk-ordering; **(c)** four-format per-item Diary/Diary_Summary exports §11.4.65 + §11.4.86-drift-proof-fingerprinted (sha256 of the sorted diary keyset) + §11.4.106 docs_chain-bound so a diary row not exported/pushed FAILs a gate (never-missed); **(d)** external-tracker SUB-TASK model — the item task’s description stays the item (§11.4.148 D2, NOT polluted with diary text), each run a CHILD sub-task (type task) with a {TODO|In-progress|Completed} lifecycle (collapsed per §11.4.33/.112), observations + action + an Evidence: line (path not raw artefact §11.4.10/.13) + a diary-entry idempotency key, reusing the §11.4.148 D5 idempotent rate-limited credential-redacted sink-side-proven push, mapper project-agnostic §11.4.28; **(e)** MINIMAL-LLM deterministic bash/Go tooling (zero LLM in the data path — the observations prose is authored by whoever ran the test; tooling only stores/renders/pushes/validates), 100% test-covered §11.4.27 (unit + integration-against-the-real-tracker with an honest §11.4.3 SKIP-when-token-absent never a faked PASS + export + HelixQA Challenge + paired §1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary §11.4.6 — the diary guarantees a complete auditable per-item test-history, NOT that any single run’s verdict is correct (rests on its own §11.4.69/.107/.123 evidence). Composes §11.4.6/.27/.45/.50/.55/.65/.69/.86/.93/.95/.106/.107/.123/.132/.148/.10/.13/.28/.33/.112/ §1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, diary doc layout, export formats, tracker sub-task field map, docs_chain context per §11.4.35. Propagation gate CM-COVENANT-114-149-PROPAGATION (literal 11.4.149) + recommended gates CM-TEST-DIARY-SYNC / CM-DIARY-PASS-REQUIRES-EVIDENCE + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.149. Non-compliance is a release blocker. No escape hatch — no --skip-testing-diary, --diary-without-evidence, --no-diary-summary, --diary-without-tracker-subtask, --llm-in-diary-data-path, --diary-export-optional flag.

§11.4.150 — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, 2026-06-11). Verbatim operator mandate: “For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don’t have a more serious problem we’re unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream!” For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN

ADDITION to tight systematic-debugging (§11.4.102), the fix, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** digging the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. Forensic case study (FACT, 2026-06-11): a subtitle-on-secondary-display goal verdicted Won't-fix: structurally-impossible (§11.4.112 secure-surface pixel-blanking) was OVERTURNED by a deep multi-angle research pass that surfaced the real OCR-of-the-PRIMARY-display path — a “we can’t” became a shipped capability; the canonical anti-pattern of a structural verdict reached for want of research. The mandate (ALL hold): **(A) No closure-as-fixed/structural WITHOUT a documented deep-research pass** — no item may be marked Fixed/Implemented/Completed (§11.4.33) OR classified structurally-impossible won't-fix (§11.4.112) until a deep multi-angle pass is documented; §11.4.150 makes §11.4.8's pass UNCONDITIONAL (every item, however trivial, at the closure AND structural-verdict gates). **(B) Multiple angles, not a single lookup** — ≥ 2 genuinely-distinct angles (official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent) so a single-source confirmation-bias miss (§11.4.145) cannot pass; a one-link drive-by is NOT deep multi-angle. **(C) Confirm-no-bigger-problem, not just find-a-fix** — explicitly seek evidence the fix does not mask a deeper defect AND the closure/verdict is genuinely safe; “we found nothing worse” requires the enumerated search (§11.4.118). **(D) Latest-source + cited** — LATEST authoritative versions per §11.4.99 (never training-data/memory), each cited by URL + access date in the research artefact AND the closure commit footer (Deep-research <date>: <urls> OR the literal NO external solution found – original work per §11.4.8). **(E) Reopen-breaking is the PURPOSE** — STOP the constant Fixed $\leftrightarrow$ Reopened churn (§11.4.34 / §11.4.55); a reopen whose root cause a deep pass would have surfaced is a §11.4.150 miss; composes §11.4.7 (demotion-evidence) + §11.4.112 (structural verdict now REQUIRES the cited-authorities pass). **(F) ALWAYS in parallel with the main stream** — background subagent-driven (§11.4.70 / §11.4.20 / §11.4.103 / §11.4.89) concurrent with main fix/build/test work, NEVER serialising it or stalling the loop (§11.4.94 / §11.4.97 / §11.4.101 / §11.4.126). **(G) Apply ASAP to every workable item** — retroactively + going forward across the §11.4.93 / §11.4.95 SSoT; items closed without a documented pass are re-audited in the §11.4.40 / §11.4.42 release-gate sweep. Honest boundary (§11.4.6): the pass reduces the unknown-unknown surface + breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 runtime-signature verification, §11.4.125 / §11.4.142 review, or §11.4.40 retest; it is the research layer every fix/closure/structural-verdict additionally crosses, and “we probably don’t have a bigger

problem” without the enumerated multi-angle search is a guess (§11.4.6), never a finding. Composes §11.4.8 / §11.4.99 / §11.4.123 / §11.4.118 / §11.4.145 / §11.4.125 / §11.4.142 / §11.4.7 / §11.4.112 / §11.4.34 / §11.4.55 / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 / §11.4.40 / §11.4.42 / §11.4.93 / §11.4.95 / §11.4.6 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Propagation gate CM-COVENANT-114-150-PROPAGATION (literal 11.4.150) + recommended gate CM-DEEP-RESEARCH-PER-ISSUE (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 mutation (strip the literal → propagation gate FAILs; close an item or reach a structural verdict with no documented multi-angle research artefact / cited footer → CM-DEEP-RESEARCH-PER-ISSUE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.150. Non-compliance is a release blocker. No escape hatch — no --skip-deep-research, --single-source-suffices, --trivial-change-no-research, --close-without-research, --structural-without-research, --serialise-research, --research-from-memory-OK flag.

§11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12). Verbatim operator mandate: “Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project’s release prefix, e.g. helix_translate-1.0.0-dev-0.0.1. The prefix MUST come from HELIX_RELEASE_PREFIX in our .env if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository.” Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (§11.4.28) MUST be prefixed with the project’s release prefix, form <PREFIX>-<version> (canonical example: helix_translate-1.0.0-dev-0.0.1), so a release is identifiable + greppable across every repository it spans (one `git tag -l '<PREFIX>-*'` enumerates the whole release surface). **Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):** (1) HELIX_RELEASE_PREFIX from the project’s .env — authoritative when set; .env is git-ignored per §11.4.30 and the variable is documented in the tracked .env.example (a §11.4.77 re-obtain mechanism, never committed); (2) fallback = the lowercased snake_case form of the project root directory name (no spaces) per §11.4.29 — used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. **The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release** — a release that tags the main repo <PREFIX>-1.2.0 while tagging an owned submodule with a different/unprefixed value is a §11.4.151 violation (the cross-repo grep no longer enumerates the release). Version codes increment monotonically within the prefix

(<PREFIX>-...-0.0.1 → <PREFIX>-...-0.0.2 → ...), never reset, never skipped. Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that its contents are correct — the tag is still created only after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1. Composes §2.1 / §11.4.29 / §11.4.30 / §11.4.40 / §11.4.113 / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Classification: universal (§11.4.17) — the consuming project supplies its concrete prefix value + the HELIX_RELEASE_PREFIX env var per §11.4.35. Propagation gate CM-COVENANT-114-151-PROPAGATION (literal 11.4.151) + recommended gate CM-RELEASE-PREFIX-NAMING (every release tag/version on the main repo + every owned submodule carries the resolved <PREFIX>- prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixed or differing-prefix release tag → CM-RELEASE-PREFIX-NAMING FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.151. Non-compliance is a release blocker. No escape hatch — no --no-release-prefix, --unprefixed-tag, --prefix-optional, --differing-submodule-prefix flag.

§11.4.152 — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, 2026-06-13).

Verbatim operator intent: “For every project that has Firebase Crashlytics enabled / wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!” Every project that has Firebase Crashlytics enabled/wired (SDK linked into a shipping artifact, crash+non-fatal+ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer (a real user hitting a broken feature while the suite reports green).

STRENGTHENS+COMPLETES §11.4.47: §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that SURFACES Crashlytics/Analytics/Performance findings; §11.4.152 owns what happens to each surfaced item AFTER — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven regression-immune closure (§11.4.47 finds it; §11.4.152 fixes it + proves it stays fixed; both mandatory, neither substitutes). The mandate (ALL hold): (1) **continuous monitoring of ALL four surfaces** — fatal crashes, ANRs, performance traces/regressions, AND non-fatals (skipping any one is a PASS-bluff; non-fatals are the silent class — every catch/fallback/recovered-error path is a real degraded UX and MUST be triaged, not ignored because the app did not crash); (2) **systematic-debugging of each (reproduce-before-fix)** per §11.4.102 Iron Law + §11.4.115 RED-baseline-on-the-broken-artifact (the recorded

stacktrace/ANR-trace/non-fatal-context IS the §11.4.5/.69 captured defect-present evidence) BEFORE the fix — a fix without a prior reproducing test is a §11.4.43/.123 violation; unreproducible ⇒ deep-research-before-declaring-untestable §11.4.123/.150; (3) **a fix/improvement for every confirmed issue** against its proven root cause (§11.4.9/.43/.108), “acknowledged in console” is not a fix, muting a recurring issue without a tracked rationale is a §11.4.6/.90 violation; (4) **validation+verification regression-test coverage per closed issue** — in the SAME commit as the fix register a permanent §11.4.135 regression guard (§11.4.115 polarity test: RED_MODE=1 captures the recorded defect on the pre-fix artifact, RED_MODE=0 standing GREEN guard asserting ABSENT) exercising the same user-reachable path the stacktrace identifies, with rock-solid captured evidence §11.4.5/.69/.107/.123 + a closure log citing the console issue id/URL + root-cause + fix commit + the validation/verification test paths; a Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN (the silent-recurrence vector, §11.4.138 operator-escape class applied to field telemetry); (5) **regular cadence** reusing the §11.4.47 five-trigger set (pre-build/pre-flash/pre-distribute/pre-tag blocking; daily + post-deployment-burn-in non-blocking) — a one-time sweep never repeated is a violation; (6) **no false results, no bluff** (§11.4/.1/.6 — “no new issues” requires the enumerated monitored-surfaces+window §11.4.118; “fixed” requires the RED → GREEN flip §11.4.115/.130; a guard whose §1.1 mutation does not FAIL it is a bluff gate). Honest boundary §11.4.6: processing every recorded issue breaks the silent-recurrence vector — it does NOT prove zero remaining field defects nor replace §11.4.108/.40; an issue is “closed” only when its guard is GREEN on a clean artifact, never when the console mark is flipped. Classification: universal (§11.4.17) — the consuming project supplies its Crashlytics handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35; the reference project-level instantiation is the Lava client’s §6.O (Crashlytics-Resolved Issue Coverage — per-issue validation test + Challenge test + .lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md closure log) + §6.AC (Comprehensive Non-Fatal Telemetry — every catch/fallback records a non-fatal with triage context).

Composes

§11.4/.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.150/
§1.1. Propagation gate CM-COVENANT-114-152-PROPAGATION (literal 11.4.152) + recommended gate CM-CRASHLYTICS-ISSUE-FULLY-COVERED (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired §1.1 meta-test mutation (gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#)
§11.4.152. Non-compliance is a release blocker. No escape hatch — no --skip-crashlytics-monitoring, --monitor-crashes-only, --skip-non-fatals, --resolve-without-regression-test, --console-mark-is-fixed, --monitor-once, --mute-without-rationale flag.

§11.4.153 — Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, 2026-06-15).

Every project MUST maintain, under docs/features/, a comprehensive feature Status document set (Status.md + its §11.4.56 Status_Summary.md companion) enumerating EVERY system component, EVERY client app/binary/surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature — including features ported from any incorporated CLI-agent / submodule catalogue (§11.4.74) — with NO single feature left out. (1) **Total categorized coverage** — per-feature table organized Component → Category, reconciled against the actual codebase (a code-present feature missing from the table, or a row with no code, is a §11.4.6/§11.4.118 bluff). (2) **Per-feature fields** — Component / Feature / Category / Implementation / Wiring (genuinely end-user-reachable, not merely compiled, §11.4.108) / Real-use / Tests-coverage (four-layer §11.4.4(b)) / Validation (PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per §11.4.45) / **Video-recording confirmation** (path to the real-use video or honest gap marker). (3) **Mandatory per-feature real-use video** — every user-visible “confirmed” claim backed by a recorded real-use video of a genuine end-user scenario (real prompts → real LLM/service responses → real results), NEVER a frozen/stale frame (§11.4.107), NEVER faked/mockd/demo-loop/bluff response or LLM error passed off as success (§11.4.2/§11.4.5), stored at the project-declared recording path (§11.4.35); a confirmed row with no real video or a bluff video = §11.4 PASS-bluff; autonomous-infeasible ⇒ honest §11.4.3/§11.4.52 SKIP + tracked migration item, NEVER a faked confirmation. (4) **Video-analysis remediation loop** — every video analysed; any defect it surfaces triggers §11.4.102 systematic-debug → fix → §11.4.146 retest → re-record → clean GO per §11.4.134 before “confirmed” (a video exposing a broken feature is a §11.4.4 test-interrupt). (5) **Always-in-sync** — §11.4.45-class roster/corpus-backed, §11.4.106 docs_chain-bound + §11.4.86 drift-proof fingerprint (sha256 of sorted feature-key roster AND sorted video-artefact roster, NOT mtime), re-syncs out-of-the-box; stale = violation. (6) **Four-format export** — HTML + PDF + **DOCX** (this doc class ADDS DOCX to the §11.4.65 HTML+PDF set; other classes unchanged), in sync per §11.4.60. (7) Follows §11.4.44/.45/.56/.57/.59/.60. Honest boundary §11.4.6 — guarantees a complete video-confirmed always-synced ledger, NOT per-feature correctness (rests on each row’s §11.4.69/.107/.123 evidence) and NOT a §11.4.40 retest substitute. Classification: universal (§11.4.17). Composes §11.4.2/.5/.44/.45/.52/.56/.57/.59/.60/.65/.86/.102/.106/.107/.108/.118/.123/.134/.146/ §1.1. Propagation gate CM-COVENANT-114-153-PROPAGATION (literal 11.4.153) + recommended gates CM-FEATURE-STATUS-COMPLETE + CM-FEATURE-STATUS-VIDEO-CONFIRMED + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.153. Non-compliance is a release blocker. No escape hatch — no --skip-feature-status, --feature-without-video, --frozen-video-OK, --bluff-response-video-

OK, --skip-video-analysis, --feature-ledger-incomplete-OK, --no-docx-export, --allow-feature-ledger-drift flag.

§11.4.154 — Window-scoped capture + fresh-corpus rotation for feature/QA recordings (User mandate, 2026-06-15). Verbatim: “recording you perform does record the window containing apps and services, not the whole desktop or monitor screen!” + “All old recording files MUST BE removed when new one starts!” Refines §11.4.2/.5/.107/.153 recording discipline with two capture-hygiene invariants. **(A) Window-scoped, NOT whole-screen** — every feature/QA video MUST capture ONLY the window/surface of the app/service under test (GUI window / CLI-TUI terminal pane / web tab-viewport / device-emulator-simulator frame), NEVER the whole desktop/monitor or unrelated windows; whole-desktop capture leaks operator-private content (§11.4.10/.83), dilutes the §11.4.107 liveness/freeze oracle, and breaks the §11.4.137 OCR/ROI content oracle. Target the window/region by stable identity (window id/title, device serial, browser context, tmux target) per §11.4.111 — never a fixed full-screen device index. Platform genuinely cannot capture below whole-screen ⇒ honest §11.4.3 SKIP + tracked migration item, never a whole-screen pass-off. **(B) Fresh-corpus rotation** — when a new recording run for a scope begins, the agent’s OWN prior in-scope stale recordings at the raw recording path MUST be removed FIRST so the live corpus reflects the current run (§11.4.107 not-stale + §11.4.86 roster-freshness). Honest boundary (§11.4.6 + §9.2): “remove old” = the agent’s own prior recordings for the SAME scope/project ONLY — NEVER another project’s/scope’s/operator-authored files; uncertain ⇒ surface, don’t delete (§11.4.122); committed docs/qa/<run-id>/evidence (§11.4.83) is the durable record, NOT rotated. Classification: universal (§11.4.17). Composes §11.4.2/.5/.10/.83/.86/.107/.111/.122/.128/.137/.153/§9.2/.6. Propagation gate CM-COVENANT-114-154-PROPAGATION (literal 11.4.154) + recommended gate CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.154. Non-compliance is a release blocker. No escape hatch — no --whole-screen-capture-OK, --skip-window-scope, --keep-stale-recordings, --no-corpus-rotation, --full-desktop-recording flag.

§11.4.155 — Project-name-prefixed feature/QA recording filenames (User mandate, 2026-06-15). Verbatim: “All recorded videos MUST START with prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME). Project name MUST be obtained according to the constitution’s own project-name resolution.” Every recorded video the project produces — every feature/QA real-use recording (§11.4.153), every window-scoped capture (§11.4.154), every always-on device recording (§11.4.128), every raw or curated artefact at the project-declared recording path (§11.4.35) + the committed docs/qa/<run-id>/trail (§11.4.83) — MUST have a filename that STARTS WITH the PROJECT-NAME prefix, ALWAYS; an unprefixed recording is a §11.4.155 violation (a multi-project/scope corpus on one

host per §11.4.128/.103 becomes un-greppable + un-attributable — the §11.4.151 identify-and-grep failure on the recording-corpus axis). **Prefix resolution (closed-set, deterministic — §11.4.6, IDENTICAL to §11.4.151):** (1) HELIX_RELEASE_PREFIX from .env (authoritative, git-ignored §11.4.30, documented in tracked .env.example §11.4.77) else (2) lowercased snake_case project-root dir name §11.4.29 — ALWAYS resolvable, zero operator input. SAME prefix for EVERY recording in a checkout so `ls '<PREFIX>--- '*` enumerates the corpus; canonical form `<PREFIX>---<feature-or-scope>---<run-id>.<ext>` (`---` delimits the prefix unambiguously); MUST equal the §11.4.151-resolved release-tag prefix (divergence is itself a §11.4.155 violation — one project, one name). Honest boundary (§11.4.6): the prefix guarantees attribution + greppability, NOT content validity (still §11.4.107/.137/.153) and does NOT relax §11.4.154's window-scope/rotation (rotation removes the agent's OWN `<PREFIX>--- *` only; foreign/operator files surfaced not deleted §11.4.122/§9.2). Classification: universal (§11.4.17). Composes §11.4.151/.128/.153/.154/.111/.83/.6/.29/.30/.35/.77/.86/§1.1. Propagation gate CM-COVENANT-114-155-PROPAGATION (literal 11.4.155) + recommended gate CM-RECORDING-PROJECT-NAME-PREFIX + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.155. Non-compliance is a release blocker. No escape hatch — no `--no-recording-prefix`, `--recording-without-project-name`, `--unprefixed-recording`, `--prefix-optional-for-recording`, `--differing-recording-prefix` flag.

§11.4.156 — All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, 2026-06-15). Verbatim operator mandate: “Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!” Every repository this Constitution governs — main repo, this constitution submodule, every owned + nested submodule we author and push — MUST ship with ALL server-side CI/CD automation DISABLED: no push to any owned upstream may trigger a GitHub Actions run, GitLab pipeline, or equivalent (Jenkins/CircleCI/Travis/Drone/Woodpecker/Bitbucket/Azure, any on: `push/schedule/workflow_dispatch`). GENERALISES + makes ABSOLUTE the §11.4.75 Layer-5 posture (remote CI DISABLED, workflow preserved at a `...disabled-local-only non-.yaml` name a provider ignores) across ALL governed repos; enforcement migrates to the LOCAL §11.4.75 git-hook ritual + §11.4.40 pre-tag sweep, never a remote runner. ALL hold: **(A)** zero active `.github/workflows/*.yaml|yml / .gitlab-ci.yaml / .gitlab/**` / equivalent at the ROOT of any governed repo/submodule (the only place a provider executes); **(B)** “disabled” = a push triggers ZERO runs — delete OR rename to a non-trigger name (§11.4.75 `.disabled/.disabled-local-only`); a live-on: `+if: false` workflow still queues runs, NOT compliant; **(C)** scope = repos we author+push — vendored/third-party nested configs below the root

(AOSP external/**, prebuilts/**, vendored submodules) are INERT (a provider never runs a non-root config), OUT of scope, MUST NOT be mass-edited (§11.4.29 vendor-exempt); test = “does a push to OUR upstream trigger a run?” yes⇒disable, inert⇒document+leave (§11.4.6 verify-not-assume); **(D)** no new CI may be added (release blocker); **(E)** pre-push verify `git ls-files | grep -E '^\.github/workflows/.*\ya?ml$|^\.gitlab-ci\.yaml$'` empty for authored repos, §11.4.109-class PreToolUse guard + gate enforce mechanically. Honest boundary (§11.4.6): file-level disabling stops FILE-triggered runs, NOT provider-side server settings (org-default required workflows, branch-protection required checks, provider scheduled exports) — the operator turns those off; the agent documents what it cannot reach, never claims unachieved completeness. Composes §11.4.75/.29/.6/.40/.42/.109/.113/§2.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-156-PROPAGATION (literal 11.4.156) + recommended gate CM-NO-ACTIVE-CI + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; add a root .github/workflows/x.yml to an authored repo → CM-NO-ACTIVE-CI FAILs). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.156. Non-compliance is a release blocker. No escape hatch — no --allow-ci, --enable-workflow, --keep-pipeline, --remote-ci-OK, --ci-exempt flag.

§11.4.157 — GEMINI.md maintained in lockstep with CLAUDE.md / AGENTS.md / QWEN.md (User mandate, 2026-06-15). Verbatim operator mandate: “Make sure with CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory fact / rule to root constitution Submodule we are inheriting / extending - CONSITUTION.md, CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!” Forensic FACT (2026-06-15): when §11.4.156/§11.4.157 were authored, Constitution/CLAUDE/AGENTS/QWEN carried the family through §11.4.155 but GEMINI.md had silently drifted to §11.4.141 — 14 mandates (§11.4.142–155) never propagated — a §11.4 propagation-bluff (a Gemini-CLI agent reads a stale Constitution). GEMINI.md is a FIRST-CLASS governance context carrier EQUAL to CLAUDE.md/AGENTS.md/QWEN.md, never optional/best-effort. ALL hold: **(A)** five-carrier lockstep — no governance change is complete until GEMINI.md carries it alongside the other three mirrors; GEMINI.md is added to the §11.4.26 propagation + cross-reference set explicitly; **(B)** no silent drift — GEMINI.md lagging the other mirrors’ highest rule is a §11.4.157 violation (§11.4.65-class), back-fill required; **(C)** equal status — GEMINI.md restates the SAME literal 11.4.N anchors the propagation gates require (§11.4.35), fleet count INCLUDES GEMINI.md; **(D)** consumer projects’ own CLAUDE/AGENTS/QWEN/GEMINI bind too (§11.4.35). Honest boundary (§11.4.6): the §11.4.142–155 GEMINI.md back-fill is a tracked release-blocking remediation; claiming GEMINI.md “in sync” while the back-fill is incomplete is itself a §11.4.157 violation. Composes §11.4.26/.35/.17/.44/.65/.140/.156/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-157-PROPAGATION (literal 11.4.157, GEMINI.md INCLUDED) + recommended gate CM-GEMINI-MD-

LOCKSTEP + paired §1.1 meta-test mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.157. Non-compliance is a release blocker. No escape hatch — no --skip-gemini-md, --gemini-optional, --gemini-lag-OK, --four-carrier-suffices flag.

§11.4.158 — Intensive all-feature/flow/edge-case video-recording + read-the-screen content-verification mandate (User mandate, 2026-06-16). Every project MUST be covered by intensive automated testing that exercises + RECORDS every feature/flow/use-case/edge-case (valid/invalid/boundary/concurrent/failure-injection §11.4.85), each recording showing the feature GENUINELY WORKING with REAL results (real prompts → real responses → real outputs §11.4.153) and NO false/simulated/stale/frozen result (§11.4.2/.5/.107) — a recording showing an error/bluff/non-working feature is a FINDING (→ §11.4.153(4)/§11.4.4 fix → retest → re-record), never a confirmation. The testing System MUST ACTUALLY READ every shown log line / message / UI label / dialog / toast / status text and VERIFY it is a genuine working result (OCR/ROI §11.4.117/.137 + confidence floor for pixel surfaces; direct capture for terminal/log; §11.4.107(10) self-validated golden-good/golden-bad analyzer) — “a video was produced” is NOT evidence, “the System read the screen + confirmed a genuine result” is. HelixQA (§11.4.27) MUST drive this exercise → record → read → score pass, PASS only on a read-confirmed genuine result + captured artefact path (§11.4.69). Default recording save path = \$HOME/Downloads (host user’s home Downloads, resolved at runtime never hardcoded) unless a project declares an override per §11.4.35; §11.4.155 project-prefix + §11.4.154 window-scope/rotation + §11.4.128 git-ignored raw corpus + §11.4.83 curated docs/qa evidence apply. Composes §11.4.2/.5/.25/.27/.52/.69/.83/.85/.107/.108/.117/.118/.128/.137/.138/.153/.154/.155/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-158-PROPAGATION (literal 11.4.158) + recommended gates CM-INTENSIVE-RECORDING-COVERAGE + CM-RECORDING-CONTENT-READ-VERIFIED + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.158. Non-compliance is a release blocker. No escape hatch — no --skip-recording-coverage, --video-without-content-read, --happy-path-recording-suffices, --recording-path-anywhere, --unread-recording-OK, --skip-helixqa-read flag. **Project instantiation (§11.4.35):** helix_code overrides the default save path to /Volumes/T7/Downloads/Recordings (recordings are NOT moved to \$HOME/Downloads).

§11.4.159 — Mandatory window-specific video recording + vision validation mandate (User mandate, 2026-06-20). Every feature test, validation, verification, challenge, and QA session that produces video evidence MUST comply with ALL of the following: **(A) Window-specific recording ONLY** — every video MUST record ONLY the target application window (Terminal pane, TUI app, browser tab, emulator frame), NEVER the whole desktop/monitor screen; use macOS screencapture -l<window_id>, Linux xdotool+ffmpeg, or equivalent; whole-desktop capture leaks operator-private content (§11.4.10), dilutes the §11.4.107

liveness oracle, and breaks §11.4.137 OCR/ROI; platform genuinely cannot capture below whole-screen => honest §11.4.3 SKIP + tracked migration item. **(B) MP4 format REQUIRED** — .mp4 (H.264, movflags +faststart, pix_fmt yuv420p); .cast files are supplementary only; auto-convert via `agg+ffmpeg` per §11.4.154(C). **(C) Project-name prefix REQUIRED** — filename MUST start with project name in snake_case from HELIX_RELEASE_PREFIX in .env (§11.4.151) or lowercased project root dir name (§11.4.29); format `<project_name>-<feature>-<timestamp>.mp4`; unprefixes = violation. **(D) Mandatory vision validation** — after EVERY recording, the agent MUST read the terminal output / video content and verify: (i) feature ACTUALLY WORKS, (ii) LLM responses are REAL, (iii) all tests show PASS, (iv) no “TODO implement” / “simulate” / “for now” patterns, (v) output demonstrates end-user working feature; MUST produce a verdict (PASS/FAIL) with evidence path (§11.4.69). **(E) Terminal window cleanup** — after each recording, dismiss/close ONLY the Terminal window used for that recording (window-specific close via `osascript/xdotool`); MUST NOT close windows belonging to other processes. **(F) Real results ONLY** — every recording MUST show REAL working features; errors/empty output/simulated responses trigger fix → retest → re-record. **(G) Re-runnable evidence** — the command shown MUST be re-runnable to produce the same results. **(H) Fresh-corpus rotation** — remove agent’s own prior in-scope stale recordings FIRST (§11.4.154). **(I) Content verification MANDATORY — not duration-based** — the value of a recording is NOT its duration but its CONTENT; a recording MUST demonstrate the ACTUAL feature being used with REAL results; before accepting ANY recording, the agent MUST verify: expected output patterns ARE present (e.g., test PASS lines, API response data, feature-specific output), the feature ACTUALLY WORKS as demonstrated (not just “something ran”), LLM responses are REAL content (not simulated, not placeholder, not empty), every claim of “working” is backed by visible evidence; a 5-second recording showing a feature working correctly is MORE valuable than a 60-second recording of empty terminal; duration is NOT a proxy for quality. **(J) Expected-content specification REQUIRED** — before recording, the agent MUST specify what content SHOULD appear (expected patterns, expected test results, expected API responses); after recording, the agent MUST verify these patterns ARE present; if expected content is MISSING, the recording is REJECTED regardless of duration. **(K) Content-verification recording workflow** — mandated workflow: (1) SPECIFY expected content patterns, (2) RECORD the feature execution, (3) EXTRACT all text from the recording, (4) VERIFY expected patterns are present, (5) CHECK for simulated/placeholder content, (6) ACCEPT only if ALL patterns found AND zero bluffs detected, (7) REJECT and re-record if ANY pattern missing or bluff detected. **(L) Root cause analysis REQUIRED for rejected recordings** — when a recording is rejected (missing expected content, bluff detected, or empty capture), the agent MUST investigate WHY before re-recording per §11.4.102; determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it; simply re-recording without understanding WHY the first attempt failed is a §11.4.159

violation. **(M) Real-time monitoring RECOMMENDED** — for complex features, use real-time monitoring that analyzes output DURING recording (not after); this catches issues immediately and allows corrective action before the recording completes. Classification: universal (§11.4.17). Composes

§11.4.2/.3/.5/.10/.29/.69/.83/.107/.111/.128/.137/.151/.153/.154/.155/.158/§1.1.

Propagation gate CM-COVENANT-114-159-PROPAGATION (literal 11.4.159) + recommended gate CM-WINDOW-VIDEO-VALIDATED + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.159. Non-compliance is a release blocker. No escape hatch — no --whole-screen-ok, --cast-only, --skip-vision-validation, --no-cleanup, --simulated-recording-ok, --unprefixed-recording flag.

§11.4.160 — Vision-verified recording + HelixQA bridge mandate (User mandate, 2026-06-21). Compact summary: every video recording for feature/QA evidence MUST be processed through a vision/OCR pipeline that reads on-screen content and confirms expected results BEFORE acceptance; the recording system MUST provide a bridge feeding captured frames to HelixQA's test infrastructure (or equivalent) for automated read-the-screen verification against SPECIFY-phase expected patterns (§11.4.159(J)); the bridge MUST capture frames at ≤5s intervals, run OCR/vision analysis with a self-validated golden-good/golden-bad analyzer (§11.4.107(10)), compare extracted text against specified patterns, produce a per-frame PASS/FAIL verdict with an evidence path to the frame, and surface failures immediately so the recording can be re-done per §11.4.159(L). The capture interval, OCR confidence floor, and pattern-matching thresholds are project-configured per §11.4.35, calibrated on the project's own fixtures (§11.4.6). Honest boundary: vision verification confirms on-screen content — does NOT replace §11.4.5 quality analysis nor §11.4.108 runtime-signature verification; FLAG_SECURE surfaces use the §11.4.117 proxy oracle per §11.4.112.

Classification: universal (§11.4.17). Composes

§11.4.5/.27/.69/.107(10)/.112/.117/.153/.158/.159(J)/§1.1. Propagation gate CM-COVENANT-114-160-PROPAGATION + recommended gate CM-VISION-VERIFIED-RECORDING-BRIDGE + paired §1.1 mutation.

§11.4.161 — Rootless container runtime mandate (User mandate, 2026-06-21). Compact summary: every project MUST use Podman in rootless mode (or equivalent rootless container runtime) for ALL containerized workloads — Docker in rootful mode, sudo, or any escalation to root for container management is STRICTLY FORBIDDEN unless the target platform has no rootless option AND that constraint is documented per §11.4.112; the vasic-digital/containers Submodule (§11.4.76) MUST be used as the sole container orchestration layer — no ad-hoc docker/podman commands outside the Submodule's pkg/boot/pkg/compose/pkg/health primitives; if a missing capability forces a raw command, the containers Submodule MUST be extended upstream per §11.4.74 rather than worked around; all container-related integration tests MUST boot infra on-demand via the Submodule (the on-demand-infra invariant). Honest boundary:

rootless execution eliminates the container-to-root privilege-escalation vector — it does NOT replace §11.4.10 credentials-handling (credentials in containers still require leak audits) nor §12.3 container hygiene (memory limits/OOM policies/restart backoff still apply). Classification: universal (§11.4.17). Composes §11.4.76/.74/.10/.112/§12.3/.6. Propagation gate CM-COVENANT-114-161-PROPAGATION + recommended gate CM-ROOTLESS-CONTAINER-RUNTIME + paired §1.1 mutation.

§11.4.162 — OpenDesign UI design system mandate (User mandate, 2026-06-21). Compact summary: every project producing user-facing interfaces (web, desktop, mobile, TUI) MUST use OpenDesign (<https://github.com/nexu-io/open-design>) as the mandatory UI design-and-refinement system — NOT ad-hoc CSS or one-off design tools; install as a project dependency and use its design tokens/themes system for: (a) all color palette definitions supporting BOTH light and dark themes from project brand assets (§11.4.35), (b) typography scale and spacing, (c) component-level design tokens; if a desired UI pattern is not supported, extend OpenDesign upstream per §11.4.74 (extend-don't-reimplement); every UI component MUST ship light + dark theme variants; elements MUST NOT overlap, fonts MUST NOT collide, labels MUST NOT overlay labels — any layout regression is a §11.4.162 violation; all UI changes MUST be covered by the project's standard test types including visual regression tests (before/after screenshots with per-pixel or perceptual-diff PASS/FAIL). Honest boundary: OpenDesign governs design tokens and themes — it does NOT replace functional testing (§11.4.27), WCAG accessibility assertions (§11.4.107), nor the §11.4.48/§11.4.49 UI-driven and dual-approach testing methodology. Classification: universal (§11.4.17). Composes §11.4.74/.25/.27/.4(b)/.107/.48/.49/.35/.69/§1.1. Propagation gate CM-COVENANT-114-162-PROPAGATION + recommended gate CM-OPENDESIGN-UI-SYSTEM + paired §1.1 mutation.

§11.4.163 — Universal Media Validation & Verification Mandate (User mandate, 2026-06-21). Compact summary: every recorded artifact (video/audio/screenshots/asciicast/text) MUST pass a MEDIA VALIDATION pipeline before acceptance — OCR for video/screenshots (§11.4.117/.107(12)), transcription for audio, text parsing; compare extracted content against SPECIFY-phase expected patterns (§11.4.159(J)); self-validated analyzer with golden-good/golden-bad fixture pair (§11.4.107(10)); produce structured verdict (PASS/FAIL + evidence path + matched/unmatched patterns + pinpoint data on FAIL — which frame/line/timestamp, expected vs actual); post-recording AND real-time trigger (§11.4.159(M)); paired §1.1 mutation ensures golden-bad fixture produces FAIL. Honest boundary: confirms artifact content matches patterns — does NOT replace §11.4.5 quality, §11.4.107 liveness, nor §11.4.108 runtime-signature. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-163-PROPAGATION + recommended gate CM-MEDIA-VALIDATION-PIPELINE + paired §1.1 mutation.

§11.4.164 — Universal Constitution Auto-Propagation & Hook System (User mandate, 2026-06-21). Compact summary: every fetch+pull of the constitution submodule MUST trigger `constitution/scripts/post_update_hook.sh` (inherited §11.4.28) that: detects changed files (governance MDs + `scripts/hooks/skills/mcp/plugins`); registers new/modified skills via consumer's `scripts/register_skills.sh`; registers MCP servers via consumer's `scripts/register_mcp.sh`; installs hooks into `.git/hooks/`; validates scripts syntax per §11.4.67; emits summary log; failures logged per §11.4.6; tested by §1.1 mutation (add skill → hook detects it). Consumer MUST invoke after every pull. Honest boundary: installs components — does NOT guarantee runtime acceptance nor replace §11.4.32 sweep. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-164-PROPAGATION + recommended gate CM-CONSTITUTION-AUTO-PROPAGATION + paired §1.1 mutation.

§11.4.165 — Universal Independent Verification Agent Mandate (User mandate, 2026-06-21). Compact summary: every code/media/docs/config output MUST pass INDEPENDENT verifier (§11.4.70/.20, structurally separate from author, §11.4.92 PRECEDES never satisfies), iterating to GO per §11.4.134. CODE: §11.4.142 review + build+test + paired §1.1 mutations + §11.4.108 runtime-signature. MEDIA: §11.4.163 pipeline + genuine-content check (§11.4.159(I)) + format check (§11.4.159(B)). DOCS: exports current §11.4.65 + revision header §11.4.44. CONFIG: syntax + schema + §11.4.10 leak check. Structured findings per class; verifier self-validated by §1.1 mutation. Honest boundary: source/artifact integrity — does NOT replace §11.4.108 runtime-signature nor §11.4.40 retest. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-165-PROPAGATION + recommended gate CM-INDEPENDENT-VERIFICATION-AGENT + paired §1.1 mutation.

§11.4.166 — REPEALED (operator decision, 2026-06-22). The Universal Semgrep static-analysis mandate is REPEALED — Semgrep is NO LONGER a mandatory requirement for projects governed by this Constitution. Static analysis / security scanning remains *encouraged* and may be re-adopted per project, but is NOT mandated, NOT a release blocker, NOT cascaded. Removed with the repeal: `scripts/semgrep/`, `scripts/hooks/semgrep_precommit.sh`, the `submodules/semgrep` submodule, the MCP/pre-commit/PATH wiring, the `.docs_chain/contexts/semgrep_status.yaml` context, and the CM-COVENANT-114-166-PROPAGATION / CM-SEMGREP-WIRED gates. Anchor number 11.4.166 is RETIRED (not reused); repealed before full fleet cascade — consuming projects never received it as a binding mandate. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.166.

§11.4.167 — Big-work-item feature work-stream lifecycle (own CoW project copy + own branch/tags + per-feature builds, no-merge-until-approved, trunk-merged-into-every-stream) (User mandate, 2026-06-23). Every BIG work item (new feature OR drastic/large fix) MUST be developed as its own isolated feature

work-stream (ALL hold): (A) one big item = one feature work-stream in a sibling project copy (<project>_<slug>, slug lowercase §11.4.29; small changes stay on the trunk); (B) disk-feasible duplication MANDATORY (reflink/CoW clone on btrfs/XFS/ZFS/APFS, or git worktree + externalized build output) — a naive full deep copy / git clone per stream that would exhaust free space is FORBIDDEN, ~0-disk cost verified as captured evidence (§11.4.69/§11.4.6); per-stream out /, one shared compiler cache; (C) own feature/<slug> branch + own feature tags (<base>-feat-<slug>[-<iter>], greppable §11.4.151), NEVER merged to trunk until the operator MANUALLY approves after §11.4.40 GREEN + §11.4.41 merge-first + §11.4.113 ff-only; (D) trunk merged INTO every live stream frequently (≥ after every trunk tag / daily, git merge never rebase) or flagged stale; (E) feature branch+tag cascades to EVERY touched submodule at the same HEAD, ff-only-pushed, untouched submodules stay trunk-pinned; (F) single-builder FIFO build queue + finite-device exclusive test queue (§11.4.119 / §11.4.58 L1 / §12.7 / §12.8), GC idle streams' regenerable output (§11.4.77); (G) every change crosses the full review + impact-research + anti-bluff gauntlet (§11.4.142 / §11.4.125 / §11.4.134 / §11.4.145) — a stream is NOT a quality carve-out; (H) resumable, data-safe lifecycle via a stream registry (§11.4.116 / §11.4.147), atomic-on-success clone, §9.2 pre-op backup on retire (unmerged work blocks auto-deletion), §12 / §12.6 bounds. Honest boundary (§11.4.6): guarantees isolation / disk-feasibility / greppable identity / controlled trunk ↔ feature merge — NOT any stream's correctness (still §11.4.108 + §11.4.40). Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-167-PROPAGATION (literal 11.4.167) + recommended gates CM-FEATURE-WORKSTREAM-COW-CLONE / CM-FEATURE-WORKSTREAM-NO-MERGE-UNTIL-APPROVED / CM-FEATURE-WORKSTREAM-TRUNK-SYNC-CADENCE / CM-FEATURE-WORKSTREAM-SUBMODULE-CASCADE + paired §1.1 mutations. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.167. Non-compliance is a release blocker. No escape hatch — no --full-clone-per-feature, --deep-copy-stream, --skip-cow-clone, --merge-feature-without-approval, --no-trunk-sync, --skip-submodule-cascade, --unbounded-parallel-builds, --shared-device-no-lock, --retire-unmerged-stream, --feature-stream-quality-carveout flag.

§11.4.168 — Exported-document independent content + textual + full-visual validation mandate (User mandate, 2026-06-23). Forensic FACT (2026-06-23): a “green” Mermaid export shipped RAW gantt-diagram source (gantt/title/dateFormat/section/:done,) as plain body text into user-facing PDFs; validation grepped the HTML for class="mermaid" but NEVER opened the exported PDF — the §11.4.108 SOURCE → ARTIFACT gap at the doc-export layer. Every generated/exported document (HTML/PDF/DOCX, §11.4.65 scope + §11.4.153 four-format set) MUST pass INDEPENDENT validation by a reviewer structurally separate from the generator (§11.4.70 / §11.4.20 / §11.4.142; NEVER generator self-check, §11.4.92 precedes never satisfies), run on BOTH source Markdown AND every exported artifact, across THREE layers ALL of which MUST hold: (1) CONTENT — nothing

dropped/truncated/reordered/garbled, every section/table/list/figure/link present + complete (the §11.4.5/§11.4.69 quality bar, not presence-only); (2) TEXTUAL — NO raw markup / diagram source (mermaid/gantt/graph/flowchart/sequenceDiagram/.../:done, task lines) / unrendered code-fence body leaking as readable text, detected by reading the exported artifact as text (pdftotext / rendered-DOM / DOCX-text); (3) FULL VISUAL — diagrams render as IMAGES not source, layout intact (no overlap/clip/cut-off/collapsed/blank), verified by RENDERING (pdftotext + pdfimages image-enumeration + pdftoppm → OCR §11.4.117 / §11.4.107(12) confidence + ROI). (4) Anti-bluff — validator self-validated golden-good/golden-bad (§11.4.107(10): a golden-bad seeded with raw gantt body text + a truncated table MUST FAIL), iterates to a clean GO per §11.4.134. Honest boundary (§11.4.6): confirms the exported artifact a reader opens is faithful + readable + visually-rendered — NOT that source content is itself correct, and does NOT replace the §11.4.65 mtime-parity freshness gate (§11.4.168 is the READABILITY/FIDELITY layer the mtime gate cannot see). Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-168-PROPAGATION (literal 11.4.168) + recommended gate CM-EXPORTED-DOC-VISUALLY-VALIDATED + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.168. Non-compliance is a release blocker. No escape hatch — no --skip-visual-validation, --raw-source-in-pdf-ok, --textual-check-suffices, --self-validate-export, --diagram-source-as-text-ok flag.

§11.4.169 — Mandatory comprehensive test-type coverage with anti-bluff captured evidence (User mandate, 2026-06-25). Verbatim (2026-06-25):

“MANDATORY test types EACH project MUST have with coverage aiming to 100%... unit, integration, e2e, full automation, challenges, helix qa, ddos, security, stress and chaos, concurrency (atomicity), race conditions or deadlocks, memory, bechmarking.” Every project MUST be covered by this CLOSED ENUMERATED test-type set, each driven to as-close-to-100% as the domain permits, every PASS backed by rock-solid captured PHYSICAL evidence (§11.4.5 / §11.4.69 / §11.4.107), zero false results / zero bluff (§11.4 / §11.4.1 / §11.4.6) — the strict enumerated expansion of §11.4.27. The set (each REQUIRED where the domain warrants; only an honest §11.4.3 SKIP-with-reason for genuinely-absent topology/hardware/credential is permitted, NEVER a silent gap): (1) Unit (ONLY layer mocks/stubs/fakes permitted §11.4.27(A)); (2) Integration (real components, infra on-demand via the containers submodule §11.4.76 / §11.4.161 rootless); (3) e2e; (4) Full-automation (autonomous, re-runnable, no manual intervention §11.4.25 / §11.4.52 / §11.4.98, deterministic §11.4.50); (5) Challenges (vasic-digital/challenges anti-bluff banks); (6) HelixQA (helix_qa submodule test banks + fully-autonomous QA sessions); (7) DDoS / load-flood; (8) Security (authn/authz, injection/taint, secret-leak §11.4.10, crypto, CVEs, default-deny; security submodule); (9) Stress + Chaos (§11.4.85); (10) Concurrency / atomicity; (11) Race-condition / deadlock (race detector + lock-order); (12) Memory (leak census / soak,

peak-RSS ceilings); (13) Benchmarking / performance (p50/p95/p99 vs baseline). Four-layer §11.4.4(b) enforcement of the suite + per-type paired §1.1 mutation; the coverage ledger (§11.4.25 / §11.4.52) classifies every feature × type × evidence-state — a missing required type or an OPERATOR_ATTENDED_ONLY row is a release blocker until promoted or honestly SKIP-justified. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-169-PROPAGATION (literal 11.4.169) + recommended gate CM-MANDATORY-TEST-TYPES-COVERED + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.169. Non-compliance is a release blocker. No escape hatch — no --skip-test-type, --partial-coverage-ok, --evidence-optional, --challenges-not-applicable, --skip-helixqa, --no-chaos, --no-ddos, --skip-memory-test, --skip-race-detector, --bench-optional flag.

§11.4.170 — Device-independent host-side rendered-UI visual-proof mandate (User mandate, 2026-06-25). Every change to any user-facing UI surface (screen/component/view/widget/layout/theme on any platform — Android-Compose, iOS-SwiftUI/UIKit, web, desktop, TUI) MUST be proven by DEVICE-INDEPENDENT host-side RENDERED PIXELS before it may be claimed correct: the actual component rendered to a PNG ON THE HOST (NO device, NO emulator, NO running app; Compose → Roborazzi/Paparazzi, web → Playwright/Storybook, SwiftUI → snapshot-testing, equiv per stack) for EVERY screen × state × {light,dark} theme, dual-validated by (i) golden image-diff AND (ii) an OCR/vision oracle reading rendered text + labels + control bounds (NO overlap / label-over-label / clipping / off-screen / collapsed-or-giant-unbounded widget — the §11.4.162 layout invariants PROVEN on real pixels). Forbidden bluff (forensic FACT 2026-06-25): UI “validated” by VALUE/token/property-equality unit tests (a hex colour string, an sp/dp integer, a design-token field == a constant) that NEVER RENDERED A PIXEL stayed GREEN while the operator opened a broken giant-button screen — value/token/property-equality UI tests are FORBIDDEN as the PROOF a UI is correct (may SUPPLEMENT, NEVER SUBSTITUTE the rendered-pixel proof); a green UI suite with no rendered-pixel artefact for the changed surface is a §11.4 PASS-bluff. ALL hold: rendered-pixel proof per screen×state×theme, run on EVERY commit with zero hardware; dual oracle both required, each PASS citing the captured PNG + diff (§11.4.5/§11.4.69); a value-only UI claim is a §11.4.4 finding → render → fix → re-render GREEN; self-validated analyzer golden-good/golden-bad (§11.4.107(10)), thresholds on the project’s OWN fixtures (§11.4.6); “device offline” is NEVER a valid skip (host-render IS the device-independent path); genuine no-host-render-harness ⇒ honest §11.4.3 SKIP + tracked migration item, never a value-only fallback. Honest boundary (§11.4.6): host-render proves the component renders correctly device-independently on every commit — it COMPLEMENTS not replaces §11.4.153/§11.4.158/§11.4.159/§11.4.160 live-device recording + §11.4.162 OpenDesign tokens + §11.4.168 exported-doc visual validation. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-170-PROPAGATION (literal 11.4.170) + recommended gate CM-HOST-RENDERED-UI-VISUAL-PROOF +

paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.170. Non-compliance is a release blocker. No escape hatch — no --value-assertion-proves-ui, --skip-rendered-pixel-proof, --token-equality-suffices, --device-offline-skip-visual, --single-theme-only, --unvalidated-image-diff-OK, --no-ocr-layout-oracle flag.

§11.4.171 — Mandatory comprehensive human-readable workable-item descriptions (User mandate, 2026-06-29). Every workable item (ATM-NNN ticket) in the project's workable-items database, Issues.md, Fixed.md, Issues_Summary.md, Fixed_Summary.md, and ALL derived documents MUST carry a comprehensive human-readable description in plain, non-technical language suitable for non-developers (project managers, stakeholders, team leads, business partners) — ≥ 5-7 sentences covering (a) WHAT it is (feature/fix/task in simple terms); (b) WHY it matters (user/business value); (c) HOW it works (mechanism, no code references); (d) WHO benefits; (e) the expected measurable outcome. Jargon, unexpanded acronyms, code references, and implementation details are FORBIDDEN as the primary explanation (technical detail may appear as supplementary context AFTER the plain-language explanation). An item without a human-readable description is a §11.4 documentation-layer violation. The description column in the SQLite DB (docs/workable_items.db) is the SINGLE SOURCE OF TRUTH; all derived documents carry the SAME description text. Gate CM-HUMAN-READABLE-DESCRIPTION (when implemented) asserts every item's description exists, non-empty, ≥ 50 characters; paired §1.1 mutation strips a description → gate FAILs. Classification: universal (§11.4.17). Composes §11.4.15 / §11.4.16 / §11.4.19 / §11.4.91 / §11.4.93 / §11.4.148 / §11.4.65. Propagation gate CM-COVENANT-114-171-PROPAGATION (literal 11.4.171) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.171. Non-compliance is a release blocker.

§11.4.172 — Mandatory production-readiness planning with realistic timeline projection (User mandate, 2026-06-29). Every project under this Constitution MUST maintain a living production-readiness planning document that: (a) provides REALISTIC timeline projections for all milestones (MVP, beta, production-ready) based on actual development velocity data (items completed per week, reopening rate, blocking dependencies); (b) identifies ALL danger zones, weaknesses, and risk areas with mitigation strategies; (c) accounts for HW-delivery timelines, licensing delays, and external-dependency risks; (d) defines the critical path and which items can slip without affecting MVP; (e) includes workstation/build-capacity analysis + optimization recommendations; (f) is updated at least monthly or whenever the item count changes by ≥ 10%; (g) is cross-referenced against the workable-items database for accuracy. The document lives at docs/research/production_planning_<YYYYMMDD>/ANALYSIS.md with HTML+PDF exports per §11.4.65. Honest boundary (§11.4.6): projections are velocity-based estimates — they guarantee sound methodology, NOT that specific dates will be met. Classification: universal (§11.4.17). Composes §11.4.6 / §11.4.40 /

§11.4.42 / §11.4.93 / §11.4.108. Propagation gate CM-COVENANT-114-172-PROPAGATION (literal 11.4.172) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.172. Non-compliance is a release blocker.

§11.4.173 — Containerized + distributed build mandate (User mandate, 2026-06-29). EVERY build of EVERY component (source compile, artifact/package/installer/container-image production, codegen, asset render — for ANY language/platform: Go, Android/Gradle, desktop, web, native, firmware) MUST run INSIDE a specialized build container provisioned via the `digital.vasic.containers` submodule (§11.4.76) — NEVER on the bare host. The build containers MUST be DISTRIBUTED to the designated remote build host(s) via the SAME containers-submodule distribution mechanism (remote compose over SSH §11.4.76, rootless §11.4.161) so the build EXECUTES on the remote host (offloading the developer/main host); once complete the produced artifacts MUST be brought BACK to the originating host (scp/rsync/volume copy) for use/flashing/distribution. Building outside a container, or on the bare host, is FORBIDDEN — a release blocker (the “works on my machine” / unreproducible-build class). Build-host target + build-container definitions are config-injected (§11.4.28, never hardcoded in the submodule); a missing build-container capability is added by EXTENDING the containers submodule upstream (§11.4.74), never an ad-hoc host build. Honest boundary (§11.4.6): the containerized + distributed build guarantees reproducibility + host-isolation + capacity offload — it does NOT replace the §11.4.40 full-suite retest, §11.4.108 artifact → runtime verification, or §11.4.38 installable-asset evidence on the brought-back artifact. Classification: universal (§11.4.17). Composes §11.4.76 / §11.4.161 / §11.4.74 / §11.4.28 / §11.4.24 / §11.4.82 / §11.4.121 / §11.4.38 / §11.4.108 / §12.6. Propagation gate CM-COVENANT-114-173-PROPAGATION (literal 11.4.173) + recommended gate CM-CONTAINERIZED-DISTRIBUTED-BUILD + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.173. Non-compliance is a release blocker. No escape hatch — no `--build-on-host`, `--skip-container-build`, `--local-build-ok`, `--no-distributed-build`, `--bare-host-build` flag.

§11.4.174 — Shared-host process-ownership verification mandate (User mandate, 2026-06-29). Verbatim operator mandate: “Make sure that processes we are checking are always OURS! ... on every host there may be other work being done!” On ANY host where concurrent work from other projects / users / agents may run (the normal case — a shared developer workstation), EVERY inspection of, conclusion about, or action on a process, build, daemon, port, lock-file, or system-state signal MUST FIRST positively VERIFY the target is OURS before drawing a diagnosis or acting. A diagnosis based on another project’s process is a §11.4.6 guessing error; KILLING / signalling / restarting / altering a process that is NOT positively ours is a §12 + §11.4.133 safety violation (operator-workload destruction) — forbidden. Forensic FACT (2026-06-29): a 200%-CPU java + gradlew on the shared host was nearly mis-attributed to our build but was a

separate catalogizer project's gradle test (cwd=Projects/catalogizer); a naive pgrep java/gradle would have matched it as "ours." ALL hold: (a) positive-ownership discriminator REQUIRED — attribute to us ONLY via a strong signal: cwd resolving inside THIS project's tree, argv naming our scripts/artifacts/repo path, a PID we launched + recorded (§11.4.147 registry), or a lock/port/socket path under our tree; a loose name match (pgrep java/gradle/node/python/podman/zig) is NEVER sufficient; (b) NEVER kill/signal/restart a process not positively ours — other projects' daemons/pods are off-limits; the §12.8 pkill -f GradleDaemon remediation MUST be scoped to OUR daemons (cwd/argv under our tree), never host-wide; (c) contention → surface, don't break (§11.4.66 / §11.4.101), keep progressing non-blocked work; (d) every pgrep/ps filter MUST exclude the checking command's OWN argv AND other-project matches, preferring our absolute repo path; (e) applies to ALL state (ports, lock files, caches, pods, tmux sessions, temp files), not just processes. Honest boundary (§11.4.6): ownership-verified inspection prevents mis-attribution + cross-project damage — it does not itself prove our process is healthy. Classification: universal (§11.4.17). Composes §11.4.6 / §11.4.101 / §11.4.66 / §11.4.147 / §12 / §12.8 / §11.4.133. Propagation gate CM-COVENANT-114-174-PROPAGATION (literal 11.4.174) + recommended gate CM-PROCESS-OWNERSHIP-VERIFIED + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.174. Non-compliance is a release blocker. No escape hatch — no --assume-process-ours, --loose-pgrep-ok, --kill-any-matching-process, --skip-ownership-check flag.

§11.4.176 — Conflict-free multi-track work-division + exactly-once claim registry + capability-aware deadlock-proof device-lock (User mandate, 2026-07-02). When a project runs multiple concurrent work-STREAM tracks (§11.4.167 / §11.4.58 / §11.4.103) sharing (a) a common pool of workable items and (b) a scarce pool of exclusive hardware/resources, TWO cooperating arbitration layers are MANDATORY, both conflict-free BY CONSTRUCTION: **(A) work-item layer — exactly-once claim registry** — every workable item OR logical GROUP of items sharing files/subsystem/feature/schema is assigned to AT MOST ONE track at a time via a lock-arbitrated append-only claim registry (extends §11.4.147 + §11.4.116); claim-before-touch, a second claim on an already-claimed item is DENIED; logical groups are claimed WHOLE (disjoint file-scope §11.4.58 L3 + per-track worktree/CoW §11.4.167 = tracks physically cannot touch the same bytes); no codebase loss (§9.2 backups / §11.4.113 ff-only-merge-onto-latest-main NEVER force-push / §11.4.84 quiescence / §11.4.41 merge-first); results ALWAYS produced (abandoned claim TTL-expires back to the pool). **(B) hardware layer — capability-aware deadlock-proof device-lock** — when N tracks outnumber M exclusive devices, exactly ONE track owns each device (§11.4.119); a track leases ONE OR MORE devices selected by required CAPABILITY (not merely identity); the arbiter is DEADLOCK-FREE by breaking all four Coffman conditions (hold-and-wait eliminated — multi-device acquire is ALL-OR-NOTHING in one critical

section; circular-wait eliminated — acquire is NON-BLOCKING, wins the whole set or returns EBUSY; no-preemption relaxed — lease TTL + heartbeat + reap so a crashed holder's lease auto-expires §11.4.147; the only lock held is a short bounded flock over the atomic check-and-claim); registry = append-only JSONL + atomically-rewritten snapshot (§11.4.116). **(C) universal decoupling + evidence-based resource-tuning** — both layers universal/decoupled/project-not-aware (§11.4.17/§11.4.28/§11.4.35); resource-tuning decisions (stream count, per-stream memory §12.6, lease breadth, queue depth §11.4.167) MUST be EVIDENCE-BASED (§11.4.24 / §11.4.5), never guessed (§11.4.6). Classification: universal (§11.4.17). Composes §11.4.6/§11.4.17/§11.4.24/§11.4.28/§11.4.35/§11.4.41/§11.4.58/§11.4.84/§11.4.85/§11.4.103/§11.4.113/§11.4.116/§11.4.119/§11.4.147/§11.4.167/§12.6/§9.2. Propagation gate CM-COVENANT-114-176-PROPAGATION (literal 11.4.176) + recommended gates CM-WORK-DIVISION-EXCLUSIVE-CLAIM / CM-NO-CROSS-TRACK-SCOPE-OVERLAP / CM-LOGICAL-GROUP-COHESION / CM-DEVICE-LOCK-DEADLOCK-FREE + paired §1.1 mutations. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.176. Non-compliance is a release blocker. No escape hatch — no --allow-cross-track-overlap, --skip-claim-registry, --split-logical-group, --blocking-device-acquire, --allow-partial-device-hold, --tune-resources-by-guess flag.

§11.4.177 — Developer-tooling project-decoupling + invocation-directory operation (User mandate, 2026-07-04). No project-specific script / hook / alias / binary may be wired — copied, symlinked, PATH-exported, or aliased — into a global or shared developer-tooling PATH; shared tooling MUST be project-agnostic and operate on the INVOCATION directory (cwd, or an explicit path/config argument), NEVER a hardcoded project path. A project-specific helper lives in + runs from its own repo; a genuinely-reusable helper is promoted to the shared toolkit ONLY after being stripped of every project-specific assumption (paths, device serials, package names, region endpoints) and taking its target from the invocation context. A project script reachable on the global/shared PATH is a decoupling violation of equal severity to importing project-specific code into a shared submodule (§11.4.28). Forensic FACT (2026-07-04): a project cwd-hook symlinked into the global Claude-Toolkit PATH re-coupled every alias to one hardcoded checkout. Honest boundary (§11.4.6): decoupling guarantees the tool is project-agnostic — NOT that it is correct (that still needs the tool's own tests). Classification: universal (§11.4.17). Composes §11.4.28 / §11.4.29 / §11.4.35 / §11.4.11 / §11.4.6. Propagation gate CM-COVENANT-114-177-PROPAGATION (literal 11.4.177) + recommended gate CM-TOOLING-PROJECT-DECOUPLED + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.177. Non-compliance is a release blocker. No escape hatch — no --global-symlink-ok, --hardcode-project-path, --wire-into-shared-path flag.

§11.4.178 — Track-qualified identity for parallel work streams (session-name collision ban) (User mandate, 2026-07-04). Parallel work streams that share hardware, a git backing store, a tmux/session namespace, a lock namespace, or

any single-owner resource MUST be addressed by a TRACK-QUALIFIED identity (<project>__<track>__<role>), NEVER a bare directory basename. Multiple checkouts/clones/worktrees sharing basename atmosphere collapse into one session/lock/log key and cross-wire tmux targets, lock files, device claims, and log directories (the observed collision, 2026-07-04). Every registry row, session name, lock path, log directory, and device claim MUST carry the track-qualified key; a bare-basename identity for a resource with ≥ 2 concurrent claimants is a violation. Honest boundary (§11.4.6): a unique identity prevents cross-wiring — it does not itself arbitrate ownership (that is §11.4.119 + §11.4.176, the multi-track work-division + exactly-once-claim + device-lock arbitration anchor). Classification: universal (§11.4.17). Composes §11.4.111 / §11.4.119 / §11.4.176 / §11.4.29 / §11.4.58. Propagation gate CM-COVENANT-114-178-PROPAGATION (literal 11.4.178) + recommended gate CM-TRACK-QUALIFIED-IDENTITY + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.178. Non-compliance is a release blocker. No escape hatch — no --bare-basename-ok, --skip-track-qualifier, --shared-session-name flag.

§11.4.179 — Corruption-isolated parallel git streams (own-.git independent clones, not shared-common-dir worktrees) (User mandate, 2026-07-04). When corruption-isolation is a requirement, parallel git work streams MUST each be an isolated repository with its OWN .git (own object store, own index, own lock namespace), NOT git worktree checkouts sharing one common .git. A shared common-dir is a single point of failure: one stale index.lock / .commit_all.lock / .push_all.lock freezes EVERY stream at once (the observed 9-h all-track freeze, 2026-07-04), and one corrupted object store loses every stream. Each stream's git rev-parse --git-common-dir MUST resolve INSIDE the stream's own tree; a common-dir resolving to a shared parent is a violation. Where CoW/reflink disk-feasibility is the constraint, use the §11.4.167 reflink-clone-with-own-.git path (a btrfs/XFS/ZFS/APFS reflink shares extents on disk while keeping a per-stream .git, so isolation and disk-feasibility both hold). Honest boundary (§11.4.6): own-.git isolation contains lock/corruption blast radius — it does NOT remove per-repo stale-lock reaping (§11.4.180) nor ff-only no-force integration (§11.4.113). Classification: universal (§11.4.17). Composes §11.4.167 / §11.4.119 / §11.4.176 / §11.4.58 / §9.2. Propagation gate CM-COVENANT-114-179-PROPAGATION (literal 11.4.179) + recommended gate CM-GIT-STREAMS-OWN-COMMON-DIR + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.179. Non-compliance is a release blocker. No escape hatch — no --shared-common-git, --worktree-when-isolation-required, --skip-git-isolation flag.

§11.4.180 — Commit/push (single-writer) wrappers MUST auto-reap provably-stale locks before acquiring (User mandate, 2026-07-04). Every commit / push / single-writer wrapper MUST, before acquiring its own lock, auto-reap any PROVABLY-stale git lock — one whose recorded holder PID is DEAD (kill -0

fails), OR (no PID recorded) older than a defined threshold AND with no live wrapper/git process holding that repo. The reap covers the wrapper's own lock plus the git-internal existence-based locks (`index.lock`, `HEAD.lock`, `refs/**/*`.lock) that git does NOT auto-release on holder death. It MUST NEVER remove a lock whose holder is ALIVE (removing a live-held lock corrupts a concurrent writer — §9.2). Each reap decision (REAPED / KEPT + reason) is logged as captured evidence (§11.4.6 — liveness PROVEN via `kill -0`, never assumed). A wrapper that blocks indefinitely on a dead-holder lock (the observed 9-h freeze, 2026-07-04) is a violation. Honest boundary (§11.4.6): reaping clears dead-holder deadlock — it does NOT substitute for the own-.git isolation of §11.4.179 (which shrinks the blast radius from all-streams to one). Classification: universal (§11.4.17). Composes §11.4.84 / §11.4.88 / §11.4.6 / §9.2 / §11.4.116 / §11.4.179. Propagation gate CM-COVENANT-114-180-PROPAGATION (literal 11.4.180) + recommended gate CM-WRAPPER-STALE-LOCK-AUTO-REAP + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.180. Non-compliance is a release blocker. No escape hatch — no `--skip-stale-lock-reap`, `--force-remove-live-lock`, `--block-on-dead-holder` flag.

§11.4.181 — Consistent controlled feature-branch naming: one feature/logic-group ⇒ exactly ONE canonical branch name across the main repo + all owned submodules (User mandate, 2026-07-05). Verbatim (2026-07-05): “System MUST pay attention not to create multiple feat. branches with different names for one same feature or group of workable items / features! Naming MUST BE consistent and carefully controlled!” One feature / logical group of workable items MUST map to EXACTLY ONE canonical branch name, used IDENTICALLY on the main repo AND every touched owned submodule — NEVER two differently-named branches for one feature/group, NEVER per-repo naming drift. ALL hold: (1) **single canonical name per group** minted ONCE per the §11.4.167 D scheme (feature/<slug> branch + <prefix>-<base>-feat-<slug> tag, <slug> lowercase snake/kebab §11.4.29), used identically across the main repo + every branched submodule (an untouched submodule stays on its base branch, NOT force-branched); (2) **single source of truth** — the canonical name is recorded ONCE in the §11.4.176 exactly-once claim registry / the §11.4.93 workable-items `logic_group` → branch mapping and LOOKED UP, never re-invented (creating a branch whose name ≠ the group's registered canonical name, OR minting a 2nd name for a group that already has one, is a §11.4.6 no-guessing violation); (3) **mechanical enforcement** (§11.4.75) — gate CM-BRANCH-NAME-CONSISTENCY FAILs on an unregistered feature-branch name / two divergent names for one group / a submodule branch diverging from the group's main-repo canonical name, paired §1.1 mutation (register one group under two divergent names → gate FAILs); (4) **risk-free reconciliation** (§9.2 / §11.4.113 / §11.4.84) — a mis-named/duplicate branch is reconciled by MERGING its commits onto the canonical branch (union, no `-s ours/rebase/reset` that drops commits, `git log <canonical>..<misnamed>` empty after merge, no conflict markers), ff-only

push to all upstreams (NEVER force-push), retiring the mis-named branch only after the merge is confirmed on all mirrors (delete-before-merge forbidden). Honest boundary (§11.4.6): “probably the same feature” is NOT a determination — bind the branch to a group via the registry, never guess. Classification: universal (§11.4.17). Composes §11.4.167 / §11.4.176 / §11.4.178 / §11.4.179 / §11.4.28 / §11.4.29 / §11.4.93 / §11.4.113 / §11.4.84 / §11.4.6 / §11.4.75 / §9.2. Propagation gate CM-COVENANT-114-181-PROPAGATION (literal 11.4.181) + recommended gate CM-BRANCH-NAME-CONSISTENCY + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.181. Non-compliance is a release blocker. No escape hatch — no --allow-divergent-branch-names, --skip-branch-registry, --rename-branch-freely, --per-repo-branch-name-OK, --two-names-one-feature-OK flag.

§11.4.182 — Track+branch work-stream identity label (T<N>/<branch> - <alias>) on every agent/subagent label + every operator-facing work-stream reference (User mandate, 2026-07-05). Operator mandate (2026-07-05): “EVERY agent/subagent/work-stream label AND every operator-facing report reference to a work stream MUST carry a (T<N>/<branch> - <alias>) prefix identifying the track + git branch + the claude-toolkit alias in use — e.g. (T1/main - claude1) ATM-312 MPV drm_prime investigation, (T2/feat/mistiq-vader - deepseek) SPK-XXX ...” In a multi-track layout (parallel checkouts each on a distinct branch and potentially driven by a distinct claude-toolkit alias — e.g. /mnt/track1..4) a bare work-stream label does not say WHICH track/branch/alias produced it, so cross-track dispatches, status reports, and progress lines collide and mislead. ALL hold: (1) **universal label form** — every agent / subagent / work-stream label AND every operator-facing reference (agent-dispatch descriptions, status reports, resume/handoff docs, progress lines) MUST START with a (T<N>/<branch> - <alias>) prefix; (2) **deterministic derivation, never guessed** (§11.4.6) — <N> from the checkout path (/mnt/track<N>/...; honest ? when not under a numbered track, never a guessed number), <branch> from git rev-parse --abbrev-ref HEAD, <alias> from the CLAUDE_CONFIG_DIR basename .claude-<alias> → <alias> (honest ? when unset/non-matching, never a guessed name) (reference labeler scripts/multitrack/track_branch_label.sh); (3) **mechanical enforcement** (§11.4.75 / §11.4.109-class) — a PreToolUse guard hook (constitution/scripts/hooks/guard-track-branch-label.sh, inherited BY REFERENCE per §11.4.177, never copied) BLOCKS any Agent/Task/TaskCreate dispatch whose description (or subagent label) does not START with $\wedge(T[0-9]^+/\wedge^+ - \wedge^+)\wedge$ (branch AND alias both required, both non-empty), non-agent tools untouched; pre-build gate CM-TRACK-BRANCH-LABEL asserts labeler + hook exist/executable/parseable, hook wired in runtime settings, convention doc exists; (4) **honest boundary** (§11.4.6) — a ? in the track field OR the alias field is a valid honest label (unknown track / no git repo / no matching CLAUDE_CONFIG_DIR), never a bluff; the enforced format surfaces (blocks) an off-track dispatch or a dispatch missing its alias

segment rather than silently mislabeling it. Classification: universal (§11.4.17). Composes §11.4.178 (the label-prefix instantiation applied to EVERY agent/subagent label + operator-facing reference) / §11.4.29 / §11.4.75 / §11.4.109 / §11.4.177 / §11.4.6 / §11.4.58 / §11.4.103 / §11.4.119 / §11.4.116 / §11.4.147. Propagation gate CM-COVENANT-114-182-PROPAGATION (literal 11.4.182) + recommended gate CM-TRACK-BRANCH-LABEL + paired §1.1 mutation. **TRUNK RULE (operator mandate, 2026-07-28):** work performed on the canonical trunk (main/master) is **ALWAYS Track 1** — the label reads (T1/main - ...), NEVER (T?/main - ...), regardless of checkout path. Precedence is deliberate: the BRANCH decides first and the checkout path decides only for non-trunk branches, so a trunk checkout living outside /mnt/track<N> — or inside a differently-numbered one — is still Track 1; the identity that matters is which STREAM the work belongs to, not which directory it happens to sit in. Rationale: §11.4.182 labels are how concurrent streams are told apart, so an unnumbered T? on trunk work makes it indistinguishable from a genuinely-unknown track in every dispatch, report and progress line — and the trunk is a well-known, singular stream that is never unknown. Reference implementation: scripts/multitrack/track_branch_label.sh derives the branch BEFORE the track number and maps main|master -> 1. Labels MUST be re-derived from that script at dispatch time and kept in real-time sync with the live branch/alias/model/effort — a remembered or stale label is itself a §11.4.182 violation (§11.4.6: live identity is not knowable from memory). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.182. Non-compliance is a release blocker. No escape hatch — no --skip-track-label, --unlabeled-agent-OK, --bare-work-stream-reference, --no-track-branch-prefix, --label-from-memory flag.

§11.4.183 — Maximal multi-agent utilization per work-stream + full-constitution-application + zero-bluff mandate (User mandate, 2026-07-07). Compact summary: every track / work-stream MUST maximize multi-agent working approaches — subagent-driven development (§11.4.20/§11.4.70), independent review agents (§11.4.125/§11.4.142/§11.4.165/§11.4.134), parallel background streams with auto-backfill (§11.4.58/§11.4.103), and every other applicable agent form — so every track produces the MOST completed work at the HIGHEST quality in the LEAST time; every track MUST apply the ENTIRE constitution (every rule, mandatory constraint, guide, and derived technology — NOTHING ignored, skipped, avoided, or deemed less-relevant); there MUST NEVER be false / faulty / untested / unverified results or bluff of any kind anywhere (§11.4 anti-bluff covenant, captured evidence §11.4.5/§11.4.69/§11.4.107); honest boundary §11.4.6 — “maximize agents” is bounded by §12.6 60% memory + §11.4.58 parallel-agent cap + genuine parallelizability + §11.4.119 single-resource-owner, spawning contending/redundant agents is NOT maximization, the bar is maximum USEFUL parallel throughput not agent count. Classification: universal (§11.4.17). Composes §11.4.20/.58/.70/.92/.94/.97/.103/.119/.125/.126/.134/.142/.165/

§12.6. Propagation gate CM-COVENANT-114-183-PROPAGATION (literal 11.4.183) + recommended gate CM-MAX-AGENTS-PER-TRACK + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.183. Non-compliance is a release blocker. No escape hatch — no --single-agent-OK, --skip-review-agents, --partial-constitution-application, --serialise-parallelisable-work flag. **§11.4.184 — Mandatory SonarQube static-analysis CLI + local tooling installed and PATH-discoverable (User mandate, 2026-07-06).** Compact summary: every project/host that runs static analysis MUST have the SonarQube scanner CLI (sonar-scanner) installed AND durably PATH-discoverable via .bashrc/.zshrc (exactly like the Firebase/GitHub/GitLab CLIs), PLUS the shared constitution/scripts/sonarqube/tooling (sonarqube_install_check.sh/sonarqube_lib.sh/sonarqube_run_scan.sh/sonarqube_container.sh+compose/) consumed BY REFERENCE (§11.4.28/§11.4.177/§11.4.80-style, NEVER copied); proven GREEN by sonarqube_install_check.sh exit 0 + command -v sonar-scanner resolving (anti-bluff §11.4.6, never assumed); the local SonarQube server runs via rootless podman (§11.4.161), never rootful docker/sudo; DISTINCT from the repealed §11.4.166 Semgrep mandate — mandates TOOLING availability, not a scan on every build. Classification: universal (§11.4.17). Composes §11.4.28/.36/.74/.75/.109/.161/.166/.177/.6. Propagation gate CM-COVENANT-114-184-PROPAGATION (literal 11.4.184) + recommended gate CM-SONARQUBE-CLI-INSTALLED + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.184. Non-compliance is a release blocker. No escape hatch — no --skip-sonarqube-install, --sonar-cli-optional, --ad-hoc-sonar-path, --rootful-sonar-OK flag. **§11.4.185 — Manual QA-team testing as the mandatory FINAL confirmation of done (User mandate, 2026-07-07).** Compact summary: no set of features / scope of work / release may be considered fully completed until it has received MANUAL TESTING confirmation by the project's QA team as the FINAL confirmation step — EVERYWHERE and ALWAYS, never forgotten. Every automated gate (§11.4 anti-bluff, §11.4.5/§11.4.69 captured evidence, §11.4.52 autonomous validation, §11.4.40 full-suite retest, §11.4.108 four-layer verification, §11.4.132 risk-ordered validation) is NECESSARY but NOT SUFFICIENT — manual QA is the FINAL human sufficiency gate. Composition with §11.4.52 ADDS, never weakens: a feature MUST have BOTH the autonomous path (primary, always-on) AND the manual-QA final confirmation — automation-only with no QA sign-off is a §11.4.185 violation; manual-QA-only with no autonomous path is still a §11.4.52 violation. Workflow: autonomous-validation-GREEN → build → flash/deploy → autonomous on-device validation → HAND OFF to the QA team → ONLY after QA-team manual confirmation is the scope “fully completed”; the §11.4.126 autonomous-loop release-tag terminal condition now REQUIRES this confirmation before the tag is cut. Honest boundary (§11.4.6): the agent MUST hand off + WAIT for QA-team confirmation, never self-certify/simulate/infer it; while waiting the agent keeps progressing every OTHER non-blocked item in parallel (§11.4.87/.94/.97/.101) —

the wait parks that one scope, not the loop. Classification: universal (§11.4.17). Composes §11.4/.5/.6/.40/.52/.69/.108/.126/.129/.130/.132. Propagation gate CM-COVENANT-114-185-PROPAGATION (literal 11.4.185) + recommended gate CM-MANUAL-QA-FINAL-CONFIRMATION + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.185. Non-compliance is a release blocker. No escape hatch — no --automation-only-is-done, --skip-manual-qa, --self-certify-qa-confirmation, --infer-qa-from-automation flag.

§12.12 — Process/thread-limit (RLIMIT_NPROC) awareness for parallel subagent/multi-process work (User mandate, 2026-07-07). Compact summary: heavy parallel subagent/multi-process work is bounded by the OS per-user process/thread limit `ulimit -u (RLIMIT_NPROC)`, which on Linux counts ALL threads AND processes owned by the real UID, not per-process; ambient tooling (MCP servers, language runtimes) can consume most of the budget before project work starts (forensic FACT: a `ulimit -u 4096` host sat at ~3900-4005 threads from the ambient MCP fleet alone, leaving under 150 threads of headroom; a 4th Go-heavy subagent hit EAGAIN/failed to create new OS thread). Thread-limit exhaustion is a §12 host-safety event that YIELDS UNCONDITIONALLY, same standing as §12.6 (memory) and §12.7 (AOSP -j cap) — an ORTHOGONAL exhaustion axis, both MUST be checked independently before scaling parallelism. Mandatory pre-scale diagnosis: `ulimit -u/ulimit -Hu (soft/hard) + live ps -L --no-headers -u "$USER" | wc -l` thread count; headroom = soft minus live; signals (EAGAIN, fork: retry, failed to create new OS thread) never guessed, always captured (§11.4.6); low headroom → SERIALIZE, avoid GOMAXPROCS=nproc-class spikes. Raising the limit — three cited techniques: (a) `prlimit --pid <PID> --nproc=<soft>:<hard>` (root, no-restart); (b) `/etc/security/limits.conf/systemd LimitNPROC` (persistent, next-session); (c) `self ulimit -u <N>` up to the existing hard ceiling (current session) — `su -c` scope is transient only. Freeing threads by stopping non-essential tooling REQUIRES explicit operator authorization (§11.4.122, no-silent-removal); cmdline-pattern kill helpers MUST exclude the caller's own `$$/$PPID` (self-kill hazard). Classification: universal (§11.4.17). Composes §12/.6/.7/§11.4.58/.101/.103/.122/§9.2. Propagation gate CM-COVENANT-12-12-PROPAGATION (literal 12.12) + recommended gate CM-NPROC-HEADROOM-CHECK + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §12.12. Non-compliance is a release blocker. No escape hatch — no --ignore-thread-limit, --dispatch-blind, --skip-headroom-check, --autonomous-tooling-kill, --assume-ulimit-scope flag.

§11.4.186 — Anti-divergence enforcement: cross-document consistency is a mandatory-before-export/commit gate, never an after-the-fact audit (research-derived, 2026-07-08). Concise summary: any project maintaining more than one representation of the same tracked data (plan + brief + workable-items DB + summaries + .md/.html/.pdf/.docx exports) MUST enforce cross-document consistency as a deterministic PASS/FAIL/SKIP gate that runs BEFORE any export

render, any doc/DB sync verify, and any doc-set commit — never as an after-the-fact human audit. Five decidable check families (DEDUP/TIMELINE/CROSS-DOC/INTEGRITY/STRUCTURAL); §11.4.86 fingerprint re-arms on any input change; self-validated analyzer (golden-good/golden-bad/negative-control) per §11.4.107(10). Honest boundary: proves internal consistency, not plan correctness. Composes §11.4.6/.12/.33/.44/.50/.53/.54/.60/.65/.73/.75/.85/.86/.91/.93/.95/.106/.107/.108/.110/.148/.176. Propagation gate CM-COVENANT-114-186-PROPAGATION (literal 11.4.186) + recommended gate CM-DOC-INTEGRITY-VALIDATION + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.186. Non-compliance is a release blocker.

§11.4.187 — Automatic multi-track ruler orchestration (User mandate, 2026-07-04; landed 2026-07-09). Concise summary: multi-track parallel-development orchestration MUST ship as a universal, automatic, out-of-the-box, inherited engine (never a per-project reimplementation) at constitution/scripts/multitrack/ — a single conductor session that spawns/resumes/monitors/controls per-track headless workers (native headless claude -p --output-format stream-json, session_id from the init event, success read from result.is_error never exit code), per-subscription auth + rate-limit fallback (unset ambient ANTHROPIC_API_KEY, rebind on 429), crash-resilient self-supervisor (durable state, write-temp-then-rename), idempotent bootstrap auto-invoked on every constitution pull. Project-agnostic engine; consumer supplies only per-host config/multitrack/<hostname>.yaml as data. Composes §11.4.20/.28/.58/.70/.101/.103/.111/.116/.119/.126/.128/.131/.147/.164/.167/.176/§12.6/§12.8. Propagation gate CM-COVENANT-114-187-PROPAGATION (literal 11.4.187) + recommended gates CM-MULTITRACK-ENGINE-IN-CONSTITUTION / CM-MULTITRACK-AUTOINSTALL-WIRED / CM-MULTITRACK-BOOTSTRAP-IDEMPOTENT / CM-MULTITRACK-BIND-ON-START / CM-MULTITRACK-FALLBACK-MONITOR / CM-MULTITRACK-HEADLESS-SPAWN / CM-MULTITRACK-RULER-SUPERVISOR + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.187. Non-compliance is a release blocker.

§11.4.188 — Regular main → feature merge cadence: every track/feature-branch MUST FREQUENTLY merge canonical main into itself DURING the work, never only at the end (User mandate, 2026-07-09). Concise summary: every long-lived feature branch AND every parallel-development track MUST regularly git merge (NEVER rebase) canonical main/master INTO itself THROUGHOUT the work — after every trunk tag, at minimum daily while live, and before starting significant new work — so the branch never drifts far and the eventual back-merge stays small and low-conflict. Fetch-first, §9.2 pre-op backup before large merges, resolve conflicts carefully (zero markers, zero dropped files), integrate only when quiescent (§11.4.84), never force-push after merge. Post-merge proof required: smoke gate GREEN + zero conflict-marker grep + no-lost-commit check. Generalises §11.4.167(D) from big feature-streams to every branch. Composes §9.9.2/.37/.41/.42/.71/.84/.88/.89/.103/.113/.167/.176/.178/.179/.181.

Propagation gate CM-COVENANT-114-188-PROPAGATION (literal 11.4.188) + recommended gate CM-REGULAR-MAIN-MERGE-INTO-FEATURE + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.188. Non-compliance is a release blocker.

§11.4.189 — Most-reopened cases get extra-depth live-testing scrutiny FIRST (User mandate, 2026-07-10). Concise summary: all live testing (on-device/on-target runtime validation) MUST give the cases with the HIGHEST §11.4.55 reopens-count extra-depth retest, re-investigation per §11.4.102/§11.4.146, and full rock-solid physical-evidence confirmation (§11.4.5/.69/.107/.123) — scrutinised FIRST, ahead of the rest of the live-test suite. Strengthens §11.4.132(d) (most-reopened as ranking factor) into its own extra-depth mandate at the live-test layer. Composes §11.4.5/.6/.55/.69/.102/.107/.108/.115/.123/.129/.130/.132/.135/.146. Propagation gate CM-COVENANT-114-189-PROPAGATION (literal 11.4.189) + recommended gate CM-MOST-REOPENED-EXTRA-DEPTH-LIVE-TEST + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.189. Non-compliance is a release blocker.

§11.4.190 — Website engineering-quality mandate: every project website MUST be fully responsive + completely SEO-optimized + uniquely OpenDesign-authored + bleeding-edge enterprise-quality, each PROVEN with captured evidence (User mandate, 2026-07-10). Concise summary: every website/web-UI surface a project ships — INCLUDING the project's own website — MUST satisfy AND PROVE with captured physical evidence (never a bare claim): (A) full responsiveness across all browser engines/OSes/device classes/screen sizes with no broken layout at any breakpoint; (B) complete SEO optimization (semantic HTML, meta/OG/Twitter cards, canonical URLs, structured data, sitemap/robots, WCAG AA, Core Web Vitals); (C) unique OpenDesign-token-authored templates (§11.4.162), never generic/off-the-shelf; (D) bleeding-edge enterprise visual quality in both light and dark themes. Proof required per facet: host-rendered cross-breakpoint screenshots + layout oracle, automated SEO audit meeting a score floor, OpenDesign-token provenance, §11.4.170 device-independent pixel proof. Strict web expansion of §11.4.162 + binds §11.4.170; sibling of §11.4.168. Composes §11.4.5/.6/.69/.107/.162/.168/.170. Propagation gate CM-COVENANT-114-190-PROPAGATION (literal 11.4.190) + recommended gates CM-WEBSITE-RESPONSIVE-PROVEN / CM-WEBSITE-SEO-OPTIMIZED / CM-WEBSITE-OPENDSIGN-UNIQUE + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.190. Non-compliance is a release blocker.

§11.4.191 — Work-to-track/branch binding enforcement: a logic-group's work can NEVER be committed OR dispatched onto the wrong track/branch (research-derived, 2026-07-10). Concise summary: every logical group of workable items binds to EXACTLY ONE canonical (track, branch) recorded ONCE in the single source of truth (the §11.4.93 workable-items DB `logic_groups.destination` + `canonical_track` + `group_paths` file-scope

manifest); no change belonging to a group may LAND on or be DISPATCHED to a different track/branch. File-scope generalisation of §11.4.181/.182 branch-name consistency to work-PLACEMENT consistency. Preventive PreToolUse guard guard-work-track-binding.sh (inherited by reference, fail-closed) BLOCKS mis-placed commits/dispatches; detective gate CM-WORK-TRACK-BINDING-ENFORCED audits added commits on the current branch. Mis-placed changes reconciled by MERGE onto the canonical branch (never force-push, never commit-loss). Composes

§11.4.6/.28/.29/.35/.58/.75/.84/.93/.95/.109/.110/.113/.119/.167/.176/.177/.178/.181/.182/
§9.2. Propagation gate CM-COVENANT-114-191-PROPAGATION (literal 11.4.191) + recommended gate CM-WORK-TRACK-BINDING-ENFORCED + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.191. Non-compliance is a release blocker.

§11.4.192 — Continuous multi-track auto-backfill: a FREE track MUST be IMMEDIATELY re-assigned its next-highest-priority domain work-set, never left idle while its domain has actionable items (User mandate, 2026-07-11).

Concise summary: the moment any track's worker COMPLETES its current work-set, that FREE track MUST be IMMEDIATELY assigned the NEXT-highest-priority actionable work-set from its own domain (§11.4.191 logic_group/canonical branch), claimed exactly-once (§11.4.176 registry), dispatched as a new §11.4.182-labelled stream — no waiting for a scheduled wake or operator prompt. When the next item's on-device validation is gated (pending build/device/manual-QA), the track still advances the item's SOURCE-SIDE work and honestly SKIPS only the device-gated validation portion. Promotes §11.4.103(B) auto-backfill from the STREAM layer to the TRACK layer. A track idles ONLY when its domain has genuinely zero remaining actionable items or every remaining item is externally blocked with no source-side advancement possible — “merges are done” is NEVER a valid idle reason. Composes

§11.4.42/.58/.72/.94/.97/.101/.103/.167/.176/.178/.179/.182/.187. Propagation gate CM-COVENANT-114-192-PROPAGATION (literal 11.4.192) + recommended gate CM-CONTINUOUS-MULTITRACK-BACKFILL + paired §1.1 mutation. **Canonical**

authority: constitution submodule [Constitution.md](#) §11.4.192. Non-compliance is a release blocker.

Cascaded anchor backfill — §11.4.18/.22/.39/.63/.100 + §11.4.193–§11.4.234 (2026-07-27)

The entries below restore five-carrier lockstep (§11.4.157) with the canonical constitution/Constitution.md, which had advanced to §11.4.234 while every consumer carrier still topped out at §11.4.192. Each entry is a short cascade reference in the established form: it carries the literal anchor token the propagation gates match on, and points at the canonical text, which remains the

sole authority (CONST-059 / §11.4.35). NOTE (§11.4.6 honest boundary): §11.4.64, §11.4.205 and §11.4.206 are CITED but never DEFINED in the canonical constitution (zero block-openers), so they are not cascadable here; that dangling-cross-reference defect belongs to the canon and is reported upstream, not papered over in the consumers.

NOTE (§11.4.6 honest boundary, cont.): §11.4.175, §11.4.203 and §11.4.204 have ZERO occurrences of any kind in the canonical constitution — neither defined nor cited. They are pure numbering gaps, not dangling cross-references: there is nothing to cascade here and nothing to report upstream.

§11.4.18 — Script documentation mandate (User mandate, 2026-05-14).

Cascade reference (short form): Forensic anchor — direct user mandate (verbatim, 2026-05-14): “Make sure that every single script inside the scripts dir (in all subdirs in depth) is fully documented and covered with full user manuals and user guides! ... Whenever is some bash script modified we MUST document it fully and create for it complete user guide(s) and manual(s) - if already exists make sure all of it is in sync and fully updated! No documentation ever can be out of sync with its codebase!” Every Bash / shell / POSIX-sh script ANYWHERE in a project (anywhere under ‘scripts/’, ... Propagation gate CM-COVENANT-114-18-PROPAGATION (literal 11.4.18) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.18 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.22 — Document-sync commit discipline (User mandate, 2026-05-14).

Cascade reference (short form): Forensic anchor — direct user mandate (verbatim, 2026-05-14): “For Issues, Issues_Summary and Fixed docs (and all its exports) we MUST commit and push (only them) as soon as they are updated (synced). Continuation doc and other similar / relevant documentation MUST be committed and pushed too! Like this we will always have up to date status of working items without need to do full commit and push of everything if that is not possible to do!” Classification: §11.4.17-classified universal — every project that tracks work items through an Issues / Fixed lif ... Propagation gate CM-COVENANT-114-22-PROPAGATION (literal 11.4.22) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.22 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.39 — Per-Feature On-Device End-User Validation Mandate (iter-76,

2026-05-17). Cascade reference (short form): Every user-facing feature in a consuming project MUST have at least one HelixQA scenario that exercises the feature from cold-launch with positive runtime evidence: - A screenshot or screen recording captured at a named checkpoint. - At least one assertion of type ‘screenshot’, ‘accessibility_count’, ‘accessibility_node’, ‘ocr_text’, ‘pixel_histogram’,

or 'editor_state'. - Evidence files written to a discoverable output directory so every PASS is cross-verifiable after the fact. ... Propagation gate CM-COVENANT-114-39-PROPAGATION (literal 11.4.39) + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.39 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.63 — Workable-items procedure docs as single source of truth (User mandate, 2026-05-19). Cascade reference (short form): Forensic anchor —

verbatim user mandate (2026-05-19 ~10:30Z): “To make workable items tracking and creation through all of the mentioned mandatory documents we MUST create proper procedure document which will be named in Constitution, AGENTS.md and CLAUDE.md as single source of truth for this procedure! Make sure it is created under: docs/procedures/issues/Creation.md which will be always exported to PDF and HTML whenever it is updated. ... Propagation gate CM-COVENANT-114-63-PROPAGATION (literal 11.4.63) + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.63 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.100 — RETIRED. Cascade reference (short form): Demoted to consumer project (a consuming project's video-color/visual-quality fidelity) per §11.4.17/ §11.4.35 — project-specific (RK3588/MPV/Arvus), not universal. See the consuming project's Constitution/CLAUDE/AGENTS/QWEN. Propagation gate CM-COVENANT-114-100-PROPAGATION (literal 11.4.100) + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.100 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.193 — Anti-blind-typing mandate: every UI interaction MUST be “seen” and “understood” via OCR/vision/screenshot proof, NEVER blind-typed (User mandate, 2026-07-13). Cascade reference (short form): Verbatim operator mandate: “No blind typing of any kind anywhere!!!” and “Login to apps and services MUST BE done with understanding and seeing the whole UI ALWAYS!!!” and “Add into the constitution that blind typing is STRICTLY FORBIDDEN, and instead everything MUST BE seen and understood in coordination with proper models and means we have incorporated in HelixQA!!!” and “No bluff of any kind anywhere!!!” Blind typing — sending keystrokes, text, or input events to a UI surface WITHOUT first capturing AND understanding what is on screen (OCR / vision / ... Propagation gate CM-COVENANT-114-193-PROPAGATION (literal 11.4.193) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.193 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.194 — Exhaustive all-scenario, all-angle code-review + verify-against-captured-runtime-evidence mandate (User mandate, 2026-07-14). Cascade reference (short form): Forensic anchor — verbatim user mandate (2026-07-14): “CRITICAL: Code review MUST TAKE into the accounts all scenarios, look at the problem or change from all angles, deepest possible analysis! No gaps, shortcomings or false results or bluff is allowed! Extend our root constitution Submodule and all relevant files properly so the review process is more precise and we do not have situations like the last one!” Forensic case study (genericised, 2026-07-14 — “the last one”). ... Propagation gate CM-COVENANT-114-194-PROPAGATION (literal 11.4.194) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.194 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.195 — Branch-structure governance: taxonomy + merge-after-live-QA + flavor/product non-merge (User mandate, 2026-07-14). Cascade reference (short form): Forensic anchor — verbatim operator rules (2026-07-14): 1. Branch taxonomy: ‘feat’, ‘product’, ‘flavor’ (extends §11.4.181 / §11.4.167(D)). 2. Tracks on FEATURE branches or ‘main’ MUST — after full confirmation + validation + verification + full LIVE MANUAL TESTING (§11.4.185) — MERGE their work to ‘main’ COMPLETELY (all submodules + main repo), ff-only, NEVER force-push (§11.4.113), no commit lost (§9.2) (extends §11.4.167(I) / §11.4.188). 3. ... Propagation gate CM-COVENANT-114-195-PROPAGATION (literal 11.4.195) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.195 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.196 — Native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery + resource-detection-by-real-identity-not-substring-match (User mandate, 2026-07-14). Cascade reference (short form): Verbatim operator mandate: use ALL operational claude-NATIVE aliases FIRST — as long as ≥ 1 native alias is operational (NOT session-rate-limited, NOT weekly-limit-reached, NOT expired-subscription) — before ANY provider alias; track per alias the session rate-limit / weekly-limit / subscription expiry+renewal from REAL signals (never faked, §11.4.6) and record when each becomes operational AGAIN; as soon as a higher-priority native recovers, rebind tracks to it; ... Propagation gate CM-COVENANT-114-196-PROPAGATION (literal 11.4.196) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.196 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.197 — Research / kicked-off-work completion mandate: every research effort MUST be driven to full, wired, verified completion or explicitly closed

— **never left un-wired in the backlog (User mandate, 2026-07-15).** Cascade reference (short form): Forensic anchor — verbatim operator mandate (2026-07-15): “all research work we start MUST BE fully finished, completely confirmed, validated and verified! It MUST NOT happen for it to stay at the backlog uncompleted and un-wired! Loss of requirements, failure to incorporate them is FORBIDDEN!!! Everything MUST BE fully completed! No exceptions!” Every research effort, kicked-off improvement, design doc, spike, prototype, proof-of-concept, or partially-landed feature that a project STARTS MUST be driven to FULL completion — OR explicitly, evidence-backed ... Propagation gate CM-COVENANT-114-197-PROPAGATION (literal 11.4.197) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.197 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.198 — Default working mechanisms: multi-alias native-first orchestration AND heavy token-optimization are ALWAYS-ON defaults for single- and multi-track work (User mandate, 2026-07-15). Cascade reference (short form): Forensic anchor — operator mandate (2026-07-15): multiple claude-toolkit aliases (native-first) AND heavy token-optimization MUST be the DEFAULT mechanisms, always used with single- and multi-track development and with single or multiple working agents. Two proven mechanisms are promoted to STANDING, always-on defaults — engaged automatically for EVERY working configuration (single-agent OR multiple-agent; single-track OR multi-track), never a per-request opt-in: (A) Multi-alias native-first orchestration is the DEFAULT. ... Propagation gate CM-COVENANT-114-198-PROPAGATION (literal 11.4.198) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.198 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.199 — Exact-reproduction-sequence mandate: when a working reproduction exists, the investigation MUST use ITS exact sequence — a deviating repro proves NOTHING (research-derived, 2026-07-15). Cascade reference (short form): Forensic anchor (FACT, 2026-07-15). An investigation concluded that a defensive “cover” mechanism “never engages” and that the defect it guards was therefore unfixed. The conclusion was WRONG: the hand-rolled reproduction drove the target with a DIFFERENT input than the existing test did (it pressed a HOME key where the test sent a media-pause key), so the hand-rolled path never reached the PRECONDITION the mechanism keys off. Re-driven with the TEST’S EXACT sequence, the mechanism provably ENGAGED. ... Propagation gate CM-COVENANT-114-199-PROPAGATION (literal 11.4.199) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.199 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.200 — Non-targetable deploy/flash tooling MUST isolate exactly ONE eligible target and VERIFY-AFTER-WRITE on the INTENDED target (research-derived, 2026-07-15). Cascade reference (short form): Forensic anchor (FACT, 2026-07-15). A flashing tool with no device-selection argument AUTO-SELECTED a device from the bus. With a sibling board still attached, it silently wrote the freshly-built image to the WRONG board, printed a SUCCESS message, and the INTENDED board kept its OLD build. Every downstream test then validated a target that had never received the artifact: the §11.4.108 ARTIFACT → RUNTIME-ON-CLEAN-TARGET layer was silently broken while every tool and gate reported green. ... Propagation gate CM-COVENANT-114-200-PROPAGATION (literal 11.4.200) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.200 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.201 — Every guard/gate MUST assert the REAL condition: a false-positive refusal is a FAIL-bluff, a false-negative pass is a PASS-bluff (research-derived, 2026-07-15). Cascade reference (short form): Forensic anchor (FACT, 2026-07-15 — two independent instances in one cycle). (a) A host-budget guard used a bare ‘pgrep -f’ and REFUSED every worker spawn on a host with ZERO real builds, because a live worker’s own command line merely QUOTED the pattern (the §11.4.196(D) / §12.12 carrier footgun). (b) A host-safety pre-flight refused a build on a “swap full” reading that did not reflect a real memory-pressure condition, and cleared only after the operator manually cycled swap. ... Propagation gate CM-COVENANT-114-201-PROPAGATION (literal 11.4.201) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.201 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.202 — Reporting directives: a report MUST auto-create a fully-populated, fully-synced workable item — never a prose acknowledgement (User mandate, 2026-07-15). Cascade reference (short form): Verbatim operator mandate: introduce reporting directives ‘ISSUE’ / ‘BUG’ / ‘TASK’ that create a proper workable item filled with all details, with the WHOLE workable-items system automatically up to date and fully in sync — the main DB, all documents, and all external tracking systems; universal, reusable on every project, fully decoupled, recognized immediately by every project that pulls the constitution. ... Propagation gate CM-COVENANT-114-202-PROPAGATION (literal 11.4.202) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.202 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.207 — Instant multi-stream resume engine: whole-fleet continuation state is a durable, content-addressed, atomically-committed snapshot resumed in O(changed) reads, not an O(history) re-read (User mandate, 2026-07-15). Cascade reference (short form): Compact summary: the continuation mechanism (§12.10 / §11.4.127 / §11.4.131 / §11.4.205) is EXTENDED — never replaced — by a mechanical engine so a fresh session resumes ALL work (every Track, main stream, and agent) at once, INSTANTLY and at token cost O(changed streams), by reading ONE snapshot manifest + one compact blob per stream, NEVER re-reading the full handoff history. ... Propagation gate CM-COVENANT-114-207-PROPAGATION (literal 11.4.207) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.207 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.208 — Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (content + timestamp+timezone + track + alias + model + effort) (User mandate, 2026-07-15). Cascade reference (short form): Verbatim operator mandate (2026-07-15): introduce a request-history document that captures, per operator request/prompt — the full request content, the accepted timestamp with an explicit timezone, the track that processed it, the alias that took the workable item, and the model + effort used — always kept in sync, with no false data and no bluff of any kind. The operator's load-bearing correction (§11.4.35 / §11.4.17): the RULE (this universal discipline) belongs in the constitution submodule; ... Propagation gate CM-COVENANT-114-208-PROPAGATION (literal 11.4.208) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.208 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.209 — Code-review MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-15). Cascade reference (short form): Verbatim operator mandate (2026-07-15): the mandatory independent code-review MUST ALWAYS be performed with the Fable model at xhigh effort; if Fable is unavailable, use Opus at xhigh effort. Every mandatory independent code-review the constitution requires — the §11.4.125 code-review-agent gate before pre-build + main build, the §11.4.142 universal every-change-review (no exception), and the §11.4.194 exhaustive all-scenario/all-angle review + verify-against-captured-runtime-evidence, iterated to a zero-finding GO per §11.4.134 — MUST be dispatched on t ... Propagation gate CM-COVENANT-114-209-PROPAGATION (literal 11.4.209) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.209 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.210 — Zero-loss request/prompt intake: every operator request MUST be mechanically captured, tracked, and processed — never skipped, ignored, avoided, or lost (User mandate, 2026-07-15). Cascade reference (short form): Forensic anchor — verbatim user mandate (2026-07-15): “We MUST CHECK which requests / prompts we have given haven’t passed through except the accumulative prompts / requests history document !!! There may be more of it! This MUST NEVER happen! We MUST make sure that every request / prompt we issues is ALWAYS respected and taken into the account, executed and processed! This is MANDATORY RULE !!! There MUST NOT BE any avoiding, ignoring, skipping or any form of loss for requests / prompts we make! Extend properly constitution Submodule and add mandatory li ... Propagation gate CM-COVENANT-114-210-PROPAGATION (literal 11.4.210) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.210 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.211 — Merge-conflict resolution during main → feature/product/flavor merges MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-16). Cascade reference (short form): Verbatim operator mandate (2026-07-16): “During the merge of main branch to all feature/product/flavor branches of tracks 2 - 4, if conflicts happen they MUST BE resolved using Fable model with xhigh effort! If Fable is not available, then with Opus with xhigh! Add this MANDATORY RULE into constitution Submodule and its files CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md, and other related relevant files!” Any merge-CONFLICT resolution performed while integrating the canonical trunk (‘main’/‘master’) into a feature / product / flavor branch — ... Propagation gate CM-COVENANT-114-211-PROPAGATION (literal 11.4.211) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.211 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.212 — Main README is the canonical starting-point / entry point for ALL project documentation — no doc may be an orphan unreachable from README (User mandate, 2026-07-16). Cascade reference (short form): Verbatim operator mandate (2026-07-16): “Manual tests catalog must have linking all through the main README as the starting point of this and any other Project documentation! This fact and mandatory rule MUST be added into the constitution.” The main README (‘README.md’ at the repository root) is the canonical STARTING POINT / entry point for ALL project documentation, for THIS project and EVERY project incorporating this constitution. ... Propagation gate CM-COVENANT-114-212-PROPAGATION (literal 11.4.212) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.212 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.213 — FEATURE research-scheduling directive: recognized via all §11.4.140 forms, SCHEDULES (never synchronously executes) a deep, enterprise-grade research + implementation-planning effort as a tracked workable item (operator-directed mandate, 2026-07-16). Cascade reference (short form): Mandate (2026-07-16): extend the §11.4.202 reporting-directive family with a fourth registered action, ‘FEATURE’, that turns a plain-language feature description into a SCHEDULED, deep, enterprise-grade research-and-planning effort — mirroring EXACTLY how ‘ISSUE’ / ‘BUG’ / ‘TASK’ turn a report into a tracked workable item, but for RESEARCH SCHEDULING rather than defect/task/feature reporting: the directive itself creates the tracked item and enqueues it durably; ... Propagation gate CM-COVENANT-114-213-PROPAGATION (literal 11.4.213) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.213 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.214 — Recurrence-links-not-mints: a defect that returns MUST reopen its existing item, never enter as a new id (research-derived, 2026-07-17). Cascade reference (short form): Forensic anchor (genericised, 2026-07-17). In a consuming project, operator-reported defects that were RECURRENCES of already-tracked, already-“fixed” defects re-entered the tracker as BRAND-NEW ids rather than reopening the originals (measured: at least five distinct recurrence chains, one three deep). ... Propagation gate CM-COVENANT-114-214-PROPAGATION (literal 11.4.214) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.214 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.215 — A doc that BINDS work MUST live, tracked, in the repository where that work happens (research-derived, 2026-07-17). Cascade reference (short form): Forensic anchor (genericised, 2026-07-17). In a consuming project, the document that DEFINED a binding acceptance gate — the gate a body of work was ordered to satisfy before proceeding — was untracked and lived in an out-of-tree scratch directory belonging to a different working checkout, i.e. it existed in NO repository and in no other agent’s or operator’s checkout. ... Propagation gate CM-COVENANT-114-215-PROPAGATION (literal 11.4.215) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.215 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.216 — Canonical machine-readable design-token source: one CSS custom-property file (:root + [data-theme="dark"]), generated bindings everywhere else (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). The mandate’s shape was proven end-to-end in the Helix Thready

design package ('docs/public/research/mvp/design/'): ONE machine-readable token file ('opendesign/tokens.css') carries every color / type / spacing / radius / elevation / motion value as CSS custom properties — light values in a 'root {}' block, dark values in a '[data-theme="dark"]' override — and every value is evidence-traceable, never invented: the brand chartreuse '#B6E376' is the recorded eyedrop MEAN of ... Propagation gate CM-COVENANT-114-216-PROPAGATION (literal 11.4.216) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.216 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.217 — OpenDesign brand contract: every UI-shipping product carries a 9-section DESIGN.md + tokens.css twin with disjoint color roles, AA-pinned accents, and a locked attribution footer (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready ships 'design/opendesign/DESIGN.md' in the OpenDesign 9-section brand-contract schema — VERIFIED against the tool's own source per §11.4.99 (spec, 'design-systems/README.md', the contributor authoring guide, and a measured 150-file corpus: two coexisting heading families, both parsed by number prefix only) — plus its machine-readable twin 'tokens.css' per the OpenDesign token schema, and an honest 'TOOLING.md' capability check. ... Propagation gate CM-COVENANT-114-217-PROPAGATION (literal 11.4.217) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.217 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.218 — Living design-library catalogue: every component × state × theme rendered self-contained + a cross-platform reusability matrix with per-cell provenance (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's 'design/library/' renders the ENTIRE component set as a living catalogue: 'components.html' — 14 groups, ~38 components, ≈ 118 component-state cells, each auditable in BOTH themes via a real toggle (≈ 236 rendered states), zero external requests, interactive behaviors REAL not mocked (theme cycling stamps 'data-theme' and persists; column sorting drives 'aria-sort'; a real 'showModal()' dialog; ... Propagation gate CM-COVENANT-114-218-PROPAGATION (literal 11.4.218) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.218 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.219 — Per-app screen catalogues: full-IA coverage per surface, per-platform chrome variants, honesty markers for unbuilt behavior (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor —

proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's 'design/screens/' carries a screen catalogue PER SURFACE — 13 web screens + an interactive prototype shell walking the four key journeys, 11 device-framed 390×844 mobile artifacts with an Android/iOS chrome TOGGLE (+ HarmonyOS/Aurora port notes) covering EVERY node of the wireframes' information architecture, a desktop shell with per-OS title bars / menus / tray / file-drop mocks, ten 80-column TUI mockups covering the full state machine, and a 3-page market ... Propagation gate CM-COVENANT-114-219-PROPAGATION (literal 11.4.219) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.219 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.220 — Open-first design tooling: self-hosted open platforms are primary; proprietary tools are import/export targets only; design sources live in-repo in open formats (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's tooling capability check ('opendesign/TOOLING.md') and export capstone prove the posture: every design source lives IN-REPO in open formats (7 SVG brand masters, 'tokens.css', JSON manifests + 28 genuine '.od-figma.json' IR captures, 6 hand-authored Lottie JSONs, self-contained HTML screens); the OpenDesign daemon/CLI status was RUNTIME-PROBED (live health endpoint, CLI round-trip, headless-vs-operator matrix), never assumed; ... Propagation gate CM-COVENANT-114-220-PROPAGATION (literal 11.4.220) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.220 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.221 — Motion discipline: token-bound durations/easings + machine-readable manifest + reduced-motion static fallbacks everywhere (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's 'design/motion/' package proves the discipline: six hand-authored vector Lottie animations bound to the token motion values, each validated by parse + structural schema check + REAL render (lottie-web under jsdom with animation-progression assertions, plus headless-Chromium light+dark contact sheets) — unverified runtimes ('lottie-compose' / 'lottie-ios') honestly '[OPEN]'-tracked; ... Propagation gate CM-COVENANT-114-221-PROPAGATION (literal 11.4.221) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.221 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.222 — Design-completion export wave: a design is not done until token codegen, the honest-output raster matrix, both-theme renders, hand-off

bundles, and doc twins land (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready’s export capstone proves the wave: an icon/asset raster matrix locked 1:1 to its generator (‘generate-raster.sh’), which SELECTS its SVG renderer by PROBE — each candidate (rsvg-convert → inkscape → resvg → headless Chromium → ImageMagick) must pass a real probe render with content verification before it is used, a failing candidate falls through with a note, and if NONE passes the script exits with instructions and writes NOTHING (“No PNGs were written (output ... Propagation gate CM-COVENANT-114-222-PROPAGATION (literal 11.4.222) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.222 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.223 — Provenance-marker discipline for design documentation + central open-item registry in the area index (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Every statement in the Helix Thready design area carries an inline provenance marker from a closed vocabulary (‘[VERIFIED]’, ‘[OPERATOR]’, ‘[RESEARCH]’, ‘[DEFAULT — adjustable]’, ‘[ASSUMED]’, ‘[BUILD-NEW]’, ‘[IN-HOUSE: module]’, ‘[CONSTITUTION \$x]’, ‘[GAP: id]’, ‘[OPEN: id]’); ... Propagation gate CM-COVENANT-114-223-PROPAGATION (literal 11.4.223) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.223 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.224 — Test-first (TDD) for ALL work, not only fixes + a minimum code-coverage floor that is NECESSARY-never-sufficient (User mandate, 2026-07-22). Cascade reference (short form): Forensic anchor — verbatim user mandate (2026-07-22): “Any work we do MUST START by writing the test! We MUST FULLY comply and follow the TDD - The test driven development! It MUST BE applied for any change, implementation, wiring, tests, challenges and any other form of codebase we write! Bash scripts as well! Coverage with TDD MUST BE close to 100%, not less than 85% !!!” The gap this closes. ... Propagation gate CM-COVENANT-114-224-PROPAGATION (literal 11.4.224) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.224 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.225 — Scheduler-quota burst-throttling telemetry + interactive-scope isolation from bursty fleets: progressive interactive degradation is diagnosed from throttling accounting over time, never from averages, and an interactive-latency-sensitive process never shares a fixed-quota no-burst scope with a bursty fleet (research-derived, 2026-07-23). Cascade reference

(short form): Forensic anchor (genericised, 2026-07-22 — two independent instances in one day). (a) Progressive interactive sluggishness — worsening with the number of live subagents — was root-caused to a scheduler CPU quota (a cgroup ‘cpu.max’ limit) with NO burst allowance on a session scope hosting BOTH the operator’s interactive terminal-multiplexer server AND a spawning agent fleet. Measured live on a 64-CPU host: the scope’s quota was 9.6 CPU per 100 ms enforcement period; ... Propagation gate CM-COVENANT-114-225-PROPAGATION (literal 11.4.225) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.225 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.226 — Evidence-class-at-closure + standing detection pressure: machinery presence does NOT predict whether a fix holds — the EVIDENCE CLASS at closure (runtime vs source) under real DETECTION PRESSURE does; prose does not bind, seams do (research-derived, 2026-07-23). Cascade reference (short form): Forensic anchor (genericised — the 2026-07-22 reopen/first-touch root-cause forensics; operator mandate 2026-07-22, verbatim: “do deep research on how we have for very long time so much reopened workable items!? Why the quality of delivered releases was so poor mostly? We mostly detect the same major problems on first touch when new releases are available, and this happens over and over again!”). The single strongest empirical result across four independent evidence lines (internal forensics; a 76-source external literature pass; ... Propagation gate CM-COVENANT-114-226-PROPAGATION (literal 11.4.226) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.226 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.227 — Governance-corpus self-custody: the rule corpus is bound by the same custody it imposes — every named gate is implemented-or-registered-deferral under a monotone ratchet, and propagation gates count anchor BLOCKS (exactly-once + lockstep-identical), never bare literals (research-derived, 2026-07-23). Cascade reference (short form): Forensic anchor (measured on THIS corpus, 2026-07-22 — commands and control needles preserved in the root-cause analysis). The canonical rule file held 10,735 lines and 220 distinct §11.4.N anchors; 413 named CM-* gates, of which 241 (58%) had no implementation in either canonical gate site, plus 85 literal “gate-code = separate work item” deferrals; the status-custody seam DESIGNED on 2026-07-17 was still prose on 2026-07-22 — the corpus’s highest-value remediation fell into its own gap within five days; ... Propagation gate CM-COVENANT-114-227-PROPAGATION (literal 11.4.227) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.227 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.228 — Cross-agent extension lifecycle mandate: per-platform compatibility declaration, source tracking, per-extension documentation with compatibility matrix, and auto-wiring of constitution extensions into all detected agent platforms on constitution pull (research-derived from Extensions Catalog forensics, 2026-07-24). Cascade reference (short form): Forensic anchor (FACT, 2026-07-24 — measured from the live Extensions Catalog at ‘docs/extensions/EXTENSIONS_CATALOG.md’). The project’s extension ecosystem — skills, MCP servers, ACPs, LSPs, plugins, and supporting infrastructure (34 skills, 1 MCP, 38 provider aliases, 21 CCR providers, 4 daemons, plus a separate HelixSkills repo) — exhibited four structural gaps: (G1) of 18 active Claude Code project skills, only Claude Code and OpenCode platforms are explicitly supported — HelixCode, Gemini CLI, and Qwen Code have zero recorded loading paths despite b ... Propagation gate CM-COVENANT-114-228-PROPAGATION (literal 11.4.228) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.228 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.229 — Live in-session task/todo tracker MUST always be up to date and fully in sync with the real work state — never stale, never showing completed work as pending nor pending as done (User mandate, 2026-07-25). Cascade reference (short form): Verbatim operator mandate (2026-07-25): “make sure it is ALWAYS up to date fully in sync! CRITICAL: This MUST BE one of mandatory rules / constraints! Add it with all relevant details into the constitution Submodule and relevant related files.” The agent’s LIVE IN-SESSION TASK/TODO TRACKER — the harness-provided task list / ‘TodoWrite’-class todo tool / any equivalent moment-to-moment execution-state surface the runtime exposes DURING a session — MUST ALWAYS be kept up to date and FULLY IN SYNC with the real live work state: never stale, never showing co ... Propagation gate CM-COVENANT-114-229-PROPAGATION (literal 11.4.229) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.229 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.230 — Parallelized-pipeline methodology: build ↔ validate overlap + affected-stages-only re-run + parallel multi-device testing/recording + fan-out-subagents-always (research-derived, 2026-07-26). Cascade reference (short form): Forensic motivation (genericised): the sequential pipeline shape — build FULLY completes → THEN validation begins → on ANY discovery, restart the WHOLE cycle from scratch → all devices tested one-at-a-time — serialises the operator’s wall-clock onto build latency + one-device-at-a-time evidence gathering, so a multi-hour containerised build followed by serial per-device validation and a from-scratch restart-on-discovery makes time-to-a-manual-QA-capable-release a multiple of what the genuinely-parallel work requires. ... Propagation gate CM-COVENANT-114-230-PROPAGATION (literal 11.4.230) + paired §1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.230 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.231 — Nano-precision model-tier selection: lightest-capable-model-first with dynamic mid-task escalation (User mandate, 2026-07-26, CRITICAL).

Cascade reference (short form): Verbatim operator mandate: every work effort MUST use the LIGHTEST CAPABLE model, escalating tier only as complexity rises — a nano-precision, per-task model-tier choice, NOT “everything on the heaviest model.” Forensic motivation (genericised, §11.4.6): dispatching every task — a trivial doc-sync, a one-line status read, a mechanical format fix — on the heaviest available model wastes capacity + token budget + wall-clock that a lighter tier would spend identically well, while the OPPOSITE failure (a hard multi-factor root-cause investigation dispatched ... Propagation gate CM-COVENANT-114-231-PROPAGATION (literal 11.4.231) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.231 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.232 — Anti-mess long-op orchestration: every long-op is registered, single-owned-per-purpose, liveness-PROVEN, handed-off-on-stop, safely-reaped, consistency-checked, and escape-bounded — the single source of truth for what is running, never inferred (research-derived, 2026-07-26).

Cascade reference (short form): Forensic anchor (genericised): in one cycle a verification sweep produced five distinct mess states — (1) multiple concurrent same-purpose sweeps ran unintentionally (no single-owner-per-purpose claim); (2) a sweep HUNG on a consumer-blocked-on-an-unclosed-producer pipeline and was undetected because liveness was inferred from process-alive (a §11.4.201(6) FALSE-NULL: a wedged op and a progressing op returned the identical “process alive” signal); ... Propagation gate CM-COVENANT-114-232-PROPAGATION (literal 11.4.232) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.232 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.233 — Standing anti-mess control plane: a level-triggered, idempotent, invariant-catalogue reconciliation loop over the WHOLE persistent System state, orchestrating every per-domain enforcer, refusing gated transitions on un-reconciled drift (research-derived, 2026-07-26).

Cascade reference (short form): Forensic anchor (genericised): every existing consistency enforcer in a multi-track / multi-agent / multi-submodule / multi-upstream System is EDGE-TRIGGERED — a PreToolUse hook fires on an action, a check runs at a seam (lock-acquisition / merge-time / build-time / commit-time), a gate runs when a build runs — so a mess that arises WITHOUT an edge firing is un-observed until it causes harm (two live examples captured 2026-07-26: a large uncommitted-batch

pile-up sat unreconciled through a session, and multiple provably-stale VCS/wrapper locks were held ... Propagation gate CM-COVENANT-114-233-PROPAGATION (literal 11.4.233) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.233 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

§11.4.234 — Dedicated hook-validation script: hook checks run as an explicit stage of a dedicated commit/push script, and the commit/push mechanism is ALWAYS unblocked (operator mandate, 2026-07-26). Cascade reference (short form): Forensic anchor (genericised): a consumer project wired its full local CI gate as an automatic git ‘pre-push’ hook (‘core.hooksPath’ → hook → multi-minute build/test/sweep stage on EVERY push). The result, captured 2026-07-26: routine pushes took many minutes each, and when one gate stage inside the hook failed or stalled, the push was rejected outright — the commit/push mechanism itself became blocked, with pending commits stranded locally and no operator-usable path forward short of a forbidden ‘-no-verify’ bypass. The gate logic was sound; ... Propagation gate CM-COVENANT-114-234-PROPAGATION (literal 11.4.234) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.234 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.

Cascaded anchor backfill — §11.4.235 (2026-07-28)

Continues the cascade above. The canonical constitution/[Constitution.md](#) advanced past §11.4.234 after that backfill closed: §11.4.235 was minted in constitution commit fce304d and its clause (B) extended in bbbdb5 (constitution HEAD at cascade time: 32d7578). This entry restores five-carrier lockstep (§11.4.157) for the new anchor. Form and authority are unchanged from the entries above — the compact cascade reference carries the literal anchor token the propagation gates match on, and the canonical text remains the sole authority (CONST-059 / §11.4.35).

§11.4.235 — Build-and-deploy the moment source fixes are proven-correct (test-side hardening is a parallel stage, never a build gate) + post-deploy version increment (operator mandate, 2026-07-27). Cascade reference (short form): (A) the moment a batch’s SOURCE fixes are done AND proven-correct — proven-correct = the mandatory independent SOURCE-correctness review has returned a clean GO (§11.4.125/§11.4.142/§11.4.194 on the source diff, on the §11.4.209 Fable-‘xhigh’ substrate, iterated to zero-finding/zero-warning GO per §11.4.134) — the build MUST be triggered, pre-build + post-build tests EXECUTED (§11.4.108 layer-1 before, layer-2 after), and the artifact deployed; gate-instrumentation hardening + new-gate authoring (§11.4.110), paired §1.1

mutations + four-layer §11.4.4(b) coverage + HelixQA Challenges, ‘gates and checks of all kinds’ (§11.4.96), and the re-review of THAT test-instrumentation work all run IN PARALLEL with the build AND the §11.4.185 manual-QA session, NEVER gating the build (§11.4.230(A)) — the build’s GO-SIGNAL is source-correctness alone, and holding the build sequentially for the test-instrumentation buildout when the source is already proven-correct is a §11.4.235(A) violation. (B) the moment image(s)/artifact(s) deploy the version code/version is INCREMENTED so every deployed artifact carries a distinct, monotonic, greppable id (§11.4.151/ §11.4.108/§11.4.200; two deployments sharing one version id defeats per-deploy artifact identity), and — extended 2026-07-28 — EVERY deployment of a release FOR MANUAL QA TESTING (§11.4.185) is a DEVELOPMENT-CYCLE BOUNDARY that CLOSES the outgoing cycle and STARTS a new one at the NEXT monotonic version, ALWAYS and by standing default, so fixes for QA findings against the deployed artifact land in the NEW cycle on the NEW version id and are never patched onto the already-QA-deployed version id. (C) STANDING DEFAULT (§11.4.126/§11.4.230/ §11.4.103), never a per-request opt-in. Honest boundary (§11.4.6): this relocates the TEST-INSTRUMENTATION re-review + gate/check hardening OFF the build critical path and weakens NO substantive gate — the SOURCE-correctness review still PRECEDES and gates the build, pre/post-build tests still run at their §11.4.108 L1-2 seams, and §11.4.185 / §11.4.129 / §11.4.40 are un-weakened; ‘test-side work is parallel, never a build gate’ is bounded to future-regression instrumentation, because a check that proves the SHIPPED artifact wrong is a source/artifact-correctness gate (§11.4.108 L1-2) and stays on the critical path. Propagation gate CM-COVENANT-114-235-PROPAGATION (literal 11.4.235) + paired §1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.235 — the full mandate text is canonical there and is inherited by reference (CONST-059 / §11.4.35). Non-compliance is a release blocker.