

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Operator-blocked	1
Bug	Queued	50
Bug	Ready for testing	1
Task	Operator-blocked	1
Task	Queued	14
TOTAL		67

Items

ATM ID	Type	Status	Severity	Description
				Reported-Via: §11.4.202 reporting directive to HelixSkills Operator What (the report, verbatim): Feature request for the HelixSkills System (git@github.com:HelixDevelopment/skills.git) and HelixAgent, covering ALL of its power-features. OPERATOR REQUIREMENTS (2026-07-29): “Full coverage of the HelixAgent the HelixSkills System for all workable power-features git@github.com:HelixDevelopment/skills.git , without power-features fully without nothing leftout! Full coverage into phases, tasks and subtasks with exhaustive test suites prepared! Full coverage with ll supported test suites (suites) is mandatory! Extend and update all existing guides, FAQs, create stunning illustrations, graphics, all materials! Any gaps, shortcomings, weak spots, inconsistencies MUST BE detected in advance and tackled in process and properly tackled with comprehensive testing, risk-free and safe! Track all progress through the system.”
HXC-159	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
				and remebering mechanisms we have! Use Git with bridge to Superpowers used for the imple development mechanism!" PRE-VERIFIED ENV not assumed — §11.4.6): - GitHub SpecKit IS in NOT yet initialized in this repo (no .specify/ and is therefore an explicit early task, not an assur skills.git IS reachable via SSH. git ls-remote - catalog-docs, feature/deep-research, feature/te branch/tag inventory MUST be re-derived at re partial listing. - The repository is NOT currentl contains no HelixDevelopment/skills entry; the codex-skills -> vasic-digital/caf-codex-skills.git, not be confused with it. - Both consumers exist (module dev.helix.code) and the submodule su dev.helix.agent). MANDATORY DELIVERY CON for ALL phases of the incorporation, bridged to executed subagent-driven (§11.4.70 / §11.4.20). tasks -> subtasks, with an exhaustive level of d where the design is load-bearing). - Complete t e2e, full-automation, Challenges (vasic-digital/ (HelixDevelopment/helix_qa) with autonomou concurrency/atomicity, race/deadlock, memory captured evidence (§11.4.5 / §11.4.69); every gu Documentation: extend AND update every existi only new ones. Produce illustrations, graphs and materials. Four-format export discipline per § the main README per §11.4.212. - Risk work is note: every gap, shortcoming, weak spot, dang detected DURING planning and each one paire free and safe solution (§11.4.102 systematic-de before rewrite (§11.4.74 / §11.4.28): survey the catalogues on GitHub AND GitLab first; extend duplicating, keeping them project-agnostic and incorporated as a properly-laid-out dependenc level or submodules/, never a nested own-org and HelixAgent must BOTH receive the full fea the other is an incompleteness to be tracked, n through the continuation docs, the memory/kr the workable-items SSoT (§12.10 / §11.4.131 / § 1. A SpecKit specification exists and is initializ feature, with an explicit feature inventory pro repository (an enumerated list, not a sample — implementation plan exists at the stated depth

ATM ID	Type	Status	Severity	Description
				mitigation. 3. Every planned capability is wired to consumers, proven by §11.4.108 runtime signature compiled. 4. The full §11.4.169 test-type matrix and HelixQA banks included. 5. All documentation produced diagrams/illustrations incorporated, Independent code review reaches a zero-finding. No capability is claimed without runtime evidence every closure (§11.4 / §11.4.1 / §11.4.123). §11.4.123 PROGRAM (scheduled, not yet executed): (a) Develop scientific papers / open-source projects+codebooks to incorporate this feature into the project (§11.4.102) of ALL data obtained to ensure no zone / inconsistency / imperfection, and design one found. (c) The design MUST ALWAYS be enhanced innovative – never less. (d) MANDATORY exhaustive documents, an implementation plan down to low-level diagrams / schemes / graphs, illustrations, SQL queries relevant material (§11.4.65 / §11.4.73). (e) Invest in + HelixDevelopment submodules/components freely where genuinely needed while keeping reusable (§11.4.28 / §11.4.74 / §11.4.177). (f) Plan one: every supported test type, the Challenges (§11.4.27 / §11.4.169). (g) Divide the work into nano-detailed enough to drive integration / implementation directly. (h) Plan explicitly for enterprise scalability: the feature has a UI/UX surface, produce full wireframe PSD / PDF, or whatever the project mandates) and §11.4.190). (j) Plan full CodeGraph integration + synchronisation (§11.4.78 / §11.4.79 / §11.4.80). items (every detail + reference + attachment) source-of-truth (§11.4.93 / §11.4.95) + every design-related project doc/component + every connection (§11.4.148 D5) – honestly SKIPPING any absent (credentials_absent / tracker_client_absent, §11.4.148 Research-doc destination (to be populated when fully_incorporate_the_helixskills_system_into_Execution model (§11.4.213 clauses 1-2): this is a multi-day research/planning/documentation work project’s standing autonomous loop (§11.4.87 / §11.4.177) when it claims this item, and MUST be driven to explicitly evidence-backed CLOSED terminal state be left sitting un-wired in the backlog. Affected research/planning – destination: docs/research/

ATM ID	Type	Status	Severity	Description
				fully_incorporate_the_helixskills_system_into_ Reproduction / context: UNKNOWN: not state established before any fix (§11.4.102 / §11.4.14 research-doc tree exists under docs/research/ fully_incorporate_the_helixskills_system_into_ every enumerated weak-spot/gap/danger-zone implementation plan down to lines-of-code exist are created, and the whole effort is validated p ADDENDUM — BINDING OPERATOR REQUIRE == VERBATIM: “Make sure that SpecKit and S and with all extensons we have created deriviv not optional colour on the SpecKit mandate — work MUST compose with rather than duplica BRIDGE EXTENSION (verified to exist, 2026-07- speckit_superpowers/implementation/ — the “ Superpowers Bridge”. Eight documents, each i IMPLEMENTATION_PLAN, NANO_TASK_ENGIN TDD_INTEGRATION, CONSTITUTION_INTEGRA SEVEN-LAYER architecture: 1 Developer Work SuperSpec bridge (orchestration — github.com (discipline enforcement — speckit-community MCP (execution) 6 Helix LLM (inference — gith Distributed host cluster (llama.cpp RPC) CONST MUST BE IN SCOPE (inherited BY REFERENCE diverges silently): constitution/skills/ (7) action reporting-workable-items, scheduled-work-qu constitution/mcp/ (2) media-validator-mcp.json plugins/ (2) helix, scheduled-work constitution system), subagent_tiering.yaml constitution/sc credential_scan_lib.sh, guard-branch-consisten guard-track-branch-label.sh, guard-work-track LOCALLY — do not re-create: .claude/skills/ 10 clarify, constitution, converge, implement, plan workflows/speckit/workflow.yml, integrations/ templates/, memory/ ADDED ACCEPTANCE CRI extension above has an explicit disposition — reason. Silence is not an answer (§11.4.118). 9. WIRED vs DESIGNED-ONLY with a captured pr evidence a layer is installed — that is precisely and reporting a documented layer as an availa namespace collision risk is addressed with a d skills system, while §11.4.164’s post_update_ho into .claude/skills/. Two systems registering int

ATM ID	Type	Status	Severity	Description
				and clash handling, not discovery at runtime. The non-own-org dependencies (\$11.4.74 vendor p supply-chain exposure must be assessed, not a them. The script that brings up the supporting service directly, rather than going through the shared maintains for exactly this purpose. That comp the same way, gets the same health checking, a repeatedly. Bypassing it means this project can and has already caused one outage of its own. service startups still hide their errors, so a fail was flagged honestly in a commit message, but obligation anyone will act on, which is why it i routing the startup through the shared compo hiding so a failed start is visible. When publishing the agent component, the ho outstanding security advisories against the lib. 204. That figure turns out to be unreliable in th drifts upward on its own as new advisories are 205 the next day, and 210 on a live re-query on be “the” count. It over-states the work by abou counted once for every dependency file it appe distinct advisories, and the five criticals collap two places this code is published; the other cop was never switched on there. The recorded br critical, 88 high, 99 moderate and 18 low. A ful completed, and it changes the picture substant the affected code can actually be reached in so them currently qualify. The Go vulnerability so actually release, reports our code is affected by against a container library are reachable only in the codebase refers to and that never enters remaining alerts, including three of the four cr third-party project whose container definition cannot start today. Most of the rest belong to a builds. Exactly one alert touches a component defence that advisory asks for has already bee alert stays open only because the dependency related piece of housekeeping was resolved wh files of downloaded third-party library code, in behind one critical advisory, had been commit removed from the tracked files, though they re
HXC-164	Task	Queued	Medium	
HXC-166	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
HXC-168	Bug	Operator-blocked	High	benefit of finishing this work is that the alarm is removed from the product, and we regain the ability to say how the project carries a reachable known vulnerability. This is a high acceptance criterion has not been done: the dependencies of the findings have not yet been upgraded or removed. We will retire the copied-in project that carries most of the work. The password used to reach the project database is stored in the setup and container files, and those files are present in the accounts. Anyone who reads the project can read the password, not introduced by recent work; the project's own password is it as something that must be replaced, and that is a high priority depends on whether the same password is used in the project has not been established either way and should be changed. The value is already public, changing the files alone is not enough; it be changed everywhere it is used, and the setup files should be a configuration source rather than containing it. The project cannot honestly claim its credentials are protected.
HXC-171	Bug	Queued	Medium	A complete copy of somebody else's web application is in the project; it alone accounts for one hundred and thirty-nine warnings, including three dependency security warnings, including three warnings that should be built or run: the startup configuration point is in the project, that folder, and its supporting packages have no warnings. The two opposite ways at once. Because it is not built, the warnings we actually run today, which is why they rank high. The references it, the setup is broken and misleading. The recipe file, all one hundred and thirty-nine warnings, is a warning to whoever does it. Anyone reading our security report is a distortion, since it currently buries the small number of warnings. The expected outcome is an explicit decision, not a properly maintain and build it, or remove it and replace it with a reference to the upstream hosted service. Before the project approval, the decision must be put to the operator.
HXC-172	Task	Operator-blocked	High	A small companion service we wrote ourselves is in the project, part of the standard deployment, and it carries one hundred and thirty-nine warnings, two of them rated high. The warnings, these have no mitigating evidence in the project, unshipped, because it demonstrably ships, and the project, whether the affected functionality is ever exercised, is a high place in the whole review where something with a high priority which makes it far more deserving of attention. The project, three is not really an upgrade problem at all: it is a high priority switched off unless explicitly enabled, so updates are not a high priority.

ATM ID	Type	Status	Severity	Description
HXC-175	Bug	Queued	Low	protection on would leave the exposure in place until the deployment running this service benefit. The error assessment of the three warnings, the two strategies and the protective setting explicitly verified as enabled and the service rebuilt and confirmed working. The cache that remembers model information is located where its clock was never configured. That falls short as the only way to build the cache always configured and unexercised. On its own that is harmless. The risk is if someone later edits it: a plausible-looking change to the current time would make every cached entry a failure in the past that the freshness check reads as a failure at expiry limit. Every entry would then be considered invalid and keep serving outdated model information indefinitely. The choice is to either add a small test that pins the cache branch entirely and document that there is one. When a caller presents an API key, the server does an ordinary text comparison. That kind of comparison is one that does not match, so a key sharing the first few characters is more reject than one that differs immediately. An attacker can exploit differences across many attempts can recover the key by guessing it outright. Network timing noise makes this attack is low rather than urgent, but it is a genuine weakness. A cheap: use the comparison function designed for timing attacks amount of time regardless of where the difference is introduced by recent work.
HXC-178	Bug	Queued	Low	The check that decides which websites may open connections deliberately allows requests that carry no origin information, because only browsers send that information and legitimately omit it. It is safe today only because it prevents away anyone without valid credentials. The two strategies nothing records or enforces that coupling: if someone logs in check while working on something unrelated, the information would become completely ungated. The work here is to add a check that specifically covers the connection upgrade, so the dependency is removed.
HXC-179	Task	Queued	Medium	Three commits in the current release range do not refer to a piece of data that was not added in the previous those three points gets a build error. The tip of the branch corrected shortly afterwards. This cannot be retracted without rewriting published history, which the project policy is for anyone hunting a regression by stepping back.
HXC-180	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				three points that fail to build for a reason unrevealed. The error may wrongly conclude the fault lies there. The error affects affected points and their shared cause somewhat, so the failure is recognised immediately as known and fixed from scratch. HXC-160 corrected the 21 stale references to scripts/generate_fixed_summary.sh across the codebase. It added the CM-SUPERSEDED-GENERATOR-NAMING to the sets were deliberately left untouched and are the same. Constitution.md still carries 14 references across the codebase. universal rules that projects other than HelixCode can't record such as the Phase 5 Go-binary-shim plan. The plan is a governance amendment that CONST-049 requires to be configured upstream, then re-cascaded and pushed to a session that closed HXC-160, and a canon amendment to drift than the stale wording, so the canon was updated. 45 owned submodules carry cascaded restatement. The correction applied in the parent repo is deliberate. because the fix names a HelixCode-specific mechanism. The wording into decoupled, project-agnostic submodules. they don't them and violate CONST-051(B). The correct set of neutral canonical wording that names the work. and push it in the constitution submodule per the plan. fleet and bump pointers. Who benefits: every contributor. the anchor text, and any agent reading a submodule. summaries are produced. Acceptance: canon remains. wording, the change is pushed to every constitution. are re-cascaded, submodule pointers are bumped. SUPERSEDED-GENERATOR-NAMING both pass. The main handbook that tells contributors and contributors organised lists several folders at the top level to when only one is present, and refers to a top-level. Anyone following it — a new contributor, or an existing orientation — will look for things that are not broken or that they have checked out the wrong of the document, which is otherwise the author's. The correction is small and purely descriptive. section to match, then keep the two in step as the plan.
HXC-181	Task	Queued	Medium	
HXC-188	Task	Queued	Low	One of the external projects this codebase depends on configuration still refers to it by its previous name. The hosting service silently forwards requests from the old forwarding is a courtesy, not a guarantee: it doesn't.
HXC-189	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				project under the old name, and it is not honoring the dependency. If it stops working, fetching the project name nobody recognises, which is unusually common for a project's current canonical name so the reference on a redirect continuing to exist. When a running command is cut short because of a timeout, it says it was not a timeout. The check that sets the caller's own deadline, which usually has not been reached, stopped the command. Investigation showed that the command either, because the internal signal it would continue to run, cancellation and never a timeout, and the timer was not a mechanism entirely. The effect is that anyone can get a failure and no indication that a time limit was reached in the wrong place. A sibling code path that runs correctly by doing it inside the timer's own callback copy.
HXC-190	Bug	Queued	Low	
HXC-191	Bug	Queued	Low	When a command is cancelled or times out, the command is not but anything it started, by signalling the whole command works if the command was placed into its own sandbox enabled — the normal configuration — the sandbox switched off, the grouping step is skipped. The operating system to signal a group that does not have an error is discarded. The result is a cancelled command that is dormant, because the only configuration that could have so no shipped path reaches it. It is worth fixing this, but would appear only in an unusual configuration and should be diagnose later.
HXC-192	Bug	Queued	Low	Output from a running command is read line by line. A line may be. When a line exceeds that limit the command is permanently for that stream, because the buffer is past it. Nothing else takes over the reading. The command is until the operating system's own small buffer fills up, someone to read. So one unusually long line, which is readable output produce quite naturally, can wait until it is to keep draining and discarding the remainder of the stream, abandoning the stream, so the producing command is not
HXC-196	Task	Queued	Medium	The handbook that describes this project's technology is one of its networking libraries. Both parts of the project because a security fix required moving off the old code deliberate; it is the handbook that is now wrong. A new contributor or an automated helper reader will confidently state a version that has not been

ATM ID	Type	Status	Severity	Description
HXC-197	Bug	Queued	Low	same document has already been found describing the same additional inaccuracy makes the whole document inconsistent with the recorded version to what is actually in use, and for this reason, so nobody later reverts it believing the document is correct. One test directory contains a duplicate of its own files. Two files appear twice in the project at two different locations. A renaming exercise that moved a tree while leaving the old one the new one. Nothing reads the nested copy, so when anyone searching the project finds two results, one is the real one which is real, and any tool that scans for duplicate files blocks a useful check from being switched on: the tool that finds duplicates cannot be enabled while this one exists. It runs every run. Removing it requires first confirming that the nested copy is genuinely the leftover and not the original.
HXC-200	Bug	Queued	Low	A detector was written to catch a specific kind of error where a worker stuck trying to hand off a result nobody is collecting. It has one particular description of what the worker is doing: a form of the same stall — where the worker was waiting for a handoff rather than a single handoff — and the detector caught it. It demonstrated rather than theorised: a deliberate test was run and went completely undetected while the detector was active. What it does catch is the historical one, so the guard is set to look for the newer structure while breaking its escape route. A detector that recognise both descriptions is a small change and a good one. A defect rather than half of it.
HXC-204	Bug	Queued	Medium	Two long-running tests — one exercising heavy load, the other one exercising leader election among worker nodes — both report failure when the machine is under load. The first test is its own. Measured: during a full test run both fail about 10% of the times consecutively, one finishing in roughly a minute and the other about a twentieth. Their verdicts are therefore not reliable for whether the product works. This matters most when a test run performed before a release is by definition a test run. These two will mark a healthy release as broken. A third test with the same problem was run and reported an inconclusive run report as skipped rather than failed. The same treatment fits here. Both should wait for a result rather than assuming a fixed time is enough.
HXC-206	Bug	Queued	Low	The newly-fixed health refresh deliberately refreshes the health provider over the network, so that a slow or hanging refresh. It then re-acquires the lock and only applies its changes if the changed underneath it. There is a narrow remnant of the old

ATM ID	Type	Status	Severity	Description
				started again entirely within one probe, the change is applied as if it were a new probe. The model provider applies a verdict formed against the previous probe, which is briefly reported unhealthy when it is fine, which is a dangerous direction — reporting a stopped probe as a new probe. Worth recording that this is cheap to close rather than expensive. The incremented on each state change is invisible to any published structure. An earlier note describing this was wrong and has been retracted. When a model provider record is created, its evidence is shared template by reference rather than copied. The record points at the same underlying settings, including the model. Nothing writes to those settings today, and the model is not corrupted. No corruption occurs at present — this was very rare. The record it is that the protection now added to the model manager cannot protect a structure shared with the model. A legitimate, properly-locked write to a provider record is a manager's view as well, and the locking will lock the record when the record is created removes the hazard. Four compose entries across two end-to-end tests, one in the Slack mock and the LLM-provider mock, both of which fail. Any attempt to build those stacks therefore fails, which means the containerised form of these tests is not exercised. This was surfaced while fixing an unrelated issue. The same two services: the fix could be proven at the time, but against a running container, precisely because the fix was produced. That gap matters beyond these two tests, as the process leaves the deployed-artifact layer unproven. The failure the four-layer verification discipline excludes the Dockerfiles, or to remove the build entries if the tests are run containerised, and then to prove the choice of the fix. A blanket rule excluding log files from version control is evidence, because evidence transcripts are written to the log, corrected so this cannot happen again, and the logs are then have been recovered and committed. Thirty-five runs exist only on the machine that produced the evidence, an evidence folder that appears empty, which is never run. Recovering them is not automatic, but it is a seventy-eight megabytes and committing those logs is a task that deserves a deliberate decision rather than a reflexive decision to keep what genuinely substantiates a closed item, and to keep ones rather than storing them whole, and recording the reason why, so the gap is visible rather than silent.
HXC-207	Bug	Queued	Low	
HXC-208	Bug	Queued	Low	
HXC-216	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
HXC-219	Task	Queued	Medium	A recent repair changed how the release gate of evidence really belongs to the change it claims to commit on a labelled line rather than merely having somewhere in the text. That new requirement have to follow it: the guide that authors read before the new labelled field, so a contributor following evidence the gate will reject without understanding describing the gate itself has printable and we page was last edited, so all three versions disagree describe. Fixing this means updating the author's citation field and regenerating the exported code who benefit are contributors recording evidence have no accurate written description of the rule can follow the written guide and produce evidence attempt. When we mark a piece of work finished, we reworks. Nobody ever checked that those pointers different kinds of rot accumulated silently and report. First, some pointers hold a whole paragraph there is no way for a machine to follow them. Second never created because the person doing the work means the proof exists but nobody can find it for ordinary progress updates rather than to the author notes look like formal proof. The effect is that a backed in every summary and dashboard which way anyone finds out is by manually hunting for make the completion pointer a checked field: if it exists, it must be attached only to real completion must be refused at the moment someone tries later. This protects everyone who relies on our claim of proof that cannot be produced on demand.
HXC-224	Bug	Queued	High	A single bundled helper component inside the roughly 208 security advisories reported against the total — and it entered the codebase without established what it does, whether we wrote it or actually uses it, or whether removing it would answered the headline advisory count is misleading carries enormous risk, when in reality the great not even be part of what we deliver. The work purpose, whether it is reachable from anything, decision about it — keep and maintain, upgrade record which was chosen and why. The decision
HXC-225	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
HXC-226	Bug	Queued	Medium	advisories, because upgrading dependencies in effort spent on nothing.
				The agent project is mirrored to two hosting and was only ever switched on for one of them. Qu hundred and ten open alerts while the other r one the main project is configured to pull from of the component we actually consume is show hundred findings sit unseen on the sibling cop the two mirrors hold the same commits. What watched. This is dangerous precisely because t findings is indistinguishable from a report of z the second case is what we have. The work is t from, confirm both then report the same figur the two so a future divergence is noticed rather statement about this component's vulnerability from.
				A design document committed forty-eight days appears to be a genuine access key for an exte placeholder. The section is headed as being co and described as already configured, and the v placeholder would have: it is fifty-one character characters, and contains no example or to-be-f the main branch and has therefore been publi day it was written. Anyone with read access to fork made at any point in those forty-eight day now would remove it from the current revisio history is forbidden, and rewriting would in ar The only action that actually withdraws the ke replacement, which requires access to the pro within the repository. Until that happens the k revoked, the document should be edited to ref of quoting its contents, and a check added that this shape.
HXC-227	Bug	Queued	—	
HXC-231	Bug	Queued	—	The HelixLLM gateway downloads large AI mo starts accepting network connections. During t the service as running and healthy, but every a on 2026-08-06 this window lasted thirty-seven actually began listening. Anyone checking the everything is fine while the service is in fact u depends on the gateway will fail for the whole because the download depends on an external one of the two downloads timed out after roug

ATM ID	Type	Status	Severity	Description
HXC-232	Bug	Queued	—	that site were unreachable the gateway might be accepting connections first and download mode might be in an honest not-ready state while the downloads are in progress. HelixAgent expects two tables to exist in its database: <code>agent_instances</code>, but neither of them has ever been mentioned anywhere in the project that creates them. As the database locks once every minute, the database refuses to accept writes and an error is written to the log. Verified on 2023-08-24 directly: neither table exists in any of them, and in a one-minute period plus a further hundred and thirty seconds, Helix otherwise runs normally, so nothing visibly broken. Helix buries genuine problems in the log and means that the support to support is silently doing nothing at all. Fixing this by creating these tables were intended to contain to write to the database is guessed.
HXC-234	Bug	Queued	—	On startup HelixAgent tries to build and launch containers known as MCP servers, using the container engine. If the containers inside the container fail during their dependency installation, the operation is abandoned. HelixAgent records a list of these servers and then continues running normally. Nothing obviously breaks. The real effect is that the servers were meant to provide is silently missing, and this is not easy for nobody to notice. Observed live on 2023-08-24. The build was attempted twice and failed both times. A large container image each time, which also means that repairing the dependency installation inside the container is slow.
HXC-235	Bug	Queued	—	The HelixLLM gateway provides a feature that allows for exact wording. On startup it reports that no relevant purpose, so it has fallen back to a simple hashing strategy. It states plainly that this fallback does not capture the full meaning and will be significantly degraded. The important part is that it will off or return an error: it keeps answering even if the question is effectively arbitrary rather than relevant. Any relevant information from the responses that it is not working properly is lost on full platform boot. Resolving this means pointing to a genuine embedding provider, and deciding what level of degradation is acceptable behaviour or whether the feature should be disabled.
HXC-236	Bug	Queued	—	The standard library routine that splits a web address into differently in two parts of this codebase, using the same logic. In the main application it rejects an address whose host is containers component it quietly accepts the same address with a colon. The cause is that each component declares its own

ATM ID	Type	Status	Severity	Description
				the newer version tightened this routine; the c one, so it keeps the older, permissive behaviour reasoning about address handling is only valid which has already caused one incorrect conclusion. Second, and more seriously, the day someone releases a newer language version, addresses it accepts that failure will appear as unreachable services rather than change. The work is to make the difference explicit, deliberately which behaviour each component will find it, and add a check that fails if not noticing. The safety check that stops half-finished experiments from an exception for captured proof files, which would satisfy the old exception rule and so could not be narrow: the file must sit in the evidence folder and be saved as non-runnable, and not begin with the word Together those make the file inert in normal use, deliberately pointing a program interpreter at a file which no file property can stop. The practical result is rather than an accident, and the alternative would be proof it is required to keep. Recording it so the file is not absolute, and so anyone widening it later. When HelixQA runs a test bank that needs to sign in, it sends a sign-in reply, so every test in that bank that needs a MECHANISM (the original wording of this item) will fail, as the reply carries the real access pass at the top level, a separate bookkeeping reference for the session object. The runner looked for the bookkeeping object, found nothing there, and stopped with a plain completion. It never handed the server a wrong value and was rejected, which failed before making a single signed-in request. The runner then handed the wrong value and being rejected, which is not what that text were nonetheless correct. A SECOND step had its own copy of the same assumption, but it could not find the pass it simply carried on with the value unsigned and the resulting rejections were recorded as failures rather than as a failure to sign in — which is what was intended. THE OBVIOUS FIX WAS WRONG: the bookkeeping object is a different family of about fourteen banks that the runner uses, and the access pass genuinely does sit at the top level under a different name. The names would have repaired one product by bringing the existing name first and adds fallbacks tried.
HXC-238	Bug	Queued	Low	
HXC-239	Bug	Ready for testing	—	

ATM ID	Type	Status	Severity	Description
				support for reaching a nested location, and rep HOW TO REPRODUCE: run a bank requiring si run stop with a missing-field complaint before OUTCOME: the bank used for measurement w passing and two failing. The two that remain a work: one posts an empty body where the serv with a leftover record created by the very mea original figures, because that bank does not cle the originally reported nine-and-one is no long banks needing sign-in authenticate against bot changes; a failure to find a usable pass is repor service faults; no credential value ever appear proves the behaviour, shown real by a paired
HXC-240	Bug	Queued	—	One of the HelixCode test banks signs up a nev uses the very same account name, 'hxc_gen_e2 succeeds and the check passes. Every run after the database, the server correctly refuses with recorded as a failure forever after. This was ob run of the bank passed that step, and a second visible in the database timestamped to the mor works under require that automated checks ca the same answer, cleaning up after themselves quietly poisons its own results, so a permanen do with whether the product works. The fix is each run, or to remove the account it created v
HXC-241	Bug	Queued	—	The checks that confirm HelixCode can genera inside the server's reply, and that comparison Two checks in the generate-e2e bank ask for w with different capitalisation, so they are mark behaved correctly. The most serious case asks reply for 'hello' in small letters; the AI answer working feature was reported as broken. This server: asking it to reply with an invented wor it what seventeen times three is returned fifty- which proves real AI generation is running an reply for 'authorization' in small letters while capital A. Expecting exact capitalisation of free nature. The consequence is false alarms that h time.
HXC-242	Bug	Queued	—	Five checks covering the built-in screenshot ar against the running HelixCode server, and eve server answers 'QA engine is disabled' and rep

ATM ID	Type	Status	Severity	Description
				server is being honest here rather than misbehaving. The whole advertised area of the product is switched off and cannot be used by anyone or confirmed as working. Listing the available screenshot engines, listing the status of a session, starting a new session, and running with nothing to run. Because the feature is off, it's unclear whether the underlying code behind it functions. The feature is meant to be available, and if so switched on, the behaviour is actually proven; if it is meant to be switched off, reporting failures. Several of the automated test collections that check for errors about what a correct answer looks like. They succeed or fail and record a pass regardless of the reply. This particular test collection was deliberately aimed at the wrong thing: an error for every request, and it reported both passes and failures against a service known to be broken. It carried one, so its green result carries no information. It was a failing test, because a suite that cannot fail is wrong, even where there is none, and it will keep producing failures on every check an explicit expectation — the status of the answer must contain — and then prove each of them wrong. Something known to be wrong and confirming it in the same run: four apparent failures in the work item harness looking for the wrong field name in the response, and that mismatch should be corrected so real failures are reported. One blocked work item (HXC-172) has its explanation. It has that same explanation written into its own description to agree. The tracker rebuilds those stored explanations each time it regenerates its records, so an explanation that the text is discarded the next time the tracker regenerates happens today: a regenerate-and-reread cycle. The tracker then correctly reports the item as blocked-with-explanation, that explanation is what makes a blocked item blocked, what was already tried, and which decision was made. It's a decision waiting on a person into an item nobody can move, keeping the context needed to make the blocking decision, why the item stalled. The expected outcome is to write the explanation again, the stored and written copies are identical, preserves it — proven by the tracker's validation. A complete run of the agent runtime's automated tests found 12 groups failing, and 74 individual test failures, more than assumed, because the early evidence pointed to a specific
HXC-243	Bug	Queued	—	
HXC-245	Bug	Queued	High	
HXC-246	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				than at the product. Three patterns account for the in-memory cache service because they look like the container while this machine publishes it on a running service to answer a web request and fail to answer correctly when asked by hand; the tests against fixed limits and overshoot them by about 10 jobs and a service rebuild were competing for the product is healthy, and none of it proves the product parallel investigations are now establishing, each rather than a plausible story. This matters because the reasons is as damaging as one that passes while everyone to stop trusting the result, and both the operator benefits from a suite whose red means a signal they can act on. The expected outcome is evidence as a genuine product defect, a test defect, contention artifact, each real defect fixed with returning to a trustworthy state. Two separate parts of the system have both been and neither knows about the other. The agent decides which service gets which port number for itself. But a different component, the model verifier names port 8100, and because the verifier starts the agent runtime is launched from a settings file the agent never even tries to claim the one it was allocated port 8100 for the agent runtime and are answered recognise the requests and rejects them with a runtime as broken when it is running perfectly because it makes a large part of the test suite successful names, and because the same collision would not be real. Whoever operates or deploys this benefits from failing, and a genuinely failing one would look like a single place that decides port allocation for everyone claiming the same number, the agent runtime is failing and the tests then passing or failing for reasons. Later evidence settles which of the two sides is correct deployment is not misconfigured at all. The model file that starts it on this machine - documents it as an agreed meeting point between the verifier and the gateway looks there by default, and, in as many cases runs on the other address here. The choice was made down at deployment time, in the verifier's favor, clearer than first described: the agent runtime
HXC-247	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				date, still claiming an address that a different, with a note that the old address was abandoned still in use. A contributing cause is that the version so the one mechanism meant to prevent two components not arbitrate a component it had never been to unchanged - the tests inherit the stale claim and the repair direction is now settled rather than the address it genuinely serves, give the verification recur, and point the test defaults at the real address establishes which side is stale and why; it does safe to execute, because that entry is consumed change still needs its impact review before it lands. A cleanup step in the automated test suite can be the only thing preventing that is a coincidence to stop any containers they started. To avoid stopping first check whether the agent service is still running <i>something</i> answers on network port 8100, with On this machine port 8100 is held by a different the check says yes, and the shutdown is skipped different port and is never actually consulted. late, or moved, that same check would say no a live platform out from under whoever was using also untrue: it reports that the agent is still running all. This matters to anyone running the test suite deployment, which is the normal case here. The the identity of what is answering - not merely makes its decision on the real condition rather logs is true. When a user asks the system to write out a real command-line coding assistant, the file it produces service. The generator fills in a network address than to the agent runtime the assistant is supplying requests to a component that does not understand reply. From the user's point of view the assistant error explaining that the address in their branch was written. Of the forty-eight assistants the problem affected; a fix landed earlier covers only two of all the rest still writes the wrong address today in the current set, because it is not an internal handed to a user with an instruction to use it. For newcomers the feature exists to help: someone following the documented path, and getting a message
HXC-248	Bug	Queued	High	
HXC-250	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
HXC-251	Bug	Queued	Medium	The expected outcome is that every generated runtime genuinely serves, that all forty-eight d than two, and that a generated config works w without editing anything by hand.
				One of our own services could not start while both had been told to use the same network ac conventional default that the wider industry u precisely what makes it the number most likel own machine it was taken, by a supporting da The result was that the service refused to start client dialling that same well-known address r which completed a perfectly healthy-looking c Because the wrong service was alive rather tha something is unreachable did not skip: they ra our service, which had never been running at twenty test failures that had been attributed to service in the project's central address book so never silently collide again, and it now fails im rather than starting somewhere unexpected. V misleading: the startup message printed for th so anyone reading the log is told the wrong pla platform benefits, and the expected outcome is the rest of the stack and that every message it bound.
				Tests in one area of the codebase interfere wit share a single piece of network machinery, so to do with the code it is testing. When each cor given its own timeout setting but no connectio library quietly falls back to one shared handle in this area start roughly a hundred and forty- around five hundred cases to run simultaneou down at the end of its own test, politely tells th and because the handler is shared, it also sever of using. The result is a failure that appears in and vanishes when the tests are run one at a ti the test setup rather than in the product. It rep one run in ten fails this way. Developers are th unpredictably teaches everyone to re-run it ra comfortably inside that noise. The expected ou own connection handler so shutting one prete test, and the suite becomes trustworthy enough genuinely wrong.
HXC-254	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
HXC-255	Task	Queued	Medium	Two things are left over in the scripts that run needs a decision from the operator rather than problem: the end-of-run summary prints “Inter whenever the log file merely exists, regardless which tests genuinely failed still shows a reass most likely to read. The second is deliberate an written to note test failures as warnings and ca status even when real tests have failed, which blocks a release. That behaviour was chosen o abort a long run and lose the rest of the inform between gathering everything in one pass and broken. Because it is a deliberate posture rather operator decision about how the project wants be altered silently by whoever happens to touc runs to judge release readiness benefits from s only that the script reached the end. The expect reflects the real result rather than the presenc explicit, recorded choice about whether these The tool that regenerates the project’s tracker happens to be run from unless it is explicitly to one optional setting quietly produces a second while the real ones stay out of date. The tool’s optional, which is what invites the omission, a folder rather than the folder where this projec found the hard way during routine work: runn setting created sixteen stray copies at the top o one directory down, kept their old contents. Th someone can look at freshly written files, see t records are current when the versions under v the tracker documents by hand is exposed, wh automated helpers working outside the standa bounded and that bounds the severity: every v destination explicitly, the tool prints the full pa visible to a careful reader, the stray copies sho separate consistency check would eventually r What is missing is the tool itself objecting. The documents without naming a destination eithe or refuses and says so, instead of succeeding s caught this does not exist today: the existing g supplies the destination explicitly, so it can nev out - adding that second case is the specified fo
HXC-256	Bug	Queued	Low	
HXC-257	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
HXC-258	Bug	Queued	Medium	The automated check that guards our workable-items database actually measures, so a reader who trusts its workable-items database computed. Internally it names the thing it inspects. The written description states that it proves the copy is correct. It never consults version control at all - it simply compares the folder, which may contain changes that have not yet been committed. The genuinely proves is narrower and still useful: the two copies right now agree with each other. The gap was not noticed because the check reported success while the database was not yet been committed - exactly the situation its workable-items database stating plainly is that the underlying behaviour is not correct. The point of the pre-release sweep, where the point is to ensure the commit is self-consistent, so inspecting the workable-items database there is a documented reason for the copy - only as a side effect. On that reading the repair is the only way to fix it. Everyone who reads that check's verdict before the difference between "what is on your disk is correct" is exactly the difference that matters with the workable-items name has already spread: the description of the workable-items database repeats the same claim. The expected outcome of the workable-items and the sentence it prints all describe only what the workable-items guarantee is wanted it is added deliberately as a side effect. The test that would have caught this discrepancy is the disagreement between the documents and the workable-items database between the committed copy and the working copy. The fix is adding that case is the specified follow-up. Our workable-items database keeps a small boolean flag for synchronisation ran between the database and documents. That note is only ever written by one of the two paths. The documents into the database - and the value it is the value that actually ran. The path that goes the other way, from the database, never updates the note at all. The path that regenerates the note still describes an older run. The path that reads it to find out what happened most recently is the path that matters because deciding the direction of a sync can destroy records if it is made backwards, and the path that judgement while being incapable of supporting the path into-database run dated the eleventh of July was the path that database-into-documents regeneration on the path that shared tooling that performs the synchronisation. The path that genuinely ran on both paths rather than assumed the path on hand would be undone by the next run and is the path that

ATM ID	Type	Status	Severity	Description
HXC-259	Task	Queued	Low	Every workable item is supposed to accumulate its status changed - opened, reopened, closed - what evidence justified it. A recent repair note no such trail. Measurement across the whole d to those two: eighty-one distinct records cover history entries at all, and all but one of them a filed away. In other words the majority of our account of how it reached that state, which we understand why something was reopened, or t cause is that items created by direct database v creation command skip the step that opens the before the trail existed never had one. This is c filling: inventing dates and authors for events trail that is worse than an absent one because instead is a decision on how to mark these rec that prevents any newly created item from join The Kubernetes manifests for HelixAgent told 9090, and they set an environment variable ca was true. No Go source in the repository reads the port the server listens on; this was proven bound :8112 while the real variable HELIXAGE 9090 is a port no HelixAgent process ever bind internal/ports registry (HelixAgentGRPC, core l by running the built cmd/grpc-server binary a The practical effect is that anyone deploying H manifests would get a Deployment exposing co gRPC traffic to targetPort 9090, and a Network artifacts agreeing with each other and all poin port the server would actually bind was blocke number because the manifest LOOKED configu and edited GRPC_PORT would see no change w This item covers correcting all four manifests t service.yaml, networkpolicy.yaml, configmap.y the dead variable with the one the server hono limitation: the image this Deployment runs is b no gRPC server at all (that is a separate binary, correct but INERT until a grpc-server containe corrected manifests make a one-line change in checking out the pre-fix tree and running the p RED_MODE=1, which asserts the dead variable Acceptance criteria: internal/ports/published_g
HXC-260	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
HXC-261	Bug	Queued	Critical	unset (asserting every manifest names the reg outside a comment) and fails on the pre-fix tre The challenge script that is supposed to prove structurally incapable of failing, and it has bee existed. Two independent faults compounded. environment variable named HELIXAGENT_G spelling that no Go source reads (the server ho number belonging to no service in the project: registry's 8112, not the HTTP port 7061. Second recorded a successful assertion whether the po check recorded true when the port was listeni when it was not; the unary-call, streaming-call The two faults together mean the challenge dia ever bound, found nothing there, and then cer most damaging class of defect in the codebase — it is an automated test actively producing fa anti-bluff covenant exists to prevent, and it ma reading the suite. It also violated the challenge framework already ships a record_skip helper record_assertion with status true to mask a mi not used. This item covers resolving the port th default, consolidating the four duplicated prob interpreted in exactly one place, making abser OK ticket rather than success, making a real fa two-state final verdict (which reported PASSED when nothing had been verified) with a three- assertion could be verified. Reproduce with th protocol_grpc_fail_open_guard.sh run with RE and releases a port to prove it is free, drives th and confirms assertions are recorded PASSED guard with RED_MODE=0 reports zero PASSED every probe SKIPPED and SKIP-OK tagged, and PASSED assertion when a real gRPC server is li The project keeps a battery of tests whose who catch problems, by deliberately breaking thing assertion in that battery — the one covering a occasionally reports a problem when nothing consecutive runs, and ran correctly six times o intermittent rather than consistently broken. T check itself: it asks the system two separate qu which process is running it, and between those move on, so the two answers describe differen
HXC-262	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				happens the check cannot find a process to examine anything. The practical harm is not a missed detection of a condition that is not present, and worse, people might think it, which quietly erodes the value of the one bad check. The fix is to read both pieces of information in the same instant. This was deliberately left unfixed because it changes the sampling behaviour of a check that is not such a change deserves its own review rather than being a side effect. The project keeps a folder of tests whose purpose is to genuinely work, by deliberately introducing a failure. A check reports a failure. A test like that only has one purpose: Counting the invocations at the current revision of the folder are ever executed by anything; the remainder of the folder never run. This is the same problem the project encountered when three safety checks turned out to exist but not be the same reason: a check nobody runs cannot be used to provide coverage. The presence of these files is a problem: counting test files sees protection that is not being provided. It looks better than it is — five of the files are merely comments, an entry in a skip list, and none of which run anything. Resolving this means either to wire it into the release sweep so it runs, or to make the choice visible rather than leaving the file sitting there to deliver. A single entry in the project's work tracker makes a mistake in different ways in three different fields, and the result says a required program is not installed on the machine but exists but cannot be found because of where the file exists but such file exists anywhere on the system. Those are three different fixes, so a reader cannot tell from the text what also proposes a remedy, giving the service a different name, not what was eventually done: the capability was added to a different service that was already running, and the file was absent from the machine today. This matters because it is a mix of what was reported and why, and an entry that is not for the purpose for anyone reviewing the history later. The entry is out of step with the database, which is a problem. The rendered documents faithfully reproduce what was in the database is inside the database entry itself. Resolving it means to get an accurate account that records both what was reported and what was resolved, without erasing the original report.
HXC-263	Bug	Queued	Medium	
HXC-264	Bug	Queued	Low	
HXC-265	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				HelixAgent keeps 36 Go modules nested inside. Most of them are never compiled by anything — a developer can rename a package, change a source file, and every build and test still reports green — the test. That silent-failure surface is what this item tracks. The documentation samples, because the project authors' documentation example works when copy-pasted, but the test cannot honour that promise. WHO IS AFFECTED? Readers who copy a documentation example expecting it to be verifiable by one command (the quotes are loose, but the pattern before git sees it and reports 2 instances: <code>git ls-files '*go.mod' \ wc -l</code> reports 36 modules: 'challenges/codebase/go_files/*/go.mod' \ wc -l' — the challenge-modules recipe walks. By directory tree: 2 under challenges/, 2 under plugins/, 2 under docker/, 2 under cli_agents/. THE CLASSIFICATION IS PER-MODULE: 25 modules: the 23 modules that are never compiled at all — nothing anywhere compiles them, the three challenge-modules only run presence checks (<code>test -f</code>, <code>grep -q</code>), and the two Dockerfiles because Go does not cross nested-module boundaries. GATED CONTAINER IMAGES — 2 modules: docker-images declared as build contexts in docker-compose, but not in their Dockerfiles, but no test suite or gate exercise. rebuilds only when the image is absent, so breaks builds deterministically rather than never. (3) FULLY BUILT: pkg/api: it is required and replaced in the root build main.go with no build tags, and built for four platforms. failures increment PHASE_FAILURES and set the test to pass and set and needs no work. Note for future readers: the test "does not contain package", which is a true output of the membership rather than whether the module is built. this item mistook the one for the other. (4) PARTIALLY BUILT: module, Toolkit: internal/verifier/startup.go imports Toolkit transitively pulls in three more Toolkit packages. compile on every root build while roughly 17 modules are EXCLUDED WITH A RECORDED REASON AND A REASON: the HXC-139 isolation marker whose header is never compiled, paired with the regression guard cli_agents_isolation_test.go. It already satisfies the test in terminal. ACCEPTANCE CRITERIA: every module in the remainder of (4) is either wired into a build or not.

ATM ID	Type	Status	Severity	Description
HXC-266	Bug	Queued	Medium	runs, or deliberately excluded with a recorded gate fails when a new nested module appears The function that builds the URL clients use to “http://%s:%d” without ever bracketing the host malformed URL. WHY THIS IS REACHABLE RAN layer deliberately treats the IPv6 loopback ::1 as a configuration that sets the host to ::1 passes even the formatter. HOW IT MANIFESTS: the emitted address is not parseable URL — the colons of the address and the host to any URL parser, so the client fails to connect with a clear configuration error. WHO IS AFFECTED: any client that uses IPv6-only or IPv6-preferred host, and any deployment that uses IPv6 literal instead of a name. HOW TO REPRODUCE: set the client base URL, and observe the returned address in the correct bracketed form http://[::1]:7061. ACCEPTED: IPv6 literals per RFC 3986 so ::1 yields http://[::1]:7061. A regression test covers IPv6 literal, bracketed IPv6, and hostname cases. When HelixQA signs in to run an API test bank but a reply the runner cannot read, the runner continues the sign-in, carries on making requests with no context, and faults in the API being tested rather than as its own. In the case where the reply IS readable but carries no context, this item covers the door that fix left open. WHY IT IS AN explanatory finding is nested inside the branch of the structured data, and that branch has no counter. WHY IT IS unreadable reply skips the reporting entirely. WHY IT IS tested service is broken when in fact the runner is not. WHY reads the report to investigate the wrong system. WHY this item originally described two silent shapes. WHY established by running every case against the test set. 10 out of eleven tested. Silent: an empty reply; an empty object; data; bytes that are not valid text; a bare quote; an object; a bare true or false; and a very large object. WHY item’s reporting: the literal null, and an empty object. WHY whether the reply PARSES, not whether it carries context. WHY narrower descriptions kept undercounting it. WHY HelixQA report for a service whose sign-in route is blocked. WHY error pages produced by a proxy or gateway in the response. REPRODUCE: point the pipeline at a test server that responds with a body of empty text or HTML, run the API test. WHY in finding is recorded while the later unauthenticated
HXC-267	Bug	Queued	Medium	When HelixQA signs in to run an API test bank but a reply the runner cannot read, the runner continues the sign-in, carries on making requests with no context, and faults in the API being tested rather than as its own. In the case where the reply IS readable but carries no context, this item covers the door that fix left open. WHY IT IS AN explanatory finding is nested inside the branch of the structured data, and that branch has no counter. WHY IT IS unreadable reply skips the reporting entirely. WHY IT IS tested service is broken when in fact the runner is not. WHY reads the report to investigate the wrong system. WHY this item originally described two silent shapes. WHY established by running every case against the test set. 10 out of eleven tested. Silent: an empty reply; an empty object; data; bytes that are not valid text; a bare quote; an object; a bare true or false; and a very large object. WHY item’s reporting: the literal null, and an empty object. WHY whether the reply PARSES, not whether it carries context. WHY narrower descriptions kept undercounting it. WHY HelixQA report for a service whose sign-in route is blocked. WHY error pages produced by a proxy or gateway in the response. REPRODUCE: point the pipeline at a test server that responds with a body of empty text or HTML, run the API test. WHY in finding is recorded while the later unauthenticated

ATM ID	Type	Status	Severity	Description
HXC-268	Bug	Queued	Medium	ACCEPTANCE CRITERIA: an unreadable sign-in finding as an unusable one, naming what was be read; the message never contains any byte c standing test covers all nine silent shapes plus paired mutation that removes the new reporting already-working shapes must produce byte-ids change cannot disturb the path that already re The model verifier builds web addresses by pa colon, without the square brackets that IPv6 ad host is an IPv6 address produces addresses no operators on IPv6-first hosts, and any deploym rather than a name. WHY IT IS REACHABLE: th with no validation, so an IPv6 value passes str MANIFESTS: the emitted address looks like htt read because the colons of the address are ind the port; the client then fails with a confusing configuration message, sending whoever debu PLACES, AND A CORRECTION WORTH RECORD contains FIFTEEN such build sites, one in each every one reachable from the single dispatch t revision of this item said three. That was wron fifteen sites use only three distinct format SHA plain, one ending in a messages segment), and upstream report were recorded as if they were places. Had the earlier figure stood, the accept with twelve sites still broken. HOW TO REPRO submodule, count occurrences of the formattin file; it reports fifteen, while listing the distinct FIX DIRECTION: do not hand-bracket fifteen pl submodule is introducing a shared endpoint h correctly using the standard library join, inclu YET part of any committed state of the submo nobody should expect to find it there today. Or it, which fixes them uniformly and prevents a Coordinate with that work before starting. SCC capability-config generator is not the whole su appears in at least nine further files elsewhere agent config writers and two unrelated subsys deliberately, but a fixer should know the capa entire problem. ACCEPTANCE CRITERIA: all fif produce properly bracketed addresses for IPv6 are unchanged; an already-bracketed value is

ATM ID	Type	Status	Severity	Description
				the IPv6 literal, already-bracketed, IPv4 and hostnames. A mutation that removes the bracketing and makes the hostnames unbracketed sites in that file is zero, verified by a test fifteen. When the model verifier builds a web address, if the address contains a question mark, a hash, or a slash, the address ends up pointing at the default port instead of the host. The address builder appends the port after whatever the address string produces something like a normal address. The query goes to the query rather than onto the host; every standard test fails. The correct, correctly, reads NO port at all, and treats the port as a host. MANIFESTS: the client connects to the wrong port, and the host is anywhere, which is the same silent-misdirection. The test eliminates. WHY IT IS REACHABLE RATHER THAN NOT: the first step accepts these shapes without complaint, so the host is an operator-facing host setting into generated code. The problem. THREE SHAPES CONFIRMED BY MEANING: a host with its port into the query; a host with a fragment of the path; and loses it into the path. WHO IS AFFECTED: any code that uses a host setting instead of the dedicated full-address. The guidance anticipates as a likely mistake. HOW TO FIX: a host containing a question mark and any port, or a host that is empty while the query carries the digits. ACTION: if the query, fragment, or path either is rejected with an unexpected shape, or is handled so the port survives. The test told rather than left with a silently wrong address. The test shapes plus the ordinary host and address cases. The test removes the handling and makes the test fail. HelixQA copies failed reply bodies straight into the test places, and one of those places is a sign-in reply body, most likely to be. WHAT HAPPENS: when a request is rejected, the code records the entire reply body as the description. The test THREE sites in the same validation routine — the test, and the media-search path. A fourth and fifth case is an error message carrying a slice of the body, and the test complains, which on the current toolchain produces a complaint. THIS MATTERS: these reports and errors are wrong. The automation, so anything the server put in that the test site is the highest risk by a wide margin — a sign-in token, returns a cookie in the body, or includes a value. The output would have that value copied verbatim. The test echo the tested service's own error bodies, which the test
HXC-269	Bug	Queued	Medium	
HXC-270	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				RELATIONSHIP TO OTHER WORK: a sibling item that it reports only a shape label, a byte count and a separate component now redacts values that likely predate that work and were deliberately left unclassified; the discipline is now inconsistent inside a single family; the failure path fifteen lines below it is not. HOW IT WAS CONFIRMED: a server whose sign-in route returns an unsuccessful response with a recognisable marker, run the validation step, add the recorded finding and in the returned error text the facade family emits raw body bytes; all report the same success path already uses; the decode-complaints family emits a shape rather than a quoted character; and a paired mutation that restores the echo and makes the facade family work. One of the two families of remote calls the age of the machine, however well configured. WHAT IS BEING DONE: ones for completing text, streaming a completion routine that is always handed an EMPTY list of providers immediately when the list is empty. So every call is successful regardless of how many providers are actually configured. A CONFIGURATION PROBLEM: the object service is not in the registry that holds the discovered providers; the facade family consults the registry into the OTHER family only, and the facade family consults it, which is why the two families behave differently from outside. HOW IT WAS CONFIRMED: capture the facade family plus real models had genuinely been discovered; the facade family in the case of an unprovisioned host — the providers were not configured; the facade family message saying no providers were configured. HOW IT WAS CONFIRMED: the facade family facade family, which is the family the publisher would naturally reach for. HOW IT WAS CONFIRMED: the facade family integrator would naturally reach for. HOW IT WAS CONFIRMED: the facade family them to use the provider registry for credentials; the facade family mistake on their side and sends them to check the facade family. ACCEPTANCE CRITERIA: the facade family either against configured providers, or is explicitly with a clear error naming its replacement; a test driver against a real configured provider and asserts of error; and a paired mutation that removes the facade family. Two automated checks that start in the same success path append to the same running log, so each is judged to be successful. WHAT HAPPENS: each run names its results folder after the second, with nothing added to make it unique, and the facade family outcome lines in that shared log to decide the outcome.
HXC-271	Bug	Queued	Critical	
HXC-272	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				BOTH WAYS, AND THE SECOND DIRECTION IS lines flip a healthy run red, which wastes inve flip a run that achieved nothing into a reporte success vector sitting in shared infrastructure uses, which places it in the same class as a che OBSERVED: a guard's positive case passed alon neighbour's failures; and two back-to-back run attempt, with the second run's results file repo plus the first run's four. A SECOND, RELATED I pointer meant to indicate the most recent run because the command used to update it follow and drops its contents inside the first run's folo folder's timestamp, so ordering the folders by . Measured live: the pointer still names the first runs existing beside it, with those later runs' p the by-recency listing therefore silently return reviewers into reading an old run as current. I within the same second and observe they shar one's verdict counting the first one's outcomes the recent-run pointer still names the first. AC share a results folder or log; the recent-run po recent run; a test starts two runs in the same s independent in BOTH directions, proven real b naming and makes the test fail. Several health checks treat any reply from a se that say the requested thing does not exist. WH address but do not ask the fetching tool to trea found reply is indistinguishable from a healthy places — two in the shared checking framewor HOW IT MANIFESTS: a service that is running on the right port with the wrong application er shape as a check that cannot fail: it consumes almost nothing. RELATIONSHIP TO OTHER WC defects closed recently where a check confirm that the RIGHT thing answered; the remedy th service could produce. WHO IS AFFECTED: any a half-broken deployment before it reaches ma health check at a server that returns a not-four check passes. ACCEPTANCE CRITERIA: all thre practical the check asserts something only the merely a successful status; and a test points ea
HXC-273	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				and asserts it fails, proven real by a paired mu behaviour and makes the test fail. Starting the remote-call server makes real outl provider APIs before it finishes starting, using in the environment. WHAT HAPPENS: start-up automatic discovery; a flag passed there disabl providers but does not disable discovery itself, provider whose credential it finds and enumer across SEVEN providers in a single start-up, dr host shell rather than supplied to the process: twice, because two different credential names eight, six, five, three and two from the rest. WH local server does not expect it to contact extern serving anything; in a sandboxed or air-gapped delay start-up; and the credentials used are wh operator deliberately supplied. IT ALSO CAUSE
HXC-274	Bug	Queued	Medium	why the server accepts connections into the op seconds before it actually begins serving them server was not speaking the expected protocol IMPORTANT NOTE FOR WHOEVER FIXES THIS wrong conclusion. The flag that disables verifi outbound behaviour, and the catalogue fetch discovery routine, so an investigation that stop outbound traffic and returned a refutation. On case where the source and the behaviour genu evidence for any claim about it. ACCEPTANCE calls unless explicitly asked to, discovery is opt are taken from deliberate configuration rather begins serving promptly after it begins accepti present and asserts no outbound calls occur, p restores eager discovery and makes the test fa