

# HelixCode — Fixed Items Tracker

Per Constitution §11.4.19 (Fixed-document column-alignment) + CONST-057 (Type-aware closure vocabulary: Bug → Fixed, Feature → Implemented, Task → Completed, all with (→ Fixed.md) suffix preserved).

This file is a **closure ledger** — items migrate here from docs/Issues.md ONLY after positive captured-evidence per §11.4.5.

**Round 189 prefix convention:** Closures originally tracked as ISSUE-NNN are now annotated with their new scope prefix and the legacy ID preserved in parentheses (e.g. HXL-001 (ex-ISSUE-003)) for git-history traceability. See docs/Issues.md “Prefix convention” for the full mapping table.

**Authoritative round-by-round narrative:** docs/CONTINUATION.md (CONST-044). Each row below points to the relevant close-out section there. Items predating the round-system are not retroactively captured (would be impractical) — they live in commit history + the docs/improvements/ evidence chain (P0-P5 phases).

| Closure    | Title                                                                                                                                           | Type | Status                 | Round              |
|------------|-------------------------------------------------------------------------------------------------------------------------------------------------|------|------------------------|--------------------|
| 2026-06-09 | HXC-044: internal/cognee — AMD-GPU rocm-smi JSON parser returns -1 sentinel instead of parsed GPU-use value                                     | Bug  | Obsolete (→ Fixed.md)  | 2026-06-09 session |
| 2026-05-29 | HXC-029: §11.4.98 full-automation compliance sweep — verify every live/integration/e2e/Challenge test is self-driving with no human-in-the-loop | Task | Completed (→ Fixed.md) | 2026-05-29 session |

| Closure    | Title                                                                                                                                                                                                | Type | Status                 | Round              |
|------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|------------------------|--------------------|
| 2026-05-29 | HXC-036: Systemic CONST-046 i18n defect — 74 packages emitted raw message-ID keys to users because the boot-time translator wiring was never implemented (a §11.9 tests-green-feature-broken defect) | Bug  | Fixed (→ Fixed.md)     | 2026-05-29 session |
| 2026-05-29 | HXC-034: Cascade constitution §11.4.102 (mandatory systematic-debugging + always-loaded using-superpowers + plugin-dependency availability) into every owned submodule + wire                        | Task | Completed (→ Fixed.md) | 2026-05-29 session |

| Closure    | Title                                                                                                                                | Type | Status                 | Round              |
|------------|--------------------------------------------------------------------------------------------------------------------------------------|------|------------------------|--------------------|
|            | the CM-COVENANT-114-102-PROPAGATION enforcement gate                                                                                 |      |                        |                    |
| 2026-05-29 | HXC-035: POST /api/v1/auth/register returned 400<br>internal_auth_failed_create_user on a fresh DB — blocked all authenticated flows | Bug  | Fixed (→ Fixed.md)     | 2026-05-29 session |
| 2026-05-29 | HXC-030: §11.4.99 latest-source documentation cross-reference sweep across all operator-facing docs                                  | Task | Completed (→ Fixed.md) | 2026-05-29 session |

| Closure    | Title                                                                                                                                                                               | Type | Status                 | Round              |
|------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|------------------------|--------------------|
| 2026-05-29 | HXC-033: codegraph 0.9.7 update dropped own-org submodules from the index + full index/sync appeared to crash (§11.4.79 regression)                                                 | Bug  | Fixed (→ Fixed.md)     | 2026-05-29 session |
| 2026-05-29 | HXC-032: LLMOrchestrator submodule had committed git conflict markers in 5 Go files (26 hunks) breaking helix_agent build — a §11.4 build-layer PASS-bluff already on origin/master | Bug  | Fixed (→ Fixed.md)     | 2026-05-29 session |
| 2026-05-28 | HXC-017: CodeGraph index excluded all own-org submodules (blanket dependencies/**) +                                                                                                | Task | Completed (→ Fixed.md) | 2026-05-28 session |

| Closure    | Title                                                                                                                                                                                       | Type | Status                | Round                 |
|------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|-----------------------|-----------------------|
|            | config.json was untracked —<br>§11.4.79/§11.4.78/§11.4.80                                                                                                                                   |      |                       |                       |
| 2026-05-28 | HXC-023: literal-true<br>Assert(true,...)/<br>AssertTrue(true,...) PASS-bluffs<br>across the e2e test banks (report<br>PASS without exercising<br>behaviour)                                | Bug  | Fixed (→<br>Fixed.md) | 2026-05-28<br>session |
| 2026-05-28 | HXC-022: test_bank<br>platform+integration packages<br>did not compile — half-written<br>stubs + root package-name<br>collision (anti-bluff: an<br>uncompilable test package never<br>runs) | Bug  | Fixed (→<br>Fixed.md) | 2026-05-28<br>session |

| Closure    | Title                                                                                                                                                     | Type | Status                 | Round              |
|------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|------|------------------------|--------------------|
| 2026-05-28 | HXC-016: §11.4.69–97 governance cascade into all owned submodules + mechanical propagation gate (CONST-047/§3, §11.4.32)                                  | Task | Completed (→ Fixed.md) | 2026-05-28 session |
| 2026-05-28 | HXC-001 (ex-ISSUE-005): CONST-052 lowercase-snake_case rename programme — all owned-org submodule leaf dirs + 57 Upstreams/ dirs renamed                  | Task | Completed (→ Fixed.md) | 2026-05-28 session |
| 2026-05-28 | HXC-021 + HXC-014a + HXC-015a: fake-skip<br>Assert(true, "...skipped") bluffs (11) + empty TestProviderStress stubs report green while exercising nothing | Bug  | Fixed (→ Fixed.md)     | 2026-05-28 session |

| Closure    | Title                                                                                                               | Type | Status                 | Round |
|------------|---------------------------------------------------------------------------------------------------------------------|------|------------------------|-------|
| 2026-05-21 | HXC-010: End-to-end Kimi CLI + Qwen Code CodeGraph verification (operator-blocked on LLM backend quota/credentials) | Task | Completed (→ Fixed.md) | 464   |

| Closure    | Title                                                                                               | Type    | Status                   | Round |
|------------|-----------------------------------------------------------------------------------------------------|---------|--------------------------|-------|
| 2026-05-20 | HXC-003: CONST-046 i18n migration backlog (no user-facing text hardcoded as static string literals) | Feature | Implemented (→ Fixed.md) | 463   |
| 2026-05-20 | HXC-008: CONST-055 G1 governance gaps surfaced by post-constitution-pull validation sweep           | Bug     | Fixed (→ Fixed.md)       | 403   |



| Closure    | Title                                                                                                                                                    | Type | Status                       | Round |
|------------|----------------------------------------------------------------------------------------------------------------------------------------------------------|------|------------------------------|-------|
| 2026-05-20 | HXC-007: Constitution<br>§11.4.68/70-74 cascade + meta-<br>pointer bump                                                                                  | Task | Completed<br>(→<br>Fixed.md) | 403   |
| 2026-05-20 | HXC-009: Owned-submodule<br>GitHub ↔ GitLab mirror-<br>divergence reconciliation                                                                         | Task | Completed<br>(→<br>Fixed.md) | 403   |
| 2026-05-20 | HXC-011: helix_qa runner emits<br>hollow sub-microsecond<br>“PASSED” rows for desktop-<br>platform bank cases instead of<br>executing the case’s action: | Bug  | Fixed (→<br>Fixed.md)        | 402   |

| Closure    | Title                                                                                                        | Type    | Status                         | Round |
|------------|--------------------------------------------------------------------------------------------------------------|---------|--------------------------------|-------|
| 2026-05-20 | HXC-012: data race in helix_code/<br>internal/llm/load_balancer.go<br>background stat-collector<br>goroutine | Bug     | Fixed (→<br>Fixed.md)          | 401   |
| 2026-05-20 | OPS-001: LLMOps 2 pre-existing<br>CreatePromptExperiment test<br>failures                                    | Bug     | Fixed (→<br>Fixed.md)          | 397   |
| 2026-05-19 | CONST-046 i18n architecture<br>design doc                                                                    | Feature | Implemented<br>(→<br>Fixed.md) | 90    |
| 2026-05-19 | pkg/i18n core foundation                                                                                     | Feature |                                | 91    |

| Closure    | Title                                                                                | Type    | Status                      | Round |
|------------|--------------------------------------------------------------------------------------|---------|-----------------------------|-------|
|            |                                                                                      |         | Implemented<br>(→ Fixed.md) |       |
| 2026-05-19 | CONST-046 audit script (soft-<br>warn)                                               | Feature | Implemented<br>(→ Fixed.md) | 92    |
| 2026-05-19 | Per-submodule i18n injection<br>wiring + i18nadapter                                 | Feature | Implemented<br>(→ Fixed.md) | 93    |
| 2026-05-19 | SelfImprove × 8 hardcoded-<br>content migration                                      | Feature | Implemented<br>(→ Fixed.md) | 94    |
| 2026-05-19 | HelixLLM × 2 CLI strings<br>migration                                                | Feature | Implemented<br>(→ Fixed.md) | 95    |
| 2026-05-19 | harmony_os × 5 CLI headers<br>migration                                              | Feature | Implemented<br>(→ Fixed.md) | 96    |
| 2026-05-19 | DocProcessor CLI × 8 migration                                                       | Feature | Implemented<br>(→ Fixed.md) | 97    |
| 2026-05-19 | Planning × 3 + VisionEngine × 4<br>migration                                         | Feature | Implemented<br>(→ Fixed.md) | 98    |
| 2026-05-19 | CONST-046 audit-gate fail-on-new<br>+ baseline                                       | Feature | Implemented<br>(→ Fixed.md) | 99b   |
| 2026-05-19 | panoptic × 5 cobra Short<br>descriptions migration                                   | Feature | Implemented<br>(→ Fixed.md) | 99a   |
| 2026-05-19 | challenges/pkg/i18n/ Phase 4<br>infrastructure + evaluators.go<br>migration          | Feature | Implemented<br>(→ Fixed.md) | 100   |
| 2026-05-19 | challenges/pkg/userflow/<br>challenge_recorded_ai_testgen.go<br>× 10 of 25 migration | Feature | Implemented<br>(→ Fixed.md) | 101   |
| 2026-05-19 | challenges/pkg/userflow/<br>challenge_desktop.go migration                           | Feature | Implemented<br>(→ Fixed.md) | 102   |

| Closure    | Title                                                                                  | Type    | Status                      | Round |
|------------|----------------------------------------------------------------------------------------|---------|-----------------------------|-------|
| 2026-05-19 | challenges/pkg/userflow/<br>challenge_ai_testgen.go × 10 user-<br>facing migration     | Feature | Implemented<br>(→ Fixed.md) | 103   |
| 2026-05-19 | challenges/pkg/userflow/<br>challenge_recorded_mobile.go × 7<br>unique × 14 call sites | Feature | Implemented<br>(→ Fixed.md) | 104   |
| 2026-05-19 | HXL-001 (ex-ISSUE-003):<br>HelixLLM analysis_test.go<br>hardcoded path                 | Bug     | Fixed (→ Fixed.md)          | 105   |
| 2026-05-19 | HXL-002 (ex-ISSUE-004):<br>HelixLLM TOON WriteTOON 500                                 | Bug     | Fixed (→ Fixed.md)          | 105   |
| 2026-05-19 | HXC-002 (ex-ISSUE-006) (partial):<br>HelixMemory LOGIC-class FAIL<br>cleanup           | Bug     | Fixed (→ Fixed.md)          | 106   |
| 2026-05-19 | HXC-002 (ex-ISSUE-006) (partial):<br>Planning LOGIC FAIL audit<br>confirms clean       | Task    | Completed<br>(→ Fixed.md)   | 107   |
| 2026-05-19 | CONST-046 i18n implemented-<br>architecture overview doc                               | Task    | Completed<br>(→ Fixed.md)   | 111   |
| 2026-05-19 | Tracker HTML + PDF exports per<br>§11.4.19                                             | Feature | Implemented<br>(→ Fixed.md) | 110   |
| 2026-05-19 | helix_code/cmd/helix_config/<br>main.go × 10 migration                                 | Feature | Implemented<br>(→ Fixed.md) | 108   |
| 2026-05-19 | helix_qa i18n kickoff (Phase 4<br>round 7)                                             | Feature | Implemented<br>(→ Fixed.md) | 112   |
| 2026-05-19 | CONST-052 rename programme<br>phased plan (HXC-001 plan, ex-<br>ISSUE-005)             | Task    | Completed<br>(→ Fixed.md)   | 113   |
| 2026-05-19 | LLMOrchestrator i18n kickoff<br>(Phase 4 round 9)                                      | Feature | Implemented<br>(→ Fixed.md) | 115   |
| 2026-05-19 | HXC-002 (ex-ISSUE-006) (final):<br>helix_agent inner LOGIC FAIL<br>cleanup             | Bug     | Fixed (→ Fixed.md)          | 109   |

| Closure    | Title                                                  | Type    | Status                      | Round |
|------------|--------------------------------------------------------|---------|-----------------------------|-------|
| 2026-05-19 | LLMsVerifier i18n kickoff (Phase 4 round 8)            | Feature | Implemented<br>(→ Fixed.md) | 114   |
| 2026-05-19 | HelixSpecifier i18n kickoff (Phase 4 round 10)         | Feature | Implemented<br>(→ Fixed.md) | 117   |
| 2026-05-19 | Storage i18n kickoff (Phase 4 round 11)                | Feature | Implemented<br>(→ Fixed.md) | 118   |
| 2026-05-19 | LLMOps i18n kickoff (Phase 4 round 12)                 | Feature | Implemented<br>(→ Fixed.md) | 119   |
| 2026-05-19 | VectorDB i18n kickoff (Phase 4 round 13)               | Feature | Implemented<br>(→ Fixed.md) | 120   |
| 2026-05-19 | Observability i18n kickoff (Phase 4 round 14)          | Feature | Implemented<br>(→ Fixed.md) | 121   |
| 2026-05-19 | MCP_Module i18n kickoff (Phase 4 round 15)             | Feature | Implemented<br>(→ Fixed.md) | 122   |
| 2026-05-19 | HXA-001 (ex-ISSUE-009):<br>helix_agent 4 handler tests | Bug     | Fixed (→ Fixed.md)          | 116   |
| 2026-05-19 | Messaging i18n kickoff (Phase 4 round 16)              | Feature | Implemented<br>(→ Fixed.md) | 123   |
| 2026-05-19 | Middleware i18n kickoff (Phase 4 round 17)             | Feature | Implemented<br>(→ Fixed.md) | 124   |
| 2026-05-19 | Plugins i18n kickoff (Phase 4 round 18)                | Feature | Implemented<br>(→ Fixed.md) | 125   |
| 2026-05-19 | Streaming i18n kickoff (Phase 4 round 19)              | Feature | Implemented<br>(→ Fixed.md) | 126   |
| 2026-05-19 | Watcher i18n kickoff (Phase 4 round 20)                | Feature | Implemented<br>(→ Fixed.md) | 127   |
| 2026-05-19 |                                                        | Feature |                             | 128   |

| Closure    | Title                                                                           | Type    | Status                      | Round |
|------------|---------------------------------------------------------------------------------|---------|-----------------------------|-------|
|            | conversation i18n kickoff (Phase 4 round 21)                                    |         | Implemented<br>(→ Fixed.md) |       |
| 2026-05-19 | containers i18n kickoff (Phase 4 round 22)                                      | Feature | Implemented<br>(→ Fixed.md) | 129   |
| 2026-05-19 | security i18n kickoff (Phase 4 round 23)                                        | Feature | Implemented<br>(→ Fixed.md) | 130   |
| 2026-05-19 | helix_code/cmd/cli/main.go × 10 migration (Phase 4 round 24)                    | Feature | Implemented<br>(→ Fixed.md) | 131   |
| 2026-05-19 | AutoTemp i18n kickoff (Phase 4 round 25)                                        | Feature | Implemented<br>(→ Fixed.md) | 132   |
| 2026-05-19 | Auth i18n kickoff (Phase 4 round 26)                                            | Feature | Implemented<br>(→ Fixed.md) | 133   |
| 2026-05-19 | helix_code/cmd/server/main.go × 10 migration (Phase 4 round 27)                 | Feature | Implemented<br>(→ Fixed.md) | 134   |
| 2026-05-19 | PliniusCommon i18n kickoff (Phase 4 round 28)                                   | Feature | Implemented<br>(→ Fixed.md) | 135   |
| 2026-05-19 | helix_code/applications/<br>terminal_ui × up to 10 migration (Phase 4 round 30) | Feature | Implemented<br>(→ Fixed.md) | 137   |
| 2026-05-19 | helix_code/applications/desktop i18n (Phase 4 round 29)                         | Feature | Implemented<br>(→ Fixed.md) | 136   |
| 2026-05-19 | helix_code/applications/ios infrastructure-only (Phase 4 round 31)              | Feature | Implemented<br>(→ Fixed.md) | 138   |
| 2026-05-19 | helix_code/applications/android infrastructure-only (Phase 4 round 32)          | Feature | Implemented<br>(→ Fixed.md) | 139   |
| 2026-05-19 | helix_code/applications/aurora_os × up to 10 migration (Phase 4 round 33)       | Feature | Implemented<br>(→ Fixed.md) | 140   |

| Closure    | Title                                                                              | Type    | Status                      | Round |
|------------|------------------------------------------------------------------------------------|---------|-----------------------------|-------|
| 2026-05-19 | helix_code/cmd/config_test × 12 migration (Phase 4 round 34)                       | Feature | Implemented<br>(→ Fixed.md) | 141   |
| 2026-05-19 | helix_code/cmd/security_test × 10 migration (Phase 4 round 35)                     | Feature | Implemented<br>(→ Fixed.md) | 142   |
| 2026-05-19 | helix_code/cmd/security_fix × 10 migration (Phase 4 round 36)                      | Feature | Implemented<br>(→ Fixed.md) | 143   |
| 2026-05-19 | helix_code/cmd/<br>performance_optimization × 10 migration (Phase 4 round 37)      | Feature | Implemented<br>(→ Fixed.md) | 144   |
| 2026-05-19 | helix_code/cmd/<br>security_fix_standalone × 10 of 27 migration (Phase 4 round 38) | Feature | Implemented<br>(→ Fixed.md) | 145   |
| 2026-05-19 | helix_code/internal/auth × up to 10 migration (Phase 4 round 39)                   | Feature | Implemented<br>(→ Fixed.md) | 146   |
| 2026-05-19 | helix_code/internal/agent × 10 of 64 migration (Phase 4 round 40)                  | Feature | Implemented<br>(→ Fixed.md) | 147   |
| 2026-05-19 | helix_code/internal/cognee × migration (Phase 4 round 41)                          | Feature | Implemented<br>(→ Fixed.md) | 148   |
| 2026-05-19 | helix_code/internal/commands × 10 migration (Phase 4 round 42)                     | Feature | Implemented<br>(→ Fixed.md) | 149   |
| 2026-05-19 | helix_code/internal/config × 10 migration (Phase 4 round 43)                       | Feature | Implemented<br>(→ Fixed.md) | 150   |
| 2026-05-19 | helix_code/internal/context × 8 sites / 5 IDs (Phase 4 round 44)                   | Feature | Implemented<br>(→ Fixed.md) | 151   |
| 2026-05-19 | helix_code/internal/database × 8 migration (Phase 4 round 45)                      | Feature | Implemented<br>(→ Fixed.md) | 152   |
| 2026-05-19 | helix_code/internal/discovery migration (Phase 4 round 47)                         | Feature |                             | 154   |

| Closure    | Title                                                            | Type    | Status                      | Round |
|------------|------------------------------------------------------------------|---------|-----------------------------|-------|
|            |                                                                  |         | Implemented<br>(→ Fixed.md) |       |
| 2026-05-19 | helix_code/internal/deployment × 10 migration (Phase 4 round 46) | Feature | Implemented<br>(→ Fixed.md) | 153   |
| 2026-05-19 | helix_code/internal/editor migration (Phase 4 round 48)          | Feature | Implemented<br>(→ Fixed.md) | 155   |
| 2026-05-19 | helix_code/internal/event migration (Phase 4 round 49)           | Feature | Implemented<br>(→ Fixed.md) | 156   |
| 2026-05-19 | helix_code/internal/focus migration (Phase 4 round 50)           | Feature | Implemented<br>(→ Fixed.md) | 157   |
| 2026-05-19 | helix_code/internal/hardware migration (Phase 4 round 51)        | Feature | Implemented<br>(→ Fixed.md) | 158   |
| 2026-05-19 | Round 74-87 release-gate stabilization                           | Task    | Completed<br>(→ Fixed.md)   | 82-87 |
| 2026-05-19 | release-gate-test.sh –skip-env-failures filter                   | Feature | Implemented<br>(→ Fixed.md) | 89    |
| 2026-05-19 | CONST-052 reference-drift sweep (73 submodules)                  | Task    | Completed<br>(→ Fixed.md)   | 88    |
| 2026-05-19 | challenges go.mod path fix ../Containers → ../containers         | Bug     | Fixed (→ Fixed.md)          | 87    |
| 2026-05-19 | LLMOrchestrator builders × 5 wired                               | Feature | Implemented<br>(→ Fixed.md) | 64-76 |
| 2026-05-19 | 4-vendor GPU telemetry chain (NVIDIA+AMD+Apple+Intel)            | Feature | Implemented<br>(→ Fixed.md) | 43-51 |
| 2026-05-19 | LLM Err coverage 100% across 17 providers                        | Feature | Implemented<br>(→ Fixed.md) | 46-63 |
| 2026-05-19 |                                                                  | Task    |                             | 188   |



| Closure                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | Title                                                                                               | Type | Status                       | Round |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|------|------------------------------|-------|
|                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | VEN-001 (ex-ISSUE-001):<br>VisionEngine helix-gitlab URL<br>fix (was misconfigured, not<br>missing) |      | Completed<br>(→<br>Fixed.md) |       |
| <p>2026-05-19   HXA-003 (ex-ISSUE-011): venice TestGetCapabilities<br/>CONST-037 model-list drift   Bug   Fixed (→ Fixed.md)   190  <br/>helix_agent (round-190) + meta pointer-bump   Replaced hardcoded<br/>venice-uncensored + llama-3.3-70b literals with structural assertion<br/>(NotEmpty + venice-uncensored* family substring scan); SKIP-OK<br/>CONST-035 marker on full-family rotation; mutation-verified<br/>revert→FAIL / restore→PASS  </p> <p>2026-05-19   HXC-004: recovery-batch 4-package under-verification<br/>(llm + logo + notification test-assertion drift + performance<br/>translator.go build break)   Bug   Fixed (→ Fixed.md)   200   (this<br/>commit)   Round-200 §11.4: 3 packages had test-assertion drift after<br/>round 161/163/167 i18n migration (tests still expected pre-i18n English<br/>literals; production now emits NoopTranslator-echoed message-IDs<br/>e.g. internal_llm_wizard_anthropic_apikey_required,<br/>internal_logo_open_source_failed,<br/>internal_notification_title_task_completed). Updated 11 test<br/>assertions (1 llm + 4 logo + 6 notification + 2 notification additional).<br/>4th package (performance) build-break already fixed inline by parent<br/>agent. All 4 packages PASS (llm 51.8s, logo 0.07s, notification 0.89s,<br/>performance 8.4s). Mutation-verified per CONST-035: 3 mutations (one<br/>per package), revert→FAIL, restore→PASS.  </p> <p>2026-05-19   PAN-001: panoptic appendString truncates multi-byte<br/>UTF-8 runes to bytes (TestResult.MarshalJSON corrupts non-ASCII)  <br/>Bug   Fixed (→ Fixed.md)   302   panoptic 24aa627 + meta pointer-<br/>bump   Replaced buf = append(buf, byte(r)) with buf =<br/>utf8.AppendRune(buf, r) in panoptic/internal/executor/<br/>executor.go:120 + added unicode/utf8 import. Evidence: go<br/>build ./... clean; go test -race -count=1 ./internal/executor/... →<br/>ok 4.470s; bash challenges/panoptic_describe_challenge.sh → 39/39<br/>PASS, 0 FAIL; UTF-8 detector flipped regression-present → fixed<br/>(literal: PASS [executor-marshal:utf8-detector:fixed] + KNOWN-ISSUE-<br/>RESOLVED: executor.appendJSONString now UTF-8 clean). Probe<br/>AppName="ü" (UTF-8 0xC3 0xBC) preserved end-to-end through<br/>MarshalJSON.  </p> <p>2026-05-20   HXC-005: cmd/performance_optimization_standalone/<br/>main.go is a CONST-035 simulation bluff   Bug   Fixed (→ Fixed.md)  <br/>318   (this commit)   Decision: DELETE (obsolete). The standalone</p> |                                                                                                     |      |                              |       |

binary fabricated improvement percentages via `improvement := 5.0 + rand.Float64()*20.0`, slept `time.Sleep(500ms)` per fake phase, and reported success for work it never performed (canonical BLUFF-001 anti-pattern, CONST-035 / Article XI §11.9). Genuinely obsolete — fully superseded by `cmd/performance_optimization/` (`snake_case` post-CONST-052) which calls the `REAL` internal/  
`performance.PerformanceOptimizer` (real `runtime.ReadMemStats`, real GOMAXPROCS tuning, real before/after measurement) + CONST-046 i18n + unit tests. `git rm -r cmd/performance_optimization_standalone/`; stale refs purged from `docs/COMPREHENSIVE_AUDIT_REPORT.md`; `gitignore` extended for auto-generated perf reports (CONST-053). Reproduce-before-fix regression test `cmd/performance_optimization/bluff_regression_test.go`:  
`TestHXC005_BluffStandaloneDirectoryDeleted` (obsolete path gone + real `cmd` survives) + `TestHXC005_RealOptimizerMeasuresActualMemory` (retained 32 MiB buffer → optimizer baseline `MemoryUsage` must track genuine `runtime.HeapAlloc`, not RNG). Evidence: `go build ./cmd/... exit 0`; `go test -count=1 -run TestHXC005 ./cmd/performance_optimization/` → both PASS (literal: optimizer baseline `MemoryUsage=33812624` bytes, `runtime.HeapAlloc=33802008` bytes — both real measurements). |  
2026-05-20 | HXV-001: LLMsVerifier 18 pre-existing tests/ failures (CLI-integration + verification/scoring) | Bug | Fixed (→ Fixed.md) | 323 | LLMsVerifier 59f542ba (github+gitlab) + meta pointer-bump | Round-323 §11.4 / CONST-035. 18 failures classified: **(A) test-build drift** — 9 CLI tests (`tests/automation_test.go`) ran `go run ../cmd/main.go` which compiles ONLY `main.go`; i18n rounds 194/308/309/312/316/319 correctly placed `tr()/trData()` in `cmd/i18n.go` (same package `main`), so single-file `go run` broke with undefined: `tr` (`go build ./cmd` was always fine). Fix: re-keyed all 14 invocations to `go run ../cmd` (whole package — exercises the real CLI binary). **(A) test-assertion drift** — 9 tests/unit/`model_verification_test.go` cases asserted the OLD §11.4 PASS-bluff (`verification.Verify()` returning `all-capabilities-true` / `all-scores-8.5`); round-17 commit `a6328629` correctly removed that fabrication → `ErrVerificationNotWired`. Fix: rewrote the unit suite to certify the HONEST contract (input validation; well-formed-but-unwired → loud `ErrVerificationNotWired` sentinel; verifier stateless + race-free). Real e2e verification stays in `llmverifier/` integration suite. **(C) environment gate** — `TestCommandFlagValidation/TestOutputFormats` list sub-tests dispatch real HTTP; tolerance only covered connection refused/dial tcp but a non-LLMsVerifier service answered 404. Fix: added `serverUnavailable()` recognising 404/5xx/no-host as a SKIP-OK: #HXV-001 env gate per CONST-035 (8 sub-tests now honest SKIP-OK). Evidence: before `go test ./tests/...` = 18 FAIL; after = ok

digital.vasic.llmsverifier/tests (17.9s), ok .../tests/integration (3.1s), ok .../tests/unit (0.003s); go build ./... clean. No production code changed — round-17 production fix was already correct; only test re-keying. |

2026-05-20 | HXQ-001 (ex-ISSUE-008): helix\_qa intermittent TestPerformance flake (host-load-sensitive) | Bug | Fixed (→ Fixed.md) | 325 | helix\_qa 649e2dd (github+gitlab) + meta pointer-bump | Round-325 \$11.4 / CONST-035. The three pkg/vision/ perf tests (TestPerformance\_DHash64\_Under5msPer1080pFrame, TestPerformance\_PHash\_Under25msPer1080pFrame, TestPerformance\_SSIM\_Under5msPer480pFrame) assert hard per-frame timing ceilings (5 ms / 25 ms / 5 ms) that flake under concurrent container/build CPU contention. **Decision: resolution path (b)** — env-var gating — chosen over path (a) (loosen tolerance). Rationale: loosening the timing tolerance would weaken the test's anti-bluff value (a genuine perf regression could then pass); path (b) preserves the strict assertions while making the flake deterministic. The three tests now gate on os.Getenv("HOST\_LOAD\_DEDICATED") — t.Skip("SKIP-OK: #HXQ-001 ...") honestly when unset (loaded/shared host), strict run when HOST\_LOAD\_DEDICATED=1 (quiescent dedicated host). NO timing tolerance loosened. docs/test-coverage.md §6.1 documents the env var. Evidence: go build ./pkg/vision/... exit 0, go vet clean; go test -count=1 -run TestPerformance ./pkg/vision/... (unset) → all 3 --- SKIP with SKIP-OK: #HXQ-001 marker; HOST\_LOAD\_DEDICATED=1 go test ... → all 3 --- PASS strict (DHash64 average 741ns, PHash average 88.969µs — well under the 5 ms / 25 ms ceilings). |

2026-05-20 | VEN-002 (ex-ISSUE-002): VisionEngine vasic-digital-github fork lineage divergent at SHA 93c830a | Bug | Fixed (→ Fixed.md) | 340 | VisionEngine merge 70c9e0c (4 remotes) + meta pointer-bump | Round-340 \$11.4.41 / CONST-061 merge-first. vasic-digital-github/master HEAD 93c830a → 256cce1 carried 15 commits absent from HelixDevelopment local master (6e3888e); local carried 60 absent from vd. **Decision: merge-first** (operator-approved) — real 2-parent merge commit 70c9e0c (parents 6e3888e HelixDevelopment + 256cce1 vasic-digital), NO force-push, NO rebase, ALL commits from both lineages preserved. Integrated vd round-48 (aaf9bda RPC lifecycle sentinels) / round-52 (7c0455b real RPC server lifecycle) / round-57 (93c830a real ProbeHosts + planning) / round-47 (5496b2d) / round-40 (1169213 SSH Shutdown) / 76452da + governance cascades. **16-file conflict resolution:** (1) 6 challenge scripts — kept HEAD canonical VISIONENGINE\_\* env names, dropped fork-specific VISIONENGINE\_VD\_\* prefix per CONST-051(B) decoupling; (2) go.mod/go.sum — kept HEAD newer testify-v1.11.1 transitive graph, go mod tidy reconciled clean; (3) pkg/analyzer/stub.go — took vd anti-bluff sentinel bodies: HEAD body

had RE-INTRODUCED the fabricated ScreenAnalysis{Title:"Unknown Screen"}+err=nil PASS-bluff that the file's own unconflicted doc comments declare removed round-27 — per CONST-035/Article XI §11.9 the bluff must not survive; removed bluff-only

i18n\_defaults.go+i18n\_callsites\_test.go (sole consumer was the removed path; pkg/i18n package untouched); (4) pkg/remote/{remote,ssh,remote\_test}.go — took HEAD: EnsureReady SSH-probe is a strict functional superset of vd's config-only EnsureReady, HEAD defines ErrBackendNotReachable+SSHConfig.BackendProbePort the merged body needs, HEAD SSHConfig is a field-superset so vd RPC code compiles unchanged; (5) pkg/remote/deployer\_test.go — modify/delete conflict, kept vd's 1026-line real anti-bluff RPC test suite per CONST-061 union-merge (deployer.go+distributed.go auto-merged as pure vd additions); (6) CONSTITUTION.md/CLAUDE.md/AGENTS.md — took HEAD (current governance lineage: 52 §11.4.x anchors + CONST-001..061 + §11.4.67; vd-unique CONST-030 superseded by CONST-035, vd-unique CONST-068 is the ID-form of §11.4.67 HEAD carries by literal). Evidence: zero conflict markers (grep -rn '<<<<<<\\|=====\\|>>>>>>' --include='\*.go' empty); go build ./... exit 0; go vet ./... clean; go test -count=1 ./... → 7 packages PASS (analyzer config graph i18n llmvision opencv remote); all 4 remotes fast-forwarded (6e3888e..70c9e0c helix-github/gitlab + vasic-digital-gitlab, 256cce1..70c9e0c vasic-digital-github — NO force on any). | 2026-05-20 | HXA-002 (ex-ISSUE-010): helix\_agent debate/llmprovider sibling-submodule API drift | Bug | Fixed (→ Fixed.md) | 342 | helix\_agent (round-342 HXA-002) + meta pointer-bump | Round-342 §11.4 / CONST-035. **Investigation finding (operator's explicit ask — moved vs deleted): GENUINELY DELETED, not moved.** git log on digital.vasic.debate (submodules/debate\_orchestrator) shows the orchestrator was rebuilt from scratch (commit 196d0ea "initial DebateOrchestrator reconstruction (Phase 1)"); git log --follow orchestrator/api.go = single entry — the slim CreateDebate/GetStatistics API is the first and only version. Tree-wide grep of dependencies/ for KnowledgeRepository/GetRecommendations/ConvertAPIRequest/GetDebateStatus/DefaultMinConsensus/MaxAgentsPerDebate/EnableAgentDiversity found ZERO surviving copies in any digital.vasic.\* package or HelixSpecifier/HelixMemory. The richer learning/knowledge/recommendations tier was a pre-reconstruction artifact that no longer exists anywhere — not relocated.

**Part 1 (mechanical import swap):** digital.vasic.llmprovider's LLMProvider interface now uses its own in-module digital.vasic.llmprovider/pkg/models; swapped digital.vasic.models → digital.vasic.llmprovider/pkg/models in 3 files (provider\_bridge.go production + mock\_test.go + provider\_bridge\_leak\_test.go unit tests).

**Part 2 (slim-API rewrite, deleted-tier path per operator):** rewrote `internal/services/debate_integration/integration_test.go` + `tests/integration/debate_full_flow_test.go` down to the slim `CreateDebate/GetStatistics/ConductDebate/CancelSession/Bank()` surface — also fixed `provider_bridge_test.go` (`NewProviderInvoker` now takes `(registry, name)`); `GetKnowledgeRepository` → `Bank()` and `service_integration_test.go` (`DebateMetrics.TotalResponses` → `ProviderCalls`); converted all `RegisterProvider` scores from the old 0-10 scale to the reconstructed `[0,1]` scale (`orchestrator` now rejects `score>1`). Lost coverage documented honestly in the rewritten files' header comments: request-conversion / knowledge-repository / recommendations / status-by-id assertions dropped (those API surfaces no longer exist). **Bonus rename-drift fix:** `helix_agent/go.mod` replace `digital.vasic.debate` still pointed at `PascalCase DebateOrchestrator` after a parallel `CONST-052` rename to `debate_orchestrator` — corrected (required precondition for verification). Evidence: `go build ./internal/services/debate_integration/... exit 0`; `go test -count=1 ./internal/services/debate_integration/... → ok dev.helix.agent/internal/services/debate_integration 0.156s (PASS)`; rewritten `TestDebateFullFlow_OrchestratorInit` verified PASS in an isolated standalone package (the `tests/integration` package itself cannot compile due to a SEPARATE pre-existing `helix_qa/pkg/autonomous` ↔ `VisionEngine` remote API drift, entirely unrelated to HXA-002 — `helix_qa` committed code at `7fa674a`). The standalone verification caught and fixed one wrong assertion in the original test (`DefaultTimeout` is 30s, not 5m). `gofmt` clean on all 8 changed files. | 2026-05-20 | VEN-002 (ex-ISSUE-002): `VisionEngine vasic-digital-github` fork lineage divergent at SHA `93c830a` | Bug | Fixed (→ `Fixed.md`) | 340 | `VisionEngine` merge `70c9e0c` (4 remotes) + meta pointer-bump | Round-340 \$11.4.41 / `CONST-061` merge-first. (See round-340 row above — duplicate header retained intentionally for migration-discipline alignment.) | 2026-05-20 | HXQ-002: `helix_qa pkg/autonomous` ↔ `VisionEngine` remote API drift blocks `helix_agent tests/integration` compile | Bug | Fixed (→ `Fixed.md`) | 344 | `helix_qa 9ef3d95` (github+gitlab) + meta pointer-bump | Round-344 \$11.4 / `CONST-035`. Drift resolved by consuming the round-340 VEN-002 merged superset remote API — drift was MECHANICAL (three changed signatures, no removed type, no renamed symbol). **Per-symbol classification: (1) ProbeHosts** — `(ctx, hosts []string, user string) []HardwareInfo` → `(ctx, hosts []SSHConfig) ([]HardwareInfo, error)`. Fix: `pipeline.go` builds a `[]visionremote.SSHConfig` from `VisionHosts` + config-injected SSH key / `known_hosts` / `port`; the joined per-host probe error is surfaced as a

warning (partial-success is normal). **(2) SelectStrongestModel** — (infos []HardwareInfo) \*ModelRecommendation → (infos []HardwareInfo, models []ModelSpec) (\*ModelRecommendation, error). Fix: a single-entry visionremote.ModelSpec catalogue is built from LlamaCppRPCModelPath + config capacity floors; error handled with single-host fallback. **(3) PlanDistribution** — (infos []HardwareInfo, path string, serverPort, rpcBasePort int) \*DistributionConfig → (infos []HardwareInfo, models []ModelSpec) (\*DistributionConfig, error). Fix: the GGUF model path and server port are no longer call arguments — they are set on the returned \*DistributionConfig after the planner's best-fit bin-pack; error handled with single-host fallback. Added five config-injected pipeline fields (VisionSSHKeyPath, VisionSSHKnownHostsPath, VisionSSHPort, VisionRPCMinGPUMemMB, VisionRPCMinRAMMB) per CONST-045 / CONST-046 — no hardcoded secrets or model metadata. VisionEngine submodule pointer was already at the round-340 merged HEAD 70c9e0c (no bump needed). Evidence: cd helix\_qa && go build ./pkg/autonomous/... exit 0; go test -count=1 ./pkg/autonomous/... → ok digital.vasic.helixqa/pkg/autonomous 14.270s (PASS); cd helix\_agent && go build ./tests/integration/... exit 0 — the original HXQ-002 symptom is resolved. | 2026-05-20 | HXV-002: LLMsVerifier verification/ package 10 pre-existing test failures | Bug | Fixed (→ Fixed.md) | 348 | LLMsVerifier (round-348 HXV-002) + meta pointer-bump | Round-348 \$11.4 / CONST-035. All 10 failures classified **(A) test-assertion drift** — every failing test asserted pre-honesty fabricated behaviour that round-17 commit a6328629 correctly removed; **no production code changed. 8 in verification\_test.go** — TestVerifier\_Verify\_Success / \_ResultScores / \_LatencyMetrics / \_CodeLanguageSupport / \_CodeCapabilities / \_ModelStatusFlags / \_ContextCancellation / \_MultipleRequests each asserted require.NoError + fabricated all-capabilities-true / all-scores-8.5 / fabricated latency+status results; Verify() now correctly returns ErrVerificationNotWired (the honest round-17 contract — verification dispatch deliberately un-wired to remove a CONST-036/037 single-source-of-truth PASS-bluff). Re-keyed each to certify the honest contract: require.Error + errors.Is(err, ErrVerificationNotWired) + nil result; \_Success renamed TestVerifier\_Verify\_NotWiredContract for accuracy. **2 in code\_verification\_test.go** — TestCodeVisibility\_Error asserted require.NoError on an HTTP 503 (the OLD bluff that swallowed API failures); production correctly propagates the error so callers distinguish API failure from negative verification — re-keyed to require.Error + 503 substring, response still carries Verified=false+Error. VerifyModelCodeVisibility\_ServerError asserted Status=="verified"+score≥0.7 on a server returning HTTP 500 for

every sample (the “relaxed verification” bluff); zero successful responses → production correctly returns `Status=="failed"` — re-keyed to assert failed + non-empty `ErrorMessage`. Evidence: before `go test ./verification/... = 10 FAIL`; after `= ok digital.vasic.llmsverifier/verification 1.635s, 0 failures`; `go build ./... clean`. Mirrors HXV-001 round-323 classification approach — production code (round-17 `ErrVerificationNotWired`, `code_verification.go` error-propagation + zero-response → failed) was already honest; only stale test assertions needed re-keying. |

2026-05-20 | HXV-003: LLMsVerifier

`ProviderAdapterForBenchmark.Complete` is a CONST-050(A) production mock-bluff | Bug | Fixed (→ `Fixed.md`) | 396 | LLMsVerifier (round-396 HXV-003) + meta pointer-bump | Round-396 \$11.4 / CONST-050(A) / CONST-035 / BLUFF-001. **Decision: WIRE, not delete.** Investigation confirmed `ProviderAdapterForBenchmark` IS live code — `BenchmarkSystem.Initialize` wires it as the runner’s `LLMProvider` (two call sites in `benchmark_test.go`), so honest deletion was not applicable. The bug: `Complete(ctx, prompt, systemPrompt)` body was `// Mock implementation - actual would call real provider + return "Response", 50, nil` — fabricated a hardcoded completion for an LLM call it never made (canonical BLUFF-001). **Fix:** the adapter’s provider field was an untyped `interface{}` it never used; retyped it to the real `benchmark.LLMProvider` interface (`Complete` + `GetName`, the same contract `HTTPBenchmarkProvider` satisfies). `NewProviderAdapterForBenchmark` NOW type-asserts the passed value to `LLMProvider` and stores it; `Complete` dispatches directly to `a.provider.Complete(ctx, prompt, systemPrompt)`, returning the underlying provider’s genuine response text + real token count + real error verbatim. When no real provider is wired, `Complete` returns the new honest sentinel `ErrProviderAdapterNotWired` — it NEVER fabricates a result (mirrors the round-28 `ErrBenchmarkProviderNotConfigured` posture). The real dispatch path is `HTTPBenchmarkProvider` (OpenAI-compatible HTTP `LLMProvider`, already in the package). **Reproduce-before-fix tests** (`benchmark_coverage_test.go`): the pre-existing `TestProviderAdapterForBenchmark_Complete` relied on the bluff (`nil provider` → `NoError` + non-empty resp) — replaced with two honest tests: `_Complete_NotWired` asserts `ErrProviderAdapterNotWired` + empty resp + zero tokens (no fabrication), and `_Complete_RealDispatch` constructs an `httptest.NewServer` OpenAI shim + `HTTPBenchmarkProvider`, asserts the adapter ACTUALLY hits the server (`hit-counter == 1`) and returns the real server payload ("`4 is the real answer`", server-reported `total_tokens 19`) — explicitly `NotEqual("Response") / NotEqual(50)`. Evidence: `go build ./... exit 0`; `go test -count=1 ./internal/benchmark/... → ok llmsverifier/`

internal/benchmark 12.334s (PASS); the 3  
TestProviderAdapterForBenchmark\* tests all --- PASS; anti-bluff smoke  
grep -rn "Mock implementation\|simulated\|for now" internal/  
benchmark/\*.go (prod only) = clean. |  
2026-05-20 | HXC-006: HelixCode Speed Programme — 3-5× faster than  
competitor AI CLI agents (6-phase / 31-task) | Feature | Implemented  
(→ Fixed.md) | 400 | P5-T04 (this commit) + 30 prior task commits |  
Round-400 / CONST-048 / CONST-035. Operator mandate 2026-05-20:  
make HelixCode + owned-submodule code 3-5× faster than competitor  
AI CLI agents without breaking any feature or weakening anti-bluff  
posture. **All 6 phases / 31 tasks landed** — P0-T01..04 (baseline harness:  
pprof + benchmarks + competitor wall-clock + scenario runner), P1-  
T01..07 (LLM & startup wins), P2-T01..07 (context-build speed), P3-  
T01..05 (interactive & agent-loop levers), P4-T01..04 (profile-gated  
tuning), P5-T01..04 (dev-experience + submodule cascade + this close-  
out). **CONST-048 coverage ledger** committed at docs/research/speed/  
05-coverage-ledger.md: **29 PASS + 2 honestly-bounded PARTIAL + 0  
DEFERRED**. PARTIALs reported truthfully — P5-T01 (98315a14) build/  
test parallelism tuning is landed and correct but the suite-wall-time  
before/after delta was not captured as a pasted benchmark; P5-T02  
(4ee771d7) is a partial single-provider internal/llm split (Cerebras  
extracted to internal/llm/providers/cerebras/) — a full 18-provider  
extraction is genuinely infeasible without an import cycle (factory.go  
in package llm constructs every provider). **Headline measured wins**  
(each carries pasted in-session evidence per CONST-035, transcribed in  
the ledger): P1-T01 HTTP/2 transport ~2×, P1-T02 lazy Ollama discovery  
67µs → 2.7ns, P1-T03 lazy CLI startup ~8.85×, P1-T06 cache pre-warm  
~7.6×, P2-T02 regexp hoist ~7.4×, P2-T03 repo-map cache ~10.6× warm,  
P2-T04 parallel repo-map ~1.67×, P2-T05 parallel search ~4.39×, P2-T06  
incremental tree-sitter ~21×, P3-T01 small-model routing ~5.87×, P3-T02  
diff edits 94-99% token cut, P3-T03 fast-apply ~516×, P3-T04 tool  
parallelism ~5.99×, P4-T01 PGO -46% CPU-bound. **Release-gate sweep  
(P5-T04):** anti-bluff smoke grep -rn "simulated\|for now\|TODO  
implement\|placeholder" helix\_code/internal helix\_code/cmd (prod) =  
clean; scripts/audit-const046-hardcoded-content.sh ran exit 0 (speed  
work added no user-facing strings, no new hardcoded content);  
scripts/verify-governance-cascade.sh = 2 pre-existing failures = the  
already-tracked HXC-008 drift (verifier stale Models path + helix\_qa/  
CONSTITUTION.md missing CONST-047..057), NOT speed-programme  
regressions. Two pre-existing defects surfaced during the programme  
filed honestly as **HXC-011** (helix\_qa runner hollow sub-µs PASSED rows  
on desktop platform — §11.4 PASS-bluff) + **HXC-012** (internal/llm/  
load\_balancer.go stat-collector data race under -race). No speed task  
introduced a regression, a new bluff pattern, or new hardcoded



content. |

2026-06-16 | HXC-113: MCP tool names use 'server:name' (colon) — OpenAI-compatible providers (DeepSeek/etc.) reject function names, breaking LLM chat with MCP enabled | Bug | Fixed (→ Fixed.md) | 2026-06-16 session | c4c88f91 | mcpToolRegisteredName sanitises MCP tool names to server\_\_name (OpenAI-compatible<sup>1</sup>); dispatch unaffected (Execute uses original server/toolName). 2 mcp\_readonly tests reconciled (§11.4.120); guard test + full internal/tools pkg pass; build exit 0. |

2026-06-15 | HXC-111: Desktop GUI shows raw i18n keys (desktop\_dashboard\_header/\_activity\_title) — CONST-046 gap | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 883c7d0f | Wired i18n.NewTranslator() in NewDesktopApp; dashboard now shows real text (verified via relaunch+AXRaise+screenshot: title 'HelixCode - Distributed AI Development Platform', 'Recent Activity', no raw keys). Desktop tests pass, build exit 0. Clean desktop video re-recorded. |

2026-06-15 | HXC-110: Extend containers submodule to launch iOS simulators (operator-directed Apple-support mechanism) | Task | Completed (→ Fixed.md) | 2026-06-15 session | 5a86067a | submodules/containers/pkg/applesim: host xcrun-simctl orchestration (Boot/Install/Launch/Record/Shutdown, by stable UDID §11.4.111), 16 tests pass incl -race, real host round-trip; cmd/applesim CLI.

Submodule a0fa823 pushed all upstreams, meta pointer bumped. | 2026-06-15 | HXC-109: Mobile apps are scaffolds — Android has no build.gradle/AndroidManifest, iOS has no Xcode project (not buildable -> not recordable) | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 5a86067a | Android: buildable Gradle project + 67MB APK runs on live Genymotion (3 videos: connect/lifecycle/tasklist, real server task list via authenticated HTTP); 2 runtime issues fixed (JWT client-mode, JSON parse). iOS: buildable Xcode project (gomobile xcframework + rewired binding) builds+runs on iPhone14 sim (helixcode-ios-launch video, Go core OK). Both committed (1ffc9b69/38caa48d). Mobile apps no longer scaffolds. |

2026-06-15 | HXC-106: helix\_agent durable memory: process-lifetime in-memory fallback is NOT disk-durable — recall lost on restart (CONTINUATION honest gap) | Task | Completed (→ Fixed.md) | 2026-06-15 session | c88e4aee | Investigated (§11.4.102): disk-durable DiskStore (sqlite, survives close+reopen) was ALREADY implemented + wired as preferred fallback (commits ac3ad237/a91faad6) via debateMemoryFallbackPath() (os.UserCacheDir, 'helixagent'-namespaced, CONST-051-decoupled). CONTINUATION 'in-memory only' gap was stale. Closed the test-coverage gap: new internal/services/debate\_memory\_fallback\_test.go proves resolver returns writable durable path + persist->Close(restart)->reopen->RECALL. ./internal/

memory + ./internal/services pass; submodule HEAD c5bdcfad pushed to upstreams. |

2026-06-15 | HXC-105: Runtime e2e for server POST /api/v1/specify — boot server -> real spec output via live provider (speckit HTTP-endpoint gap) | Task | Completed (→ Fixed.md) | 2026-06-15 session | c88e4aee | tests/integration/specify\_server\_e2e\_test.go boots the real server + POSTs /api/v1/specify against live ollama qwen2.5:3b: HTTP 200 status:success provider:ollama qualityScore:0.9808, real 3-round 2-agent debate, provider\_calls=6 total\_tokens=806; output non-empty + NOT the ‘awaiting provider wiring’ stub. PASS 75.93s, vet clean, build exit 0. Evidence docs/qa/web-llm-e2e-20260615/. |

2026-06-15 | HXC-103: Web-client runtime e2e proof — live browser/ HTTP -> server -> LLM provider round-trip for /api/v1/llm/generate + /llm/stream (CONTINUATION honest gap) | Task | Completed (→ Fixed.md) | 2026-06-15 session | 3b7e692f | All 3 web LLM paths proven e2e against live ollama qwen2.5:3b: /generate (HTTP 200 content:4 provider:ollama), /stream (9 SSE frames + [DONE], streamed 1..5, >1 frame proves streaming), browser->server->provider (chromedp: #output DOM=4, #meta provider=ollama, screenshot). Exposed+fixed production hang HXC-104. Evidence + README docs/qa/web-llm-e2e-20260615/. SKIP-OK when provider down (\$11.4.3). |

2026-06-15 | HXC-104: streamLLM /api/v1/llm/stream hangs forever — chunkChan never closed, [DONE] never emitted (production defect found by web e2e) | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 3b7e692f | Fixed via ‘defer close(chunkChan)’ in streamLLM goroutine. Regression guard tests/integration/llm\_stream\_e2e\_test.go (TestLLMStreamE2E): post-fix streams 9 SSE frames + [DONE] in 1.1s, deterministic -count=3; server unit pkg ok; build exit 0. Evidence docs/qa/web-llm-e2e-20260615/. |

2026-06-15 | HXC-102: harmony\_os main\_nogui.go — 2 user-facing strings (‘Goodbye!’, ‘Error: %v’) bypass i18n (CONST-046, low severity) | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 127b8ccd | Routed cmdInteractive Goodbye!/Error: %v through cliApp.tr with new bundle keys (error binds {{.Error}}, no ). Guard hxc102\_interactive\_i18n\_test.go (nogui): GREEN, RED-without-key (raw-key leak exit 1), restored GREEN. Full harmony\_os pkg ok both tag variants; build exit 0; gofmt clean. |

2026-06-15 | HXC-099: Systemic i18n raw-key sweep redo (CONST-046) — corrected, regression-free, with default-translator contract decision | Task | Completed (→ Fixed.md) | 2026-06-15 session | e4bdd0d0 | Corrected redo per operator decision (preserve loud raw-key NoopTranslator default — NO default swap; 9 guards pass unchanged). GOAL A: WireAll() was only called from cmd/cli; added entry-path init() wiring to cmd/server + 4 apps so internal-package strings resolve for

real users. GOAL B: removed redundant {{.Err}} placeholder from 8 internal/project messages (nil-data -> ""). RED → GREEN guards captured (non-tautology proven); 9 guards green -count=1; all touched pkgs pass; build exit 0; vet clean; no mutation residue. Commit a02a8aa8. (B's rejected default-swap approach preserved in git stash@{0}.) |

2026-06-15 | HXC-101: security/security\_test.go TestTLSConfiguration — external-network dependency + nil-deref panic crashes the whole security test binary | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 61a33986 | Replaced live httpbin.org call with hermetic httptest.NewTLSServer + t.Fatalf on error path (no nil-deref fall-through). Verified: TestTLSConfiguration PASS 3/3 (-count=3) deterministic, no external net; full security pkg ok (0.223s); go build ./... exit 0; gofmt clean. |

2026-06-15 | HXC-100: Resync docs/CONTINUATION.md to current HEAD + de-bloat the 32k-token line-1 header (CONST-044/\$12.10 + CONST-064 hygiene) | Task | Completed (→ Fixed.md) | 2026-06-15 session | afbafeea | CONTINUATION.md resynced to HEAD 3aacfa9f + line-1 header de-bloated (max line 54856->2931 chars). History preserved: 4 mega-prose lines replaced by CONST-064 metadata table + ToC + condensed close-out 143-169 table; close-outs 131-142 dedicated sections intact; this session's 5 rows added. git diff +143/-8 (additive, no content loss); exports rendered=2 failed=0 fresh. Gated independently by conductor. |

2026-06-15 | HXC-078: T1.6 SKILL.md precedence resolution (partial) | Task | Completed (→ Fixed.md) | 2026-06-15 session | 3aacfa9f | FindMatching/List already resolve overlapping matches deterministically by lexicographic name (sort.Strings, documented contract markdown\_skills.go:194) — gap was missing coverage, not a bug. Added TestSkillRegistry\_FindMatching\_OverlappingPatternsDeterministic (insertion-order independence + 50-iter stability + lexicographic order), PASS; full internal/commands pkg ok; build exit 0. Commit 6e03fe15. |

2026-06-15 | HXC-077: T1.5 context-window percentage indicator (partial) | Feature | Implemented (→ Fixed.md) | 2026-06-15 session | 3aacfa9f | TUI status bar now renders honest context-window USED-% (commit 6e03fe15). sessionUsedTokens accumulates real per-turn tokens; window from catalogue ContextSize->GetContextWindow; OMITS when unknown (CONST-035); label i18n-routed via new terminal\_ui\_chat\_context\_usage key (CONST-046). Guard context\_usage\_test.go: GREEN with key, RED without (raw-key echo exit 1), restored GREEN. Full terminal\_ui pkg ok; build exit 0. |

2026-06-15 | HXC-098: out-of-box config fails 'version required' validation — blocks client status/system commands | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 80e62afa | Reproduced via

LoadHelixConfig path (RED: version-less config.json -> Version="" + server.port=0 -> validateConfig rejects). Fixed in internal/config/config.go loadConfigLocked: decode JSON ON TOP of getDefaultConfig() so all viper defaults merge. Guard hxc098\_version\_default\_test.go: RED pre-fix (exit 1), GREEN post-fix (Version=1.0.0,Port=8080). config.yaml ships explicit version. Full config pkg ok. |

2026-06-15 | HXC-093: helix\_code module graph has phantom digital.vasic.\* requires + private transitive blocking go list -m all / gomobile bind | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 5690dc58 | helix\_code/go.mod: 27 replace directives -> local submodules; go list -m all no-such-host 26->0 (462 modules); REAL .aar produced (classes.jar + jni/libgojni.so x4 ABIs); go build ./... green throughout |

2026-06-15 | HXC-090: panoptic tracks test-generated audit.json users.json (CONST-053 hygiene) | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 36c9ee42 | panoptic a77228e: enterprise TestMain runs from temp dir; tree stays clean post-test x2, tests PASS, no production change. Guard via §11.4.135 (HXC-096) committed separately. |

2026-06-15 | HXC-097: SYSTEMIC: standalone binaries + internal/config + internal/database never wire i18n Translator -> raw keys at runtime | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | cdf4fa81 | Systemic unwired-translator fixed across aurora\_os (842551d3) + harmony\_os (6dcf64aa) + internal/config + internal/database: real bundle translator wired (binaries: SetTranslator in main(); libs: init() default). Before/after raw-key->prose captured each; §11.4.115 guards. Broader follow-up: other CONST-046 pkgs may share the WireAll-only-on-CLI-path class (init()-default pattern is the fleet fix). |

2026-06-15 | HXC-096: desktop nogui prints raw i18n keys + %!(EXTRA) format mismatch in status/help | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | a843cc0f | cmd/cli/main.go: known-provider-prefix guard on model parsing; live qwen2.5:3b -> real answer '4', zero 404; build/test 0 |

2026-06-15 | HXC-095: CLI binary generate/debate/specify return 404 against live local ollama | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 89d0cd3c | desktop i18n/bundle.go + main() SetTranslator wiring; before raw keys/%!(EXTRA) -> after resolved prose 'HelixCode Desktop CLI (nogui mode)'; build/test 0, paired-mutation guard |

2026-06-15 | HXC-094: F12 workspace checkpoints — file snapshot + restore/undo safety net | Feature | Implemented (→ Fixed.md) | 2026-06-15 session | 3811508b | internal/checkpoint + cmd/cli / checkpoint; restore-bytes round-trip + survives-restart tests PASS, go build ./... 0 |

2026-06-15 | HXC-092: debate\_orchestrator 30s DefaultTimeout too

short for capable models on multi-round /specify | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | f8763c8a | debate\_orchestrator 659559e: DefaultTimeout 30s→180s (justified ~96s worst-case + headroom) + WithTimeout option; build/vet 0, 15-pkg suite ok. Live capable-model /specify re-verify is follow-up. |

2026-06-15 | HXC-091: containers custom health-check duration can be 0 (timer-resolution flake) | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | e95cd735 | containers a36a435: duration floor 1ns; test 0/20 fail post-fix, pkg/health ok |

2026-06-15 | HXC-089: panoptic web Element infinite-retry hang plus recorder zero-frames | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 2b8f46df | panoptic fcc6322: bounded Element timeout + real initial Screenshot frame; platforms PASS x3, full suite green |

2026-06-15 | HXC-088: llm\_orchestrator opencode cancel path hangs 30s — cmd.WaitDelay unset | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 2b8f46df | llm\_orchestrator c791f02: cmd.WaitDelay=2s; ContextCancel ok -count=3, pkg/agent ok |

2026-06-15 | HXC-087: skill\_registry randomString UnixNano same-tick produces identical chars and colliding IDs | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 2b8f46df | skill\_registry 5e5dc75: crypto/rand; tests pass -count=5 |

2026-06-15 | HXC-086: SSE broker client-ID UnixNano collision under concurrent connect | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | e4aa9a26 | streaming 3e15904: SSE fallback client-ID atomic counter; RED 134/1000 lost -> GREEN 1000/1000 unique under -race |

2026-06-15 | HXC-085: 14 LLM providers HealthCheck hardcodes production URL ignoring injected baseURL | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | e4aa9a26 | helix\_agent 8b622c7a (13) + llm\_provider 18108f4 (14): HealthCheck derives from injected baseURL (kimi pattern); existing tests were bluffs -> added real TestHealthCheck\_HonorsInjectedBaseURL, RED-proven; both trees build 0, suites ok |

2026-06-15 | HXC-084: challenge scripts use GNU-only grep -P backslash-K — breaks on macOS BSD grep | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | e4aa9a26 | challenges 280c2d2 + helix\_agent 8b622c7a: all grep -oP/-P/-> portable sed -nE/grep -oE/grep -F (incl. android\_save/cognee/runtime\_debate/mcps/helixmemory/output\_formatting), each proven on sample input + edge cases, bash -n 0 |

2026-06-15 | HXC-083: production\_deployer fabricates rollback env-prep server-validation and strategy differentiation | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | fde6bc0f | production\_deployer.go: 5 bluff sites honest (rollback/env/validate/strategy/monitoring); RED/GREEN polarity; build/test 0, go build ./... 0, smoke clean |

2026-06-15 | HXC-082: performance optimizer fabricates success — 8  
apply methods sleep and return Success true | Bug | Fixed (→  
Fixed.md) | 2026-06-15 session | fde6bc0f | optimizer.go: 8 methods  
Success:false+ErrOptimizationNotWired (honest), GC kept real;  
§11.4.120 reconcile + RED-polarity; build/test 0, go build ./... 0, smoke  
clean |

2026-06-15 | HXC-081: helix\_specifier speckit topic i18n key unresolved  
plus format-verb mismatch | Bug | Fixed (→ Fixed.md) | 2026-06-15  
session | ce097488 | helix\_specifier 188f9bc: BundleTranslator resolves  
phase-topic keys; GREEN prose ‘Create a detailed specification ...  
Request: ’, no raw key/no %!(EXTRA), 22 pkgs no-regression; /debate  
binary CONCLUSION now resolved prose |

2026-06-15 | HXC-080: /debate and /specify broken at runtime — single  
agent vs 2-min | Bug | Fixed (→ Fixed.md) | 2026-06-15 session |  
63b035fc | cmd/cli+TUI register 2 agents; specify\_e2e\_test.go LIVE-  
PROVEN Success=true qualityScore=0.86 real output PASS 14.69s  
(conductor podman ollama) |

2026-06-15 | HXC-079: debate\_orchestrator consensus emits unresolved  
i18n key | Bug | Fixed (→ Fixed.md) | 2026-06-15 session | 09b72280 |  
debate\_orchestrator 4df3874: embed-bundle translator resolves  
consensus key; GREEN test ‘Debate on completed across N round(s).’,  
RED/GREEN polarity §11.4.115 |

2026-06-14 | HXC-076: Web /llm/generate + /llm/stream endpoints with  
frontend (partial — e2e pending) | Feature | Implemented (→  
Fixed.md) | 2026-06-14 session | c850d9d6 | eafdda36 tests/integration/  
llm\_generate\_e2e\_test.go: real GIN 200 content:4 provider:ollama  
qwen2.5:0.5b, conductor-reproduced |

2026-06-14 | HXC-075: Phase-1 CLI-Agent Fusion plan reconciliation  
with delivered state | Task | Completed (→ Fixed.md) | 2026-06-14  
session | 7fc09724 | commit:e3063af1 Phase-1 plan reconciliation |

2026-06-14 | HXC-074: Mobile gomobile Generate binding for on-device  
LLM calls | Feature | Implemented (→ Fixed.md) | 2026-06-14 session  
| 7fc09724 | commit:28465071 mobile gomobile Generate binding |

2026-06-14 | HXC-073: Autocommit git substrate backing CLI edit  
history | Feature | Implemented (→ Fixed.md) | 2026-06-14 session |  
7fc09724 | commit:61d7167e autocommit git substrate |

2026-06-14 | HXC-072: CLI /undo and /diff slash commands over  
autocommit substrate | Feature | Implemented (→ Fixed.md) |  
2026-06-14 session | 7fc09724 | commit:bd5228f8 CLI /undo + /diff  
commands |

2026-06-14 | HXC-071: Web LLM handler httptest coverage for generate  
and stream | Task | Completed (→ Fixed.md) | 2026-06-14 session |  
7fc09724 | commit:6f382b95 web handler httptest coverage |

2026-06-14 | HXC-070: HelixMemory persist log no longer misreports

success on failure | Bug | Fixed (→ Fixed.md) | 2026-06-14 session | 7fc09724 | commit:a0239f52 honest HelixMemory persist log | 2026-06-14 | HXC-069: HelixMemory default-on durable persistence with graceful fallback | Feature | Implemented (→ Fixed.md) | 2026-06-14 session | 7fc09724 | commit:ac3ad237 HelixMemory default-on persist+fallback | 2026-06-14 | HXC-068: speckit debate adapter wireable into agentic debate flow | Feature | Implemented (→ Fixed.md) | 2026-06-14 session | 7fc09724 | commit:95b7385c speckit debate adapter wireable | 2026-06-09 | HXC-067: inner internal/redis stress suite reads TEST\_REDIS\_HOST/PORT (default :6379) not HELIX\_REDIS\_HOST/PORT | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 5f98eae9 | docs/qa/HXC-067/evidence.md | 2026-06-09 | HXC-066: inner internal/database integration tests hardcode localhost:5433/helix\_test, never read HELIX\_DATABASE\_\* env | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 5f98eae9 | docs/qa/HXC-066/evidence.md | 2026-06-09 | HXC-065: cache/pkg/postgres: finite-TTL Set invisible to immediate Get (expires\_at clock/timezone skew vs real PG) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 5f98eae9 | docs/qa/HXC-065/evidence.md | 2026-06-09 | HXC-064: cognee AMD-GPU parser tests flake under parallel load (rocm-smi fake subprocess signal-killed before 2s timeout, \$11.4.50) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 43521764 | docs/qa/HXC-064/evidence.md | 2026-06-09 | HXC-063: panoptic StartRecording: unreachable recording-bootstrap after early return nil — recorder never starts | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 6c7fead9 | docs/qa/HXC-063/evidence.md | 2026-06-09 | HXC-062: helix\_specifier pkg/metrics copies sync.RWMutex by value (vet lock-copy, concurrency hazard) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 6c7fead9 | docs/qa/HXC-062/evidence.md | 2026-06-09 | HXC-061: helix\_agent legacy unit-test calls memory.GetRelevant with stale 2-arg signature (won't compile) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 6c7fead9 | docs/qa/HXC-061/evidence.md | 2026-06-09 | HXC-060: debate\_orchestrator challenges/runner/main.go:516 context cancel not called on all return paths (vet leak) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 6c7fead9 | docs/qa/HXC-060/evidence.md | 2026-06-09 | HXC-059: debate\_orchestrator sandbox: ctx-cancel/timeout fails to kill child process tree on non-Linux (\$11.4.81) | Bug | Fixed (→

Fixed.md) | 2026-06-09 session | 6c7fead9 | docs/qa/HXC-059/  
evidence.md |  
2026-06-09 | HXC-031: Deferred long-tail: CONST-052 renames  
(RESOLVED — none remain) + Codex/Cline reference-agent ports |  
Task | Completed (→ Fixed.md) | 2026-06-09 session | fe4539f2 | docs/  
qa/HXC-031/evidence.md |  
2026-06-09 | HXC-039: G7 §11.4.83 docs/qa evidence gap: 8 past feature/  
fix commits lack docs/qa// directories | Task | Completed (→ Fixed.md)  
| 2026-06-09 session | 45362598 | docs/qa/HXC-039/evidence.md |  
2026-06-09 | HXC-058: helix\_agent go build fails on vendored third-  
party cli\_agents/continue test fixture (quarantine) | Bug | Fixed (→  
Fixed.md) | 2026-06-09 session | f4c5d7d6 | docs/qa/HXC-058/  
evidence.md |  
2026-06-09 | HXC-057: recovery go.mod missing require+replace for  
digital.vasic.concurrency (pkg/breaker import unwired) | Bug | Fixed  
(→ Fixed.md) | 2026-06-09 session | 11104a05 | docs/qa/HXC-057/  
evidence.md |  
2026-06-09 | HXC-056: 7 submodules: CONST-052 capitalised replace  
=> ../PliniusCommon (dir is plinius\_common) | Bug | Fixed (→  
Fixed.md) | 2026-06-09 session | 11104a05 | docs/qa/HXC-056/  
evidence.md |  
2026-06-09 | HXC-055: formatters brittle cat -version probe + go\_hello  
fixtures break go build | Bug | Fixed (→ Fixed.md) | 2026-06-09  
session | 94b46bfe | docs/qa/HXC-055/evidence.md |  
2026-06-09 | HXC-054: leak\_detector parallel test flake — §11.4.50  
determinism | Bug | Fixed (→ Fixed.md) | 2026-06-09 session |  
94b46bfe | docs/qa/HXC-054/evidence.md |  
2026-06-09 | HXC-051: helix\_llm + helix\_memory go.mod replace  
directives point to non-existent .././vasic-digital/\* sibling layout  
(CONST-051(C) dependency-layout) | Bug | Fixed (→ Fixed.md) |  
2026-06-09 session | 94b46bfe | docs/qa/HXC-051/evidence.md |  
2026-06-09 | HXC-038: docs\_chain G14: fixed.yaml fixed\_summary  
transform-contract mismatch + stale state.json baseline false-  
CONFLICT on issues context | Bug | Fixed (→ Fixed.md) | 2026-06-09  
session | ac4ce917 | docs/qa/HXC-038/evidence.md |  
2026-06-09 | HXC-053: conversation go.mod build break — capitalised  
replace path ../Messaging | Bug | Fixed (→ Fixed.md) | 2026-06-09  
session | ac4ce917 | docs/qa/HXC-053/evidence.md |  
2026-06-09 | HXC-052: background\_tasks go.mod build break —  
capitalised replace paths | Bug | Fixed (→ Fixed.md) | 2026-06-09  
session | ac4ce917 | docs/qa/HXC-052/evidence.md |  
2026-06-09 | HXC-050: event\_bus NATS env-gated integration skips lack  
SKIP-OK markers required by the no-silent-skips gate (§11.4.98) | Task  
| Completed (→ Fixed.md) | 2026-06-09 session | f0b0329d | docs/qa/



HXC-050/evidence.md |  
2026-06-09 | HXC-049: doc\_processor TestAutomation\_UpstreamsExist reads capital 'Upstreams' but canonical dir is lowercase 'upstreams' (CONST-052 case drift, deterministic FAIL) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | f0b0329d | docs/qa/HXC-049/evidence.md |  
2026-06-09 | HXC-048: helixcode-system.yaml HelixQA bank: 11 self-driving http cases for the non-auth server surface (health/server-info/system-status/llm-providers + negatives) | Feature | Implemented (→ Fixed.md) | 2026-06-09 session | 3e6d0830 | docs/qa/HXC-048/evidence.md |  
2026-06-09 | HXC-047: internal/redis TestNewClient\_WithDatabase needs-live-Redis with no SKIP-OK guard (§11.4.98) + i18n error no longer contains literal Redis | Task | Completed (→ Fixed.md) | 2026-06-09 session | 644019e5 | docs/qa/HXC-047/evidence.md |  
2026-06-09 | HXC-045: internal/hooks: cancelled hook ExecutionResult leaves duration unset (should always be populated) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | f1bf436f | docs/qa/HXC-045/evidence.md |  
2026-06-09 | HXC-046: internal/memory/providers: generateThreadID() non-unique under fast back-to-back calls (timestamp-only, same-nanosecond collision) | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 44a4d43f | docs/qa/HXC-046/evidence.md |  
2026-06-09 | HXC-043: auth Login nil-DB panic causes HTTP 500: server advertises graceful no-DB operation but first /api/v1/auth/login dereferences nil s.db | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | d6a37173 | docs/qa/HXC-043/evidence.md |  
2026-06-09 | HXC-042: CONST-050(B) challenge-coverage gap: 12 missing challenge scripts in debate\_orchestrator + helix\_agent (ddos/stress/chaos/scaling/ui/ux) | Task | Completed (→ Fixed.md) | 2026-06-09 session | 54ab4e95 | docs/qa/HXC-042/evidence.md |  
2026-06-09 | HXC-041: helixqa standalone HTTP bank-runner subcommand (helixqa http) drives http: banks vs live server without Playwright or LLM | Feature | Implemented (→ Fixed.md) | 2026-06-09 session | 7bc10191 | docs/qa/HXC-041/evidence.md |  
2026-06-09 | HXC-040: CLAUDE.md §9/§3.4 anti-bluff smoke command false-alarm (527 i18n/test hits) + case-sensitivity miss of BLUFF-001 | Bug | Fixed (→ Fixed.md) | 2026-06-09 session | 7bc10191 | docs/qa/HXC-040/evidence.md |  
2026-06-09 | HXC-037: §11.4.103-141 + CONST-048..060 anchor-cascade backfill into 7 owned submodules (verify-governance-cascade.sh 30 → 0) | Task | Completed (→ Fixed.md) | 2026-06-09 session | 7bc10191 | docs/qa/HXC-037/evidence.md |

*Last regenerated: 2026-05-20 (round 463 — HXC-003 closure: CONST-046 i18n migration campaign concluded — the genuine user-facing (C) string-literal surface is exhausted across all 7 scope areas (helix\_code internal /+cmd/+applications/, LLMsVerifier, helix\_qa, all owned vasic-digital/\*+HelixDevelopment/\* submodules); ~91-462 rounds migrated tens of thousands of literals through i18n seams with paired-mutation anti-bluff tests; the remaining ~55k audit-baseline hits are all out of CONST-046 scope per docs/audits/2026-05-20-internal-const046-classification.md. Closed Implemented (→ Fixed.md) per CONST-057). Previous round 403 — HXC-008/HXC-007/HXC-009 closures. Round 400 — HXC-006 closure (HelixCode Speed Programme — 6 phases / 31 tasks). Earlier closures (P0-P5 phases) tracked via docs/improvements/PROGRESS.md + docs/improvements/\*evidence\*.md. ## HXC-037 — §11.4.103-141 + CONST-048..060 anchor-cascade backfill into 7 owned submodules (verify-governance-cascade.sh 30 → 0)*

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/HXC-037/evidence.md **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

verify-governance-cascade.sh reported 30 FAIL: debate\_orchestrator/doc\_processor/event\_bus/helix\_agent/llm\_ops/llm\_orchestrator/llm\_provider each missing §11.9 + CONST-048..060 + §11.4.103-121 heading anchors in CONSTITUTION/CLAUDE/AGENTS.md. Authored deterministic additive scripts/backfill\_anchor\_cascade.sh (verbatim golden helix\_qa cascade, idempotent, §11.4.122 no-removal); backfilled+committed+pushed all 7 to origin; verifier now 0 FAIL PASS.

## **HXC-040 — CLAUDE.md §9/§3.4 anti-bluff smoke command false-alarm (527 i18n/test hits) + case-sensitivity miss of BLUFF-001**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-040/evidence.md **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

The documented anti-bluff grep one-liner prints BLUFF FOUND on a clean codebase (527 hits = 218 \_test.go + 123 i18n message-keys + 306 placeholder/template infra; 0 real production bluffs) AND was case-sensitive so it would miss the canonical ‘// For now, simulate generation’ (capital F). Refined to word-bounded case-insensitive

markers with `_test.go/i18n/comment-citation` exclusions; clean on real tree; §1.1 mutation-proven (planted 3 bluffs caught then reverted byte-identical). Both §3.4 and §9 updated.

## **HXC-041 — helixqa standalone HTTP bank-runner subcommand (helixqa http) drives http: banks vs live server without Playwright or LLM**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** docs/qa/HXC-041/evidence.md **Severity:** Medium **Created-By:** Claude  
**Assigned-To:** Claude

HelixQA could only run banks via Playwright (absent) or Ollama LLM (absent); the LLM-free HTTPExecutor that drives helixcode-auth.yaml's 16 http: cases against the live server was wired only internally. Built 'helixqa http -bank -base-url ' (cmd/helixqa/http.go +281, http\_test.go +285, 2 mutation tests); build+tests green; live run vs booted helixcode = 15/16 PASS exit 1. helix\_qa commit d6c084d6.

## **HXC-042 — CONST-050(B) challenge-coverage gap: 12 missing challenge scripts in debate\_orchestrator + helix\_agent (ddos/stress/chaos/scaling/ui/ux)**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/HXC-042/evidence.md **Severity:** Medium **Created-By:** Claude  
**Assigned-To:** Claude

verify-cascade-coverage.sh required 6 challenge scripts each in debate\_orchestrator + helix\_agent. Authored 12 REAL scripts (no stubs): concurrent flood w/ p50/p95, sustained-load degradation budget, /dev/tcp malformed+slowloris chaos, multi-replica sha256 body-identity, CLI panic/leak detection; honest SKIP-OK when no env target. bash -n 12/12 PASS; real DDoS run 200/200 ok. Committed debate\_orchestrator 19bd8e5b + helix\_agent 6eee57e1.

## **HXC-043 — auth Login nil-DB panic causes HTTP 500: server advertises graceful no-DB operation but first /api/v1/auth/login dereferences nil s.db**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-043/evidence.md **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

Found by HXC-041 live helixqa-http run (HXC-AUTH-003 expected 401 got 500 empty body). With db=nil (server's graceful no-DB path), helix\_code/internal/auth/auth.go:156 (\*AuthService).Login calls s.db.GetUserByUsername on nil s.db then nil-pointer panic then Gin Recovery then HTTP 500. Fix: guard nil s.db in Login and sibling db-touching auth paths, return clean 401/503. RED test exists: helixcode-auth.yaml HXC-AUTH-003 via helixqa http.

## **HXC-046 — internal/memory/providers: generateThreadID() non-unique under fast back-to-back calls (timestamp-only, same-nanosecond collision)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-046/evidence.md **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

Found by isolated-worktree full unit sweep (go test ./internal/... HEAD 54ab4e95, hermetic test, no infra). See docs/qa/HXC-046/evidence.md for the exact failing test + file:line + message. Genuine product defect reproducible deterministically.

## **HXC-045 — internal/hooks: cancelled hook ExecutionResult leaves duration unset (should always be populated)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-045/evidence.md **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

Found by isolated-worktree full unit sweep (go test ./internal/... HEAD 54ab4e95, hermetic test, no infra). See docs/qa/HXC-045/evidence.md for the exact failing test + file:line + message. Genuine product defect reproducible deterministically.

## **HXC-044 — internal/cognee: AMD-GPU rocm-smi JSON parser returns -1 sentinel instead of parsed GPU-use value**

**Status:** Obsolete (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-044/evidence.md **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude **Obsolete-Details:** Since: 2026-06-09; Reason: not-reproducible; Superseding-item: none; Triple-check evidence: docs/qa/HXC-044/evidence.md

Found by isolated-worktree full unit sweep (go test ./internal/... HEAD 54ab4e95, hermetic test, no infra). See docs/qa/HXC-044/evidence.md for the exact failing test + file:line + message. Genuine product defect reproducible deterministically.

## **HXC-047 — internal/redis TestNewClient\_WithDatabase needs-live-Redis with no SKIP-OK guard (§11.4.98) + i18n error no longer contains literal Redis**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/HXC-047/evidence.md **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

Hermetic unit run found this test silently depends on a live Redis at 127.0.0.1:6379 (no SKIP-OK §11.4.3/§11.4.98) AND asserts the error contains literal 'Redis' which the i18n-keyed error (internal\_redis\_failed\_connect) no longer contains. Fix: SKIP-OK guard when no Redis + reconcile assertion. Evidence docs/qa/HXC-047/evidence.md (HEAD 54ab4e95).

## **HXC-048 — helixcode-system.yaml HelixQA bank: 11 self-driving http cases for the non-auth server surface (health/server-info/system-status/llm-providers + negatives)**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** docs/qa/HXC-048/evidence.md **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

Authored + parse-validated a new LLM-free http: bank (banks/helixcode-system.yaml, 11 cases) covering /health, /api/v1/server/info, /api/v1/system/status(401), /api/v1/llm/providers + 404/405 negatives, using only helixqa-http runner-consumed fields. helixqa list → 11 cases; dry run fired real requests. Confident body asserts from captured responses; status-only/\_skip where unverified (§11.4.6). helix\_qa f18a5d3b. Live-run vs booted server queued.

## **HXC-049 — doc\_processor TestAutomation\_UpstreamsExist reads capital ‘Upstreams’ but canonical dir is lowercase ‘upstreams’ (CONST-052 case drift, deterministic FAIL)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-049/evidence.md **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

Owned-submodule health sweep found internal automation\_test.go:140 os.ReadDir(“Upstreams”) failing every run because the on-disk dir is lowercase ‘upstreams’ per CONST-052. Test-side fix “Upstreams” → “upstreams”. GREEN: ok digital.vasic.docprocessor. Commit ecb384f.

## **HXC-050 — event\_bus NATS env-gated integration skips lack SKIP-OK markers required by the no-silent-skips gate (\$11.4.98)**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/HXC-050/evidence.md **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

Owned-submodule health sweep found pkg/nats/integration\_test.go:23,120 env-gated t.Skip (legitimately runs vs real NATS when NATS\_URL is set) without literal SKIP-OK markers. Added 'SKIP-OK: #HXC-050 ...' to both; build+skip clean. Commit 1cae683.

## **HXC-052 — background\_tasks go.mod build break — capitalised replace paths**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-052/evidence.md **Severity:** Medium

submodules/background\_tasks/go.mod replace directives pointed at ../Concurrency and ../Models which no longer exist after the CONST-052 lowercase rename; go build ./... failed until corrected to ../concurrency and ../models.

## **HXC-053 — conversation go.mod build break — capitalised replace path ../Messaging**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-053/evidence.md **Severity:** Medium

submodules/conversation/go.mod line 24 replace digital.vasic.messaging pointed at ../Messaging (capitalised) which broke go build ./... after the CONST-052 lowercase rename; corrected to ../messaging.

## **HXC-038 — docs\_chain G14: fixed.yaml fixed\_summary transform-contract mismatch + stale state.json baseline false- CONFLICT on issues context**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-038/evidence.md **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

verify-all-constitution-rules.sh G14 (docs\_chain verify -all) fails: (1) fixed.yaml fixed\_summary node perpetually STALE because generate\_fixed\_summary.sh writes the file as a side-effect and prints a 'wrote ...' status line to stdout, while docs\_chain captures stdout as content — content is correct+deterministic, the node I/O contract needs a stdout mode or a writes-file declaration; (2) issues context reports a §11.4.6 CONFLICT (both issues\_md + items\_db dirty vs stale state.json baseline) though MD $\approx$ DB are verified byte-identical via db-to-md — needs a baseline refresh. governance context already in-sync this session.

## **HXC-051 — helix\_llm + helix\_memory go.mod replace directives point to non- existent ../../vasic-digital/\* sibling layout (CONST-051(C) dependency-layout)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-051/evidence.md **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

Owned-submodule health sweep (D-2): helix\_llm (internal/knowledge/embedding\_providers.go) + helix\_memory (pkg/provider/adapter.go) fail to build standalone because their go.mod 'replace' directives target ../../vasic-digital/ sibling dirs not materialized in the HelixCode layout. CONST-051(C) requires deps resolvable from the project root (submodules/). The other 8/24 packages build+test ok; only the replace-dep-needing packages fail. Investigation: confirm whether HelixCode's build actually compiles these (vs reference/standalone), then rewire replace paths to the root submodule layout OR document the standalone-build requirement. Not breaking helix\_code's own build. Found by D-2 health sweep.



## **HXC-054 — leak\_detector parallel test flake — §11.4.50 determinism**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-054/evidence.md **Severity:** Low

Discovery sweep: leak\_detector test exhibits non-deterministic PASS/FAIL under parallel execution (timing-sensitive), violating §11.4.50 determinism; needs forensic root-cause before a deterministic fix. Open.

## **HXC-055 — formatters brittle cat -version probe + go\_hello fixtures break go build**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-055/evidence.md **Severity:** Low

Discovery sweep: formatters test relies on brittle ‘cat -version’ (§11.4.81 cross-platform), and committed go\_hello fixture sources break a tree-wide go build; both need isolation/fixing. Open.

## **HXC-056 — 7 submodules: CONST-052 capitalised replace => ../PliniusCommon (dir is plinius\_common)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-056/evidence.md **Severity:** Medium

auto\_temp, claritas, gandalf\_solutions, hyper\_tune, leak\_hub, ouroborous, veritas each have go.mod line replace digital.vasic.pliniuscommon => ../PliniusCommon; capitalised dir absent, lowercase sibling plinius\_common exists; go build ./... fails on all 7.

## **HXC-057 — recovery go.mod missing require+replace for digital.vasic.concurrency (pkg/breaker import unwired)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-057/evidence.md **Severity:** Medium

recovery/pkg/breaker/breaker.go imports digital.vasic.concurrency/pkg/breaker but recovery/go.mod has no require/replace for the concurrency sibling; go build ./... fails: no required module provides package. Sibling submodules/concurrency provides pkg/breaker.

## **HXC-058 — helix\_agent go build fails on vendored third-party cli\_agents/continue test fixture (quarantine)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-058/evidence.md **Severity:** Low

helix\_agent go build ./... fails ONLY in vendored third-party cli\_agents/continue/ subtree (upstream continue project test fixture with bogus relative import); no owned dev.helix.agent package fails; needs build-exclusion/quarantine of the vendored fixture subtree, not an owned-code fix.

## **HXC-039 — G7 §11.4.83 docs/qa evidence gap: 8 past feature/fix commits lack docs/qa// directories**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/HXC-039/evidence.md **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

verify-all-constitution-rules.sh G7 (enforcing) reports 8 feature/fix commits since baseline ed84f90e without a docs/qa// evidence dir (81f3c482 deployment/perf, 83b2690a config var-expansion, d985e3ae worker consensus W6B, cee5cdae Phase-2 cascade, 5c5c44bc, c63c8963, 3ce30285). Retro-adding to those commits needs history-rewrite which

§11.4.113 forbids — operator decision on remediation (baseline reset vs documented exception). New work this session (HXC-037) ships its docs/qa evidence.

## **HXC-031 — Deferred long-tail: CONST-052 renames (RESOLVED — none remain) + Codex/Cline reference-agent ports**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/HXC-031/evidence.md

Deferred long-tail: CONST-052 renames (RESOLVED — none remain) + Codex/Cline reference-agent ports

## **HXC-059 — debate\_orchestrator sandbox: ctx-cancel/timeout fails to kill child process tree on non-Linux (§11.4.81)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-059/evidence.md **Severity:** Medium

testing/sandbox\_other.go (!linux) killProcessGroup is a no-op so Setpgid is never set; on macOS cmd.Cancel SIGKILLs only the direct child and the sleep-30 grandchild survives. TestSandboxExecute\_CtxCancel + TestSandboxExecute\_TimeoutEnforced FAIL deterministically (elapsed ~30s vs 100ms cap). Linux process-group kill has no functioning non-Linux equivalent (§11.4.81 parity gap).

## **HXC-060 — debate\_orchestrator challenges/runner/main.go:516 context cancel not called on all return paths (vet leak)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-060/evidence.md **Severity:** Low

go vet: challenges/runner/main.go:516 the cancel function is not used on all paths (possible context leak); 571 return may be reached without using the cancel var defined on line 516. Owned-code vet finding.

## **HXC-061 — helix\_agent legacy unit-test calls memory.GetRelevant with stale 2-arg signature (won't compile)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-061/evidence.md **Severity:** Medium

tests/unit/debate\_security\_legacy/debate\_security\_test.go:335 calls memory.GetRelevant(string, number) but the current signature is (context.Context, string, int); go vet of the owned test tree fails to compile. Stale API call in test code.

## **HXC-062 — helix\_specifier pkg/metrics copies sync.RWMutex by value (vet lock-copy, concurrency hazard)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-062/evidence.md **Severity:** Medium

go vet: pkg/metrics/metrics.go:143 assignment copies lock value to cp; :163 return copies lock value — Metrics struct contains sync.RWMutex copied by value. Genuine owned-code concurrency hazard; build+tests pass but the copied mutex does not protect the original.

## **HXC-063 — panoptic StartRecording: unreachable recording-bootstrap after early return nil — recorder never starts**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-063/evidence.md **Severity:** Medium

internal/platforms/desktop.go:304 unconditional return nil makes lines 305+ dead (go vet: 305:2 unreachable code): os.MkdirAll video-dir creation + background recorder process startup never execute, so StartRecording returns success without recording. Latent correctness defect; investigate per §11.4.124 (likely restore by removing the early return, not delete the block).

## **HXC-064 — cognee AMD-GPU parser tests flake under parallel load (rocm-smi fake subprocess signal-killed before 2s timeout, §11.4.50)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-064/evidence.md **Severity:** Low

internal/cognee TestProbeAMDGPU\_HandlesAltKeyName + \_GpuUtilization call queryAMDGPUUsage() which execs a fake rocm-smi via a 2s const timeout; under heavy batch/parallel host load the echo subprocess is signal-killed before completing → product correctly returns sentinel -1 but the parser tests assert 33/77 → non-deterministic FAIL. Product correct; test timeout load-fragile. Fix: make rocmSmiQueryTimeout an overridable var (prod default unchanged) + parser tests raise it.

## **HXC-065 — cache/pkg/postgres: finite-TTL Set invisible to immediate Get (expires\_at clock/timezone skew vs real PG)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-065/evidence.md **Severity:** Medium

digital.vasic.cache pkg/postgres integration\_test.go:195 — a value Set with a finite TTL (200ms) returns empty on an immediate Get even before expiry, deterministic across -count=3 vs real booted PG. Siblings TestSetGet/Exists/ZeroTTL pass, so isolated to the finite-TTL expires\_at WHERE-clause: likely Go-process time.Now() vs PG server now() timezone/clock skew making the just-written row appear already-expired. Real cache-backend correctness defect.

## **HXC-066 — inner internal/database integration tests hardcode localhost:5433/helix\_test, never read HELIX\_DATABASE\_\* env**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-066/evidence.md **Severity:** Low

helix\_code/internal/database/database\_integration\_test.go hardcodes Config{Host:localhost,Port:5433,User:helix\_test,DBName:helix\_test} (lines 19-20,43-46,330-333) with zero env sourcing; port 5433 closed → internal\_database\_ping\_failed against booted PG (15432). DB layer sound (persistence passes). Harness defect: should read DB\_/HELIX\_DATABASE\_ env.

## **HXC-067 — inner internal/redis stress suite reads TEST\_REDIS\_HOST/PORT (default :6379) not HELIX\_REDIS\_HOST/PORT**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/HXC-067/evidence.md **Severity:** Low

helix\_code/internal/redis/redis\_stress\_test.go:38-39 reads TEST\_REDIS\_HOST/TEST\_REDIS\_PORT (default localhost:6379) instead of the standard HELIX\_REDIS\_HOST/HELIX\_REDIS\_PORT contract; causes false 100%-error FAIL against booted Redis on 16379. Pointed at TEST\_REDIS\_PORT=16379 it's GREEN. Env-var-contract inconsistency (harness).

## **HXC-068 — speckit debate adapter wireable into agentic debate flow**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** commit:95b7385c speckit debate adapter wireable **Severity:** High  
**Created-By:** Claude **Assigned-To:** Claude

Wire the speckit debate adapter so the agentic debate flow can invoke it end-to-end; adapter is constructable and dispatchable from the orchestrator (commit 95b7385c).

## **HXC-069 — HelixMemory default-on durable persistence with graceful fallback**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** commit:ac3ad237 HelixMemory default-on persist+fallback **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

HelixMemory persists cross-session memory by default and falls back gracefully when the backend store is unavailable, so recall works out of the box (commit ac3ad237).

## **HXC-070 — HelixMemory persist log no longer misreports success on failure**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** commit:a0239f52 honest HelixMemory persist log **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

The HelixMemory persistence log now reports the real outcome instead of an unconditional success line, removing a \$11.4 honest-logging bluff (commit a0239f52).

## **HXC-071 — Web LLM handler httpstest coverage for generate and stream**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** commit:6f382b95 web handler httpstest coverage **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

Add httpstest-based handler tests exercising the web /llm/generate and /llm/stream endpoints with real request/response round-trips (commit 6f382b95).

## **HXC-072 — CLI /undo and /diff slash commands over autocommit substrate**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** commit:bd5228f8 CLI /undo + /diff commands **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

Implement CLI /undo and /diff commands that revert and show changes against the git autocommit substrate, giving users real edit history control (commit bd5228f8).

## **HXC-073 — Autocommit git substrate backing CLI edit history**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Severity:** Medium  
**Created-By:** Claude **Assigned-To:** Claude

The CLI keeps an automatic git-backed history of every edit it makes so users can review and safely roll back changes; this item established that autocommit substrate underneath the edit history.

## **HXC-074 — Mobile gomobile Generate binding for on-device LLM calls**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** commit:28465071 mobile gomobile Generate binding **Severity:** Medium  
**Created-By:** Claude **Assigned-To:** Claude

Expose a gomobile-bound Generate entrypoint so the mobile applications can invoke the LLM provider through the shared client core (commit 28465071).

## **HXC-075 — Phase-1 CLI-Agent Fusion plan reconciliation with delivered state**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** commit:e3063af1 Phase-1 plan reconciliation **Severity:** Low  
**Created-By:** Claude **Assigned-To:** Claude

Reconcile the Phase-1 implementation plan against actually-delivered state so the programme plan reflects reality, not aspiration (commit e3063af1).



## **HXC-076 — Web /llm/generate + /llm/stream endpoints with frontend (partial — e2e pending)**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** eafdda36 tests/integration/llm\_generate\_e2e\_test.go: real GIN 200 content:4 provider:ollama qwen2.5:0.5b, conductor-reproduced **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

Backend /llm/generate and /llm/stream endpoints plus frontend wiring landed (commit 32e7e5b8); REMAINING: full browser-driven end-to-end test proving real streamed output renders in the UI is not yet captured — keep open until e2e evidence lands.

## **HXC-079 — debate\_orchestrator consensus emits unresolved i18n key**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** debate\_orchestrator 4df3874: embed-bundle translator resolves consensus key; GREEN test ‘Debate on completed across N round(s).’, RED/GREEN polarity §11.4.115 **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

Live /debate e2e (debate\_e2e\_test.go) shows CONCLUSION/Summary print literal ‘debate.orchestrator.consensus\_conclusion’ i18n message-key instead of resolved prose; per-agent LLM content is real, consensus synthesis layer in submodules/debate\_orchestrator does not resolve the key. §11.4.118 discovery finding.

## **HXC-080 — /debate and /specify broken at runtime — single agent vs 2-min**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** cmd/cli+TUI register 2 agents; specify\_e2e\_test.go LIVE-PROVEN Success=true qualityScore=0.86 real output PASS 14.69s (conductor podman ollama) **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

handleDebate/handleSpecify (cmd/cli) + TUI registered ONE AgentSpec but orchestrator MinAgentsPerDebate=2; users hit ‘insufficient agents (have 1, need 2)’. Round-7 debate proof used a 2-agent test, masking the 1-agent production gap (§11.4.108).

## HXC-081 — helix\_specifier speckit topic i18n key unresolved plus format-verb mismatch

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** helix\_specifier 188f9bc: BundleTranslator resolves phase-topic keys; GREEN prose 'Create a detailed specification ... Request: ', no raw key/no %!(EXTRA), 22 pkgs no-regression; /debate binary CONCLUSION now resolved prose **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

Live /specify run shows 'Topic: helixspecifier\_speckit\_topic\_specify%!(EXTRA string=...)' — the speckit phase prompt emits a raw i18n message-key with a Go Sprintf arg-count mismatch instead of resolved prose. Same class as HXC-079, in submodules/helix\_specifier. Captured in specify\_e2e\_test.go output.

## HXC-082 — performance optimizer fabricates success — 8 apply methods sleep and return Success true

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** optimizer.go: 8 methods Success:false+ErrOptimizationNotWired (honest), GC kept real; §11.4.120 reconcile + RED-polarity; build/test 0, go build ./... 0, smoke clean **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

internal/performance/optimizer.go:540-760: applyCPU/Memory/Concurrency/Cache/Network/Database/Worker/LLM Optimization each time.Sleep(200ms) doing NO real work then return Success:true with fabricated MetricsChange. User-reachable via cmd/performance\_optimization/main.go:89. Rule 2 / §11.4 bluff. Fix: real tuning OR honest ErrOptimizationNotWired sentinel.

## HXC-083 — production\_deployer fabricates rollback env-prep server-validation and strategy differentiation

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:**

production\_deployer.go: 5 bluff sites honest (rollback/env/validate/strategy/monitoring); RED/GREEN polarity; build/test 0, go build ./... 0, smoke clean **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

internal/deployment/production\_deployer.go: triggerRollback(1028) sleeps+logs success no real rollback;  
prepareEnvironment(810)+validateTargetServers(820) sleep+log success; executeBlueGreen/Canary/Rolling/Recreate(962) all just call executeProductionDeploy (no-op differentiation);  
executeMonitoring(758) ends with success notification contradicting its honest gap-log. §11.4 bluffs. Fix: real work or honest sentinels.

## HXC-084 — challenge scripts use GNU-only grep -P backslash-K — breaks on macOS BSD grep

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** challenges 280c2d2 + helix\_agent 8b622c7a: all grep -oP/-P/-> portable sed -nE/grep -oE/grep -F (incl. android\_save/cognee/runtime\_debate/mcps/helixmemory/output\_formatting), each proven on sample input + edge cases, bash -n 0 **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

29 owned challenge/CI shell scripts use grep -oP / grep -P / (GNU/PCRE-only). Stock macOS /usr/bin/grep rejects -P (invalid option) -> empty evidence capture -> corrupted PASS/FAIL gates (false FAIL or silently-masked). §11.4.81. Fix: portable sed -E/awk/perl. Highest: challenges android\_save, helix\_agent cognee/runtime\_debate/mcps.

## HXC-085 — 14 LLM providers HealthCheck hardcodes production URL ignoring injected baseURL

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** helix\_agent 8b622c7a (13) + llm\_provider 18108f4 (14): HealthCheck derives from injected baseURL (kimi pattern); existing tests were bluffs -> added real TestHealthCheck\_HonorsInjectedBaseURL, RED-proven; both trees build 0, suites ok **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

openai/groq/cohere/fireworks/ai21/chutes/nvidia/publicai/replicate/together/cerebras/deepseek/mistral/claude HealthCheck hits a hardcoded \*ModelsURL const/literal while Generate uses p.baseURL — config-injection / CONST-051(B), breaks httptest + proxy/Azure endpoints. Duplicated in helix\_agent/internal/llm/providers + llm\_provider/pkg/providers. Fix: mirror the existing kimi.go derive-from-baseURL fix.

## HXC-086 — SSE broker client-ID UnixNano collision under concurrent connect

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** streaming 3e15904: SSE fallback client-ID atomic counter; RED 134/1000 lost -> GREEN 1000/1000 unique under -race **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

submodules/streaming/pkg/sse/sse.go:140 clientID=fmt.Sprintf('client-%d',UnixNano()) used as b.clients map key, generated per concurrent SSE connect — same-tick collision overwrites/loses a client. Fix: crypto/rand or atomic counter suffix.

## HXC-087 — skill\_registry randomString UnixNano same-tick produces identical chars and colliding IDs

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** skill\_registry 5e5dc75: crypto/rand; tests pass -count=5 **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

skill\_registry/types.go:173 randomString used charset[UnixNano%len] in a tight loop -> all-identical chars + colliding execution IDs. Fixed with crypto/rand. Proven class.

## **HXC-088 — llm\_orchestrator opencode cancel path hangs 30s — cmd.WaitDelay unset**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** llm\_orchestrator c791f02: cmd.WaitDelay=2s; ContextCancel ok -count=3, pkg/agent ok **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

llm\_orchestrator pkg/agent/opencode\_agent.go runCapture set cmd.Cancel but WaitDelay==0 -> on ctx-cancel Wait() blocks on pipe drainage when a grandchild holds stdout (30s hang vs 5s test). Fixed cmd.WaitDelay=2s.

## **HXC-089 — panoptic web Element infinite-retry hang plus recorder zero-frames**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** panoptic fcc6322: bounded Element timeout + real initial Screenshot frame; platforms PASS x3, full suite green **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

panoptic internal/platforms: web.go Fill/Click/Submit page.Element on unbounded ctx -> missing selector retries forever (9m hang); screencast.go relied only on async CDP events -> 0 frames on immediate start/stop. Fixed: bounded page.Timeout + synchronous initial Screenshot frame.

## **HXC-091 — containers custom health-check duration can be 0 (timer-resolution flake)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** containers a36a435: duration floor 1ns; test 0/20 fail post-fix, pkg/health ok **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

containers pkg/health/custom.go NewCustomCheckFunc  
duration=time.Since(start) returns 0 for a no-op check ->  
TestNewCustomCheckFunc\_Success NotZero flake. Fixed: floor to  
time.Nanosecond when <=0 (no fabricated delay).

## **HXC-092 — debate\_orchestrator 30s DefaultTimeout too short for capable models on multi-round /specify**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** debate\_orchestrator  
659559e: DefaultTimeout 30s->180s (justified ~96s worst-case +  
headroom) + WithTimeout option; build/vet 0, 15-pkg suite ok. Live  
capable-model /specify re-verify is follow-up. **Severity:** Medium  
**Created-By:** Claude **Assigned-To:** Claude

debate\_orchestrator DefaultTimeout=30s x DefaultMaxRounds=3  
(types.go:41-42) is tuned for fast qwen2.5:0.5b. A capable qwen2.5:3b  
(~16s/round) blows the 30s cap on the 3-round Specify pillar -> context  
deadline exceeded. /debate works (WithMaxRounds(1), rich quality  
0.875 proven). Fix: raise per-debate timeout for the speckit Specify use  
case (adapter WithTimeout or orchestrator default). Tunable, not a  
code defect; surfaced honestly (no fabrication).

## **HXC-094 — F12 workspace checkpoints — file snapshot + restore/undo safety net**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** internal/  
checkpoint + cmd/cli /checkpoint; restore-bytes round-trip + survives-  
restart tests PASS, go build ./... 0 **Severity:** Medium **Created-By:** Claude  
**Assigned-To:** Claude

internal/checkpoint Manager (git-plumbing + file-copy backends) + /  
checkpoint create/list/restore CLI command: snapshot working-tree file  
contents and restore real bytes later (the cli\_agents F12 oops-revert  
net). Existing checkpoints were task-DB rows, not file snapshots.

## **HXC-095 — CLI binary generate/debate/ specify return 404 against live local ollama**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** desktop i18n/bundle.go + main() SetTranslator wiring; before raw keys/%!(EXTRA) -> after resolved prose 'HelixCode Desktop CLI (nogui mode)'; build/test 0, paired-mutation guard **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

helixcli -prompt/-stream + /debate + /specify hit 'API request failed: API returned status 404' against a working local ollama (qwen2.5:3b on :11434). The web POST /llm/generate + integration tests work with the same NewOllamaProvider — so the CLI's default-local-provider path uses a wrong endpoint/model-name (§11.4.108 different-path gap). AI features are broken for the end user via the binary. Found while recording feature videos.

## **HXC-096 — desktop nogui prints raw i18n keys + %!(EXTRA) format mismatch in status/help**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** cmd/cli/main.go: known-provider-prefix guard on model parsing; live qwen2.5:3b -> real answer '4', zero 404; build/test 0 **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

applications/desktop/main\_nogui.go status/help output prints raw message keys (desktop\_cli\_status\_header, desktop\_cli\_help\_body) + a Printf arg-count mismatch (%!(EXTRA int=0...)) in status. Same i18n-resolution class as HXC-079/081. CLI binary unaffected. Found while assessing desktop for video.

## **HXC-097 — SYSTEMIC: standalone binaries + internal/config + internal/database never wire i18n Translator -> raw keys at runtime**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Systemic unwired-translator fixed across aurora\_os (842551d3) + harmony\_os (6dcf64aa) + internal/config + internal/database: real bundle translator wired (binaries: SetTranslator in main(); libs: init() default). Before/after raw-

key->prose captured each; §11.4.115 guards. Broader follow-up: other CONST-046 pkgs may share the WireAll-only-on-CLI-path class (init()-default pattern is the fleet fix). **Severity:** High **Created-By:** Claude **Assigned-To:** Claude

Same unwired-translator bug as HXC-095 found across the fleet: aurora\_os standalone nogui (aurora\_os\_cli\_version\_banner + %!(EXTRA) at runtime) — round-7's aurora/harmony 'i18n fix' added KEYS but never wired SetTranslator in main(), so keys still echo raw (§11.4.108 fixed-in-source-not-at-runtime). Also internal/config (internal\_config\_info\_using\_config\_file) + internal/database (internal\_database\_ping\_failed) echo raw keys in CLI output. Fix: wire a real Translator (embed-bundle pattern) at each binary's main()/package init; add runtime guards that assert resolved prose (not just key-presence).

## **HXC-090 — panoptic tracks test-generated audit.json users.json (CONST-053 hygiene)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** panoptic a77228e: enterprise TestMain runs from temp dir; tree stays clean post-test x2, tests PASS, no production change. Guard via §11.4.135 (HXC-096) committed separately. **Severity:** Low **Created-By:** Claude **Assigned-To:** Claude

panoptic internal/enterprise/{audit,users}.json are version-tracked but overwritten by the enterprise test suite every run (timestamps/random IDs) -> perpetual dirty tree. CONST-053: test-generated data should be gitignored + a fixture template used. Pre-existing, low severity.

## **HXC-093 — helix\_code module graph has phantom digital.vasic.\* requires + private transitive blocking go list -m all / gomobile bind**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** helix\_code/go.mod: 27 replace directives -> local submodules; go list -m all no-such-host 26->0 (462 modules); REAL .aar produced (classes.jar + jni/libgojni.so x4 ABIs); go build ./... green throughout **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude



gomobile bind fails because go list -m -json all errors: ~22 digital.vasic. {cache,database,eventbus,...,vectordb} module paths are required with NO replace + NO remote ('no such host'), and github.com/HelixDevelopment/helix\_agent/Toolkit (private, separate from the replaced dev.helix.agent) needs interactive git creds. go build ./... works (imported subset only); full-graph tooling (gomobile, go list -m all) is blocked. Fix (repo-side, careful — editing go.mod risks the build): add replace directives for the phantom paths OR prune them, make Toolkit resolvable (replace/GOPRIVATE+SSH), persist x/mobile as a tool directive. Toolchain+NDK+Xcode all present; gobind codegen already works -> artifact achievable once graph resolves.

## HXC-098 — out-of-box config fails 'version required' validation — blocks client status/system commands

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Reproduced via LoadHelixConfig path (RED: version-less config.json -> Version="" + server.port=0 -> validateConfig rejects). Fixed in internal/config/config.go loadConfigLocked: decode JSON ON TOP of getDefaultConfig() so all viper defaults merge. Guard hxc098\_version\_default\_test.go: RED pre-fix (exit 1), GREEN post-fix (Version=1.0.0,Port=8080). config.yaml ships explicit version. Full config pkg ok. **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

Default config/config.yaml + the operator's ~/.config/helixcode/config.json have no top-level Version -> config validation fails with 'version is required' (internal\_config\_validate\_version\_required), blocking desktop/aurora/harmony status/system/version (and CLI subsystems) for a fresh user. Found while recording client videos (had to use a throwaway minimal valid config). Fix: ship a valid default Version in config/config.yaml (+ docs), or default it in config.Load when absent.

## HXC-077 — T1.5 context-window percentage indicator (partial)

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** TUI status bar now renders honest context-window USED-% (commit 6e03fe15). sessionUsedTokens accumulates real per-turn tokens; window from catalogue ContextSize->GetContextWindow; OMITS when unknown

(CONST-035); label i18n-routed via new terminal\_ui\_chat\_context\_usage key (CONST-046). Guard context\_usage\_test.go: GREEN with key, RED without (raw-key echo exit 1), restored GREEN. Full terminal\_ui pkg ok; build exit 0. **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

Context-window usage percentage indicator scaffolded (commit 3c9d3495, task T1.5); REMAINING: live token-accounting wiring + UI verification across TUI/desktop still pending — keep open until the indicator reflects real per-session usage.

## HXC-078 — T1.6 SKILL.md precedence resolution (partial)

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** FindMatching/List already resolve overlapping matches deterministically by lexicographic name (sort.Strings, documented contract markdown\_skills.go:194) — gap was missing coverage, not a bug. Added TestSkillRegistry\_FindMatching\_OverlappingPatternsDeterministic (insertion-order independence + 50-iter stability + lexicographic order), PASS; full internal/commands pkg ok; build exit 0. Commit 6e03fe15. **Severity:** Medium **Created-By:** Claude **Assigned-To:** Claude

SKILL.md precedence ordering partially implemented (commit 51302bf8, task T1.6); REMAINING: full precedence-conflict resolution + tests covering overlapping skill definitions still outstanding — keep open until precedence is fully resolved and tested.

## HXC-100 — Resync docs/CONTINUATION.md to current HEAD + de-bloat the 32k-token line-1 header (CONST-044/\$12.10 + CONST-064 hygiene)

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** CONTINUATION.md resynced to HEAD 3aacfa9f + line-1 header de-bloated (max line 54856->2931 chars). History preserved: 4 mega-prose lines replaced by CONST-064 metadata table + ToC + condensed close-out 143-169 table; close-outs 131-142 dedicated sections intact; this

session's 5 rows added. git diff +143/-8 (additive, no content loss); exports rendered=2 failed=0 fresh. Gated independently by conductor.  
**Created-By:** Claude **Assigned-To:** Claude

CONTINUATION.md is stale (Last updated 2026-06-14, refs HEAD e3063af1; current is 80e62afa after HXC-098 fix + HXC-099 i18n-sweep finding). Per CONST-044/§12.10 out-of-sync CONTINUATION is a CRITICAL DEFECT. ALSO: line 1 ('Last updated' header) has accreted ~32k tokens into a SINGLE line across rounds — pathological; Read/Edit of lines 1-12 alone exceeds 25k tokens, making safe edits hard and risking corruption with blind sed. Overnight zero-risk policy => queued for careful daytime fix: (1) refactor the bloated header into a normal metadata table + a round-by-round table row, (2) add rounds for this session (B i18n sweep discarded+stashed -> HXC-099; HXC-098 config-default fix), (3) restore CONST-064 ToC parity, (4) regen .html/.pdf. Live resumption currently served by the up-to-date .remember/remember.md (§11.4.131).

## **HXC-101 — security/security\_test.go**

### **TestTLSConfiguration — external-network dependency + nil-deref panic crashes the whole security test binary**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Replaced live httpbin.org call with hermetic httpptest.NewTLSServer + t.Fatalf on error path (no nil-deref fall-through). Verified: TestTLSConfiguration PASS 3/3 (-count=3) deterministic, no external net; full security pkg ok (0.223s); go build ./... exit 0; gofmt clean. **Created-By:** Claude **Assigned-To:** Claude

Discovery sweep (§11.4.118) found the only failing package in a ~270-pkg sweep. TestTLSConfiguration called live https://httpbin.org/get (non-deterministic, §11.4.98) and on the error path did t.Errorf without return -> defer resp.Body.Close() with resp==nil -> SIGSEGV panic that crashed the entire security test binary (took every other test in the package down). Fixed: hermetic httpptest.NewTLSServer + t.Fatalf on error path.

## **HXC-099 — Systemic i18n raw-key sweep redo (CONST-046) — corrected, regression-free, with default-translator contract decision**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** Corrected redo per operator decision (preserve loud raw-key NoopTranslator default — NO default swap; 9 guards pass unchanged). GOAL A: WireAll() was only called from cmd/cli; added entry-path init() wiring to cmd/server + 4 apps so internal-package strings resolve for real users. GOAL B: removed redundant {{.Err}} placeholder from 8 internal/project messages (nil-data -> “). RED → GREEN guards captured (non-tautology proven); 9 guards green -count=1; all touched pkgs pass; build exit 0; vet clean; no mutation residue. Commit a02a8aa8. (B’s rejected default-swap approach preserved in git stash@{0}.) **Created-By:** Claude **Assigned-To:** Claude

First mechanical sweep (stash@{0}, agent a55802ad) wired a real embedded-English bundle translator as package default across 36 internal/ packages but INTRODUCED 13 regressions vs green HEAD d85f6962: (1) real Go-template bug — internal/project messages render “ (error-detail param dropped); (2) defeats intentional NoopTranslator-echoes-raw-key anti-bluff guards in 9 pkgs (tools,voice,plantree,repomap,context,hardware,persistence,mcp,template); (3) autocommit+project tests assert real message text now broken. Redo MUST fix templating + needs operator decision: missing i18n key echoes raw key (loud default) OR falls back to embedded English (polished, risks hiding missing translations). Work in git stash; green tree restored (build exit 0, 13/13 pass).

## **HXC-102 — harmony\_os main\_nogui.go — 2 user-facing strings (‘Goodbye!’, ‘Error: %v’) bypass i18n (CONST-046, low severity)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Routed cmdInteractive Goodbye!/Error: %v through cliApp.tr with new bundle keys (error binds {{.Error}}, no ). Guard hxc102\_interactive\_i18n\_test.go (nogui): GREEN, RED-without-key (raw-

key leak exit 1), restored GREEN. Full harmony\_os pkg ok both tag variants; build exit 0; gofmt clean. **Created-By:** Claude **Assigned-To:** Claude

Discovery sweep: applications/harmony\_os/main\_nogui.go uses its i18n bundle heavily (~95 refs) but lines 876 ('Goodbye!') + 882 ('Error: %v') are user-facing UI text printed via fmt without the translator (CONST-046 localization gap). Lines 784/789-793 are developer-facing diagnostics (arguably out of scope). Low severity; may be folded into the HXC-099 entry-path i18n work.

## **HXC-104 — streamLLM /api/v1/llm/stream hangs forever — chunkChan never closed, [DONE] never emitted (production defect found by web e2e)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Fixed via 'defer close(chunkChan)' in streamLLM goroutine. Regression guard tests/integration/llm\_stream\_e2e\_test.go (TestLLMStreamE2E): post-fix streams 9 SSE frames + [DONE] in 1.1s, deterministic -count=3; server unit pkg ok; build exit 0. Evidence docs/qa/web-llm-e2e-20260615/. **Created-By:** Claude **Assigned-To:** Claude

streamLLM goroutine (internal/server/llm\_generate.go) called provider.GenerateStream(ctx, llmReq, chunkChan) but never closed chunkChan, so streamProviderToSSE never saw ok=false, never wrote the terminal data:[DONE] frame, and EVERY /stream request blocked until the 120s ctx deadline — a real user-facing hang that handler/http test tests missed. Exposed by the new TestLLMStreamE2E runtime e2e. Fixed: defer close(chunkChan).

## **HXC-103 — Web-client runtime e2e proof — live browser/HTTP -> server -> LLM provider round-trip for /api/v1/llm/generate + /llm/stream (CONTINUATION honest gap)**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** All 3 web LLM paths proven e2e against live ollama qwen2.5:3b: /generate (HTTP 200 content:4 provider:ollama), /stream (9 SSE frames + [DONE], streamed 1..5, >1 frame proves streaming), browser->server->provider

(chromedp: #output DOM=4, #meta provider=ollama, screenshot).  
Exposed+fixed production hang HXC-104. Evidence + README docs/qa/  
web-llm-e2e-20260615/. SKIP-OK when provider down (§11.4.3).  
**Created-By:** Claude **Assigned-To:** Claude

The web POST /llm/generate + /llm/stream endpoints + minimal web  
frontend are httptest-handler-verified only; NO full runtime e2e (no  
live client->server->provider round-trip captured) per §11.4.108  
layer-3/4. Deliver a real automated e2e (boot server, real HTTP/  
chromedp client hits the endpoints, REAL provider responds, capture  
the round-trip evidence). SKIP-OK per §11.4.3 when no real provider  
reachable (CONST-050(A) real-infra mandate) — never a fake PASS.

## **HXC-105 — Runtime e2e for server POST / api/v1/specify — boot server -> real spec output via live provider (speckit HTTP- endpoint gap)**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** tests/  
integration/specify\_server\_e2e\_test.go boots the real server + POSTs /  
api/v1/specify against live ollama qwen2.5:3b: HTTP 200 status:success  
provider:ollama qualityScore:0.9808, real 3-round 2-agent debate,  
provider\_calls=6 total\_tokens=806; output non-empty + NOT the  
‘awaiting provider wiring’ stub. PASS 75.93s, vet clean, build exit 0.  
Evidence docs/qa/web-llm-e2e-20260615/. **Created-By:** Claude  
**Assigned-To:** Claude

specify\_e2e\_test.go + debate\_e2e\_test.go exercise the speckit path  
provider-direct (pillar.ExecutePhase / responder.Generate); NO e2e  
boots the real HTTP server and POSTs to /api/v1/specify (server  
specifyHandler). Add one mirroring llm\_generate\_e2e\_test.go: boot  
server.New, POST a real request, assert a genuine 200 + non-fabricated  
phase output from live ollama; honest 502/SKIP otherwise.

## **HXC-106 — helix\_agent durable memory: process-lifetime in-memory fallback is NOT disk-durable — recall lost on restart (CONTINUATION honest gap)**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** Investigated (§11.4.102): disk-durable DiskStore (sqlite, survives close+reopen) was ALREADY implemented + wired as preferred fallback (commits ac3ad237/a91faad6) via debateMemoryFallbackPath() (os.UserCacheDir, 'helixagent'-namespaced, CONST-051-decoupled). CONTINUATION 'in-memory only' gap was stale. Closed the test-coverage gap: new internal/services/debate\_memory\_fallback\_test.go proves resolver returns writable durable path + persist->Close(restart)->reopen->RECALL. ./internal/memory + ./internal/services pass; submodule HEAD c5bdcfad pushed to upstreams. **Created-By:** Claude **Assigned-To:** Claude

helix\_agent memory persists by default with a local fallback, but the fallback is process-lifetime in-memory only: recall survives within a process but is LOST on restart unless a durable backend is configured. Investigate the fallback in the helix\_agent submodule; make it disk-durable (e.g. local file/sqlite-backed store) so recall survives restart out-of-the-box, OR document precisely why not + the required backend. Submodule work: own commit + push discipline.

## **HXC-109 — Mobile apps are scaffolds — Android has no build.gradle/AndroidManifest, iOS has no Xcode project (not buildable -> not recordable)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Android: buildable Gradle project + 67MB APK runs on live Genymotion (3 videos: connect/lifecycle/tasklist, real server task list via authenticated HTTP); 2 runtime issues fixed (JWT client-mode, JSON parse). iOS: buildable Xcode project (gomobile xcfamework + rewired binding) builds+runs on iPhone14 sim (helixcode-ios-launch video, Go core OK). Both committed (1ffc9b69/38caa48d). Mobile apps no longer scaffolds. **Created-By:** Claude **Assigned-To:** Claude

Inventory found applications/android + applications/ios are single-screen scaffolds with hardcoded localhost and NO build system. Must create Android Gradle build + manifest and an iOS Xcode project (gomobile or native) before the apps build/run on the Genymotion emulator / iOS simulators for recording.

## **HXC-110 — Extend containers submodule to launch iOS simulators (operator-directed Apple-support mechanism)**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** submodules/containers/pkg/applesim: host xcrun-simctl orchestration (Boot/Install/ Launch/Record/Shutdown, by stable UDID §11.4.111), 16 tests pass incl -race, real host round-trip; cmd/applesim CLI. Submodule a0fa823 pushed all upstreams, meta pointer bumped. **Created-By:** Claude **Assigned-To:** Claude

Operator (2026-06-15): create a proper mechanism for starting mandatory iOS simulators through the containers submodule, extending it to support anything Apple-related. NOTE: iOS simulators run natively via xcrun simctl on macOS (cannot run inside Linux containers) — the containers submodule mechanism must orchestrate the host-native simctl lifecycle (boot/install/record) under its unified API. Investigate + extend containers (§11.4.76).

## **HXC-111 — Desktop GUI shows raw i18n keys (desktop\_dashboard\_header/\_activity\_title) — CONST-046 gap**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Wired i18n.NewTranslator() in NewDesktopApp; dashboard now shows real text (verified via relaunch+AXRaise+screenshot: title ‘HelixCode - Distributed AI Development Platform’, ‘Recent Activity’, no raw keys). Desktop tests pass, build exit 0. Clean desktop video re-recorded. **Created-By:** Claude **Assigned-To:** Claude

After fixing the launch crash, the desktop dashboard renders raw message-ID keys (desktop\_dashboard\_header, desktop\_dashboard\_activity\_title) instead of localized text. Likely the



desktop i18n bundle is missing those keys OR Fyne locale-parse error ('subtag at unknown') broke bundle loading. Real CONST-046 defect visible in helixcode-desktop-dashboard-20260615.mp4.

## **HXC-113 — MCP tool names use 'server:name' (colon) — OpenAI-compatible providers (DeepSeek/etc.) reject function names, breaking LLM chat with MCP enabled**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** mcpToolRegisteredName sanitises MCP tool names to server\_\_name (OpenAI-compatible <sup>2</sup>); dispatch unaffected (Execute uses original server/toolName). 2 mcp\_readonly tests reconciled (§11.4.120); guard test + full internal/tools pkg pass; build exit 0. **Created-By:** Claude **Assigned-To:** Claude

internal/tools/registry.go:897 RegisterMCPManager names MCP tools server+'.'+name (e.g. fs:read\_file). OpenAI/DeepSeek function-calling requires names matching <sup>3</sup>+\$, so a chat turn with MCP tools enabled returns HTTP 400. Found while recording the TUI (had to disable .helixcode/mcp.yml to record). Fix: sanitize MCP tool names (e.g. server\_\_name or server-name) at registration + map back when dispatching.

## **HXC-114 — Wire facade (/v1/chat/completions, /v1/messages) had NO authentication — unauthenticated, paid-provider-consuming surface reachable on 0.0.0.0**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Severity:** Critical (production/release blocker — unauthenticated real, paid LLM-provider generation, reachable on every interface per every shipped config's server.address: "0.0.0.0") **Discovered-By:** AI (independent security review + the dual-wire OpenAI/Anthropic facade review, both flagged it against commit 51c058b1) **Evidence:** §11.4.115 RED->GREEN regression guard internal/server/wire\_facade\_auth\_test.go. RED (captured by

temporarily reverting the middleware wiring and running

```
RED_WIRE_FACADE_AUTH=1 go test -run
```

TestWireFacadeEndpoints\_NoKeyRejected ./internal/server/...): an unauthenticated POST to either route was NOT rejected 401 — it fell straight through to a real provider. Generate call (observed dialing localhost:11434, i.e. it would have reached a real backend and billed a real request had one been reachable). GREEN (post-fix, default mode):

```
go test -count=1 ./internal/server/... — both routes now reject a
```

no-key request 401 (TestWireFacadeEndpoints\_NoKeyRejected), reject EVERY request when no key is configured even with a bearer token supplied — fail-closed by design

(TestWireFacadeEndpoints\_NoKeyConfiguredStillRejected), and pass a request bearing a key that matches the operator-configured list via either Authorization: Bearer <key> OR x-api-key: <key>

(TestWireFacadeEndpoints\_ValidKeyPassesMiddleware).go build ./internal/server/... ./internal/config/... + go vet clean; full internal/server + internal/config package suites pass. **Root cause:** The routes were registered directly on the bare gin router (s.router.POST(...), server.go) with no middleware group at all, and the file-level doc-comment in wire\_facade.go explicitly (and honestly) documented the gap as a tracked follow-up rather than a silent omission — but it had not yet been closed. **Fix:** Added s.wireFacadeAuthMiddleware() (server.go) — a dedicated, wire-compatible API-key check DISTINCT from the internal-user JWT authMiddleware()/VerifyJWTWithDB (genuine OpenAI/Anthropic SDK clients send an API key, never this server's session JWT). Pattern reused, not reinvented (§11.4.74): mirrors submodules/helix\_llm's internal/gateway/middleware.APIKeyAuth (the DZ-05 fix, commit 60b8e4a) — Bearer/x-api-key token compared against a configured, comma-separated key list — with one deliberate strengthening: fails CLOSED (empty cfg.Auth.WireFacadeAPIKeys, the zero-value default in every shipped config, rejects ALL requests) rather than APIKeyAuth's "empty keys => open access" convenience mode, since these routes drive real paid-provider calls. New config field Auth.WireFacadeAPIKeys (internal/config/config.go, env HELIX\_WIRE\_FACADE\_API\_KEYS) —

§11.4.10: no key is hardcoded anywhere; an operator must explicitly configure at least one key before either route accepts traffic.

**Composes-with:** §11.4.115 (RED-baseline + polarity switch), §11.4.135 (this guard is now the standing regression test for this defect class), §11.4.74 (reuse of the helix\_llm/DZ-05 APIKeyAuth pattern), CONST-035/BLUFF-001 (real provider calls, no simulation). **Created-By:** Claude **Assigned-To:** Claude

Commit 51c058b1 (“feat(server): dual OpenAI + Anthropic wire facade on HelixCode’s own server”) added POST /v1/chat/completions and POST /v1/messages to HelixCode’s own server with no auth middleware attached, while every shipped config profile (config/production-config.yaml included) binds server.address to 0.0.0.0. Flagged as a production blocker by both an independent security review and the dual-wire review. Fixed by wiring a dedicated wireFacadeAuthMiddleware (fail-closed API-key check accepting Authorization: Bearer or x-api-key) onto both routes; see wire\_facade\_auth\_test.go for the RED->GREEN regression guard.

## **HXC-112 — Desktop GUI feature-recording: Fyne OpenGL canvas ignores osascript synthetic clicks — need cliclick/real-event automation to record LLM-chat in-GUI**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc\_residuals\_operator\_closure\_20260711/DECISION.md **Created-By:** Claude **Assigned-To:** Claude

Desktop GUI feature-recording: Fyne OpenGL canvas ignores osascript synthetic clicks — need cliclick/real-event automation to record LLM-chat in-GUI

## **HXC-108 — Video-QA program: record all clients x all features with strongest models + ensemble -> /Volumes/T7/Downloads/Recordings, analyze + fix**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc\_residuals\_operator\_closure\_20260711/DECISION.md **Created-By:** Claude **Assigned-To:** Claude

Video-QA program: record all clients x all features with strongest models + ensemble -> /Volumes/T7/Downloads/Recordings, analyze + fix

## **HXC-107 — Feature Status docs program (docs/features) — comprehensive per-feature inventory across all components/clients/submodules/ported-cli\_agents, docs\_chain-synced**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc\_residuals\_operator\_closure\_20260711/DECISION.md **Created-By:** Claude **Assigned-To:** Claude

Feature Status docs program (docs/features) — comprehensive per-feature inventory across all components/clients/submodules/ported-cli\_agents, docs\_chain-synced

## **HXC-115 — CONST-046 anti-bluff gate is non-functional on any non-original checkout**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/discovery\_hardening\_20260711T212548Z/S3\_const046\_gate\_evidence.md **Severity:** High **Created-By:** Claude

The automated check that scans the codebase for forbidden hardcoded user-facing text stores its list of known-good files using full disk paths tied to one specific machine. On any other computer or checkout location those paths no longer match, so the check wrongly flags every file as a brand-new violation and exits with an error. This silently disables a governance safety-net everywhere except its original machine. The fix stores the known-good list using paths relative to the project folder so the check works anywhere. This restores real enforcement for every developer and automated run.

## **HXC-116 — Missing multi-track development runbook referenced by config and tooling**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/discovery\_hardening\_20260711T212548Z/S4\_docs\_evidence.md **Severity:** Medium **Created-By:** Claude

The multi-track development system's configuration file and its command-line tool both direct readers to an operating guide that does not exist in the repository. Anyone trying to bring up or operate the parallel development tracks has no authoritative instructions. This risks mistakes with the shared storage drives. The task is to write the missing guide covering the track layout, how to start tracks safely, and the storage-drive precautions. Operators can then run the multi-track system correctly from documented steps.

## **HXC-120 — Regression test asserted insecure wildcard CORS the hardened server no longer returns**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/discovery\_hardening\_20260711T212548Z/S1\_cors\_evidence.md  
**Severity:** Medium **Created-By:** Claude

A safety test for the web server checked that cross-origin browser requests are answered with a permissive wildcard allowing any website to call the API. The server was correctly hardened to allow only an explicit list of trusted sites, so the old test failed and encoded an insecure expectation. The fix updates the test to verify the secure behavior (only listed sites allowed, others refused) without reintroducing the wildcard. The test suite now protects the secure behavior instead of demanding an insecure one.

## **HXC-121 — Two production model-provider integrations (HuggingFace, Together) had no automated tests**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/discovery\_hardening\_20260711T212548Z/S2\_provider\_tests\_evidence.md **Severity:** Medium **Created-By:** Claude

The code connecting the platform to the HuggingFace and Together AI model services shipped with no automated tests. A future change could silently break how requests are built or responses parsed and nothing would catch it. The work adds real tests exercising the actual client

code against a simulated provider endpoint, checking the outgoing request, the parsed reply, and error handling. These two integrations become protected against silent regressions.

## **HXC-123 — The security-scan command tool has no automated tests**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/discovery\_hardening\_20260711T212548Z/S5\_security\_scan\_evidence.md **Severity:** Low **Created-By:** Claude

The command-line security-scan helper ships without automated tests covering its behavior, so bugs could go unnoticed until a real scan produces wrong results. The work adds tests verifying the tool's core scanning and reporting logic. The security-scan tool then becomes protected against silent breakage.

## **HXC-130 — Building the desktop and mobile GUIs fails without documented X11/OpenGL system libraries**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/discovery\_hardening\_20260711T212548Z/S4\_docs\_evidence.md **Severity:** Medium **Created-By:** Claude

A full build fails on a clean machine because the desktop and mobile graphical apps need system graphics libraries (X11 and OpenGL) that are neither installed nor documented, so newcomers hit a confusing build error with no guidance. The work documents the exact system packages required, the command to install them, and the headless no-graphics build path used for development and continuous integration. Anyone can then build the project or knowingly choose the headless path without surprise failures.

## **HXC-125 — Integration-tagged tests are invisible to the default test run**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/discovery\_hardening\_wave2src\_20260711T220325Z/W2B\_integration\_tag\_evidence.md **Severity:** Low **Created-By:** Claude

A large set of integration tests is hidden behind a build tag, so an ordinary test run never compiles or executes them and their pass/fail signal is absent from routine checks. They pass when the tag is supplied, so this is a visibility gap, not a broken feature. The work makes the standard test command or a documented target include them so their status is always visible. Everyday test results then reflect integration coverage.

## **HXC-127 — The obsolete-items detail table is empty, so retired items lack required justification**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/HXC-044/evidence.md **Severity:** Low **Created-By:** Claude

When an item is retired as obsolete, governance requires a recorded reason, date, and superseding reference, but the table holding these details is completely empty, including for the one currently obsolete item. This leaves retirements unexplained and non-compliant. The work populates the required justification details for obsolete items so every retired item carries an auditable reason.

## **HXC-129 — Severity is blank on 79 closed feature items, blocking risk-based prioritization**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/fdbtool\_hygiene\_20260712T071056Z/W2C\_hygiene\_proposals.jsonl **Severity:** Low **Created-By:** Claude

Seventy-nine finished feature items have no severity recorded, so risk-ordered validation and reporting cannot rank them by importance. The work backfills an appropriate severity for each item based on its content and impact. Risk-based ordering and reporting then work correctly across the full item set.

## **HXC-128 — Thirty-six work items have descriptions too short to be understandable**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/fdbtool\_hygiene\_20260712T071056Z/W2C\_hygiene\_proposals.jsonl  
**Severity:** Low **Created-By:** Claude

Thirty-six tracked items have descriptions shorter than the minimum length needed to convey what they are about, so a reader cannot understand them without reading code. The work rewrites these into clear plain-language descriptions explaining what each item is and why it matters. Every item then becomes understandable to non-developers and stakeholders.

## **HXC-133 — Azure provider factory test assumes an ambient endpoint environment variable**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/followup\_fixes\_20260712T085616Z/HXC133\_evidence.md **Severity:** Low **Created-By:** Claude

One Azure-provider test quietly depends on an endpoint value being present in the environment; when that value is injected by the full-test setup the test's assumptions no longer hold. It does not crash, but it is fragile and can give misleading results depending on the environment. The work is to make the test hermetic (control its own environment) like the sibling test already fixed. This makes the Azure test suite reliable regardless of how it is run.



## HXC-124 — HelixQA task-workflow bank cannot fully self-run due to a JWT token-minting gap

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/followup\_fixes\_20260712T085616Z/HXC124\_evidence.md **Severity:** Medium **Created-By:** Claude

One automated quality-assurance test bank that drives real server workflows has a documented gap: it cannot mint the authentication token needed to finish the workflow end-to-end without manual help, so it does not meet the fully-automated no-human-intervention standard. The work provides an automated way to obtain a valid token during the test so the workflow runs unattended. The QA bank then becomes fully automated and re-runnable.

## HXC-131 — Cognee client authentication and bearer-token caching regressed

**Status:** Obsolete (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/followup\_fixes\_20260712T085616Z/HXC131\_evidence.md **Severity:** Medium **Created-By:** Claude **Obsolete-Details:** Since: 2026-07-12; Reason: duplicate-of; Superseding-item: b058c7c2; Triple-check evidence: docs/qa/followup\_fixes\_20260712T085616Z/HXC131\_evidence.md

The client that talks to the Cognee memory service stopped completing its login and caching the access token — its tests show the login endpoint is never called and no bearer token is stored, so authenticated calls would fail. This means memory features that rely on Cognee cannot authenticate reliably. The work is to restore the login-then-cache-token flow and prove it with the existing auth tests. Users regain dependable access to Cognee-backed memory.

## **HXC-132 — Full-test container stack cannot build the HelixCode server and tests hardcoded port 8080**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc132\_20260712T090048Z/HXC132\_evidence.md **Severity:** Medium **Created-By:** Claude

The full-infrastructure test stack fails to build the HelixCode server container because its build recipe points at a Dockerfile path that does not exist, and many memory and security tests skip themselves because they look for a server on a fixed port 8080 with no way to override it. As a result a whole class of tests never run against a real server. The work is to fix the container build path and make the server URL configurable so those tests execute. This unlocks real end-to-end coverage of server-dependent features.

## **HXC-137 — Re-verify every owned code module builds, checks, and tests cleanly**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/submodule\_health\_20260712/results.csv **Severity:** Medium **Created-By:** Claude

The project depends on many owned code modules, and a full health check of all of them (does each build, pass static checks, and pass its tests) did not finish in the latest session. The work is to run that complete health sweep and record the result for every module. This assures that the whole codebase, not just the main application, is in good shape.

## **HXC-139 — A vendored reference-agent fixture breaks the helix\_agent module build**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** /home/milos/Factory/projects/tools\_and\_research/helix\_code/docs/qa/hxc139\_20260712T102230Z/EVIDENCE.md **Severity:** High **Created-By:** Claude

A vendored copy of a third-party reference coding-agent (the Continue project) includes a Go source file that imports a path that does not exist, and because that file has no separate module marker it gets swept into the helix\_agent module's build — breaking the build and static checks for the whole module. This blocks reliable building and testing of the agent module. The work is to isolate those vendored reference files so they are not compiled as part of our module (a build-ignore or nested module marker). Developers regain a clean, buildable agent module.

## **HXC-141 — mcp\_module Docker adapter crashes when stopping a container that was never started**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** /home/milos/Factory/projects/tools\_and\_research/helix\_code/docs/qa/hxc141\_20260712T111849Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

The MCP module's Docker adapter crashes with a null-pointer error when asked to stop a container that was never started or does not exist, instead of returning cleanly. This can bring down callers that expect a safe no-op. The work is to guard the stop path so a not-started or missing container is handled gracefully. The adapter becomes robust against stop-before-start and missing-container situations.

## **HXC-134 — Model verifier returns the model id as a number but the platform expects text**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** /home/milos/Factory/projects/tools\_and\_research/helix\_code/docs/qa/hxc134\_20260712T124602Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

The central model-verifier service reports each model's id as a numeric value, while HelixCode expects the id as text — a type mismatch that can break how verified models are matched and displayed. The work is to align the two so the id is consistently text end to end. Correct model identity keeps verification, listing, and status accurate for users.

## **HXC-140 — helix\_qa copies a lock by value and one test-bank test fails**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** /home/milos/Factory/projects/tools\_and\_research/helix\_code/docs/qa/hxc140\_20260712T124724Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

The quality-assurance module has code that copies a value containing a lock (a mutex) instead of sharing it, which the Go checker flags as unsafe and can cause subtle concurrency bugs; separately, one test that loads real test banks is failing. The work is to pass the lock-bearing value by reference (pointer) instead of copying it, and to fix or reconcile the failing test-bank test. This makes the QA module concurrency-safe and its tests green.

## **HXC-135 — Model verifier should publish the six advanced-capability flags so the platform can show them**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** docs/qa/hxc135\_20260712T130921Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

HelixCode is now wired to read six advanced capability indicators (tool protocols, code intelligence, retrieval, skills, plugins) from the central verifier, but the verifier's live responses do not yet include those fields, so the flags always read as unsupported. The work is to have the verifier publish these capability values it already computes. Then users see accurate per-model capability information across the product.

## **HXC-117 — Model capability flags (MCP, LSP, ACP, RAG, Skills, Plugins) are not sourced from the verifier as required**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc117\_20260712T140900Z/EVIDENCE.md **Severity:** High **Created-By:** Claude

Governance rule CONST-040 requires that every advanced capability a model supports be reported by the central verifier component rather than hardcoded. Today the verifier only records whether a model supports embeddings; the other six capabilities are documented as verifier-sourced but are not implemented there. As a result the product cannot truthfully tell users which models support which capabilities. The work adds these capability fields to the verifier's results and has the product read them from there. Users then receive accurate, single-source-of-truth capability information.

## HXA-001 (ex-ISSUE-009) — helix\_agent handler tests surfaced after round-109 fix

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 109) **Discovered-By:** AI subagent (helix\_agent LOGIC audit) **Evidence:** Mid-run panic in `TestIsProviderAvailable_NotAvailable` aborted test binary; round 109's fix unblocked execution, surfacing 4 pre-existing FAILs: `TestFormattersHandler_FormatCode_UnsupportedLanguage`, `TestEmbeddingHandler_WithRealManager`, `TestGetTaskResources`, `TestGetTaskLogs`. Out of round-109's 5-fix cap. **Resolution path:** Per-handler investigation, similar to round 109's test-side bluff pattern. **Closure:** Fixed round 116 (commit da782d4) — all 4 handler tests fixed; closure recorded in `docs/Fixed.md`. This `**Status:**` line was stale (Queued) and is corrected here to match `Fixed.md + Issues_Summary.md` (§11.4.12 sync).

---

## HXA-002 (ex-ISSUE-010) — helix\_agent debate/llmprovider sibling-submodule API drift

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 109) **Discovered-By:** AI subagent **Investigated:** 2026-05-20 (round 324) — split into a mechanical part and a design-decision part. **Closed:** 2026-05-20 (round 342) **Closure-Ref:** helix\_agent commit (round-342 HXA-002 debate API drift) + meta-repo `.gitmodules` pointer-bump **Investigation finding (operator's explicit ask — moved vs deleted):** The learning/knowledge/recommendations capability tier was **genuinely DELETED, not moved**. `git log` on the `digital.vasic.debate` submodule (submodules/debate\_orchestrator, renamed from `DebateOrchestrator` per CONST-052) shows the orchestrator was rebuilt

from scratch — commit 196d0ea “feat: initial DebateOrchestrator reconstruction (Phase 1)”. orchestrator/api.go has carried only the slim CreateDebate/GetStatistics surface since that very first commit (git log --follow orchestrator/api.go = single entry). A tree-wide grep of dependencies/ for KnowledgeRepository, GetRecommendations, ConvertAPIRequest, GetDebateStatus, DefaultMinConsensus, MaxAgentsPerDebate, EnableAgentDiversity found **zero** surviving copies in any digital.vasic.\* package or in HelixSpecifier/HelixMemory. The slim API is the first and only version — the richer tier was a pre-reconstruction artifact that no longer exists anywhere. Per the operator’s chosen direction for the deleted case, the helix\_agent tests were rewritten down to the slim API. **Resolution:** See docs/Fixed.md row for the full closure narrative (Part-1 import swap + Part-2 slim-API rewrite + score-scale + go.mod rename-drift fix + captured evidence).

---

## HXA-003 (ex-ISSUE-011) — venice TestGetCapabilities model-list drift (CONST-037)

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 109) **Discovered-By:** AI subagent **Closed:** 2026-05-19 (round 190) **Closure-Ref:** helix\_agent commit (round-190 venice CONST-037 model-list drift) + meta-repo pointer-bump **Evidence:** Test hardcoded venice-uncensored; Venice API returned 75 models with the family rotated to venice-uncensored-1-2 / venice-uncensored-role-play. Per CONST-037 (LLMsVerifier is the single source of truth for model metadata) the assertion violated the no-hardcoded-list rule. **Resolution:** helix\_agent/internal/llm/providers/venice/venice\_test.go::TestGetCapabilities — replaced assert.Contains(..., "venice-uncensored") and assert.Contains(..., "llama-3.3-70b") with structural assertion: NotEmpty(SupportedModels) plus a substring scan for the venice-uncensored\* family. SKIP-OK marker per CONST-035 fires if the entire family disappears (avoids false-positive PASS). Mutation-verified (revert → FAIL with the original drift, restore → PASS).

---

**HXC-001 (ex-ISSUE-005) — CONST-052**  
**rename programme: meta-repo directories**  
**still PascalCase — CLOSED ( → Fixed.md)**

**Status:** Completed ( → Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Task **Closure (2026-05-28):** all owned-org submodule LEAF dirs renamed to lowercase snake\_case (Phases 1–4: 1-A..1-D Upstreams; 2-A..2-D / 3 / 4 leaf dirs) + all 57 Upstreams/ → upstreams/ dirs. Phase 5 (org-grouping dirs dependencies/vasic-digital/ + dependencies/HelixDevelopment/) resolved as a NO-OP per operator decision 2026-05-28 (AskUserQuestion): kept as GitHub-org namespace carve-outs. Section retained as a migration tombstone per §11.4.19; round-343 detail below preserved for history. **Discovered:** 2026-05-15 (CONST-052 cascade landed) **Discovered-By:** Constitution **Evidence:** Meta-repo top-level dirs already snake\_case (round-88 sweep). Remaining non-compliance is two layers deeper: dependencies/HelixDevelopment/\* + dependencies/vasic-digital/\* owned-org submodule dirs (PascalCase), and 59 Upstreams/ dirs inside submodule trees. **Resolution path:** Phased migration per CONST-052 §11.4.29. Round 113 produced the phased plan (f666410, docs/superpowers/specs/2026-05-19-const052-rename-programme-plan.md). Round 343 executed the safe (zero-submodule-go.mod-entanglement) batches.

**Round-343 12 chosen snake\_case names (operator “agent defaults”):** D-1 sequential phases; D-2 helix\_development (parent dir, deferred — touches every consumer go.mod); D-3 vasic-digital kept (GitHub-org handle, proper-noun carve-out); D-4 n/a (helix\_code already snake\_case); D-5 LLMsVerifier/Assets+Website deferred (deployment-wire audit); D-6 mcp\_module; D-7 i\_llm; D-8 toon; D-9 rag; D-10 cluster-C upstreams strict; D-11 yes co-authored; D-12 one approval per batch.

**Round-343 per-batch status:**

| Batch | Renamed                                                             | Status            | Evidence                                                                |
|-------|---------------------------------------------------------------------|-------------------|-------------------------------------------------------------------------|
| 1     | HelixDevelopment/<br>Models → models                                | LANDED<br>a1ea3c8 | submodule resolves;<br>go build ./<br>internal/... ./<br>cmd/... exit 0 |
| 2     | HelixDevelopment/<br>DebateOrchestrator<br>→<br>debate_orchestrator | LANDED<br>416fe8e | go list -m<br>digital.vasic.debate<br>→ new path; build<br>exit 0       |

| Batch    | Renamed                                                                                                                                                                                                                | Status            | Evidence                                                                                                                                                         |
|----------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 3        | 11 vasic-digital/*<br>zero-go.mod-<br>consumer leaves<br>(auto_temp, claritas,<br>doc_processor,<br>gandalf_solutions,<br>hyper_tune, i_llm,<br>leak_hub, ouroborous,<br>plinius_common,<br>veritas,<br>vision_engine) | LANDED<br>e813b5c | 11 submodule<br>statuses resolve;<br>build exit 0                                                                                                                |
| Deferred | ~37 owned-org<br>leaves consumed by<br>helix_agent/<br>helix_qa/HelixLLM<br>go.mod                                                                                                                                     | DEFERRED          | renaming requires<br>submodule-internal<br>go.mod commits<br>entangled with pre-<br>existing<br>uncommitted work<br>— needs dedicated<br>per-submodule<br>rounds |
| Deferred | parent dirs<br>HelixDevelopment/<br>→ helix_development/<br>(D-2), vasic-digital/<br>kept (D-3)                                                                                                                        | DEFERRED          | parent rename<br>touches every<br>consumer go.mod<br>atomically                                                                                                  |
| Deferred | 59 Upstreams/<br>→ upstreams/ (cluster<br>C, D-10)                                                                                                                                                                     | DEFERRED          | live inside<br>submodule trees —<br>separate-repo<br>commits                                                                                                     |

13 of ~58 owned-org leaf renames done this round (zero build breakage). HXC-001 stays In progress pending the deferred submodule-entangled and parent-dir batches.

## HXC-002 (ex-ISSUE-006) — Round-74 residual LOGIC-class FAILs (CLOSED)

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 74 — release-gate-test.sh creation; classified by round 89) **Discovered-By:** AI release-gate sweep **Closure progress:** - ✓ **HelixMemory:** closed



round 106 (commit 69016df — single-line `go.mod` fix; 6 FAIL → 0 FAIL) -  
✓ vasic-digital/Planning: round 107 NO-OP — 275 PASS / 0 FAIL / 20  
SKIP-OK; likely incidentally fixed by round 98 i18n migration - ✓  
helix\_agent inner: closed round 109 (commit 0f492e98 — 5 test-side  
bluff fixes, zero production changes) **Evidence:** Round 74 surfaced 26  
FAILs across submodules; rounds 82-87 closed 19; this Issue tracked the  
residual 7 across 3 submodules. All 3 components closed by rounds 106  
+ 107 + 109. **Follow-ups surfaced (NEW issues filed):** 4 helix\_agent  
handler tests previously masked by mid-run panic (now visible) + 3  
build-failed packages depending on sibling submodule API drift  
(digital.vasic.debate) + 2 LOGIC FAILs reclassified as cross-cutting  
work (venice CONST-037 model-wiring + compliance CONST-051  
architectural reconciliation). See HXA-001 through HXA-003 (filed  
below).

---

## HXC-003 (ex-ISSUE-007) — CONST-046 i18n migration backlog — CLOSED (migrated to docs/Fixed.md)

**Status:** Implemented (→ Fixed.md) — see docs/Fixed.md for the full  
closure record. **Type:** Feature The genuine user-facing (C) string-literal  
surface is exhausted across all 7 CONST-046 scope areas — helix\_code  
internal/ + cmd/ + applications/ (exhausted rounds 461/462),  
LLMsVerifier (round 452), helix\_qa, and every owned vasic-digital/\*  
+ HelixDevelopment/\* submodule (rounds 413/441). Hundreds of rounds  
(~91-462) migrated tens of thousands of literals through i18n seams  
with paired-mutation anti-bluff tests. The remaining ~55k audit-  
baseline hits are all OUT of CONST-046 scope (LLM prompt templates,  
wrapped-error tech strings, identifier tokens, struct-tag keys, format-  
spec tokens, test fixtures) — documented in docs/audits/2026-05-20-  
internal-const046-classification.md. Closed by round 463. Section  
retained as a migration tombstone per §11.4.19 — the authoritative  
closure narrative is in docs/Fixed.md.

---

## HXC-004 — Recovery-batch under-verification (40% FAIL rate per round-193 audit)

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 193 — recovery-batch verification audit) **Discovered-By:** AI subagent **Fixed:** 2026-05-19 (round 200 — per-package test-assertion repair) **Evidence:** Round-193 audit of 10 recovery-batch-landed packages (recovery commits b7f8672 + 5c94696) found 6 PASS / 4 FAIL: - internal/llm (round 161): test-assertion drift — tests still expected pre-i18n English literal “api\_key”, production emits message-ID internal\_llm\_wizard\_anthropic\_apikey\_required - internal/logo (round 163): same drift — tests expected “failed to open” / “failed to decode”, production emits internal\_logo\_open\_source\_failed / internal\_logo\_decode\_source\_failed - internal/notification (round 167): same drift — tests expected Title literals “Task Completed”, “Task Failed”, “Workflow Completed”, “Workflow Failed”, “Worker Disconnected”, “System Error”, “System Started”; production emits internal\_notification\_title\_\* IDs - internal/performance (round 168): build break — translator.go stdctx.Context vs plain context — fixed inline by parent agent **Root cause:** Recovery-batch commits captured stalled-agent file content but did NOT re-run consuming-test updates + did NOT verify build/test green per-package. **Resolution:** Round-200 subagent updated test assertions in all 3 drifted packages to expect message-ID echoes (internal\_<pkg>\_\* prefix). Per-package PASS confirmed (llm: 51.8s, logo: 0.07s, notification: 0.89s, performance: 8.4s). Per CONST-035 mutation-verified one assertion per package: revert to literal → FAIL (production emits the ID), restore → PASS. **Audit reference:** docs/audits/2026-05-19-recovery-batch-verification.md (commit 1bade1).

---

## HXC-005 — cmd/performance\_optimization\_standalone/main.go is a CONST-035 simulation bluff

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-20 (round 317 — cmd i18n migration subagent) **Discovered-By:** AI subagent — refused to localize the file because doing so would polish a bluff **Fixed:** 2026-05-20 (round 318) **Evidence:** helix\_code/cmd/performance\_optimization\_standalone/main.go was a package main that

printed “ Starting HelixCode Production Performance Optimization” then *simulated* every optimization phase: // Simulate production optimization phases, time.Sleep(500 \* time.Millisecond) per phase, and improvement := 5.0 + rand.Float64()\*20.0 — fabricated improvement percentages from a random number generator. No real profiling, no real optimization, no real measurement. The canonical BLUFF-001-class anti-pattern (CLAUDE.md §3.3 / §6 ANTI-PATTERN 1) — a binary that reports success for work it never performed. Violated CONST-035 / Article XI §11.9. **Resolution:** Decision — **DELETE** (resolution path b). The standalone tool was genuinely obsolete: fully superseded by cmd/performance\_optimization/ (snake\_case post-CONST-052), which calls the REAL dev.helix.code/internal/performance.PerformanceOptimizer — real runtime.ReadMemStats, real GOMAXPROCS tuning, real before/after measurement — and carries CONST-046 i18n + a unit-test file. git rm -r cmd/performance\_optimization\_standalone/ removed the dead bluff; stale references purged from docs/COMPREHENSIVE\_AUDIT\_REPORT.md. Reproduce-before-fix Challenge added at cmd/performance\_optimization/bluff\_regression\_test.go: TestHXC005\_BluffStandaloneDirectoryDeleted asserts the obsolete path is gone + the real command survives; TestHXC005\_RealOptimizerMeasuresActualMemory allocates a retained 32 MiB buffer and asserts the optimizer’s baseline MemoryUsage tracks a genuine runtime.HeapAlloc reading (not an RNG single-digit). Post-fix evidence: go build ./cmd/... exit 0; go test -count=1 -run TestHXC005 ./cmd/performance\_optimization/ → both PASS (literal log: optimizer baseline MemoryUsage=33812624 bytes, runtime.HeapAlloc=33802008 bytes – both real measurements, same order of magnitude. No RNG-fabricated improvement percentages.); anti-bluff smoke on the deleted path returns N/A (directory gone).

---

## HXC-006 — HelixCode Speed Programme — CLOSED (migrated to docs/Fixed.md)

**Status:** Implemented (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Feature All 6 phases / 31 tasks landed; CONST-048 coverage ledger at docs/research/speed/05-coverage-ledger.md (29 PASS + 2 PARTIAL + 0 DEFERRED). Closed by P5-T04 round 400. Section retained as a migration tombstone per §11.4.19 — the authoritative closure narrative is in docs/Fixed.md.

---

## **HXC-007 — Constitution §11.4.68/70-74 cascade + meta-pointer bump — CLOSED (migrated to docs/Fixed.md)**

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Task Cascade verified complete (round 403): constitution 584b3ee → 34a82b3, all 6 rules cascaded to the meta-repo + 67 owned submodules, meta pointer confirmed at 34a82b3. Section retained as a migration tombstone per §11.4.19.

---

## **HXC-008 — CONST-055 G1 governance gaps surfaced by post-constitution-pull validation sweep — CLOSED (migrated to docs/ Fixed.md)**

**Status:** Fixed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Bug Both gaps fixed (round 403): (a) submodule\_owned.txt corrected HelixDevelopment/Models → models; (b) CONST-047..057 cascaded into helix\_qa/CONSTITUTION.md, §11.4.69 cascaded into VisionEngine/CONSTITUTION.md. verify-governance-cascade.sh → 0 failures; verify-all-constitution-rules.sh → 6 gates / 0 failures. Section retained as a migration tombstone per §11.4.19.

---

## **HXC-009 — Owned-submodule GitHub ↔ GitLab mirror-divergence reconciliation — CLOSED (migrated to docs/Fixed.md)**

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Task Reconciliation verified complete (round 403): helix\_qa, VisionEngine, LLMProvider, challenges, containers, DocProcessor all reconciled via merge-first (CONST-061 / §11.4.71), all owned submodules converged + pushed to all upstreams. Section retained as a migration tombstone per §11.4.19.

---

## HXC-010 — End-to-end Kimi CLI + Qwen Code CodeGraph verification — CLOSED (migrated to docs/Fixed.md)

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Task Operator supplied OpenAI-compatible router credentials (2026-05-21). Both `cg-challenge-05-kimi.sh` and `cg-challenge-07-qwen.sh` re-run produce **true tier-1 PASS** — each agent genuinely invoked `codegraph_search` and received real graph data from the scanned HelixCode code-graph. Kimi driven via an `openai_legacy` provider, Qwen via `--auth-type openai`, both against SiliconFlow; API keys injected via environment variables only, never written to any tracked file (CONST-042). Section retained as a migration tombstone per §11.4.19.

---

## HXC-013 — Adopt SQLite-backed single-source-of-truth for workable items (§11.4.93/95)

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Closure (2026-05-29, subagent-driven §11.4.70 + conductor integration):** the constitution submodule's non-functional `workable-items` scaffold is now a FUNCTIONAL Go binary (built UPSTREAM per §11.4.74 in `constitution/scripts/workable-items/`, pushed to its origin at `e460a5d`): `sync md-to-db` parses `Issues.md+Fixed.md` (`## <atm-id> - <title> blocks → items rows: atm_id/type/status/severity/title/description + the full raw markdown block in body_md`), `sync db-to-md` regenerates both docs **byte-identically modulo trailing whitespace** (verified: 0 diff lines on the real 137-item tracker), `validate` enforces closed-set status/type + §11.4.91 description floor + no-dup-atm\_id, `diff` reports `md ↔ db` drift. `go build + go test` (round-trip + validate paired-mutation) pass. Per §11.4.95 the tracked DB lives at `docs/workable_items.db` (137 items: 31 Issues + 106 Fixed), WAL-checkpointed before commit, with `.gitignore` gaining `!docs/workable_items.db` (negating the blanket `*.db` while keeping the transient `*.db-wal/*.db-shm` sidecars ignored). HelixCode's constitution pin bumped `6017af9 → e460a5d`. **Operator-decisions resolved with the scope's stated-recommendation safe-defaults (§11.4.101):** `id-strategy = keep ATM/HXC-NNN free-text atm_id`; `round-trip-tolerance = byte-identical-via-body_md-blob`; `upstream-first = built+pushed upstream`

per §11.4.74. **Follow-up (not yet wired):** auto-invoking `sync md-to-db` from a commit hook / `commit_all.sh` + a pre-build `CM-WORKABLE-ITEMS-MD-DB-IN-SYNC` gate (the binary + tracked DB foundation is in place; the auto-sync wiring is the remaining incremental step). **Discovered:** 2026-05-28 (constitution pull 7f738df → 15cd4bc) **Discovered-By:** AI (post-pull awareness review) **Scope:** §11.4.93 + §11.4.95 require a tracked `docs/workable_items.db` as authoritative, with a Go binary doing bidirectional `md ↔ db` sync. Critical findings: the constitution's `constitution/scripts/workable-items/` binary is a NON-FUNCTIONAL Phase-2 scaffold (every subcommand prints “not yet implemented”) — adoption requires building the parser/renderer UPSTREAM per §11.4.74 (triggers §11.4.26); `.gitignore:162` blanket `*.db` must gain `!docs/workable_items.db`; HelixCode lacks `generate_issues_summary.sh/sync_issues_docs.sh`. Effort ~5-9 subagent-days. **Operator decisions:** id strategy (recommend keep HXC-NNN in free-text `atm_id` col); `.gitignore` negation; upstream-first vs wait; round-trip tolerance.

---

## HXC-014 — Stress + chaos test coverage (§11.4.85)

**Status:** Completed (→ Fixed.md) **Type:** Task **Closure (2026-05-29):** the retroactive §11.4.85 sweep is complete across every package with a real resilience surface — 31 in-process packages (batches 1-11) + 5 real-infra packages (redis/database/server/verifier against real podman PG+Redis+server, ollama against real local Ollama). **35 packages covered; 34 real production bugs surfaced + fixed** (systemic classes: panic-in-goroutine-no-recover process crashes, non-reentrant-RWMutex deadlocks, unguarded/declared-unused-mutex data races, plus an auth forged-token DoS and the Ollama CONST-035 empty-generation bluff). The harness (`tests/stresschaos/`), `make stress-chaos` (29 in-process targets), `make stress-chaos-meta` (§1.1 harness self-tests), and `make stress-chaos-infra` (integration-tagged real-infra) are all in place — §11.4.85 is now a per-change discipline going forward (every new fix ships its stress+chaos suite). **Honest out-of-scope remainder (NOT bluffed as covered):** (a) cloud LLM provider endpoints (OpenAI/Anthropic/Gemini/...) need real API keys + incur real cost → operator/external-gated, not stress-loopable safely; (b) low-concurrency stateless utility packages (`version/logo/hardware-probe/adapters/fix`) have NO meaningful resilience surface — a “stress test” of a pure function is a benchmark, not a §11.4.85 resilience test, so forcing tests there would itself be a bluff. The systemic translator `i18n` hardening is tracked + closed as HXC-014b; the pre-existing `llm` integration-build breakage

surfaced during the infra batch is tracked + closed as HXC-024.

**Discovered:** 2026-05-28 (constitution pull) **Discovered-By:** AI **Scope:** §11.4.85 requires every fix/improvement to ship stress (sustained-load  $N \geq 100 / \geq 30s$  p50/p95/p99 + concurrency  $N \geq 10$  + boundary) + chaos (process-death/network-fault/input-corruption/resource-exhaustion/state-corruption) suites with captured evidence. **Batch 1 DONE (2026-05-28, commits 76586014 + a9f883c6):** Go-native helper `helix_code/tests/stresschaos/` (RunSustainedLoad p50/p95/p99 → latency.json; RunConcurrent  $\geq 10$  goroutines + deadlock/leak guards; ChaosKill/CorruptInput/ResourcePressure injectors  $\leq 128MB$  §12.6-safe; evidence → qa-results/ gitignored CONST-053) + 7 paired §1.1 meta-tests proving the harness can't bluff (deadlock/leak/error-rate/below-floor/panic all DETECTED). First 2 targets done: internal/worker + internal/llm/load\_balancer under -race. **The chaos tests SURFACED + FIXED 2 REAL production bugs:** (1) WorkerPool.AssignTask RWMutex-reentrancy deadlock (double-RLock), (2) GetPoolStats data race on PoolWorker.Status. make stress-chaos + make stress-chaos-meta added. Conductor independently re-ran meta-tests → PASS. **Batch 2 DONE (2026-05-28, commit 0505c337):** internal/task (TaskManager/TaskQueue — sustained p50=0.21ms, concurrent 16×120 gDelta=0, state-corruption + input-corruption chaos) + internal/session (TranscriptStore/ResumeFinder, real disk I/O — sustained, concurrent 12×60 no-loss, input-corruption + process-death-mid-append crash-consistent recovery). 8 tests PASS under -race, captured evidence. No new bugs (components robust). All 4 in-process targets (worker/load\_balancer/task/session) now covered. **Batch 3 DONE (2026-05-28, subagent-driven §11.4.70):** internal/event + internal/memory + internal/context (3 parallel disjoint-scope subagents). **2 MORE REAL production bugs SURFACED + FIXED:** (1) internal/event/bus.go — a panicking event handler crashed the whole process (in async dispatch the handler runs in its own goroutine; an unrecovered panic took down every goroutine). Fix: EventBus.invokeHandler recover-wrapper routed through all 3 dispatch sites (Publish async/sync + PublishAndWait) → graceful degradation. (2) internal/memory/manager.go — GetConversation/GetActive/GetAll/GetBySession/GetRecent/Search returned the LIVE stored \*Conversation pointer out from under the RLock; callers read conv.Messages/MessageCount concurrently with AddMessage/Clear writers → genuine DATA RACE the map-only RWMutex didn't cover. Fix: return conv.Clone() snapshots + repaired Conversation.Clone() which silently dropped CharacterID/UserID/CharMessages/Version/Status. internal/context: no bug (RWMutex discipline held), full coverage added. All consumers of the memory read methods confirmed read-only (context builder / persistence snapshot / listing) → no regression. Each fix proven anti-

bluff by reverting → test FAILs → restore → PASS. make stress-chaos target extended to include task/session/event/memory/context; full target PASSES under -race. 7 in-process targets now covered. **Batch 4 DONE (2026-05-28, subagent-driven §11.4.70):** internal/rules + internal/repomap + internal/discovery (3 parallel disjoint-scope subagents). No new production bugs — all three components already correctly RWMutex-guarded and survived concurrent churn + Clear/cancel races + malformed/corrupt input + bounded memory pressure with no panic/race/deadlock/leak. Each suite proven genuinely fail-capable (anti-bluff): rules → strip AddProjectRule lock → DATA RACE FAIL; repomap → write results[0] in parseFilesParallel → DATA RACE FAIL; discovery → strip port-range check → boundary FAIL; all restored byte-identical. discovery deliberately targets in-process registry/config/JSON-decode paths (not flaky UDP-multicast/net.Listen). 10 in-process targets now covered; make stress-chaos extended to all 10, full target PASSES under -race. **Batch 5 DONE (2026-05-28, subagent-driven §11.4.70):** internal/tools + internal/workflow + internal/hooks (3 parallel disjoint-scope subagents). **3 MORE REAL production bugs SURFACED + FIXED, all in internal/hooks:** (1) `executor.go` `executeSync/executeAsync` called `hook.Execute` with no `recover()` — a panicking async hook handler crashes the whole process (same bug class as the event-bus fix); fixed via `Executor.runHandler` `recover-wrapper` → graceful degradation. (2) `manager.go` `TriggerEvent/TriggerEventAndWait/TriggerEventSync` held `m.mu.RLock()` then called `GetByType()` which re-locks the non-reentrant RWMutex — a writer (Register/Unregister) queuing between the two RLock acquisitions deadlocks both (same bug class as the worker batch-1 deadlock); fixed by dropping the redundant outer RLock (`GetByType` already returns a defensive copy under its own lock — verified). (3) `Register(nil)` nil-pointer panic (Validate dereferences nil receiver); fixed with a nil guard + CONST-046-compliant i18n key `internal_hooks_nil`. internal/tools (ToolRegistry + real fs/shell tools — command-injection + CONST-033 power-mgmt commands confirmed REFUSED by the security gate before `os/exec`) and internal/workflow (state-machine + executor + BackgroundManager) had no bugs, full coverage added. Each fix anti-bluff-proven: revert → FAIL (panic / 30s deadlock timeout / FATAL) → restore → PASS. 13 in-process targets now covered; make stress-chaos extended to all 13, full target PASSES under -race. **Batch 6 DONE (2026-05-28, subagent-driven §11.4.70):** internal/agent + internal/monitoring + internal/commands (3 parallel disjoint-scope subagents). **3 MORE REAL production bugs SURFACED + FIXED:** (1) internal/agent/agent.go — `AgentRegistry.agents` map had NO mutex; Register/Unregister write it while Get/List/Count/GetByType/GetByCapability read it, and the Coordinator delegates without holding



c.mu → unsynchronised concurrent map access = guaranteed fatal error: concurrent map read and map write under load. Fix: added sync.RWMutex (writers Lock, readers RLock, never reentrant). (2) internal/monitoring/monitor.go:48 — CollectMetrics called collector.Collect() under the write lock with no recover(); a panicking collector (HTTP scrape / /proc read fault) crashes the process. Fix: safeCollect recover-wrapper + i18n key

internal\_monitoring\_collector\_panic. (3) internal/commands/executor.go — Executor.Execute called cmd.Execute() with no recover(); a panicking slash/markdown/third-party command crashes the host CLI/server. Fix: executeRecovered wrapper + i18n key

internal\_commands\_command\_panicked. All CONST-046-compliant (no hardcoded English). Each fix anti-bluff-proven: revert → DATA RACE / concurrent-map-crash / panic → restore → PASS. **16 in-process targets now covered; 8 real production bugs fixed this session** (worker/load\_balancer batch 1; event/memory batch 3; hooks×3 batch 5; agent/monitoring/commands batch 6); make stress-chaos extended to all 16, full target PASSES under -race. **Batch 7 DONE (2026-05-28, subagent-driven §11.4.70):** internal/mcp + internal/notification + internal/performance (3 parallel disjoint-scope subagents). **4 MORE REAL production bugs SURFACED + FIXED:** (1+2) internal/mcp/server.go handleCallTool invoked tool.Handler with no recover() (handleSession dispatches each message via go handleMessage, so a panicking tool handler crashes the process) + internal/mcp/lifecycle.go Client.setState called the caller-supplied onEvent callback inline with no recover (panic crashes the Connect/Close/recvLoop goroutine); fixed via invokeToolHandler recover-wrapper + i18n key

internal\_mcp\_server\_tool\_handler\_panicked + callback recover. (3) internal/notification/engine.go sendToChannels called channel.Send with no recover — a panicking channel crashes the process when dispatched from a NotificationQueue worker goroutine; fixed via sendOne recover-wrapper. (4) internal/performance/optimizer.go declared po.mutex but NEVER used it — StartProductionOptimization wrote po.optimizations while getOptimizationsByType iterated it unsynchronised → fatal error: concurrent map iteration and map write; fixed with Lock on write + RLock on read. Each anti-bluff-proven: revert → panic/FATAL/DATA RACE → restore → PASS. **19 in-process targets now covered; 12 real production bugs fixed this session** (+mcp×2, notification, performance in batch 7); make stress-chaos extended to all 19, full target PASSES under -race. **Batch 8 DONE (2026-05-28, subagent-driven §11.4.70):** internal/security + internal/providers + internal/llm-manager (3 parallel disjoint-scope subagents). **6 MORE REAL production bugs SURFACED + FIXED:** (1+2) internal/security/translator.go — package-level translator var read by tr() +

written by SetTranslator() with NO mutex (data race; scans run in background goroutines) + tr() called translator.T() with no recover (panicking injected translator crashes the scan-worker goroutine); fixed with RWMutex + recover. (3+4) internal/providers/ai\_integration.go — Initialize held ai.mu.Lock() then called NewConversationManager → GetProvider → ai.mu.RLock() = non-reentrant self-deadlock hanging the process (fixed: release write lock before building managers, re-acquire to store) + vector\_integration.go NewVectorIntegration(nil) left config nil → Initialize SIGSEGV (fixed: nil-config default mirroring NewMemoryIntegration; codifying-the-crash test assertion corrected). (5) internal/llm/model\_manager.go:341 — calculateHardwareCompatibility called hardwareDetector.Detect() (mutates detector state) under only RLock → data race across concurrent SelectOptimalModel callers; fixed with sync.Once (host hardware static). (6) **load\_balancer follow-up (conductor-fixed, surfaced by the llm subagent):** internal/llm/load\_balancer.go Stop() never cancelled the collectStats goroutine (leak when Start got context.Background()) + collectPerformanceStats deref'd lb.manager with no nil-guard → nil-pointer panic after the 10s ticker; fixed with a cancelStats CancelFunc cancelled in Stop() + nil-guard. Also resolved the file's pre-existing CWE-338 (math/rand → crypto/rand uniform provider selection, dropping the deprecated rand.Seed) flagged by the semgrep gate. Full internal/llm package now PASSES under -race with NO skip (previously the nil-manager test had to be skipped). Each fix anti-bluff-proven: revert → DATA RACE / 15s-deadlock-timeout / SIGSEGV / FATAL → restore → PASS. **22 in-process targets now covered; 18 real production bugs fixed this session. Batch 9 DONE (2026-05-28, subagent-driven §11.4.70):** internal/editor + internal/template + internal/cognee (3 parallel disjoint-scope subagents). **2 MORE REAL production bugs SURFACED + FIXED in internal/template/manager.go:** Register/Update/Delete invoked user OnCreate/OnUpdate/OnDelete callbacks (a) with no recover() → a panicking callback crashes the process (esp. when Register runs on a goroutine), and (b) the On\* registrars appended to the callback slices with NO lock while the trigger paths iterated them under lock → data race. Fix: snapshotCallbacks under lock + fireCallbacks panic-isolated AFTER unlock (also prevents re-entrant-callback deadlock); On\* registrars now lock-guarded. internal/editor (CodeEditor/WholeEditor/SearchReplaceEditor/LineEditor/DiffEditor over real t.TempDir files) and internal/cognee (ServiceCache/ServiceStatistics/event-handler fan-out in-process; network Client paths honestly out of scope) had no bugs, full coverage added. Each anti-bluff-proven: revert → panic / DATA RACE → restore → PASS. **25 in-process targets now covered; 20 real production bugs fixed this session. Batch 10 DONE (2026-05-28,**

**subagent-driven §11.4.70):** internal/focus + internal/config + internal/persistence (3 parallel disjoint-scope subagents). **5 MORE REAL production bugs SURFACED + FIXED:** (1+2) internal/focus/manager.go — CreateChain/SetActiveChain/DeleteChain invoked onCreate/onActivate/onDelete callbacks in a loop with no recover() WHILE holding m.mu.Lock() → panicking callback crashes the process + a callback re-entering the Manager deadlocks the non-reentrant RWMutex; fixed via invokeCallbacks (recover) + snapshot-then-release-lock-before-dispatch. (3) internal/config/config.go — ConfigManager had NO mutex at all; GetConfig/loadConfig/ImportConfig/UpdateConfig/ResetToDefaults/saveConfig read+wrote m.config unguarded while the ConfigAPI HTTP handlers drive GET /config concurrently with POST /config/reload → data race; fixed with RWMutex + copy-on-write UpdateConfig + load/save split into public(locking)+Locked(caller-holds) helpers (non-reentrant-safe) + decode-into-fresh-struct-swap-on-success. (4+5) internal/persistence/store.go — SaveAll/LoadAll/triggerError invoked user callbacks with no recover (process crash) + On\* registrars appended to callback slices with no lock while Save/Load iterated under lock (data race); fixed with invoke\*Callback recover-wrappers + lock-guarded registration + snapshot-dispatch-outside-lock + 3 i18n keys. (Note: a transient build-break appeared mid-batch when config's in-flight edit referenced its not-yet-added Locked helpers; the persistence subagent correctly verified in an isolated git worktree per §11.4.84/§11.4.96; combined final state builds + passes -race, conductor re-verified.) internal/focus, config, persistence all CONST-046-compliant. Each anti-bluff-proven: revert → panic / deadlock / DATA RACE → restore → PASS. **28 in-process targets now covered; 25 real production bugs fixed this session. Batch 11 DONE (2026-05-28, subagent-driven §11.4.70):** internal/auth + internal/deployment + internal/project. **4 MORE REAL production bugs SURFACED + FIXED:** (1) **SECURITY** internal/auth/auth.go VerifyJWT — unchecked type assertions claims["username"].(string)/claims["email"].(string); a validly-signed token carrying a missing/numeric/null username/email claim PANICS the process (crash-on-untrusted-input DoS — a single forged request). Fixed with comma-ok assertions → ErrTokenInvalid. (2+3) internal/deployment/production\_deployer.go addNotification appended to shared status with a declared-but-UNUSED mutex (data race) + deployment/translator.go package-var race; fixed with the mutex + translatorMu. (4) internal/project/manager.go GetActiveProject wrote m.activeProject under only RLock (lazy-scan write under read-lock) → data race; fixed with exclusive Lock + re-check. Each anti-bluff-proven: revert → panic / DATA RACE → restore → PASS. **31 in-process targets now covered; 29 real production bugs fixed this session.**

## HXC-014b — Systemic unguarded i18n translator.go data-race + panic-crash (cross-package)

**Status:** Fixed (→ Fixed.md) **Type:** Bug Closure (2026-05-28, subagent-driven §11.4.70): all 52 remaining unguarded internal/\*/  
translator.go files fixed (3 parallel disjoint-group subagents, 18+17+17) to match the proven internal/security + internal/deployment guarding pattern — added translatorMu sync.RWMutex (Lock in SetTranslator, RLock-snapshot in tr) + a recover() in tr (named return) degrading a panicking translator to the message ID, and removed the inline if translator==nil { translator = Noop } write-on-read-path. 54/54 translator.go files now guarded; whole internal/...+cmd/... tree builds clean; gofmt clean. Anti-bluff proof: new internal/logging/translator\_race\_test.go hammers SetTranslator concurrently with tr (16×300) + a panicking translator; §1.1 paired mutation (strip the guard) → WARNING: DATA RACE + panic + FAIL → restore → PASS.

**Regression follow-up (2026-05-28, same day):** the sweep's claim of "behaviour-preserving" was incomplete — removing the inline if translator==nil { translator = Noop } write-on-read-path (the race) broke 10 pre-existing TestTr\_RecoversFromNilPackageTranslator unit tests (adapters/kilocode/verifier/plantree/repomap/logging/projectmemory/quality/render/voice) that asserted tr() SELF-HEALS the package-level var back to non-nil — i.e. they asserted exactly the racy write the sweep correctly removed. The HXC-014b verification gap: it ran go build + make stress-chaos (-run 'Stress|Chaos' filtered) but NOT the swept packages' full unit suites, so these unit failures were masked. Fixed by removing the stale self-heal assertion from all 10 (the msgID-echo assertion already proves graceful degradation via the local Noop snapshot — no package-var mutation, no race). All 10 packages' full unit suites + gofmt now green (verified per-package); full go test ./internal/... blast-radius sweep re-run → exit 0, 0 FAILs (no other HXC-014b regression). Net behaviour: tr() with a nil package translator returns the msgID via a race-free local Noop, never mutating shared state. **Discovered:** 2026-05-28 (HXC-014 batch 8+11 — same defect found independently in internal/security AND internal/deployment translator.go) **Severity:** Medium (latent: requires SetTranslator() concurrent with tr()); existing -race tests pass because SetTranslator is boot-only in practice) **Scope:** ~50 packages under helix\_code/internal/\*/translator.go share a copy-pasted CONST-046 resolver seam with (a) an UNGUARDED package-level var translator <pkg>i18n.Translator — SetTranslator writes it + tr() reads it (and self-heals translator = Noop{} on the read path, a write) with NO

mutex → data race if SetTranslator is ever called concurrently with tr(); (b) NO recover() around translator.T() → a panicking injected/buggy Translator crashes the emitting goroutine (process-wide). The FIX PATTERN is already proven in internal/security/translator.go + internal/deployment/translator.go: add translatorMu sync.RWMutex (Lock in SetTranslator, RLock-snapshot in tr) + a recover() in tr degrading to the message ID (named return). Mechanical, identical per file; ~50 files remain (event/memory/context/rules/repomap/discovery/tools/workflow/hooks/agent/monitoring/commands/mcp/notification/performance/providers/llm/editor/template/cognee/focus/config/persistence/auth/project + ~25 more incl. approval/clarification/planner/plantree/plugins/quality/redis/render/roocode/secrets/session/verifier/voice/worker/workspace/etc.). NOTE: several of those packages' OTHER concurrency bugs were already fixed in batches 3-11, but most still carry the unguarded translator seam. **Best fixed as a single carefully-verified mechanical sweep (build + -race after) rather than rushed — tracked separately so it is not lost. INFRA BATCH (partial) DONE (2026-05-28, subagent-driven \$11.4.70, operator-authorized heavy-infra):** brought up real Postgres + Redis via podman (lightweight — only the services the tests need, not the full 17-service docker stack). **internal/redis** (real Redis localhost:6379) — 13 integration-tagged stress+chaos tests PASS under -race (sustained SET/GET/DEL p50=0.138ms; concurrent INCR atomicity; pub/sub; pipeline; cancel/corrupt/pressure/connection-churn chaos; translator-panic isolation). **internal/database** (real PG) — 10 integration-tagged stress+chaos tests PASS under -race (sustained p50=3.09ms; 16-goroutine no-lost-writes; tx contention; cancel-mid-query; SQL-injection-safe param binding verified; pool-exhaustion clean recovery; connection-churn). No new stress/chaos bugs (both layers robust; translator already fixed by HXC-014b). **Also fixed 2 pre-existing database\_integration\_test.go failures surfaced by running the long-dormant integration suite:** (1) TestNew\_InvalidHost asserted the pre-CONST-046 English “failed to ping database” — updated to the i18n message-ID internal\_database\_ping\_failed; (2) TestPoolSizing asserted exactly-0 CLI idle conns but New()'s mandatory ping legitimately leaves 1 warm idle (CLI MinConns=0 confirmed — no pre-warm) — corrected to the true invariant (≤1 post-ping AND strictly < server's pre-warmed pool). New make stress-chaos-infra target (integration-tagged, real PG+Redis). Each anti-bluff-proven. **INFRA BATCH (part 2) DONE (2026-05-28): internal/server** — 8 integration-tagged DDoS-class tests PASS under -race against the REAL Gin server booted via server.New(cfg, real-PG, real-Redis) + httptest.NewServer(srv.router) (full middleware chain): sustained health p50=0.54ms, 16-goroutine flood, 64KB → 8MiB bodies (all 400, no

OOM), malformed-request + handler-panic-isolation + slowloris chaos, server stays up (post-chaos /health 200). No new bug; anti-bluff proven (remove gin.Recovery → panic escapes → FAIL → restore → pass).

**internal/verifier** — 14 stress+chaos tests PASS under -race (HealthMonitor circuit-breaker 16×4800 concurrent, Cache tiered 20×4000, EventPublisher pubsub, Poller lifecycle). **1 REAL bug fixed:** internal/verifier/adapter.go EventPublisher.Publish launched each subscriber via go fn(event) with NO recover() → a panicking subscriber crashes the process (the Poller publishes to these); fixed with a per-subscriber recover guard. Anti-bluff proven. **Remaining (long-tail):** internal/llm provider endpoints (need live LLM/Ollama), low-concurrency utility packages (version/hardware/adapters/fix/logo — minimal concurrent state). internal/persistence confirmed NO live-DB path (file-based, covered batch 10). **INFRA BATCH (part 3 — Ollama) DONE (2026-05-28, real local Ollama qwen2.5:0.5b via podman):** internal/llm Ollama provider — 9 integration-tagged stress+chaos tests PASS under -race against REAL Ollama (fast-path GetHealth/GetModels sustained+concurrent; bounded real-generation concurrency proving non-empty completions e.g. “Nonce 64000.”; cancel-mid-generate, hostile-prompt, closed-port chaos). **3 REAL production bugs fixed in internal/llm/ollama\_provider.go:** (1) **CONST-035 / Article XI §11.9 GENERATION BLUFF** — the provider POSTs /api/chat (completion in message.content) but OllamaAPIResponse only parsed the top-level response field → real Ollama generation returned EMPTY text to the end user; unit tests masked it by mocking response. Fixed: decode message.content via completionText() (preferred, response fallback). (2) data race on isRunning plain bool (read by Generate/IsAvailable/GetHealth, written by Close) → atomic.Bool. (3) connection/goroutine leak (unbounded transport + bodies closed-without-drain; GetHealth never closed on success) → bounded Transport + drain-then-close (800 → 2 goroutines). Each anti-bluff-proven (revert → empty/FAIL/race → restore → PASS). **NOTE:** the internal/llm package’s -tags=integration build is PRE-EXISTINGLY BROKEN (stale integration\_test.go/local\_providers\_integration\_test.go/cloud\_providers\_integration\_test.go/etc. reference ~9 deleted provider constructors + a duplicate MockProvider) — filed as **HXC-024**; the ollama integration tests pass in isolation and land once HXC-024 restores the package integration build. **Session total: 34 real production bugs fixed.**

## HXC-015 — Cross-platform parity (§11.4.81)

**Status:** Completed (→ Fixed.md) **Type:** Task **Closure (2026-05-28, subagent-driven §11.4.70):** supported-platforms manifest docs/platforms/supported\_platforms.yaml (+README) declaring linux/macOS/windows host-shell targets + ios/android/aurora\_os/harmony\_os cross-compile (macOS/windows ci\_test\_hardware: false — honest, no hardware enrolled). scripts/gates/cross\_platform\_parity\_gate.sh (CM-CROSS-PLATFORM-PARITY) scans for case "\$(uname -s)" dispatch + asserts no multi-platform script silently drops a manifest platform without a # PARITY-GAP: honest-kernel-gap citation; wired into the CONST-055 sweep as **G10** (PASS). Real finding caught+fixed: scripts/install.sh did linux+darwin dispatch but silently dropped Windows\_NT → added honest gap citation. Paired-mutation meta-test scripts/tests/cross\_platform\_parity\_meta\_test.sh (3 assertions: missing-Darwin bluff FAILs → add-branch PASS → honest-gap-citation PASS), exit 0. **Doc-fix done:** root CLAUDE.md §3.2.1 said the inner module is at helix\_code/helix\_code/ (nonexistent — test -d helix\_code/helix\_code ABSENT) + mislabelled “submodule” — corrected to helix\_code/ “tracked subdirectory” with go.mod evidence inline (inner module dev.helix.code go 1.26 at helix\_code/go.mod; thin root go.mod go 1.25.2). **sandbox.go rlimit re-assessed — NOT the “never-applied stub” the original scope claimed:** internal/tools/sandbox/native\_backend.go (linux) applies real syscall.Setrlimit(RLIMIT\_AS) (cgroup-v2 preferred, rlimit fallback); native\_backend\_other.go (!linux) is an HONEST fail-closed stub (returns “unavailable on non-Linux, deferred to F14.5”, not a silent no-op). A real per-OS impl (macOS Seatbelt+RLIMIT\_CPU/SIGXCPU proxy per the §11.4.81(C) XNU-RLIMIT\_AS kernel gap; Windows Job Object) is spec-deferred to F14.5 — recommend a dedicated F14.5 ticket. bash -n clean. **Operator-gated remainder (honest OPERATOR-BLOCKED):** actually running per-OS test branches on real macOS/Windows/mobile hardware needs that hardware enrolled (none currently); the gate + manifest + honest-gap discipline are the enforceable core and are complete. **Discovered:** 2026-05-28 (constitution pull) **Discovered-By:** AI **Scope:** §11.4.81 requires per-OS equivalents (runtime dispatch) for platform-specific primitives + honest kernel-gap citations. Gaps: shell scripts rarely uname-dispatch, no CM-CROSS-PLATFORM-PARITY gate, no supported-platforms manifest, shell/sandbox.go rlimits a never-applied stub. Sub-defect **HXC-015a (7 platform Assert(true, "...skipped") fake-skips) already FIXED in commit f464adb0** (honest v.Skip on runtime.GOOS). This item tracks the remaining parity gate + manifest + rlimit work. Open decision: which

non-Linux platforms have test hardware (else honest OPERATOR-BLOCKED). **Doc-fix:** root CLAUDE.md §3.2.1 prose says inner module is helix\_code/helix\_code/; actual is helix\_code/ — correct the prose.

---

## HXC-016 — §11.4.69–97 governance cascade into owned submodules (CONST-047/§3) — CLOSED (→ Fixed.md)

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Task **Discovered:** 2026-05-28 (constitution pull) **Discovered-By:** AI **Closure (2026-05-28):** all 24 anchors (§11.4.69 + §11.4.75–97) cascaded into the 5 root govfiles (27929ae1) AND all ~68 owned-submodule CONSTITUTION/CLAUDE/AGENTS/QWEN files (batches 1-7: ef4b3986/a864039d/e4046668/3adb2e63/464b2401/b4ad4f50/053fd731). A loose-grep false-match (the §11.4.93 body cites §11.4.95) had skipped the §11.4.95 heading in batch-1-6 submodules — repaired by fix-up A/B/C (79478ed5/903b9225/a9a1a6a1) + the gate tightened to the ## §11.4.NN – heading marker (d2165bf7). 4 HelixDevelopment submodules regressed to the wrong branch (main vs canonical master) by the cascade reset — repaired (b4b790ea). Gate submodule-scope enabled (9031368d). Final verify-governance-cascade.sh → **0 failures**, 204 submodule covenant-114 PASS lines; paired §1.1 mutation proven (strip §11.4.95 → 1 failure → restore → 0). Section retained as a migration tombstone per §11.4.19.

---

## HXC-017 — CodeGraph own-org submodule indexing + update automation (§11.4.79/80) — CLOSED (→ Fixed.md)

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Task **Discovered:** 2026-05-28 (constitution pull) **Discovered-By:** AI **Closure (2026-05-28, commits 176fe07b + 551552f7 + 876b3b36):** .codegraph/config.json blanket dependencies/\*\* exclude replaced with 3 specific third-party excludes (LLama\_CPP/Ollama/HuggingFace\_Hub) so own-org dependencies/vasic-digital/\*\* + dependencies/HelixDevelopment/\*\* are now INCLUDED; credential excludes (/env, key, pem, secrets) added. Root .gitignore fixed so .codegraph/config.json is TRACKED (§11.4.78 — it had been blanket-ignored). Re-index (codegraph index . exit 0): Files 39,024 → 76,044,



**Nodes 624,103 → 1,255,974, Edges 1.64M → 3.96M.** §11.4.79 anti-bluff probe (independently re-verified by conductor):\*\* codegraph query EventBus → submodules/event\_bus/pkg/bus/bus.go:85; helix\_memory → submodules/helix\_memory/...; third-party llama filtered to LLama\_CPP → empty. docs/codegraph/Status.md + Status\_Summary.md created (§11.4.80; weekly automation inherited by reference from constitution codegraph\_update.sh/codegraph\_sync.sh). Section retained as a migration tombstone per §11.4.19.

---

## HXC-018 — Obsolete status (§11.4.90) + summary-doc clarity (§11.4.91) tracker tooling

**Status:** Completed (→ Fixed.md) **Type:** Task **Discovered:** 2026-05-28 (constitution pull) **Discovered-By:** AI **Scope:** §11.4.90 adds terminal Obsolete (→ Fixed.md) status + Obsolete-Details line + colorizer cell-status-obsolete; §11.4.91 forbids anti-pattern summary one-liners (“Composes with” etc.) and requires generators to refuse them. Extend HelixCode’s tracker + summary generators + colorizer. **Closure (2026-05-28, subagent-driven §11.4.70):** §11.4.90 — docs/\_progress-style.css adds the tr.cell-status-obsolete rule (light-gray #E0E0E0 + strikethrough); scripts/gates/obsolete\_colorize.sh tags Obsolete-status <tr>s post-render; scripts/regenerate-tracker-exports.sh wires --css docs/\_progress-style.css --embed-resources + the colorizer into the HTML render; scripts/gates/obsolete\_details\_gate.sh asserts every Obsolete (→ Fixed.md) heading carries a valid \*\*Obsolete-Details:\*\* line (Since/Reason-from-closed-vocab/Superseding-item/Triple-check). §11.4.91 — scripts/gates/summary\_clarity\_gate.sh FAILs on anti-pattern one-liners + descriptions <6 words AND <40 chars; found + fixed 1 real violation (HXA-001 Issues\_Summary row). Both wired into the CONST-055 sweep as G8/G9. Anti-bluff: scripts/tests/{obsolete\_details,summary\_clarity}\_meta\_test.sh paired-mutation (plant violation → gate FAIL → remove → PASS), both exit 0; gates run clean against real docs (0 violations). bash -n clean (CONST-068). The §11.4.90 *terminal-status DB column* couples with HXC-013 (SQLite) and remains future; the MD-tracker gates + colorizer are complete now.

---

## HXC-019 — docs/qa/ end-user evidence tree (§11.4.83)

**Status:** Completed (→ Fixed.md) **Type:** Task **Discovered:** 2026-05-28 (constitution pull) **Discovered-By:** AI **Scope:** §11.4.83 requires every shipped feature to carry a recorded e2e transcript + materials under docs/qa/<run-id>/; release gate refuses feature commits lacking it. Establish the tree + gate. **Closure (2026-05-28, subagent-driven §11.4.70 — operator authorised hard-gate promotion):** the tree + advisory scanner already existed; promoted scripts/verify\_qa\_evidence.sh to an ENFORCING release gate with --enforce + mandatory --since <baseline> (baseline = the docs/qa/README.md-introducing commit ed84f90e, so pre-convention history is exempt) + a [no-qa-evidence] per-commit opt-out for non-feature changes. Wired into scripts/release-gate-test.sh via scripts/gates/qa\_evidence\_gate.sh AND into the CONST-055 sweep as G7. Also fixed a latent git-2.50 bug in the original scanner (git show -s --name-only matched zero commits → silent false-clean; replaced with git diff-tree). The enforcing gate correctly flagged this session's own 10 HXC-014 feature commits as lacking evidence → resolved honestly by adding docs/qa/HXC-014/transcript.md (real captured stress/chaos evidence + anti-bluff mutation proofs); gate now PASSES (10/10 matched). Anti-bluff: scripts/tests/verify\_qa\_evidence\_meta\_test.sh paired-mutation (6 assertions: no-evidence → 1, add-evidence → 0, opt-out → 0, untagged → 1, no-since → 2, pre-baseline-exempt → 0), exit 0. bash -n clean (CONST-068). NOT wired into pre-commit/pre-push (release-gate-only per the mandate wording).

---

## HXC-022 — test\_bank platform + integration packages do not compile (pre-existing) — CLOSED (→ Fixed.md)

**Status:** Fixed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Bug **Discovered:** 2026-05-28 (anti-bluff sweep during HXC-021 fix) **Discovered-By:** AI (captured: go build ./platform/... ./integration/... in helix\_code/tests/e2e/test\_bank) **Closure (2026-05-28, commit 02b3081c):** all ~11 named declared and not used half-written stubs COMPLETED with real assertions (created-resource IDs → assert non-empty; metric values → assert non-nil; 2 vestigial unsent request bodies removed); root-dir package collision resolved by git mv

performance\_security\_tests.go → performance/ subpackage; unmasked pre-existing core/ defects (duplicate GetCoreTests, unused imports) fixed too. go build ./... exit 0 (whole module), go vet ./... clean — independently re-verified. HXC-021 runtime-verified through the now-compiling banks: platform SKIP=3 (honest “not running on macOS/Windows/ARM”), integration SKIP=2 (honest “Ollama not reachable”/“OPENAI\_API\_KEY not configured”), no fake PASS/green-empty. Section retained as a migration tombstone per §11.4.19.

---

## **HXC-023 — Assert(true,...) / AssertTrue(true,...) literal-true bluffs across test\_bank — CLOSED (→ Fixed.md)**

**Status:** Fixed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Type:** Bug **Discovered:** 2026-05-28 (surfaced while fixing HXC-022) **Discovered-By:** AI **Closure (2026-05-28, commits 8e80e0c0 + b514f8bb):** ALL literal-true PASS-bluffs across the e2e test banks replaced with real assertions or honest skips — batch 1 (core/ additional\_tests.go, 41 fixed) + batch 2 (distributed 12, integration 11, platform 5, core/tests.go 4, performance 1 = 33). Pattern: mislabelled “X succeeded” branches that fired on non-2xx → assert the expected 2xx; 401/403/429 branches → assert that exact status; feature-404 branches → honest v.Skip(reason) (§11.4.3); the 4 legitimate “Running on ” positive-platform asserts left untouched. Verification: go build ./... + go vet ./... exit 0; full-tree grep for literal-true bluffs = 0; runtime harness (down server) → all changed cases HONEST-FAIL, 0 green-empty. Section retained as a migration tombstone per §11.4.19.

---

## **HXC-024 — internal/llm -tags=integration build broken (stale tests reference deleted providers)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Closure (2026-05-29, subagent-driven §11.4.70):** go test -tags=integration ./internal/llm/ now compiles + runs (ok 86s, real PG/Redis/Ollama). Fixes: (1) deduped MockProvider — renamed the integration-only field-based one → integrationMockProvider (the func-based tool\_provider\_test.go one is canonical) + added the now-required GetContextWindow()/CountTokens() interface methods. (2) LlamaConfig.ModelPath → Model rename fixed

(struct + NewLlamaCplusplusProvider still exist) + corrected degenerate 30ns timeouts to \*time.Second. (3) local\_providers\_integration\_test.go: 10 constructors (NewVLLMProvider/NewLocalAIProvider/NewFastChatProvider/NewTextGenProvider/NewLMStudioProvider/NewJanProvider/NewGPT4AllProvider/NewTabbyAPIProvider/NewMLXProvider/NewMistralRSPProvider) grep-verified DELETED from production → their dead tests removed; NewKoboldAIProvider survives → TestKoboldAIProviderIntegration + trimmed helper retained (honest SKIP, no endpoint). (4) test-only provider-type collision bug fixed (TestProviderHealthIntegration registered 2 mocks both keyed ProviderTypeLocal → distinct types). The new ollama integration stress/chaos tests now run live + PASS; make stress-chaos-infra extended to server + llm. Test-only changes (no production code touched). Cloud-provider tests honest-SKIP when API keys unset (env-dependent failures with keys present are pre-existing + out of scope). gofmt + build + vet clean. **Discovered:** 2026-05-28 (HXC-014 Ollama infra batch — surfaced running go test -tags=integration ./internal/llm/) **Severity:** Medium (CONST-050(B): the llm integration suite cannot compile/run; masks integration regressions + blocks the new ollama integration stress/chaos tests from running via the normal path) **Scope:** Under -tags=integration, internal/llm fails to compile: tool\_provider\_test.go/integration\_test.go redeclare MockProvider; local\_providers\_integration\_test.go + cloud\_providers\_integration\_test.go + integrated\_model\_manager\_test.go + cross\_provider\_test.go + model\_download\_manager\_test.go reference ~9 deleted constructors (NewVLLMProvider/NewLocalAIProvider/NewLlamaCplusplusProvider/... + a removed ModelPath field). Fix: dedup MockProvider + update/remove the stale tests to the current provider API (the providers were removed from production, so their dead tests should be deleted or rewritten against extant constructors), WITHOUT dropping legitimate coverage. Then go test -tags=integration ./internal/llm/ compiles and the new ollama\_provider\_{stress,chaos}\_test.go run live. **HXC-014a** (empty TestProviderStress stub) already FIXED (f464adb0). **Operator decision deferred:** promoting tests/stresschaos/ into the constitution submodule for cross-project reuse (triggers §11.4.26 cross-project workflow) — interim home is project-local.

---

## HXC-025 — Constitution §11.4.98/99/101 cascade (CONST-047/§3/§11.4.26)

**Status:** Completed (→ Fixed.md) **Type:** Task **Discovered:** 2026-05-29 (read-only fetch of constitution submodule surfaced new upstream commits 15cd4bc → 6017af9) **Closure (2026-05-29, subagent-driven §11.4.70):** the constitution submodule advanced with 3 new UNIVERSAL anchors — §11.4.98 (Full-Automation Anti-Bluff: live tests self-driving e2e), §11.4.99 (Latest-Source Documentation Cross-Reference), §11.4.101 (Autonomous-decision-over-blocking) — plus §11.4.100 (video-color) which was DEMOTED to the ATMOSphere project (project-specific, NOT cascaded). Per §11.4.26: pinned HelixCode's constitution submodule to 6017af9; cascaded the 3 universal anchors into HelixCode's 5 root govfiles (Phase 1, commit 901e7a55) AND all 68 owned submodules' CONSTITUTION/CLAUDE/AGENTS/QWEN) govfiles (Phase 2, 6 parallel subagent waves + a pilot — each submodule committed + pushed to its own org remotes; append-only; non-ff rejections union-merged preserving all governance content; no force-push). Bumped HelixCode's 68 submodule pointers + extended `verify-governance-cascade.sh` COVENANT114\_ANCHORS to enforce §11.4.98/99/101 (24 → 27). **Cascade verifier** → **0 failures** (root + all 68 submodules carry all 27 anchors); post-pull CONST-055 sweep (G1-G10) → 0 failures. §11.4.101's decision rule (reversible + evidence-determinable + bounded) authorised proceeding autonomously through the org-wide cascade. NOTE: the constitution submodule's own new mandate §11.4.98 (live-test full-automation) + §11.4.99 (latest-source doc xref) imply future HelixCode work (audit live tests for manual-action dependencies; add Sources-verified footers to operator docs) — tracked for follow-up.

---

## HXC-026 — workable-items md ↔ db sync gate (§11.4.93/95 follow-up)

**Status:** Completed (→ Fixed.md) **Type:** Task **Discovered:** 2026-05-29 (HXC-013 follow-up — the binary + tracked DB existed but no gate enforced md ↔ db sync) **Closure (2026-05-29):** `scripts/gates/workable_items_sync_gate.sh` (CM-WORKABLE-ITEMS-MD-DB-IN-SYNC) builds the constitution workable-items binary and asserts three invariants on TEMP COPIES (never opening the tracked DB in-place — SQLite WAL-mode dirties the header even on read): (1) the committed docs/workable\_items.db validates; (2) Issues.md/Fixed.md round-trip

md → db → md byte-identically; (3) the committed DB's md projection matches the live docs (DB not stale). Honest SKIP-OK when the CGO/sqlite binary can't build in-env (never a fake pass). Wired into the CONST-055 sweep as **G11** (full sweep G1-G11 → 0 failures). Paired-mutation meta-test scripts/tests/workable\_items\_sync\_meta\_test.sh (plant a phantom item in Issues.md → gate FAILs → trap-restore byte-identical → gate PASSes; 3/3 assertions). bash -n clean. Completes the HXC-013 §11.4.93 enforcement loop (the remaining auto-invoke-from-commit\_all.sh wiring stays a forward note).

---

## HXC-027 — §11.4.98 live-test full-automation compliance audit

**Status:** Completed (→ Fixed.md) **Type:** Task **Discovered:** 2026-05-29 (constitution §11.4.98 cascaded in HXC-025 — mandate requires classifying every test COMPLIANT vs NON-COMPLIANT) **Closure (2026-05-29):** audited HelixCode's \*\_test.go suite for §11.4.98 manual-action anti-patterns — **RESULT: COMPLIANT**. Scans: (a) stdin/manual-input dependency (bufio.NewReader(os.Stdin)/fmt.Scan/os.Stdin.Read) → ZERO; (b) operator-prompt waits (“operator must”/“please type”/“press enter”/“manually run”) → ZERO; (c) silent t.Skip() without a reason/SKIP-OK marker → ZERO (all skips cite a reason per §11.4.3); (d) human-response-window waits → the only long time.Sleeps (discovery health-timeout 10s, token\_budget 61s, e2e propagation 10-15s, hardware/load 10-15s) are DETERMINISTIC machine-timing waits (waiting for a timeout/budget-window to elapse, then continuing automatically) — fully self-driving, NOT human-response windows, so §11.4.98-compliant. The suite reports PASS/FAIL/SKIP-with-reason without human action after startup. **Forward note (§11.4.82, not a §11.4.98 violation):** internal/llm/token\_budget\_test.go:436 sleeps a real 61s to exercise a 60s budget window — a configurable/shorter window would speed iteration; non-blocking. No NON-COMPLIANT tests found → no remediation campaign needed (unlike the CONST-046 i18n campaign).

---

## HXC-028 — §11.4.99 latest-source documentation cross-reference (README)

**Status:** Completed (→ Fixed.md) **Type:** Task **Discovered:** 2026-05-29 (constitution §11.4.99 cascaded in HXC-025 — operator-facing instruction docs MUST be verified vs latest official sources) **Closure (2026-05-29):** applied §11.4.99 to the primary operator-facing doc README.md — cross-referenced its setup/build instructions against the latest official sources (WebFetch <https://go.dev/doc/devel/release>) + the repo's actual state. **3 real instruction defects found + fixed:** (1) prerequisite “Go 1.24.0+” → “Go 1.26+” (the inner helix\_code/go.mod is go 1.26 — it does NOT build below 1.26; verified Go 1.26.0 released 2026-02-10 / latest 1.26.3, and Go 1.24 is now past support per go.dev's two-newer-majors policy); (2) clone URL <https://github.com/your-org/helixcode.git> → SSH git@github.com:HelixDevelopment/HelixCode.git (Rule 3 SSH-only / CONST-038); (3) build step cd HelixCode → cd helix\_code (lowercase per CONST-052; on-disk dir verified). Also fixed the Support section's your-org placeholder links → in-repo tracker pointers. Added a ## Sources verified footer (date + go.dev URL + repo cross-ref + a negative finding: README still uses the legacy §11.4.44 bold header, §11.4.61 table migration deferred). PostgreSQL 15+/Redis 7+ confirmed consistent with CLAUDE.md §3.1. **4th finding (follow-up):** the README Documentation section's 4 links (helix\_code/docs/{ARCHITECTURE,DEVELOPMENT,USER\_GUIDE,API}.md) were ALL broken — the canonical docs live under helix\_code/docs/general/ (and API → API\_DOCUMENTATION.md); all 4 links repointed + verified to resolve. README.html + README.pdf regenerated (CONST-062/066 sync). **Extended to helix\_code/docs/general/DEVELOPMENT.md:** fixed “Go 1.21+” → “Go 1.26+” (verified vs go.dev + go.mod); replaced the docker build/docker run section (Rule 4 violation — host uses podman) with the real repo-root ./helix facade (subcommands start/status/logs/restart/stop, verified against the script); all 8 cited make targets verified present; added a ## Sources verified footer. **§11.4.99 negative finding fixed in CLAUDE.md §3.4:** the listed make container-builder-image/container-build/container-test/container-shell/container-release targets do NOT exist in any Makefile — replaced with the real ./helix facade + a note. **Systemic sweep (2026-05-29):** scanned the doc tree for the recurring stale defects + fixed: helix\_code/docs/general/ARCHITECTURE.md (“Go 1.24+” → “Go 1.26+”), USER\_MANUAL.md (“Go 1.24.0+” → “Go 1.26+”), USER\_GUIDE.md (Docker-install docker run/docker-compose → ./helix facade + SSH clone + cd helix\_code build path), COMPLETE\_CLI\_REFERENCE.md (“Docker Commands” recommending non-existent helixcode-security-scanner/-performance-optimizer/-config-

validator images → real make scan-all + go run ./cmd/  
performance\_optimization + go run ./cmd/config\_test). Every  
replacement command was verified to exist before citing (the point of  
§11.4.99). Total §11.4.99 pass: 7 operator-facing docs corrected, all  
systemic stale-Go-version + Rule-4-docker + broken-path/link + non-  
existent-command defects fixed. The §11.4.99 discipline now binds  
future operator-doc edits.

---

## **HXC-029 — §11.4.98 full-automation compliance sweep of every live/integration/ e2e/Challenge test (no human-in-the-loop) — CLOSED (→ Fixed.md)**

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full  
closure record. **Closure (2026-05-29):** §11.4.98 manual-intervention  
sweep COMPLETE — 0 remaining human-in-the-loop violations. Static  
audit (docs/qa/HXC-029/compliance-ledger.md) found exactly 2 NON-  
COMPLIANT → both FIXED (server\_timeout\_test.go manual-skip → real  
self-driving net/http test -count=3 green; clean.sh interactive read -p →  
--force/tty-gated). All 7 **HelixCode-scope HelixQA banks** verified  
self-driving vs the live :8080 server (4 API + 3 CLI, each 3× deterministic  
+ flip-mutation-proof, real bin/cli via os/exec, grep -c manual-review-  
required=0, honest\_skip for absent tools). **31/31 integration files**  
runtime-verified self-driving (0 manual deps; the 3 FAILs were a real  
product defect → HXC-036, now fixed). The **20 browser/Android/  
capture banks are OUT of scope** (HelixQA is shared — they target  
Catalogizer/Yole/HelixQA-engine; HelixCode has no web UI (API-only) +  
no Android app; the 2 connected devices are ATMOSphere hardware)  
— per §11.4.79/§11.4.51 not converted. e2e suites static-clean (only  
config-bootstrap skips = permitted §11.4.98(B) exception; no manual  
signal). Evidence docs/qa/HXC-029/{compliance-ledger/  
run\_,integration-classification,playwright-android}. Closed Completed  
(→ Fixed.md) per CONST-057 (Type Task).

---



## HXC-030 — §11.4.99 forward: latest-source documentation cross-reference sweep across all operator-facing docs — CLOSED (→ Fixed.md)

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Closure (2026-05-29):** §11.4.99 operator-instruction sweep COMPLETE — **38/38 (100%)** operator-facing instruction/guide/manual/setup/troubleshooting/tutorial docs now carry a WebFetch-verified `## Sources verified` footer (8 batches: Go 1.24 → 1.26.3, golang:1.21 → 1.26 Dockerfiles, go1.21.5 → 1.26.3, postgres 14 → 15, stale Anthropic/OpenAI doc-redirect URLs all corrected against live official sources; honest negative findings recorded where sources were 403/unreachable; per CONST-036 model IDs flagged-not-guessed). New scripts/gates/sources\_verified\_gate.sh wired as **G13 (now ENFORCING)** in verify-all-constitution-rules.sh — any future operator doc lacking a footer FAILs the sweep. Out of §11.4.99 scope (evidence/internal, not operator instructions): docs/qa\_evidence/ (93 QA reports), docs/helix\_qa/, docs/architecture/, docs/coverage/, docs/materials/. Ongoing 90-day staleness re-verification is steady-state discipline (gate --check-stale), not an open task. Section retained as §11.4.19 tombstone. **Type:** Task **Discovered:** 2026-05-29 (constitution §11.4.99 cascaded via HXC-025; HXC-028 applied it to README only) **Discovered-By:** AI **Forensic-anchor:** §11.4.99 — “ALWAYS check against latest versions of services we use web / online docs before creating instructions”. **Scope:** Extend the HXC-028 README treatment to every operator-facing guide/manual/setup/troubleshooting doc under docs/; for each, WebFetch the latest official source of every referenced service/library, cross-reference each instruction, add a `## Sources verified <date>` footer + commit-message footer; flag negative findings; classify docs >6 months stale (>90 days for risk-classified services) for re-verify or §11.4.90 Obsolete (Reason=stale-documentation). **Closure criteria:** Every operator-facing doc carries a `## Sources verified` footer with URLs+date; release-gate check for the footer added; stale-beyond-grace docs triaged. **Composes-with:** §11.4.99, §11.4.92, HXC-028.

---

## HXC-032 — LLMOrchestrator submodule: committed merge-conflict markers break helix\_agent build — CLOSED (→ Fixed.md)

**Status:** Fixed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Closure (2026-05-29):** all 26 conflict hunks across 5 LLMOrchestrator Go files resolved to the HEAD (i18n-migrated) side; bundle.go BundleTranslator gained an honest TPlural; automation\_test.go aligned to the lowercase upstreams/ rename.go go build/go vet/go test ./... (10/10 pkgs) PASS; downstream helix\_agent go build ./... exit 0. Submodule d3956ad pushed origin/master (FF); meta pointer bumped. Section retained as a §11.4.19 migration tombstone. **Type:** Bug **Severity:** High (breaks helix\_agent go build ./...; a §11.4 PASS-bluff at the build layer — tracked source does not compile) **Discovered:** 2026-05-29 (surfaced by scripts/const052\_verify\_refs.sh CHECK 3 while investigating HXC-031)

**Discovered-By:** AI **Evidence:** submodules/llm\_orchestrator (digital.vasic.llmorchestrator) has UNRESOLVED git conflict markers committed into 5 tracked Go files — pkg/i18n/translator.go (1 hunk), pkg/i18n/translator\_test.go (2), pkg/agent/multi\_pool.go (1), cmd/orchestrator/main.go (7), challenges/runner/main.go (16) = 26 hunks. go build ./... fails expected 'package', found '<<'. The markers are present in commit 5d9f5fc and the current HEAD merge 1e198e3 "Merge branch 'master'", and are **already pushed to origin/master**.

**Root cause (forensic, no guessing per §11.4.6):** a merge between the i18n-migration lineage (clean at 8032035 / a7fda2a round-383 CONST-046) and the CONST-052 rename commit 4350384 "fix(const052): rename Upstreams/→upstreams/ (HXC-001)" was committed with conflict markers unresolved. **Resolution direction (verified-correct side):** the <<<<<< HEAD side is the canonical one — consumer analysis proves it: cmd/orchestrator/i18n\_msg.go, pkg/agent/claudecode\_agent.go, and the package tests call i18n.Pkg().T(), i18n.SetPkgTranslator(), and Translator.TPlural, which ONLY the HEAD (i18n-migrated) side defines; the 4350384 side is the older pre-TPlural/pre-Pkg() variant. Tr()/Trf()/Global() live in bundle.go (no collision). Fix recipe: restore the 5 files from the last-clean i18n-lineage commit (8032035 for translator.go; locate equivalents per file), go build ./... + go test ./... GREEN as the verifier, then go build ./... in helix\_agent GREEN, commit to the submodule, push origin/master per §11.4.71 (merge-first per §11.4.41 since the broken state is already on the remote), bump the meta .gitmodules pointer in the same meta commit. **Why not fixed in the discovering session:** 26 hunks across 5 files on an already-pushed branch is irreversible high-blast-radius

work whose per-hunk safe path needs build+test verification (§11.4.101) — deferred to a dedicated verified pass rather than rush-pushed. `translator.go` resolution was proven correct but reverted to keep the submodule in a single coherent committed state until all 5 are fixed together. **Composes-with:** §11.4 (build-layer bluff), §11.4.41 (merge-first), §11.4.71, HXC-001, CONST-046, CONST-052.

---

## HXC-033 — codegraph 0.9.7 update: full index/sync crashes + own-org submodules dropped from the index (§11.4.79 regression) — CLOSED (→ Fixed.md)

**Status:** Fixed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Closure (2026-05-29):** ROOT CAUSE confirmed = codegraph 0.9.7 requires an explicit `codegraph init` before `index` (data-compat change; old DB incompatible) — exactly the operator’s hypothesis. Fix (operator-directed): full wipe of the gitignored DB + `codegraph init` (tracked `config.json` preserved) + `codegraph index .. Result Files 75,663 / Nodes 1,272,492` (edges finalize async). §11.4.79 probe PASSES via `codegraph query: NewBundleTranslator→submodules/llm_orchestrator/...+vasic-digital/...` (10 own-org hits), third-party `LLama_CPP` excluded. Two earlier mis-diagnoses corrected per §11.4.6 (the “crash” was a faulty `pgrep` pattern; “own-org unreachable” used the wrong verb `search` vs `query` + a stale MCP DB). Also cleaned `Status.md` 3.66 MB → 8 KB (ANSI-spinner bloat from `codegraph_sync.sh`). Follow-ups (non-blocking, noted in `Status.md`): restart the codegraph MCP server to serve the fresh DB; fix `codegraph_sync.sh` to strip ANSI. Section retained as §11.4.19 tombstone. **Type:** Bug **Severity:** High (§11.4.79 release-blocker — AI agents querying the code-graph get NO own-org submodule symbols; index also unbuildable) **Discovered:** 2026-05-29 (operator installed codegraph 0.9.7; surfaced during §11.4.80 post-update sync) **Discovered-By:** AI **Operator-Blocked-Details:** By: AI; Since: 2026-05-29; Reason: external-tool-state (operator-installed codegraph 0.9.7 crashes with no actionable diagnostic — not determinable/fixable from captured evidence per §11.4.6); Unblock: operator decision required — (a) downgrade codegraph to the last version that indexed this repo cleanly, (b) file an upstream bug with the maintainer, or (c) accept a degraded/partial index temporarily. §11.4.80 mandates latest-installed, so an autonomous downgrade would itself violate it — hence operator-gated. **Evidence:** `codegraph --version→0.9.7`. Full codegraph

index . KILLED mid-run twice (no exit code/diagnostic; left 54,207 then --force 4,630 files vs HXC-017's 76,044 baseline); codegraph sync . exit 1 (8,461 files). MCP codegraph\_search BundleTranslator resolves ONLY helix\_code/..., NOT the own-org submodules/llm\_orchestrator/pkg/i18n/bundle.go → own-org unreachable. Host memory ample (51 GiB free — not §12.6 OOM). Tracked .codegraph/config.json intact (own-org includes + §11.4.10 credential excludes present — NOT a config regression). Logs: qa-results/codegraph\_index\_\*.log, codegraph\_recover\_\*.log; docs/codegraph/Status.md 2026-05-29 entry. **Root cause:** UNCONFIRMED (§11.4.6) — 0.9.7 index/sync terminate without diagnostic on this 76k-file repo; whether stability bug / submodule-traversal change / config-schema change is undetermined. **Composes-with:** §11.4.78, §11.4.79, §11.4.80, §11.4.6, §11.4.101, HXC-017.

---

## **HXC-034 — Cascade constitution §11.4.102 into owned submodules + implement CM-COVENANT-114-102-PROPAGATION gate — CLOSED (→ Fixed.md)**

**Status:** Completed (→ Fixed.md) — see docs/Fixed.md for the full closure record.

---

## **HXC-035 — POST /api/v1/auth/register returns 400 internal\_auth\_failed\_create\_user on the live server (blocks all authenticated flows) — CLOSED (→ Fixed.md)**

**Status:** Fixed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Closure (2026-05-29, systematic-debugging per §11.4.102):** ROOT CAUSE (confirmed via direct psql INSERT → ERROR: column "display\_name" of relation "users" does not exist): createSchemaSQL's CREATE TABLE users omitted display\_name, while the compensating ALTER TABLE ... ADD display\_name migration in InitializeSchema() runs ONLY in the if schemaExists branch — so a FRESH DB takes the else path, creates users without display\_name, and auth\_db.CreateUser's INSERT fails (error swallowed by the i18n translator into the generic 400). FIX: added display\_name VARCHAR(255)

to `helix_code/internal/database/database.go createSchemaSQL`.  
VERIFIED: fixed server → register HTTP 201 (was 400) + login → valid session token (evidence docs/qa/HXC-035/fix-verification.txt). Unblocks the HXC-029 banks' authenticated-positive paths. **Type:** Bug **Severity:** High (no user can register → no JWT mintable → every authenticated API path is undrivable; blocks the positive-path coverage of all 4 HXC-029 API banks) **Discovered:** 2026-05-29 (surfaced by HXC-029 bank verification against the live server) **Discovered-By:** AI **Evidence:** Against the live HelixCode server (real PG+Redis on :8080, schema freshly created), POST `/api/v1/auth/register` with a well-formed body returns HTTP 400 `internal_auth_failed_create_user`. Reproduced across all 4 HXC-029 API-bank verification runs (entity-management/admin-operations/security-validation/full-qa-api each had to `_skip` their authenticated-positive paths with `#HXC-029-REGISTER-BROKEN`). Captured in docs/qa/HXC-029/\*/endpoint-probes.txt. **Root cause:** PENDING\_FORENSICS — being root-caused now via `superpowers:systematic-debugging` per §11.4.102 (the create-user DB write fails; need the auth handler + DB-layer forensic to determine why: schema mismatch / constraint / missing migration / password-hash error). **Composes-with:** §11.4.102 (mandatory systematic-debugging), §11.4.98 (it blocks live-test positive paths), HXC-029.

---

\*Last updated: 2026-05-29 — filed HXC-035 (POST `/auth/register` 400 `internal_auth_failed_create_user` — High, systematic-debugging in progress per §11.4.102); HXC-029 now 4/18 banks verified (full-qa-api + entity-management + admin-operations + security-validation, each 3×+mutation vs live server); filed HXC-034 (cascade constitution §11.4.102 into 68 owned submodules + gate — Task); constitution submodule §11.4.102 added+pushed (656b43a), meta pointer bumped; HXC-029 full-qa-api bank verified (§11.4.98); HXC-030 §11.4.99 sweep COMPLETE (38/38). Prior: filed HXC-033 (codegraph 0.9.7 index/sync crash + §11.4.79 own-org regression — Operator-blocked); HXC-032 FIXED+closed (LLMOrchestrator conflict markers; submodule d3956ad, `helix_agent` builds); reclassified HXC-031 (CONST-052 renames RESOLVED/none-remain, only Codex/Cline ports remain); HXC-029 §11.4.98 2 confirmed violations fixed; HXC-030 §11.4.99 Go 1.24 → 1.26.3 + PG 14 → 15 doc reconciliation. Prior: filed HXC-029 (§11.4.98 forward sweep), HXC-030 (§11.4.99 forward sweep), HXC-031 (deferred rename/port long-tail) per operator “do it all”; added scripts/`generate_{issues,fixed}_summary.sh` + G12 summary-freshness gate (§11.4.91/12). Previously: 2026-05-28 — constitution submodule pulled 7f738df → 15cd4bc (§11.4.79–97); HXC-013..019,022 filed (open: SQLite-DB / stress+chaos / cross-platform / submodule-cascade / codegraph-own-org / obsolete+summary-tooling / docs-qa / test\_bank-noncompile);

HXC-021 + HXC-014a + HXC-015a FIXED → Fixed.md (commit f464adb0 — fake-skip Assert(true) bluffs + empty stress stub → honest SKIP); CONST-052/HXC-001 leaf-rename programme COMPLETE (Phases 1-4), Phase 5 org-grouping dirs kept as namespace carve-outs per operator decision 2026-05-28 → HXC-001 closeable. Prior: 2026-05-20 (round 463 — HXC-003 closed Implemented (→ Fixed.md) and migrated to docs/Fixed.md: the CONST-046 i18n migration campaign is concluded — the genuine user-facing (C) string-literal surface is exhausted across all 7 scope areas (helix\_code internal/+cmd/+applications/, LLMsVerifier, helix\_qa, all owned vasic-digital/\*+HelixDevelopment/\* submodules); ~91-462 rounds migrated tens of thousands of literals with paired-mutation anti-bluff tests; remaining ~55k audit hits are all out of CONST-046 scope per docs/audits/2026-05-20-internal-const046-classification.md. Open set is now HXC-001 (CONST-052 renames — Task, In progress) + HXC-010 (Kimi/Qwen codegraph e2e — Operator-blocked Task)). Previous round 402 — HXC-011 closed Fixed (→ Fixed.md): the helix\_qa runner's run path on the desktop platform now genuinely executes a bank case's shell: action via os/exec. Round 400 — speed-programme close-out: HXC-006 closed Implemented (→ Fixed.md). To update Issues\_Summary.md mechanically, run scripts/generate\_issues\_summary.sh (TODO: create — currently this Issues.md is the source of truth and Summary is hand-maintained).\* ## HXC-036 — Systemic CONST-046 i18n defect: 74 packages emitted raw message-ID keys because boot-time translator wiring was never implemented — CLOSED (→ Fixed.md)

**Status:** Fixed (→ Fixed.md) — see docs/Fixed.md for the full closure record. **Closure (2026-05-29, Option A boot-wiring, 4 phases):** the CONST-046 migration built 74 SetTranslator seams + bundles but never wired a real translator at boot (0 .SetTranslator( call sites module-wide). Fix: per-package bundle.go (//go:embed active.en.yaml → i18n.NewBundle/Localizer → i18nadapter.New) + central internal/i18nwiring.WireAll() (63 internal pkgs incl. shared internal/workflow/i18n) called at boot + in integration TestMain; 9 package main binaries self-wire via cmd/<m>/i18n\_boot\_wire.go init(). VERIFIED: the 3 originally-failing integration tests PASS with resolved interpolated text; WireAll() returns nil at 74/74; resolved-text captured for askuser (“Enter choice [1-3]:”), approval, auth, llm, config, cli (“Inspect or run user-defined Markdown slash commands”), autonomy (“Full Auto (Fully Autonomous)”), planmode (“Score: 87.5/100”); paired-mutation proven (no WireAll → raw keys → FAIL). Evidence docs/qa/HXC-036/phase{1,2,3,4}/. Commits f3b864f4 (P1) + 31c57a2a (P2, 70 pkgs) + 1ea79fd2 (P3, 9 mains) + d570b05e (P4, autonomy/planmode).

---

## **HXL-001 (ex-ISSUE-003) — HelixLLM**

### **internal/agents/tools/analysis\_test.go**

### **hardcoded absolute path**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 95 — HelixLLM migration; surfaced as pre-existing failure)  
**Discovered-By:** AI subagent during HelixLLM standalone test run  
**Closed-By:** Round 105 (commit a5e56d4 in HelixLLM; meta pointer fedd152) **Attribution correction:** Originally documented as helix\_agent; actual location is HelixLLM submodule (submodules/helix\_llm/internal/agents/tools/). Commit SHAs 0a84310 resolved there. **Resolution:** Replaced hardcoded path with t.TempDir() + 2 synthesised fixture files. Bonus: same bug-pattern discovered in git\_test.go (constant helixLLMRoot + 7 tests) — refactored gitSandbox() signature. 6 tests now PASS on any host. Mutation verified.

---

## **HXL-002 (ex-ISSUE-004) — HelixLLM**

### **internal/gateway/middleware TOON WriteTOON**

### **returns 500**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 95) **Discovered-By:** AI subagent **Closed-By:** Round 105 (commit a5e56d4) **Attribution correction:** Originally documented as helix\_agent; actual location is HelixLLM submodule. Commit 6f11c56 resolved there. **Resolution:** Root cause was vasic-digital/TOON's round-27 anti-bluff change (Marshal returns ErrTOONEncodingNotImplemented unconditionally) combined with WriteTOON treating ANY Marshal error as 500. Fix: fall back to json.Marshal while preserving application/toon Content-Type (matches ContentNegotiation middleware). 500 still returned for genuinely unmarshallable values (channels). 19 middleware tests now PASS. Mutation verified.

---

## HXQ-001 (ex-ISSUE-008) — helix\_qa intermittent TestPerformance flake (host-load-sensitive)

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 82) **Discovered-By:** AI subagent **Fixed:** 2026-05-20 (round 325)  
**Evidence:** helix\_qa TestPerformance (three perf tests in pkg/vision/ — TestPerformance\_DHash64\_Under5msPer1080pFrame, TestPerformance\_PHash\_Under25msPer1080pFrame, TestPerformance\_SSIM\_Under5msPer480pFrame) fails intermittently under high host load (concurrent containers + builds). Not a code bug per se; a sensitivity issue — the hard per-frame timing ceilings (5 ms / 25 ms / 5 ms) are only meaningful on a quiescent host. **Resolution:** Decision — **path (b)** (env-var gating) chosen over path (a) (loosen tolerance). Rationale: loosening the timing tolerance would weaken the test's anti-bluff value — a genuine perf regression could then pass. Path (b) preserves the strict assertions while making the flake deterministic: the three tests now check `os.Getenv("HOST_LOAD_DEDICATED")` and `t.Skip("SKIP-OK: #HXQ-001 ...")` honestly when unset, running strict only on a quiescent dedicated host (`HOST_LOAD_DEDICATED=1`). This is the CONST-035-compliant choice. Landed in helix\_qa submodule commit 649e2dd + meta .gitmodules pointer bump. docs/test-coverage.md §6.1 documents the env var. Post-fix evidence: `go build ./pkg/vision/...` exit 0, `go vet clean`; `go test -count=1 -run TestPerformance ./pkg/vision/...` (unset) → all 3 --- SKIP with SKIP-OK: #HXQ-001 marker; `HOST_LOAD_DEDICATED=1 go test -count=1 -run TestPerformance ./pkg/vision/...` → all 3 --- PASS strict (DHash64 average 741ns, PHash average 88.969µs — well under the 5 ms / 25 ms ceilings).

---

## HXV-002 — LLMsVerifier verification/ package 10 pre-existing test failures

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-20 (round 345 — LLMsVerifier i18n round-12 subagent) **Discovered-By:** AI subagent — `git stash test` confirmed the 10 failures reproduce at submodule HEAD 582ae9c7 (round-336) *without* the round-345 i18n change, proving pre-existing and unrelated **Resolved:** 2026-05-20 (round 348) **Evidence:** `go test ./verification/...` in submodules/llms\_verifier/llm-verifier/ reported 10 failures; after round-348 fix ok digital.vasic.llmsverifier/verification (1.635s), 0 failures, go



build ./... clean. **Resolution:** All 10 failures classified **(A) test-assertion drift** — every failing test asserted pre-honesty fabricated behaviour that round-17 commit a6328629 correctly removed. **No production code changed.** Per-failure classification table:

| #  | Test                                                              | File         |
|----|-------------------------------------------------------------------|--------------|
| 1  | TestVerifier_Verify_Success                                       | verification |
| 2  | TestVerifier_Verify_ResultScores                                  | verification |
| 3  | TestVerifier_Verify_LatencyMetrics                                | verification |
| 4  | TestVerifier_Verify_CodeLanguageSupport                           | verification |
| 5  | TestVerifier_Verify_CodeCapabilities                              | verification |
| 6  | TestVerifier_Verify_ModelStatusFlags                              | verification |
| 7  | TestVerifier_Verify_ContextCancellation                           | verification |
| 8  | TestVerifier_Verify_MultipleRequests                              | verification |
| 9  | TestCodeVerificationService_TestCodeVisibility_Error              | code_verif   |
| 10 | TestCodeVerificationService_VerifyModelCodeVisibility_ServerError | code_verif   |

Mirrors HXV-001 round-323’s classification approach. The production code (verification.go round-17 ErrVerificationNotWired, code\_verification.go error propagation + zero-response → failed) was already honest; only the stale test assertions needed re-keying to certify it.

## PAN-001 — panoptic appendString truncates multi-byte UTF-8 runes to bytes (TestResult.MarshalJSON corrupts non-ASCII)

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Discovered:** 2026-05-19 (round 298 — panoptic enrichment subagent) **Discovered-By:** AI subagent runner-detector against real executor.TestResult.MarshalJSON **Fixed:** 2026-05-19 (round 302 — panoptic submodule commit 24aa627 + meta pointer bump) **Evidence:** panoptic/internal/executor/executor.go:120 — buf = append(buf, byte(r)) in the else branch of appendString casts a rune to a single byte. Multi-byte UTF-8 codepoints (German umlauts, Spanish accents, Japanese CJK, Serbian Cyrillic, Chinese Han) get truncated to one byte each, producing corrupted JSON output. Honestly tracked via the round-298 Challenge runner's executor-marshall:utf8-detector:regression-present PASS line + KNOWN-ISSUE entry in panoptic/docs/test-coverage.md §7. Affects every consumer that JSON-marshals a TestResult containing non-ASCII text. **Resolution:** Replaced buf = append(buf, byte(r)) with buf = utf8.AppendRune(buf, r) (Go 1.21+) + added unicode/utf8 import. Single-line functional fix. Post-fix evidence: go test -race -count=1 ./internal/executor/... → ok 4.470s; bash challenges/panoptic\_describe\_challenge.sh → 39/39 PASS, 0 FAIL; runner UTF-8 detector flipped from regression-present → fixed (literal log: PASS [executor-marshall:utf8-detector:fixed] + KNOWN-ISSUE-RESOLVED: executor.appendJSONString now UTF-8 clean). Closed in this round.

---

## VEN-001 (ex-ISSUE-001) — VisionEngine helix-gitlab remote repo missing (404) — CLOSED (→ Fixed.md)

**Status:** Completed (→ Fixed.md) **Type:** Task **Discovered:** 2026-05-19 (round 98 — Planning + VisionEngine i18n migration) **Discovered-By:** AI subagent during 4-remote push attempt **Closed-By:** Round 188 (subagent repo-inventory sweep) **Root cause:** The helix-gitlab remote URL in submodules/vision\_engine/.git/config pointed at git@gitlab.com:HelixDevelopment/visionengine.git — a non-existent group path. The actual GitLab group is helixdevelopment1 (path) / HelixDevelopment (display name). The repository helixdevelopment1/VisionEngine (id 80411994) already existed since 2026-03-19. NOT a missing-repo issue — a URL-misconfiguration issue. **Fix:** git remote

set-url helix-gitlab git@gitlab.com:helixdevelopment1/  
VisionEngine.git in the VisionEngine submodule, then git push helix-  
gitlab master (FF-safe: local was 46 commits ahead, remote 0 ahead).  
Push landed at SHA 2d0c35b (verified via git ls-remote helix-gitlab  
master). The Upstreams recipe push-helix-gitlab.sh references the  
remote by name (not URL), so it continues to work unchanged.  
**Evidence:** - glab api projects/helixdevelopment1%2Fvisionengine → id  
80411994 OK - git ls-remote helix-gitlab HEAD →  
2d0c35bebb199a9a199fbf899eae292e38eaf17 (matches local HEAD) -  
Original broken URL still 404s when probed directly (proves URL was  
the issue, not perms)

---

## HXC-126 — Eleven closed items still appear in the open-issues tracker instead of the resolved tracker

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/  
hxc126\_20260712T142600Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

Eleven work items marked finished still appear in the open-issues document and are missing from the resolved-items document, so the two views disagree about their state and become untrustworthy. The work regenerates the tracker documents from the authoritative database so finished items appear only in the resolved view. The human-readable trackers then accurately reflect the true state.

## HXC-142 — Automation test suite (test/automation, -tags=automation) does not compile

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/  
hxc142\_143\_20260712T150000Z/EVIDENCE.md **Severity:** High **Created-By:** Claude

The automation-tagged test package fails to build, so an entire mandated test type cannot execute at all. Two real causes were found during the 2026-07-12 real-infra retest: a duplicate symbol (isRateLimitError/contains declared in both xai and qwen automation test files) and deeper API drift where the tests reference

llm.ProviderConfig and NewProviderManager which no longer exist in the current provider package. The work is to reconcile the automation tests with the current provider API and remove the duplicate helpers so the suite compiles and runs against real infrastructure. Evidence: docs/qa/infra\_retest\_20260712\_hxc122\_138/automation\_tests.log.

## **HXC-143 — E2E test suite (test/e2e, -tags=e2e) does not compile due to redeclared getEnvOrDefault**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc142\_143\_20260712T150000Z/EVIDENCE.md **Severity:** High **Created-By:** Claude

The e2e-tagged test package fails to build because getEnvOrDefault is declared more than once in the package, so another mandated test type cannot execute. The work is to remove the duplicate declaration (consolidate to a single shared helper) so the e2e suite compiles and can run end-to-end against a real server. Discovered during the 2026-07-12 real-infra retest. Evidence: docs/qa/infra\_retest\_20260712\_hxc122\_138/EVIDENCE.md.

## **HXC-118 — Retrieval-Augmented-Generation (RAG) module exists but is not connected to the application**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** docs/qa/hxc118\_20260712T151500Z/EVIDENCE.md **Severity:** High **Created-By:** Claude

A dedicated Retrieval-Augmented-Generation component is maintained as its own reusable module, but the main application does not import or use it anywhere. A capability the product is expected to offer (answering using retrieved documents) is therefore effectively unavailable to end users despite the code existing. The work integrates the existing RAG module into the application, wires it into the request flow, and exposes its capability flag. Users gain working document-grounded answers instead of an orphaned, unused component.

## **HXC-144 — Server leaks goroutines under sustained DDoS-flood load (chaos test)**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc144\_20260712T152500Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

Under the sustained request-flood chaos test against the real running server, the goroutine count grew by 5 which exceeds the tolerance of 4, signalling a goroutine leak in a request-handling path when the server is hammered. Left unaddressed this degrades long-running server stability under load. The work is to find the leaking goroutine (likely an unclosed channel, context, or connection in a hot handler) and fix it so the count stays within tolerance under flood. Evidence: docs/qa/infra\_retest\_20260712\_hxc122\_138/EVIDENCE.md (Server 7/8).

## **HXC-122 — Memory and automation test suites mostly skip themselves without a running server**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/infra\_retest\_20260712\_hxc122\_138/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

Two categories of tests, memory-usage and end-to-end automation, skip most of their cases by default because they require a live server or special environment flags not set in normal runs. In practice these areas are largely unverified even though the tests appear to exist. The work provides a documented, repeatable way to run them against real infrastructure so they actually execute and prove the behavior. Memory and automation behavior then becomes genuinely tested rather than merely scaffolded.

## **HXC-136 — Verify the remaining automated test types run with real captured evidence**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/infra\_retest2\_20260712\_hxc136/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

Several mandated automated test categories — load/denial-of-service, scaling, stress and chaos, and user-interface/experience — were not exercised in the latest real-infrastructure run, so their current health is unconfirmed. The work is to run each of these test types against real infrastructure and capture proof of the results. This completes the promised full test-type coverage and confirms the product holds up under load and adverse conditions.

## **HXC-138 — Run the end-to-end challenge suite against a running server**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/infra\_retest\_20260712\_hxc122\_138/EVIDENCE.md **Severity:** Low **Created-By:** Claude

The end-to-end challenge runner can now launch all its scenarios (a missing option was just fixed), but the scenarios still need to be executed against a live server with a real model to confirm the complete user journeys work. The work is to stand up a server and run the challenges, capturing the results. This provides real proof that the headline user workflows function end to end.

## **HXC-150 — tests/ddos env-var names (TEST\_PG\_/TEST\_REDIS\_) mismatch .env.full-test causing false SKIP**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc150\_20260712T172000Z/EVIDENCE.md **Severity:** Low **Created-By:** Claude

The load/DDoS test harness (tests/ddos, -tags=integration) reads TEST\_PG\_\* and TEST\_REDIS\_\* environment variables to reach the real Postgres and Redis, but the projects .env.full-test only exports HELIX\_DATABASE\_\* and HELIX\_REDIS\_\*, and the defaults do not match the container credentials. So under the normal make test-load-full workflow the suite falsely SKIPS even when infra is fully healthy (it only ran once the correct credentials were passed manually during the 2026-07-12 retest). The work is to align the harness env-var names/defaults with .env.full-test so load/DDoS executes out of the box. Low severity, test-harness only.

## HXC-148 — Wire RAG retrieval-augmentation into the OpenAI/Anthropic wire-facade endpoints

**Status:** Fixed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc148\_20260712T173000Z/EVIDENCE.md **Severity:** Low **Created-By:** Claude

HXC-118 wired Retrieval-Augmented-Generation into the native server generate and stream endpoints and the CLI, but the OpenAI-compatible and Anthropic-compatible wire-facade endpoints (/v1/chat/completions and /v1/messages) still bypass RAG entirely, so clients using those compatibility surfaces do not get retrieval-augmentation even when it is enabled. The work is to apply the same applyRAGContext wiring to those facade handlers so RAG behaves consistently across every generate surface. This is a smaller secondary surface than the native path already fixed. Found during HXC-118 review 2026-07-12.

## HXC-149 — Stale git gitlink at pre-rename path containers breaks git submodule walk

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc149\_20260712T173500Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

The main repository git index carries approximately 70 stale submodule gitlinks at pre-rename paths (dependencies/HelixDevelopment/, *dependencies/vasic-digital/*, and top-level helix\_agent/helix\_qa/panoptic/security) with no .gitmodules mapping, left over from a historical path-rename to the submodules/ layout. As a result git submodule status and git submodule foreach abort mid-walk with no submodule mapping found in .gitmodules for path ..., so any release or maintenance script that walks all submodules unfiltered fails partway. The work is to remove ALL stale cached gitlinks (git rm –cached on each) so the submodule set is consistent with .gitmodules and submodule-walking tooling completes. This is a git-index-only change, reversible via git reset. Found by the 2026-07-12 release-readiness survey. Low runtime risk but blocks release automation.

## HXC-145 — Configured Xiaomi model mimo-v2-flash is rejected by the real Xiaomi API

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc145\_147\_20260712T190000Z/EVIDENCE.md **Severity:** Low **Created-By:** Claude

During the real-infra retest the Xiaomi provider chaos tests failed 2 of 5 because the model id configured for Xiaomi (mimo-v2-flash) is rejected by the live Xiaomi API, indicating the configured model name is stale or wrong. Users selecting the Xiaomi provider with that model would get errors. The work is to determine the correct current Xiaomi model id (from the provider or the verifier as single source of truth) and update the configuration so Xiaomi requests succeed. Evidence: docs/qa/infra\_retest\_20260712\_hxc122\_138/EVIDENCE.md (Xiaomi 3/5).

## HXC-147 — OpenRouter provider automation test nil-pointer panics on stale model deepseek-r1-free

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc145\_147\_20260712T190000Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

Running the (now-compilable) automation test binary against the live OpenRouter API, TestAllFreeProvidersAutomation Provider\_OpenRouter BasicGeneration panics with a nil-pointer dereference: the configured free model id deepseek-r1-free is stale/rejected and the code path is missing a nil-check on the error before using the response. Users of the OpenRouter free provider with that model would hit the same crash. The work is to correct the free-provider model id (sourced from the verifier as single source of truth) and add the missing nil-check so a rejected model degrades gracefully instead of panicking. NOTE this environment has live provider API keys set so provider tests spend real money; guard/skip accordingly. Found 2026-07-12.



## **HXC-146 — E2E challenge runner interfaces (cli/rest/tui/websocket) do not drive the real server HTTP API**

**Status:** Fixed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc146\_20260712T191000Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

The e2e challenge runner advertises multiple interface modes (cli, rest, tui, websocket) but none of them actually exercises the HelixCode server's real HTTP API during a run, so the challenges validate the runner's own logic rather than the shipped server endpoints. This is a documentation-versus-implementation gap that weakens the end-to-end proof. The work is to wire the challenge runner's interfaces to genuinely call the running server's HTTP API so the challenges prove the real user-facing endpoints work. Discovered 2026-07-12 real-infra retest. Evidence: docs/qa/infra\_retest\_20260712\_hxc122\_138/hxc138\_challenge\_report.json.

## **HXC-119 — Agent Client Protocol (ACP) support is absent from the platform**

**Status:** Implemented (→ Fixed.md) **Type:** Feature **Evidence:** docs/qa/hxc119\_20260712T193000Z/EVIDENCE.md **Severity:** High **Created-By:** Claude

Governance rule CONST-040 lists the Agent Client Protocol among required capabilities, but there is no implementation of it anywhere in the codebase. Any user or integration expecting ACP connectivity currently cannot use it. The work is to design and implement real ACP support, or, if it proves structurally infeasible, to document that with cited evidence. The platform will then either genuinely support ACP or hold an honest, evidenced position instead of an unmet claim.

## **HXC-152 — llm\_working\_funnel\_test.go defaults to RED mode, so its fix-verifying assertions never run and it still reports success**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc151\_hxc152\_replica\_red\_falsifiability\_20260729T0010Z/EVIDENCE.md **Severity:** High **Created-By:** Claude

One test file in the server package chooses between ‘reproduce the old bug’ mode and ‘confirm the fix works’ mode by reading an environment variable, but its default is backwards: when the variable is unset — which is how the suite runs normally and in every automated run — it selects the old-bug mode. The practical effect is that the checks confirming the fixed behaviour never execute at all, while the test still reports success, so the team sees a green result for behaviour nobody actually verified. The behaviour left unguarded is the model-listing funnel that is supposed to hide models which failed verification, scored below the quality threshold, or have no API key configured, which is precisely the anti-bluff filtering the original fix was written to provide. A second test in the same file simply skips itself outright under the same default, so the provider key-gate is unguarded too. Every other polarity switch in this package defaults the opposite way, so this file is also inconsistent with its neighbours and will mislead anyone who assumes the shared convention. To reproduce: run the package normally and observe that the guarded assertions are never reached while both tests report PASS. Done means the default is flipped so the fix-verifying assertions run in ordinary runs, the reproduce-the-old-bug mode still works when explicitly requested, and captured evidence shows those assertions genuinely execute and genuinely fail when the filtering is removed.

## **HXC-151 — Residual replica-RED branches in internal/server regression guards cannot fail, so they guard nothing**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc151\_hxc152\_replica\_red\_falsifiability\_20260729T0010Z/EVIDENCE.md **Severity:** Medium **Created-By:** Claude

Several regression tests in the server package carry a RED\_MODE switch whose job is to prove the test can actually catch the bug it guards. In these cases the RED half checks a private copy of the old code that lives inside the test file itself, instead of the real shipped code. Because that copy is part of the test and never changes when the product changes, the check produces the same result on every build ever made — the working one, the broken one, and any future one — so it can never fail and proves nothing about the product. This matters because everyone reads a green result as evidence that the guard is real, when the guard is in fact inert: the original defect could return and the suite would stay green. The confirmed cases are `handlers_systemstatus_guard_test.go` (calls a test-local `preFixSystemStatusHealthCheck` helper instead of the real `getSystemStatus` handler, so the nil-database guard in the product is bypassed by construction), `handlers_qastatus_deadlock_guard_test.go` (hand-rolls the locking sequence in `reproduceRecursiveRLockDeadlock` rather than calling the shipped handler, so it only re-proves a standard Go library property), `llm_rag_test.go` (deliberately skips the real handler and its retrieval step, so the ‘prompt was not augmented’ result is guaranteed by the test’s own omission), and `llm_working_funnel_test.go` (one case has no RED branch at all, and another calls a real function that the fix never touched). To reproduce: revert the matching product fix and re-run the test with `RED_MODE=1`; the branch still passes, which it must not. The work is to rewrite each RED branch so it drives the real shipped code path, exactly as was done for the three sibling guards in `llm_default_model_regression_test.go` and `llm_generate_regression_test.go`. Done means that for every listed test, `RED_MODE=1` passes on an artifact with only its own fix reverted and fails on the current artifact, that this four-way result is captured as evidence, and that a repeat of the enumerated search finds no remaining test-local stand-ins.

## **HXC-156 — harmony\_os background system monitor used an unsynchronised on/off flag as its stop signal, racing application shutdown**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/fyneui\_uithread\_siblings\_20260728T192405Z/03\_MUTATION\_harmony\_race\_returns.txt **Severity:** High **Created-By:** Claude

The Harmony OS desktop application runs a background system monitor that samples processor and memory usage every few seconds. To decide whether to keep going, that monitor repeatedly read a plain on/off marker, while the application's shutdown routine switched the very same marker off from a different execution thread — with nothing coordinating the two. Two threads touching the same value at the same moment with no coordination is a data race: the shutdown signal can be missed entirely, leaving the monitor running and still writing to shared state after the application believes it has torn itself down, and on some machines the value read is not guaranteed to be either the old or the new one. This was not theoretical; it surfaced as a reproducible test failure, with the Harmony OS test suite failing on its cleanup test on every single run once race detection was enabled. To reproduce it, run the Harmony OS test package with the Go race detector switched on; the report names the monitor loop and the cleanup routine as the two conflicting parties. The same audit found a second, quieter instance of the same problem in the same component: the measured processor, memory, temperature and power figures were written by the monitor thread while the system-monitoring screen read them for display, again with no coordination, which could paint a screen mixing numbers from two different sampling rounds. The fix makes the monitor stop on the same shutdown channel the rest of the application already uses, protects the shared figures and the status marker with a lock, and makes shutdown wait for the monitor to confirm it has genuinely finished before proceeding. Acceptance: the Harmony OS test package completes with zero races over repeated consecutive runs under the race detector, and deliberately restoring the old code makes the identical race reappear, proving the guard is real and not merely agreeable.

## **HXC-157 — harmony\_os and aurora\_os changed on-screen elements directly from background threads, which the Fyne UI toolkit forbids**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/fyneui\_uithread\_siblings\_20260728T192405Z/04\_GREEN\_harmony\_all\_fixes\_count3.txt **Severity:** High **Created-By:** Claude

The Harmony OS and Aurora OS desktop front-ends both perform slow work in the background — asking a language model for an answer, polling provider health, refreshing dashboard statistics, starting the server — and each of those background workers wrote its result straight into an on-screen element. The user interface toolkit this product uses does not permit that: on-screen elements may only be changed from the single thread that draws them, and any other thread must hand its update over to be applied there. Doing it directly means the drawing code can be reading an element at the exact moment a background worker rewrites it, which is a data race that can corrupt what the user sees or crash the application outright. The problem was disguised rather than hidden: in both applications the offending line sat directly beneath a comment stating ‘Update UI on main thread’, which was simply untrue — it is the identical false comment already removed from the desktop front-end when this same defect class was fixed there. Two further aggravating factors were found by auditing every background worker in both files rather than only the reported line: several of these workers also read on-screen elements to decide what to do, which races just as surely as writing them, and several ran on endless timers with no way to stop, so they kept modifying elements belonging to a window that had already been closed and leaked for the lifetime of the process. The fix routes every such update through the shared dispatch helper introduced for the desktop front-end, keeps each read-and-then-write pair inside a single handover so the read cannot be left behind, moves the reads that happen at dispatch time onto the drawing thread where they belong, corrects the false comments in place so a later edit is not invited to undo the fix on their authority, and makes the endless timers stop when the application shuts down. Acceptance: both packages build and pass with the race detector enabled over repeated runs, no background worker in either file touches an on-screen element without going through the helper, and the desktop front-end’s behaviour is unchanged.

## **HXC-153 —**

### **TestGuard\_GetSystemStatus\_WithDB\_StillReports is named for a database-present case it never exercises**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc153\_systemstatus\_withdb\_20260728T194000Z/README.md  
**Severity:** Medium **Created-By:** Claude

A regression test in the server package is named as though it proves the system-status endpoint behaves correctly when a real database IS attached, but the test actually builds a server with no database at all and then only checks that the response code is 200. That means it silently re-runs the same no-database case a neighbouring test already covers, and the database-present path it advertises has no coverage from it whatsoever. This matters because the name is what a reader trusts when deciding whether a behaviour is guarded: anyone scanning the suite would reasonably conclude the with-database path is protected when it is not. It is the same documentation-versus-reality defect class as HXC-152, where a file's comment disagreed with its code, and it was surfaced by the independent reviewer of that batch. The test's own comment does honestly defer coverage to another file, so nothing is being falsely asserted about the product — the defect is confined to the misleading name. Done means the test either is renamed to describe what it genuinely checks, or is given a fake-database seam so it earns the name it already has, with the choice backed by a captured run.

## **HXC-155 — Standardise the divergent polarity-switch environment variable convention across the server test package**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc155\_red\_polarity\_convention\_20260728T194815Z/README.md  
**Severity:** Low **Created-By:** Claude

The server package's regression guards each support a switch that flips them between reproducing an old defect and guarding against its return, but the switch is spelled four different ways across the files and is read through three different helper styles. Having several names is genuinely useful, because each guard corresponds to a different historical fix and testers need to flip them one at a time rather than all at once, so the goal is not to collapse them into one variable. The problem is that the inconsistency has already caused real defects twice: one file shipped with its default set the wrong way round so its verification checks never ran, and another shipped with a comment describing the opposite of what its code did. Both were fixed, but the drift that produced them remains. This work is to agree one naming pattern and one helper shape, apply it to every guard in the package, and record the convention in a single place so the next guard author copies something correct. Done means every switch follows the

documented pattern, every default is the safe standing-guard direction, and a check exists that would catch a future file breaking the convention.

## **HXC-154 — TestStartQASession\_Success asserts on a session field a background goroutine is concurrently changing**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc154\_qa\_session\_created\_status\_20260729T005900Z **Severity:** Medium **Created-By:** Claude

A QA-session test checks that a freshly started session reports the status ‘pending’, but the value it reads is one that a background worker is racing to change to ‘running’ at that very moment. The server hands the JSON encoder the same live session object the worker mutates, so whether the response says ‘pending’ or ‘running’ depends purely on which goroutine wins, and under load or with the race detector enabled the worker can win. The result is a test that usually passes and occasionally fails for reasons unrelated to any real product defect, which erodes trust in the suite and trains readers to dismiss red runs. It was observed failing once during unrelated work, then passed eleven consecutive full-package race runs afterwards, so it is genuinely rare rather than broken outright. The underlying product behaviour is not necessarily wrong — a caller may legitimately see either state — so the likely correct outcome is to make the assertion accept either value or to observe the status through a synchronised read instead. Done means the test no longer depends on goroutine scheduling, proven by repeated runs under the race detector, with the reproduction captured first.

## **HXC-158 — No test builds the harmony\_os or aurora\_os screens, so their widget-threading fixes rest on code review rather than captured runtime proof**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** helix\_code/applications/harmony\_os/gui\_thread\_race\_test.go + helix\_code/applications/aurora\_os/gui\_thread\_race\_test.go (commit 9afc3da2); RED on unwrapped-fyneui artifact: harmony exit 1 / 415 DATA RACE, aurora

exit 1 / 495 DATA RACE, reports landing on the covered sites (harmony main.go:1583,:1631; aurora main.go:1756,:678,:1803). GREEN at HEAD: harmony exit 0 / 0 races, aurora exit 0 / 0 races, independently re-run by the conductor (harmony 8.597s, aurora 12.387s). Harness self-validated golden-bad per 11.4.107(10): dropping the renderer UI lock on the FIXED artifact still provokes 1296 (harmony) / 902 (aurora) races, proving the harness is not blind. Suite wiring landed in 8f072866. **Severity:** Medium **Created-By:** Claude

The Harmony OS and Aurora OS front-ends were just corrected so that background workers no longer modify on-screen elements directly. That correction is, however, only partly backed by evidence. No test in either package ever constructs the application's screens, so the background workers that paint them are never started while the tests run, and the race detector therefore never gets the opportunity to observe the very code paths that were changed. Concretely, running the Aurora OS package under the race detector reported zero races both before and after the fix — not because the defect was absent, but because nothing under test reaches it; the corresponding Harmony OS failure that was reproduced and fixed came from a different component, the background system monitor, which tests do start. The practical risk is that a future edit could silently reintroduce a direct on-screen write from a background worker in either file and every test would stay green, exactly the situation the anti-bluff policy exists to prevent, since a green suite would then be certifying behaviour nobody has actually exercised. Closing this requires a test that builds the tabs and drives those workers — the chat worker, the provider-health poller and the dashboard and resource timers — with the race detector on, so the fix is proven by observation rather than by inspection, together with a deliberately-broken variant proving the new test genuinely fails when the direct write is put back. Acceptance: a test in each package starts the previously-unexercised background workers against real widgets, passes with zero races over repeated runs, and demonstrably fails when the dispatch helper is removed from any one of those sites. Until that exists, the fixes for those specific sites should be described as review-justified, not runtime-proven.



## **HXC-162 — Skill auto-trigger is dead in the interactive CLI: the dispatcher is built then thrown away**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc162\_hxc163\_skills\_wiring\_20260728T205054Z **Created-By:** Claude

When a user types a request in the HelixCode interactive command-line client, the feature that is supposed to notice the request matches a known skill and run it automatically never fires. The code that decides which skill applies is created correctly and then immediately discarded without ever being connected to anything, so it can never act on what the user types. To the user this looks like the skills feature simply does not exist in the CLI, even though the underlying machinery is present, correct and documented. This matters because skills are how the product is meant to turn a plain-language request into real work without the user learning any commands. Anyone reading the source would reasonably conclude the feature is implemented, which is why this went unnoticed. The fix is to connect the dispatcher to the place where user input is handled, and prove it with a test that types a matching request and observes the skill actually run. Reported at cmd/cli/main.go:1060 by SOURCE-layer inspection; still needs runtime confirmation per 11.4.108.

## **HXC-163 — None of the seven shared governance skills are registered, so none of them can ever run**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc162\_hxc163\_skills\_wiring\_20260728T205054Z **Created-By:** Claude

The project ships seven ready-made skills that come from the shared governance component and are meant to be available to every project that uses it. In practice none of the seven is registered with the running system, so not one of them can be triggered by a user or by an agent. The effect is that work the organisation already paid for and maintains sits unused, and each project quietly re-solves problems these skills already solve. Separately, one entry in the skills folder points at a storage location that only exists on a Mac, so on this Linux machine it is a broken link that resolves to nothing. These two problems compound: the catalogue looks populated from the outside while being

empty in practice. Closing this means registering the seven skills through the supported mechanism, repairing or removing the broken link, and adding a check that fails if a shipped skill is present on disk but unreachable at runtime.

## **HXC-161 — Fixed.md H2-section-to-pipe-row parity gate fails on 58 items because the DB emits each item in only one representation**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** commit 34afaf11. Root cause was NOT the ticket's stated premise (that premise is refuted: items PK is (atm\_id,current\_location,representation) and 105 Fixed items carry BOTH representations). Real cause: the consumer changed — 8494380a superseded the pipe-table-only generate\_fixed\_summary.sh (now exit 2, refuses to run) with a representation-agnostic DB-derived renderFixedSummary. Of 56 failing ids, 0 lost a row; all 56 never had one. Reconciled per 11.4.120, not retired: CM-FIXED-H2-PIPE-ROW-PARITY -> CM-FIXED-SUMMARY-ITEM-VISIBILITY asserting every H2 closure, every ticket-id pipe row and every Obsolete item appears in Fixed\_Summary.md, bounded-matched so ATM-97 is not satisfied by ATM-970. STRENGTHENING PROOF against 0a4df699 (the real 156-item data-loss commit): retired invariant flagged 0 summary-visibility failures, reconciled invariant flags 246. Mutation proof both directions: GREEN unmutated exit 0 (153 sections/99 rows/1 Obsolete); delete section-only HXC-157 -> exit 1; delete pipe-row HXC-044 -> exit 1; empty summary -> exit 1 (253 failures); Fixed.md with no closures -> exit 1 via anti-tautology tripwire refusing a vacuous PASS. **Created-By:** Claude

The guard at scripts/gates/fixed\_h2\_pipe\_row\_parity\_gate.sh checks that every closed item written in docs/Fixed.md as a detail section also appears as a row in that file's summary table. It currently reports 58 failures, and it was already failing before this batch — the same failures are reproducible against the last committed version of the file, so this is a standing red gate rather than something a recent change broke. The cause is a change in how the file is produced. The tracker is now generated from the workable-items database, and the database records for each item which single surface form it takes: some items are written as table rows and others as detail sections, never both. The guard was written earlier, when the table was maintained by hand and every item was expected to appear in both places, so it is now asserting a relationship the data model deliberately no longer provides. Two

outcomes are possible and they must not be guessed between: either the guard is stale and should be rewritten to assert what the generator actually guarantees, or the generator is genuinely dropping rows that ought to exist and the guard is correctly catching real data loss. Deciding which requires reading the exporter's representation handling and confirming against the database, so this is deliberately not resolved here — weakening or deleting a failing guard without that investigation is exactly the suppression the constitution forbids. Who benefits: anyone relying on this guard to catch items vanishing from the closed-items summary, which is the precise failure it was created to prevent after one item went missing twice over. Reproduce by running `bash scripts/gates/fixed_h2_pipe_row_parity_gate.sh` from the repository root and observing 58 failures naming sections such as HXC-013 and HXC-119. Acceptance: the correct branch is chosen with captured evidence, the guard passes for the right reason rather than by being softened, its paired mutation still makes it fail when the invariant is broken, and if the generator was at fault the affected items are restored.

## **HXC-160 — Governance manuals still name the superseded shell summary generators as the canonical tracker-summary generators**

**Status:** Completed (→ Fixed.md) **Type:** Task **Evidence:** docs/qa/hxc160\_superseded\_generator\_naming\_20260728T210702Z/verification.md **Created-By:** Claude

The project's agent manuals — CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md and CRUSH.md at the repository root, their copies under `helix_code/` and `github_pages_website/`, and the cascaded copies inside the formatters, plugins and models submodules — still tell every reader that `docs/Issues_Summary.md` and `docs/Fixed_Summary.md` are produced by the shell scripts `scripts/generate_issues_summary.sh` and `scripts/generate_fixed_summary.sh`. That has not been true since the workable-items SQLite database became the single source of truth: both summaries are now generated from that database by the constitution submodule's Go exporter, while the shell scripts derive them from the Markdown trackers instead, which silently drops every item recorded as a heading section rather than a table row. The practical harm is that an agent doing exactly what the manual says destroys tracked state: on 2026-07-29 a hand invocation rewrote the closed-item tally from 344 to 188, corrupted

every per-type count, and replaced the full per-item table with aggregate counts only, before the sync gate caught it. The scripts themselves have now been changed to refuse to run, so the immediate danger is closed, but the manuals still point at them, so every new agent is still given the wrong instruction and discovers the refusal only by tripping over it. Fixing this means updating the CONST-057 and section 11.4.91 anchor text in every carrier so it names the database exporter as the generator, which spans roughly twenty files including copies inside submodules that other repositories own, and must keep all five carriers in lockstep per section 11.4.157. Who benefits: any agent or engineer who reads a manual to learn how to regenerate a tracker summary, and everyone relying on those summaries to see true project state. Reproduce by opening any listed manual and searching for `generate_fixed_summary.sh` — the surrounding sentence presents it as the current generator with no mention of supersession. Acceptance: every carrier names the DB exporter, no carrier presents the shell scripts as current, the five root carriers stay in lockstep, and a check exists that would fail if a manual reintroduced the stale naming.

## **HXC-182 — A corrupted performance reference file is sitting uncommitted and any broad commit would make it permanent**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** Baseline restored and independently verified by the coordinator: `worktree sha256 a13ff412b10ac88d == HEAD sha256 a13ff412b10ac88d`, and `'git status --porcelain reports/latency/'` is empty. The corrupted working copy (73f8ef85..., holding this contended host's 901ns/1.156929ms figures instead of the committed 982ns/1.237002ms reference) is gone, so no broad `'git add'` can sweep a corrupted yardstick in. Independently corroborated by the HXC-165 round-2 reviewer, which additionally confirmed the fix's runtime signature by a stronger method than hash-equality: 3x consecutive isolated runs left the baseline's MTIME (1785273093) unchanged, proving the file is never opened for writing at all rather than merely rewritten with identical bytes. A §9.2 pre-op backup of the dirty bytes was verified to exist at `/tmp/hxc165_backup_1785270412/baseline.dirty.bak` hashing to 73f8ef85..., so nothing was discarded unrecoverably. Root cause is fixed separately under HXC-165 (commits b43e69d5 + fbcc034c). **Severity:** High  
**Created-By:** Claude

The file that records the reference performance figures new measurements are compared against has been overwritten by test runs and now holds this machine's numbers instead of the agreed ones. It has not been committed, so the correct values are still safely recorded in the project history and can be restored with one command. The danger is that many people and automated helpers work in this same copy of the project at once, and any of them running a command that commits everything at once would sweep the corrupted file in permanently. At that point the reference would silently become a busy shared machine's timings, every later comparison against it would be meaningless, and nobody would get a warning. The underlying cause is being fixed separately so future runs stop writing the file, but that fix does not clean up the damage already present. This needs the file restored now, and it needs to be visible to everyone working in this copy rather than mentioned only in a commit message where no one will see it.

## **HXC-165 — A stress test overwrites its own committed reference baseline every time it runs**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** helix\_agent commits b43e69d5 + fbcc034c + 5a7a819f (unpushed at closure). RED on the prefix artifact changed the tracked baseline sha256 73f8ef85 -> 88c8b1a9, proving the write was real. DESIGN DECISION, evidence-led: NO comparison consumer exists (the only references were the write site and the plan doc that authored it), so no comparison was invented — doing so would have baked in a bogus yardstick, since the committed figures are a single contended-host run and the author's own runs measured P99 82ms against them. The project's genuine baseline-comparison mechanism is elsewhere and untouched (tests/performance/baseline\_regression\_test.go over benchmarks/baselines/\*.txt). Git history per 11.4.124 shows the file is an accidentally-committed run artefact: five post-creation commits, all pure numbers churn, two via blanket Auto-commit, while its own frozen Date: 2026-03-16 header stayed fixed as contents drifted months later. It was NOT deleted — that is its own commit and it is now harmless. RUNTIME SIGNATURE (11.4.108): against the restored baseline, content AND git status AND mtime all unchanged across a run (a13ff412, status empty, mtime 1785273093), stable across 4 consecutive runs — the file is not even opened for writing, which is strictly stronger than hash-equality. git check-ignore -v resolves the new run-output path

to .gitignore:101. Six paired 1.1 mutations each proven to FAIL then restored. Two independent guards: a package-level TestMain hashing around the whole run (guards the DEFECT SITE, not merely the new helper) plus an isolation test that self-validates its own hash oracle on a temp copy. Independent review: GO on correctness, safety and anti-bluff, zero blocking findings, both Important findings landed. Baseline restored 73f8ef85 -> a13ff412 with the corrupted bytes preserved out-of-tree per 9.2. HONEST LIMITS recorded: the probe-mutation evidence rests on the author's capture alone (the reviewer's constraints forbade writing the tracked baseline); a first commit landed mid-review contrary to 11.4.142, recorded in the amended message rather than hidden; and the guard's blind spots (a byte-identical rewrite, or a writer in the sibling tests/stress/verifier package) are documented boundaries, not silent gaps. **Created-By:** Claude

One of the performance stress tests writes its results directly over a reference file that is stored in version control, on every single run. The reference file is supposed to be the fixed yardstick that new measurements are compared against, so overwriting it means the yardstick silently becomes whatever the machine happened to measure most recently. That destroys the comparison the file exists to support, and on a busy shared machine the numbers are not even representative. It also leaves the project permanently showing unsaved changes, which trains everyone to ignore that signal and makes accidental commits of the corrupted baseline likely. A reviewer had to actively argue against committing it during unrelated work. The fix is to write run results to a location that is not version controlled and treat the committed baseline as read-only.

## **HXC-184 — Running a command that leaves a background process behind now hangs forever and permanently consumes a slot**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Severity:** Critical **Created-By:** Claude

A recent change made the command-runner wait for its output pipes to close rather than for the command itself to finish. Those are normally the same moment, but not when the command leaves something running in the background: the background process holds the pipes open, so the wait never ended. The request hung indefinitely and permanently consumed one of a small fixed number of execution slots, so ten such commands stopped the facility accepting any work at all.

FIXED IN TWO ROUNDS. Round one added a bounded grace period after the command is reaped, then stopped reading and released everything. An independent review accepted the mechanism but found it had moved the boundary somewhere unmeasured: an adversarial probe using a perfectly healthy command with no background process at all, consumed slowly by a caller following the documented pattern, delivered 490 of 5000 lines where the previous code delivered all 5000. The original reasoning held only when the consumer kept pace, not when the operating system's own buffer did. Round two replaced the fixed grace with a progress-aware one: the timer restarts whenever a line was actually handed to a consumer, so only a window in which nothing at all was delivered counts as stuck. The author then tested an objection the review had not raised — whether that could revive the original hang — and showed it cannot, because delivery requires a consumer that is still listening; an abandoned one stops consuming within two windows by construction. The one case that remains genuinely unbounded is a caller still actively reading output from a background process that keeps producing, which is documented rather than hidden, because cutting it off would discard output someone is reading. Three smaller findings were also closed, one of them by deriving the truncation flag from a recorded error rather than from which timing branch ran, so no scheduling coincidence can set it falsely.

## **HXC-174 — The terminal interface reads live QA session values while another thread is changing them**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc174\_qa\_dashboard\_session\_race\_20260728T221246Z/README.md  
**Severity:** High **Created-By:** Claude

The text-based terminal interface shows a page listing quality-assurance sessions and their progress. It reads those values straight from the shared session objects at the moment it draws the screen, while a background worker is simultaneously updating the very same values. Nothing coordinates the two, so the screen can read a half-updated state. The sharpest case is the elapsed-time display: the code checks that a finish time has been recorded and then immediately uses it, but between those two steps the background worker may still be finishing the write, so it can use a value that is not fully there yet. In practice this shows up as wrong or briefly nonsensical timings, and on

an unlucky run it can crash the interface. This is not new and was not introduced by recent work, but it is now the only place in the project that ignores the locking rule the rest of the code follows and that the session code explicitly documents. The fix is to take a consistent copy of each session while holding the lock, then draw the screen from that copy.

## **HXC-176 — The shared skill installer writes machine-specific paths that break on any other computer**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc176\_portable\_skill\_symlink\_20260729T031600Z/README.md  
**Severity:** Medium **Created-By:** Claude

The shared governance component installs each skill by creating a shortcut that points at wherever that component happened to sit on the machine doing the install. Because it records the full machine-specific location rather than a location relative to the project, the shortcut only works on the computer that created it. One such shortcut was already broken in this project: it had been created on a Mac pointing at an external drive, so on this Linux machine it pointed at nothing at all, and had never once worked here. That single broken shortcut has been repaired, but the installer that produced it is unchanged, so re-running it recreates the same problem. A guard now catches the broken result, which is a safety net rather than a cure. The proper fix belongs in the shared component: record the location relative to the project so the shortcut works on every machine that checks the project out.

## **HXC-177 — A shipped skill uses a placeholder the system cannot fill, so raw placeholder text is sent to the AI**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc177\_skill\_placeholder\_guard\_20260729T031600Z/README.md  
**Severity:** Medium **Created-By:** Claude

One of the skills shipped with the product writes a fill-in-the-blank placeholder in a form the system does not recognise. The system only substitutes placeholders written in one specific form, so this one is



never replaced. The consequence is that the literal placeholder text is sent onward to the AI model as part of the request, where it reads as meaningless noise rather than as the value it was supposed to stand for. The user sees a lower-quality answer with no indication of why. This was found while working on unrelated skill wiring and is not caused by that work. The fix is to correct the placeholder to the recognised form, and to add a check that refuses to ship a skill containing a placeholder the system cannot resolve, so the same mistake cannot recur silently.

## **HXC-170 — Small, low-risk library updates that close eleven of the agent's eighteen confirmed weak spots**

**Status:** Completed (→ Fixed.md) **Type:** Task **Severity:** High

Eight of the agent's supporting libraries have newer releases that fix known weaknesses, and every one of these updates is a minor or bug-fix-level release within the same major version, meaning the library authors are committing to not breaking how the code is called. Together they close eleven of the eighteen known weaknesses in the agent's own compiled program, including four we have confirmed the program genuinely reaches during normal operation. This matters because it is the single highest-value change available: a small, reviewable set of version bumps that measurably shrinks our exposure without redesigning anything. One of the eight also fixes a text-handling flaw that our automated safety checks route untrusted user input through, and which no advisory service had flagged at all. Everyone deploying or depending on the agent benefits, and the reviewer benefits from a change small enough to actually read. The expected outcome is the updates applied, the project rebuilt, its tests run, and a fresh scan confirming those eleven weaknesses no longer appear, with the before-and-after scan output kept as evidence. Two of the eight are minor rather than bug-fix releases and deserve a slightly closer look at build and test results.

## HXC-187 — Two different pieces of code claim the same identity, so which one gets used depends on where the build starts

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** commit 09a086a6, go.mod only (1 file, +1/-1). Root module renamed dev.helix.code -> dev.helix.code/meta; ZERO import updates were needed because the root module genuinely had no importers. TWO-METHOD verification, and method 1 alone would have produced two FALSE POSITIVES: a git ls-files sweep hit an ASCII-art string literal ('dev.helix.code v1.0.0') and prose in doc comments, neither an import. Method 2 (filesystem, sees submodules + untracked) found all 947 quoted imports under helix\_code/ and zero elsewhere; sweep validity proven rather than assumed (9944 .go files under submodules/, 1909 under cli\_agents/ — it had real content to see). Reversal conditions each cleared: no go.mod requires/replaces the root; scripts/audit\_const046 is an independent module not a consumer; and D-7's 'no go.work exists' was imprecise — two do exist but neither names dev.helix.code. SEMANTIC PROOF the collision is gone: dev.helix.code/internal/theme now resolves ONLY from the inner module; from the root it reports 'no required module provides package'. Build/vet for the root module's own 5 packages: 0/0. The agent deliberately did NOT delete the root stub (removal is its own 11.4.124 decision) and did NOT touch skill\_registry/U-12 (it cannot be fixed in isolation — reconciling it repairs helix\_agent's replace but BREAKS helix\_llm's currently-correct one). Gate deliberately left UNREGISTERED and unweakened: its premise 'exactly one duplicate' proved stale — 118 live modules, 8 duplicate groups, 7 surviving the rename (1 ours: a self-nested tests/e2e/orchestrator copy from refactor cc339fc0; 6 in another repo's nested submodule content). Registering a failing gate would break the sweep for all live agents, and narrowing its exclusions would destroy its ability to detect an accidentally-nested duplicate in our own tree (assertion-weakening under 11.4.120). Both refusals correct. **Severity:** High **Created-By:** Claude

The project declares the same module identity in two places: once at the top level and once for the real application inside it. Because of that, one particular internal name refers to two completely different pieces of code — a five-line placeholder at the top level, and the real fifty-kilobyte subsystem inside. They share no files at all. Which one any given piece of code actually receives depends entirely on which directory the build was started from, which is not something a developer would ever expect or notice. Nothing is broken today only because the top-level placeholder has no users, but that is luck rather

than design, and the first time someone adds one the behaviour will differ between two builds of the same source with no error to explain it. The fix is to give the top-level module a distinct identity so the collision cannot occur. Removing the placeholder is a separate decision and is deliberately not bundled with this.

## **HXC-173 — An `aurora_os` screen-refresh test reports failure when the machine is merely busy**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** docs/qa/hxc173\_gui\_race\_overlap\_floor\_20260728T224308Z/README.md  
**Severity:** Medium **Created-By:** Claude

One of the tests that checks the `aurora_os` dashboard updates its on-screen values safely will report a failure when the machine running it is under heavy load, and passes cleanly when run on its own. The underlying feature is fine — this was confirmed by running the same test in isolation, where it passes — so the failure is an artefact of timing, not a defect in the product. This still matters, because a test that fails for reasons unrelated to the code teaches everyone to ignore its result, and the next time it fails for a real reason nobody will believe it. It is also a direct problem for release preparation, where the full test run happens while many other jobs are active. The fix is to make the check wait for the condition it cares about rather than assuming the screen refreshes within a fixed window, so that a slow machine produces a slow pass instead of a false failure.

## **HXC-183 — One AI provider still leaks a background worker when the caller stops listening**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Severity:** High **Created-By:** Claude

When the product streams a reply from an AI provider, a background worker hands results back to the caller piece by piece. If the caller gives up early — the user cancels, or a timeout fires — that worker must notice and stop. Six providers were recently fixed to do exactly that, and each got a test proving it. One provider was missed: it still hands results back without checking whether anyone is still listening,

so the worker blocks forever and both it and its open network connection are never released. Every cancelled request against that provider leaks a little more of the machine's capacity until the service is restarted. The fix is mechanical, because the correct pattern already exists in its six siblings and there is a shared test harness it can be added to. This was found by an independent review noticing that the round of fixes stopped one directory short of covering every affected provider.

## **HXC-185 — The fix for modern network addresses was applied in two places but not to fourteen others that need it**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Severity:** High **Created-By:** Claude

Modern internet addresses contain colons, so when one is combined with a port number it has to be wrapped in brackets or the result is unreadable to the software that parses it. Two places were recently corrected to do this properly. A review found fourteen more places that build addresses the same unsafe way and can receive such an address from configuration or from service discovery. The clearest illustration sits inside a single file that checks service health two different ways: the connection check carefully adds the brackets and works, while the web-request check a few lines away does not and fails outright for the same service. The others cover the local model server, the memory service, the notification service, worker connections and the test infrastructure. Each is a one-line change using the standard helper that handles this correctly. Left alone, any deployment using these newer addresses sees confusing partial failures where some checks pass and others fail against the same healthy service.

## **HXC-167 — The tool that strips secrets from recorded test transcripts can fail without anyone noticing**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Severity:** High **Created-By:** Claude

Before a recorded test conversation is saved into the project, a cleanup step is supposed to remove any secret values it contains. That cleanup has two weaknesses. First, it can only remove secrets the test script itself already knew about, so anything the remote service sends back to us — a session cookie, a returned access token, a pre-signed link — is unknown to the cleanup and is saved exactly as received. Nothing of that kind has actually leaked so far, because none of the services exercised happen to return one, but the protection simply is not there. Second, and worse, the cleanup can fail silently: the secret is fed into a text-replacement tool in a way that treats certain punctuation as instructions, so a secret containing one of those characters makes the step fail, and the failure is deliberately discarded rather than reported. The result is a secret that stays in the file with no warning to anyone. The fix is to make the replacement treat secrets as plain text, make any failure loud rather than silent, and add a scan of the finished file that also catches secrets the remote side sent us.

## **HXC-203 — Asking for local model status corrupts that status when two requests arrive at once**

**Status:** Fixed (→ Fixed.md) **Type:** Bug **Evidence:** commits b2215793 (fix + guard), be3e6605 (captures), 1c2e7d07 (round-2 review record); unpushed. RED on the pre-fix artifact reproduced the exact signature — write/write at local\_llm\_manager.go:613, one arm via GetRunningProviders:621 — captured on a QUIET host (load 2.04 / 64 cores, no competing builds), so the verdict is not contention noise. Defect 2 independently confirmed: NO mutex of any kind existed on the type at HEAD. GREEN: same test exit 0; both suites -race -count=2 exit 0; guard default polarity -count=3 exit 0. LOAD-BEARING PROOF: the NEW guard run against pre-fix source (restored from HEAD, applied and reverted in one command under a restore trap, tree verified byte-identical md5 79f3ee7f) exits 1 with BOTH halves firing — and the second half is INDEPENDENT of the race detector, because a caller's delete() on the returned map genuinely deregistered a provider from the manager, proving the live-internal-map defect on its own terms. RED\_MODE=1 self-validation exits 1, so the guard is not blind. LOCKING DESIGN under the network-I/O constraint I set: the loop is NOT wrapped. Three phases — snapshot probe targets under RLock with the port copied BY VALUE so the probe dereferences nothing shared; probe with NO lock held; write back under Lock with a compare-and-set. The CAS is required PRECISELY BECAUSE the lock is

dropped: a verdict formed against a provider that StopProvider has since stopped must not clobber newer state. A new `isEndpointHealthy(ctx, port)` ensures the probe path touches no shared memory. Reporting is separated from refreshing and callers receive a deep copy (Process nil'd, maps and slices copied), verified against every caller. SIBLING SWEEP within the type: fixing only :613 would have left the real-world case (status display during a provider start) racing, so Status/Process writes in StartProvider, StopProvider and installProvider now go through accessors holding the lock ONLY around the field write — never across exec, Wait, or the 60s readiness poll. StopProvider now claims the process handle under lock, which additionally prevents two concurrent stops double-Wait'ing the same child. Deadlock-safety is mechanical: every accessor takes and releases mu itself, so no critical section calls another locking method. Independent review returned GO with zero blocking findings and zero warnings, having traced the whole lock sequence itself and separately verified the read placement was unchanged from pre-delta so no behaviour change was smuggled in. HONEST LIMIT, recorded in the evidence README rather than implied closed: no provider was ever 'running' in these runs (the test manager skips installs and starts nothing), so refresh phases 2-3 INCLUDING THE CAS rest on reasoning and review, not executed evidence; closing that needs a stubbed provider endpoint. Residuals filed as HXC-205 (the same defect class in AutoLLMManager — eighth sibling instance this cycle), HXC-206 (ABA window in the CAS), HXC-207 (shared Environment map aliasing, latent not live). **Severity:** Critical **Created-By:** Claude

The component that tracks locally-running AI model servers offers a method to report their status. Despite being named as a query, it also writes to the shared records it reports on: it overwrites each provider's health and stamps a last-checked time on every call. Nothing coordinates those writes, and the component has no locking of any kind, so two callers asking at the same moment write over each other. This was caught by the race detector during a full test run, which does not produce false alarms, so it is a real defect rather than a timing coincidence. Two further problems compound it: the method hands back the internal collection itself instead of a copy, so every caller receives live references it can also modify; and the health check that decides the value performs network calls while holding no lock, so slow or hanging providers widen the window. The consequence for a user is a status display that reports a healthy provider as unhealthy or the reverse, and in the worst case corrupted records shared across the

whole component. The fix is to protect the shared records with a lock, return a copy rather than the live collection, and separate the act of refreshing health from the act of reporting it.

---

1. A-Za-z0-9\_-<

2. A-Za-z0-9\_-<

3. a-zA-Z0-9\_-<