

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Operator-blocked	1
Bug	Queued	33
Task	Operator-blocked	1
Task	Queued	11
TOTAL		46

Items

ATM ID	Type	Status	Severity	Description
				Reported-Via: §11.4.202 reporting directive to Operator What (the report, verbatim): Feature the HelixSkills System (git@github.com:HelixD and HelixAgent, covering ALL of its power-fea OPERATOR REQUIREMENTS (2026-07-29): “Full HelixAgent the HelixSkills System for all work git@github.com:HelixDevelopment/skills.git , v power-features fully without nothing leftout! F into phases, tasks and subtasks with exhaustiv prepared! Full coverage with ll supported test (suites) is mandatory! Extend and update all ex guides, FAQs, create stunning illustrations, gra all materials! Any gaps, shortcomings, weak sp inconsistencies MUST BE detected in advance o process and properly tackled with comprehen risk-free and safe! Track all progress through t and remebering mechanisms we have! Use Git
HXC-159	Task	Queued	High	

ATM ID	Type	Status	Severity	Description
				with bridge to Superpowers used for the implementation of the development mechanism!" PRE-VERIFIED ENVIRONMENT (not assumed — §11.4.6): - GitHub SpecKit IS initialized and NOT yet initialized in this repo (no .specify/ and no .specify/ is therefore an explicit early task, not an assumption). - skills.git IS reachable via SSH. git ls-remote -p https://github.com/catalog-docs, feature/deep-research, feature/teaching, feature/branch/tag inventory MUST be re-derived at repository update or partial listing. - The repository is NOT currently initialized and contains no HelixDevelopment/skills entry; the repository is codex-skills -> vasic-digital/caf-codex-skills.git, not be confused with it. - Both consumers exist (module dev.helix.code) and the submodule subagent (dev.helix.agent). MANDATORY DELIVERY CONSTRUCTION for ALL phases of the incorporation, bridged to the subagent executed subagent-driven (§11.4.70 / §11.4.20). - tasks -> subtasks, with an exhaustive level of detail where the design is load-bearing). - Complete the e2e, full-automation, Challenges (vasic-digital/Challenges) (HelixDevelopment/helix_qa) with autonomous concurrency/atomicity, race/deadlock, memory management, captured evidence (§11.4.5 / §11.4.69); every gap in Documentation: extend AND update every existing and only new ones. Produce illustrations, graphs and materials. Four-format export discipline per §11.4.212. the main README per §11.4.212. - Risk work is a note: every gap, shortcoming, weak spot, danger detected DURING planning and each one paired with a free and safe solution (§11.4.102 systematic-deadlock before rewrite (§11.4.74 / §11.4.28): survey the catalogues on GitHub AND GitLab first; extend and duplicating, keeping them project-agnostic and incorporated as a properly-laid-out dependency level or submodules/, never a nested own-org and HelixAgent must BOTH receive the full feature, the other is an incompleteness to be tracked, not through the continuation docs, the memory/knowledge the workable-items SSoT (§12.10 / §11.4.131 / §11.4.131). 1. A SpecKit specification exists and is initialized as a feature, with an explicit feature inventory provided in the repository (an enumerated list, not a sample — implementation plan exists at the stated depth of implementation mitigation. 3. Every planned capability is wired

ATM ID	Type	Status	Severity	Description
				consumers, proven by §11.4.108 runtime signals compiled. 4. The full §11.4.169 test-type matrix and HelixQA banks included. 5. All documentation produced diagrams/illustrations incorporated, Independent code review reaches a zero-finding. No capability is claimed without runtime evidence every closure (§11.4 / §11.4.1 / §11.4.123). §11.4. PROGRAM (scheduled, not yet executed): (a) Do scientific papers / open-source projects+codebases to incorporate this feature into the project (§11.4.102) of ALL data obtained to ensure zone / inconsistency / imperfection, and design one found. (c) The design MUST ALWAYS be innovative – never less. (d) MANDATORY exhaustive documents, an implementation plan down to diagrams / schemes / graphs, illustrations, SQL relevant material (§11.4.65 / §11.4.73). (e) Invest + HelixDevelopment submodules/components freely where genuinely needed while keeping reusable (§11.4.28 / §11.4.74 / §11.4.177). (f) Plan one: every supported test type, the Challenges (§11.4.27 / §11.4.169). (g) Divide the work into nano-detailed enough to drive integration / implementation directly. (h) Plan explicitly for enterprise scalability the feature has a UI/UX surface, produce full wireframe PSD / PDF, or whatever the project mandates) and §11.4.190). (j) Plan full CodeGraph integration + synchronisation (§11.4.78 / §11.4.79 / §11.4.80). items (every detail + reference + attachment) source-of-truth (§11.4.93 / §11.4.95) + every design-related project doc/component + every connection (§11.4.148 D5) – honestly SKIPPING any absent (credentials_absent / tracker_client_absent, §11.4.148 D5) Research-doc destination (to be populated when fully_incorporate_the_helixskills_system_into_Execution model (§11.4.213 clauses 1-2): this item multi-day research/planning/documentation work project's standing autonomous loop (§11.4.87 / when it claims this item, and MUST be driven to explicitly evidence-backed CLOSED terminal state be left sitting un-wired in the backlog. Affected research/planning – destination: docs/research fully_incorporate_the_helixskills_system_into_

ATM ID	Type	Status	Severity	Description
				Reproduction / context: UNKNOWN: not state established before any fix (§11.4.102 / §11.4.14 research-doc tree exists under docs/research/ fully_incorporate_the_helixskills_system_into every enumerated weak-spot/gap/danger-zone implementation plan down to lines-of-code exist are created, and the whole effort is validated p ADDENDUM — BINDING OPERATOR REQUIRE === VERBATIM: “Make sure that SpecKit and S and with all extensons we have created deriviv not optional colour on the SpecKit mandate — work MUST compose with rather than duplica BRIDGE EXTENSION (verified to exist, 2026-07- speckit_superpowers/implementation/ — the “ Superpowers Bridge”. Eight documents, each i IMPLEMENTATION_PLAN, NANO_TASK_ENGIN TDD_INTEGRATION, CONSTITUTION_INTEGRA SEVEN-LAYER architecture: 1 Developer Work SuperSpec bridge (orchestration — github.com (discipline enforcement — speckit-community MCP (execution) 6 Helix LLM (inference — gith Distributed host cluster (llama.cpp RPC) CONST MUST BE IN SCOPE (inherited BY REFERENCE , diverges silently): constitution/skills/ (7) action reporting-workable-items, scheduled-work-qu constitution/mcp/ (2) media-validator-mcp.json plugins/ (2) helix, scheduled-work constitution system), subagent_tiering.yaml constitution/sc credential_scan_lib.sh, guard-branch-consisten guard-track-branch-label.sh, guard-work-track LOCALLY — do not re-create: .claude/skills/ 10 clarify, constitution, converge, implement, plan workflows/speckit/workflow.yml, integrations/ templates/, memory/ ADDED ACCEPTANCE CRI extension above has an explicit disposition — reason. Silence is not an answer (§11.4.118). 9. WIRED vs DESIGNED-ONLY with a captured pr evidence a layer is installed — that is precisely and reporting a documented layer as an availa namespace collision risk is addressed with a d skills system, while §11.4.164’s post_update_ho into .claude/skills/. Two systems registering int and clash handling, not discovery at runtime. .

ATM ID	Type	Status	Severity	Description
				non-own-org dependencies (§11.4.74 vendor p supply-chain exposure must be assessed, not a them. The script that brings up the supporting service directly, rather than going through the shared maintains for exactly this purpose. That comp the same way, gets the same health checking, a repeatedly. Bypassing it means this project car and has already caused one outage of its own. service startups still hide their errors, so a fail was flagged honestly in a commit message, but obligation anyone will act on, which is why it i routing the startup through the shared compo hiding so a failed start is visible. When publishing the agent component, the ho outstanding security advisories against the lib 204. That figure turns out to be unreliable in th drifts upward on its own as new advisories are 205 the next day, and 210 on a live re-query on be “the” count. It over-states the work by abou counted once for every dependency file it appe distinct advisories, and the five criticals collap two places this code is published; the other cop was never switched on there. The recorded br critical, 88 high, 99 moderate and 18 low. A ful completed, and it changes the picture substant the affected code can actually be reached in so them currently qualify. The Go vulnerability so actually release, reports our code is affected by against a container library are reachable only in the codebase refers to and that never enters remaining alerts, including three of the four c third-party project whose container definition cannot start today. Most of the rest belong to a builds. Exactly one alert touches a component defence that advisory asks for has already bee alert stays open only because the dependency related piece of housekeeping was resolved wh files of downloaded third-party library code, in behind one critical advisory, had been commit removed from the tracked files, though they re benefit of finishing this work is that the alarm
HXC-164	Task	Queued	Medium	
HXC-166	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				the product, and we regain the ability to say he carries a reachable known vulnerability. This is an acceptance criterion has not been done: the de findings have not yet been upgraded or removed, retire the copied-in project that carries most of the project. The password used to reach the project database is in the setup and container files, and those files are present in all accounts. Anyone who reads the project can read the password, not introduced by recent work; the project's own password is it as something that must be replaced, and that replacement depends on whether the same password is used elsewhere. It has not been established either way and should be. The password value is already public, changing the files alone is not enough; it be changed everywhere it is used, and the setup files should be a configuration source rather than containing it. The project cannot honestly claim its credentials are protected. A complete copy of somebody else's web application is in it alone accounts for one hundred and thirty-nine warnings, dependency security warnings, including three warnings that be built or run: the startup configuration point is in that folder, and its supporting packages have not been built two opposite ways at once. Because it is not built the way we actually run today, which is why they rank it as a warning references it, the setup is broken and misleading. The recipe file, all one hundred and thirty-nine warnings, to whoever does it. Anyone reading our security report is a distortion, since it currently buries the small number of warnings. The expected outcome is an explicit decision, not to properly maintain and build it, or remove it and replace it with a reference to the upstream hosted service. Before any approval, the decision must be put to the operator. A small companion service we wrote ourselves is part of the standard deployment, and it carries a warning. Its supporting packages, two of them rated high. Under the warnings, these have no mitigating evidence in the project, unshipped, because it demonstrably ships, and that is whether the affected functionality is ever exercised. It has place in the whole review where something was not mentioned, which makes it far more deserving of attention than it is. Three is not really an upgrade problem at all: it is not switched off unless explicitly enabled, so updating it without protection on would leave the exposure in place.
HXC-168	Bug	Operator-blocked	High	
HXC-171	Bug	Queued	Medium	
HXC-172	Task	Operator-blocked	High	

ATM ID	Type	Status	Severity	Description
				deployment running this service benefit. The e assessment of the three warnings, the two stra protective setting explicitly verified as enabled and the service rebuilt and confirmed working The cache that remembers model information where its clock was never configured. That fal the only way to build the cache always configu unexercised. On its own that is harmless. The n if someone later edits it: a plausible-looking ch current time would make every cached entry a in the past that the freshness check reads as a expiry limit. Every entry would then be consid keep serving outdated model information inde The choice is to either add a small test that pin branch entirely and document that there is on
HXC-175	Bug	Queued	Low	When a caller presents an API key, the server c ordinary text comparison. That kind of compa that does not match, so a key sharing the first f reject than one that differs immediately. An att differences across many attempts can recover guessing it outright. Network timing noise mak is low rather than urgent, but it is a genuine w cheap: use the comparison function designed f amount of time regardless of where the differe introduced by recent work.
HXC-178	Bug	Queued	Low	The check that decides which websites may op deliberately allows requests that carry no orig behaviour, because only browsers send that in legitimately omit it. It is safe today only becaus away anyone without valid credentials. The tw nothing records or enforces that coupling: if so login check while working on something unrel information would become completely ungate work here is to add a check that specifically co the connection upgrade, so the dependency is
HXC-179	Task	Queued	Medium	Three commits in the current release range do referred to a piece of data that was not added those three points gets a build error. The tip of corrected shortly afterwards. This cannot be r rewriting published history, which the project is for anyone hunting a regression by stepping three points that fail to build for a reason unre
HXC-180	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				may wrongly conclude the fault lies there. The affected points and their shared cause somewhat so the failure is recognised immediately as known from scratch. HXC-160 corrected the 21 stale references to scripts/generate_fixed_summary.sh across the added the CM-SUPERSEDED-GENERATOR-NAMING sets were deliberately left untouched and are the Constitution.md still carries 14 references across universal rules that projects other than HelixCode records such as the Phase 5 Go-binary-shim plan is a governance amendment that CONST-049 reconfigured upstream, then re-cascaded and postponed session that closed HXC-160, and a canon amendment drift than the stale wording, so the canon was 45 owned submodules carry cascaded restatement correction applied in the parent repo is deliberate because the fix names a HelixCode-specific mechanism wording into decoupled, project-agnostic submodules them and violate CONST-051(B). The correct set neutral canonical wording that names the word and push it in the constitution submodule per fleet and bump pointers. Who benefits: every of the anchor text, and any agent reading a submodule summaries are produced. Acceptance: canon rewording, the change is pushed to every constitution are re-cascaded, submodule pointers are bumped SUPERSEDED-GENERATOR-NAMING both pass The main handbook that tells contributors and organised lists several folders at the top level to when only one is present, and refers to a top-level Anyone following it — a new contributor, or an orientation — will look for things that are not broken or that they have checked out the wrong of the document, which is otherwise the author. The correction is small and purely descriptive. section to match, then keep the two in step as to One of the external projects this codebase depends configuration still refers to it by its previous name hosting service silently forwards requests from forwarding is a courtesy, not a guarantee: it does project under the old name, and it is not honored
HXC-181	Task	Queued	Medium	
HXC-188	Task	Queued	Low	
HXC-189	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				dependency. If it stops working, fetching the project's name nobody recognises, which is unusually common for the project's current canonical name so the reference to it on a redirect continuing to exist. When a running command is cut short because of a timeout, it says it was not a timeout. The check that sets the deadline is the caller's own deadline, which usually has not been reached, so the command stopped the command. Investigation showed that the command was either, because the internal signal it would complete was cancelled, or cancellation and never a timeout, and the timeout was not a mechanism entirely. The effect is that anyone who uses the command fails and no indication that a time limit was reached. The bug is in the wrong place. A sibling code path that runs commands correctly by doing it inside the timer's own callback, but not a copy.
HXC-190	Bug	Queued	Low	
				When a command is cancelled or times out, the command is not cancelled, but anything it started, by signalling the whole command, works if the command was placed into its own sandbox. If the sandbox is enabled — the normal configuration — the sandbox is switched off, the grouping step is skipped, and the operating system to signal a group that does not exist. The error is discarded. The result is a cancelled command that is dormant, because the only configuration that could have been so no shipped path reaches it. It is worth fixing this, but it would appear only in an unusual configuration and is not a diagnose later.
HXC-191	Bug	Queued	Low	
				Output from a running command is read line by line. The limit of a line may be. When a line exceeds that limit the command is permanently for that stream, because the buffer is full and past it. Nothing else takes over the reading. The command is not until the operating system's own small buffer is full, and someone to read. So one unusually long line, which is not readable output produce quite naturally, can be read. The fix is to keep draining and discarding the remainder of the command, abandoning the stream, so the producing command is not abandoned.
HXC-192	Bug	Queued	Low	
				The handbook that describes this project's technology is one of its networking libraries. Both parts of the handbook are because a security fix required moving off the old handbook. It is deliberate; it is the handbook that is now wrong. It is a new contributor or an automated helper read the handbook. It will confidently state a version that has not been found. The same document has already been found described in the handbook.
HXC-196	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				additional inaccuracy makes the whole document recorded version to what is actually in use, and for this reason, so nobody later reverts it believing the version recorded. One test directory contains a duplicate of its own files appear twice in the project at two different times during a renaming exercise that moved a tree while leaving the old one. Nothing reads the nested copy, so anyone searching the project finds two results, one of which is real, and any tool that scans for duplicates blocks a useful check from being switched on: the tool blocks duplicates cannot be enabled while this one exists every run. Removing it requires first confirming the nested copy is genuinely the leftover and not the original.
HXC-197	Bug	Queued	Low	
				A detector was written to catch a specific kind of error: stuck trying to hand off a result nobody is collecting. It has one particular description of what the worker is doing: in form of the same stall — where the worker was waiting for rather than a single handoff — and the detector was demonstrated rather than theorised: a deliberate test case went completely undetected while the detector was running. What it does catch is the historical one, so the guard is set to ignore the newer structure while breaking its escape condition. Recognising both descriptions is a small change and a good one, a defect rather than half of it.
HXC-200	Bug	Queued	Low	
				Two long-running tests — one exercising heavy load, the other one exercising leader election among worker nodes — both report failure when the machine is under load. The failure is on its own. Measured: during a full test run both fail about once in ten times consecutively, one finishing in roughly a minute, the other about a twentieth. Their verdicts are therefore unreliable.
HXC-204	Bug	Queued	Medium	Whether the product works. This matters most when a test run performed before a release is by definition successful: these two will mark a healthy release as broken. The test goes through. A third test with the same problem was performed: an inconclusive run report as skipped rather than as failed. The same treatment fits here. Both should wait for a verdict rather than assuming a fixed time is enough.
				The newly-fixed health refresh deliberately refreshes the health provider over the network, so that a slow or hanging refresh. It then re-acquires the lock and only applies its changes if the changed underneath it. There is a narrow race condition: if started again entirely within one probe, the change is lost.
HXC-206	Bug	Queued	Low	

ATM ID	Type	Status	Severity	Description
				applies a verdict formed against the previous i briefly reported unhealthy when it is fine, whi dangerous direction — reporting a stopped pr Worth recording that this is cheap to close rath incremented on each state change is invisible t to any published structure. An earlier note des was wrong and has been retracted. When a model provider record is created, its e shared template by reference rather than copi points at the same underlying settings, includi Nothing writes to those settings today, and the no corruption occurs at present — this was ver record it is that the protection now added to th manager cannot protect a structure shared wi legitimate, properly-locked write to a provider manager's view as well, and the locking will lo when the record is created removes the hazard Four compose entries across two end-to-end te the Slack mock and the LLM-provider mock, b Any attempt to build those stacks therefore fai which means the containerised form of these t exercised. This was surfaced while fixing an u same two services: the fix could be proven at t against a running container, precisely because produced. That gap matters beyond these two process leaves the deployed-artifact layer unpr failure the four-layer verification discipline ex Dockerfiles, or to remove the build entries if th run containerised, and then to prove the choic A blanket rule excluding log files from version evidence, because evidence transcripts are wr corrected so this cannot happen again, and the have been recovered and committed. Thirty-fi the runs exist only on the machine that produ an evidence folder that appears empty, which never run. Recovering them is not automatic, l seventy-eight megabytes and committing those deserves a deliberate decision rather than a re keep what genuinely substantiates a closed ite ones rather than storing them whole, and reco why, so the gap is visible rather than silent.
HXC-207	Bug	Queued	Low	
HXC-208	Bug	Queued	Low	
HXC-216	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
HXC-219	Task	Queued	Medium	A recent repair changed how the release gate of evidence really belongs to the change it claims to commit on a labelled line rather than merely having somewhere in the text. That new requirement have to follow it: the guide that authors read before the new labelled field, so a contributor following evidence the gate will reject without understanding describing the gate itself has printable and we page was last edited, so all three versions disagree describe. Fixing this means updating the author's citation field and regenerating the exported code who benefit are contributors recording evidence have no accurate written description of the rule can follow the written guide and produce evidence attempt. When we mark a piece of work finished, we reworks. Nobody ever checked that those pointers different kinds of rot accumulated silently and report. First, some pointers hold a whole paragraph there is no way for a machine to follow them. Second never created because the person doing the work means the proof exists but nobody can find it for ordinary progress updates rather than to the author notes look like formal proof. The effect is that a backed in every summary and dashboard which way anyone finds out is by manually hunting for make the completion pointer a checked field: if it exists, it must be attached only to real completion must be refused at the moment someone tries later. This protects everyone who relies on our claim of proof that cannot be produced on demand.
HXC-224	Bug	Queued	High	A single bundled helper component inside the roughly 208 security advisories reported against the total — and it entered the codebase without established what it does, whether we wrote it or actually uses it, or whether removing it would answered the headline advisory count is misleading carries enormous risk, when in reality the great not even be part of what we deliver. The work purpose, whether it is reachable from anything, decision about it — keep and maintain, upgrade record which was chosen and why. The decision
HXC-225	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				advisories, because upgrading dependencies in effort spent on nothing. The agent project is mirrored to two hosting and was only ever switched on for one of them. Qu hundred and ten open alerts while the other r one the main project is configured to pull from of the component we actually consume is show hundred findings sit unseen on the sibling cop the two mirrors hold the same commits. What watched. This is dangerous precisely because t findings is indistinguishable from a report of z the second case is what we have. The work is t from, confirm both then report the same figur the two so a future divergence is noticed rather statement about this component's vulnerability from. A design document committed forty-eight days appears to be a genuine access key for an exte placeholder. The section is headed as being co and described as already configured, and the v placeholder would have: it is fifty-one character characters, and contains no example or to-be-f the main branch and has therefore been publi day it was written. Anyone with read access to fork made at any point in those forty-eight day now would remove it from the current revisio history is forbidden, and rewriting would in ar The only action that actually withdraws the ke replacement, which requires access to the pro within the repository. Until that happens the k revoked, the document should be edited to ref of quoting its contents, and a check added that this shape. The multi-provider routing gateway starts in it release mode, and says so in its own startup lo this should be switched for production use. De prints every registered route at startup, produ request, and in this framework is documented extra output and the internal detail it exposes launched by the system supervisor as a long-ru the production case the warning is about. Noth failing, which is why this has gone unnoticed, .
HXC-226	Bug	Queued	Medium	
HXC-227	Bug	Queued	—	
HXC-229	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
HXC-231	Bug	Queued	—	logging at development verbosity and revealing the service unit is the natural release mode when the service runs under the production unit ever starts in debug mode again. The HelixLLM gateway downloads large AI models before it starts accepting network connections. During this time, the service as running and healthy, but every so often on 2026-08-06 this window lasted thirty-seven minutes before actually began listening. Anyone checking the status of everything is fine while the service is in fact unresponsive. It depends on the gateway will fail for the whole service because the download depends on an external service. If one of the two downloads timed out after roughly 30 minutes that site were unreachable the gateway might be stuck accepting connections first and download mode later, in an honest not-ready state while the downloads are in progress. HelixAgent expects two tables to exist in its database: agent_instances, but neither of them has ever existed anywhere in the project that creates them. As a result, it locks once every minute, the database refuses to accept the lock and an error is written to the log. Verified on 2026-08-06 directly: neither table exists in any of them, and the error occurs every minute period plus a further hundred and thirty seconds, otherwise runs normally, so nothing visibly broken. It buries genuine problems in the log and means that the service to support is silently doing nothing at all. Fixing this by creating these tables were intended to contain to write to the database guessed. When the HelixLLM gateway starts it tries to launch a program called llama-server and fails because the llama-server machine confirms the situation is worse than a llama-server exists anywhere on the host, so the service is pointing the service at the right directory. The service continues and carries on without it, so nothing appears broken. It is absent for the entire life of the service. This is a problem I had just spent thirty-seven minutes downloading llama-server could serve it; the download succeeded and the service live on 2026-08-06 during a full platform boot. The fix is installing the program, then making sure the service is aware of whether the gateway should be downloading llama-server would consume them is not present.
HXC-232	Bug	Queued	—	
HXC-233	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
HXC-234	Bug	Queued	—	On startup HelixAgent tries to build and launch containers known as MCP servers, using the container engine. If the containers inside the container fail during their dependency installation, the operation is abandoned. HelixAgent records a list of these servers and then continues running normally. Nothing obviously breaks. The real effect is that the servers we were meant to provide are silently missing, and it's not easy for anybody to notice. Observed live on 2023-10-20. The build was attempted twice and failed both times. A large container image each time, which also means repairing the dependency installation inside the container.
HXC-235	Bug	Queued	—	The HelixLLM gateway provides a feature that allows for exact wording. On startup it reports that no response for this purpose, so it has fallen back to a simple hashing strategy. It states plainly that this fallback does not capture the full meaning and will be significantly degraded. The important part is that it will off or return an error: it keeps answering even when the response is effectively arbitrary rather than relevant. Any deviation from the responses that it is not working properly is a full platform boot. Resolving this means pointing to a genuine embedding provider, and deciding what level of degradation is acceptable behaviour or whether the feature should be disabled.
HXC-236	Bug	Queued	—	The standard library routine that splits a web address into components differently in two parts of this codebase, using the same logic. In the main application it rejects an address whose host component contains a colon. In the containers component it quietly accepts the same address. The cause is that each component declared its own logic. The newer version tightened this routine; the containers component, so it keeps the older, permissive behaviour. The reasoning about address handling is only valid in the containers component, which has already caused one incorrect conclusion. Second, and more seriously, the day someone releases a newer language version, addresses it accepts that currently fail will appear as unreachable services rather than a configuration change. The work is to make the difference explicit in the code, deliberately which behaviour each component has. The reader will find it, and add a check that fails if the behaviour is noticed.
HXC-238	Bug	Queued	Low	The safety check that stops half-finished experiments from throwing an exception for captured proof files, which would otherwise satisfy the old exception rule and so could not be caught. It's narrow: the file must sit in the evidence folder.

ATM ID	Type	Status	Severity	Description
				saved as non-runnable, and not begin with the Together those make the file inert in normal use. The bug is a deliberate pointing a program interpreter at a file which no file property can stop. The practical result is a crash rather than an accident, and the alternative was to keep a proof it is required to keep. Recording it so the bug is not absolute, and so anyone widening it later. When HelixQA runs a test bank that needs to stop, it gives out the wrong piece of it. The HelixCode server gives out different values: a real access pass (the field named 'access' for reference for the session (the field named 'session' for proof of identity on later requests. HelixQA is supposed to so it hands the server a value the server does not expect, or expired token'. The effect is that every test that fails even when the feature it is testing works perfectly. The worker-management bank scored 6 passed and 1 failed when pointed at the correct reason nobody noticed is that HelixQA's own implementation of the access pass in the field HelixQA expects, so the test against the real product. Anyone reading the results of the parts of the product are broken. One of the HelixCode test banks signs up a new account and uses the very same account name, 'hxc_gen_e2e'. The first run succeeds and the check passes. Every run after that, the database, the server correctly refuses with a failure recorded as a failure forever after. This was observed when the first run of the bank passed that step, and a second run was visible in the database timestamped to the moment the first works under require that automated checks can't be the same answer, cleaning up after themselves. The bug quietly poisons its own results, so a permanent failure to do with whether the product works. The fix is to reset each run, or to remove the account it created with the same name. The checks that confirm HelixCode can generate a valid response inside the server's reply, and that comparison of the two. Two checks in the generate-e2e bank ask for a response with different capitalisation, so they are marked as failed. The checks behaved correctly. The most serious case asks for a reply for 'hello' in small letters; the AI answered with a working feature was reported as broken. This was a bug in the server: asking it to reply with an invented word, and the answer it what seventeen times three is returned fifty-
HXC-239	Bug	Queued	—	
HXC-240	Bug	Queued	—	
HXC-241	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
				which proves real AI generation is running and a reply for ‘authorization’ in small letters while capital A. Expecting exact capitalisation of free nature. The consequence is false alarms that h time. Five checks covering the built-in screenshot against the running HelixCode server, and even server answers ‘QA engine is disabled’ and rep server is being honest here rather than misbeh whole advertised area of the product is switch cannot be used by anyone or confirmed as wor listing the available screenshot engines, listing status of a session, starting a new session, and with nothing to run. Because the feature is off, whether the underlying code behind it function feature is meant to be available, and if so switch behaviour is actually proven; if it is meant to s reporting failures. Several of the automated test collections that c about what a correct answer looks like. They s and record a pass regardless of the reply. This collection was deliberately aimed at the wrong an error for every request, and it reported bot passes against a service known to be broken ca one, so its green result carries no information failing test, because a suite that cannot fail is v where there is none, and it will keep producin every check an explicit expectation — the statu answer must contain — and then prove each o something known to be wrong and confirming the same run: four apparent failures in the wo harness looking for the wrong field name in th and that mismatch should be corrected so real One blocked work item (HXC-172) has its expla has that same explanation written into its own to agree. The tracker rebuilds those stored exp time it regenerates its records, so an explanati the text is discarded the next time the tracker happens today: a regenerate-and-reread cycle then correctly reports the item as blocked-with that explanation is what makes a blocked item what was already tried, and which decision w
HXC-242	Bug	Queued	—	
HXC-243	Bug	Queued	—	
HXC-245	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				decision waiting on a person into an item nobo keeping the context needed to make the blocki why the item stalled. The expected outcome is explanation again, the stored and written copi preserves it — proven by the tracker’s validati A complete run of the agent runtime’s automa 12 groups failing, and 74 individual test failure than assumed, because the early evidence poin than at the product. Three patterns account fo the in-memory cache service because they loo container while this machine publishes it on a running service to answer a web request and f and answering correctly when asked by hand; against fixed limits and overshoot them by abou jobs and a service rebuild were competing for product is healthy, and none of it proves the pr parallel investigations are now establishing, ea rather than a plausible story. This matters beca reasons is as damaging as one that passes whil everyone to stop trusting the result, and both h operator benefits from a suite whose red mean signal they can act on. The expected outcome i evidence as a genuine product defect, a test de contention artifact, each real defect fixed with returning to a trustworthy state. Two separate parts of the system have both be and neither knows about the other. The agent decides which service gets which port number itself. But a different component, the model ve names port 8100, and because the verifier star agent runtime is launched from a settings file t never even tries to claim the one it was allocat port 8100 for the agent runtime and are answe recognise the requests and rejects them with a runtime as broken when it is running perfectly because it makes a large part of the test suite s names, and because the same collision would n real. Whoever operates or deploys this benefit failing, and a genuinely failing one would look a single place that decides port allocation for e claiming the same number, the agent runtimes and the tests then passing or failing for reason
HXC-246	Bug	Queued	High	
HXC-247	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
HXC-248	Bug	Queued	High	A cleanup step in the automated test suite can be the only thing preventing that is a coincidence to stop any containers they started. To avoid this first check whether the agent service is still running. <i>something</i> answers on network port 8100, with <i>something</i>. On this machine port 8100 is held by a different process. If the check says yes, and the shutdown is skipped, the check is never actually consulted. If the agent is late, or moved, that same check would say no and the live platform out from under whoever was using it. This is also untrue: it reports that the agent is still running. This matters to anyone running the test suite. This deployment, which is the normal case here. The test suite checks the identity of what is answering - not merely whether the logs is true.