Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Operator-blocked	1
Bug	Queued	50
Bug	Ready for testing	1
Task	Operator-blocked	1
Task	Queued	14
TOTAL		**67**

Items

ATM ID	Type	Status	Severity	Description
HXC-159	Task	Queued	High	**Reported-Via:** §11.4.202 reporting directive `ta` Operator **What (the report, verbatim):** Featu the HelixSkills System (git@github.com:HelixD and HelixAgent, covering ALL of its power-fea OPERATOR REQUIREMENTS (2026-07-29): “Ful HelixAgent the HelixSkills System for all work git@github.com:HelixDevelopment/skills.git , w power-features fully without nothing leftout! E into phases, tasks and subtasks with exhaustiv prepared! Full coverage with ll supported test (suites) is mandatory! Extend and update all ex guides, FAQs, create stunning illustrations, gra all materials! Any gaps, shortcomings, weak sp inconsistencies MUST BE detected in advance c process and properly tackled with comprehens risk-free and safe! Track all progress through t

ATM ID	Type	Status	Severity	Description
				and remebering mechanisms we have! Use Git with bridge to Superpowers used for the imple development mechanism!" PRE-VERIFIED ENV not assumed — §11.4.6): - GitHub SpecKit IS in NOT yet initialized in this repo (no .specify/ an is therefore an explicit early task, not an assur skills.git IS reachable via SSH. `git ls-remote -` catalog-docs, feature/deep-research, feature/te branch/tag inventory MUST be re-derived at re partial listing. - The repository is NOT currentl contains no HelixDevelopment/skills entry; the codex-skills -> vasic-digital/caf-codex-skills.git, not be confused with it. - Both consumers exist (module dev.helix.code) and the submodule su dev.helix.agent). MANDATORY DELIVERY CONS for ALL phases of the incorporation, bridged to executed subagent-driven (§11.4.70 / §11.4.20). tasks -> subtasks, with an exhaustive level of d where the design is load-bearing). - Complete t e2e, full-automation, Challenges (vasic-digital/ (HelixDevelopment/helix_qa) with autonomou concurrency/atomicity, race/deadlock, memory captured evidence (§11.4.5 / §11.4.69); every gu Documentation: extend AND update every exis only new ones. Produce illustrations, graphs a materials. Four-format export discipline per §1 the main README per §11.4.212. - Risk work is note: every gap, shortcoming, weak spot, dang detected DURING planning and each one paire free and safe solution (§11.4.102 systematic-de before rewrite (§11.4.74 / §11.4.28): survey the catalogues on GitHub AND GitLab first; extend duplicating, keeping them project-agnostic and incorporated as a properly-laid-out dependenc level or submodules/, never a nested own-org and HelixAgent must BOTH receive the full fea the other is an incompleteness to be tracked, n through the continuation docs, the memory/kn the workable-items SSoT (§12.10 / §11.4.131 / § 1. A SpecKit specification exists and is initialize feature, with an explicit feature inventory pro repository (an enumerated list, not a sample — implementation plan exists at the stated depth

ATM ID	Type	Status	Severity	Description
				mitigation. 3. Every planned capability is wire consumers, proven by §11.4.108 runtime signa compiled. 4. The full §11.4.169 test-type matrix and HelixQA banks included. 5. All documenta produced diagrams/illustrations incorporated, Independent code review reaches a zero-findi No capability is claimed without runtime evid every closure (§11.4 / §11.4.1 / §11.4.123). §11.4 PROGRAM (scheduled, not yet executed): (a) D scientific papers / open-source projects+codeb to incorporate this feature into the project (§1 debug (§11.4.102) of ALL data obtained to enu zone / inconsistency / imperfection, and desig one found. (c) The design MUST ALWAYS be en innovative – never less. (d) MANDATORY exha documents, an implementation plan down to l diagrams / schemes / graphs, illustrations, SQL relevant material (§11.4.65 / §11.4.73). (e) Inve + HelixDevelopment submodules/components freely where genuinely needed while keeping reusable (§11.4.28 / §11.4.74 / §11.4.177). (f) Pla one: every supported test type, the Challenges (§11.4.27 / §11.4.169). (g) Divide the work into nano-detailed enough to drive integration / im directly. (h) Plan explicitly for enterprise scala the feature has a UI/UX surface, produce full w PSD / PDF, or whatever the project mandates) a §11.4.190). (j) Plan full CodeGraph integration synchronisation (§11.4.78 / §11.4.79 / §11.4.80). items (every detail + reference + attachment) s source-of-truth (§11.4.93 / §11.4.95) + every de related project doc/component + every connec (§11.4.148 D5) – honestly SKIPPING any absent (credentials_absent / tracker_client_absent, §1 Research-doc destination (to be populated whe fully_incorporate_the_helixskills_system_into_ Execution model (§11.4.213 clauses 1-2): this it multi-day research/planning/documentation w project's standing autonomous loop (§11.4.87 / when it claims this item, and MUST be driven explicitly evidence-backed CLOSED terminal s be left sitting un-wired in the backlog. **Affecte** research/planning – destination: docs/research

ATM ID	Type	Status	Severity	Description
				fully_incorporate_the_helixskills_system_into_ **Reproduction / context:** UNKNOWN: not state established before any fix (§11.4.102 / §11.4.14 research-doc tree exists under docs/research/ fully_incorporate_the_helixskills_system_into_ every enumerated weak-spot/gap/danger-zone implementation plan down to lines-of-code exi are created, and the whole effort is validated p ADDENDUM — BINDING OPERATOR REQUIRE === VERBATIM: “Make sure that SpecKit and S and with all extensons we have created derriv not optional colour on the SpecKit mandate — work MUST compose with rather than duplica BRIDGE EXTENSION (verified to exist, 2026-07 speckit_superpowers/implementation/ — the “ Superpowers Bridge”. Eight documents, each i IMPLEMENTATION_PLAN, NANO_TASK_ENGIN TDD_INTEGRATION, CONSTITUTION_INTEGRA SEVEN-LAYER architecture: 1 Developer Work SuperSpec bridge (orchestration — github.com (discipline enforcement — speckit-community MCP (execution) 6 Helix LLM (inference — git Distributed host cluster (llama.cpp RPC) CONS MUST BE IN SCOPE (inherited BY REFERENCE diverges silently): constitution/skills/ (7) action reporting-workable-items, scheduled-work-qu constitution/mcp/ (2) media-validator-mcp.json plugins/ (2) helix, scheduled-work constitution system), subagent_tiering.yaml constitution/sc credential_scan_lib.sh, guard-branch-consiste guard-track-branch-label.sh, guard-work-track LOCALLY — do not re-create: .claude/skills/ 10 clarify, constitution, converge, implement, pla workflows/speckit/workflow.yml, integrations/ templates/, memory/ ADDED ACCEPTANCE CR extension above has an explicit disposition — reason. Silence is not an answer (§11.4.118). 9. WIRED vs DESIGNED-ONLY with a captured p evidence a layer is installed — that is precisely and reporting a documented layer as an availa namespace collision risk is addressed with a d skills system, while §11.4.164’s post_update_h into .claude/skills/. Two systems registering in

ATM ID	Type	Status	Severity	Description
				and clash handling, not discovery at runtime. I non-own-org dependencies (§11.4.74 vendor p supply-chain exposure must be assessed, not a them.
HXC-164	Task	Queued	Medium	The script that brings up the supporting servic directly, rather than going through the shared maintains for exactly this purpose. That compo the same way, gets the same health checking, a repeatedly. Bypassing it means this project car and has already caused one outage of its own. service startups still hide their errors, so a fail was flagged honestly in a commit message, bu obligation anyone will act on, which is why it i routing the startup through the shared compo hiding so a failed start is visible.
HXC-166	Bug	Queued	Medium	When publishing the agent component, the ho outstanding security advisories against the lib 204. That figure turns out to be unreliable in th drifts upward on its own as new advisories are 205 the next day, and 210 on a live re-query on be "the" count. It over-states the work by abou counted once for every dependency file it appe distinct advisories, and the five criticals collap two places this code is published; the other cop was never switched on there. The recorded bre critical, 88 high, 99 moderate and 18 low. A ful completed, and it changes the picture substant the affected code can actually be reached in so them currently qualify. The Go vulnerability sc actually release, reports our code is affected by against a container library are reachable only in the codebase refers to and that never enters remaining alerts, including three of the four cr third-party project whose container definition cannot start today. Most of the rest belong to a builds. Exactly one alert touches a component defence that advisory asks for has already bee alert stays open only because the dependency related piece of housekeeping was resolved wh files of downloaded third-party library code, in behind one critical advisory, had been commit removed from the tracked files, though they re

ATM ID	Type	Status	Severity	Description
				benefit of finishing this work is that the alarmi the product, and we regain the ability to say ho carries a reachable known vulnerability. This i acceptance criterion has not been done: the de findings have not yet been upgraded or remov retire the copied-in project that carries most o
HXC-168	Bug	Operator-blocked	High	The password used to reach the project databa setup and container files, and those files are pu accounts. Anyone who reads the project can re not introduced by recent work; the project's ov it as something that must be replaced, and tha depends on whether the same password is use has not been established either way and shoul value is already public, changing the files alon be changed everywhere it is used, and the setu configuration source rather than containing it cannot honestly claim its credentials are prote
HXC-171	Bug	Queued	Medium	A complete copy of somebody else's web applic it alone accounts for one hundred and thirty-n dependency security warnings, including thre be built or run: the startup configuration poin that folder, and its supporting packages have r two opposite ways at once. Because it is not bu we actually run today, which is why they rank references it, the setup is broken and misleadi recipe file, all one hundred and thirty-nine wa to whoever does it. Anyone reading our securi distortion, since it currently buries the small n The expected outcome is an explicit decision, r properly maintain and build it, or remove it an a reference to the upstream hosted service. Be approval, the decision must be put to the opera
HXC-172	Task	Operator-blocked	High	A small companion service we wrote ourselve part of the standard deployment, and it carries supporting packages, two of them rated high. U warnings, these have no mitigating evidence in unshipped, because it demonstrably ships, and whether the affected functionality is ever exer place in the whole review where something w which makes it far more deserving of attentio three is not really an upgrade problem at all: i switched off unless explicitly enabled, so upda

ATM ID	Type	Status	Severity	Description
				protection on would leave the exposure in plac deployment running this service benefit. The e assessment of the three warnings, the two stra protective setting explicitly verified as enabled and the service rebuilt and confirmed working
HXC-175	Bug	Queued	Low	The cache that remembers model information where its clock was never configured. That fal the only way to build the cache always configu unexercised. On its own that is harmless. The r if someone later edits it: a plausible-looking ch current time would make every cached entry a in the past that the freshness check reads as a expiry limit. Every entry would then be consid keep serving outdated model information inde The choice is to either add a small test that pin branch entirely and document that there is on
HXC-178	Bug	Queued	Low	When a caller presents an API key, the server c ordinary text comparison. That kind of compa that does not match, so a key sharing the first f reject than one that differs immediately. An at differences across many attempts can recover guessing it outright. Network timing noise mal is low rather than urgent, but it is a genuine w cheap: use the comparison function designed f amount of time regardless of where the differe introduced by recent work.
HXC-179	Task	Queued	Medium	The check that decides which websites may op deliberately allows requests that carry no orig behaviour, because only browsers send that in legitimately omit it. It is safe today only becau away anyone without valid credentials. The tw nothing records or enforces that coupling: if so login check while working on something unrel information would become completely ungate work here is to add a check that specifically co the connection upgrade, so the dependency is
HXC-180	Task	Queued	Low	Three commits in the current release range do referred to a piece of data that was not added those three points gets a build error. The tip of corrected shortly afterwards. This cannot be re rewriting published history, which the project is for anyone hunting a regression by stepping

ATM ID	Type	Status	Severity	Description
				three points that fail to build for a reason unre may wrongly conclude the fault lies there. The affected points and their shared cause somewh so the failure is recognised immediately as kno from scratch.
HXC-181	Task	Queued	Medium	HXC-160 corrected the 21 stale references to sc scripts/generate_fixed_summary.sh across the added the CM-SUPERSEDED-GENERATOR-NAM sets were deliberately left untouched and are t Constitution.md still carries 14 references acro universal rules that projects other than HelixC records such as the Phase 5 Go-binary-shim pla is a governance amendment that CONST-049 re configured upstream, then re-cascaded and po session that closed HXC-160, and a canon amen drift than the stale wording, so the canon was 45 owned submodules carry cascaded restatem correction applied in the parent repo is delibe because the fix names a HelixCode-specific me wording into decoupled, project-agnostic subm them and violate CONST-051(B). The correct se neutral canonical wording that names the wor and push it in the constitution submodule per fleet and bump pointers. Who benefits: every c the anchor text, and any agent reading a subm summaries are produced. Acceptance: canon r wording, the change is pushed to every constit are re-cascaded, submodule pointers are bump SUPERSEDED-GENERATOR-NAMING both pass
HXC-188	Task	Queued	Low	The main handbook that tells contributors and organised lists several folders at the top level t when only one is present, and refers to a top-le Anyone following it — a new contributor, or an orientation — will look for things that are not t broken or that they have checked out the wron of the document, which is otherwise the autho The correction is small and purely descriptive: section to match, then keep the two in step as t
HXC-189	Task	Queued	Low	One of the external projects this codebase depe configuration still refers to it by its previous na hosting service silently forwards requests from forwarding is a courtesy, not a guarantee: it dis

ATM ID	Type	Status	Severity	Description
				project under the old name, and it is not honou dependency. If it stops working, fetching the pr name nobody recognises, which is unusually c project's current canonical name so the referer on a redirect continuing to exist.
HXC-190	Bug	Queued	Low	When a running command is cut short becaus says it was not a timeout. The check that sets th the caller's own deadline, which usually has no stopped the command. Investigation showed th either, because the internal signal it would cor cancellation and never a timeout, and the time mechanism entirely. The effect is that anyone c failure and no indication that a time limit was wrong place. A sibling code path that runs com correctly by doing it inside the timer's own cal copy.
HXC-191	Bug	Queued	Low	When a command is cancelled or times out, th but anything it started, by signalling the whole works if the command was placed into its own sandbox enabled — the normal configuration sandbox switched off, the grouping step is skip operating system to signal a group that does no error is discarded. The result is a cancelled com dormant, because the only configuration that c so no shipped path reaches it. It is worth fixing would appear only in an unusual configuration diagnose later.
HXC-192	Bug	Queued	Low	Output from a running command is read line b line may be. When a line exceeds that limit the permanently for that stream, because the buffe past it. Nothing else takes over the reading. Th until the operating system's own small buffer f someone to read. So one unusually long line, w readable output produce quite naturally, can w is to keep draining and discarding the remaind abandoning the stream, so the producing comm
HXC-196	Task	Queued	Medium	The handbook that describes this project's tech one of its networking libraries. Both parts of th because a security fix required moving off the deliberate; it is the handbook that is now wron a new contributor or an automated helper rea will confidently state a version that has not be

ATM ID	Type	Status	Severity	Description
				same document has already been found descri additional inaccuracy makes the whole docum recorded version to what is actually in use, and reason, so nobody later reverts it believing the
HXC-197	Bug	Queued	Low	One test directory contains a duplicate of its ov files appear twice in the project at two differer renaming exercise that moved a tree while lea the new one. Nothing reads the nested copy, so anyone searching the project finds two results which is real, and any tool that scans for dupli blocks a useful check from being switched on: duplicates cannot be enabled while this one ex every run. Removing it requires first confirmin nested copy is genuinely the leftover and not t
HXC-200	Bug	Queued	Low	A detector was written to catch a specific kind stuck trying to hand off a result nobody is colle one particular description of what the worker form of the same stall — where the worker wa rather than a single handoff — and the detecto demonstrated rather than theorised: a delibera went completely undetected while the detecto does catch is the historical one, so the guard is the newer structure while breaking its escape recognise both descriptions is a small change a defect rather than half of it.
HXC-204	Bug	Queued	Medium	Two long-running tests — one exercising heav one exercising leader election among worker r report failure when the machine is under load own. Measured: during a full test run both fail times consecutively, one finishing in roughly a about a twentieth. Their verdicts are therefore whether the product works. This matters most test run performed before a release is by defin these two will mark a healthy release as broke through. A third test with the same problem w inconclusive run report as skipped rather than the same treatment fits here. Both should wait than assuming a fixed time is enough.
HXC-206	Bug	Queued	Low	The newly-fixed health refresh deliberately rel provider over the network, so that a slow or ha It then re-acquires the lock and only applies its changed underneath it. There is a narrow rem

ATM ID	Type	Status	Severity	Description
				started again entirely within one probe, the ch applies a verdict formed against the previous i briefly reported unhealthy when it is fine, whi dangerous direction — reporting a stopped pro Worth recording that this is cheap to close rath incremented on each state change is invisible t to any published structure. An earlier note des was wrong and has been retracted.
HXC-207	Bug	Queued	Low	When a model provider record is created, its e shared template by reference rather than copi points at the same underlying settings, includi Nothing writes to those settings today, and the no corruption occurs at present — this was ver record it is that the protection now added to th manager cannot protect a structure shared wi legitimate, properly-locked write to a provider manager's view as well, and the locking will lo when the record is created removes the hazard
HXC-208	Bug	Queued	Low	Four compose entries across two end-to-end te the Slack mock and the LLM-provider mock, b Any attempt to build those stacks therefore fai which means the containerised form of these t exercised. This was surfaced while fixing an u same two services: the fix could be proven at t against a running container, precisely because produced. That gap matters beyond these two process leaves the deployed-artifact layer unpr failure the four-layer verification discipline ex Dockerfiles, or to remove the build entries if th run containerised, and then to prove the choic
HXC-216	Task	Queued	Low	A blanket rule excluding log files from version evidence, because evidence transcripts are wr corrected so this cannot happen again, and the have been recovered and committed. Thirty-fi the runs exist only on the machine that produc an evidence folder that appears empty, which never run. Recovering them is not automatic, b seventy-eight megabytes and committing those deserves a deliberate decision rather than a re keep what genuinely substantiates a closed ite ones rather than storing them whole, and reco why, so the gap is visible rather than silent.

ATM ID	Type	Status	Severity	Description
HXC-219	Task	Queued	Medium	A recent repair changed how the release gate c evidence really belongs to the change it claims commit on a labelled line rather than merely h somewhere in the text. That new requirement have to follow it: the guide that authors read b the new labelled field, so a contributor followi evidence the gate will reject without understar describing the gate itself has printable and we page was last edited, so all three versions disag describe. Fixing this means updating the autho citation field and regenerating the exported co who benefit are contributors recording eviden have no accurate written description of the ru can follow the written guide and produce evid attempt.
HXC-224	Bug	Queued	High	When we mark a piece of work finished, we re works. Nobody ever checked that those pointe different kinds of rot accumulated silently and report. First, some pointers hold a whole parag there is no way for a machine to follow them. S never created because the person doing the wo means the proof exists but nobody can find it f ordinary progress updates rather than to the a notes look like formal proof. The effect is that a backed in every summary and dashboard whi way anyone finds out is by manually hunting f make the completion pointer a checked field: i exists, it must be attached only to real complet must be refused at the moment someone tries later. This protects everyone who relies on our claim of proof that cannot be produced on den
HXC-225	Task	Queued	Medium	A single bundled helper component inside the roughly 208 security advisories reported agair the total — and it entered the codebase withou established what it does, whether we wrote it c actually uses it, or whether removing it would answered the headline advisory count is misle carries enormous risk, when in reality the grea not even be part of what we deliver. The work purpose, whether it is reachable from anythin decision about it — keep and maintain, upgrad record which was chosen and why. The decisio

ATM ID	Type	Status	Severity	Description
				advisories, because upgrading dependencies i effort spent on nothing.
HXC-226	Bug	Queued	Medium	The agent project is mirrored to two hosting a was only ever switched on for one of them. Qu hundred and ten open alerts while the other r one the main project is configured to pull from of the component we actually consume is show hundred findings sit unseen on the sibling cop the two mirrors hold the same commits. What watched. This is dangerous precisely because t findings is indistinguishable from a report of z the second case is what we have. The work is t from, confirm both then report the same figur the two so a future divergence is noticed rathe statement about this component's vulnerabilit from.
HXC-227	Bug	Queued	—	A design document committed forty-eight day appears to be a genuine access key for an exte placeholder. The section is headed as being co and described as already configured, and the v placeholder would have: it is fifty-one charact characters, and contains no example or to-be-f the main branch and has therefore been publi day it was written. Anyone with read access to fork made at any point in those forty-eight day now would remove it from the current revisio history is forbidden, and rewriting would in a The only action that actually withdraws the ke replacement, which requires access to the pro within the repository. Until that happens the k revoked, the document should be edited to ref of quoting its contents, and a check added tha this shape.
HXC-231	Bug	Queued	—	The HelixLLM gateway downloads large AI mo starts accepting network connections. During t the service as running and healthy, but every a on 2026-08-06 this window lasted thirty-seven actually began listening. Anyone checking the everything is fine while the service is in fact u depends on the gateway will fail for the whole because the download depends on an externa one of the two downloads timed out after roug

ATM ID	Type	Status	Severity	Description
				that site were unreachable the gateway might accepting connections first and download mod honest not-ready state while the downloads ar
HXC-232	Bug	Queued	—	HelixAgent expects two tables to exist in its da agent_instances, but neither of them has ever anywhere in the project that creates them. As locks once every minute, the database refuses and an error is written to the log. Verified on 2 directly: neither table exists in any of them, an minute period plus a further hundred and thi otherwise runs normally, so nothing visibly br buries genuine problems in the log and means to support is silently doing nothing at all. Fixi these tables were intended to contain to write guessed.
HXC-234	Bug	Queued	—	On startup HelixAgent tries to build and launc known as MCP servers, using the container en inside the container fail during their depende operation is abandoned. HelixAgent records a these servers and then continues running nor nothing obviously breaks. The real effect is tha were meant to provide is silently missing, and easy for nobody to notice. Observed live on 20 the build was attempted twice and failed both large container image each time, which also m repairing the dependency installation inside t
HXC-235	Bug	Queued	—	The HelixLLM gateway provides a feature that exact wording. On startup it reports that no re purpose, so it has fallen back to a simple hashi states plainly that this fallback does not captu will be significantly degraded. The important p off or return an error: it keeps answering ever effectively arbitrary rather than relevant. Any from the responses that it is not working prop full platform boot. Resolving this means pointi genuine embedding provider, and deciding wh is acceptable behaviour or whether the featur
HXC-236	Bug	Queued	—	The standard library routine that splits a web differently in two parts of this codebase, using the main application it rejects an address who containers component it quietly accepts the sa colon. The cause is that each component decla

ATM ID	Type	Status	Severity	Description
				the newer version tightened this routine; the c one, so it keeps the older, permissive behaviou reasoning about address handling is only valid which has already caused one incorrect conclu Second, and more seriously, the day someone r newer language version, addresses it accepts t failure will appear as unreachable services rat change. The work is to make the difference exp deliberately which behaviour each component reader will find it, and add a check that fails if noticing.
HXC-238	Bug	Queued	Low	The safety check that stops half-finished exper an exception for captured proof files, which w satisfy the old exception rule and so could not narrow: the file must sit in the evidence folder saved as non-runnable, and not begin with the Together those make the file inert in normal us deliberately pointing a program interpreter at which no file property can stop. The practical r rather than an accident, and the alternative wa proof it is required to keep. Recording it so the than absolute, and so anyone widening it later
HXC-239	Bug	Ready for testing	—	When HelixQA runs a test bank that needs to s sign-in reply, so every test in that bank that ne MECHANISM (the original wording of this item the reply carries the real access pass at the top separate bookkeeping reference for the session session object. The runner looked for the book nothing there, and stopped with a plain compl never handed the server a wrong value and w failed before making a single signed-in request the wrong value and being rejected, which is n that text were nonetheless correct. A SECOND step had its own copy of the same assumption could not find the pass it simply carried on wit unsigned and the resulting rejections were rec rather than as a failure to sign in — which is w THE OBVIOUS FIX WAS WRONG: the bookkeep different family of about fourteen banks that t access pass genuinely does sit at the top level u names would have repaired one product by br the existing name first and adds fallbacks tried

ATM ID	Type	Status	Severity	Description
				support for reaching a nested location, and rep HOW TO REPRODUCE: run a bank requiring si run stop with a missing-field complaint before OUTCOME: the bank used for measurement w passing and two failing. The two that remain a work: one posts an empty body where the serv with a leftover record created by the very mea original figures, because that bank does not cl the originally reported nine-and-one is no long banks needing sign-in authenticate against bot changes; a failure to find a usable pass is repo service faults; no credential value ever appear proves the behaviour, shown real by a paired
HXC-240	Bug	Queued	—	One of the HelixCode test banks signs up a nev uses the very same account name, 'hxc_gen_e2 succeeds and the check passes. Every run afte the database, the server correctly refuses with recorded as a failure forever after. This was o run of the bank passed that step, and a second visible in the database timestamped to the mo works under require that automated checks ca the same answer, cleaning up after themselves quietly poisons its own results, so a permanen do with whether the product works. The fix is each run, or to remove the account it created
HXC-241	Bug	Queued	—	The checks that confirm HelixCode can genera inside the server's reply, and that comparison Two checks in the generate-e2e bank ask for w with different capitalisation, so they are mark behaved correctly. The most serious case asks reply for 'hello' in small letters; the AI answer working feature was reported as broken. This server: asking it to reply with an invented wor it what seventeen times three is returned fifty- which proves real AI generation is running an reply for 'authorization' in small letters while capital A. Expecting exact capitalisation of fre nature. The consequence is false alarms that h time.
HXC-242	Bug	Queued	—	Five checks covering the built-in screenshot a against the running HelixCode server, and eve server answers 'QA engine is disabled' and rep

ATM ID	Type	Status	Severity	Description
				server is being honest here rather than misbeh whole advertised area of the product is switch cannot be used by anyone or confirmed as wor listing the available screenshot engines, listing status of a session, starting a new session, and with nothing to run. Because the feature is off, whether the underlying code behind it functio feature is meant to be available, and if so switc behaviour is actually proven; if it is meant to s reporting failures.
HXC-243	Bug	Queued	—	Several of the automated test collections that c about what a correct answer looks like. They s and record a pass regardless of the reply. This collection was deliberately aimed at the wrong an error for every request, and it reported bot passes against a service known to be broken ca one, so its green result carries no information failing test, because a suite that cannot fail is w where there is none, and it will keep producin every check an explicit expectation — the statu answer must contain — and then prove each o something known to be wrong and confirming the same run: four apparent failures in the wo harness looking for the wrong field name in th and that mismatch should be corrected so real
HXC-245	Bug	Queued	High	One blocked work item (HXC-172) has its expla has that same explanation written into its own to agree. The tracker rebuilds those stored exp time it regenerates its records, so an explanati the text is discarded the next time the tracker happens today: a regenerate-and-reread cycle then correctly reports the item as blocked-with that explanation is what makes a blocked item what was already tried, and which decision wo decision waiting on a person into an item nobo keeping the context needed to make the blocki why the item stalled. The expected outcome is explanation again, the stored and written copi preserves it — proven by the tracker's validati
HXC-246	Bug	Queued	High	A complete run of the agent runtime's automa 12 groups failing, and 74 individual test failure than assumed, because the early evidence poi

ATM ID	Type	Status	Severity	Description
				than at the product. Three patterns account fo the in-memory cache service because they lool container while this machine publishes it on a running service to answer a web request and f and answering correctly when asked by hand; against fixed limits and overshot them by abou jobs and a service rebuild were competing for product is healthy, and none of it proves the pr parallel investigations are now establishing, ea rather than a plausible story. This matters beca reasons is as damaging as one that passes whil everyone to stop trusting the result, and both l operator benefits from a suite whose red mear signal they can act on. The expected outcome i evidence as a genuine product defect, a test de contention artifact, each real defect fixed with returning to a trustworthy state.
HXC-247	Bug	Queued	High	Two separate parts of the system have both be and neither knows about the other. The agent decides which service gets which port number itself. But a different component, the model ve names port 8100, and because the verifier star agent runtime is launched from a settings file never even tries to claim the one it was allocat port 8100 for the agent runtime and are answe recognise the requests and rejects them with a runtime as broken when it is running perfectly because it makes a large part of the test suite s names, and because the same collision would r real. Whoever operates or deploys this benefit failing, and a genuinely failing one would look a single place that decides port allocation for e claiming the same number, the agent runtimes and the tests then passing or failing for reason Later evidence settles which of the two sides is deployment is not misconfigured at all. The mo file that starts it on this machine - documents i agreed meeting point between the verifier and gateway looks there by default, and, in as man runs on the other address here. The choice wa down at deployment time, in the verifier's favo clearer than first described: the agent runtime

ATM ID	Type	Status	Severity	Description
				date, still claiming an address that a different, with a note that the old address was abandone still in use. A contributing cause is that the ver so the one mechanism meant to prevent two c not arbitrate a component it had never been to unchanged - the tests inherit the stale claim ar the repair direction is now settled rather than the address it genuinely serves, give the verifie recur, and point the test defaults at the real ad establishes which side is stale and why; it does safe to execute, because that entry is consume change still needs its impact review before it l
HXC-248	Bug	Queued	High	A cleanup step in the automated test suite can the only thing preventing that is a coincidence to stop any containers they started. To avoid st first check whether the agent service is still ru *something* answers on network port 8100, with On this machine port 8100 is held by a differen the check says yes, and the shutdown is skippe different port and is never actually consulted. late, or moved, that same check would say no a live platform out from under whoever was usi also untrue: it reports that the agent is still run all. This matters to anyone running the test sui deployment, which is the normal case here. Th the identity of what is answering - not merely makes its decision on the real condition rather logs is true.
HXC-250	Bug	Queued	High	When a user asks the system to write out a rea command-line coding assistant, the file it prod service. The generator fills in a network addre than to the agent runtime the assistant is supp requests to a component that does not underst reply. From the user's point of view the assista error explaining that the address in their bran was written. Of the forty-eight assistants the p affected; a fix landed earlier covers only two o all the rest still writes the wrong address today in the current set, because it is not an internal handed to a user with an instruction to use it. newcomers the feature exists to help: someone following the documented path, and getting a

ATM ID	Type	Status	Severity	Description
				The expected outcome is that every generated runtime genuinely serves, that all forty-eight d than two, and that a generated config works w without editing anything by hand.
HXC-251	Bug	Queued	Medium	One of our own services could not start while both had been told to use the same network ad conventional default that the wider industry u precisely what makes it the number most likel own machine it was taken, by a supporting dat The result was that the service refused to start client dialling that same well-known address r which completed a perfectly healthy-looking c Because the wrong service was alive rather tha something is unreachable did not skip: they ra our service, which had never been running at twenty test failures that had been attributed to service in the project's central address book so never silently collide again, and it now fails im rather than starting somewhere unexpected. W misleading: the startup message printed for th so anyone reading the log is told the wrong pla platform benefits, and the expected outcome is the rest of the stack and that every message it bound.
HXC-254	Bug	Queued	Medium	Tests in one area of the codebase interfere wit share a single piece of network machinery, so to do with the code it is testing. When each con given its own timeout setting but no connectio library quietly falls back to one shared handle in this area start roughly a hundred and forty- around five hundred cases to run simultaneou down at the end of its own test, politely tells th and because the handler is shared, it also seve of using. The result is a failure that appears in and vanishes when the tests are run one at a ti the test setup rather than in the product. It rep one run in ten fails this way. Developers are th unpredictably teaches everyone to re-run it ra comfortably inside that noise. The expected ou own connection handler so shutting one prete test, and the suite becomes trustworthy enoug genuinely wrong.

ATM ID	Type	Status	Severity	Description
HXC-255	Task	Queued	Medium	Two things are left over in the scripts that run needs a decision from the operator rather thar problem: the end-of-run summary prints "Inte whenever the log file merely exists, regardless which tests genuinely failed still shows a reass most likely to read. The second is deliberate ar written to note test failures as warnings and ca status even when real tests have failed, which blocks a release. That behaviour was chosen or abort a long run and lose the rest of the inform between gathering everything in one pass and broken. Because it is a deliberate posture rathe operator decision about how the project wants be altered silently by whoever happens to touc runs to judge release readiness benefits from s only that the script reached the end. The expec reflects the real result rather than the presenc explicit, recorded choice about whether these
HXC-256	Bug	Queued	Low	The tool that regenerates the project's tracker happens to be run from unless it is explicitly to one optional setting quietly produces a second while the real ones stay out of date. The tool's optional, which is what invites the omission, a folder rather than the folder where this projec found the hard way during routine work: runn setting created sixteen stray copies at the top o one directory down, kept their old contents. Th someone can look at freshly written files, see t records are current when the versions under w the tracker documents by hand is exposed, wh automated helpers working outside the standa bounded and that bounds the severity: every w destination explicitly, the tool prints the full pa visible to a careful reader, the stray copies sho separate consistency check would eventually r What is missing is the tool itself objecting. The documents without naming a destination eithe or refuses and says so, instead of succeeding so caught this does not exist today: the existing gu supplies the destination explicitly, so it can nev out - adding that second case is the specified fo
HXC-257	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				The automated check that guards our workabl actually measures, so a reader who trusts its w computed. Internally it names the thing it insp written description states that it proves the co it never consults version control at all - it simp folder, which may contain changes that have n genuinely proves is narrower and still useful: right now agree with each other. The gap was the check reported success while the database yet been committed - exactly the situation its o stating plainly is that the underlying behaviou of the pre-release sweep, where the point is to commit is self-consistent, so inspecting the wo there is a documented reason for the copy - op as a side effect. On that reading the repair is th Everyone who reads that check's verdict befor the difference between "what is on your disk i correct" is exactly the difference that matters w name has already spread: the description of th repeats the same claim. The expected outcome and the sentence it prints all describe only wh guarantee is wanted it is added deliberately as control. The test that would have caught this d disagreement between the documents and the between the committed copy and the working adding that case is the specified follow-up.
HXC-258	Bug	Queued	Medium	Our workable-items database keeps a small bo synchronisation ran between the database and That note is only ever written by one of the tw documents into the database - and the value it that actually ran. The path that goes the other database, never updates the note at all. The pr regenerated the note still describes an older ru reading it to find out what happened most rec matters because deciding the direction of a syn can destroy records if it is made backwards, an judgement while being incapable of supportin into-database run dated the eleventh of July w database-into-documents regeneration on the shared tooling that performs the synchronisat genuinely ran on both paths rather than assun hand would be undone by the next run and is

ATM ID	Type	Status	Severity	Description
HXC-259	Task	Queued	Low	Every workable item is supposed to accumulat its status changed - opened, reopened, closed - what evidence justified it. A recent repair note no such trail. Measurement across the whole d to those two: eighty-one distinct records cover history entries at all, and all but one of them a filed away. In other words the majority of our account of how it reached that state, which we understand why something was reopened, or t cause is that items created by direct database creation command skip the step that opens the before the trail existed never had one. This is c filling: inventing dates and authors for events trail that is worse than an absent one because instead is a decision on how to mark these rec that prevents any newly created item from joi
HXC-260	Bug	Queued	High	The Kubernetes manifests for HelixAgent told 9090, and they set an environment variable ca was true. No Go source in the repository reads the port the server listens on; this was proven bound :8112 while the real variable HELIXAGE 9090 is a port no HelixAgent process ever bind internal/ports registry (HelixAgentGRPC, core by running the built cmd/grpc-server binary a The practical effect is that anyone deploying H manifests would get a Deployment exposing co gRPC traffic to targetPort 9090, and a Network artifacts agreeing with each other and all poin port the server would actually bind was block number because the manifest LOOKED configu and edited GRPC_PORT would see no change w This item covers correcting all four manifests t service.yaml, networkpolicy.yaml, configmap.y the dead variable with the one the server hono limitation: the image this Deployment runs is no gRPC server at all (that is a separate binary, correct but INERT until a grpc-server containe corrected manifests make a one-line change in checking out the pre-fix tree and running the p RED_MODE=1, which asserts the dead variable Acceptance criteria: internal/ports/published_

ATM ID	Type	Status	Severity	Description
				unset (asserting every manifest names the reg outside a comment) and fails on the pre-fix tre
HXC-261	Bug	Queued	Critical	The challenge script that is supposed to prove structurally incapable of failing, and it has bee existed. Two independent faults compounded. environment variable named HELIXAGENT_G spelling that no Go source reads (the server ho number belonging to no service in the project: registry's 8112, not the HTTP port 7061. Secon recorded a successful assertion whether the po check recorded true when the port was listeni when it was not; the unary-call, streaming-call The two faults together mean the challenge dia ever bound, found nothing there, and then cer most damaging class of defect in the codebase — it is an automated test actively producing fa anti-bluff covenant exists to prevent, and it ma reading the suite. It also violated the challenge framework already ships a record_skip helper record_assertion with status true to mask a mi not used. This item covers resolving the port th default, consolidating the four duplicated prob interpreted in exactly one place, making abser OK ticket rather than success, making a real fa two-state final verdict (which reported PASSEI when nothing had been verified) with a three- assertion could be verified. Reproduce with th protocol_grpc_fail_open_guard.sh run with RE and releases a port to prove it is free, drives th and confirms assertions are recorded PASSED guard with RED_MODE=0 reports zero PASSED every probe SKIPPED and SKIP-OK tagged, and PASSED assertion when a real gRPC server is li
HXC-262	Bug	Queued	Medium	The project keeps a battery of tests whose who catch problems, by deliberately breaking thing assertion in that battery — the one covering a occasionally reports a problem when nothing consecutive runs, and ran correctly six times o intermittent rather than consistently broken. T check itself: it asks the system two separate qu which process is running it, and between those move on, so the two answers describe differen

ATM ID	Type	Status	Severity	Description
				happens the check cannot find a process to exa anything. The practical harm is not a missed d a condition that is not present, and worse, peo it, which quietly erodes the value of the one ba The fix is to read both pieces of information in same instant. This was deliberately left unfixed changes the sampling behaviour of a check tha such a change deserves its own review rather
HXC-263	Bug	Queued	Medium	The project keeps a folder of tests whose purp genuinely work, by deliberately introducing a check reports a failure. A test like that only has Counting the invocations at the current revisio folder are ever executed by anything; the rema never run. This is the same problem the projec when three safety checks turned out to exist b the same reason: a check nobody runs cannot provide coverage. The presence of these files is counting test files sees protection that is not be look better than it is — five of the files are men mentions are a comment, an entry in a skip lis none of which run anything. Resolving this me to wire it into the release sweep so it runs, or t choice visible rather than leaving the file sittin deliver.
HXC-264	Bug	Queued	Low	A single entry in the project's work tracker ma different ways in three different fields, and the says a required program is not installed on the exists but cannot be found because of where th such file exists anywhere on the system. Those different fixes, so a reader cannot tell from the also proposes a remedy, giving the service a di not what was eventually done: the capability w different service that was already running, and absent from the machine today. This matters b of what was reported and why, and an entry th purpose for anyone reviewing the history later out of step with the database, which is a proble rendered documents faithfully reproduce wha is inside the database entry itself. Resolving it accurate account that records both what was c resolved, without erasing the original report.
HXC-265	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				HelixAgent keeps 36 Go modules nested inside Most of them are never compiled by anything a developer can rename a package, change a s and every build and test still reports green — t That silent-failure surface is what this item tra documentation samples, because the project a documentation example works when copy-pas cannot honour that promise. WHO IS AFFECTE readers who copy a documentation example e verifiable by one command (the quotes are loa the pattern before git sees it and reports 2 inst `git ls-files '*/go.mod' | wc -l` reports 36 r `'challenges/codebase/go_files/*/go.mod' | w` challenge-modules recipe walks. By directory t challenges/, 2 under plugins/, 2 under docker/, cli_agents/. THE CLASSIFICATION IS PER-MODU revisions of this item stated a single wrong tot answered by any one command, so each class COMPILE PATH AT ALL — 25 modules: the 23 u nothing anywhere compiles them, the three ch only run presence checks (test -f, grep -q), and because Go does not cross nested-module bou GATED CONTAINER IMAGES — 2 modules: doc declared as build contexts in docker-compose. their Dockerfiles, but no test suite or gate exer rebuilds only when the image is absent, so bre deterministically rather than never. (3) FULLY pkg/api: it is required and replaced in the root main.go with no build tags, and built for four p failures increment PHASE_FAILURES and set t set and needs no work. Note for future readers does not contain package", which is a true out membership rather than whether the module this item mistook the one for the other. (4) PAR module, Toolkit: internal/verifier/startup.go im transitively pulls in three more Toolkit packag compile on every root build while roughly 17 r EXCLUDED WITH A RECORDED REASON AND is the HXC-139 isolation marker whose header be compiled, paired with the regression guard cli_agents_isolation_test.go. It already satisfies terminal. ACCEPTANCE CRITERIA: every modu remainder of (4) is either wired into a build or

ATM ID	Type	Status	Severity	Description
				runs, or deliberately excluded with a recorded gate fails when a new nested module appears
HXC-266	Bug	Queued	Medium	The function that builds the URL clients use to "http://%s:%d" without ever bracketing the ho malformed URL. WHY THIS IS REACHABLE RA layer deliberately treats the IPv6 loopback ::1 configuration that sets the host to ::1 passes ev formatter. HOW IT MANIFESTS: the emitted a parseable URL — the colons of the address and to any URL parser, so the client fails to connec clear configuration error. WHO IS AFFECTED: IPv6-only or IPv6-preferred host, and any depl IPv6 literal instead of a name. HOW TO REPRO the client base URL, and observe the returned correct bracketed form http://[::1]:7061. ACCEP literals per RFC 3986 so ::1 yields http://[::1]:70 a regression test covers IPv6 literal, bracketed IPv4, and hostname cases.
HXC-267	Bug	Queued	Medium	When HelixQA signs in to run an API test bank but a reply the runner cannot read, the runne sign-in, carries on making requests with no cr faults in the API being tested rather than as its the case where the reply IS readable but carri this item covers the door that fix left open. WH explanatory finding is nested inside the branc structured data, and that branch has no count unreadable reply skips the reporting entirely. tested service is broken when in fact the runn reads the report to investigate the wrong syste this item originally described two silent shape established by running every case against the out of eleven tested. Silent: an empty reply; an data; bytes that are not valid text; a bare quot an object; a bare true or false; and a very larg item's reporting: the literal null, and an empty whether the reply PARSES, not whether it carr narrower descriptions kept undercounting it. HelixQA report for a service whose sign-in rou error pages produced by a proxy or gateway in REPRODUCE: point the pipeline at a test serve with a body of empty text or HTML, run the A in finding is recorded while the later unauthe

ATM ID	Type	Status	Severity	Description
				ACCEPTANCE CRITERIA: an unreadable sign-ir finding as an unusable one, naming what was be read; the message never contains any byte c standing test covers all nine silent shapes plus paired mutation that removes the new reporti already-working shapes must produce byte-ide change cannot disturb the path that already re
HXC-268	Bug	Queued	Medium	The model verifier builds web addresses by pa colon, without the square brackets that IPv6 ac host is an IPv6 address produces addresses no operators on IPv6-first hosts, and any deploym rather than a name. WHY IT IS REACHABLE: th with no validation, so an IPv6 value passes stra MANIFESTS: the emitted address looks like htt read because the colons of the address are ind the port; the client then fails with a confusing configuration message, sending whoever debu PLACES, AND A CORRECTION WORTH RECORI contains FIFTEEN such build sites, one in each every one reachable from the single dispatch t revision of this item said three. That was wron fifteen sites use only three distinct format SHA plain, one ending in a messages segment), and upstream report were recorded as if they were places. Had the earlier figure stood, the accept with twelve sites still broken. HOW TO REPROI submodule, count occurrences of the formattir file; it reports fifteen, while listing the distinct FIX DIRECTION: do not hand-bracket fifteen pl submodule is introducing a shared endpoint h correctly using the standard library join, inclu YET part of any committed state of the submoc nobody should expect to find it there today. Or it, which fixes them uniformly and prevents a Coordinate with that work before starting. SCC capability-config generator is not the whole su appears in at least nine further files elsewhere agent config writers and two unrelated subsys deliberately, but a fixer should know the capab entire problem. ACCEPTANCE CRITERIA: all fif produce properly bracketed addresses for IPv6 are unchanged; an already-bracketed value is

ATM ID	Type	Status	Severity	Description
				the IPv6 literal, already-bracketed, IPv4 and ho mutation that removes the bracketing and mal unbracketed sites in that file is zero, verified b fifteen.
HXC-269	Bug	Queued	Medium	When the model verifier builds a web address contains a question mark, a hash, or a slash, th ends up pointing at the default port instead of address builder appends the port after whatev string produces something like a normal addre query rather than onto the host; every standar correctly, reads NO port at all, and treats the p MANIFESTS: the client connects to the wrong p anywhere, which is the same silent-misdirectic eliminate. WHY IT IS REACHABLE RATHER TH step accepts these shapes without complaint, s operator-facing host setting into generated cor problem. THREE SHAPES CONFIRMED BY MEA its port into the query; a host with a fragment l loses it into the path. WHO IS AFFECTED: any c host setting instead of the dedicated full-addre guidance anticipates as a likely mistake. HOW host containing a question mark and any port, is empty while the query carries the digits. AC query, fragment, or path either is rejected with expected shape, or is handled so the port survi told rather than left with a silently wrong addr shapes plus the ordinary host and address case removes the handling and makes the test fail.
HXC-270	Bug	Queued	Medium	HelixQA copies failed reply bodies straight into places, and one of those places is a sign-in repl most likely to be. WHAT HAPPENS: when a req code records the entire reply body as the descr THREE sites in the same validation routine — t and the media-search path. A fourth and fifth c error message carrying a slice of the body, and complaint, which on the current toolchain quo THIS MATTERS: these reports and errors are w automation, so anything the server put in that site is the highest risk by a wide margin — a si token, returns a cookie in the body, or includes output would have that value copied verbatim echo the tested service's own error bodies, whi

ATM ID	Type	Status	Severity	Description
				RELATIONSHIP TO OTHER WORK: a sibling ite it reports only a shape label, a byte count and separate component now redacts values that l predate that work and were deliberately left u discipline is now inconsistent inside a single fu failure path fifteen lines below it is not. HOW T server whose sign-in route returns an unsucce recognisable marker, run the validation step, a recorded finding and in the returned error tex sites emits raw body bytes; all report the same success path already uses; the decode-complai a shape rather than a quoted character; and a failure body never appears in any finding or e paired mutation that restores the echo and ma
HXC-271	Bug	Queued	Critical	One of the two families of remote calls the age machine, however well configured. WHAT IS E ones for completing text, streaming a completi routine that is always handed an EMPTY list of immediately when the list is empty. So every c regardless of how many providers are actually A CONFIGURATION PROBLEM: the object servi to the registry that holds the discovered provid the registry into the OTHER family only, and th consult it, which is why the two families behav from outside. HOW IT WAS CONFIRMED: captu plus real models had genuinely been discovere case of an unprovisioned host — the providers message saying no providers were configured. facade family, which is the family the publishe integrator would naturally reach for. HOW IT them to use the provider registry for credentia mistake on their side and sends them to check ACCEPTANCE CRITERIA: the facade family eith against configured providers, or is explicitly w a clear error naming its replacement; a test dr against a real configured provider and asserts of error; and a paired mutation that removes t
HXC-272	Bug	Queued	High	Two automated checks that start in the same s append to the same running log, so each is jud HAPPENS: each run names its results folder af second, with nothing added to make it unique, outcome lines in that shared log to decide the o

ATM ID	Type	Status	Severity	Description
				BOTH WAYS, AND THE SECOND DIRECTION IS lines flip a healthy run red, which wastes inves flip a run that achieved nothing into a reporte success vector sitting in shared infrastructure uses, which places it in the same class as a che OBSERVED: a guard's positive case passed alor neighbour's failures; and two back-to-back rur attempt, with the second run's results file repo plus the first run's four. A SECOND, RELATED I pointer meant to indicate the most recent run because the command used to update it follow and drops its contents inside the first run's fold folder's timestamp, so ordering the folders by Measured live: the pointer still names the first runs existing beside it, with those later runs' p the by-recency listing therefore silently return reviewers into reading an old run as current. I within the same second and observe they shar one's verdict counting the first one's outcomes the recent-run pointer still names the first. AC share a results folder or log; the recent-run po recent run; a test starts two runs in the same s independent in BOTH directions, proven real b naming and makes the test fail.
HXC-273	Bug	Queued	Medium	Several health checks treat any reply from a se that say the requested thing does not exist. WH address but do not ask the fetching tool to trea found reply is indistinguishable from a health places — two in the shared checking framewo HOW IT MANIFESTS: a service that is running on the right port with the wrong application e shape as a check that cannot fail: it consumes almost nothing. RELATIONSHIP TO OTHER WO defects closed recently where a check confirm that the RIGHT thing answered; the remedy th service could produce. WHO IS AFFECTED: an a half-broken deployment before it reaches m health check at a server that returns a not-fou check passes. ACCEPTANCE CRITERIA: all thre practical the check asserts something only the merely a successful status; and a test points ea

ATM ID	Type	Status	Severity	Description
				and asserts it fails, proven real by a paired mu behaviour and makes the test fail.
HXC-274	Bug	Queued	Medium	Starting the remote-call server makes real outb provider APIs before it finishes starting, using in the environment. WHAT HAPPENS: start-up automatic discovery; a flag passed there disabl providers but does not disable discovery itself, provider whose credential it finds and enumer across SEVEN providers in a single start-up, dr host shell rather than supplied to the process: twice, because two different credential names eight, six, five, three and two from the rest. WH local server does not expect it to contact extern serving anything; in a sandboxed or air-gapped delay start-up; and the credentials used are wh operator deliberately supplied. IT ALSO CAUSE why the server accepts connections into the op seconds before it actually begins serving them server was not speaking the expected protocol IMPORTANT NOTE FOR WHOEVER FIXES THIS wrong conclusion. The flag that disables verific outbound behaviour, and the catalogue fetche discovery routine, so an investigation that stop outbound traffic and returned a refutation. On case where the source and the behaviour genu evidence for any claim about it. ACCEPTANCE calls unless explicitly asked to, discovery is opt are taken from deliberate configuration rather begins serving promptly after it begins accepti present and asserts no outbound calls occur, p restores eager discovery and makes the test fai