

Issues_Summary

Open workable items (current_location = Issues), regenerated from the SQLite single-source-of-truth (§11.4.12).

Counts by Type × Status

Type	Status	Count
Bug	Operator-blocked	1
Bug	Queued	41
Task	Operator-blocked	1
Task	Queued	13
TOTAL		56

Items

ATM ID	Type	Status	Severity	Description
				Reported-Via: §11.4.202 reporting directive to the HelixSkills System (git@github.com:HelixDevelopment/skills.git , v1.0.0) and HelixAgent, covering ALL of its power-features fully without nothing leftout! Full coverage with ll supported test (suites) is mandatory! Extend and update all existing guides, FAQs, create stunning illustrations, graphics, all materials! Any gaps, shortcomings, weak spots, inconsistencies MUST BE detected in advance and properly tackled with comprehensive process and properly tackled with comprehensive risk-free and safe! Track all progress through the and remebering mechanisms we have! Use Git
HXC-159	Task	Queued	High	Operator What (the report, verbatim): Feature the HelixSkills System (git@github.com:HelixDevelopment/skills.git , v1.0.0) and HelixAgent, covering ALL of its power-features fully without nothing leftout! Full coverage with ll supported test (suites) is mandatory! Extend and update all existing guides, FAQs, create stunning illustrations, graphics, all materials! Any gaps, shortcomings, weak spots, inconsistencies MUST BE detected in advance and properly tackled with comprehensive process and properly tackled with comprehensive risk-free and safe! Track all progress through the and remebering mechanisms we have! Use Git

ATM ID	Type	Status	Severity	Description
				with bridge to Superpowers used for the implementation of the development mechanism!" PRE-VERIFIED ENVIRONMENT not assumed — §11.4.6): - GitHub SpecKit IS initialized, NOT yet initialized in this repo (no .specify/ and .specify/ is therefore an explicit early task, not an assumption). skills.git IS reachable via SSH. git ls-remote -p https://github.com/catalog-docs, feature/deep-research, feature/telemetry, branch/tag inventory MUST be re-derived at repository update, partial listing. - The repository is NOT currently initialized, contains no HelixDevelopment/skills entry; the repository is codex-skills -> vasic-digital/caf-codex-skills.git, not be confused with it. - Both consumers exist (module dev.helix.code) and the submodule subagent (dev.helix.agent). MANDATORY DELIVERY CONSTRUCTION for ALL phases of the incorporation, bridged to the executed subagent-driven (§11.4.70 / §11.4.20). tasks -> subtasks, with an exhaustive level of detail where the design is load-bearing). - Complete the e2e, full-automation, Challenges (vasic-digital/Challenges (HelixDevelopment/helix_qa) with autonomous concurrency/atomicity, race/deadlock, memory management, captured evidence (§11.4.5 / §11.4.69); every gap in Documentation: extend AND update every existing only new ones. Produce illustrations, graphs and materials. Four-format export discipline per §11.4.212. the main README per §11.4.212. - Risk work is a note: every gap, shortcoming, weak spot, danger detected DURING planning and each one paired with a free and safe solution (§11.4.102 systematic-deadlock before rewrite (§11.4.74 / §11.4.28): survey the catalogues on GitHub AND GitLab first; extend the duplicating, keeping them project-agnostic and incorporated as a properly-laid-out dependency level or submodules/, never a nested own-org and HelixAgent must BOTH receive the full feature, the other is an incompleteness to be tracked, not through the continuation docs, the memory/knowledge the workable-items SSoT (§12.10 / §11.4.131 / §11.4.131). 1. A SpecKit specification exists and is initialized feature, with an explicit feature inventory project repository (an enumerated list, not a sample — implementation plan exists at the stated depth mitigation. 3. Every planned capability is wired

ATM ID	Type	Status	Severity	Description
				consumers, proven by §11.4.108 runtime signals compiled. 4. The full §11.4.169 test-type matrix and HelixQA banks included. 5. All documentation produced diagrams/illustrations incorporated, Independent code review reaches a zero-finding. No capability is claimed without runtime evidence every closure (§11.4 / §11.4.1 / §11.4.123). §11.4. PROGRAM (scheduled, not yet executed): (a) Do scientific papers / open-source projects+codebases to incorporate this feature into the project (§11.4.102) of ALL data obtained to ensure zone / inconsistency / imperfection, and design one found. (c) The design MUST ALWAYS be innovative – never less. (d) MANDATORY exhaustive documents, an implementation plan down to diagrams / schemes / graphs, illustrations, SQL relevant material (§11.4.65 / §11.4.73). (e) Invest + HelixDevelopment submodules/components freely where genuinely needed while keeping reusable (§11.4.28 / §11.4.74 / §11.4.177). (f) Plan one: every supported test type, the Challenges (§11.4.27 / §11.4.169). (g) Divide the work into nano-detailed enough to drive integration / implementation directly. (h) Plan explicitly for enterprise scalability the feature has a UI/UX surface, produce full work PSD / PDF, or whatever the project mandates) at §11.4.190). (j) Plan full CodeGraph integration + synchronisation (§11.4.78 / §11.4.79 / §11.4.80). items (every detail + reference + attachment) source-of-truth (§11.4.93 / §11.4.95) + every design-related project doc/component + every connection (§11.4.148 D5) – honestly SKIPPING any absent (credentials_absent / tracker_client_absent, §11.4.148 D5) Research-doc destination (to be populated when fully_incorporate_the_helixskills_system_into_Execution model (§11.4.213 clauses 1-2): this is a multi-day research/planning/documentation work project's standing autonomous loop (§11.4.87 / §11.4.148 D5) when it claims this item, and MUST be driven to explicitly evidence-backed CLOSED terminal state be left sitting un-wired in the backlog. Affected research/planning – destination: docs/research/fully_incorporate_the_helixskills_system_into_

ATM ID	Type	Status	Severity	Description
				Reproduction / context: UNKNOWN: not state established before any fix (§11.4.102 / §11.4.14 research-doc tree exists under docs/research/ fully_incorporate_the_helixskills_system_into every enumerated weak-spot/gap/danger-zone implementation plan down to lines-of-code exist are created, and the whole effort is validated p ADDENDUM — BINDING OPERATOR REQUIRE === VERBATIM: “Make sure that SpecKit and S and with all extensons we have created deriviv not optional colour on the SpecKit mandate — work MUST compose with rather than duplica BRIDGE EXTENSION (verified to exist, 2026-07- speckit_superpowers/implementation/ — the “ Superpowers Bridge”. Eight documents, each i IMPLEMENTATION_PLAN, NANO_TASK_ENGIN TDD_INTEGRATION, CONSTITUTION_INTEGRA SEVEN-LAYER architecture: 1 Developer Work SuperSpec bridge (orchestration — github.com (discipline enforcement — speckit-community MCP (execution) 6 Helix LLM (inference — gith Distributed host cluster (llama.cpp RPC) CONST MUST BE IN SCOPE (inherited BY REFERENCE , diverges silently): constitution/skills/ (7) action reporting-workable-items, scheduled-work-qu constitution/mcp/ (2) media-validator-mcp.json plugins/ (2) helix, scheduled-work constitution system), subagent_tiering.yaml constitution/sc credential_scan_lib.sh, guard-branch-consisten guard-track-branch-label.sh, guard-work-track LOCALLY — do not re-create: .claude/skills/ 10 clarify, constitution, converge, implement, plan workflows/speckit/workflow.yml, integrations/ templates/, memory/ ADDED ACCEPTANCE CRI extension above has an explicit disposition — reason. Silence is not an answer (§11.4.118). 9. WIRED vs DESIGNED-ONLY with a captured pr evidence a layer is installed — that is precisely and reporting a documented layer as an availa namespace collision risk is addressed with a d skills system, while §11.4.164’s post_update_ho into .claude/skills/. Two systems registering int and clash handling, not discovery at runtime. .

ATM ID	Type	Status	Severity	Description
				non-own-org dependencies (§11.4.74 vendor p supply-chain exposure must be assessed, not a them. The script that brings up the supporting service directly, rather than going through the shared maintains for exactly this purpose. That comp the same way, gets the same health checking, a repeatedly. Bypassing it means this project car and has already caused one outage of its own. service startups still hide their errors, so a fail was flagged honestly in a commit message, but obligation anyone will act on, which is why it i routing the startup through the shared compo hiding so a failed start is visible. When publishing the agent component, the ho outstanding security advisories against the lib 204. That figure turns out to be unreliable in th drifts upward on its own as new advisories are 205 the next day, and 210 on a live re-query on be “the” count. It over-states the work by abou counted once for every dependency file it appe distinct advisories, and the five criticals collap two places this code is published; the other cop was never switched on there. The recorded br critical, 88 high, 99 moderate and 18 low. A ful completed, and it changes the picture substant the affected code can actually be reached in so them currently qualify. The Go vulnerability so actually release, reports our code is affected by against a container library are reachable only in the codebase refers to and that never enters remaining alerts, including three of the four c third-party project whose container definition cannot start today. Most of the rest belong to a builds. Exactly one alert touches a component defence that advisory asks for has already bee alert stays open only because the dependency related piece of housekeeping was resolved wh files of downloaded third-party library code, in behind one critical advisory, had been commit removed from the tracked files, though they re benefit of finishing this work is that the alarm
HXC-164	Task	Queued	Medium	
HXC-166	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
HXC-168	Bug	Operator-blocked	High	the product, and we regain the ability to say he carries a reachable known vulnerability. This is acceptance criterion has not been done: the de findings have not yet been upgraded or removed retire the copied-in project that carries most of The password used to reach the project database setup and container files, and those files are private accounts. Anyone who reads the project can read not introduced by recent work; the project's own it as something that must be replaced, and that depends on whether the same password is used has not been established either way and should value is already public, changing the files alone be changed everywhere it is used, and the setup configuration source rather than containing it. cannot honestly claim its credentials are protected
				A complete copy of somebody else's web application it alone accounts for one hundred and thirty-nine dependency security warnings, including three be built or run: the startup configuration point that folder, and its supporting packages have not two opposite ways at once. Because it is not built we actually run today, which is why they rank references it, the setup is broken and misleading recipe file, all one hundred and thirty-nine ways to whoever does it. Anyone reading our security distortion, since it currently buries the small number The expected outcome is an explicit decision, not properly maintain and build it, or remove it and a reference to the upstream hosted service. Before approval, the decision must be put to the operator
HXC-171	Bug	Queued	Medium	A small companion service we wrote ourselves part of the standard deployment, and it carries supporting packages, two of them rated high. Under warnings, these have no mitigating evidence in unshipped, because it demonstrably ships, and whether the affected functionality is ever exercised place in the whole review where something was which makes it far more deserving of attention three is not really an upgrade problem at all: it is switched off unless explicitly enabled, so updating protection on would leave the exposure in place
HXC-172	Task	Operator-blocked	High	

ATM ID	Type	Status	Severity	Description
				deployment running this service benefit. The e assessment of the three warnings, the two stra protective setting explicitly verified as enabled and the service rebuilt and confirmed working The cache that remembers model information where its clock was never configured. That fal the only way to build the cache always configu unexercised. On its own that is harmless. The n if someone later edits it: a plausible-looking ch current time would make every cached entry a in the past that the freshness check reads as a expiry limit. Every entry would then be consid keep serving outdated model information inde The choice is to either add a small test that pin branch entirely and document that there is on
HXC-175	Bug	Queued	Low	When a caller presents an API key, the server c ordinary text comparison. That kind of compa that does not match, so a key sharing the first f reject than one that differs immediately. An att differences across many attempts can recover guessing it outright. Network timing noise mak is low rather than urgent, but it is a genuine w cheap: use the comparison function designed f amount of time regardless of where the differe introduced by recent work.
HXC-178	Bug	Queued	Low	The check that decides which websites may op deliberately allows requests that carry no orig behaviour, because only browsers send that in legitimately omit it. It is safe today only becaus away anyone without valid credentials. The tw nothing records or enforces that coupling: if so login check while working on something unrel information would become completely ungate work here is to add a check that specifically co the connection upgrade, so the dependency is
HXC-179	Task	Queued	Medium	Three commits in the current release range do referred to a piece of data that was not added those three points gets a build error. The tip of corrected shortly afterwards. This cannot be r rewriting published history, which the project is for anyone hunting a regression by stepping three points that fail to build for a reason unre
HXC-180	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				may wrongly conclude the fault lies there. The affected points and their shared cause somewhat so the failure is recognised immediately as known from scratch. HXC-160 corrected the 21 stale references to scripts/generate_fixed_summary.sh across the added the CM-SUPERSEDED-GENERATOR-NAMING sets were deliberately left untouched and are the Constitution.md still carries 14 references across universal rules that projects other than HelixCode records such as the Phase 5 Go-binary-shim plan is a governance amendment that CONST-049 reconfigured upstream, then re-cascaded and postponed session that closed HXC-160, and a canon amendment drift than the stale wording, so the canon was 45 owned submodules carry cascaded restatement correction applied in the parent repo is deliberate because the fix names a HelixCode-specific meaning wording into decoupled, project-agnostic submodules them and violate CONST-051(B). The correct set neutral canonical wording that names the word and push it in the constitution submodule per fleet and bump pointers. Who benefits: every of the anchor text, and any agent reading a submodule summaries are produced. Acceptance: canon rewording, the change is pushed to every constitution are re-cascaded, submodule pointers are bumped SUPERSEDED-GENERATOR-NAMING both pass The main handbook that tells contributors and organised lists several folders at the top level to when only one is present, and refers to a top-level Anyone following it — a new contributor, or an orientation — will look for things that are not broken or that they have checked out the wrong of the document, which is otherwise the author. The correction is small and purely descriptive. section to match, then keep the two in step as to One of the external projects this codebase depends configuration still refers to it by its previous name hosting service silently forwards requests from forwarding is a courtesy, not a guarantee: it does project under the old name, and it is not honored
HXC-181	Task	Queued	Medium	
HXC-188	Task	Queued	Low	
HXC-189	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
				dependency. If it stops working, fetching the project's name nobody recognises, which is unusually common for the project's current canonical name so the reference to it on a redirect continuing to exist. When a running command is cut short because of a timeout, it says it was not a timeout. The check that sets the deadline is the caller's own deadline, which usually has not been reached, so the command stopped the command. Investigation showed that the command was either, because the internal signal it would complete was not cancelled, or cancellation and never a timeout, and the timeout mechanism entirely. The effect is that anyone using the command will see a failure and no indication that a time limit was reached. The bug is in the wrong place. A sibling code path that runs commands correctly by doing it inside the timer's own callback will not have this problem.
HXC-190	Bug	Queued	Low	
				When a command is cancelled or times out, the command is not cancelled, but anything it started, by signalling the whole command, so it works if the command was placed into its own sandbox. If the sandbox is enabled — the normal configuration — the sandbox is switched off, the grouping step is skipped, and the operating system to signal a group that does not exist. The error is discarded. The result is a cancelled command that is dormant, because the only configuration that could have been so no shipped path reaches it. It is worth fixing this, because it would appear only in an unusual configuration and is hard to diagnose later.
HXC-191	Bug	Queued	Low	
				Output from a running command is read line by line, and the line may be. When a line exceeds that limit the command is permanently for that stream, because the buffer is full and past it. Nothing else takes over the reading. The command is not until the operating system's own small buffer is full, and then someone to read. So one unusually long line, which is not readable output produce quite naturally, can be read. The fix is to keep draining and discarding the remainder of the command, abandoning the stream, so the producing command can continue.
HXC-192	Bug	Queued	Low	
				The handbook that describes this project's technology is one of its networking libraries. Both parts of the handbook are because a security fix required moving off the old handbook. It is deliberate; it is the handbook that is now wrong. It is a new contributor or an automated helper reading the handbook will confidently state a version that has not been found. The same document has already been found described in the handbook.
HXC-196	Task	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				additional inaccuracy makes the whole document recorded version to what is actually in use, and for this reason, so nobody later reverts it believing the version recorded. One test directory contains a duplicate of its own files appear twice in the project at two different times during a renaming exercise that moved a tree while leaving the new one. Nothing reads the nested copy, so anyone searching the project finds two results, one of which is real, and any tool that scans for duplicates blocks a useful check from being switched on: the duplicates cannot be enabled while this one exists every run. Removing it requires first confirming the nested copy is genuinely the leftover and not the original.
HXC-197	Bug	Queued	Low	
				A detector was written to catch a specific kind of error: stuck trying to hand off a result nobody is collecting. One particular description of what the worker was doing: form of the same stall — where the worker was waiting rather than a single handoff — and the detector was demonstrated rather than theorised: a deliberate test went completely undetected while the detector was running. What it does catch is the historical one, so the guard is still using the newer structure while breaking its escape strategy. To recognise both descriptions is a small change and a bug, not a defect rather than half of it.
HXC-200	Bug	Queued	Low	
				Two long-running tests — one exercising heavy load, one exercising leader election among worker nodes — report failure when the machine is under load. The failure is its own. Measured: during a full test run both fail about 10 times consecutively, one finishing in roughly a minute, the other about a twentieth. Their verdicts are therefore unreliable.
HXC-204	Bug	Queued	Medium	whether the product works. This matters most when a test run performed before a release is by definition successful: these two will mark a healthy release as broken. The test goes through. A third test with the same problem was an inconclusive run report as skipped rather than failed. The same treatment fits here. Both should wait for a result rather than assuming a fixed time is enough.
				The newly-fixed health refresh deliberately re-contacts the provider over the network, so that a slow or hanging probe. It then re-acquires the lock and only applies its changes if the changed underneath it. There is a narrow room for error: started again entirely within one probe, the change is lost.
HXC-206	Bug	Queued	Low	

ATM ID	Type	Status	Severity	Description
				applies a verdict formed against the previous i briefly reported unhealthy when it is fine, whi dangerous direction — reporting a stopped pr Worth recording that this is cheap to close rath incremented on each state change is invisible t to any published structure. An earlier note des was wrong and has been retracted. When a model provider record is created, its e shared template by reference rather than copi points at the same underlying settings, includi Nothing writes to those settings today, and the no corruption occurs at present — this was ver record it is that the protection now added to th manager cannot protect a structure shared wi legitimate, properly-locked write to a provider manager's view as well, and the locking will lo when the record is created removes the hazard Four compose entries across two end-to-end te the Slack mock and the LLM-provider mock, b Any attempt to build those stacks therefore fai which means the containerised form of these t exercised. This was surfaced while fixing an u same two services: the fix could be proven at t against a running container, precisely because produced. That gap matters beyond these two process leaves the deployed-artifact layer unpr failure the four-layer verification discipline ex Dockerfiles, or to remove the build entries if th run containerised, and then to prove the choic A blanket rule excluding log files from version evidence, because evidence transcripts are wr corrected so this cannot happen again, and the have been recovered and committed. Thirty-fi the runs exist only on the machine that produ an evidence folder that appears empty, which never run. Recovering them is not automatic, l seventy-eight megabytes and committing those deserves a deliberate decision rather than a re keep what genuinely substantiates a closed ite ones rather than storing them whole, and reco why, so the gap is visible rather than silent.
HXC-207	Bug	Queued	Low	
HXC-208	Bug	Queued	Low	
HXC-216	Task	Queued	Low	

ATM ID	Type	Status	Severity	Description
HXC-219	Task	Queued	Medium	A recent repair changed how the release gate of evidence really belongs to the change it claims to commit on a labelled line rather than merely having somewhere in the text. That new requirement now have to follow it: the guide that authors read before the new labelled field, so a contributor following evidence the gate will reject without understanding describing the gate itself has printable and we saw the page was last edited, so all three versions disagree to describe. Fixing this means updating the author's citation field and regenerating the exported code for who benefit are contributors recording evidence that have no accurate written description of the rule. They can follow the written guide and produce evidence without attempt. When we mark a piece of work finished, we recheck our works. Nobody ever checked that those pointers to different kinds of rot accumulated silently and without report. First, some pointers hold a whole paragraph where there is no way for a machine to follow them. Second, they were never created because the person doing the work never means the proof exists but nobody can find it for ordinary progress updates rather than to the author's notes look like formal proof. The effect is that a pointer backed in every summary and dashboard which is the way anyone finds out is by manually hunting for it. To make the completion pointer a checked field: if it exists, it must be attached only to real completion. It must be refused at the moment someone tries to add it later. This protects everyone who relies on our claim of proof that cannot be produced on demand.
HXC-224	Bug	Queued	High	A single bundled helper component inside the codebase roughly 208 security advisories reported against the total — and it entered the codebase without having established what it does, whether we wrote it or whether it actually uses it, or whether removing it would be safe. The answer to the headline advisory count is misleading because it carries enormous risk, when in reality the greatest risk is not even be part of what we deliver. The work is done for the purpose, whether it is reachable from anything. The decision about it — keep and maintain, upgrade or remove — is a record which was chosen and why. The decision is not a record.
HXC-225	Task	Queued	Medium	A single bundled helper component inside the codebase roughly 208 security advisories reported against the total — and it entered the codebase without having established what it does, whether we wrote it or whether it actually uses it, or whether removing it would be safe. The answer to the headline advisory count is misleading because it carries enormous risk, when in reality the greatest risk is not even be part of what we deliver. The work is done for the purpose, whether it is reachable from anything. The decision about it — keep and maintain, upgrade or remove — is a record which was chosen and why. The decision is not a record.

ATM ID	Type	Status	Severity	Description
				advisories, because upgrading dependencies in effort spent on nothing. The agent project is mirrored to two hosting and was only ever switched on for one of them. Qu hundred and ten open alerts while the other r one the main project is configured to pull from of the component we actually consume is show hundred findings sit unseen on the sibling cop the two mirrors hold the same commits. What watched. This is dangerous precisely because t findings is indistinguishable from a report of z the second case is what we have. The work is t from, confirm both then report the same figur the two so a future divergence is noticed rather statement about this component's vulnerability from. A design document committed forty-eight days appears to be a genuine access key for an exte placeholder. The section is headed as being co and described as already configured, and the v placeholder would have: it is fifty-one character characters, and contains no example or to-be-f the main branch and has therefore been publi day it was written. Anyone with read access to fork made at any point in those forty-eight day now would remove it from the current revisio history is forbidden, and rewriting would in ar The only action that actually withdraws the ke replacement, which requires access to the pro within the repository. Until that happens the k revoked, the document should be edited to ref of quoting its contents, and a check added that this shape. The multi-provider routing gateway starts in i release mode, and says so in its own startup lo this should be switched for production use. De prints every registered route at startup, produ request, and in this framework is documented extra output and the internal detail it exposes launched by the system supervisor as a long-ru the production case the warning is about. Noth failing, which is why this has gone unnoticed, .
HXC-226	Bug	Queued	Medium	
HXC-227	Bug	Queued	—	
HXC-229	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
HXC-231	Bug	Queued	—	logging at development verbosity and revealing the service unit is the natural release mode when the service runs under the production unit ever starts in debug mode again. The HelixLLM gateway downloads large AI models before it starts accepting network connections. During this time, the service as running and healthy, but every so often on 2026-08-06 this window lasted thirty-seven minutes before actually began listening. Anyone checking the status of everything is fine while the service is in fact unresponsive. It depends on the gateway will fail for the whole service because the download depends on an external service. If one of the two downloads timed out after roughly 30 minutes that site were unreachable the gateway might be stuck accepting connections first and download mode later, in an honest not-ready state while the downloads are in progress. HelixAgent expects two tables to exist in its database: agent_instances, but neither of them has ever existed anywhere in the project that creates them. As a result, it locks once every minute, the database refuses to accept the lock and an error is written to the log. Verified on 2026-08-06 directly: neither table exists in any of them, and the error occurs every minute period plus a further hundred and thirty seconds, otherwise runs normally, so nothing visibly broken. It buries genuine problems in the log and means that the service to support is silently doing nothing at all. Fixing this by creating these tables were intended to contain to write the data I guessed. When the HelixLLM gateway starts it tries to launch a program called llama-server and fails because the llama-server machine confirms the situation is worse than a llama-server exists anywhere on the host, so the service is pointing the service at the right directory. The service continues and carries on without it, so nothing appears broken. It is absent for the entire life of the service. This is a problem I had just spent thirty-seven minutes downloading llama-server could serve it; the download succeeded and the service live on 2026-08-06 during a full platform boot. The fix is installing the program, then making sure the service is aware of whether the gateway should be downloading llama-server would consume them is not present.
HXC-232	Bug	Queued	—	
HXC-233	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
HXC-234	Bug	Queued	—	On startup HelixAgent tries to build and launch containers known as MCP servers, using the container engine. If the containers inside the container fail during their dependency installation, the operation is abandoned. HelixAgent records a list of these servers and then continues running normally. Nothing obviously breaks. The real effect is that the servers we were meant to provide are silently missing, and this is easy for nobody to notice. Observed live on 2023-10-20. The build was attempted twice and failed both times. A large container image each time, which also means repairing the dependency installation inside the container.
HXC-235	Bug	Queued	—	The HelixLLM gateway provides a feature that allows for exact wording. On startup it reports that no response for this purpose, so it has fallen back to a simple hashing strategy. It states plainly that this fallback does not capture the full meaning and will be significantly degraded. The important part is that it will off or return an error: it keeps answering even though the response is effectively arbitrary rather than relevant. Any deviation from the responses that it is not working properly is a full platform boot. Resolving this means pointing to a genuine embedding provider, and deciding what level of degradation is acceptable behaviour or whether the feature should be disabled.
HXC-236	Bug	Queued	—	The standard library routine that splits a web address into components differently in two parts of this codebase, using the same logic. In the main application it rejects an address whose host component contains a colon. In the containers component it quietly accepts the same address. The cause is that each component declared its own logic. The newer version tightened this routine; the containers component, so it keeps the older, permissive behaviour. Reasoning about address handling is only valid if the logic is consistent, which has already caused one incorrect conclusion. Second, and more seriously, the day someone releases a newer language version, addresses it accepts that currently fail will appear as unreachable services rather than a configuration change. The work is to make the difference explicit, to document deliberately which behaviour each component has, and ensure the reader will find it, and add a check that fails if the logic is inconsistent, noticing.
HXC-238	Bug	Queued	Low	The safety check that stops half-finished experiments from throwing an exception for captured proof files, which would otherwise satisfy the old exception rule and so could not be handled. The narrow: the file must sit in the evidence folder.

ATM ID	Type	Status	Severity	Description
				saved as non-runnable, and not begin with the Together those make the file inert in normal use. The bug is a deliberate pointing a program interpreter at a file which no file property can stop. The practical result is a crash rather than an accident, and the alternative was to keep a proof it is required to keep. Recording it so the bug is not than absolute, and so anyone widening it later. When HelixQA runs a test bank that needs to stop, it gives out the wrong piece of it. The HelixCode server gives out different values: a real access pass (the field named 'access' for reference for the session (the field named 'session' for proof of identity on later requests. HelixQA is supposed to so it hands the server a value the server does not expect, or expired token'. The effect is that every test that fails even when the feature it is testing works perfectly. The worker-management bank scored 6 passed and 1 failed when passed and 1 failed when pointed at the correct value. The reason nobody noticed is that HelixQA's own implementation of the access pass in the field HelixQA expects, so it fails against the real product. Anyone reading the results of the parts of the product are broken. One of the HelixCode test banks signs up a new account and uses the very same account name, 'hxc_gen_e2e'. The first run succeeds and the check passes. Every run after that, the database, the server correctly refuses with a failure recorded as a failure forever after. This was observed when a run of the bank passed that step, and a second run was visible in the database timestamped to the moment the first works under require that automated checks can't be the same answer, cleaning up after themselves. HelixQA quietly poisons its own results, so a permanent record of do with whether the product works. The fix is to reset each run, or to remove the account it created with the same name. The checks that confirm HelixCode can generate a valid response inside the server's reply, and that comparison of the two. Two checks in the generate-e2e bank ask for a response with different capitalisation, so they are marked as failed. behaved correctly. The most serious case asks for a reply for 'hello' in small letters; the AI answered with a working feature was reported as broken. This was the server: asking it to reply with an invented word, and it what seventeen times three is returned fifty-
HXC-239	Bug	Queued	—	
HXC-240	Bug	Queued	—	
HXC-241	Bug	Queued	—	

ATM ID	Type	Status	Severity	Description
				which proves real AI generation is running and a reply for ‘authorization’ in small letters while capital A. Expecting exact capitalisation of free nature. The consequence is false alarms that h time. Five checks covering the built-in screenshot against the running HelixCode server, and even server answers ‘QA engine is disabled’ and rep server is being honest here rather than misbeh whole advertised area of the product is switch cannot be used by anyone or confirmed as wor listing the available screenshot engines, listing status of a session, starting a new session, and with nothing to run. Because the feature is off, whether the underlying code behind it function feature is meant to be available, and if so switch behaviour is actually proven; if it is meant to s reporting failures. Several of the automated test collections that c about what a correct answer looks like. They s and record a pass regardless of the reply. This collection was deliberately aimed at the wrong an error for every request, and it reported bot passes against a service known to be broken ca one, so its green result carries no information failing test, because a suite that cannot fail is v where there is none, and it will keep producin every check an explicit expectation — the statu answer must contain — and then prove each o something known to be wrong and confirming the same run: four apparent failures in the wo harness looking for the wrong field name in th and that mismatch should be corrected so real One blocked work item (HXC-172) has its expla has that same explanation written into its own to agree. The tracker rebuilds those stored exp time it regenerates its records, so an explanati the text is discarded the next time the tracker happens today: a regenerate-and-reread cycle then correctly reports the item as blocked-with that explanation is what makes a blocked item what was already tried, and which decision w
HXC-242	Bug	Queued	—	
HXC-243	Bug	Queued	—	
HXC-245	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				decision waiting on a person into an item nobo keeping the context needed to make the blocki why the item stalled. The expected outcome is explanation again, the stored and written copi preserves it — proven by the tracker’s validati A complete run of the agent runtime’s automa 12 groups failing, and 74 individual test failure than assumed, because the early evidence poin than at the product. Three patterns account fo the in-memory cache service because they loo container while this machine publishes it on a running service to answer a web request and f and answering correctly when asked by hand; against fixed limits and overshoot them by abou jobs and a service rebuild were competing for product is healthy, and none of it proves the pr parallel investigations are now establishing, ea rather than a plausible story. This matters beca reasons is as damaging as one that passes whil everyone to stop trusting the result, and both h operator benefits from a suite whose red mean signal they can act on. The expected outcome i evidence as a genuine product defect, a test de contention artifact, each real defect fixed with returning to a trustworthy state. Two separate parts of the system have both be and neither knows about the other. The agent decides which service gets which port number itself. But a different component, the model ve names port 8100, and because the verifier star agent runtime is launched from a settings file t never even tries to claim the one it was allocat port 8100 for the agent runtime and are answer recognise the requests and rejects them with a runtime as broken when it is running perfectly because it makes a large part of the test suite s names, and because the same collision would n real. Whoever operates or deploys this benefit failing, and a genuinely failing one would look a single place that decides port allocation for e claiming the same number, the agent runtimes and the tests then passing or failing for reason
HXC-246	Bug	Queued	High	
HXC-247	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				Later evidence settles which of the two sides is deployment is not misconfigured at all. The machine file that starts it on this machine - documents it as an agreed meeting point between the verifier and the gateway looks there by default, and, in as many cases, runs on the other address here. The choice was made down at deployment time, in the verifier's favor. It is clearer than first described: the agent runtime at deployment date, still claiming an address that a different, newer, with a note that the old address was abandoned but still in use. A contributing cause is that the verifier is so the one mechanism meant to prevent two components not arbitrate a component it had never been to the unchanged - the tests inherit the stale claim and the repair direction is now settled rather than the address it genuinely serves, give the verifier to recur, and point the test defaults at the real address establishes which side is stale and why; it does not safe to execute, because that entry is consumed and change still needs its impact review before it is late. A cleanup step in the automated test suite can be the only thing preventing that is a coincidence to stop any containers they started. To avoid stale first check whether the agent service is still running <i>something</i> answers on network port 8100, with the On this machine port 8100 is held by a different the check says yes, and the shutdown is skipped different port and is never actually consulted. If late, or moved, that same check would say no and a live platform out from under whoever was using also untrue: it reports that the agent is still running all. This matters to anyone running the test suite deployment, which is the normal case here. The the identity of what is answering - not merely the makes its decision on the real condition rather than logs is true. When a user asks the system to write out a real command-line coding assistant, the file it produces service. The generator fills in a network address than to the agent runtime the assistant is supposed requests to a component that does not understand reply. From the user's point of view the assistant
HXC-248	Bug	Queued	High	
HXC-250	Bug	Queued	High	

ATM ID	Type	Status	Severity	Description
				error explaining that the address in their branch was written. Of the forty-eight assistants the problem affected; a fix landed earlier covers only two of them, all the rest still writes the wrong address today. This is in the current set, because it is not an internal error, but handed to a user with an instruction to use it. For newcomers the feature exists to help: someone following the documented path, and getting a message. The expected outcome is that every generated config works at runtime genuinely serves, that all forty-eight do more than two, and that a generated config works without editing anything by hand. One of our own services could not start while others both had been told to use the same network address. A conventional default that the wider industry uses is precisely what makes it the number most likely to be on our own machine it was taken, by a supporting data set. The result was that the service refused to start, and the client dialling that same well-known address received a message which completed a perfectly healthy-looking connection. Because the wrong service was alive rather than the right one, something is unreachable did not skip: they ran our service, which had never been running at all. Twenty test failures that had been attributed to the service in the project's central address book so far, never silently collide again, and it now fails immediately rather than starting somewhere unexpected. Very misleading: the startup message printed for the wrong platform, so anyone reading the log is told the wrong platform benefits, and the expected outcome is that the rest of the stack and that every message it receives is bound. Tests in one area of the codebase interfere with others, share a single piece of network machinery, so one test has to do with the code it is testing. When each component is given its own timeout setting but no connection, the library quietly falls back to one shared handler. In this area start roughly a hundred and forty-five tests, around five hundred cases to run simultaneously, and down at the end of its own test, politely tells the others and because the handler is shared, it also serves the rest of using. The result is a failure that appears in
HXC-251	Bug	Queued	Medium	
HXC-254	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
				and vanishes when the tests are run one at a time in the test setup rather than in the product. It repeats one run in ten fails this way. Developers are then unpredictably teaches everyone to re-run it rather than comfortably inside that noise. The expected outcome of own connection handler so shutting one preter test, and the suite becomes trustworthy enough to be genuinely wrong. Two things are left over in the scripts that run the tests: needs a decision from the operator rather than a simple problem: the end-of-run summary prints “Intermediate” whenever the log file merely exists, regardless of whether which tests genuinely failed still shows a reassurance, the most likely to read. The second is deliberate and is written to note test failures as warnings and cautions, and status even when real tests have failed, which is a problem that blocks a release. That behaviour was chosen over the option to abort a long run and lose the rest of the information gathered between gathering everything in one pass and the information broken. Because it is a deliberate posture rather than an operator decision about how the project wants to proceed, it can be altered silently by whoever happens to touch the scripts. The runs to judge release readiness benefits from showing only that the script reached the end. The expected result reflects the real result rather than the presence of a test, an explicit, recorded choice about whether these tests are The tool that regenerates the project’s tracker code happens to be run from unless it is explicitly told to regenerate one optional setting quietly produces a second set of files while the real ones stay out of date. The tool’s code is optional, which is what invites the omission, and it is in a folder rather than the folder where this project was found the hard way during routine work: running the optional setting created sixteen stray copies at the top of one directory down, kept their old contents. The tool, if someone can look at freshly written files, see that the records are current when the versions under version control the tracker documents by hand is exposed, when the automated helpers working outside the standard boundaries are bounded and that bounds the severity: every version of the destination explicitly, the tool prints the full path of the version visible to a careful reader, the stray copies show
HXC-255	Task	Queued	Medium	
HXC-256	Bug	Queued	Low	

ATM ID	Type	Status	Severity	Description
				separate consistency check would eventually n What is missing is the tool itself objecting. The documents without naming a destination either or refuses and says so, instead of succeeding so caught this does not exist today: the existing gr supplies the destination explicitly, so it can nev out - adding that second case is the specified fo The automated check that guards our workabl actually measures, so a reader who trusts its w computed. Internally it names the thing it insp written description states that it proves the cor it never consults version control at all - it simp folder, which may contain changes that have n genuinely proves is narrower and still useful: right now agree with each other. The gap was the check reported success while the database yet been committed - exactly the situation its o stating plainly is that the underlying behaviou of the pre-release sweep, where the point is to commit is self-consistent, so inspecting the wo there is a documented reason for the copy - op as a side effect. On that reading the repair is th Everyone who reads that check's verdict befor the difference between "what is on your disk is correct" is exactly the difference that matters v name has already spread: the description of th repeats the same claim. The expected outcome and the sentence it prints all describe only wh guarantee is wanted it is added deliberately as control. The test that would have caught this d disagreement between the documents and the between the committed copy and the working adding that case is the specified follow-up. Our workable-items database keeps a small bo synchronisation ran between the database and That note is only ever written by one of the tw documents into the database - and the value it that actually ran. The path that goes the other database, never updates the note at all. The pr regenerated the note still describes an older ru reading it to find out what happened most rec matters because deciding the direction of a syn
HXC-257	Bug	Queued	Medium	
HXC-258	Bug	Queued	Medium	

ATM ID	Type	Status	Severity	Description
HXC-259	Task	Queued	Low	can destroy records if it is made backwards, and judgement while being incapable of supporting into-database run dated the eleventh of July with database-into-documents regeneration on the shared tooling that performs the synchronisation genuinely ran on both paths rather than assuming hand would be undone by the next run and is Every workable item is supposed to accumulate its status changed - opened, reopened, closed - what evidence justified it. A recent repair note no such trail. Measurement across the whole database to those two: eighty-one distinct records covering history entries at all, and all but one of them are filed away. In other words the majority of our account of how it reached that state, which we understand why something was reopened, or the cause is that items created by direct database creation command skip the step that opens the trail before the trail existed never had one. This is a filling: inventing dates and authors for events a trail that is worse than an absent one because instead is a decision on how to mark these records that prevents any newly created item from joining The Kubernetes manifests for HelixAgent told 9090, and they set an environment variable called was true. No Go source in the repository reads the port the server listens on; this was proven bound :8112 while the real variable HELIXAGENT 9090 is a port no HelixAgent process ever binds internal/ports registry (HelixAgentGRPC, core 1 by running the built cmd/grpc-server binary and The practical effect is that anyone deploying Helix manifests would get a Deployment exposing core gRPC traffic to targetPort 9090, and a NetworkPolicy artifacts agreeing with each other and all pointing port the server would actually bind was blocked number because the manifest LOOKED configured and edited GRPC_PORT would see no change with This item covers correcting all four manifests to service.yaml, networkpolicy.yaml, configmap.yaml the dead variable with the one the server honours limitation: the image this Deployment runs is based
HXC-260	Bug	Queued	High	The Kubernetes manifests for HelixAgent told 9090, and they set an environment variable called was true. No Go source in the repository reads the port the server listens on; this was proven bound :8112 while the real variable HELIXAGENT 9090 is a port no HelixAgent process ever binds internal/ports registry (HelixAgentGRPC, core 1 by running the built cmd/grpc-server binary and The practical effect is that anyone deploying Helix manifests would get a Deployment exposing core gRPC traffic to targetPort 9090, and a NetworkPolicy artifacts agreeing with each other and all pointing port the server would actually bind was blocked number because the manifest LOOKED configured and edited GRPC_PORT would see no change with This item covers correcting all four manifests to service.yaml, networkpolicy.yaml, configmap.yaml the dead variable with the one the server honours limitation: the image this Deployment runs is based

ATM ID	Type	Status	Severity	Description
				no gRPC server at all (that is a separate binary, correct but INERT until a grpc-server contained corrected manifests make a one-line change in checking out the pre-fix tree and running the p RED_MODE=1, which asserts the dead variable Acceptance criteria: internal/ports/published_g unset (asserting every manifest names the reg outside a comment) and fails on the pre-fix tre The challenge script that is supposed to prove structurally incapable of failing, and it has bee existed. Two independent faults compounded. environment variable named HELIXAGENT_G spelling that no Go source reads (the server ho number belonging to no service in the project: registry's 8112, not the HTTP port 7061. Second recorded a successful assertion whether the po check recorded true when the port was listeni when it was not; the unary-call, streaming-call The two faults together mean the challenge dia ever bound, found nothing there, and then cer most damaging class of defect in the codebase — it is an automated test actively producing fa anti-bluff covenant exists to prevent, and it ma reading the suite. It also violated the challenge framework already ships a record_skip helper record_assertion with status true to mask a mi not used. This item covers resolving the port th default, consolidating the four duplicated prob interpreted in exactly one place, making absen OK ticket rather than success, making a real fa two-state final verdict (which reported PASSED when nothing had been verified) with a three- assertion could be verified. Reproduce with th protocol_grpc_fail_open_guard.sh run with RE and releases a port to prove it is free, drives th and confirms assertions are recorded PASSED guard with RED_MODE=0 reports zero PASSED every probe SKIPPED and SKIP-OK tagged, and PASSED assertion when a real gRPC server is li
HXC-261	Bug	Queued	Critical	