

03_open_clarifications

Open Clarifications — operator-gated ambiguities (§11.4.6 / §11.4.66 / §11.4.105)

These are names/targets the research could NOT resolve to a single unambiguous real API/spec. Per §11.4.6 I will NOT invent an endpoint. Per §11.4.101 these are PARKED (they do not block the running research); they will be asked as ONE batched question at the Phase-R → P synthesis boundary. Confirmed items proceed regardless.

#	Name (as given)	Research finding	OPERATOR RESOLUTION (2026-07-06)
C1	“GPT SOL”	OpenAI GPT-5.6 “Sol” (rides existing OpenAI adapter) vs Upstage Solar (new adapter).	OpenAI GPT-5.6 “Sol” — NO new adapter; add model IDs to the existing OpenAI provider when it reaches GA (track as preview-pending).
C2	“Google OKF”	Google Cloud Open Knowledge Format v0.1 — a knowledge-packaging spec, not a model/provider.	RAG/Skills knowledge format via MCP resources — implement OKF as the on-disk knowledge format served through an MCP resources server (context layer, not provider layer).
C3	“ACP”	Zed Agent Client Protocol (editor ↔ agent) vs IBM/BeeAI ACP vs Google A2A (agent ↔ agent).	Google A2A (Agent-to-Agent) — implement A2A interop, NOT Zed ACP. (Note: A2A is agent ↔ agent, so “editor auto-recognition” is served separately via the /v1/models + MCP server/discover

#	Name (as given)	Research finding	OPERATOR RESOLUTION (2026-07-06)
			capability surface; A2A adds cross-agent interop.)
C4	Subquadratic (SubQ)	Private beta only — no public base URL/auth.	BLOCKED-until-GA — tracked as a deferred item; no adapter work possible until a public API (or operator-supplied beta credentials) exists.

Resolved WITHOUT needing the operator (recorded for transparency)

- **Poe** → <https://api.poe.com/v1> (OpenAI-compatible aggregator) — adapter-ready.
- **Perplexity** → <https://api.perplexity.ai> Sonar models — adapter-ready.
- **Sakana Fugu** → <https://api.sakana.ai/v1> (fugu, fugu-ultra) — adapter-ready.
- **Xiaomi MiMo** → <https://api.xiaomimimo.com/v1> (+ Anthropic surface) mimo-v2.5-pro — adapter-ready.
- **Tencent / Yuanbao** → **Yuanbao is the consumer app; the API target is Hunyuan** <https://api.hunyuan.cloud.tencent.com/v1> (hunyuan-t1, hunyuan-t1-vision) — adapter-ready.
- **Qwythos 9B** → [empero-ai/Qwythos-9B-...](#) (Qwen3.5-9B, 1M ctx, Apache-2.0) — **self-host** via the existing local/Ollama/HF path, no new hosted adapter.
- **GOT family** → **GOT-OCR2.0**, an OCR model (not a chat LLM) — belongs to the OCR capability (stream 07), self-host.
- **High-value additions to also implement** (all confirmed public APIs): xAI Grok, Moonshot Kimi, Zhipu GLM, Fireworks, DeepInfra, Novita, AI21 Jamba, Reka (+ verify Hyperbolic/Baseten base URLs at build).

Protocol decisions recorded (stream 05, ready)

- **MCP** — official Go SDK now exists (github.com/modelcontextprotocol/go-sdk, v1.6.1, maintained with Google) → closes the old “no Go SDK” gap. Design the gateway **stateless-first** (2026-07-28 RC drops the session handshake). Gateway = MCP server (expose Helix tools) + MCP host (consume CodeGraph/OpenDesign).
- **Capability advertisement (CONST-040)** — flags sourced from LLMsVerifier, surfaced on 3 channels: `/v1/models[].capabilities` (HelixCode), MCP

server/discover (Claude Code), ACP initialize (editors). Tool-calling normalized to one canonical JSON-Schema + grammar-constrained decoding per local-model parser (anti-bluff).

- **OKF** — used as the on-disk RAG/Skills knowledge format served via an MCP resources server (pending C2 confirmation).