

06_providers_coverage

Provider-Coverage Research — LLMsVerifier + HelixAgent

Scope: Deep multi-angle web research to enable an adapter to be written for every provider/model the operator named, plus the highest-value additional providers. **Grounding fact:** HelixAgent already implements ~45 provider adapters under `submodules/helix_agent/internal/llm/providers/*` (openai, anthropic, claude, gemini, deepseek, groq, mistral, ollama, lmstudio, cerebras, chutes, together, nvidia, sambanova, openrouter, huggingface, helixllm, vertex, azure, cohere, ...). LLMsVerifier is the single source of truth for model/provider metadata (CONST-036/037). **Author:** T1/main deep-research subagent. **Access date for ALL citations:** 2026-07-06. **Anti-bluff (§11.4.6):** nothing below was invented. Any endpoint/model I could not confirm from an official/authoritative source is marked UNCONFIRMED: or NEEDS-CLARIFICATION.

Revision: 1 · **Last modified:** 2026-07-06 · **Status summary:** initial research pass.

Legend

- **Status** = confidence of the provider having a real, public, callable API:
 - CONFIRMED — public API + base URL verified from an official/authoritative source.
 - UNCONFIRMED — provider/model is real, but a specific API detail (base URL, auth, compat) could NOT be verified from an official source this pass.
 - NEEDS-CLARIFICATION — the operator's name is ambiguous or does not map to a callable LLM API; operator input required before an adapter can be scoped.
 - **OAI** = OpenAI-compatible /v1/chat/completions. **Anth** = Anthropic-compatible /v1/messages.
-

GROUP A — CONFIRMED, public API, adapter-ready (operator-named targets)

Provider	Base URL	Auth	OAI	Anth
Poe (Quora)	https://api.poe.com/v1 (OpenAI-compatible interface)	Bearer API key (Poe account)	Yes	Partial (proxies Anthropic models <i>through</i> OAI iface)
Perplexity (Sonar)	https://api.perplexity.ai	Bearer API key	Yes	No
Sakana AI (Fugu)	https://api.sakana.ai/v1 (per-account base URL shown in console)	Bearer API key (console.sakana.ai)	Yes	No (native OAI; also supports Responses API)
Xiaomi (MiMo)	OAI: https://api.xiaomimimo.com/v1 · Anthropic: https://api.xiaomimimo.com/anthropic	Bearer API key	Yes	Yes
Tencent Hunyuan	https://api.hunyuan.cloud.tencent.com/v1	Tencent Cloud API key (Bearer)	Yes	No (OAI- only verified)

Provider	Base URL	Auth	OAI	Anth

Notes on Group A operator targets

- **Poe** is an **aggregator/gateway** (like OpenRouter), not a first-party model house. Billing is Poe **points** tied to subscription tiers (\$4 . 99 → \$249 . 99/mo) + \$30/1M extra points — NOT a per-token API key with usage billing. An adapter treats it as an OpenAI-compatible gateway; model discovery returns 100+ bot IDs.
- **Tencent Yuanbao** = the **consumer assistant app** (first released 2024-05-30) front-ending Hunyuan; it is **NOT a separate developer API**. The programmatic surface is the **Hunyuan** API above. A community “yuanbao-free-api” reverse-proxy exists ([chenwr727/yuanbao-free-api](#)) but is unofficial and **MUST NOT** be used as an official adapter target. **Adapter target** = **Hunyuan**.

GROUP B — NEEDS-CLARIFICATION (operator name ambiguous / not a callable LLM API)

Operator term	What it actually is (verified)	Callable LLM API?	Recommended action	Source
Google “OKF”	Open Knowledge Format — a Google Cloud open specification (markdown + YAML frontmatter) for representing agent context/ knowledge, published	No	Do not scope a model adapter. OKF is relevant to HelixAgent’s <i>context/ knowledge</i> layer, not the provider layer. If the operator meant a Google <i>model</i> API, the target is Gemini (already	Google Cloud Tech (X) · explainx OKF

Operator term	What it actually is (verified)	Callable LLM API?	Recommended action	Source
	2026-06-12. It is a <i>data/metadata format</i> , NOT a model and NOT an LLM provider/API.		implemented) / Vertex (already implemented). Operator clarification required.	
“GPT SOL”	Most-probable match: OpenAI GPT-5.6 “Sol” (Sol/Terra/Luna family, announced 2026-06-26) — limited preview to select partners only , no broad/public API yet. (Alt reading “Solar” = Upstage Solar Pro/Mini, a <i>different</i> vendor with its own API.)	Not yet public (Sol) / Yes (if Upstage Solar)	Clarify which is meant. If GPT-5.6 Sol : it will be served via the existing OpenAI adapter once GA — no new adapter, just a model-id add. If Upstage Solar : that is a distinct provider needing its own adapter.	DataCamp GPT-5.6 Sol · Upstage Solar
GOT family	GOT-OCR2.0 — open-source OCR vision-text model (580M params, VitDet encoder + Qwen-0.5B	Self-host only (no first-party hosted API)	Adapter is feasible ONLY as a self-hosted / HF-inference target (like the existing huggingface/ollama path),	HF GOT-OCR2 docs · arXiv 2409.01704 · GitHub Ucas-HaoranWei/GOT-OCR2.0

Operator term	What it actually is (verified)	Callable LLM API?	Recommended action	Source
	decoder, 8K ctx), Apache-style research release by Haoran Wei et al.		and it is an OCR model, not a general chat LLM. Confirm the operator wants OCR coverage.	

GROUP C — CONFIRMED provider, but distribution model needs an operator decision

Provider	What it is	API shape	Status	Source
Subquadratic (SubQ)	Miami frontier-AI startup (founded 2026, CEO Justin Dangel, CTO Alexander Whedon). Model SubQ 1M-Preview — “fully subquadratic” SSA attention, 12M-token research context.	Three products in private beta : an API exposing the full context window, SubQ Code CLI, SubQ Search. Base URL / auth NOT publicly documented yet.	UNCONFIRMED (real company + model; no public API endpoint verified — private beta only; independent verification of claims still disputed)	subq.ai/introducing-subq · MIT Tech Review · VentureBeat

Recommendation: Subquadratic cannot be adapter-scoped until it exposes a public API with a documented base URL/auth. Track for GA; the API is expected OpenAI-shaped per the beta description but this is **UNCONFIRMED**.

GROUP D — CONFIRMED single models to run via existing self-host paths (no new hosted adapter)

Model	What it is	How to call	Status	Source
Qwythos 9B (empero-ai / Qwythos-9B-Claude-Mythos-5-1M)	Full-parameter reasoning model on a Qwen3.5-9B base, 1M-token ctx, native Qwen function-calling, Apache-2.0, released by Empero AI June 2026. No first-party hosted API.	Self-host via vLLM / SGLang / Transformers ; GGUF quant runs on llama.cpp / Ollama / LM Studio / Jan / KoboldCpp → reachable through HelixAgent’s existing ollama / lmstudio / huggingface adapters .	CONFIRMED (as a downloadable model; not a hosted provider)	HF empero-ai / Qwythos-9B · deployment guide

Recommendation: No new adapter needed — register Qwythos-9B as a model behind the Ollama/LM Studio/HF-local providers already implemented. Confirm the “Claude-Mythos / uncensored” provenance is acceptable per policy before shipping.

GROUP E — Highest-value ADDITIONAL providers to add (not in the ~45 list) — all CONFIRMED

Provider	Base URL	Auth	OAI	Anth	Models endpoint
xAI (Grok)	https://api.x.ai/v1	Bearer API key	Yes	Yes (Anthropic SDK supported)	GET /v1/models

Provider	Base URL	Auth	OAI	Anth	Models endpoint
Moonshot AI (Kimi)	https://api.moonshot.ai/v1	Bearer API key	Yes	Yes (OpenAI <i>and</i> Anthropic compatible)	OAI /v1/ models
Zhipu AI (GLM / Z.ai)	OAI: https://open.bigmodel.cn/api/paas/v4/ · Anthropic: https://open.bigmodel.cn/api/anthropic (intl alias https://api.z.ai/...)	Bearer API key	Yes	Yes	OAI v4 models
Fireworks AI	https://api.fireworks.ai/inference/v1	Bearer API key	Yes	No	OAI /v1/ models
DeepInfra	https://api.deepinfra.com/v1/openai	Bearer API key	Yes	No	OAI /v1/ openai/ models
Novita AI	https://api.novita.ai/v3/		Yes	No	OAI models

Provider	Base URL	Auth	OAI	Anth	Models endpoint
	openai (legacy .../openai)	Bearer API key			
AI21 Labs	https://api.ai21.com/studio/v1 (also on Bedrock/Azure Foundry)	Bearer API key	Partial (native REST + OAI-ish)	No	REST /v1/*
Reka AI	https://api.reka.ai/v1 (v0 legacy https://api.reka.ai/)	Bearer API key	Yes (fully OAI-compatible)	No	OAI models
Hyperbolic	https://api.hyperbolic.xyz/v1	Bearer API key	Yes	No	OAI models
Baseten	Model-specific / dedicated-deploy endpoints (Model APIs https://inference.baseten.co/v1 per model)	Bearer API key	Yes (OAI-compatible Model APIs)	No	per-deployment

Recommendation summary — what to build

1. **Add first (CONFIRMED, mainstream, OAI+Anth): xAI Grok, Moonshot Kimi, Zhipu GLM** — all three are dual OpenAI+Anthropic compatible with documented base URLs and are widely used. Highest ROI.
 2. **Add next (CONFIRMED open-model gateways, OAI): Fireworks, DeepInfra, Novita** — thin OpenAI-compatible adapters, large catalogs, cheap.
 3. **Add (CONFIRMED first-party model houses): AI21 Jamba, Reka** — smaller but distinctive (long-context SSM; multimodal).
 4. **Operator-named CONFIRMED, adapter-ready: Perplexity Sonar, Xiaomi MiMo, Tencent Hunyuan, Sakana Fugu, Poe** (as an OAI gateway). Note Poe = points billing; Yuanbao = the Hunyuan API.
 5. **Self-host path, no new hosted adapter: Qwythos 9B** (Ollama/LM Studio/HF), **GOT-OCR2.0** (HF/self-host, OCR-specific).
 6. **Verify base URL before coding: Hyperbolic, Baseten** (3rd-party sourced this pass).
 7. **Blocked pending public API: Subquadratic SubQ** (private beta, no public endpoint).
 8. **Not a provider — clarify with operator: Google “OKF”** (a knowledge *format*, not a model), **“GPT SOL”** (likely OpenAI GPT-5.6 Sol → existing OpenAI adapter once GA; or Upstage Solar → separate adapter).
-

Sources verified 2026-07-06

- Poe/Quora API: <https://techcrunch.com/2025/07/31/quoras-poe-is-releasing-an-api-for-developers-to-easily-access-a-boquet-of-models/> , <https://hyper.ai/en/stories/cd3cd6a39b8b109839aa301d2a150f76>
- Perplexity Sonar: <https://docs.perplexity.ai/docs/sonar/quickstart> , <https://www.perplexity.ai/api-platform>
- Sakana Fugu: <https://sakana.ai/fugu/> , <https://console.sakana.ai/get-started> , <https://console.sakana.ai/models>
- Xiaomi MiMo: <https://mimo.mi.com/> , https://docs.litellm.ai/docs/providers/xiaomi_mimo , <https://github.com/xiaomimimo/mimo>
- Tencent Hunyuan / Yuanbao: <https://hunyuan.tencent.com/> , <https://intl.cloud.tencent.com/ind/document/product/1290/79463> , https://aiwiki.ai/wiki/tencent_yuanbao
- Google OKF: <https://x.com/GoogleCloudTech/status/2067012903337664886> , <https://www.explainx.ai/blog/google-open-knowledge-format-okf-ai-agents-2026>
- GOT-OCR2.0: https://huggingface.co/docs/transformers/en/model_doc/got_ocr2 , <https://arxiv.org/pdf/2409.01704> , <https://github.com/Ucas-HaoranWei/GOT-OCR2.0>

- “GPT SOL” (GPT-5.6 Sol / Upstage Solar): <https://www.datacamp.com/blog/gpt-5-6-sol-luna-terra> , <https://www.upstage.ai/solar-llm>
- Subquadratic SubQ: <https://subq.ai/introducing-subq> , <https://www.technologyreview.com/2026/06/19/1139313/a-startup-claims-it-broke-through-a-bottleneck-thats-holding-back-llms/> , <https://venturebeat.com/technology/miami-startup-subquadratic-claims-1-000x-ai-efficiency-gain-with-subq-model-researchers-demand-independent-proof>
- Qwythos 9B: <https://huggingface.co/empero-ai/Qwythos-9B-Claude-Mythos-5-1M> , <https://knightli.com/en/2026/06/24/qwythos-9b-claude-mythos-1m-context-guide/>
- xAI Grok: <https://docs.x.ai/developers/models> , <https://x.ai/api>
- Moonshot Kimi: <https://platform.moonshot.ai/> , <https://github.com/moonshotai/kimi-k2>
- Zhipu GLM / Z.ai: <https://pypi.org/project/zhipuai/> , <https://ai-sdk.dev/providers/community-providers/zhipu> , <https://docs.roocode.com/providers/zai>
- Fireworks: <https://docs.fireworks.ai/tools-sdks/openai-compatibility>
- DeepInfra: https://deepinfra.com/docs/openai_api
- Novita: <https://novita.ai/docs/guides/llm-api>
- AI21 Jamba: <https://docs.ai21.com/docs/jamba-foundation-models> , <https://docs.litellm.ai/docs/providers/ai21>
- Reka: <https://docs.reka.ai/> , <https://docs.reka.ai/quickstart>
- Hyperbolic: <https://docs.litellm.ai/docs/providers/hyperbolic>
- Baseten: <https://apimart.ai/blog/mai-models-fireworks-ai-baseten-open-router>